
PSyclone Documentation

Release 1.9.0

**Andrew Coughtrie, Rupert Ford, Joerg Henrichs, Iva Kavcic,
Andrew Porter and Sergi Siso**

Jan 20, 2021

TABLE OF CONTENTS

1	Introduction	1
2	Getting Going	3
2.1	Download	3
2.2	Dependencies	3
2.3	Environment	6
2.4	Configuration	6
2.5	Test	7
2.6	Run	7
3	The psyclone command	11
3.1	Running	11
3.2	Basic Use	12
3.3	Choosing the API	12
3.4	File output	13
3.5	Algorithm files with no invokes	13
3.6	Kernel search directory	13
3.7	Transformation script	14
3.8	Fortran line length	14
3.9	Distributed memory	14
3.10	Automatic Profiling Instrumentation	15
3.11	Outputting of Transformed Kernels	15
3.12	Fortran INCLUDE Files	15
4	Tutorial	17
5	Examples	19
5.1	GOcean	19
5.2	LFRic	21
5.3	NEMO	24
5.4	PSyIR	25
6	System-specific Set-up for Users	27
6.1	Installing dependencies	27
6.2	Installing PSyclone	28
6.3	Common installation	28
7	Kernel layer	29
7.1	API	29
7.2	Metadata	30

8	Algorithm layer	33
8.1	API	33
8.2	Limitations	34
9	PSy layer	37
9.1	Code Generation	37
9.2	API	38
10	Built-ins	43
10.1	Example	43
10.2	Supported Built-in operations	47
10.3	Adding new Built-in operations	47
11	Dynamo0.3 API	49
11.1	Algorithm	49
11.2	PSy-layer	53
11.3	Kernel	54
11.4	Built-ins	74
11.5	Boundary Conditions	84
11.6	Conventions	84
11.7	Configuration	85
11.8	Transformations	85
12	GOcean1.0 API	95
12.1	Introduction	95
12.2	The GOcean Library	95
12.3	Algorithm	99
12.4	Kernel	100
12.5	Built-ins	106
12.6	Conventions	106
12.7	Configuration	107
12.8	Transformations	109
13	NEMO API	115
13.1	Algorithm	115
13.2	Constructing the PSyIR	115
13.3	Example	116
13.4	Transformations	117
13.5	Limitations	120
14	PSyIR : The PSyclone Internal Representation	121
14.1	PSyIR Nodes	121
14.2	Node Descriptions	122
14.3	Text Representation	123
14.4	Tree Navigation	124
14.5	DataTypes	126
14.6	Symbols and Symbol Tables	127
14.7	Creating PSyIR	135
15	Transformations	139
15.1	Finding	139
15.2	Standard Functionality	140
15.3	Available transformations	140
15.4	Kernels	158
15.5	Applying	161

15.6	OpenMP	164
15.7	OpenCL	164
15.8	OpenACC	166
15.9	SIR	167
16	Distributed Memory	169
16.1	Correctness	169
16.2	Performance	169
16.3	Implementation	170
16.4	Control	170
16.5	Status	170
17	Stub Generation	171
17.1	Quick Start	171
17.2	Introduction	171
17.3	Use	172
17.4	Kernels	172
17.5	Example	173
17.6	Errors	176
18	Line length	179
18.1	Script	179
18.2	Interactive	179
18.3	Limitations	180
19	Fortran Naming Conventions	181
20	API	183
20.1	The generator module	183
20.2	The parse module	184
20.3	The transformations module	184
20.4	The psyGen module	211
20.5	The alg_gen module	212
20.6	The line_length module	213
21	PSyData API	215
21.1	Read-Only Verification	215
21.2	NAN Test	217
22	Profiling	219
22.1	Interface to Third Party Profiling Tools	219
22.2	Required Modifications to the Program	220
22.3	Profiling Command Line Options	221
22.4	Profiling in Scripts - ProfileTrans	223
22.5	Naming Profiling Regions	224
23	PSy Kernel Extractor (PSyKE)	227
23.1	Introduction	227
23.2	Use	229
23.3	NetCDF Extraction Example	232
24	Configuration	233
24.1	Options	233
	Bibliography	237

Python Module Index	239
Index	241

INTRODUCTION

PSyclone, the PSy code generator, is being developed for use in finite element, finite volume and finite difference codes. PSyclone development started with the aim to support the emerging API in the GungHo project for a finite element dynamical core.

The [GungHo project](#) was initiated in 2011 to address challenges of weather and climate prediction on the next generation of supercomputers. The project ran for 5 years as a collaboration between the Met Office, NERC (via NERC funded academics) and STFC. It laid a foundation for redesign of the heart of the Met Office's Unified Model, known as the dynamical core, from the choices of numerical methods and model grids to the implementation of parallel algorithms that will scale to millions of cores.

The software infrastructure based on the GungHo project recommendations is now being developed in the [LFRic project](#) and is expected to be operational in 2022. Its development is led by the requirements to support multiple meshes and element types, thus allowing for future model development.

GungHo also proposed a novel separation of concerns for the software implementation of the dynamical core. This approach distinguishes between three layers: the Algorithm layer, the Kernel layer and the Parallelisation System (PSy) layer. Together this separation is termed PSyKAL.

The Algorithm layer specifies the algorithm that the scientist would like to run (in terms of calls to kernel routines and built-in operations) and logically operates on full fields.

The Kernel layer provides the implementation of the code kernels as subroutines. These subroutines operate on local fields (a set of elements, a vertical column, or a set of vertical columns, depending on the kernel).

The PSy layer sits in-between the algorithm and kernel layers and its primary role is to provide node-based parallel performance for the target architecture. The PSy layer can be optimised for a particular hardware architecture, such as multi-core, many-core, GPGPUs, or some combination thereof with no change to the algorithm or kernel layer code. This approach therefore offers the potential for portable performance.

Rather than writing the PSy layer manually, the PSyclone code generation system can help a user to optimise the code for a particular architecture (by providing optimisations such as blocking, loop merging, inlining etc), or alternatively, generate the PSy layer automatically.

PSyclone is also being extended to support an API being developed in the [GOcean project](#) for two finite difference ocean model benchmarks, one of which is based on the [NEMO](#) ocean model.

GETTING GOING

2.1 Download

The following instructions are intended for a PSyclone user who wants to work with a released version of the code. If you are a developer or wish to test a specific branch of PSyclone from the GitHub repository please see [Installation](#) in the [Developer Guide](#).

PSyclone is available on the Python Package Index (pypi.org) and is hosted on GitHub:

```
https://github.com/stfc/PSyclone
```

The latest release is 1.9.0 and the latest stable version is on the master branch.

PSyclone can be installed using pip:

```
> pip install psyclone
```

or downloaded from github - either see 1.9.0 in the `releases` tab on the PSyclone page or download and extract the latest release of PSyclone directly, e.g.

```
> wget https://github.com/stfc/PSyclone/archive/1.9.0.tar.gz
> tar xzf 1.9.0.tar.gz
> ls
PSyclone-1.9.0
```

Hereon the location where you download or clone PSyclone (including the PSyclone directory itself) will be referred to as `<PSYCLONEHOME>`.

2.2 Dependencies

PSyclone is written in Python so needs Python to be installed on the target machine. PSyclone has been tested under Python 2.7 and 3.6.

Warning: Release 1.6.0 of PSyclone requires version 0.0.7 of `fparser` and will fail on more recent versions. However simply installing this version using `pip` will install a later version of `fparser`. The suggested solution is to use release 1.6.1 which fixes this problem and is otherwise identical to 1.6.0. However, if you want to use release 1.6.0 you need to ensure that you have `fparser` 0.0.7. This can be achieved by downgrading an existing `fparser` installation or by installing version 0.0.7 of `fparser` *before* installing PSyclone.

PSyclone immediately relies on four external Python packages; `six`, `configparser`, `fparser` and `pyparsing`. The easiest way to satisfy the Python dependencies is to use the Python Package Index (pypi.org) and `pip`. See <https://packaging.python.org/installing/> for more information.

If everything is working correctly then using pip to install PSyclone:

```
> pip install psyclone
```

will automatically install the Python dependencies (`fparser` and `pyparsing`).

In addition to the mandatory dependencies just described, PSyclone also has optional dependencies on both `graphviz` and `termcolor`. PSyclone can use `graphviz` to produce a visualisation of a schedule's dependency graph. If this is desired then the Python package `graphviz` (for the Python bindings) as well as the `graphviz` package itself must be installed. If the `graphviz` package is not available then the associated PSyclone routines will return silently and no visualisations will be produced. The Python package `termcolor` is used for pretty-printing a schedule in terminals that support coloured text. If the package is not available then the schedule is simply printed in plain text without colour highlighting.

2.2.1 System-specific set-up

System-specific Set-up for Users instructions are available for Ubuntu 14.04.2 and OpenSUSE 42.2.

2.2.2 fparser

The `fparser` package (<https://github.com/stfc/fparser>) is a Fortran parser originally developed as a part of the `f2py` project.

The minimum version of `fparser` required by PSyclone is currently 0.0.9 but we strongly recommend you install the latest version to reduce the chance of encountering problems when parsing existing algorithm or kernel code. (Note that for older versions of PSyclone up to and including 1.5.1 you must use version 0.0.6 of `fparser`.)

`fparser` is available from the Python Package Index and thus may be installed using `pip` (<https://packaging.python.org/installing/#requirements-for-installing-packages>):

```
> pip install fparser
```

If you do not have sufficient permissions to perform a system-wide install then you can instruct pip to do a user-local install:

```
> pip install --user fparser
```

Should you wish to remove `fparser` then simply do:

```
> pip uninstall fparser
```

If you have already installed `fparser` and want to upgrade to the latest version simply do:

```
> pip install fparser --upgrade
```

2.2.3 pyparsing

PSyclone requires `pyparsing`, a library designed to allow parsers to be built in Python. PSyclone uses `pyparsing` to parse fortran regular expressions as `fparser` does not fully parse these, (see <http://pyparsing.wikispaces.com> for more information).

PSyclone has been tested with `pyparsing` versions 1.5.2, 2.0.1 and 2.2.0.

You can test whether `pyparsing` is already installed on your machine by typing `import pyparsing` from the python command line. If `pyparsing` is installed, this command will complete successfully. If `pyparsing` is installed you can check its version by typing `pyparsing.__version__` after successfully importing it.

If `pyparsing` is not installed on your system then it may be installed from the Python Package Index using `pip`:

```
> pip install pyparsing
```

Should you wish to, uninstalling is simply performed by doing:

```
> pip uninstall pyparsing
```

If you do not have sufficient privileges for a system-wide install then you can instruct `pip` to do a user-local install:

```
> pip install --user pyparsing
```

Alternatively, you could follow the instructions here <http://pyparsing.wikispaces.com/Download+and+Installation>.

2.2.4 graphviz

The data dependencies of a PSyIR schedule determine the validity of changes to this schedule. PSyclone supports the visualisation of these dependencies as a graph using `graphviz`. This visualisation is not needed to use PSyclone.

If the Python bindings to `graphviz` are not installed on your system then it may be installed from the Python Package Index using `pip`:

```
> sudo pip install graphviz
```

Should you wish to, uninstalling is simply performed by doing:

```
> sudo pip uninstall graphviz
```

If you do not have sufficient privileges for a system-wide install then you can instruct `pip` to do a user-local install:

```
> pip install --user graphviz
```

If `graphviz` itself is not installed on your system and your system supports the `apt` package manager then see below, otherwise please refer to the download and install instructions which are available here <http://www.graphviz.org/Download..php>.

If your system supports the `apt` package manager then it can be installed and removed in the following way:

```
> sudo apt install graphviz
> sudo apt remove graphviz
```

2.2.5 termcolor

By default, the `view()` method available on any PSyIR (PSyclone Internal Representation) object prints a plain-text representation to standard-out. However, if the `termcolor` package is available then PSyclone uses this to add colour highlighting to the output text.

Installation (and uninstallation) of this package can be done via `pip` in exactly the same way as for `graphviz`, as described above.

2.3 Environment

In order to use PSyclone (including running the test suite and building documentation) you will need to install it. The simplest way to do this is to use pip with the supplied `setup.py` file:

```
> cd <PSYCLONEHOME>
> pip install .
```

By default pip will attempt a system-wide install. If you wish to do a user-local install instead then supply the `--user` flag:

```
> pip install --user .
```

This installs the PSyclone modules in `~/.local/lib/pythonX.Y/site-packages` (where `X.Y` is the version of Python that you are using) and the `'psyclone'` script in `~/.local/bin`. Depending on your linux distribution, you may need to add the latter location to your `$PATH`. The examples are installed in `~/.local/share/psyclone/examples`.

If for some reason you'd rather not use pip then you can run the setup manually:

```
> python setup.py install
```

or, if you don't have root access:

```
> python setup.py install --user
```

or,

```
> python setup.py install --prefix /my/install/path
```

If using the latter method then it will be necessary to take further action to ensure PSyclone can find the configuration file installed as a part of this process - see below.

2.4 Configuration

Various aspects of PSyclone are configured through a configuration file, `psyclone.cfg`. The default version of this file is installed to `<python-base-prefix>/shared/psyclone/` during the installation process. If a system-wide installation is being performed then this will be something like `/usr/share/psyclone/`. If a user-local installation is performed (`--user` flag to `pip install`) then the location will be something like `~/.local/share/psyclone/`.

Warning: if PSyclone is installed to a non-standard location (e.g. by specifying the `--prefix=/some/path` option to `pip install`) then PSyclone will not be able to find the configuration file at execution time. There are two solutions to this: 1. copy the configuration file to a location where PSyclone will find it (see [Configuration](#)) or 2. set the `PSYCLONE_CONFIG` environment variable to the full-path to the configuration file, e.g.:

```
> export PSYCLONE_CONFIG=/some/path/PSyclone/config/psyclone.cfg
```

Warning: when installing in 'editable' mode (`-e` flag to pip), pip does *not* install the configuration file. You will have to take one of the two actions described above.

See [Configuration](#) for details of the settings contained within the config file.

2.5 Test

PSyclone contains an extensive test suite, but this test suite is not part of a standard installation. If you want to run the full test suite, you need to install PSyclone from source, see [Installation](#) in the [Developer Guide](#).

2.6 Run

You are now ready to try running PSyclone on the examples. One way of doing this is to use the `psyclone` driver script. Assuming it is on your PATH:

```
> psyclone
usage: psyclone [-h] [-oalg OALG] [-opsy OPSY] [-okern OKERN] [-api API]
               [-s SCRIPT] [-d DIRECTORY] [-I INCLUDE] [-l {off,all,output}]
               [-dm] [-nodm] [--kernel-renaming {multiple,single}]
               [--profile {invokes,kernels}] [--config CONFIG] [-v]
               filename
psyclone: error: too few arguments
```

As indicated above, the `psyclone` script takes the name of the Fortran source file containing the algorithm specification (in terms of calls to `invoke()`). It parses this, finds the necessary kernel source files and produces two Fortran files. The first contains the PSy, middle layer and the second a re-write of the algorithm code to use that layer. These files are named according to the user-supplied arguments (options `-oalg` and `-opsy`). If those arguments are not supplied then the script writes the generated/re-written Fortran to the terminal. For details of the other command-line arguments please see the [The `psyclone` command](#) Section.

Examples are provided in the `examples` directory of the PSyclone git repository - if you have cloned the repository then `EGS_HOME` in what follows is the root PSyclone directory. Alternatively, if you have installed PSyclone using `pip` then they may be found in the `share/psyclone` directory under your Python installation. In this case you should copy the whole `examples` directory to some convenient location (hereafter called `EGS_HOME`) before attempting to carry out the following instructions. Depending on your precise setup, you may also need to set `PSYCLONE_CONFIG` to the full-path to the PSyclone configuration file (see [Configuration](#)).

There are 3 subdirectories (`lfric`, `gocean` and `nemo`) corresponding to different APIs that are supported by PSyclone. Note, the `lfric` directory corresponds to the `dynamo0.1` and `dynamo0.3` APIs. In this case we are going to use one of the LFRic examples:

```
> cd <EGS_HOME>/examples/lfric/egl
> psyclone -api dynamo0.1 \
> -oalg dynamo_alg.f90 -opsy dynamo_psy.f90 dynamo.F90
```

You should see two new files created called `dynamo_alg.f90` (containing the re-written algorithm layer) and `dynamo_psy.f90` (containing the generated PSy- or middle-layer). Since this is an LFRic example the Fortran source code has dependencies on the LFRic system and therefore cannot be compiled stand-alone.

You can also use the `runme.py` example to see the interactive API in action. This script contains:

```
from psyclone.parse.algorithm import parse
from psyclone.psyGen import PSyFactory

# This example uses version 0.1 of the Dynamo API
api = "dynamo0.1"
```

(continues on next page)

(continued from previous page)

```

# Parse the file containing the algorithm specification and
# return the Abstract Syntax Tree and invokeInfo objects
ast, invokeInfo = parse("dynamo.F90", api=api)

# Create the PSy-layer object using the invokeInfo
psy = PSyFactory(api).create(invokeInfo)
# Generate the Fortran code for the PSy layer
print(psy.gen)

# List the invokes that the PSy layer has
print(psy.invokes.names)

# Examine the 'schedule' (e.g. loop structure) that each
# invoke has
schedule = psy.invokes.get('invoke_0_v3_kernel_type').schedule
schedule.view()

schedule = psy.invokes.get('invoke_1_v3_solver_kernel_type').schedule
schedule.view()

```

It can be run non-interactively as follows:

```

> cd <EGS_HOME>/example/lfric/egl
> python runme.py

```

However, to understand this example in more depth it is instructive to cut-and-paste from the `runme.py` file into your own, interactive python session:

```

> cd <EGS_HOME>/example/lfric/egl
> python

```

In addition to the `runme.py` script, there is also `runme_openmp.py` which illustrates how one applies an OpenMP transform to a loop schedule within the PSy layer. The initial part of this script is the same as that of `runme.py` (above) and is therefore omitted here:

```

# List the various invokes that the PSy layer contains
print(psy.invokes.names)

# Get the loop schedule associated with one of these
# invokes
schedule = psy.invokes.get('invoke_v3_kernel_type').schedule
schedule.view()

# Get the list of possible loop transformations
from psyclone.psyGen import TransInfo
t = TransInfo()
print(t.list)

# Create an OpenMPLoop-transformation object
ol = t.get_trans_name('OMPLoopTrans')

# Apply it to the loop schedule of the selected invoke
new_schedule, memento = ol.apply(schedule.children[0])
new_schedule.view()

```

(continues on next page)

(continued from previous page)

```
# Replace the original loop schedule of the selected invoke
# with the new, transformed schedule
psy.invokes.get('invoke_v3_kernel_type').__schedule = new_schedule
# Generate the Fortran code for the new PSy layer
print(psy.gen)
```


THE PSYCLONE COMMAND

The simplest way to run PSyclone is to use the `psyclone` command. If you installed PSyclone using `pip` then this command should be available on your `PATH` (see [Environment](#) for more details). Alternatively it can be found in the `<PSYCLONEHOME>/bin` directory. The command takes an algorithm file as input and outputs modified algorithm code and generated PSy code. This section walks through its functionality.

3.1 Running

The `psyclone` command is an executable script designed to be run from the command line, e.g.:

```
> psyclone <args>
```

The optional `-h` argument gives a description of the options provided by the command:

```
> psyclone -h
```

```
usage: psyclone [-h] [-oalg OALG] [-opsy OPSY] [-okern OKERN] [-api API]
               [-s SCRIPT] [-d DIRECTORY] [-I INCLUDE] [-l {off,all,output}]
               [-dm] [-nodm] [--kernel-renaming {multiple,single}]
               [--profile {invokes,kernels}] [--config CONFIG] [-v]
               filename
```

Run the PSyclone code generator on a particular file

positional arguments:

filename	algorithm-layer source code
----------	-----------------------------

optional arguments:

<code>-h, --help</code>	show this help message and exit
<code>-oalg OALG</code>	filename of transformed algorithm code
<code>-opsy OPSY</code>	filename of generated PSy code
<code>-okern OKERN</code>	directory in which to put transformed kernels
<code>-api API</code>	choose a particular api from ['dynamo0.1', 'dynamo0.3', 'gocean0.1', 'gocean1.0', 'nemo'], default 'dynamo0.3'.
<code>-s SCRIPT, --script SCRIPT</code>	filename of a PSyclone optimisation script
<code>-d DIRECTORY, --directory DIRECTORY</code>	path to root of directory structure containing kernel source code
<code>-I INCLUDE, --include INCLUDE</code>	

```
                                path to Fortran INCLUDE files (nemo API only)
-l {off,all,output}, --limit {off,all,output}
                                limit the Fortran line length to 132 characters
                                (default 'off'). Use 'on' to apply limit to both input
                                and output Fortran. Use 'output' to apply line-length
                                limit to output Fortran only.
-dm, --dist_mem                generate distributed memory code
-nodm, --no_dist_mem           do not generate distributed memory code
--kernel-renaming {single,multiple}
                                Naming scheme to use when re-naming transformed
                                kernels.
--profile {invokes,kernels}, -p {invokes,kernels}
                                Add profiling hooks for either 'kernels' or 'invokes'
--config CONFIG                Config file with PSyclone specific options.
-v, --version                  Display version information (1.9.0)
```

3.2 Basic Use

The simplest way to use `psyclone` is to provide it with an algorithm file:

```
> psyclone alg.f90
```

If the algorithm file is invalid for some reason, the command should return with an appropriate error. For example, if we use the Python `genkernelstub` script as an algorithm file we get the following:

```
> psyclone <PSYCLONEHOME>/bin/genkernelstub
...
1:#!/usr/bin/env python <== no parse pattern found for "#" in 'BeginSource' block.
'Parse Error: Fatal error in external fparser tool'
```

If the algorithm file is valid then the modified algorithm code and the generated PSy code will be output to the terminal screen.

3.3 Choosing the API

In the previous section we relied on PSyclone using the default API. The default API, along with the supported APIs can be seen by running the `psyclone` command with the `-h` option.

If you use a particular API frequently and it is not the default then you can change the default by creating a copy of the default `psyclone.cfg` file and editing it. See [Configuration](#) for more details.

If your code uses an API that is different to the default then you can specify this as an argument to the `psyclone` command.

```
> psyclone -api dynamo0.1 alg.f90
```

3.4 File output

By default the modified algorithm code and the generated PSy code are output to the terminal. These can instead be output to files by using the `-oalg <file>` and `-opsy <file>` options, respectively. For example, the following will output the generated PSy code to the file ‘psy.f90’ but the algorithm code will be output to the terminal:

```
> psyclone -opsy psy.f90 alg.f90
```

If PSyclone is being used to transform Kernels then the location to write these to is specified using the `-okern <directory>` option. If this is not supplied then they are written to the current working directory. By default, PSyclone will overwrite any kernel of the same name in that directory. To change this behaviour, the user can use the `--no_kernel_clobber` option. This causes PSyclone to re-name any transformed kernel that would clash with any of those already present in the output directory.

3.5 Algorithm files with no invokes

If `psyclone` is provided with a file that contains no `invoke` calls then the command outputs a warning to `stdout` and copies the input file to `stdout`, or to the specified algorithm file (if the `-oalg <file>` option is used). No PSy code will be output. If a file is specified using the `-opsy <file>` option this file will not be created.

```
> psyclone -opsy psy.f90 -oalg alg_new.f90 empty_alg.f90
Warning: 'Algorithm Error: Algorithm file contains no invoke() calls: refusing to
generate empty PSy code'
```

3.6 Kernel search directory

When an algorithm file is parsed, the parser looks for the associated kernel files. The way in which this is done requires that any user-defined kernel routine (as opposed to *Built-ins*) called within an `invoke` must have an explicit `use` statement. For example, the following code gives an error:

```
> cat no_use.f90
program no_use
  call invoke(testkern_type(a,b,c,d,e))
end program no_use
> psyclone -api gocean1.0 no_use.f90
"Parse Error: kernel call 'testkern_type' must either be named in a use statement or
↳be a recognised built-in (one of '[]' for this API)"
```

(If the chosen API has any *Built-ins* defined then these will be listed within the `[]` in the above error message.) If the name of the kernel is provided in a `use` statement then the parser will look for a file with the same name as the module in the `use` statement. In the example below, the parser will look for a file called “testkern.f90” or “testkern.F90”:

```
> cat use.f90
program use
  use testkern, only : testkern_type
  call invoke(testkern_type(a,b,c,d,e))
end program use
```

Therefore, for PSyclone to find kernel files, the module name of a kernel file must be the same as its filename. By default the parser looks for the kernel file in the same directory as the algorithm file. If this file is not found then an error is reported.

```
> psyclone use.f90
Kernel file 'testkern.[fF]90' not found in <location>
```

The `-d` option can be used to tell `psyclone` where to look for kernel files by supplying it with a directory. The execution will recurse from the specified directory path to look for the required file. There must be only one instance of the specified file within (or below) the specified directory:

```
> cd <PSYCLONEHOME>/src/psyclone
> psyclone -d . use.f90
More than one match for kernel file 'testkern.[fF]90' found!
> psyclone -d tests/test_files/dynamo0p3 -api dynamo0.3 use.f90
[code output]
```

Note: The `-d` option is limited to a single directory. Therefore a current limitation in PSyclone is that all kernel files required by an algorithm file must exist within a directory hierarchy where their file names are unique.

3.7 Transformation script

By default the `psyclone` command will generate ‘vanilla’ PSy layer code. The `-s` option allows a Python script to be specified which can transform the PSy layer. This option is discussed in more detail in the [Script](#) section.

3.8 Fortran line length

By default the `psyclone` command will generate Fortran code with no consideration of Fortran line-length limits. As the line-length limit for free-format Fortran is 132 characters, the code that is output may be non-conformant.

Line length is not an issue for many compilers as they allow compiler flags to be set which allow lines longer than the Fortran standard. However this is not the case for all compilers.

When either the `-l all` or `-l output` option is specified to the `psyclone` command, the output will be line wrapped so that the output lines are always within the 132 character limit.

The `-l all` additionally checks the parsed algorithm and kernel files for conformance and raises an error if they do not conform.

Line wrapping is not performed by default. There are two reasons for this. This first reason is that most compilers are able to cope with long lines. The second reason is that the line wrapping implementation could fail in certain pathological cases. The implementation and limitations of line wrapping are discussed in the [Limitations](#) section.

3.9 Distributed memory

By default the `psyclone` command will generate distributed memory (DM) code (i.e. parallelised using MPI). As with the choice of API, this default may be configured by editing `psyclone.cfg` - see [Configuration](#). Alternatively, whether or not to generate DM code can be specified as an argument to the `psyclone` command using the `-dm/--dist_mem` or `-nodm/--no_dist_mem` flags, respectively.

For details of PSyclone’s support for generating DM code see [Distributed Memory](#).

3.10 Automatic Profiling Instrumentation

The `--profile` option allows the user to instruct PSyclone to automatically insert profiling calls within the generated PSy code. Two options are provided, `invokes` and `kernels`. The first of these causes PSyclone to insert profiling-start and -stop calls at the beginning and end of every generated invoke routine. The second puts profiling calls around every kernel call (including the associated loops). The generated code must be linked against the PSyclone profiling interface and the profiling tool itself. The application that calls the PSyclone-generated code is responsible for initialising and finalising the profiling library that is being used. For full details on the use of this profiling functionality please see the [Profiling](#) section.

3.11 Outputting of Transformed Kernels

When transforming kernels there are two use-cases to consider:

1. a given kernel will be transformed only once and that version then used from multiple, different Invokes and Algorithms;
2. a given kernel is used from multiple, different Invokes and Algorithms and is transformed differently, depending on the Invoke.

Whenever PSyclone is used to transform a kernel, the new kernel must be re-named in order to avoid clashing with other possible calls to the original. By default (`--kernel-renaming multiple`), PSyclone generates a new, unique name for each kernel that is transformed. Since PSyclone is run on one Algorithm file at a time, it uses the chosen kernel output directory (`-okern`) to ensure that names created by different invocations do not clash. Therefore, when building a single application, the same kernel output directory must be used for each separate invocation of PSyclone.

Alternatively, in order to support use case 1, a user may specify `--kernel-renaming single`: now, before transforming a kernel, PSyclone will check the kernel output directory and if a transformed version of that kernel is already present then that will be used. Note, if the kernel file on disk does not match with what would be generated then PSyclone will raise an exception.

3.12 Fortran INCLUDE Files

For the NEMO API, if the source code to be processed by PSyclone contains `INCLUDE` statements (other than those for libraries such as MPI) then the location of any `INCLUDE`'d files must be supplied to PSyclone via the `-I` or `--include` option. (This is necessary because `INCLUDE` lines are a part of the Fortran language and must therefore be parsed - they are not handled by any pre-processing step.) Multiple locations may be specified by using multiple `-I` flags, e.g.:

```
> psyclone api "nemo" -I /some/path -I /some/other/path alg.f90
```

If no include paths are specified then the directory containing the source file currently being parsed is searched by default. If the specified include file is not found then ideally the `INCLUDE` line would be left unchanged. However, `fparses` currently treats any such `INCLUDE` lines as comments which results in them being lost (`fparses` issue #138). The workaround for this is to ensure that the location of *all* `INCLUDE` files is supplied to PSyclone.

Attempting to specify `-I/--include` for any API other than NEMO will be rejected by PSyclone.

TUTORIAL

PSyclone provides a tutorial, part of which uses Jupyter notebooks. This can be launched from a browser using binder: <https://mybinder.org/v2/gh/stfc/psyclone/master?filepath=tutorial%2Fnotebooks%2Fintroduction.ipynb/>.

If PSyclone is installed on your system then you can run the tutorial locally. First find the tutorial. If you have installed PSyclone using pip then the examples may be found in `share/psyclone/tutorial/notebooks` under your Python installation (e.g. `~/local` for a user-local installation). Next copy the tutorial to a local writable space (as the notebooks will be modified as you go through the tutorial) and change directory to this tutorial. Lastly, start up the tutorial's introduction *jupyter-notebook introduction.ipynb*.

EXAMPLES

Various examples of the use of PSyclone are provided under the `examples` directory in the git repository. If you have installed PSyclone using `pip` then the examples may be found in `share/psyclone/examples` under your Python installation (e.g. `~/local` for a user-local installation).

Running any of these examples requires that PSyclone be installed on the host system, see [Getting Going](#). This section is intended to provide an overview of the various examples so that a user can find one that is appropriate to them. For details of how to run each example please see the `README.md` files in the associated directories.

Alternatively, some of the examples have associated Jupyter notebooks that may be launched with Binder on [MyBinder](#). This is most easily done by following the links from the top-level [README](#).

For the purposes of correctness checking, the whole suite of examples may be executed using `Gnu make` (this functionality is used by GitHub Actions alongside the test suite). The default target is `transform` which just performs the PSyclone code transformation steps for each example. For those examples that support it, the `compile` target also requests that the generated code be compiled. The `notebook` target checks the various Jupyter notebooks using `nbconvert`.

Note: if you have copied the examples directory to some other location but still wish to use `make` then you will also have to set the `PSYCLONE_CONFIG` environment variable to the full path to the PSyclone configuration file, e.g. `$ PSYCLONE_CONFIG=/some/path/psyclone.cfg make`

Each of the Makefiles also provides the `clean` and `allclean` targets. The former simply removes all generated files while the latter also cleans each of the infrastructure libraries used when compiling the examples (`dl_timer`, `dl_esm_inf` and the various PSyData wrapper libraries).

5.1 GOcean

5.1.1 Example 1: Loop transformations

Examples of applying various transformations (loop fusion, OpenMP, OpenACC, OpenCL) to the semi-PSyKAI'd version of the Shallow benchmark. ("semi" because not all kernels are called from within `invoke()`'s.) Also includes an example of generating a DAG from an `InvokeSchedule`.

5.1.2 Example 2: OpenACC

This is a simple but complete example of using PSyclone to enable an application to run on a GPU by adding OpenACC directives. A Makefile is included which will use PSyclone to generate the PSy code and transformed kernels and then compile the application. This compilation requires that the `dl_esm_inf` library (github.com/stfc/dl_esm_inf) be installed/available - it is provided as a git submodule of the PSyclone project (see [Installation](#) in the Developers' Guide for details on working with submodules).

The supplied Makefile also provides a second, "profile" target which performs the same OpenACC transformations but then encloses the whole of the resulting PSy layer in a profiling region. By linking this with the PSyclone NVTX profiling wrapper (and the NVTX library itself), the resulting application can be profiled using NVIDIA's *nvprof* or *nvvp* tools.

5.1.3 Example 3: OpenCL

Example of the use of PSyclone to generate an OpenCL driver version of the PSy layer and OpenCL kernels. The Makefile in this example provides a target (*make compile-ocl*) to compile the generated OpenCL code. This requires an OpenCL implementation installed in the system. Read the README provided in the example folder for more details about how to compile and execute the generated OpenCL code.

5.1.4 Example 4: Kernels containing use statements

Transforming kernels for use with either OpenACC or OpenCL requires that we handle those that access data and/or routines via module `use` statements. This example shows the various forms for which support is being implemented. Although there is support for converting global-data accesses into kernel arguments, PSyclone does not yet support nested `use` of modules (i.e. data accessed via a module that in turn imports that symbol from another module) and kernels that call other kernels (Issue #342).

5.1.5 Example 5: Profiling

This example shows how to use the profiling support in PSyclone. It instruments two `invoke` statements and can link in with any of the following profiling wrapper libraries: `template`, `simple_timer`, `apeg-dl_timer`, and `DrHook` (see [Interface to Third Party Profiling Tools](#)). The `README.md` file contains detailed instructions on how to build the different executables. By default (i.e. just using `make` without additional parameters) it links in with the `template` profiling library included in PSyclone. This library just prints out the name of the module and region before and after each `invoke` is executed. This example can actually be executed to test the behaviour of the various profiling wrappers, and is also useful if you want to develop your own wrapper libraries.

5.1.6 Example 6: Kernel data extraction

This example shows the use of kernel data extraction in PSyclone. It instruments each of the two `invoke`s in the example program with the PSyData-based kernel extraction code. It uses the `dl_esm_inf`-specific extraction library '`netcdf`' (`lib/extract/netcdf/dl_esm_inf`), and needs NetCDF to be available (including `nf-config` to detect installation-specific paths). You need to compile the NetCDF extraction library (see [NetCDF Extraction Example](#)). The makefile in this example will link with the compiled NetCDF extraction library and NetCDF. You can execute the created binary and it will create two output netcdf files, one for each of the two `invoke`s.

It will also create two stand-alone driver programs (one for each `invoke`), that will read the corresponding NetCDF file, and then executes the original code.

Note: At this stage the driver program will not compile (see issue #644).

5.1.7 Example 7: Read-only-verification

This example shows the use of read-only-verification with PSyclone. It instruments each of the two invokes in the example program with the PSyData-based read-only-verification code. It uses the `dl_esm_inf`-specific read-only-verification library (`lib/read_only/dl_esm_inf/`).

Note: The `update_field_mod` subroutine contains some very buggy and non-standard code to change the value of some read-only variables and fields, even though the variables are all declared with `intent(in)`. It uses the addresses of variables and then out-of-bound writes to a writeable array to actually overwrite the read-only variables. Using array bounds checking at runtime will be triggered by these out-of-bound writes.

The makefile in this example will link with the compiled read-only-verification library. You can execute the created binary and it will print two warnings about modified read-only variables:

```
-----
Double precision field b_fld has been modified in main : update
Original checksum:  4611686018427387904
New checksum:      4638355772470722560
-----
-----
Double precision variable z has been modified in main : update
Original value:    1.0000000000000000
New value:        123.00000000000000
-----
```

5.2 LFRic

Examples 1 and 2 are for the (deprecated) Dynamo 0.1 API. The remaining examples are all for the Dynamo 0.3 API.

5.2.1 Example 1

Basic operation of PSyclone with `invoke()`'s containing just one kernel for the Dynamo 0.1 API. Also includes an example of transforming for OpenMP.

5.2.2 Example 2

A more complex example for the Dynamo 0.1 API containing multi-kernel invokes. Provides examples of OpenMP and loop fusion transformations.

5.2.3 Example 3

Shows the use of colouring and OpenMP for the Dynamo 0.3 API. Includes multi-kernel, named invokes with both user-supplied and built-in kernels. Also shows the use of `Wchi` function space metadata for coordinate fields in LFRic.

5.2.4 Example 4

Demonstrates the use of the special `enforce_bc_kernel` which PSyclone recognises as a boundary-condition kernel.

5.2.5 Example 5

Example of kernels which require stencil information.

5.2.6 Example 6

Example of applying OpenMP to an `InvokeSchedule` containing kernels that perform reduction operations. Two scripts are provided, one of which demonstrates how to request that PSyclone generate code for a reproducible OpenMP reduction. (The default OpenMP reduction is not guaranteed to be reproducible from one run to the next on the same number of threads.)

5.2.7 Example 7

Example of kernels requiring Column-Matrix Assembly operators.

5.2.8 Example 8

Example of the use of the redundant-computation and move transformations to eliminate and re-order halo exchanges.

5.2.9 Example 9

Demonstrates the behaviour of PSyclone for kernels that read and write quantities on horizontally-discontinuous function spaces. In addition, this example demonstrates how to write a PSyclone transformation script that only colours loops over continuous spaces.

5.2.10 Example 10

Demonstrates the use of “inter-grid” kernels that prolong or restrict fields (map between grids of different resolutions), as well as the use of `ANY_DISCONTINUOUS_SPACE` function space metadata.

5.2.11 Example 11

Example of the use of transformations to introduce redundant computation, split synchronous halo exchanges into asynchronous exchanges (start and stop) and move the starts of those exchanges in order to overlap them with computation.

5.2.12 Example 12

Example of applying code extraction to Nodes in an Invoke Schedule:

```
> psyclone -nodm -s ./extract_nodes.py \
    gw_mixed_schur_preconditioner_alg_mod.x90
```

or to a Kernel in an Invoke after applying transformations:

```
> psyclone -nodm -s ./extract_kernel_with_transformations.py \
    gw_mixed_schur_preconditioner_alg_mod.x90
```

For now it only inserts comments in appropriate locations while the the full support for code extraction is being developed.

This example also contains a Python helper script `find_kernel.py` which displays the names and Schedules of Invokes containing call(s) to the specified Kernel:

```
> python find_kernel.py
```

5.2.13 Example 13 : Kernel transformation

Demonstrates how an LFRic kernel can be transformed. The example transformation makes Kernel values constant where appropriate. For example, the number of levels is usually passed into a kernel by argument but the transformation allows a particular value to be specified which the transformation then sets as a parameter in the kernel. Hard-coding values in a kernel helps the compiler to do a better job when optimising the code.

5.2.14 Example 14: OpenACC

Example of adding OpenACC directives in the dynamo0.3 API. This is a work in progress so the generated code may not work as expected. However it is never-the-less useful as a starting point. Three scripts are provided.

The first script (`acc_kernels.py`) shows how to add OpenACC Kernels directives to the PSy-layer. This example only works with distributed memory switched off as the OpenACC Kernels transformation does not yet support halo exchanges within an OpenACC Kernels region.

The second script (`acc_parallel.py`) shows how to add OpenACC Loop, Parallel and Enter Data directives to the PSy-layer. Again this example only works with distributed memory switched off as the OpenACC Parallel transformation does not support halo exchanges within an OpenACC Parallel region.

The third script (`acc_parallel_dm.py`) is the same as the second except that it does support distributed memory being switched on by placing an OpenACC Parallel directive around each OpenACC Loop directive, rather than having one for the whole invoke. This approach avoids having halo exchanges within an OpenACC Parallel region.

The generated code has a number of problems including 1) it does not modify the kernels to include the OpenACC Routine directive, 2) a loop's upper bound is computed via a derived type (this should be computed beforehand) 3) `set_dirty` and `set_clean` calls are placed within an OpenACC Parallel directive and 4) there are no checks on whether loops are parallel or not, it is just assumed they are - i.e. support for colouring or locking is not yet implemented.

5.2.15 Example 15: CPU Optimisation of Matvec

Example of optimising the LFRic matvec kernel for CPUs. This is work in progress with the idea being that PSyclone transformations will be able to reproduce hand-optimised code.

There is one script which, when run:

```
> psyclone ./matvec_opt.py ../code/gw_mixed_schur_preconditioner_alg_mod.x90
```

will print out the modified matvec kernel code. At the moment no transformations are included (as they are work-in-progress) so the code that is output is the same as the original (but looks different as it has been translated to PSyIR and then output by the PSyIR fortran back-end).

5.2.16 Example 16: Kernel Extraction

The subdirectory `full_example_extract` contains a runnable example of LFRic code that creates a NetCDF file with the input and output parameters of a simple kernel. This example requires the installation of a Fortran compiler and NetCDF development environment. After compilation of the code, you can run the example, which will produce a NetCDF file:

```
cd full_example_extraction
make
./extract
ncdump ./main-update.nc | less
```

5.3 NEMO

These examples may all be found in the `examples/nemo` directory.

5.3.1 Example 1: OpenMP parallelisation of tra_adv

Demonstrates the use of PSyclone to parallelise the loops over vertical levels in a NEMO tracer-advection benchmark using OpenMP.

5.3.2 Example 2: OpenMP parallelisation of traldf_iso

Demonstrates the use of PSyclone to parallelise the loops over vertical levels in some NEMO tracer-diffusion code using OpenMP.

5.3.3 Example 3: OpenACC parallelisation of tra_adv

Demonstrates the introduction of simple OpenACC parallelisation (using the `data` and `kernels` directives) for a NEMO tracer-advection benchmark.

5.3.4 Example 4: Transforming Fortran code to the SIR

Demonstrates that simple Fortran code examples which conform to the NEMO API can be transformed to the Stencil Intermediate Representation (SIR). The SIR is the front-end language to DAWN (<https://github.com/MeteoSwiss-APN/dawn>), a tool which generates optimised cuda, or gridtools code. Thus these simple Fortran examples can be transformed to optimised cuda and/or gridtools code by using PSyclone and then DAWN.

5.4 PSyIR

Examples may all be found in the `example/psyir` directory. Read the `README.md` file in this directory for full details.

5.4.1 Example 1: Constructing PSyIR and Generating Code

`create.py` is a Python script that demonstrates the use of the various `create` methods to build a PSyIR tree from scratch.

5.4.2 Example 2: Creating PSyIR for Aggregate Types

`create_aggregate_types.py` demonstrates the representation of aggregate types (i.e. Fortran derived types or C structs) in the PSyIR.

SYSTEM-SPECIFIC SET-UP FOR USERS

This chapter describes the setup for a user of PSyclone. It includes all steps necessary to be able to use PSyclone. And while you could obviously do some development, none of the required tools for testing or documentation creation will be installed.

The [System-specific Developer Set-up](#) in the Developers' Guide describes the additional installation of all required tools to run tests and create documentation.

Detailed instructions are provided for Ubuntu 16.04.2 and OpenSUSE 42.2 - if you are working with a different Linux distribution some adjustments will be necessary. You will need a terminal window open in which to enter the commands.

6.1 Installing dependencies

Most required dependencies are installed from the Python Package Index (<https://packaging.python.org/installing/>) using the program `pip` ("PIP Installs Packages"). Besides `pip` it is also recommended to install the `graphviz` package to be able to visualise dependency graphs. This is optional and the associated routine will silently return if the `graphviz` bindings are not installed.

6.1.1 Installing dependencies on Ubuntu

On Ubuntu `pip` and `graphviz` are installed using `apt-get`. Remember that `graphviz` is optional and that you'll need to install the `graphviz` package in addition to the Python bindings.

```
> sudo apt-get install python-pip graphviz
```

6.1.2 Installing dependencies on OpenSUSE

The vanilla OpenSUSE installation includes `pip` for Python 3. Note that the `graphviz` package is installed by default.

```
> sudo zypper install python-pip
```

6.2 Installing PSyclone

Change your working directory to where you would like to place the code and download the latest stable release of PSyclone.

```
> cd <PSYCLONEHOME>
> wget https://github.com/stfc/PSyclone/archive/1.9.0.tar.gz
> gunzip 1.9.0.tar.gz
> tar xf 1.9.0.tar
> rm 1.9.0.tar
> cd PSyclone-1.9.0
> export PYTHONPATH=`pwd`/src:${PYTHONPATH}
> export PATH=`pwd`/bin:${PATH}
```

This sets up your python path and path appropriately. You may want to set these paths permanently (e.g. by editing your `${HOME}/.bashrc` file if you run the BASH shell). You can also use the latest version using git, as described in [Installing PSyclone From GitHub](#).

6.3 Common installation

To avoid warnings during the dependency installation, it is recommended to update `pip` to the latest version:

```
> sudo pip install --upgrade pip
```

Next you need to install the `fparser` and `pyparsing` packages:

```
> sudo pip install fparser pyparsing
```

Tip: With `pip` it is possible to install packages either system-wide (which requires root privileges) as above, or for a single user only (in `~/.local`). While the latter is only useful for one particular user, it means that PSyclone can be installed using `pip` without needing root privileges. In order to install a package for a user, add the `--user` command line option to all `pip` commands. This flag requests that the packages be installed locally for the current user rather than requiring root access:

```
> pip install --user fparser pyparsing
```

You may remove the use of `sudo` and add the `--user` option to all `pip` commands described in this document.

Uninstalling is simply a matter of doing:

```
> sudo pip uninstall fparser pyparsing
```

PSyclone supports the ability to output a schedule dependency graph using the `graphviz` package. This is optional and the associated routine will silently return if the `graphviz` bindings are not installed. If you have the `graphviz` package installed (see especially section [Installing dependencies on Ubuntu](#) if you are on Ubuntu), you also need to install the python bindings to the `graphviz` package:

```
> sudo pip install graphviz
```

If you just want to use PSyclone then you've installed all you need and you are ready to go to the getting-going [Run](#) section.

KERNEL LAYER

In the PSyKAl separation of concerns, Kernel code (code which is created to run within the Kernel layer), operates on a subset of a field (such as a column of cells). The reason for doing this is that it gives the PSy layer the responsibility of calling the Kernel over the spatial domain which is where parallelism is typically exploited in finite element and finite difference codes. The PSy layer is therefore able to call the kernel layer in a flexible way (blocked and/or in parallel for example). Kernel code in the kernel layer is not allowed to include any parallelisation calls or directives and works on raw Fortran arrays (to allow the compiler to optimise the code).

Since a Kernel is called over the spatial domain (by the PSy layer) it must take at least one field or operator as an argument.

7.1 API

Kernels in the kernel layer are implemented as subroutines within Fortran modules. One or more kernel modules are allowed, each of which can contain one or more kernel subroutines. In the example below there is one module `integrate_one_module` which contains one kernel subroutine `integrate_one_code`. The kernel subroutines contain the code that operates over a subset of the field (such as a column).

Metadata describing the kernel subroutines is required by the PSyclone system to generate appropriate PSy layer code. The metadata is written by the kernel developer and is kept with the kernel code in the same module using a sub-type of the `kernel_type` type. In the example below the `w3_solver_kernel_type` type specifies the appropriate metadata information describing the kernel code for the `dynamo0.3` api:

```
module w3_solver_kernel_mod

  use kernel_mod,          only : kernel_type
  use constants_mod,       only : r_def, i_def
  use fs_continuity_mod,   only : W3, Wchi
  use argument_mod,        only : arg_type, func_type, &
                             GH_FIELD, GH_SCALAR, &
                             GH_REAL, GH_READ, GH_WRITE, &
                             GH_BASIS, GH_DIFF_BASIS, &
                             GH_QUADRATURE_XYZ, CELLS

  implicit none

  private

  type, public, extends(kernel_type) :: w3_solver_kernel_type
    private
    type(arg_type) :: meta_args(4) = (/ &
      arg_type(GH_FIELD, GH_REAL, GH_WRITE, W3), &
```

(continues on next page)

(continued from previous page)

```

        arg_type(GH_FIELD,    GH_REAL, GH_READ,  W3),    &
        arg_type(GH_FIELD*3, GH_REAL, GH_READ,  Wchi),   &
        arg_type(GH_SCALAR,   GH_REAL, GH_READ)         &
    /)
    type(func_type) :: meta_funcs(2) = (/               &
        func_type(W3,    GH_BASIS),                     &
        func_type(Wchi, GH_DIFF_BASIS)                   &
    /)
    integer :: gh_shape = GH_QUADRATURE_XYzZ
    integer :: operates_on = CELL_COLUMN
contains
    procedure, nopass :: solver_w3_code
end type

contains

    subroutine solver_w3_code(nlayers,                    &
                             x, rhs,                     &
                             chi_1, chi_2, chi_3, ascalar, &
                             ndf_w3, undf_w3, map_w3, w3_basis, &
                             ndf_w0, undf_w0, map_w0, w0_diff_basis, &
                             nqp_h, nqp_v, wqp_h, wqp_v)
    ...
    end subroutine solver_w3_code
end module w3_solver_kernel_mod

```

7.2 Metadata

Kernel metadata is not required if the PSy layer is going to be written manually - its sole purpose is to let PSyclone know how to generate the PSy layer. The content of Kernel metadata differs depending on the particular API and this information can be found in the API-specific sections of this document.

In all APIs the kernel metadata is implemented as an extension of the *kernel_type* type. The reason for using a type to specify metadata is that it allows the metadata to be kept with the code and for it to be compilable. In addition, currently all APIs will contain information about the arguments in an array called *meta_args*, a specification of what data the kernel code expects in a variable called *operates_on* and a reference to the kernel code itself as a type-bound procedure:

```

type, extends(kernel_type) :: integrate_one_kernel
...
type(...) :: meta_args(...) = (/ ... /)
...
integer :: operates_on = ...
...
contains
...
procedure ...
...
end type integrate_one_kernel

```

If no type-bound procedure is declared then a named interface with module procedures must be included in the module:

```
type, extends(kernel_type) :: integrate_one_kernel
...
type(...) :: meta_args(...) = (/ ... /)
...
integer :: operates_on = ...
...
end type integrate_one_kernel

interface ...
  module procedure ...
end interface
```

These module procedures provide alternative implementations (using different precisions) of the kernel code. They are selected as appropriate by the Fortran compiler, depending on the precision of the fields being passed to them.

ALGORITHM LAYER

In the PSyKAl separation of concerns, the Algorithm layer specifies the algorithm that the scientist would like to run (in terms of calls to kernel routines and Built-in operations) and logically operates on full fields. Algorithm code in the algorithm layer is not allowed to include any parallelisation calls or directives and passes datatypes specified by the particular API.

8.1 API

The Algorithm layer is forbidden from calling the Kernel layer directly. In PSyclone, if the programmer would like to call a Kernel routine or a Built-in operation from the algorithm layer they must use the `invoke` call (which is common to all APIs). The `invoke` call is not necessary (and indeed will not work) if the PSy layer is written manually.

To make an `invoke` call, the algorithm layer developer adds one or more `call invoke()` statements to their code and within the content of the `invoke` call they add a reference to the required Kernel/Built-in and the data to pass to it. For example,

```
...  
call invoke(integrate_one_kernel(arg1,arg2))  
...
```

For more information on the concept of Built-in operations see the [Built-ins](#) Section. Details of which operations are supported for a specific API are given in the documentation of that API.

The algorithm layer can consist of an arbitrary number of files containing fortran code, any of which may contain as many `invoke()` calls as is required. PSyclone is applied to an individual algorithm layer file and must therefore be run multiple times if multiple files containing `invoke()` calls exist in the algorithm layer.

The algorithm developer is also able to reference more than one Kernel/Built-in within an `invoke`. In fact this feature is encouraged for performance reasons. **As a general guideline the developer should aim to use as few invokes as possible with as many Kernel references within them as is possible.** The reason for this is that it allows for greater freedom for optimisation in the PSy layer as PSy layer optimisations are limited to the contents of individual `invoke` calls - PSyclone currently does not attempt to optimise the PSy layer over multiple `invoke` calls.

As well as generating the PSy layer code, PSyclone modifies the Algorithm layer code, replacing `invoke` calls with calls to the generated PSy layer so that the algorithm code is compilable and linkable to the PSy layer and adding in the appropriate `use` statement. For example, the above `integrate_one_kernel` `invoke` is translated into something like the following:

```
...  
use psy, only : invoke_0_integrate_one_kernel  
...  
call invoke_0_integrate_one_kernel(arg1,arg2)  
...
```

You may have noticed from other examples in this guide that an algorithm specification in an invoke call references the metadata `type` in an invoke call, not the `code` directly; this is by design.

For example, in the invoke call below, `integrate_one_kernel` is used.

```
...  
call invoke(integrate_one_kernel(arg1,arg2))  
...
```

`integrate_one_kernel` is the name of the metadata type in the module, not the name of the subroutine in the Kernel...

```
module integrate_one_module  
  ...  
  type, extends(kernel_type) :: integrate_one_kernel  
  ...  
end type  
...  
contains  
  ...  
  subroutine integrate_one_code(...)  
  ...  
end subroutine integrate_one_code  
...  
end module integrate_one_module
```

8.1.1 Named Invokes

PSyclone permits the user to optionally specify a label for an invoke call like so:

```
...  
call invoke(integrate_one_kernel(arg1,arg2), &  
           name="compute something")  
...
```

The `name` argument to the invoke call is optional. If supplied it must be a string literal. The content of this string (with any spaces replaced by ‘`_`’ characters) is used in naming the corresponding PSy-layer routine generated by PSyclone. So, for the above example, the generated PSy-layer subroutine will be named “`invoke_compute_something`.” Each invoke label must currently be unique within an Algorithm source file. Note that, in keeping with the Fortran language, labels are not case sensitive and, after having any spaces replaced by underscores, must be valid Fortran names (e.g. `name="compute(1)"` is invalid). In the future it is intended that the labelling of invokes will help to support invoke-specific optimisations to be applied as well as enabling more readable profiling output. It may also be used to instruct PSyclone to just generate a single subroutine to implement all invokes that share the same label.

8.2 Limitations

In order to re-write the Algorithm layer, as just described, PSyclone must obviously be able to parse the invoke calls. Since the Fortran expression parser used by PSyclone is relatively simple, this means there are limitations on what Fortran may be used when specifying kernel arguments in an invoke call. Since these limitations can have a direct impact on the natural science code, the PSyclone developers endeavour to keep them to a minimum.

The current list of known limitations/restrictions on the form of kernel arguments within an invoke is:

- No arithmetic expressions (e.g. `kernel_type(a+b)` or `kernel_type(-a)`)

- No named (optional) arguments (e.g. `kernel_type(fn(my_arg=a))`)

If you encounter any other limitations (or have a burning desire to use one of the above forms) then please contact the PSyclone developers.

PSY LAYER

In the PSyKAI separation of concerns, the PSy layer is responsible for linking together the Algorithm and Kernel layers and for providing the implementation of any Built-in operations used. Its functional responsibilities are to

1. map the arguments supplied by an Algorithm `invoke` call to the arguments required by a Built-in or Kernel call (as these will not have a one-to-one correspondance).
2. call any Kernel routines such that they cover the required iteration space and
3. perform any Built-in operations (either by including the necessary code directly in the PSy layer or by e.g. calling a maths library) and
4. include any required distributed memory operations such as halo swaps and reductions.

Its other role is to allow the optimisation expert to optimise any required distributed memory operations, include and optimise any shared memory parallelism and optimise for single node (e.g. cache and vectorisation) performance.

9.1 Code Generation

The PSy layer can be written manually but this is error prone and potentially complex to optimise. The PSyclone code generation system generates the PSy layer so there is no need to write the code manually.

To generate correct PSy layer code, PSyclone needs to understand the arguments and datatypes passed by the algorithm layer and the arguments and datatypes expected by the Kernel layer; it needs to know the name of the Kernel subroutine(s); it needs to know the iteration space that the Kernel(s) is/are written to iterate over; it also needs to know the ordering of Kernels and Built-ins as specified in the algorithm layer. Finally, it needs to know where to place any distributed memory operations.

PSyclone determines the above information by being told the API in question (by the user), by reading the appropriate Kernel and Built-in metadata and by reading the order of Kernels and Built-ins in an `invoke` call (as specified in the algorithm layer).

PSyclone has an API-specific parsing stage which reads the algorithm layer and all associated Kernel metadata. This information is passed to a PSy-generation stage which creates a high level view of the PSy layer. From this high level view the PSy-generation stage can generate the required PSy code.

For example, the following Python code shows a code being parsed, a PSy-generation object being created using the output from the parser and the PSy layer code being generated by the PSy-generation object.

```
from psyclone.parse.algorithm import parse
from psyclone.psyGen import PSyFactory

# This example uses version 0.1 of the Dynamo API
api = "dynamo0.1"
```

(continues on next page)

(continued from previous page)

```
# Parse the file containing the algorithm specification and  
# return the Abstract Syntax Tree and invokeInfo objects  
ast, invokeInfo = parse("dynamo.F90", api=api)  
  
# Create the PSy-layer object using the invokeInfo  
psy = PSyFactory(api).create(invokeInfo)  
# Generate the Fortran code for the PSy layer  
print psy.gen
```

9.2 API

The PSy-layer of a single algorithm file is represented by the **PSy** class. The PSy class has an **Invokes** object which contain one or more **Invoke** instances (one for each invoke in the algorithm layer). Each **Invoke** has an **InvokeSchedule** object with the PSyIR tree that describes the PSy layer invoke subroutine. This subroutine is called by the Algorithm layer and itself calls one or more kernels and/or implements any required Built-in operations.

All this classes can be specialised in each PSyclone API to support the specific features of the APIs. The class diagram for the above base classes is shown below using the dynamo0.3 API as an illustration. This class diagram was generated from the source code with pyreverse and edited with inkscape.



The InvokeSchedule can currently contain nodes of type: **Loop**, **Kernel**, **Built-in** (see the *Built-ins* section), **Directive** (of various types), **HaloExchange**, or **GlobalSum** (the latter two are only used if distributed memory is supported and is switched on; see the *Distributed Memory* section). The order of the tree (depth first) indicates the order of the associated Fortran code.

PSyclone will initially create a “vanilla” (functionally correct but not optimised) InvokeSchedule. This “vanilla” InvokeSchedule can be modified by changing the objects within it. For example, the order that two Kernel calls appear in the generated code can be changed by changing their order in the tree. The ability to modify this high level view of a InvokeSchedule allows the PSy layer to be optimised for a particular architecture (by applying optimisations such as blocking, loop merging, inlining, OpenMP parallelisation etc.). The tree could be manipulated directly, however, to simplify optimisation, a set of transformations are supplied. These transformations are discussed in the next section.

9.2.1 InvokeSchedule visualisation

PSyclone supports visualising an InvokeSchedule (or any other PSyIR node) in two ways. First the *view()* method outputs textual information about the contents of a PSyIR node. If we were to look at the LFRic eg6 example we would see the following output:

```
>>> schedule.view()
InvokeSchedule[invoke='invoke_0', dm=True]
  0: Directive[OMP parallel do]
    Schedule[]
      0: Loop[type='dofs', field_space='any_space_1', it_space='dofs', 'upper_'
↳bound='ndofs']
        Literal[value:'NOT_INITIALISED']
        Literal[value:'NOT_INITIALISED']
        Literal[value:'1']
        Schedule[]
          0: BuiltIn setval_X_code(p, z)
          1: BuiltIn X_innerproduct_Y_code(rs_old, res, z)
      1: GlobalSum[scalar='rs_old']
```

The above output tells us that the invoke name for the InvokeSchedule we are looking at is *invoke_0* and that the distributed_memory option has been switched on. Within the InvokeSchedule is an OpenMP parallel directive containing a loop which itself contains two builtin calls. As the latter of the two builtin calls requires a reduction and distributed memory is switched on, PSyclone has added a GlobalSum call for the appropriate scalar.

Second, the *dag()* method (standing for directed acyclic graph), outputs the PSyIR nodes and its data dependencies. By default a file in dot format is output with the name *dag* and a file in svg format is output with the name *dag.svg*. The file name can be changed using the *file_name* optional argument and the output file format can be changed using the *file_format* optional argument. The *file_format* value is simply passed on to graphviz so the graphviz documentation should be consulted for valid formats if svg is not required.

```
>>> schedule.dag(file_name="lovely", file_format="png")
```

Note: The dag method can be called from any node and will output the dag for that node and all of its children.

If we were to look at the LFRic eg6 example we would see the following image:



In the image, all PSyIR nodes with children are split into a start vertex and an end vertex (for example the InvokeSchedule node has both *schedule_start* and *schedule_end* vertices). Blue arrows indicate that there is a parent to child relationship (from a start node) or a child to parent relationship (to an end node). Green arrows indicate that a Node depends on another Node later in the schedule (which we call a forward dependence). Therefore the OMP parallel loop must complete before the globalsum is performed. Red arrows indicate that a Node depends on another Node that is earlier in the schedule (which we call a backward dependence). However the direction of the red arrows are reversed to improve the flow of the dag layout. In this example the forward and backward dependence is the same, however this is not always the case. The two built-ins do not depend on each other, so they have no associated green or red arrows.

The dependence graph output gives an indication of whether nodes can be moved within the InvokeSchedule. In this case it is valid to run the builtin's in either order. The underlying dependence analysis used to create this graph is used to determine whether a transformation of a Schedule is valid from the perspective of data dependencies.

BUILT-INS

Built-ins (named by analogy with the native functionality provided by Python) are operations which can be specified within an invoke call in the algorithm layer but do not require an associated kernel to be implemented as they are provided directly by the infrastructure.

One use of Built-ins is for commonly used operations. In this case Built-ins simplify the use of the system as users do not need to write kernel routines. Built-ins also offer a potential performance advantage as they provide a specification of what is required without an implementation. Therefore the PSy layer is free to implement these operations in whatever way it chooses.

Note: In general, PSyclone will need to know the types of the arguments being passed to any Built-in. The parser obtains this information from an API-specific file that contains the metadata for all Built-in operations supported for that API.

10.1 Example

In the following example, the invoke call includes a call to two Built-ins (`setval_c` and `X_divideby_Y`) and a user-supplied kernel (`matrix_vector_kernel_mm_type`). The `setval_c` Built-in sets all values in the field `Ax` to 1.0 and the `X_divideby_Y` Built-in divides values in the field `rhs` by their equivalent (per degree of freedom) values in the field `lumped_weight` (see [supported Dynamo0.3 API Built-ins](#)). Notice that, unlike the kernel call, no use association is required for the Built-ins since they are provided as part of the environment (c.f. Fortran intrinsics such as `sin()`).

```

module solver_mod
...
use matrix_vector_mm_mod, only: matrix_vector_kernel_mm_type
...

subroutine jacobi_solver_algorithm(lhs, rhs, mm, mesh, n_iter)

  integer(kind=i_def), intent(in)      :: n_iter
  type(field_type),   intent(inout)    :: lhs
  type(field_type),   intent(in)       :: rhs
  type(operator_type), intent(in)       :: mm
  type(mesh_type),    intent(in)       :: mesh
  type(field_type)    :: Ax, lumped_weight, res

  real(kind=r_def), parameter :: MU = 0.9_r_def
...

```

(continues on next page)

(continued from previous page)

```

! Compute mass lump
call invoke( name = "Jacobi_mass_lump",           &
             setval_c(Ax, 1.0_r_def),           &
             matrix_vector_kernel_mm_type(lumped_weight, Ax, mm), &
             X_divideby_Y(lhs, rhs, lumped_weight) )

end subroutine jacobi_solver_algorithm
...
end module solver_mod

```

Below is an example of a kernel that is consistent with the `matrix_vector_kernel_mm_type` kernel specified in the example above.

```

module matrix_vector_mm_mod
  type, public, extends(kernel_type) :: matrix_vector_kernel_mm_type
  private
  type(arg_type) :: meta_args(3) = (/
    arg_type(GH_FIELD,    GH_REAL, GH_INC,  ANY_SPACE_1),
    arg_type(GH_FIELD,    GH_REAL, GH_READ, ANY_SPACE_1),
    arg_type(GH_OPERATOR, GH_REAL, GH_READ, ANY_SPACE_1, ANY_SPACE_1) &
  /)
  integer :: operates_on = CELL_COLUMN
contains
  procedure, nopass :: matrix_vector_mm_code
end type
contains
  subroutine matrix_vector_mm_code(cell,      &
                                   nlayers,   &
                                   lhs, x,    &
                                   ncell_3d,  &
                                   mass_matrix, &
                                   ndf, undf, map)
  end subroutine matrix_vector_mm_code
end module matrix_vector_mm_mod

```

We now translate the algorithm layer code and generate the PSy layer code. The algorithm code is assumed to be in a file called `solver_mod.x90` (see Example 3 in *LFric examples* section). In this case we use the top level Python interface. See the *API* section for different ways to translate/generate code.

```

> psyclone -nodm -oalg solver_mod.f90 -opsy solver_mod_psy.f90 \
> solver_mod.x90

```

The resultant generated algorithm code is given below.

Ignoring the difference in case (which is due to the output format of the code parser) the differences between the original algorithm code and the translated algorithm code are:

- The generic calls to `invoke` have been replaced by specific `CALL invoke_<xx>`. The kernel calls within the original `invoke` are removed, as are duplicate arguments and any literals, leaving the five fields and one operator as arguments;
- A `use` statement is added for each of the new `CALL invoke_<xx>` which will call the generated PSy layer code.

The existence of calls to Built-ins has made no difference at this point:

```

SUBROUTINE jacobi_solver_algorithm(lhs, rhs, mm, mesh, n_iter)
  USE solver_mod_psy, ONLY: invoke_jacobi_iterloop
  USE solver_mod_psy, ONLY: invoke_2l
  USE solver_mod_psy, ONLY: invoke_jacobi_mass_lump

  IMPLICIT NONE

  INTEGER(KIND = i_def), INTENT(IN) :: n_iter
  TYPE(field_type), INTENT(INOUT) :: lhs
  TYPE(field_type), INTENT(IN) :: rhs
  TYPE(operator_type), INTENT(IN) :: mm
  TYPE(mesh_type), INTENT(IN) :: mesh
  TYPE(field_type) :: Ax, lumped_weight, res

  REAL(KIND = r_def), PARAMETER :: MU = 0.9_r_def

  INTEGER(KIND = i_def) :: iter
  INTEGER(KIND = i_def) :: rhs_fs
  TYPE(function_space_type) :: fs

  ...
  CALL invoke_jacobi_mass_lump(ax, lumped_weight, mm, lhs, rhs)
  ...

END SUBROUTINE jacobi_solver_algorithm

```

A vanilla (with no distributed and shared-memory optimisations) version of the generated PSy layer is given below. As expected, the kernel code is called from the PSy layer. However, in the case of the Built-ins, the code for these has been written directly into the PSy layer:

- `setval_c` translates to the loop setting `ax_proxy%data(df) = 1.0_r_def;`
- `X_divideby_Y` translates to the loop setting `lhs_proxy%data(df) = rhs_proxy%data(df) / lumped_weight_proxy%data(df).`

This example illustrates that Built-ins may be implemented in whatever way PSyclone sees fit with no change to the algorithm and kernel layers.

```

MODULE solver_mod_psy
  ...

  SUBROUTINE invoke_jacobi_mass_lump(ax, lumped_weight, mm, lhs, rhs)
    USE matrix_vector_mm_mod, ONLY: matrix_vector_mm_code
    TYPE(field_type), intent(in) :: ax, lumped_weight, lhs, rhs
    TYPE(operator_type), intent(in) :: mm
    ...
    !
    ! Initialise field and/or operator proxies
    !
    ax_proxy = ax%get_proxy()
    lumped_weight_proxy = lumped_weight%get_proxy()
    mm_proxy = mm%get_proxy()
    lhs_proxy = lhs%get_proxy()
    rhs_proxy = rhs%get_proxy()
    !
    ! Initialise number of layers
    !
    nlayers = ax_proxy%vspace%get_nlayers()
  
```

(continues on next page)

(continued from previous page)

```

!
! Look-up dofmaps for each function space
!
map_aspc1_lumped_weight => lumped_weight_proxy%vspace%get_whole_dofmap()
!
! Initialise number of DoFs for aspc1_ax
!
ndf_aspc1_ax = ax_proxy%vspace%get_ndf()
undf_aspc1_ax = ax_proxy%vspace%get_undf()
!
! Initialise number of DoFs for aspc1_lumped_weight
!
ndf_aspc1_lumped_weight = lumped_weight_proxy%vspace%get_ndf()
undf_aspc1_lumped_weight = lumped_weight_proxy%vspace%get_undf()
!
! Initialise number of DoFs for aspc1_lhs
!
ndf_aspc1_lhs = lhs_proxy%vspace%get_ndf()
undf_aspc1_lhs = lhs_proxy%vspace%get_undf()
!
! Call our kernels
!
DO df=1,undf_aspc1_ax
  ax_proxy%data(df) = 1.0_r_def
END DO
DO cell=1,lumped_weight_proxy%vspace%get_ncell()
  !
  CALL matrix_vector_mm_code(cell, nlayers, &
                             lumped_weight_proxy%data, &
                             ax_proxy%data, &
                             mm_proxy%ncell_3d, &
                             mm_proxy%local_stencil, &
                             ndf_aspc1_lumped_weight, &
                             undf_aspc1_lumped_weight, &
                             map_aspc1_lumped_weight(:,cell))

  END DO
DO df=1,undf_aspc1_lhs
  lhs_proxy%data(df) = rhs_proxy%data(df) / lumped_weight_proxy%data(df)
END DO
!
END SUBROUTINE invoke_jacobi_mass_lump
...
END MODULE solver_mod_psy

```

This example is distributed with PSyclone and can be found in <PSYCLONEHOME>/examples/lfric/eg3.

10.2 Supported Built-in operations

The list of supported Built-ins is API-specific and therefore is described under the documentation of each API.

10.3 Adding new Built-in operations

1. Identify the PSyclone source file for the API to be extended. *e.g.* for Dynamo0.3 (LFRic) API it is `src/psyclone/dynamo0p3_builtins.py`.
2. Edit this source file to create the class for this new call. It must inherit from the API-specific parent class for Built-in operations (`DynBuiltInKern` for Dynamo0.3).
3. Implement `__str__` and `gen_code()` methods for this new class.
4. Add the name of the new Built-in operation and its corresponding class to the `BUILTIN_MAP` dictionary in that source file.
5. Add metadata describing this call to the appropriate file specified in the `BUILTIN_DEFINITIONS_FILE` in that source file. For Dynamo0.3 this is `src/psyclone/parse/dynamo0p3_builtins_mod.f90`.
6. Add relevant tests to the PSyclone test file for the API to be extended. *e.g.* for Dynamo0.3 it is `src/psyclone/tests/dynamo0p3_builtins_test.py`. The tests rely on `single_invoke` Fortran examples in the relevant `src/psyclone/tests/test_files/` subfolder.
7. Add an appropriate Fortran `single_invoke` example for the new Built-in in the relevant `src/psyclone/tests/test_files/` subfolder. *e.g.* for Dynamo0.3 it is `src/psyclone/tests/test_files/dynamo0p3/`. Names of examples follow the template `<category.number>.<subcategory.number>_<built-in_name>.f90`. *e.g.* for Dynamo0.3 API `<category.number>` is 15 and `<built-in_name>` follows the *Dynamo0.3 API Built-in naming scheme*.
8. Document the new Built-in in the documentation of the relevant API (*e.g.* `doc/dynamo0p3.rst` for Dynamo0.3 API).

If the API being extended does not currently support any Built-ins then the `BUILTIN_MAP` and `BUILTIN_DEFINITIONS_FILE` module variables must be added to the source file for the API. A Fortran module file must be created in the PSyclone `src` directory (with the name specified in `BUILTIN_DEFINITIONS_FILE`) containing metadata describing the Built-in operations. Finally, `parse.get_builtin_defs()` must be extended to import `BUILTIN_MAP` and `BUILTIN_DEFINITIONS_FILE` for this API.

DYNAMO0.3 API

This section describes the Dynamo0.3 application programming interface (API). This API explains what a user needs to write in order to make use of the Dynamo0.3 API in PSyclone.

As with the majority of PSyclone APIs, the Dynamo0.3 API specifies how a user needs to write the algorithm layer and the kernel layer to allow PSyclone to generate the PSy layer. These algorithm and kernel APIs are discussed separately in the following sections.

The Dynamo0.3 API supports the Met Office's finite element (hereafter FEM) based GungHo dynamical core (see [Introduction](#)). The Met Office Collaboration Wiki (login required) provides more detailed information about the dynamical core's [formulation](#) and [data model](#).

The GungHo dynamical core with atmospheric physics parameterisation schemes is a part of the Met Office LFRic modelling system [AFH+19], currently being developed in preparation for exascale computing in the 2020s. The LFRic repository and the associated wiki are hosted at the [Met Office Science Repository Service](#). The code is BSD-licensed, however browsing the [LFRic wiki](#) and [code repository](#) requires login access to MOSRS.

11.1 Algorithm

The general requirements for the structure of an Algorithm are explained in the [Algorithm layer](#) section. This section explains the Dynamo0.3-API-specific specialisations and extensions.

11.1.1 Example

An example Dynamo0.3 API invoke call is given below with various different types of objects supported by the API. These different objects and their use are discussed in the following sections.

```
real(kind=r_def)          :: scalar1
integer(kind=i_def)       :: stencil_extent
type(field_type)          :: field1, field2, field3
type(field_type)          :: field5(3), field6(3)
type(quadrature_type)     :: qr
type(operator_type)       :: operator1
type(columnwise_operator_type) :: cma_op1
...
call invoke( kernel1(field1, field2, operator1, qr),          &
             builtin1(scalar1, field2, field3),             &
             kernel2(field1, stencil_extent, field3, scalar1), &
             assembly_kernel(cma_op1, operator1),           &
             name="some_calculation"                        &
             )
call invoke( prolong_kernel_type(field1, field4),            &
```

(continues on next page)

(continued from previous page)

```

        restrict_kernel_type(field5, field6)
    )

```

Please see the *Algorithm layer* section for a description of the `name` argument.

Objects in the Dynamo0.3 API can be categorised by their functionality as data structures and information that specifies supported operations on a particular data structure. These data structures are represented by the five Dynamo0.3 API argument types: *scalar*, *field*, *field vector*, *operator* and *column-wise operator*. All of them except the field vector are represented in the above example. `qr` represents a quadrature object which provides information required by a kernel to operate on fields (see section *Quadrature* for more details).

11.1.2 Scalar

In the Dynamo0.3 API a scalar is a single-valued argument that can be *either real or integer*. Scalars are identified with `GH_SCALAR` metadata.

11.1.3 Field

Dynamo0.3 API fields, identified with `GH_FIELD` metadata, represent FEM discretisations of various dynamical core prognostic and diagnostic variables. In FEM, variables are discretised by placing them into a function space (see *Supported Function Spaces*) from which they inherit a polynomial expansion via the basis functions of that space. Field values at points within a cell are evaluated as the sum of a set of basis functions multiplied by coefficients which are the data points. Points of evaluation are determined by a quadrature object (*Quadrature*) and are independent of the function space the field is on. Placement of field data points, also called degrees of freedom (hereafter “DoFs”), is determined by the function space the field is on.

11.1.4 Field Vector

Depending on the function space a field lives on, the field data at a point can be a scalar or a vector (see *Supported Function Spaces* for the list of scalar and vector function spaces). There is an additional option which specifies whether the data itself is vector valued and it usually refers to bundles of scalar valued fields. Field vectors are represented as `GH_FIELD*N` where `N` is the size of the vector. The 3D coordinate field, for example, has (`x`, `y`, `z`) scalar values at the nodes and therefore has a vector size of 3.

11.1.5 Operator

Represents a matrix constructed on a per-cell basis using Local Matrix Assembly (LMA) and is identified with `GH_OPERATOR` metadata.

11.1.6 Column-Wise Operator

The Dynamo0.3 API has support for the construction and use of column-wise/Column Matrix Assembly (CMA) operators whose metadata identifier is `GH_COLUMNWISE_OPERATOR`. As the name suggests, these are operators constructed for a whole column of the mesh. These are themselves constructed from the Local Matrix Assembly (LMA) operators of each cell in the column. The rules governing Kernels that have CMA operators as arguments are given in the *Kernel* section below.

There are three recognised Kernel types involving CMA operations; construction, application (including inverse application) and matrix-matrix. The following example sketches-out what the use of such kernels might look like in the Algorithm layer:


```

use field_mod, only: field_type
use operator_mod, only : operator_type, columnwise_operator_type
type(field_type) :: field1, field2, field3
type(operator_type) :: lma_op1, lma_op2
type(columnwise_operator_type) :: cma_op1, cma_op2, cma_op3
real(kind=r_def) :: alpha
...
call invoke(
    assembly_kernel(cma_op1, lma_op1, lma_op2),
    assembly_kernel2(cma_op2, lma_op1, lma_op2, field3),
    apply_kernel(field1, field2, cma_op1),
    matrix_matrix_kernel(cma_op3, cma_op1, alpha, cma_op2),
    apply_kernel(field3, field1, cma_op3),
    name="cma_example")

```

The above invoke uses two LMA operators to construct the CMA operator `cma_op1`. A second CMA operator, `cma_op2`, is assembled from the same two LMA operators but also uses a field. The first of these CMA operators is then applied to `field2` and the result stored in `field1` (assuming that the metadata for `apply_kernel` specifies that it is the first field argument that is written to). The two CMA operators are then combined to produce a third, `cma_op3`. This is then applied to `field1` and the result stored in `field3`.

Note that PSyclone identifies the type of kernels performing Column-Wise operations based on their arguments as described in metadata (see [Rules for Kernels that work with CMA Operators](#) below). The names of the kernels in the above example are purely illustrative and are not used by PSyclone when determining kernel type.

A full example of CMA operator construction is available in `examples/lfric/eg7`.

11.1.7 Quadrature

Kernels conforming to the Dynamo0.3 API may require quadrature information (specified using e.g. `gh_shape = gh_quadrature_XYoZ` in the kernel metadata - see Section [gh_shape and gh_evaluator_targets](#)). This information must be passed to the kernel from the Algorithm layer in the form of one or more `quadrature_type` objects. These must be the last arguments passed to the kernel and must be provided in the same order that they are specified in the kernel metadata, e.g. if the metadata for kernel `pressure_gradient_kernel_type` specified `gh_shape = gh_quadrature_XYoZ` and that for kernel `geopotential_gradient_kernel` had `gh_shape(2) = (\ gh_quadrature_XYoZ, gh_quadrature_face \)` then the corresponding invoke would look something like:

```

...
qr_xyoz = quadrature_xyoz_type(nqp_exact, rule)
qr_face = quadrature_face_type(nqp_exact, ..., rule)
call invoke(pressure_gradient_kernel_type(rhs_tmp(igh_u), rho, theta, qr_xyoz), &
    geopotential_gradient_kernel_type(rhs_tmp(igh_u), geopotential, &
    qr_xyoz, qr_face))

```

These quadrature objects specify the set(s) of points at which the basis/differential-basis functions required by the kernel are to be evaluated.

11.1.8 Stencils

The metadata for a Kernel which operates on a cell-column may specify that a Kernel performs a stencil operation on a field. Any such metadata must provide a stencil type. See the *meta_args* section for more details. The supported stencil types are X1D, Y1D, XORY1D, CROSS or CROSS2D.

If a stencil operation is specified by the Kernel metadata the algorithm layer must provide the `extent` of the stencil (the maximum distance from the central cell that the stencil extends). The Dynamo0.3 API expects this information to be added as an additional `integer` argument immediately after the relevant field when specifying the Kernel via an `invoke`.

For example:

```
integer(kind=i_def) :: extent = 2
call invoke(kernel(field1, field2, extent))
```

where `field2` has kernel metadata specifying that it has a stencil access.

`extent` may also be passed as a literal. For example:

```
call invoke(kernel(field1, field2, 2))
```

where, again, `field2` has kernel metadata specifying that it has a stencil access.

Note: The stencil extent specified in the Algorithm layer is not the same as the stencil size passed in to the Kernel. The latter contains the number of cells in the stencil which is dependent on both the stencil type and extent.

If the Kernel metadata specifies that the stencil is of type XORY1D (which means X1D or Y1D) then the algorithm layer must specify whether the stencil is X1D or Y1D for that particular kernel call. The Dynamo0.3 API expects this information to be added as an additional argument immediately after the relevant stencil extent argument. The argument should be an `integer` with valid values being `x_direction` or `y_direction`, both being supplied by the LFRic infrastructure via the `flux_direction_mod` fortran module

For example:

```
use flux_direction_mod, only : x_direction
integer(kind=i_def) :: direction = x_direction
integer(kind=i_def) :: extent = 2
! ...
call invoke(kernel(field1, field2, extent, direction))
```

`direction` may also be passed as a literal. For example:

```
use flux_direction_mod, only : x_direction
integer(kind=i_def) :: extent = 2
! ...
call invoke(kernel(field1, field2, extent, x_direction))
```

If the stencil is of type CROSS2D then the arrays passed to the kernel are of different dimensions to those of other stencils. The CROSS2D stencil is designed for use when it is necessary for a kernel to know where the stencil cells are, relative to the current cell. For this reason, the `stencil_size` passed to the kernel is an array of length 4 containing sizes for each branch of the stencil. The `stencil_size` array is always ordered: West, South, East, North. This branch dimension is also part of the `stencil_dofmap` array making it possible to loop over each branch of the stencil individually. The `invoke` call for the CROSS2D stencil remains of the same form as for other stencils.

If certain fields use the same value of extent and/or direction then the same variable, or literal value can be provided.

For example:

```
call invoke(kernel1(field1, field2, extent, field3, extent, direction), &
           kernel2(field1, field2, extent2, field4, extent, direction))
```

In the above example `field2` and `field3` in `kernel1` and `field4` in `kernel2` will have the same extent value but `field2` in `kernel2` may have a different value. Similarly, `field3` in `kernel1` and `field4` in `kernel2` will have the same direction value.

An example of the use of stencils is available in `examples/lfric/eg5`.

There is currently no attempt to perform type checking in PSyclone so any errors in the type and/or position of arguments will not be picked up until compile time. However, PSyclone does check for the correct number of algorithm arguments. If the wrong number of arguments is provided then an exception is raised.

For example, running test 19.2 from the Dynamo0.3 API test suite gives:

```
cd <PSYCLONEHOME>/src/psyclone/tests
psyclone test_files/dynamo0p3/19.2_single_stencil_broken.f90
"Generation Error: error: expected '5' arguments in the algorithm layer but found '4'.
Expected '4' standard arguments, '1' stencil arguments and '0' qr_arguments"
```

11.1.9 Inter-grid

From the Algorithm layer, an `Invoke` for inter-grid kernels (those that map fields between grids of different resolution) looks much like an `Invoke` containing general-purpose kernels. The only restrictions to be aware of are that inter-grid kernels accept only field or field-vectors as arguments and that an `Invoke` may not mix inter-grid kernels with any other kernel type. (Hence the second, separate `Invoke` in the example Algorithm code given at the beginning of this Section.)

11.2 PSy-layer

The general details of the PSy-layer are explained in the *PSy layer* section. This section describes any dynamo0p3 specific issues.

11.2.1 Module name

The PSy-layer code is contained within a Fortran module. The name of the module is determined from the algorithm-layer name with “_psy” appended. The algorithm-layer name is the algorithm’s module name if it is a module, its subroutine name if it is a subroutine that is not within a module, or the program name if it is a program.

So, for example, if the algorithm code is contained within a module called “fred” then the PSy-layer module name will be “fred_psy”.

Argument Intents

LFRic *fields*, *field vectors*, *operators* and *column-wise operators* are objects that contain pointers to data rather than data. The data are accessed by proxies of these objects and modified in *kernels*. As the objects themselves are not modified in the PSy layer, their Fortran intents there are always `intent(in)`.

The Fortran intent of *scalars* is still defined by their *access metadata* as they are actual data. This means `intent(in)` for `GH_READ` and `intent(out)` for `GH_SUM` (more details in *meta_args* section below).

The intent of other data structures is mandated by the relevant Dynamo0.3 API rules described in sections below.

11.3 Kernel

The general requirements for the structure of a Kernel are explained in the *Kernel layer* section. In the LFRic API there are five different Kernel types; general purpose, CMA, inter-grid, domain and *Built-ins*. In the case of built-ins, PSyclone generates the source of the kernels. This section explains the rules for the other four, user-supplied kernel types and then goes on to describe their metadata and subroutine arguments. Domain kernels are distinct from the other three user-supplied kernel types in that they must be passed data for the whole domain rather than a single cell-column.

11.3.1 Rules for all User-Supplied Kernels that Operate on Cell-Columns

In the following, ‘operator’ refers to both LMA and CMA operator types.

- 1) A Kernel must have at least one argument that is a field, field vector, or operator. This rule reflects the fact that a Kernel operates on some subset of the whole domain (e.g. a cell-column) and is therefore designed to be called from within a loop that iterates over those subsets of the domain.
- 2) The continuity of the iteration space of the Kernel is determined from the function space of the modified argument (see Section *Supported Function Spaces* below). If more than one argument is modified then the iteration space is taken to be the largest required by any of those arguments. E.g. if a Kernel writes to two fields, the first on W3 (discontinuous) and the second on W1 (continuous), then the iteration space of that Kernel will be determined by the field on the continuous space.
- 3) If any of the modified arguments are declared with the generic function space metadata (e.g. ANY_SPACE_n, see *Supported Function Spaces*) and their actual space cannot be determined statically then the iteration space is assumed to be
 - 1) discontinuous for ANY_DISCONTINUOUS_SPACE_n;
 - 2) continuous for ANY_SPACE_n and ANY_W2. This assumption is always safe but leads to additional computation if the quantities being updated are actually on discontinuous function spaces.
- 4) Operators do not have halo operations operating on them as they are either cell- (LMA) or column-based (CMA) and therefore act like discontinuous fields.
- 5) Any Kernel that writes to an operator will have its iteration space expanded such that valid values for the operator are computed in the level-1 halo.
- 6) Any Kernel that reads from an operator must not access halos beyond level 1. In this case PSyclone will check that the Kernel does not require values beyond the level-1 halo. If it does then PSyclone will abort.

11.3.2 Rules specific to General-Purpose Kernels without CMA Operators

- 1) General-purpose kernels with `operates_on = CELL_COLUMN` accept arguments of any of the following types: field, field vector, LMA operator, scalar integer, scalar real.
- 2) A Kernel is permitted to write to more than one quantity (field or operator) and these quantities may be on the same or different function spaces.
- 3) A Kernel may not write to a scalar argument. (Only *built-ins* are permitted to do this.) Any scalar arguments must therefore be declared in the metadata as `GH_READ` - see *below*.

11.3.3 Rules for Kernels that work with CMA Operators

The Dynamo0.3 API has support for kernels that assemble, apply (or inverse-apply) column-wise/Column Matrix Assembly (CMA) operators. Such operators may also be used by matrix-matrix kernels. There are thus three types of CMA-related kernels. Since, by definition, CMA operators only act on data within a column, they have no horizontal dependencies. Therefore, kernels that write to them may be parallelised without colouring.

All three CMA-related kernel types must obey the following rules:

- 1) Since a CMA operator only acts within a single column of data, stencil operations are not permitted.
- 2) No vector quantities (e.g. `GH_FIELD*3` - see below) are permitted as arguments.
- 3) The kernel must operate on cell-columns.

There are then additional rules specific to each of the three CMA kernel types. These are described below.

Assembly

CMA operators are themselves constructed from Local-Matrix-Assembly (LMA) operators. Therefore, any kernel which assembles a CMA operator must obey the following rules:

- 1) Have one or more LMA operators as read-only arguments.
- 2) Have exactly one CMA operator argument which must have write access.
- 3) Other types of argument (e.g. scalars or fields) are permitted but must be read-only.

Application and Inverse Application

Column-wise operators can only be applied to fields. CMA-Application kernels must therefore:

- 1) Have a single CMA operator as a read-only argument.
- 2) Have exactly two field arguments, one read-only and one that is written to.
- 3) The function spaces of the read and written fields must match the from and to spaces, respectively, of the supplied CMA operator.

Matrix-Matrix

A kernel that has just column-wise operators as arguments and zero or more read-only scalars is identified as performing a matrix-matrix operation. In this case:

- 1) Arguments must be CMA operators and, optionally, one or more scalars.
- 2) Exactly one of the CMA arguments must be written to while all other arguments must be read-only.

11.3.4 Rules for Inter-Grid Kernels

- 1) An inter-grid kernel is identified by the presence of a field or field-vector argument with the optional *mesh_arg* metadata element (see *Inter-Grid Metadata*).
- 2) An invoke that contains one or more inter-grid kernels must not contain any other kernel types. (This restriction is an implementation decision and could be lifted in future if there is a need.)
- 3) An inter-grid kernel is only permitted to have field or field-vector arguments.
- 4) All inter-grid kernel arguments must have the *mesh_arg* metadata entry.
- 5) An inter-grid kernel (and metadata) must have at least one field on each of the fine and coarse meshes. Specifying all fields as coarse or fine is forbidden.
- 6) Fields on different meshes must always live on different function spaces.
- 7) All fields on a given mesh must be on the same function space.
- 8) An inter-grid kernel must operate on cell-columns.

A consequence of Rules 5-7 is that an inter-grid kernel will only involve two function spaces.

11.3.5 Rules for User-Supplied Kernels that Operate on the Domain

The rules for kernels that have `operates_on = DOMAIN` are a subset of *those* for kernels that operate on a `CELL_COLUMN` without CMA Operators. Specifically:

- 1) Only scalar, field and field vector arguments are permitted.
- 2) All fields must be on discontinuous function spaces.
- 3) Stencil accesses are not permitted.

11.3.6 Metadata

The code below outlines the elements of the Dynamo0.3 API Kernel metadata, 1) '*meta_args*', 2) '*meta_funcs*', 3) '*meta_reference_element*', 4) '*meta_mesh*', 5) '*gh_shape*' (*gh_shape* and *gh_evaluator_targets*), 6) '*operates_on*' and 7) '*procedure*':

```
type, public, extends(kernel_type) :: my_kernel_type
  type(arg_type) :: meta_args(...) = (/ ... /)
  type(func_type) :: meta_funcs(...) = (/ ... /)
  type(reference_element_data_type) :: meta_reference_element(...) = (/ ... /)
  type(mesh_data_type) :: meta_mesh(...) = (/ ... /)
  integer :: gh_shape = gh_quadrature_XYoZ
  integer :: operates_on = cell_column
contains
  procedure, nopass :: my_kernel_code
end type
```

These various metadata elements are discussed in order in the following sections.

meta_args

The `meta_args` array specifies information about data that the kernel code expects to be passed to it via its argument list. There is one entry in the `meta_args` array for each **scalar**, **field**, or **operator** passed into the Kernel and the order that these occur in the `meta_args` array must be the same as they are expected in the kernel code argument list. The entry must be of `arg_type` which itself contains metadata about the associated argument. The size of the `meta_args` array must correspond to the number of **scalars**, **fields** and **operators** passed into the Kernel.

Note: It makes no sense for a Kernel to have only **scalar** arguments (because the PSy layer will call a Kernel for each point in the spatial domain) and PSyclone will reject such Kernels.

For example, if there are a total of 2 **scalar / field / operator** entities being passed to the Kernel then the `meta_args` array will be of size 2 and there will be two `arg_type` entries:

```
type(arg_type) :: meta_args(2) = (/  
    arg_type( ... ),  
    arg_type( ... )  
/)
```

Argument metadata (information contained within the brackets of an `arg_type` entry), describes either a **scalar**, a **field** or an **operator** (either LMA or CMA).

The first argument-metadata entry describes whether the data that is being passed is for a scalar (`GH_SCALAR`), a field (`GH_FIELD`) or an operator (either `GH_OPERATOR` for LMA or `GH_COLUMNWISE_OPERATOR` for CMA). This information is mandatory.

Additionally, argument metadata can be used to describe a vector of fields (see the *Field Vector* section for more details).

As an example, the following `meta_args` metadata describes 4 entries, the first is a scalar, the next two are fields and the fourth is an operator. The third entry is a field vector of size 3.

```
type(arg_type) :: meta_args(4) = (/  
    arg_type(GH_SCALAR, GH_REAL, ...),  
    arg_type(GH_FIELD, GH_INTEGER, ... ),  
    arg_type(GH_FIELD*3, GH_REAL, ... ),  
    arg_type(GH_OPERATOR, GH_REAL, ... )  
/)
```

The second item in a metadata entry describes the Fortran primitive (intrinsic) type of the data of a kernel argument. The currently supported values are `GH_REAL` and `GH_INTEGER` for real and integer data, respectively. This information is mandatory. Valid data types for each Dynamo0.3 API argument type are specified later in this section (see *Valid Data Types*).

The third component of argument metadata describes how the Kernel makes use of the data being passed into it (the way it is accessed within a Kernel). This information is mandatory. There are currently 5 possible values of this metadata `GH_WRITE`, `GH_READ`, `GH_INC`, `GH_READWRITE` and `GH_SUM`. However, not all combinations of metadata entries are valid and PSyclone will raise an exception if an invalid combination is specified. Valid combinations are specified later in this section (see *Valid Access Modes*).

- `GH_WRITE` indicates the data is modified in the Kernel before (optionally) being read.
- `GH_READ` indicates that the data is read and is unmodified.
- `GH_INC` indicates that different iterations of a Kernel make contributions to shared values. For example, values at cell faces may receive contributions from cells on either side of the face. This means that such a Kernel needs appropriate synchronisation (or colouring) to run in parallel.

- `GH_READWRITE` indicates that different iterations of a Kernel update quantities which do not share DoFs, such as operators and fields over discontinuous function spaces. If a Kernel modifies only discontinuous fields and/or operators there is no need for synchronisation or colouring when running such Kernels in parallel. However, modifying another field with a `GH_INC` access in a Kernel means that synchronisation or colouring is required for parallel runs.
- `GH_SUM` is an example of a reduction and is the only reduction currently supported in PSyclone. This metadata indicates that values are summed over calls to Kernel code.

For example:

```
type(arg_type) :: meta_args(4) = (/
    arg_type(GH_SCALAR, GH_REAL, GH_SUM),
    arg_type(GH_FIELD, GH_INTEGER, GH_INC, ... ),
    arg_type(GH_FIELD*3, GH_REAL, GH_WRITE, ... ),
    arg_type(GH_OPERATOR, GH_REAL, GH_READ, ... )
/)
```

Note: In the Dynamo0.3 API only *Built-ins* are permitted to write to scalar arguments (and hence perform reductions). Furthermore, this permission is currently restricted to real scalars (`GH_SCALAR`, `GH_REAL`) as the LFRic infrastructure does not yet support integer reductions.

For a scalar the argument metadata contains only these three entries. However, fields and operators require further entries specifying function-space information. The meaning of these further entries differs depending on whether a field or an operator is being described.

In the case of an operator, the fourth and fifth arguments describe the `to` and `from` function spaces respectively. In the case of a field the fourth argument specifies the function space that the field lives on. More details about the supported function spaces are in subsection *Supported Function Spaces*.

For example, the metadata for a kernel that applies a Column-wise operator to a field might look like:

```
type(arg_type) :: meta_args(3) = (/
    arg_type(GH_FIELD, GH_REAL, GH_INC, W1),
    arg_type(GH_FIELD, GH_REAL, GH_READ, W2H),
    arg_type(GH_COLUMNWISE_OPERATOR, GH_REAL, GH_READ, W1, W2H)
/)
```

In some cases a Kernel may be written so that it works for fields and/or operators from any type of a vector `W2*` space (all `W2*` spaces except for the `W2*trace` spaces, see Section *Supported Function Spaces* below). In this case the metadata should be specified as being `ANY_W2`.

Warning: In the current implementation it is assumed that all fields and/or operators specifying `ANY_W2` within a kernel will use the **same** function space. It is up to the user to ensure this is the case as otherwise invalid code would be generated.

It may be that a Kernel is written such that a field and/or operators may be on/map-between any function space(s). In this case the metadata should be specified as being one of `ANY_SPACE_1, ..., ANY_SPACE_10` (see *Supported Function Spaces*). If the generic function spaces are known to be discontinuous the metadata may be specified as being one of `ANY_DISCONTINUOUS_SPACE_1, ..., ANY_DISCONTINUOUS_SPACE_10` in order to avoid unnecessary computation into the halos (see rules for *user-supplied kernels* above). The reason for having different names is that a Kernel might be written to allow 2 or more arguments to be able to support any function space but for a particular call the function spaces may have to be the same as each other.

In the example below, the first field entry supports any function space but it must be the same as the operator's `to` function space. Similarly, the second field entry supports any function space but it must be the same as the operator's `from` function space. Note, the metadata does not forbid `ANY_SPACE_1` and `ANY_SPACE_2` from being the same.

```
type(arg_type) :: meta_args(3) = (/
    arg_type(GH_FIELD,    GH_REAL, GH_INC,  ANY_SPACE_1),
    arg_type(GH_FIELD*3,  GH_REAL, GH_INC,  ANY_SPACE_2),
    arg_type(GH_OPERATOR, GH_REAL, GH_READ, ANY_SPACE_1, ANY_SPACE_2) &
    /)
```

Note also that the scope of this naming of any-space function spaces is restricted to the argument list of individual kernels. I.e. if an `Invoke` contains say, two kernel calls that each support arguments on any function space, e.g. `ANY_SPACE_1`, there is no requirement that these two function spaces be the same. Put another way, if an `Invoke` contained two calls of a kernel with arguments described by the above metadata then the first field argument passed to each kernel call need not be on the same space.

Valid Data Types

As mentioned earlier, the currently supported Fortran primitive (intrinsic) types for kernel argument data are `real` and `integer`, described by the `GH_REAL` and `GH_INTEGER` metadata descriptors. Supported data types for each argument type are given in the table below (please note that *field vectors* follow the same rules as the *LFRic fields*):

Argument Type	Data Type
<code>GH_SCALAR</code>	<code>GH_REAL</code> , <code>GH_INTEGER</code>
<code>GH_FIELD</code>	<code>GH_REAL</code> , <code>GH_INTEGER</code>
<code>GH_OPERATOR</code>	<code>GH_REAL</code>
<code>GH_COLUMNWISE_OPERATOR</code>	<code>GH_REAL</code>

Valid Access Modes

As mentioned earlier, not all combinations of metadata are valid. Valid combinations for each argument type in user-defined Kernels are summarised here. All argument types (`GH_SCALAR`, `GH_FIELD`, `GH_OPERATOR` and `GH_COLUMNWISE_OPERATOR`) may be read within a Kernel and this is specified in metadata using `GH_READ`. At least one kernel argument must be listed as being modified. When data is *modified* in a user-supplied Kernel (i.e. a Kernel that operates on a `CELL_COLUMN`, see *iteration space metadata*) then the permitted access modes depend upon the argument type and the function space it is on:

Argument Type	Function Space	Access Type
<code>GH_SCALAR</code>	n/a	<code>GH_READ</code>
<code>GH_FIELD</code>	Discontinuous	<code>GH_READ</code> , <code>GH_WRITE</code> , <code>GH_READWRITE</code>
<code>GH_FIELD</code>	Continuous	<code>GH_READ</code> , <code>GH_INC</code>
<code>GH_OPERATOR</code>	Any for both 'to' and 'from'	<code>GH_READ</code> , <code>GH_WRITE</code> , <code>GH_READWRITE</code>
<code>GH_COLUMNWISE_OPERATOR</code>	Any for both 'to' and 'from'	<code>GH_READ</code> , <code>GH_WRITE</code> , <code>GH_READWRITE</code>

Note that scalar arguments to user-defined Kernels must be read-only. Only *Built-ins* are permitted to modify scalar arguments. In practice this means that the only allowed access for the scalars in user-defined Kernels is `GH_READ` (see the allowed accesses for arguments in Built-ins in the *section below*).

Note also that a `GH_FIELD` argument that has `GH_WRITE` or `GH_READWRITE` as its access pattern must be on a horizontally-discontinuous function space (see *Supported Function Spaces* for the list of discontinuous function spaces). Parallelisation of the loop over the horizontal domain for a kernel that updates such a field will not require colouring for either of the above cases (since there are no shared entities).

If a field is described as being on `ANY_SPACE`, there is currently no way to determine its continuity from the metadata (unless we can statically determine the space of the field being passed in). At the moment this type of a user-supplied Kernel is always treated as if it is updating a field that is on a function space that is continuous in the horizontal, even if it is not (see rules for *user-supplied kernels* above).

There is no restriction on the number and function spaces of other quantities that a general-purpose kernel can modify other than that it must modify at least one. The rules for kernels involving CMA operators, however, are stricter and only one argument may be modified (the CMA operator itself for assembly, a field for CMA-application and a CMA operator for matrix-matrix kernels). If a kernel writes to quantities on different function spaces then PSyclone generates loop bounds appropriate to the largest iteration space. This means that if a single kernel updates one quantity on a continuous function space and one on a discontinuous space then the resulting loop will include cells in the level-1 halo since they are required for a quantity on a continuous space. As a consequence, any quantities on a discontinuous space will then be computed redundantly in the level-1 halo. Currently PSyclone makes no attempt to take advantage of this (by e.g. setting the appropriate level-1 halo to ‘clean’).

PSyclone ensures that both CMA and LMA operators are computed (redundantly) out to the level-1 halo cells. This permits their use in kernels which modify quantities on continuous function spaces and also in subsequent redundant computation of other quantities on discontinuous function spaces. In conjunction with this, PSyclone also checks (when generating the PSy layer) that any kernels which read operator values do not do so beyond the level-1 halo. If any such accesses are found then PSyclone aborts.

Supported Function Spaces

As mentioned in the *Field* and *Field Vector* sections, the function space of an argument specifies how it maps onto the underlying topology and, additionally, whether the data at a point is a vector. In Dynamo0.3 API the dimension of the basis function set for the scalar function spaces is 1 and for the vector function spaces is 3 (see the table in *Rules for General-Purpose Kernels* for the dimensions of the basis and differential basis functions).

Function spaces can share DoFs between cells in the horizontal, vertical or both directions. Depending on the function space and FEM order, the shared DoFs can lie on one or more cell entities (faces, edges and vertices) in each direction. This property is referred to as the **continuity** of a function space (horizontal, vertical or full). Alternatively, if there are no shared DoFs a function space is described as **discontinuous** (fully or in a particular direction).

The mixed FEM formulation is built on a foundation set of four function spaces described below.

- `W0` is the space of scalar functions with full continuity. The shared DoFs lie on cell vertices in the lowest order FEM and on all three entities in higher order FEM.
- `W1` is the space of vector functions with full continuity in the tangential direction only. In the lowest order FEM the shared DoFs lie on cell edges for each component, whereas in higher order they also lie on cell faces.
- `W2` is the space of vector functions with full continuity in the normal direction only. The shared DoFs lie on cell faces for each component.
- `W3` is the space of scalar functions with full discontinuity. All DoFs lie within the cell volume and are not shared across the cell boundaries.

Other spaces required for representation of scalar or component-wise vector variables are:

- `Wtheta` is the space of scalar functions based on the vertical part of `W2`, discontinuous in the horizontal and continuous in the vertical;
- `W2H` is the space of vector functions based on the horizontal part of `W2`, continuous in the horizontal and discontinuous in the vertical;
- `W2V` is the space of vector functions based on the vertical part of `W2`, discontinuous in the horizontal and continuous in the vertical;

- `W2broken` is the space of vector functions, locally identical to the `W2` space. However, DoFs are topologically discontinuous in all directions despite their placement on cell faces;
- `W2trace` is the space of scalar functions defined only on cell faces, resulting from taking the trace of a `W2` space. DoFs are shared between faces, hence making this space fully continuous;
- `W2Htrace` is the space of scalar functions defined only on cell faces in the horizontal, resulting from taking the trace of a `W2H` space. DoFs are shared between horizontal faces, hence making this space continuous in the horizontal and discontinuous in the vertical;
- `W2Vtrace` is the space of scalar functions defined only on cell faces in the vertical, resulting from taking the trace of a `W2V` space. DoFs are shared between vertical faces, hence making this space discontinuous in the horizontal and continuous in the vertical;
- `Wchi` is the space of scalar functions used to store coordinates in LFRic. It is fully discontinuous except for the coordinate order 0 when it becomes the `W0` space (i.e. fully continuous). Please see the next section for more details on this function space.

In addition to the specific function space metadata, there are also three generic function space metadata descriptors mentioned in sections above:

- `ANY_SPACE_n`, $n = 1, 2, \dots, 10$, for when the function space of the modified argument(s) cannot be determined and/or for when a Kernel has been written so that it works with fields on any of the available spaces;
- `ANY_DISCONTINUOUS_SPACE_n`, $n = 1, 2, \dots, 10$, for when the function space of the modified argument(s) cannot be determined but is known to be discontinuous and/or for when a Kernel has been written so that it works with fields on any of the discontinuous spaces;
- `ANY_W2` for any type of a vector `W2*` function space, i.e. `W2`, `W2H`, `W2V` and `W2broken` but not `W2*trace` spaces.

As mentioned *previously*, `ANY_SPACE_n` and `ANY_W2` function space types are treated as continuous while `ANY_DISCONTINUOUS_SPACE_n` spaces are treated as discontinuous.

Note: The name and use of `ANY_W2` metadata (e.g. continuity and vector or/and scalar basis of `W2*` spaces the metadata can represent) are being reviewed in PSyclone issue #540.

Since the Dynamo0.3 API operates on columns of data, function spaces are categorised as continuous or discontinuous with regard to their **continuity in the horizontal**. For example, a `GH_FIELD` that specifies `GH_INC` as its access pattern (see :ref:dynamo0.3-kernel-valid-access: above) may be continuous in the vertical (and discontinuous in the horizontal), continuous in the horizontal (and discontinuous in the vertical), or continuous in both. In each case the code is the same. This principle of horizontal continuity also applies to the three generic `ANY_***` function space identifiers above. The valid metadata values for continuous and discontinuous function spaces are summarised in the table below.

Function Space Continuity	Function Space Name
Continuous	<code>W0</code> , <code>W1</code> , <code>W2</code> , <code>W2H</code> , <code>W2trace</code> , <code>W2Htrace</code> , <code>ANY_W2</code> , <code>ANY_SPACE_n</code>
Discontinuous	<code>W2broken</code> , <code>W2V</code> , <code>W2Vtrace</code> , <code>W3</code> , <code>Wtheta</code> , <code>ANY_DISCONTINUOUS_SPACE_n</code>

Horizontally discontinuous function spaces and fields over them will not need colouring so PSyclone does not perform it. If such attempt is made, PSyclone will raise a `Generation Error` in the **Dynamo0p3ColourTrans** transformation (see [Transformations](#) for more details on transformations). An example of fields iterating over a discontinuous function space `Wtheta` is given in `examples/lfric/eg9`, with the `GH_READWRITE` access descriptor denoting an update to the relevant fields. This example also demonstrates how to only colour loops over continuous function spaces when transformations are applied.

Read-Only Function Spaces

LFRic supports the concept of a **read-only function space**. A field on such a function space must not be modified by any kernels contained within `invoke` calls (i.e. within any code that PSyclone is responsible for). Further, a field on a read-only function space must contain clean halos in order to avoid any halo exchanges that would occur if the field is read within a kernel where redundant computation is performed.

The primary reason for including a read-only function space is that it does not need any halo-exchange support e.g. it does not require a routing table, which can reduce the memory footprint.

Currently `Wchi` is the only read-only function space in LFRic.

As a read-only function space is not modified, it does not matter whether it is classified as continuous or discontinuous. LFRic therefore treats read-only as a third category of function space.

Optional Field Metadata

A field entry in the `meta_args` array may have an optional fifth element. This element describes either a stencil access or, for inter-grid kernels, which mesh the field is on. Since an inter-grid kernel is not permitted to have stencil accesses, these two options are mutually exclusive. The metadata for each case is described in the following sections.

Stencil Metadata

Stencil metadata specifies that the corresponding field argument is accessed as a stencil operation within the Kernel. Stencil metadata only makes sense if the associated field is read within a Kernel i.e. it only makes sense to specify stencil metadata if the first entry is `GH_FIELD` and the second entry is `GH_READ`.

Stencil metadata is written in the following format:

```
STENCIL (type)
```

where `type` may be one of `X1D`, `Y1D`, `XORY1D`, `CROSS` or `CROSS2D`. As the stencil `extent` (the maximum distance from the central cell that the stencil extends) is not provided in the metadata, it is expected to be provided by the algorithm writer as part of the `invoke` call (see Section [Stencils](#)). As there is currently no way to specify a fixed extent value for stencils in the Kernel metadata, Kernels must therefore be written to support different values of extent (i.e. stencils with a variable number of cells).

The `XORY1D` stencil type indicates that the Kernel can accept either `X1D` or `Y1D` stencils. In this case it is up to the algorithm developer to specify which of these it is from the algorithm layer as part of the `invoke` call (see Section [Stencils](#)).

For example, the following stencil (with `extent=2`):

```
| 3 | 2 | 1 | 4 | 5 |
```

would be declared as:

```
STENCIL (X1D)
```

and the following stencil (with `extent=2`):

```
|   |   | 9 |   |   |
|   |   | 8 |   |   |
| 3 | 2 | 1 | 6 | 7 |
|   |   | 4 |   |   |
|   |   | 5 |   |   |
```

would be declared as:

```
STENCIL(CROSS)
```

Below is an example of stencil information within the full kernel metadata:

```
type(arg_type) :: meta_args(3) = (/  
    arg_type(GH_FIELD,      GH_REAL, GH_INC,  W1),      &  
    arg_type(GH_FIELD,      GH_REAL, GH_READ, W2H, STENCIL(CROSS)), &  
    arg_type(GH_OPERATOR, GH_REAL, GH_READ, W1, W2H)    &  
    /)
```

There is a full example of this distributed with PSyclone. It may be found in `examples/lfric/eg5`.

Inter-Grid Metadata

The alternative form of the optional fifth metadata argument for a field specifies which mesh the associated field is on. This is required for inter-grid kernels which perform prolongation or restriction operations on fields (or field vectors) existing on grids of different resolutions.

Mesh metadata is written in the following format:

```
mesh_arg=type
```

where `type` may be one of `GH_COARSE` or `GH_FINE`. Any kernel having a field argument with this metadata is assumed to be an inter-grid kernel and, as such, all of its other arguments (which must also be fields) must have it specified too. An example of the metadata for such a kernel is given below:

```
type(arg_type) :: meta_args(2) = (/  
    arg_type(GH_FIELD, GH_REAL, GH_READWRITE, ANY_DISCONTINUOUS_SPACE_1, &  
                                                    mesh_arg=GH_COARSE), &  
    arg_type(GH_FIELD, GH_REAL, GH_READ,      ANY_DISCONTINUOUS_SPACE_2, &  
                                                    mesh_arg=GH_FINE ) &  
    /)
```

Note that an inter-grid kernel must have at least one field (or field- vector) argument on each mesh type. Fields that are on different meshes cannot be on the same function space while those on the same mesh must also be on the same function space.

Column-wise Operators (CMA)

In this section we provide example metadata for each of the three recognised kernel types involving CMA operators.

Column-wise operators are constructed from cell-wise (local) operators. Therefore, in order to **assemble** a CMA operator, a kernel must have at least one read-only LMA operator, e.g.:

```
type(arg_type) :: meta_args(2) = (/  
    arg_type(GH_OPERATOR,      GH_REAL, GH_READ,  ANY_SPACE_1, ANY_SPACE_2), &  
    arg_type(GH_COLUMNWISE_OPERATOR, GH_REAL, GH_WRITE, ANY_SPACE_1, ANY_SPACE_2) &  
    /)
```

CMA operators (and their inverse) are **applied** to fields. Therefore any kernel of this type must have one read-only CMA operator, one read-only field and a field that is updated, e.g.:

```
type(arg_type) :: meta_args(3) = (/                                &
    arg_type(GH_FIELD,          GH_REAL, GH_INC,  ANY_SPACE_1),    &
    arg_type(GH_FIELD,          GH_REAL, GH_READ, ANY_SPACE_2),    &
    arg_type(GH_COLUMNWISE_OPERATOR, GH_REAL, GH_READ, ANY_SPACE_1, ANY_SPACE_2) &
    /)
```

Matrix-matrix kernels compute the product/linear combination of CMA operators. They must therefore have one such operator that is updated while the rest are read-only. They may also have read-only scalar arguments, e.g.:

```
type(arg_type) :: meta_args(3) = (/                                &
    arg_type(GH_COLUMNWISE_OPERATOR, GH_REAL, GH_WRITE, ANY_SPACE_1, ANY_SPACE_2), &
    arg_type(GH_COLUMNWISE_OPERATOR, GH_REAL, GH_READ, ANY_SPACE_1, ANY_SPACE_2), &
    arg_type(GH_COLUMNWISE_OPERATOR, GH_REAL, GH_READ, ANY_SPACE_1, ANY_SPACE_2), &
    arg_type(GH_SCALAR,            GH_REAL, GH_READ) /)
```

Note: The order with which arguments are specified in metadata for CMA kernels does not affect the process of identifying the type of kernel (whether it is assembly, matrix-matrix etc.)

meta_funcs

The (optional) second component of kernel metadata specifies whether any quadrature or evaluator data is required for a given function space. (If no quadrature or evaluator data is required then this metadata should be omitted.) Consider the following kernel metadata:

```
type, extends(kernel_type) :: testkern_operator_type
  type(arg_type), dimension(3) :: meta_args =                &
    (/ arg_type(gh_operator, gh_real,    gh_write, w0, w0), &
       arg_type(gh_field*3,  gh_real,    gh_read,  w1),    &
       arg_type(gh_scalar,   gh_integer, gh_read)          &
    /)
  type(func_type) :: meta_funcs(2) =                          &
    (/ func_type(w0, gh_basis, gh_diff_basis)                &
       func_type(w1, gh_basis)                                &
    /)
  integer :: gh_shape = gh_quadrature_XYoZ
  integer :: operates_on = cell_column
contains
  procedure, nopass :: code => testkern_operator_code
end type testkern_operator_type
```

The `arg_type` component of this metadata describes a kernel that takes three arguments (an operator, a field and an integer scalar). Following the `meta_args` array we now have a `meta_funcs` array. This allows the user to specify that the kernel requires basis functions (`gh_basis`) and/or the differential of the basis functions (`gh_diff_basis`) on one or more of the function spaces associated with the arguments listed in `meta_args`. In this case we require both for the `W0` function space but only basis functions for `W1`.

Note: Basis and differential basis functions for both `real`- and `integer`-valued field arguments have `real` values on the points on which these functions are *required*.

meta_reference_element

A kernel that requires properties of the reference element in LFRic specifies those properties through the `meta_reference_element` metadata entry. (If no reference element properties are required then this metadata should be omitted.) Consider the following example kernel metadata:

```
type, extends(kernel_type) :: testkern_type
  type(arg_type), dimension(2) :: meta_args =      &
    (/ arg_type(gh_field, gh_real, gh_read, w1), &
       arg_type(gh_field, gh_real, gh_inc,  w0) /)
  type(reference_element_data_type), dimension(2) ::      &
    meta_reference_element =                          &
    (/ reference_element_data_type(normals_to_horizontal_faces), &
       reference_element_data_type(normals_to_vertical_faces) /)
contains
  procedure, nopass :: code => testkern_code
end type testkern_type
```

This metadata specifies that the `testkern_type` kernel requires two properties of the reference element. The supported properties are listed below:

Name	Description
<code>normals_to_horizontal_faces</code>	Array of normals pointing in the positive (x, y, z) axis direction for each horizontal face indexed as (component, face).
<code>normals_to_vertical_faces</code>	Array of normals pointing in the positive (x, y, z) axis direction for each vertical face indexed as (component, face).
<code>normals_to_faces</code>	Array of normals pointing in the positive (x, y, z) axis direction for each face indexed as (component, face).
<code>outward_normals_to_horizontal_faces</code>	Array of outward-pointing normals for each horizontal face indexed as (component, face).
<code>outward_normals_to_vertical_faces</code>	Array of outward-pointing normals for each vertical face indexed as (component, face).
<code>outward_normals_to_faces</code>	Array of outward-pointing normals for each face indexed as (component, face).

meta_mesh

A kernel that requires properties of the LFRic mesh object specifies those properties through the `meta_mesh` metadata entry. (If no mesh properties are required then this metadata should be omitted.) Consider the following example kernel metadata:

```
type, extends(kernel_type) :: testkern_type
  type(arg_type), dimension(2) :: meta_args =      &
    (/ arg_type(gh_field, gh_real, gh_read, w1), &
       arg_type(gh_field, gh_real, gh_inc,  w0) /)
  type(mesh_data_type), dimension(1) ::      &
    meta_mesh =                                    &
    (/ mesh_data_type(adjacent_face) /)
contains
  procedure, nopass :: code => testkern_code
end type testkern_type
```

This metadata specifies that the `testkern_type` kernel requires one property of the mesh. There is currently one supported property:

Name	Description
adjacent_face	Local ID of a neighbouring face in each horizontally-adjacent cell indexed as (face).

gh_shape and gh_evaluator_targets

If a kernel requires basis or differential-basis functions then the metadata must also specify the set of points on which these functions are required. This information is provided by the `gh_shape` component of the metadata. Currently PSyclone supports four shapes; `gh_quadrature_XYoZ` for Gaussian quadrature points, `gh_quadrature_face` for quadrature points on cell faces, `gh_quadrature_edge` for quadrature points on cell edges and `gh_evaluator` for evaluation at nodal points. If a kernel requires just one of these then `gh_shape` is a scalar integer. However, if more than one is required then `gh_shape` becomes a one-dimensional, integer array, e.g.:

```
integer :: gh_shape(2) = (/ gh_quadrature_face, gh_quadrature_edge /)
```

If a kernel requires an evaluator then there are two options: if an evaluator is required for multiple function spaces then these can be specified using the additional `gh_evaluator_targets` metadata entry. This entry is a one-dimensional, integer array containing the desired function spaces. For example, to request basis/differential-basis functions evaluated on both W0 and W1, the metadata would be:

```
integer :: gh_shape = gh_evaluator
integer :: gh_evaluator_targets(2) = (/W0, W1/)
```

The kernel must have an argument (field or operator) on each of the function spaces listed in `gh_evaluator_targets`. The default behaviour if `gh_evaluator_targets` is not specified is to provide evaluators for each function space associated with the quantities that the kernel is updating. All necessary data is extracted in the PSy layer and passed to the kernel(s) as required - nothing is required from the Algorithm layer. If a kernel requires quadrature on the other hand, the Algorithm writer must supply a `quadrature_type` object for each specified quadrature as the last argument(s) to the kernel (see Section [Quadrature](#)).

Note that it is an error for kernel metadata to specify a value for `gh_shape` if no basis or differential-basis functions are required. It is also an error to specify `gh_evaluator_targets` if the kernel does not require an evaluator (i.e. `gh_shape != gh_evaluator`).

operates_on

The fourth type of metadata provided is `OPERATES_ON`. This specifies that the Kernel has been written with the assumption that it is supplied with the specified data for each field/operator argument. For user-supplied kernels this is currently only permitted to be `CELL_COLUMN` or `DOMAIN`. The possible values for `OPERATES_ON` and their interpretation are summarised in the following table:

operates_on	Data passed for each field/operator argument
cell_column	Single column of cells
dof	Single DoF (currently <i>Built-ins</i> only)
domain	All columns of cells

procedure

The fifth and final type of metadata is `procedure` metadata. This specifies the name of the Kernel subroutine that this metadata describes.

For example:

```
procedure, nopass :: my_kernel_subroutine
```

11.3.7 Subroutine

Rules for General-Purpose Kernels

The arguments to general-purpose kernels (those that do not involve either CMA operators or prolongation/restriction operations) that operate on cell-columns follow a set of rules which have been specified for the LFRic API. These rules are encoded in the `generate()` method within the `ArgOrdering` abstract class in the `dynamo0p3.py` file. The rules, along with PSyclone's naming conventions, are:

- 1) If an LMA operator is passed then include the `cells` argument. `cells` is an integer and has intent `in`.
- 2) Include `nlayers`, the number of layers in a column. `nlayers` is an integer of type `i_def` and has intent `in`.
- 3) For each scalar/field/vector_field/operator in the order specified by the `meta_args` metadata:
 - 1) If the current entry is a scalar quantity then include the Fortran variable in the argument list. The intent is determined from the metadata (see [meta_args](#) for an explanation).
 - 2) If the current entry is a field then include the field array. The field array name is currently specified as being `"field_"<argument_position>"_"<field_function_space>`. A field array is a real array of type `r_def` and dimensioned as the unique degrees of freedom for the space that the field is on. This value is passed in separately. Again, the intent is determined from the metadata (see [meta_args](#)).
 - 1) If the field entry has a stencil access then add an integer (or if the stencil is of type `CROSS2D`, an integer array of dimension(4)) stencil-size argument with intent `in`. This will supply the number of cells in the stencil or, in the case of the `CROSS2D` stencil, the number of cells in each branch of the stencil.
 - 2) If the stencil is of type `CROSS2D` then an integer of intent `in` for the max branch length is needed. This is used in defining the dimensions of the stencil dofmap array and is required due to the varying length of the branches of the stencil when used on planar meshes.
 - 3) Also needed is a stencil dofmap array of type integer and intent `in` in either 2 or 3 dimensions. For a `CROSS2D` stencil the array needs dimensions of (number-of-dofs-in-cell, max-branch-length, 4). All other stencils need dimensions of (number-of-dofs-in-cell, stencil-size).
 - 4) If the field entry stencil access is of type `XORY1D` then add an additional integer direction argument with intent `in`.
 - 3) If the current entry is a field vector then for each dimension of the vector, include a field array. The field array name is specified as being using `"field_"<argument_position>"_"<field_function_space>"_v"<vector_position>`. A field array in a field vector is declared in the same way as a field array (described in the previous step).
 - 4) If the current entry is an operator then first include a dimension size. This is an integer of type `i_def`. The name of this size is `<operator_name>"_ncell_3d"`. Next include the operator. This is a real array of type `r_def` and is 3 dimensional. The first two dimensions are the local degrees of freedom for the `to` and `from` function spaces respectively. The third dimension is the dimension size mentioned before. The name of the operator is `"op_"<argument_position>`. Again the intent is determined from the metadata (see [meta_args](#)).

- 4) For each function space in the order they appear in the metadata arguments (the `to` function space of an operator is considered to be before the `from` function space of the same operator as it appears first in lexicographic order)
 - 1) Include the number of local degrees of freedom (i.e. number per-cell) for the function space. This is an integer of type `i_def` and has intent `in`. The name of this argument is `"ndf_"<field_function_space>`.
 - 2) If there is a field on this space
 - 1) Include the unique number of degrees of freedom for the function space. This is an integer of type `i_def` and has intent `in`. The name of this argument is `"undf_"<field_function_space>`.
 - 2) Include the **dofmap** for this function space. This is an integer array of type `i_def` with intent `in`. It has one dimension sized by the local degrees of freedom for the function space.
 - 3) For each operation on the function space (`basis`, `diff_basis`, `orientation`) in the order specified in the metadata
 - 1) If it is a basis or differential basis function then we must pass real arrays of kind `r_def` with intent `in`. For each shape specified in the `gh_shape` metadata entry:
 - 1) If shape is `gh_quadrature_*` then the arrays are of rank four and are named `"basis_"<field_function_space>_<quadrature_arg_name>` or `"diff_basis_"<field_function_space>_<quadrature_arg_name>`, as appropriate:
 - 1) If shape is `gh_quadrature_xyoz` then the arrays have extent (dimension, number_of_dofs, np_xy, np_z).
 - 2) If shape is `gh_quadrature_face` or `gh_quadrature_edge` then the arrays have extent (dimension, number_of_dofs, np_xyz, nfaces or nedges).
 - 2) If shape is `gh_evaluator` then we pass one array for each target function space (i.e. as specified by `gh_evaluator_targets`). Each of these arrays are of rank three with extent (dimension, number_of_dofs, ndf_<target_function_space>). The name of the argument is `"basis_"<field_function_space>_"on_"<target_function_space>` or `"diff_basis_"<field_function_space>_"on_"<target_function_space>`, as appropriate.

where `<quadrature_arg_name>` is the name of the corresponding quadrature object being passed to the `Invoke`. `dimension` is 1 or 3 and depends upon the function space (see [Supported Function Spaces](#) above for more information) and whether or not it is a basis or a differential basis function (see the table below). `number_of_dofs` is the number of degrees of freedom (DoFs) associated with the function space and `np_*` are the number of points to be evaluated: i) `*_xyz` in all directions (3D); ii) `*_xy` in the horizontal plane (2D); iii) `*_x`, `*_y` in the horizontal (1D); and iv) `*_z` in the vertical (1D). `nfaces` and `nedges` are the number of horizontal faces/edges obtained from the appropriate quadrature object supplied to the `Invoke`.

Function Type	Dimension	Function Space Name
Basis	1	W0, W2trace, W2Htrace, W2Vtrace, W3, Wtheta, Wchi
	3	W1, W2, W2H, W2V, W2broken, ANY_W2
Differential Basis	1	W2, W2H, W2V, W2broken, ANY_W2
	3	W0, W1, W2trace, W2Htrace, W2Vtrace, W3, Wtheta, Wchi

- 2) If it is an orientation array, include the associated argument. The argument is an integer array of type `i_def` with intent `in`. There is one dimension of size the local degrees of freedom for the function space. The name of the array is `"orientation_"<field_function_space>`.

- 5) If either the `normals_to_horizontal_faces` or `outward_normals_to_horizontal_faces` properties of the reference element are required then pass the number of horizontal faces of the reference element (`nfaces_re_h`). Similarly, if either the `normals_to_vertical_faces` or `outward_normals_to_vertical_faces` are required then pass the number of vertical faces (`nfaces_re_v`). This also holds for the `normals_to_faces` and `outward_normals_to_faces` where the number of all faces of the reference element (`nfaces_re`) is passed to the kernel. (All of these quantities are integers of kind `i_def`.) Then, in the order specified in the `meta_reference_element` metadata:
- 1) For the `normals_to_horizontal/vertical_faces`, pass a rank-2 integer array of type `i_def` with dimensions `(3, nfaces_re_h/v)`.
 - 2) For the `outward_normals_to_horizontal/vertical_faces`, pass a rank-2 integer array of type `i_def` with dimensions `(3, nfaces_re_h/v)`.
 - 3) For `normals_to_faces` or `outward_normals_to_faces` pass a rank-2 integer array of type `i_def` with dimensions `(3, nfaces_re)`.
- 6) If the `adjacent_face` mesh property is required then:
- 1) If the number of horizontal cell faces obtained from the reference element (`nfaces_re_h`) is not already being passed to the kernel (due to rule 5 above) then supply it here. This is an integer of kind `i_def`.
 - 2) Pass a rank-1, integer array of kind `i_def` and extent `nfaces_re_h`.
- 7) If Quadrature is required (`gh_shape = gh_quadrature_*`) then, for each shape in the order specified in the `gh_shape` metadata:
- 1) Include integer, scalar arguments of kind `i_def` with intent `in` that specify the extent of the basis/diff-basis arrays:
 - 1) If `gh_shape` is `gh_quadrature_XYoZ` then pass `np_xy_<quadrature_arg_name>` and `np_z_<quadrature_arg_name>`.
 - 2) If `gh_shape` is `gh_quadrature_face/edge` then pass `nfaces/nedges_<quadrature_arg_name>` and `np_xyz_<quadrature_arg_name>`.
 - 2) Include weights which are real arrays of kind `r_def`:
 - 1) If `gh_quadrature_XYoZ` pass in `weights_xz_<quadrature_arg_name>` (rank one, extent `np_xy_<quadrature_arg_name>`) and `weights_z_<quadrature_arg_name>` (rank one, extent `np_z_<quadrature_arg_name>`).
 - 2) If `gh_quadrature_face/edge` pass in `weights_xyz_<quadrature_arg_name>` (rank two with extents `[np_xyz_<quadrature_arg_name>, nfaces/nedges_<quadrature_arg_name>]`).

Examples

For instance, if a kernel has only one written argument and requires an evaluator then its metadata might be:

```

type, extends(kernel_type) :: testkern_operator_type
  type(arg_type), dimension(2) :: meta_args =          &
    (/ arg_type(gh_operator, gh_real, gh_write, w0, w1), &
      arg_type(gh_field*3, gh_real, gh_read, w0) /)
  type(func_type) :: meta_funcs(1) =                  &
    (/ func_type(w0, gh_basis) /)
  integer :: operates_on = cell_column
  integer :: gh_shape = gh_evaluator
contains

```

(continues on next page)

(continued from previous page)

```

procedure, nopass :: code => testkern_operator_code
end type testkern_operator_type

```

then we only pass the basis functions evaluated on W0 (the space of the written kernel argument). The subroutine arguments will therefore be:

```

subroutine testkern_operator_code(cell, nlayers, ncell_3d,      &
    local_stencil, xdata, ydata, zdata, ndf_w0, undf_w0, map_w0, &
    basis_w0_on_w0, ndf_w1)

```

where `local_stencil` is the operator, `xdata`, `ydata` etc. are the three components of the field vector and `map_w0` is the dofmap for the W0 function space.

If instead, `gh_evaluator_targets` is specified in the metadata:

```

type, extends(kernel_type) :: testkern_operator_type
  type(arg_type), dimension(2) :: meta_args =      &
    (/ arg_type(gh_operator, gh_real, gh_write, w0, w1), &
    arg_type(gh_field*3, gh_real, gh_read, w0) /)
  type(func_type) :: meta_funcs(1) =      &
    (/ func_type(w0, gh_basis) /)
  integer :: operates_on = cell_column
  integer :: gh_shape = gh_evaluator
  integer :: gh_evaluator_targets(2) = (/W0, W1/)
contains
  procedure, nopass :: code => testkern_operator_code
end type testkern_operator_type

```

then we will need to pass two sets of basis functions (evaluated at W0 and at W1):

```

subroutine testkern_operator_code(cell, nlayers, ncell_3d,      &
    local_stencil, xdata, ydata, zdata, ndf_w0, undf_w0, map_w0, &
    basis_w0_on_w0, basis_w0_on_w1, ndf_w1)

```

If the metadata specifies that a kernel requires both an evaluator and quadrature:

```

type, extends(kernel_type) :: testkern_operator_type
  type(arg_type), dimension(2) :: meta_args =      &
    (/ arg_type(gh_operator, gh_real, gh_write, w0, w1), &
    arg_type(gh_field*3, gh_real, gh_read, w0) /)
  type(func_type) :: meta_funcs(1) =      &
    (/ func_type(w0, gh_basis) /)
  integer :: operates_on = cell_column
  integer :: gh_shape(2) = (/ gh_evaluator, gh_quadrature_face /)
contains
  procedure, nopass :: code => testkern_operator_code
end type testkern_operator_type

```

then we will need to pass basis functions for both the evaluator and the quadrature (where `qr_face` is the name of the face-quadrature object passed to the `Invoke`):

```

subroutine testkern_operator_code(cell, nlayers, ncell_3d,      &
    local_stencil, xdata, ydata, zdata, ndf_w0, undf_w0, map_w0, &
    basis_w0_on_w0, basis_w0_qr_face, ndf_w1,      &
    np_xyz_qr_face, nfaces_qr_face, weights_xyz_qr_face)

```

If the metadata specifies that the kernel requires a property of the reference element:

```

type, extends(kernel_type) :: testkern_operator_type
  type(arg_type), dimension(2) :: meta_args =          &
    (/ arg_type(gh_operator, gh_real, gh_write, w0, w1), &
      arg_type(gh_field*3, gh_real, gh_read, w0) /)
  type(reference_element_data_type) :: meta_reference_element(1) = &
    (/ reference_element_data_type(normals_to_horizontal_faces) /)
  integer :: operates_on = cell_column
contains
  procedure, nopass :: code => testkern_operator_code
end type testkern_operator_type

```

then the kernel must be passed the number of faces of the reference element and the array of face normals in the specified direction (here horizontal):

```

subroutine testkern_operator_code(cell, nlayers, ncell_3d,          &
  local_stencil, xdata, ydata, zdata, ndf_w0, undf_w0, map_w0, &
  nfaces_re_h, normals_face_h)

```

Rules for CMA Kernels

Kernels involving CMA operators are restricted to just three types; assembly, application/inverse-application and matrix-matrix. We give the rules for each of these in the sections below.

Assembly

An assembly kernel requires the column-banded dofmap for both the to- and from-function spaces of the CMA operator being assembled as well as the number of DoFs for each of the dofmaps. The full set of rules is:

- 1) Include the `cell` argument. `cell` is an integer and has intent `in`.
- 2) Include `nlayers`, the number of layers in a column. `nlayers` is an integer of type `i_def` and has intent `in`.
- 3) Include the number of cells in the 2D mesh, `ncell_2d`, which is an integer of type `i_def` with intent `in`.
- 4) Include the total number of cells, `ncell_3d`, which is an integer of type `i_def` with intent `in`.
- 5) For each argument in the `meta_args` metadata array:
 - 1) If it is a LMA operator, include a real, 3-dimensional array of type `r_def`. The first two dimensions are the local degrees of freedom for the `to` and `from` spaces, respectively. The third dimension is `ncell_3d`;
 - 2) If it is a CMA operator, include a real, 3-dimensional array of type `r_def`. The first dimension is `"bandwidth_"<operator_name>`, the second is `"nrow_"<operator_name>`, and the third is `ncell_2d`.
 - 1) Include the number of rows in the banded matrix. This is an integer of type `i_def` with intent `in` and is named as `"nrow_"<operator_name>`.
 - 2) If the from-space of the operator is *not* the same as the to-space then include the number of columns in the banded matrix. This is an integer of type `i_def` with intent `in` and is named as `"ncol_"<operator_name>`.
 - 3) Include the bandwidth of the banded matrix. This is an integer of type `i_def` with intent `in` and is named as `"bandwidth_"<operator_name>`.
 - 4) Include banded-matrix parameter `alpha`. This is an integer of type `i_def` with intent `in` and is named as `"alpha_"<operator_name>`.

- 5) Include banded-matrix parameter `beta`. This is an integer of type `i_def` with intent `in` and is named as `"beta_"<operator_name>`.
- 6) Include banded-matrix parameter `gamma_m`. This is an integer of type `i_def` with intent `in` and is named as `"gamma_m_"<operator_name>`.
- 7) Include banded-matrix parameter `gamma_p`. This is an integer of type `i_def` with intent `in` and is named as `"gamma_p_"<operator_name>`.
- 3) If it is a field or scalar argument then include arguments following the same rules as for general-purpose kernels.
- 6) For each unique function space in the order they appear in the metadata arguments (the `to` function space of an operator is considered to be before the `from` function space of the same operator as it appears first in lexicographic order):
 - 1) Include the number of degrees of freedom per cell for the space. This is an integer of type `i_def` with intent `in`. The name of this argument is `"ndf_"<arg_function_space>`.
 - 2) If there is a field on this space then:
 - 1) Include the unique number of degrees of freedom for the function space. This is an integer of type `i_def` and has intent `in`. The name of this argument is `"undf_"<field_function_space>`.
 - 2) Include the dofmap for this space. This is an integer array of type `i_def` with intent `in`. It has one dimension sized by the local degrees of freedom for the function space.
 - 3) If the CMA operator has this space as its `to/from` space then include the column-banded dofmap, the list of offsets for the `to/from`-space. This is an integer array of rank 2 and type `i_def`. The first dimension is `"ndf_"<arg_function_space>` and the second is `nlayers`.

Application/Inverse-Application

A kernel applying a CMA operator requires the column-indirection dofmap for both the `to`- and `from`-function spaces of the CMA operator. Since it does not have any LMA operator arguments it does not require the `ncell_3d` and `nlayers` scalar arguments. (Since a column-wise operator is, by definition, assembled for a whole column, there is no loop over levels when applying it.) The full set of rules is then:

- 1) Include the `cell` argument. `cell` is an integer and has intent `in`.
- 2) Include the number of cells in the 2D mesh, `ncell_2d`, which is an integer of type `i_def` with intent `in`.
- 3) For each argument in the `meta_args` metadata array:
 - 1) If it is a field, include the field array. This is a real array of kind `r_def` and is of rank 1. The field array name is currently specified as being `"field_"<argument_position>"_"<field_function_space>`. The extent of the array is the number of unique degrees of freedom for the function space that the field is on. This value is passed in separately. The intent of the argument is determined from the metadata (see [meta_args](#));
 - 2) If it is a CMA operator, include it and its associated parameters (see Rule 5 of CMA Assembly kernels).
- 4) For each of the unique function spaces encountered in the metadata arguments (the `to` function space of an operator is considered to be before the `from` function space of the same operator as it appears first in lexicographic order):
 - 1) Include the number of degrees of freedom per cell for the associated function space. This is an integer of type `i_def` with intent `in`. The name of this argument is `"ndf_"<field_function_space>`;
 - 2) Include the number of unique degrees of freedom for the associated function space. This is an integer of type `i_def` with intent `in`. The name of this argument is `"undf_"<field_function_space>`;

- 3) Include the dofmap for this function space. This is a rank-1 integer array of type `i_def` with extent equal to the number of degrees of freedom of the space ("`ndf_`"<`field_function_space`>).
- 5) Include the indirection map for the to-space of the CMA operator. This is a rank-1 integer array of type `i_def` with extent `nrow`.
- 6) If the from-space of the operator is *not* the same as the to-space then include the indirection map for the from-space of the CMA operator. This is a rank-1 integer array of type `i_def` with extent `ncol`.

Matrix-Matrix

Does not require any dofmaps and also does not require the `nlayers` and `ncell_3d` scalar arguments. The full set of rules are then:

- 1) Include the `cell` argument. `cell` is an integer and has intent `in`.
- 2) Include the number of cells in the 2D mesh, `ncell_2d`, which is an integer of type `i_def` with intent `in`.
- 3) For each CMA operator or scalar argument specified in metadata:
 - 1) If it is a CMA operator, include it and its associated parameters (see Rule 5 of CMA Assembly kernels);
 - 2) If it is a scalar argument include the corresponding Fortran variable in the argument list with intent `in`.

Rules for Inter-Grid Kernels

As already specified, inter-grid kernels are only permitted to take fields and/or field-vectors as arguments. Fields (and field-vectors) that are on different meshes must be on different function spaces. Fields on the same mesh must also be on the same function space.

Argument ordering follows the general pattern used for ‘normal’ kernels with field data being followed by dofmap data. The rules for arguments to inter-grid kernels are as follows:

- 1) Include `nlayers`, the number of layers in a column. `nlayers` is an integer of type `i_def` and has intent `in`.
- 2) Include the `cell_map` for the current cell (column). This is an integer array of rank one, type `i_def` and intent `in` which provides the mapping from the coarse to the fine mesh. It has extent `ncell_f_per_c`.
- 3) Include `ncell_f_per_c`, the number of fine cells per coarse cell. This is an integer of type `i_def` and has intent `in`.
- 4) Include `ncell_f`, the number of cells (columns) in the fine mesh. This is an integer of type `i_def` and has intent `in`.
- 5) For each argument in the `meta_args` metadata array (which must be a field or field-vector):
 - 1) Pass in field data as done for a regular kernel.
- 6) For each unique function space (of which there will currently be two) in the order in which they are encountered in the `meta_args` metadata array, include dofmap information:

If the dofmap is associated with an argument on the fine mesh:

- 1) Include `ndf_fine`, the number of DoFs per cell for the FS of the field on the fine mesh;
- 2) Include `undf_fine`, the number of unique DoFs per cell for the FS of the field on the fine mesh;
- 3) Include `dofmap_fine`, the *whole* dofmap for the fine mesh. This is an integer array of rank two and type `i_def` with intent `in`. The extent of the first dimension is `ndf_fine` and that of the second is `ncell_f`.

else, the dofmap is associated with an argument on the coarse mesh:

- 1) Include `undf_coarse`, the number of unique DoFs for the coarse field. This is an integer of type `i_def` with intent `in`;
- 2) Include `dofmap_coarse`, the dofmap for the current cell (column) in the coarse mesh. This is an integer array of rank one, type `i_def` and has intent `in`.

Rules for Domain Kernels

The rules for kernels that have `operates_on = DOMAIN` are identical to those for general-purpose kernels (described [above](#)), allowing for the fact that they are not permitted any type of operator argument or any argument with a stencil access.

Argument Intents

As described [above](#), LFRic kernels read and/or update the data pointed to by objects such as *fields* or *operators*. This data is passed to the kernels as *subroutine arguments* and their Fortran intents usually follow the logic determined by their *access modes*.

- `GH_READ` indicates intent (`in`) as the argument is only ever read from.
- `GH_WRITE` (for discontinuous function spaces) indicates that the argument is only written to in a kernel. When the argument is defined inside a kernel its intent is `intent(out)` and when it is defined outside of a kernel its intent is `intent(inout)`.
- `GH_INC` and `GH_READWRITE` indicate `intent(inout)` as the arguments are updated (albeit in a different way due to different access to DoFs, see [meta_args](#) for more details).

11.4 Built-ins

The basic concept of a PSyclone Built-in is described in the [Built-ins](#) section. In the Dynamo0.3 API, calls to Built-ins generally follow a convention that the field/scalar written to comes first in the argument list. Dynamo0.3 Built-ins must conform to the following rules:

- 1) They must have one and only one modified (i.e. written to) argument.
- 2) They must operate on a DoF (`operates_on = DOF` metadata).
- 3) There must be at least one field in the argument list. This is so that we know the number of DoFs to iterate over in the PSy layer.
- 4) Kernel arguments must be either fields or scalars.
- 5) All field arguments to a given Built-in must be on the same function space. This is because all current Built-ins operate on DoFs and therefore all fields should have the same number. It also means that we can determine the number of DoFs uniquely when a scalar is written to.

The Built-ins supported for the Dynamo0.3 API are listed in the related subsections, grouped by the mathematical operation they perform. For clarity, the calculation performed by each Built-in is described using Fortran array syntax; this does not necessarily reflect the actual implementation of the Built-in (*e.g.* it could be implemented by PSyclone generating a call to an optimised maths library).

As described in the PSy-layer [Argument Intents](#) section, the Fortran intent of LFRic *field* objects is always `in`. The field or scalar that has its data modified by a Built-in is marked in **bold**.

11.4.1 Metadata

The code below outlines the elements of the LFRic API Built-in metadata, 1) ‘meta_args’, 2) ‘operates_on’ and 3) ‘procedure’:

```

type, public, extends(kernel_type) :: aX_plus_bY
  private
  type(arg_type) :: meta_args(5) = (
    arg_type(GH_FIELD, GH_REAL GH_WRITE, ANY_SPACE_1), &
    arg_type(GH_SCALAR, GH_REAL, GH_READ), &
    arg_type(GH_FIELD, GH_REAL, GH_READ, ANY_SPACE_1), &
    arg_type(GH_SCALAR, GH_REAL, GH_READ), &
    arg_type(GH_FIELD, GH_REAL GH_READ, ANY_SPACE_1) &
  )
  integer :: operates_on = DOF
contains
  procedure, nopass :: aX_plus_bY_code
end type aX_plus_bY

```

As can be seen, the metadata for a Built-in kernel is a subset of that for a *user-defined Kernel* with the exception that `operates_on` must be `DOF` instead of `CELL_COLUMN`.

Valid Access Modes

The allowed accesses for arguments in Built-in kernels are a bit different than for the *user-defined Kernels* and are listed in the table below.

Argument Type	Data Type	Function Space	Access Type
GH_SCALAR	GH_INTEGER	n/a	GH_READ
GH_SCALAR	GH_REAL	n/a	GH_READ, GH_SUM
GH_FIELD	GH_REAL	ANY_SPACE_n	GH_READ, GH_WRITE, GH_READWRITE

Note: Since the LFRic infrastructure does not currently support integer reductions, integer scalar arguments in Built-ins are restricted to having read-only access.

Note: Built-ins are currently implemented only for real-valued fields and will be extended to update integer-valued fields in issue #853.

11.4.2 Naming scheme

The supported Built-ins in the Dynamo0.3 API are named according to the scheme presented below. Any new Built-in needs to comply with these rules.

- 1) Ordering of arguments in Built-ins calls follows *LHS (result) <- RHS (operation on arguments)* direction, except where a Built-in returns the *LHS* result to one of the *RHS* arguments. In that case ordering of arguments remains as in the *RHS* expression, with the returning *RHS* argument written as close to the *LHS* as it can be without affecting the mathematical expression.
- 2) Field names begin with upper case in short form (e.g. **X**, **Y**, **Z**) and any case in long form (e.g. **Field1**, **field**).

- 3) Scalar names begin with lower case: e.g. **a**, **b**, are **scalar1**, **scalar2**. Special names for scalars are: **constant** (or **c**), **innprod** (inner/scalar product of two fields) and **sumfld** (sum of a field).
- 4) Arguments in Built-ins variable declarations and constructs (PSyclone Fortran and Python definitions):
 - 1) Are always written in long form and lower case (e.g. **field1**, **field2**, **scalar1**, **scalar2**);
 - 2) *LHS* result arguments are always listed first;
 - 3) *RHS* arguments are listed in order of appearance in the mathematical expression, except when one of them is the *LHS* result.
- 5) Built-ins names in Fortran consist of:
 - 1) *RHS* arguments in short form (e.g. **X**, **Y**, **a**, **b**) only;
 - 2) Descriptive name of mathematical operation on *RHS* arguments in the form `<operationname>_<RHSarg>` for one *RHS* argument or `<RHSargs>_<operationname>_<RHSargs>` for more;
 - 3) Prefix `"inc_"` where the result is returned to one of the *RHS* arguments (i.e. `"inc_"<RHSargs>_<operationname>_<RHSargs>`).
- 6) Built-ins names in Python definitions are similar to their Fortran counterparts, with a few differences:
 - 1) Operators and *RHS* arguments are all in upper case (e.g. **X**, **Y**, **A**, **B**, **Plus**, **Minus**);
 - 2) There are no underscores;
 - 3) Common prefix is `"Dyn"`, common suffix is `"Kern"`.

11.4.3 Addition

Built-ins which add (scaled) fields are denoted with the keyword **plus**.

X_plus_Y

X_plus_Y (*field3*, *field1*, *field2*)

Sums two fields ($Z = X + Y$):

```
field3(:) = field1(:) + field2(:)
```

where:

- `type(field_type), intent(in) :: field3, field1, field2`

inc_X_plus_Y

inc_X_plus_Y (*field1*, *field2*)

Adds the second field to the first and returns it ($X = X + Y$):

```
field1(:) = field1(:) + field2(:)
```

where:

- `type(field_type), intent(in) :: field1, field2`

aX_plus_Y

aX_plus_Y (*field3, scalar, field1, field2*)

Performs $Z = aX + Y$:

```
field3(:) = scalar*field1(:) + field2(:)
```

where:

- `real(r_def), intent(in) :: scalar`
- `type(field_type), intent(in) :: field3, field1, field2`

inc_aX_plus_Y

inc_aX_plus_Y (*scalar, field1, field2*)

Performs $X = aX + Y$ (increments the first field):

```
field1(:) = scalar*field1(:) + field2(:)
```

where:

- `real(r_def), intent(in) :: scalar`
- `type(field_type), intent(in) :: field1, field2`

inc_X_plus_bY

inc_X_plus_bY (*field1, scalar, field2*)

Performs $X = X + bY$ (increments the first field):

```
field1(:) = field1(:) + scalar*field2(:)
```

where:

- `real(r_def), intent(in) :: scalar`
- `type(field_type), intent(in) :: field1, field2`

aX_plus_bY

aX_plus_bY (*field3, scalar1, field1, scalar2, field2*)

Performs $Z = aX + bY$:

```
field3(:) = scalar1*field1(:) + scalar2*field2(:)
```

where:

- `real(r_def), intent(in) :: scalar1, scalar2`
- `type(field_type), intent(in) :: field3, field1, field2`

inc_aX_plus_bY

inc_aX_plus_bY (*scalar1*, *field1*, *scalar2*, *field2*)

Performs $X = aX + bY$ (increments the first field):

```
field1(:) = scalar1*field1(:) + scalar2*field2(:)
```

where:

- `real(r_def), intent(in) :: scalar1, scalar2`
- `type(field_type), intent(in) :: field1, field2`

11.4.4 Subtraction

Built-ins which subtract (scaled) fields are denoted with the keyword **minus**.

X_minus_Y

X_minus_Y (*field3*, *field1*, *field2*)

Subtracts the second field from the first and stores the result in the third ($Z = X - Y$):

```
field3(:) = field1(:) - field2(:)
```

where:

- `type(field_type), intent(in) :: field3, field1, field2`

inc_X_minus_Y

inc_X_minus_Y (*field1*, *field2*)

Subtracts the second field from the first and returns it ($X = X - Y$):

```
field1(:) = field1(:) - field2(:)
```

where:

- `type(field_type), intent(in) :: field1, field2`

aX_minus_Y

aX_minus_Y (*field3*, *scalar*, *field1*, *field2*)

Performs $Z = aX - Y$:

```
field3(:) = scalar*field1(:) - field2(:)
```

where:

- `real(r_def), intent(in) :: scalar`
- `type(field_type), intent(in) :: field3, field1, field2`

X_minus_bY

X_minus_bY (*field3, field1, scalar, field2*)

Performs $Z = X - bY$:

```
field3(:) = field1(:) - scalar*field2(:)
```

where:

- `real(r_def), intent(in) :: scalar`
- `type(field_type), intent(in) :: field3, field1, field2`

inc_X_minus_bY

inc_X_minus_bY (*field1, scalar, field2*)

Performs $X = X - bY$ (increments the first field):

```
field1(:) = field1(:) - scalar*field2(:)
```

where:

- `real(r_def), intent(in) :: scalar`
- `type(field_type), intent(in) :: field1, field2`

11.4.5 Multiplication

Built-ins which multiply (scaled) fields are denoted with the keyword **times**.

X_times_Y

X_times_Y (*field3, field1, field2*)

Multiplies two fields together and returns the result in a third field ($Z = X*Y$):

```
field3(:) = field1(:)*field2(:)
```

where:

- `type(field_type), intent(in) :: field3, field1, field2`

inc_X_times_Y

inc_X_times_Y (*field1, field2*)

Multiplies the first field by the second and returns it ($X = X*Y$):

```
field1(:) = field1(:)*field2(:)
```

where:

- `type(field_type), intent(in) :: field1, field2`

inc_aX_times_Y

inc_aX_times_Y (*scalar*, *field1*, *field2*)

Performs $X = a * X * Y$ (increments the first field):

```
field1(:) = scalar*field1(:)*field2(:)
```

where:

- `real(r_def), intent(in) :: scalar`
- `type(field_type), intent(in) :: field1, field2`

11.4.6 Scaling

Built-ins which scale fields are technically cases of multiplying a field by a scalar and are hence also denoted with the keyword **times**.

a_times_X

a_times_X (*field2*, *scalar*, *field1*)

Multiplies a field by a scalar and stores the result in a second field ($Y = a * X$):

```
field2(:) = scalar*field1(:)
```

where:

- `real(r_def), intent(in) :: scalar`
- `type(field_type), intent(in) :: field2, field1`

inc_a_times_X

inc_a_times_X (*scalar*, *field*)

Multiplies a field by a scalar value and returns the field ($X = a * X$):

```
field(:) = scalar*field(:)
```

where:

- `real(r_def), intent(in) :: scalar`
- `type(field_type), intent(in) :: field`

11.4.7 Division

Built-ins which divide (scaled) fields are denoted with the keyword **divideby**.

X_divideby_Y

X_divideby_Y (*field3, field1, field2*)

Divides the first field by the second and returns the result in the third ($Z = X/Y$):

```
field3(:) = field1(:)/field2(:)
```

where:

- `type(field_type), intent(out) :: field3, field1, field2`

inc_X_divideby_Y

inc_X_divideby_Y (*field1, field2*)

Divides the first field by the second and returns it ($X = X/Y$):

```
field1(:) = field1(:)/field2(:)
```

where:

- `type(field_type), intent(in) :: field1, field2`

11.4.8 Setting to value

Built-ins which set field elements to some value and hence are denoted with the keyword **setval**.

setval_c

setval_c (*field, constant*)

Sets all elements of the field *field* to the value *constant* ($X = c$):

```
field(:) = constant
```

where:

- `type(field_type), intent(in) :: field`
- `real(r_def), intent(in) :: constant`

Note: The field may be on any function space.

setval_X

setval_X (*field2*, *field1*)

Sets a field *field2* equal to field *field1* ($Y = X$):

```
field2(:) = field1(:)
```

where:

- `type(field_type), intent(in) :: field2, field1`

11.4.9 Raising to power

Built-ins which raise field elements to an exponent are denoted with the keyword **powreal** for real exponent or **powint** for integer exponent.

inc_X_powreal_a

inc_X_powreal_a (*field*, *rscalar*)

Raises a field to a real scalar value and returns the field ($X = X**a$):

```
field(:) = field(:)**rscalar
```

where:

- `type(field_type), intent(in) :: field`
- `real(r_def), intent(in) :: rscalar`

inc_X_powint_n

inc_X_powint_n (*field*, *iscalar*)

Raises a field to an integer scalar value and returns the field ($X = X**n$):

```
field(:) = field(:)**iscalar
```

where:

- `type(field_type), intent(in) :: field`
- `integer(i_def), intent(in) :: iscalar`

11.4.10 Inner product

Built-ins which calculate the inner product of two fields or of a field with itself are denoted with the keyword **inner-product**.

X_innerproduct_Y

X_innerproduct_Y (*innprod*, *field1*, *field2*)

Computes the inner product of the fields *field1* and *field2*, *i.e.*:

```
innprod = SUM(field1(:)*field2(:))
```

where:

- `real(r_def), intent(out) :: innprod`
- `type(field_type), intent(in) :: field1, field2`

Note: When used with distributed memory this Built-in will trigger the addition of a global sum which may affect the performance and/or scalability of the code.

X_innerproduct_X

X_innerproduct_X (*innprod*, *field*)

Computes the inner product of the field *field1* by itself, *i.e.*:

```
innprod = SUM(field(:)*field(:))
```

where:

- `real(r_def), intent(out) :: innprod`
- `type(field_type), intent(in) :: field`

Note: When used with distributed memory this Built-in will trigger the addition of a global sum which may affect the performance and/or scalability of the code.

11.4.11 Sum of elements

Built-in which sums the elements of a field is denoted with the keyword *sum*.

sum_X

sum_X (*sumfld*, *field*)

Sums all of the elements of the field *field* and returns the result in the scalar variable *sumfld*:

```
sumfld = SUM(field(:))
```

where:

- `real(r_def), intent(out) :: sumfld`
- `type(field_type), intent(in) :: field`

Note: When used with distributed memory this Built-in will trigger the addition of a global sum which may affect the performance and/or scalability of the code.

11.5 Boundary Conditions

In the Dynamo0.3 API, boundary conditions for a field or LMA operator can be enforced by the algorithm developer by calling the Kernels `enforce_bc_type` or `enforce_operator_bc_type`, respectively. These kernels take a field or operator as input and apply boundary conditions. For example:

```
call invoke( kernel_type(field1, field2),      &
             enforce_bc_type(field1),         &
             kernel_with_op_type(field1, op1), &
             enforce_operator_bc_type(op1)     &
             )
```

The particular boundary conditions that are applied are not known by PSyclone, PSyclone simply recognises these kernels by their names and passes pre-specified dofmap and boundary_value arrays into the kernel implementations, the contents of which are set by the LFRic infrastructure.

Up to and including version 1.4.0 of PSyclone, boundary conditions were applied automatically after a call to `matrix_vector_type` if the field arguments were on a vector function space (one of W1, W2, W2H, W2V or W2broken). With the subsequent introduction of the ability to apply boundary conditions to operators this functionality is no longer required and has been removed.

Example `eg4` in the `examples/lfric` directory includes a call to `enforce_bc_kernel_type` so can be used to see the boundary condition code that is added by PSyclone. See the `README` in the `examples/lfric` directory for instructions on how to run this example.

An example of applying boundary conditions to an operator is the kernel `enforce_operator_bc_kernel_mod.F90` in the `<PSYCLONEHOME>/src/psyclone/tests/test_files/dynamo0p3` directory. Since operators are discontinuous quantities, updating their values can be safely performed in parallel (see Section [Kernel](#)). The `GH_READWRITE` access is used for updating discontinuous operators (see subsection [Valid Access Modes](#) for more details).

11.6 Conventions

The naming of Dynamo0.3 API kernels and associated entities (types, subroutines and modules) follows the PSyclone Fortran naming conventions (see [Fortran Naming Conventions](#)). However, PSyclone does not need this convention to be followed apart from the stub generator (see the [Stub Generation](#) Section) where the name of the metadata to be parsed is determined from the module name.

The contents of the metadata is also usually declared private but this does not affect PSyclone.

Finally, the `procedure` metadata (located within the kernel metadata) usually has `nopass` specified but again this is ignored by PSyclone.

11.7 Configuration

11.7.1 Annexed DoFs

When a kernel operates on DoFs (rather than cell-columns) for a continuous field using distributed memory (see the *Distributed Memory* Section), then PSyclone need only ensure that DoFs owned by a processor are computed. However, for continuous fields, shared DoFs at the boundary between processors must be replicated (as different cells share the same DoF). Only one processor can own a DoF, therefore processors will have continuous fields which contain DoFs that the processor does not own. These unowned DoFs are called *annexed* in the Dynamo0.3 API and are a separate, but related, concept to field halos.

When a kernel that operates on a cell-column needs to read a continuous field then the annexed DoFs must be up-to-date on all processors. If they are not then a halo exchange must be added. Currently PSyclone defaults, for kernels which iterate over DoFs, to iterating over only owned DoFs. This behaviour can be changed by setting `COMPUTE_ANNEXED_DOFS` to `true` in the `dynamo0.3` section of the configuration file (see the *Configuration* section). PSyclone will then generate code to iterate over both owned and annexed DoFs, thereby reducing the number of halo exchanges required (at the expense of redundantly computing annexed DoFs). For more details please refer to the LFRic (Dynamo0.3) developers section.

11.7.2 Run-time Checks

PSyclone performs static consistency checks where possible. When this is not possible PSyclone can generate run-time checks. As there may be performance costs associated with run-time checks they may be switched on or off by the `RUN_TIME_CHECKS` option in the configuration file.

Currently run-time checks can be generated to:

- 1) Check that a field with a read-only function space (see section *Read-Only Function Spaces*) is not modified by a kernel. This is enforced by checking that all fields that are marked (in kernel metadata) as being updated by a kernel are not on a read-only function space. A second check that is required for fields on read-only function spaces is to ensure that the halo is clean before it is accessed. This check is currently implemented within the LFRic infrastructure halo exchange call (that the PSyclone LFRic API places at appropriate locations). If the halo is clean then the halo exchange will not be called. However, if the halo is not clean then the resulting halo exchange call will cause the infrastructure to raise an error (because the field is on a read-only space).
- 2) Check that the function space of a field is consistent with the kernel function space metadata that the field's data is passed into. For example, if kernel metadata specifies that a field is on the W2 function space then a run-time check is added to ensure that the field object passed into the PSy layer is indeed on that space. For more general kernel function space metadata, such as `ANY_DISCONTINUOUS_SPACE_*` then a run-time check is added to ensure that the field is on one of the discontinuous function spaces supported in the LFRic API.

11.8 Transformations

This section describes the Dynamo0.3-API-specific transformations. In cases, excepting `Dynamo0p3RedundantComputationTrans`, `Dynamo0p3AsyncHaloExchangeTrans` and `Dynamo0p3KernelConstTrans`, these transformations are specialisations of generic transformations described in the *Transformations* section. The difference between these transformations and the generic ones is that these perform Dynamo0.3-API-specific checks to make sure the transformations are valid. In practice these transformations perform the required checks then call the generic ones internally.

The use of the Dynamo0.3-API-specific transformations is exactly the same as the equivalent generic ones in all cases excepting `DynamoLoopFuseTrans`. In this case an additional optional argument `same_space` can be set after creating an instance of the transformation. The reason for this is to allow loop fusion when one or more of the

iteration spaces is determined by a function space that is unknown by PSyclone at compile time. This is the case when the `ANY_SPACE_n` function space is specified in the Kernel metadata. Setting `ftrans.same_space = True` allows the user to specify that the spaces are the same. This option should therefore be used with caution. PSyclone will raise an error if **same_space** is used when at least one of the function spaces is not `ANY_SPACE_n` or both spaces are not the same. As a general PSyclone will not allow loop fusion if it does not know the spaces are the same. The exception are loops over discontinuous spaces (see [Supported Function Spaces](#) for list of discontinuous function spaces) for which loop fusion is allowed (unless the loop bounds become different due to a prior transformation).

The **Dynamo0p3RedundantComputationTrans** and **Dynamo0p3AsyncHaloExchange** transformations are only valid for the Dynamo0.3 API. This is because this API is currently the only one that supports distributed memory. An example of redundant computation can be found in `examples/lfric/eg8` and an example of asynchronous halo exchanges can be found in `examples/lfric/eg11`.

The **Dynamo0p3KernelConstTrans** transformation is only valid for the Dynamo0.3 API. This is because the properties that it makes constant are API specific.

The Dynamo0.3-API-specific transformations currently available are given below. If the name of a transformation includes “Dynamo0p3” it means that the transformation is only valid for this particular API. If the name of the transformation includes “Dynamo” then it should work with all versions of the Dynamo API.

Note: Only the loop-colouring and OpenMP transformations are currently supported for loops that contain inter-grid kernels. Attempting to apply other transformation types will result in PSyclone raising an error.

```
class psyclone.domain.lfric.transformations.LFRicExtractTrans (node_class=<class
                                                                    'psy-
                                                                    clone.psyir.nodes.extract_node.ExtractNode>
```

Dynamo0.3 API application of ExtractTrans transformation to extract code into a stand-alone program. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>>
>>> API = "dynamo0.3"
>>> FILENAME = "solver_alg.x90"
>>> ast, invokeInfo = parse(FILENAME, api=API)
>>> psy = PSyFactory(API, distributed_memory=False).create(invoke_info)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>>
>>> from psyclone.domain.lfric.transformations import LFRicExtractTrans
>>> etrans = LFRicExtractTrans()
>>>
>>> # Apply LFRicExtractTrans transformation to selected Nodes
>>> newsched, _ = etrans.apply(schedule.children[0:3])
>>> newsched.view()
```

validate (*node_list*, *options=None*)

Perform Dynamo0.3 API specific validation checks before applying the transformation.

Parameters

- **node_list** (list of `psyclone.psyir.nodes.Node`) – the list of Node(s) we are checking.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises TransformationError – if transformation is applied to a Loop over cells in a colour without its parent Loop over colours.

class `psyclone.transformations.DynamoLoopFuseTrans` (*same_space=False*)

Dynamo0.3 API specialisation of the *base class* in order to fuse two Dynamo loops after performing validity checks. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>>
>>> API = "dynamo0.3"
>>> FILENAME = "alg.x90"
>>> ast, invokeInfo = parse(FILENAME, api=API)
>>> psy = PSyFactory(API, distributed_memory=False).create(invoke_info)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>>
>>> from psyclone.transformations import DynamoLoopFuseTrans
>>> ftrans = DynamoLoopFuseTrans()
>>>
>>> new_schedule, memento = ftrans.apply(schedule[0], schedule[1])
>>> new_schedule.view()
```

The optional argument *same_space* can be set as

```
>>> ftrans.same_space = True
```

after the instance of the transformation is created.

apply (*node1, node2, options=None*)

Fuses two `psyclone.dynamo0p3.DynLoop` loops after performing validity checks by calling `LoopFuseTrans.apply()` method of the base class.

Parameters

- **node1** (`psyclone.dynamo0p3.DynLoop`) – the first Loop to fuse.
- **node2** (`psyclone.dynamo0p3.DynLoop`) – the second Loop to fuse.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Returns two-tuple of the modified Schedule and a record of the transformation.

Return type (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

property same_space

Returns the *same_space* flag that is specified when applying this transformation. The default value is *False*. This optional flag, set to *True*, asserts that an unknown iteration space (i.e. *ANY_SPACE*) matches the other iteration space. This is set at the user's own risk. If both iteration spaces are discontinuous the loops can be fused without having to use the *same_space* flag.

validate (*node1, node2, options=None*)

Performs various checks to ensure that it is valid to apply the `DynamoLoopFuseTrans` transformation to the supplied loops.

Parameters

- **node1** (`psyclone.dynamo0p3.DynLoop`) – the first Loop to fuse.
- **node2** (`psyclone.dynamo0p3.DynLoop`) – the second Loop to fuse.

- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises

- **TransformationError** – if either of the supplied loops contains an inter-grid kernel.
- **TransformationError** – if one or both function spaces have invalid names.
- **TransformationError** – if the *same_space* flag was set, but does not apply because neither field is on *ANY_SPACE* or the spaces are not the same.
- **TransformationError** – if one or more of the iteration spaces is unknown (*ANY_SPACE*) and the *same_space* flag is not set to *True*.
- **TransformationError** – if the loops are over different spaces that are not both discontinuous and the loops both iterate over cells.
- **TransformationError** – if the loops' upper bound names are not the same.
- **TransformationError** – if the halo-depth indices of two loops are not the same.
- **TransformationError** – if each loop already contains a reduction.
- **TransformationError** – if the first loop has a reduction and the second loop reads the result of the reduction.

class `psyclone.transformations.DynamoOMPParallelLoopTrans` (*omp_schedule='static'*)
 Dynamo-specific OpenMP loop transformation. Adds Dynamo specific validity checks. Actual transformation is done by the *base class*.

apply (*node, options=None*)

Perform Dynamo specific loop validity checks then call the *apply()* method of the *base class*.

Parameters

- **node** (`psyclone.psyir.nodes.Node`) – the Node in the Schedule to check
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises **TransformationError** – if the associated loop requires colouring.

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

class `psyclone.transformations.Dynamo0p3AsyncHaloExchangeTrans`

Splits a synchronous halo exchange into a halo exchange start and halo exchange end. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "dynamo0.3"
>>> ast, invokeInfo = parse("file.f90", api=api)
>>> psy=PSyFactory(api).create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import Dynamo0p3AsyncHaloExchangeTrans
>>> trans = Dynamo0p3AsyncHaloExchangeTrans()
```

(continues on next page)

(continued from previous page)

```
>>> new_schedule, memento = trans.apply(schedule.children[0])
>>> new_schedule.view()
```

apply (*node*, *options*=None)

Transforms a synchronous halo exchange, represented by a HaloExchange node, into an asynchronous halo exchange, represented by HaloExchangeStart and HaloExchangeEnd nodes.

Parameters

- **node** (`psyclone.psygen.HaloExchange`) – a synchronous haloexchange node.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Returns tuple of the modified schedule and a record of the transformation.

Return type (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

Return type `str`

validate (*node*, *options*)

Internal method to check whether the node is valid for this transformation.

Parameters

- **node** (`psyclone.psygen.HaloExchange`) – a synchronous Halo Exchange node
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises TransformationError – if the node argument is not a HaloExchange (or subclass thereof)

class `psyclone.transformations.Dynamo0p3ColourTrans`

Split a Dynamo 0.3 loop over cells into colours so that it can be parallelised. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> import transformations
>>> import os
>>> import pytest
>>>
>>> TEST_API = "dynamo0.3"
>>> _, info=parse(os.path.join(os.path.dirname(os.path.abspath(__file__)),
>>>                             "tests", "test_files", "dynamo0p3",
>>>                             "4.6_multikernel_invokes.f90"),
>>>               api=TEST_API)
>>> psy = PSyFactory(TEST_API).create(info)
>>> invoke = psy.invokes.get('invoke_0')
>>> schedule = invoke.schedule
>>>
>>> ctrans = Dynamo0p3ColourTrans()
>>> otrans = DynamoOMPParallelLoopTrans()
>>>
>>> # Colour all of the loops
>>> for child in schedule.children:
```

(continues on next page)

(continued from previous page)

```

>>> cschedule, _ = ctrans.apply(child)
>>>
>>> # Then apply OpenMP to each of the colour loops
>>> schedule = cschedule
>>> for child in schedule.children:
>>>     newsched, _ = otrans.apply(child.children[0])
>>>
>>> newsched.view()

```

Colouring in the Dynamo 0.3 API is subject to the following rules:

- Only kernels with an iteration space of CELLS and which modify a continuous field require colouring. Any other type of loop will cause this transformation to raise an exception.
- A kernel may have at most one field with ‘INC’ access
- A separate colour map will be required for each field that is coloured (if an invoke contains >1 kernel call)

apply (*node*, *options=None*)

Performs Dynamo0.3-specific error checking and then uses the parent class to convert the Loop represented by *node* into a nested loop where the outer loop is over colours and the inner loop is over cells of that colour.

Parameters

- **node** (`psyclone.dynamo0p3.DynLoop`) – the loop to transform.
- **options** – a dictionary with options for transformations. :type options: dictionary of string:values or None

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

class `psyclone.transformations.Dynamo0p3KernelConstTrans`

Modifies a kernel so that the number of dofs, number of layers and number of quadrature points are fixed in the kernel rather than being passed in by argument.

```

>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "dynamo0.3"
>>> ast, invokeInfo = parse("file.f90", api=api)
>>> psy=PSyFactory(api).create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import Dynamo0p3KernelConstTrans
>>> trans = Dynamo0p3KernelConstTrans()
>>> for kernel in schedule.coded_kernels():
>>>     new_schedule, _ = trans.apply(kernel, number_of_layers=150)
>>>     kernel_schedule = kernel.get_kernel_schedule()
>>>     kernel_schedule.symbol_table.view()

```

apply (*node*, *options=None*)

Transforms a kernel so that the values for the number of degrees of freedom (if a valid value for the `element_order` arg is provided), the number of quadrature points (if the `quadrature` arg is set to True) and

the number of layers (if a valid value for the `number_of_layers` arg is provided) are constant in a kernel rather than being passed in by argument.

The “`cellshape`”, “`element_order`” and “`number_of_layers`” arguments are provided to mirror the namelist values that are input into an LFRic model when it is run.

Quadrature support is currently limited to XYoZ in this transformation. In the case of XYoZ the number of quadrature points (for horizontal and vertical) are set to the `element_order` + 3 in the LFRic infrastructure so their value is derived.

Parameters

- **node** (`psyclone.psygen.DynKern`) – a kernel node.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["cellshape"]** (*str*) – the shape of the cells. This is provided as it helps determine the number of dofs a field has for a particular function space. Currently only “quadrilateral” is supported which is also the default value.
- **options["element_order"]** (*int*) – the order of the cell. In combination with `cellshape`, this determines the number of dofs a field has for a particular function space. If it is set to `None` (the default) then the dofs values are not set as constants in the kernel, otherwise they are.
- **options["number_of_layers"]** (*int*) – the number of vertical layers in the LFRic model mesh used for this particular run. If this is set to `None` (the default) then the `nlayers` value is not set as a constant in the kernel, otherwise it is.
- **options["quadrature"]** (*bool*) – whether the number of quadrature points values are set as constants in the kernel (`True`) or not (`False`). The default is `False`.

Returns tuple of the modified schedule and a record of the transformation.

Return type (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

Return type *str*

validate (*node, options=None*)

This method checks whether the input arguments are valid for this transformation.

Parameters

- **node** (`psyclone.psygen.DynKern`) – a dynamo 0.3 kernel node.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["cellshape"]** (*str*) – the shape of the elements/cells.
- **options["element_order"]** (*int*) – the order of the elements/cells.
- **options["number_of_layers"]** (*int*) – the number of layers to use.
- **options["quadrature"]** (*bool*) – whether quadrature dimension sizes should or shouldn't be set as constants in a kernel.

Raises TransformationError – if the node argument is not a dynamo 0.3 kernel, the `cellshape` argument is not set to “quadrilateral”, the `element_order` argument is not a 0 or a

positive integer, the number of layers argument is not a positive integer, the quadrature argument is not a boolean, neither element order nor number of layers arguments are set (as the transformation would then do nothing), or the quadrature argument is True but the element order is not provided (as the former needs the latter).

class `psyclone.transformations.Dynamo0p3OMPLoopTrans` (*omp_schedule='static'*)
 Dynamo 0.3 specific orphan OpenMP loop transformation. Adds Dynamo-specific validity checks. Actual transformation is done by *base class*.

apply (*node, options=None*)

Perform Dynamo 0.3 specific loop validity checks then call `OMPLoopTrans.apply()`.

Parameters

- **node** (`psyclone.psyir.nodes.Node`) – the Node in the Schedule to check
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations and validation.
- **options["reprod"]** (*bool*) – indicating whether reproducible reductions should be used. By default the value from the config file will be used.

Raises `TransformationError` – if an OMP loop transform would create incorrect code.

property name

Returns the name of this transformation as a string.

class `psyclone.transformations.Dynamo0p3RedundantComputationTrans`

This transformation allows the user to modify a loop's bounds so that redundant computation will be performed. Redundant computation can result in halo exchanges being modified, new halo exchanges being added or existing halo exchanges being removed.

- This transformation should be performed before any parallelisation transformations (e.g. for OpenMP) to the loop in question and will raise an exception if this is not the case.
- This transformation can not be applied to a loop containing a reduction and will again raise an exception if this is the case.
- This transformation can only be used to add redundant computation to a loop, not to remove it.
- This transformation allows a loop that is already performing redundant computation to be modified, but only if the depth is increased.

apply (*loop, options=None*)

Apply the redundant computation transformation to the loop `loop`. This transformation can be applied to loops iterating over 'cells' or 'dofs'. If `depth` is set to a value then the value will be the depth of the field's halo over which redundant computation will be performed. If `depth` is not set to a value then redundant computation will be performed to the full depth of the field's halo.

Parameters

- **loop** (`psyclone.psyGen.DynLoop`) – the loop that we are transforming.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["depth"]** (*int*) – the depth of the stencil. Defaults to None.

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

validate (*node*, *options=None*)

Perform various checks to ensure that it is valid to apply the RedundantComputation transformation to the supplied node

Parameters

- **node** (`psyclone.psyir.nodes.Node`) – the supplied node on which we are performing validity checks
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["depth"]** (*int*) – the depth of the stencil if the value is provided and None if not.

Raises

- **TransformationError** – if the node is not a `psyclone.psyir.nodes.Loop`.
- **TransformationError** – if the parent of the loop is a `psyclone.psyGen.Directive`.
- **TransformationError** – if the parent of the loop is not a `psyclone.psyir.nodes.Loop` or a `psyclone.psyGen.DynInvokeSchedule`.
- **TransformationError** – if the parent of the loop is a `psyclone.psyir.nodes.Loop` but the original loop does not iterate over ‘colour’.
- **TransformationError** – if the parent of the loop is a `psyclone.psyir.nodes.Loop` but the parent does not iterate over ‘colours’.
- **TransformationError** – if the parent of the loop is a `psyclone.psyir.nodes.Loop` but the parent’s parent is not a `psyclone.psyGen.DynInvokeSchedule`.
- **TransformationError** – if this transformation is applied when distributed memory is not switched on.
- **TransformationError** – if the loop does not iterate over cells, dofs or colour.
- **TransformationError** – if the transformation is setting the loop to the maximum halo depth but the loop already computes to the maximum halo depth.
- **TransformationError** – if the transformation is setting the loop to the maximum halo depth but the loop contains a stencil access (as this would result in the field being accessed beyond the halo depth).
- **TransformationError** – if the supplied depth value is not an integer.
- **TransformationError** – if the supplied depth value is less than 1.
- **TransformationError** – if the supplied depth value is not greater than 1 when a continuous loop is modified as this is the minimum valid value.
- **TransformationError** – if the supplied depth value is not greater than the existing depth value, as we should not need to undo existing transformations.
- **TransformationError** – if a depth value has been supplied but the loop has already been set to the maximum halo depth.

GOCEAN1.0 API

12.1 Introduction

The GOcean 1.0 application programming interface (API) was originally designed to support ocean models that use the finite-difference scheme for two-dimensional domains. However, the approach is not specific to ocean models and can potentially be applied to any finite-difference code.

As with all PSyclone APIs, the GOcean 1.0 API specifies how a user must write the Algorithm Layer and the Kernel Layer to allow PSyclone to generate the PSy Layer. These Algorithm and Kernel APIs are discussed separately in the sections below. Before these we describe the functionality provided by the GOcean Library.

12.2 The GOcean Library

The use of PSyclone and the GOcean 1.0 API implies the use of a standard set of data types and associated infrastructure. This is provided by version 1.0 of the GOcean Library (GOLib v.1.0). Currently this library is distributed separately from PSyclone and is available from <http://puma.nerc.ac.uk/trac/GOcean>.

12.2.1 Grid

The GOLib contains a `grid_mod` module which defines a `grid_type` and associated constructor:

```
use grid_mod
...
!> The grid on which our fields are defined
type(grid_type), target :: model_grid
...
! Create the model grid
model_grid = grid_type(GO_ARAKAWA_C,                                &
                      (/GO_BC_EXTERNAL, GO_BC_EXTERNAL, GO_BC_NONE/), &
                      GO_OFFSET_NE)
```

Note: The grid object itself must be declared with the `target` attribute. This is because each field object will contain a pointer to it.

The `grid_type` constructor takes three arguments:

1. The type of grid (only `GO_ARAKAWA_C` is currently supported)
2. The boundary conditions on the domain for the *x*, *y* and *z* dimensions (see below). The value for the *z* dimension is currently ignored.

3. The ‘index offset’ - the convention used for indexing into offset fields.

Three types of boundary condition are currently supported:

Name	Description
GO_BC_NONE	No boundary conditions are applied.
GO_BC_EXTERNAL	Some external forcing is applied. This must be implemented by a kernel. The domain must be defined with a T-point mask (see The grid_init Routine).
GO_BC_PERIODIC	Periodic boundary conditions are applied.

The infrastructure requires this information in order to determine the extent of the model grid.

The index offset is required because a model (kernel) developer has choice in how they actually implement the staggering of variables on a grid. This comes down to a choice of which grid points in the vicinity of a given T point have the same array (i, j) indices. In the diagram below, the image on the left corresponds to choosing those points to the South and West of a T point to have the same (i, j) index. That on the right corresponds to choosing those points to the North and East of the T point (this is the offset scheme used in the NEMO ocean model):



The GOcean 1.0 API supports these two different offset schemes, which we term `GO_OFFSET_SW` and `GO_OFFSET_NE`.

Note that the constructor does not specify the extent of the model grid. This is because this information is normally obtained by reading a file (a namelist file, a netcdf file etc.) which is specific to an application. Once this information has been obtained, a second routine, `grid_init`, is provided with which to ‘load’ a grid object with state. This is discussed below.

The `grid_init` Routine

Once an application has determined the details of the model configuration, it must use this information to populate the grid object. This is done via a call to the `grid_init` subroutine:

```

subroutine grid_init(grid, m, n, dxarg, dyarg, tmask)
  !> The grid object to configure
  type(grid_type), intent(inout) :: grid
  !> Dimensions of the model grid
  integer,          intent(in)    :: m, n
  !> The (constant) grid spacing in x and y (m)
  real(wp),         intent(in)    :: dxarg, dyarg
  !> Optional T-point mask specifying whether each grid point is
  !! wet (1), dry (0) or external (-1).
  integer, dimension(m,n), intent(in), optional :: tmask

```

If no T-mask is supplied then this routine configures the grid appropriately for an all-wet domain with periodic boundary conditions in both the x - and y -dimensions. It should also be noted that currently only grids with constant resolution in x and y are supported by this routine.

12.2.2 Fields

Once a model has a grid defined it will require one or more fields. The GOLib contains a `field_mod` module which defines an `r2d_field` type (real, 2-dimensional field) and associated constructor:

```
use field_mod
...
!> Current ('now') sea-surface height at different grid points
type(r2d_field) :: sshn_u_fld, sshn_v_fld, sshn_t_fld
...

! Sea-surface height now (current time step)
sshn_u = r2d_field(model_grid, GO_U_POINTS)
sshn_v = r2d_field(model_grid, GO_V_POINTS)
sshn_t = r2d_field(model_grid, GO_T_POINTS)
```

The constructor takes two arguments:

1. The grid on which the field exists
2. The type of grid point at which the field is defined (GO_U_POINTS, GO_V_POINTS, GO_T_POINTS or GO_F_POINTS)

Note that the grid object need not have been fully configured (by a call to `grid_init` for instance) before it is passed into this constructor.

12.2.3 Example

PSyclone is distributed with a full example of the use of the GOcean Library. See <PSYCLONEHOME>/examples/gocean/shallow_alg.f90. In what follows we will walk through a slightly cut-down example for a different program.

The following code illustrates the use of the GOLib in constructing an application:

```
program gocean2d
  use grid_mod   ! From dl_esm_inf
  use field_mod  ! From dl_esm_inf
  use model_mod
  use boundary_conditions_mod

  !> The grid on which our fields are defined. Must have the 'target'
  !! attribute because each field object contains a pointer to it.
  type(grid_type), target :: model_grid

  !> Current ('now') velocity component fields
  type(r2d_field) :: un_fld, vn_fld
  !> 'After' velocity component fields
  type(r2d_field) :: ua_fld, va_fld
  ...

  ! time stepping index
  integer :: istp

  ! Create the model grid. We use a NE offset (i.e. the U, V and F
  ! points immediately to the North and East of a T point all have the
  ! same i,j index). This is the same offset scheme as used by NEMO.
  model_grid = grid_type(GO_ARAKAWA_C,
```

(continues on next page)

(continued from previous page)

```

                                (/GO_BC_EXTERNAL,GO_BC_EXTERNAL,GO_BC_NONE/), &
                                GO_OFFSET_NE)

!! read in model parameters and configure the model grid
CALL model_init(model_grid)

! Create fields on this grid

! Velocity components now (current time step)
un_fld = r2d_field(model_grid, GO_U_POINTS)
vn_fld = r2d_field(model_grid, GO_V_POINTS)

! Velocity components 'after' (next time step)
ua_fld = r2d_field(model_grid, GO_U_POINTS)
va_fld = r2d_field(model_grid, GO_V_POINTS)

...

!! time stepping
do istp = nit000, nitend, 1

    call step(istp,
              ua_fld, va_fld, un_fld, vn_fld,
              ...)
end do
...
end program gocean2d

```

The `model_init` routine is application specific since it must determine details of the model configuration being run, e.g. by reading a namelist file. An example might look something like:

```

subroutine model_init(grid)
  type(grid_type), intent(inout) :: grid

  !> Problem size, read from namelist
  integer :: jpiglo, jpjglo
  real(wp) :: dx, dy
  integer, dimension(:,,:), allocatable :: tmask

  ! Read model configuration from namelist
  call read_namelist(jpiglo, jpjglo, dx, dy, &
                    nit000, nitend, irecord, &
                    jphgr_msh, dep_const, rdt, cbfr, visc)

  ! Set-up the T mask. This defines the model domain.
  allocate(tmask(jpiglo, jpjglo))

  call setup_tpoints_mask(jpiglo, jpjglo, tmask)

  ! Having specified the T points mask, we can set up mesh parameters
  call grid_init(grid, jpiglo, jpjglo, dx, dy, tmask)

  ! Clean-up. T-mask has been copied into the grid object.
  deallocate(tmask)

end subroutine model_init

```


Here, only `grid_type` and the `grid_init` routine come from the GOLib. The remaining code is all application specific.

Once the grid object is fully configured and all fields have been constructed, a simulation will proceed by performing calculations with those fields. In the example program given above, this calculation is performed in the time-stepping loop within the `step` subroutine. The way in which this routine uses Invoke calls is described in the [Invokes](#) Section.

12.3 Algorithm

The Algorithm is the top-level specification of the natural science implemented in the software. Essentially it consists of mesh setup, field declarations, initialisation of fields and (a series of) Kernel calls. Infrastructure to support these tasks is provided in version 1.0 of the GOcean library (see [The GOcean Library](#)).

12.3.1 Invokes

The Kernels to call are specified through the use of Invokes, e.g.:

```
call invoke( kernel1(field1, field2),      &
             kernel2(field1, field3)      &
             )
```

The location and number of these `call invoke(...)` statements within the source code is entirely up to the user. The only requirement is that PSyclone must be run on every source file that contains one or more Invokes. The body of each Invoke specifies the kernels to be called, the order in which they are to be applied and the fields (and scalars) that they work with.

Note that the kernel names specified in an Invoke are the names of the corresponding kernel *types* defined in the kernel metadata (see the [Kernel](#) Section). These are not the same as the names of the Fortran subroutines which contain the actual kernel code. The kernel arguments are typically field objects, as described in the [Fields](#) Section, but they may also be scalar quantities (real or integer).

In the example `gocean2d` program shown earlier, there is only one Invoke call and it is contained within the `step` subroutine:

```
subroutine step(istp,      &
               ua, va, un, vn,      &
               sshn_t, sshn_u, sshn_v, &
               ssha_t, ssha_u, ssha_v, &
               hu, hv, ht)
  use kind_params_mod ! From dl_esm_inf
  use grid_mod        ! From dl_esm_inf
  use field_mod        ! From dl_esm_inf
  use model_mod, only: rdt ! The model time-step
  use continuity_mod, only: continuity
  use momentum_mod, only: momentum_u, momentum_v
  use boundary_conditions_mod, only: bc_ssh, bc_solid_u
  !> The current time step
  integer,          intent(inout) :: istp
  type(r2d_field), intent(inout) :: un, vn, sshn_t, sshn_u, sshn_v
  type(r2d_field), intent(inout) :: ua, va, ssha_t, ssha_u, ssha_v
  type(r2d_field), intent(inout) :: hu, hv, ht

  call invoke(
    continuity(ssha_t, sshn_t, sshn_u, sshn_v,      &
```

(continues on next page)

(continued from previous page)

```

        hu, hv, un, vn, rdt),          &
        momentum_u(ua, un, vn, hu, hv, ht,          &
        ssha_u, sshn_t, sshn_u, sshn_v),          &
        momentum_v(va, un, vn, hu, hv, ht,          &
        ssha_v, sshn_t, sshn_u, sshn_v),          &
        bc_ssh(istp, ssha_t),          &
        bc_solid_u(ua),          &
        ...
    )
end subroutine step

```

Note that in this example the grid was constructed for a model with ‘external’ boundary conditions. These boundary conditions are applied through several user-supplied kernels, two of which (`bc_ssh` and `bc_solid_u`) are include in the above code fragment.

12.4 Kernel

The general requirements for the structure of a Kernel are explained in the [Kernel layer](#) section. This section explains the metadata and subroutine arguments that are specific to the GOcean 1.0 API.

12.4.1 Metadata

The metadata for a GOcean 1.0 API kernel has four components:

- 1) ‘meta_args’,
- 2) ‘iterates_over’,
- 3) ‘index_offset’ and
- 4) ‘procedure’:

These are illustrated in the code below:

```

type, extends(kernel_type) :: my_kernel_type
  type(go_arg), dimension(...) :: meta_args = (/ ... /)
  integer :: iterates_over = ...
  integer :: index_offset = ...
contains
  procedure, nopass :: code => my_kernel_code
end type my_kernel_type

```

These four metadata elements are discussed in order in the following sections.

Argument Metadata: meta_args

The `meta_args` array specifies information about data that the kernel code expects to be passed to it via its argument list. There is one entry in the `meta_args` array for each **scalar**, **field**, or **grid-property** passed into the Kernel. Their ordering in the `meta_args` array must be the same as that in the kernel code argument list. The entry must be of type `go_arg` which itself contains metadata about the associated argument. The size of the `meta_args` array must correspond to the total number of **scalars**, **fields** and **grid properties** passed into the Kernel.

For example, if there are a total of two **field** entities being passed to the Kernel then the `meta_args` array will be of size 2 and there will be two entries of type `GO_arg`:

```
type(GO_arg) :: meta_args(2) = (/                                &
    go_arg( ... ),                                              &
    go_arg( ... )                                              &
    /)
```

Argument-metadata (metadata contained within the brackets of an `go_arg` entry), describes either a **scalar**, a **field** or a **grid property**.

The first argument-metadata entry describes how the kernel will access the corresponding argument. As an example, the following `meta_args` metadata describes four entries, the first one is written to by the kernel while the remaining three are only read:

```
type(go_arg) :: meta_args(4) = (/                                &
    go_arg(GO_WRITE, ... ),                                    &
    go_arg(GO_READ, ... ),                                    &
    go_arg(GO_READ, ... ),                                    &
    go_arg(GO_READ, ... )                                     &
    /)
```

The second entry to argument-metadata (information contained within the brackets of an `go_arg` type) describes the type of data represented by the argument. This type falls into three categories; field data, scalar data and grid properties. For field data the metadata entry consists of the type of grid-point that field values are defined on. Since the GOcean API supports fields on an Arakawa C grid, the possible grid-point types are `GO_CU`, `GO_CV`, `GO_CF` and `GO_CT`. GOcean Kernels can also take scalar quantities as arguments. Since these do not live on grid-points they are specified as either `GO_R_SCALAR` or `GO_I_SCALAR` depending on whether the corresponding Fortran variable is a real or integer quantity. Finally, grid-property entries are used to specify any properties of the grid required by the kernel (*e.g.* the area of cells at U points or whether T points are wet or dry).

For example:

```
type(go_arg) :: meta_args(4) = (/                                &
    go_arg(GO_WRITE, GO_CT, ... ),                            &
    go_arg(GO_READ,  GO_CU, ... ),                            &
    go_arg(GO_READ,  GO_R_SCALAR, ... ),                      &
    go_arg(GO_READ,  GO_GRID_AREA_U)                          &
    /)
```

Here, the first argument is a field on T points, the second is a field on U points, the fourth is a real scalar and the fifth is a property of the grid (cell area at U points).

The full list of supported grid properties in the GOcean 1.0 API is:

Table 12.1: Grid Properties Table

Name	Description	Type
go_grid_area_t	Cell area at T point	Real array, rank=2
go_grid_area_u	Cell area at U point	Real array, rank=2
go_grid_area_v	Cell area at V point	Real array, rank=2
go_grid_mask_t	T-point mask (1=wet, 0=dry)	Integer array, rank=2
go_grid_dx_t	Grid spacing in x at T points	Real array, rank=2
go_grid_dx_u	Grid spacing in x at U points	Real array, rank=2
go_grid_dx_v	Grid spacing in x at V points	Real array, rank=2
go_grid_dy_t	Grid spacing in y at T points	Real array, rank=2
go_grid_dy_u	Grid spacing in y at U points	Real array, rank=2
go_grid_dy_v	Grid spacing in y at V points	Real array, rank=2
go_grid_lat_u	Latitude of U points (gphiu)	Real array, rank=2
go_grid_lat_v	Latitude of V points (gphiv)	Real array, rank=2
go_grid_dx_const	Grid spacing in x if constant	Real, scalar
go_grid_dy_const	Grid spacing in y if constant	Real, scalar
go_grid_x_min_index	Minimum X index	Integer, scalar
go_grid_x_max_index	Maximum X index	Integer, scalar
go_grid_y_min_index	Minimum Y index	Integer, scalar
go_grid_y_max_index	Maximum Y index	Integer, scalar

These are defined in the psyclone config file (see [Configuration](#)), and the user or infrastructure library developer can provide additional entries if required. PSyclone will query PSyclone's Configuration class to get the properties required. All of the rank-two arrays have the first rank as longitude (x) and the second as latitude (y).

Scalars and fields contain a third argument-metadata entry which describes whether the kernel accesses the corresponding argument with a stencil. The value `GO_POINTWISE` indicates that there is no stencil access. Metadata for a scalar field is limited to this value. Grid-property arguments have no third metadata argument. If there are no stencil accesses then the full argument metadata for our previous example will be:

```
type(go_arg) :: meta_args(4) = (/
    go_arg(GO_WRITE, GO_CT,      GO_POINTWISE),
    go_arg(GO_READ,  GO_CU,      GO_POINTWISE),
    go_arg(GO_READ,  GO_R_SCALAR, GO_POINTWISE),
    go_arg(GO_READ,  GO_GRID_AREA_U)
/)
```

If a kernel accesses a field using a stencil then the third argument metadata entry should take the form `go_stencil(...)`. Note, a stencil access is only allowed for a field that is READ by a kernel.

In the GOcean API, fields are implemented as two-dimensional arrays. In Fortran, a standard 5-point stencil would look something like the following:

```
a(i,j) + a(i+1,j) + a(i-1,j) + a(i,j+1) + a(i,j-1)
```

If we view the above accesses as co-ordinates relative to the `a(i,j)` access we get `(0,0)`, `(1,0)`, `(-1,0)`, `(0,1)`, `(0,-1)`. If we then view these accesses in graphical form with `i` being in the horizontal direction and `j` in the vertical and with a 1 indicating a (depth-1) access and a 0 indicating there is no access we get the following:

```
010
111
010
```

In the GOcean API a stencil access is captured as a triplet of integers (one row at a time from top to bottom) using the above view i.e.

```
go_stencil(010,111,010)
```

So far we have only considered depth-1 stencils. In our notation the depth of access is captured by the integer value (0 for no access, 1 for depth 1, 2 for depth 2 etc). For example:

```
a(i,j) + a(i,j+1) + a(i,j+2)
```

would be captured as:

```
go_stencil(020,010,000)
```

All forms of stencil can be **summarised** using this triplet notation up to a depth of 9 apart from the central $a(i, j)$ value which can either be 0 (not accessed) or 1 (accessed). Note, the central value is not currently used by PSyclone. The notation is a **summary** in two ways

- 1) it only captures the depth of the stencil in a particular direction, not the actual accesses. Therefore, there is no way to distinguish between the stencil $a(i+2, j)$ and the stencil $a(i+1, j) + a(i+2, j)$.
- 2) when there are offsets for both i and j e.g. $a(i+1, j+1)$ it only captures whether there is an access in that direction at a particular depth, not the details of the access. For example, there is no way to distinguish between $a(i+2, j+2)$ and $a(i+2, j+2) + a(i+1, j+2) + a(i+2, j+1)$.

Whilst the description is a summary, it is accurate enough for PSyclone as this information is primarily used to determine which grid partitions must communicate with which for the purposes of placing halo exchange calls. In this case, it is the depth and direction information that is most important.

Iterates Over

The second element of kernel metadata is `ITERATES_OVER`. This specifies that the Kernel has been written with the assumption that it is iterating over grid points of the specified type. By default the supported values are: `GO_INTERNAL_PTS`, `GO_EXTERNAL_PTS` and `GO_ALL_PTS`. These may be understood by considering the following diagram of an example model configuration:



GO_INTERNAL_PTS are then those points that are within the Model domain (fuchsia box), GO_EXTERNAL_PTS are

those outside the domain and `GO_ALL_PTS` encompasses all grid points in the model. The chosen value is specified in the kernel-meta data like so:

```
integer :: iterates_over = GO_INTERNAL_PTS
```

A user can use a config file (see [Configuration](#)) to add additional iteration spaces to PSyclone.

Index Offset

The third element of kernel metadata, `INDEX_OFFSET`, specifies the index-offset that the kernel uses. This is the same quantity as supplied to the grid constructor (see the [Grid](#) Section for a description).

The GOcean 1.0 API supports two different offset schemes; `GO_OFFSET_NE`, `GO_OFFSET_SW`. The scheme used by a kernel is specified in the metadata as, e.g.:

```
integer :: index_offset = GO_OFFSET_NE
```

Currently all kernels used in an application must use the same offset scheme which must also be the same as passed to the grid constructor.

Procedure

The fourth and final type of metadata is `procedure` metadata. This specifies the name of the Kernel Fortran subroutine that this metadata describes.

For example:

```
procedure :: my_kernel_code
```

12.4.2 Subroutine

Rules

Kernel arguments follow a set of rules which have been specified for the GOcean 1.0 API. These rules are encoded in the `gen_code()` method of the `GOKern` class in the `goceanlp0.py` file. The rules, along with PSyclone's naming conventions, are:

- 1) Every kernel has the indices of the current grid point as the first two arguments, `i` and `j`. These are integers and have intent `in`.
- 2) For each field/scalar/grid property in the order specified by the `meta_args` metadata:
 - 1) For a field; the field array itself. A field array is a real array of kind `go_wp` and rank two. The first rank is longitude (`x`) and the second latitude (`y`).
 - 2) For a scalar; the variable itself. A real scalar is of kind `go_wp`.
 - 3) For a grid property; the array or variable (see the earlier table) containing the specified property.

Note: Grid properties are not passed from the Algorithm Layer. PSyclone generates the necessary lookups in the PSy Layer and includes the resulting references in the arguments passed to the kernel.

As an example, consider the `bc_solid_u` kernel that is used in the `gocean2d` program shown earlier. The metadata for this kernel is:

```
type, extends(kernel_type) :: bc_solid_u
  type(go_arg), dimension(2) :: meta_args = &
    (/ go_arg(GO_WRITE, GO_CU, GO_POINTWISE),      &
      go_arg(GO_READ,      GO_GRID_MASK_T)        &
    /)

  !> This is a boundary-conditions kernel and therefore
  !! acts on all points of the domain rather than just
  !! those that are internal
  integer :: ITERATES_OVER = GO_ALL_PTS

  integer :: index_offset = GO_OFFSET_NE

contains
  procedure, nopass :: code => bc_solid_u_code
end type bc_solid_u
```

The interface to the subroutine containing the implementation of this kernel is:

```
subroutine bc_solid_u_code(ji, jj, ua, tmask)
  integer,          intent(in)    :: ji, jj
  integer, dimension(:, :), intent(in) :: tmask
  real(wp), dimension(:, :), intent(inout) :: ua
```

As described above, the first two arguments to this subroutine specify the grid-point at which the computation is to be performed. The third argument is the field that this kernel updates and the fourth argument is the T-point mask. The latter is a property of the grid and is provided to the kernel call from the PSy Layer.

Comparing this interface definition with the use of the kernel in the Invoke call:

```
call invoke ( ...,      &
             bc_solid_u(ua), &
             ... )
```

we see that in the Algorithm Layer the user need only provide the field(s) (and possibly scalars) that a kernel operates on. The index of the grid point and any grid properties are provided in the (generated) PSy Layer where the kernel subroutine proper is called.

12.5 Built-ins

The GOcean 1.0 API does not support any built-in operations.

12.6 Conventions

The GOcean 1.0 API kernel code conforms to the PSyclone Fortran naming conventions (see *Fortran Naming Conventions*). However, PSyclone's support for the GOcean 1.0 API does not rely on this convention.

The contents of the kernel metadata is usually declared private but this does not affect PSyclone.

Finally, the `procedure` metadata (located within the kernel metadata) usually has `nopass` specified but again this is ignored by PSyclone.

12.7 Configuration

The configuration file (see [Configuration](#)) used by PSyclone can contain GOcean 1.0 specific options. For example, after the default section the GOcean 1.0 specific section looks like this:

```
[gocean1.0]
iteration-spaces=offset_sw:ct:test_only:1:2:3:4
                offset_sw:ct:internal_ns_halo:{start}-1:{stop}+1:{start}:{stop}
```

The supported keys are listed in the next section.

12.7.1 Iteration-spaces

This section lists additional iteration spaces that can be used in a kernel metadata declaration to allow PSyclone to create a loop with different loop boundaries. Each line of the `iteration-spaces` declaration contains 7 values, separated by `:`. The fields are:

Field	Description	Details
1	Index Offset	See <i>Index Offset</i> .
2	grid-point types	See <i>Grid point types</i> .
3	Iterates Over	See <i>Iterates Over</i> .
4	Start index of outer loop	Start index of North-South loop.
5	End index of outer loop	End index of North-South loop.
6	Start index of inner loop	Start index of East-West loop.
7	End index of inner loop	End index of East-West loop.

Two special variables can be used in an iteration space: `{start}` and `{stop}`. These values will be replaced by PSyclone with the correct loop boundaries for the inner points of a grid (i.e. the non-halo area). This means that the depth-1 halo region can be specified using `{start}-1` and `{stop}+1`.

For example, given the `iteration-spaces` declaration above, a kernel declared with `iterates_over=internal_ns_halo` for a field type `ct` and index offset `offset_sw` would create the following loop boundaries:

```
DO j=2-1, jstop+1
  DO i=2, istop
    CALL (i, j, ...)
  END DO
END DO
```

Warning: With user defined iteration spaces it is possible that PSyclone will create code that does not compile: if you specify syntactically correct, but semantically incorrect boundary definitions, the PSyclone internal tests will accept the new iteration space, but the compiler will not. For example if one of the loop boundaries contains the name of a variable that is not defined, compilation will fail. It is the responsibility of the user to make sure that valid loop boundaries are specified in a new iteration space definition.

12.7.2 Grid Properties

Various grid properties can be specified as parameters to a kernel. The actual names and meaning of these properties depend on the infrastructure library used. By default PSyclone provides settings for the `dl_esm_inf` infrastructure library. But the user or a library developer can change or add definitions to the configuration file as required.

The grid properties are specified as values for the key `grid-properties`. They consist of three entries, separated by “.”.

- The first entry is the name of the property as used in kernel metadata.
- The next entry is the way of dereferencing the corresponding value in Fortran. The expression `{0}` is replaced with the field name that is used. Note that any `%` must be replaced with `%%` (due to the way Python reads in configuration files).
- The last entry specifies whether the value is an array or a scalar.

Below an excerpt from the configuration file that is distributed with PSyclone:

```
grid-properties = go_grid_xstop: {0}%%grid%%subdomain%%internal%%xstop: scalar,
                  go_grid_ystop: {0}%%grid%%subdomain%%internal%%ystop: scalar,
                  go_grid_data: {0}%%data: array,
                  go_grid_internal_inner_stop: {0}%%internal%%xstop: scalar,
                  go_grid_internal_outer_stop: {0}%%internal%%ystop: scalar,
                  go_grid_whole_inner_stop: {0}%%whole%%xstop: scalar,
                  go_grid_whole_outer_stop: {0}%%whole%%ystop: scalar,
                  ...
```

Most of the property names can be set arbitrarily by the user (to match whatever infrastructure library is being used), but PSyclone relies on a small number of properties that must be defined with the right name:

Key	Description
<code>go_grid_data</code>	This property gives access to the raw 2d-field.
<code>go_grid_xstop, go_grid_ystop</code>	These values specify the upper loop boundary when computing the constant loop boundaries.
<code>go_grid_{internal,whole}_ _{inner,outer}_{start,stop}</code>	These eight values are required to specify the loop boundaries depending on the field space.
<code>go_grid_nx, go_grid_ny</code>	These properties are only required when OpenCL is enabled. They specify the overall array size (including any padding that the infrastructure library might implement).

12.7.3 Debug Mode

The GOcean configuration also includes a boolean parameter to enable or disable the generation of additional code which may impact performance but is useful for debugging the application. By default it is set to `False`, but it can be changed by updating the following line in the configuration file:

```
[gocean1.0]
DEBUG_MODE = true
```

Currently, only the OpenCL Invokes generate additional debugging code.

12.8 Transformations

In this section we describe the transformations that are specific to the GOcean 1.0 API. For an overview of transformations in general see *Transformations*.

class `psyclone.domain.gocean.transformations.GOceanExtractTrans`

GOcean1.0 API application of ExtractTrans transformation to extract code into a stand-alone program. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>>
>>> API = "gocean1.0"
>>> FILENAME = "shallow_alg.f90"
>>> ast, invokeInfo = parse(FILENAME, api=API)
>>> psy = PSyFactory(API, distributed_memory=False).create(invoke_info)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>>
>>> from psyclone.domain.gocean.transformations import GOceanExtractTrans
>>> etrans = GOceanExtractTrans()
>>>
>>> # Apply GOceanExtractTrans transformation to selected Nodes
>>> newsched, _ = etrans.apply(schedule.children[0])
>>> newsched.view()
```

apply (*nodes*, *options=None*)

Apply this transformation to a subset of the nodes within a schedule - i.e. enclose the specified Nodes in the schedule within a single PSyData region. Note that this implementation just calls the base class, it is only added here to provide the documentation for this function, since it accepts different options to the base class (e.g. `create_driver`, which is passed to the `GOceanExtractNode` instance that will be inserted.).

Parameters

- **nodes** (`psyclone.psyir.nodes.Node` or list of `psyclone.psyir.nodes.Node`) – can be a single node or a list of nodes.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["prefix"]** (*str*) – a prefix to use for the PSyData module name (`prefix_psy_data_mod`) and the PSyDataType (`prefix_PSyDataType`) - a “_” will be added automatically. It defaults to “extract”, resulting in e.g. `extract_psy_data_mod`.
- **options["create_driver"]** (*bool*) – whether or not to create a driver program at code-generation time. If set, the driver will be created in the current working directory with the name “driver-MODULE-REGION.f90” where MODULE and REGION will be the corresponding values for this region. Defaults to False.
- **options["region_name"]** (*((str, str))*) – an optional name to use for this PSyData area, provided as a 2-tuple containing a location name followed by a local name. The pair of strings should uniquely identify a region unless aggregate information is required (and is supported by the runtime library).

Returns Tuple of the modified schedule and a record of the transformation.

Return type (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`)

validate (*node_list*, *options=None*)

Perform GOcean1.0 API specific validation checks before applying the transformation.

Parameters

- **node_list** (list of `psyclone.psyir.nodes.Node`) – the list of Node(s) we are checking.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["create_driver"]** (*bool*) – whether or not to create a driver program at code-generation time. If set, the driver will be created in the current working directory with the name “driver-MODULE-REGION.f90” where MODULE and REGION will be the corresponding values for this region. This flag is forwarded to the `GoceanExtractNode`. Its default value is `False`.
- **options["region_name"]** (*((str, str))*) – an optional name to use for this data-extraction region, provided as a 2-tuple containing a module name followed by a local name. The pair of strings should uniquely identify a region unless aggregate information is required (and is supported by the runtime library). This option is forwarded to the `PSyDataNode` (where it changes the region names) and to the `GOceanExtractNode` (where it changes the name of the created output files and the name of the driver program).

Raises TransformationError – if transformation is applied to an inner Loop without its parent outer Loop.

class `psyclone.transformations.GOceanLoopFuseTrans`

GOcean API specialisation of the *base class* in order to fuse two GOcean loops after performing validity checks (e.g. that the loops are over the same grid-point type). For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast, invokeInfo = parse("shallow_alg.f90")
>>> psy = PSyFactory("gocean1.0").create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import GOceanLoopFuseTrans
>>> ftrans = GOceanLoopFuseTrans()
>>> new_schedule, memento = ftrans.apply(schedule[0], schedule[1])
>>> new_schedule.view()
```

apply (*node1*, *node2*, *options=None*)

Fuses two `psyclone.gocean1p0.GOLoop` loops after performing validity checks by calling `LoopFuseTrans.apply()` method of the base class.

Parameters

- **node1** (`psyclone.gocean1p0.GOLoop`) – the first Node representing a GOLoop.
- **node2** (`psyclone.gocean1p0.GOLoop`) – the second Node representing a GOLoop.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Returns two-tuple of the modified Schedule and a record of the transformation.

Return type (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`)

Raises

- **TransformationError** – if the supplied loops are over different grid-point types.
- **TransformationError** – if there is an unexpected exception.

property name

Returns the name of this transformation as a string.

validate (*node1*, *node2*, *options=None*)

Checks if it is valid to apply the GOceanLoopFuseTrans transform. It ensures that the fused loops are over the same grid-point types, before calling the normal LoopFuseTrans validation function.

Parameters

- **node1** (`psyclone.gocean1p0.GOLoop`) – the first Node representing a GOLoop.
- **node2** (`psyclone.gocean1p0.GOLoop`) – the second Node representing a GOLoop.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises

- **TransformationError** – if the supplied loops are over different grid-point types.
- **TransformationError** – if invalid parameters are passed in.

class `psyclone.transformations.GOceanOMPParallelLoopTrans` (*omp_schedule='static'*)
GOcean specific OpenMP Do loop transformation. Adds GOcean specific validity checks (that supplied Loop is an inner or outer loop). Actual transformation is done by *base class*.

Parameters *omp_schedule* – the omp schedule to be created. Must be one of ‘runtime’, ‘static’, ‘dynamic’, ‘guided’ or ‘auto’.

apply (*node*, *options=None*)

Perform GOcean-specific loop validity checks then call `OMPParallelLoopTrans.apply()`.

Parameters

- **node** (`psyclone.psyir.nodes.Loop`) – a Loop node from an AST.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations and validation.

Raises **TransformationError** – if the supplied node is not an inner or outer loop.

Returns 2-tuple of new schedule and memento of transform

Return type (`psyclone.dynamo0p3.GOInvokeSchedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

class `psyclone.transformations.GOceanOMPLoopTrans` (*omp_schedule='static'*)
GOcean-specific orphan OpenMP loop transformation. Adds GOcean specific validity checks (that the node is either an inner or outer Loop). Actual transformation is done by *base class*.

Parameters `omp_schedule` – the omp schedule to be created. Must be one of ‘runtime’, ‘static’, ‘dynamic’, ‘guided’ or ‘auto’.

apply (*node*, *options=None*)

Perform GOcean specific loop validity checks then call `OMPLoopTrans.apply()`.

Parameters

- **node** (`psyclone.psyir.nodes.Loop`) – the loop to parallelise using OMP Do.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

property name

Returns the name of this transformation as a string.

class `psyclone.transformations.GOConstLoopBoundsTrans`

Switch on (or off) the use of constant loop bounds within a GOInvokeSchedule. In the absence of constant loop bounds, PSyclone will generate loops where the bounds are obtained by de-referencing a field object, e.g.:

```
DO j = my_field%grid%internal%ystart, my_field%grid%internal%ystop
```

Some compilers are able to produce more efficient code if they are provided with information on the relative trip-counts of the loops within an Invoke. With constant loop bounds switched on, PSyclone generates code like:

```
ny = my_field%grid%subdomain%internal%ystop
...
DO j = 1, ny-1
```

In practice, the application of the constant loop bounds looks something like, e.g.:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> import os
>>> TEST_API = "gocean1.0"
>>> _, info = parse(os.path.join("tests", "test_files", "gocean1p0",
>>>                               "single_invoke.f90"),
>>>                  api=TEST_API)
>>> psy = PSyFactory(TEST_API).create(info)
>>> invoke = psy.invokes.get('invoke_0_compute_cu')
>>> schedule = invoke.schedule
>>>
>>> from psyclone.transformations import GOConstLoopBoundsTrans
>>> clbtrans = GOConstLoopBoundsTrans()
>>>
>>> newsched, _ = clbtrans.apply(schedule)
>>> # or, to turn off const. loop bounds:
>>> # newsched, _ = clbtrans.apply(schedule, const_bounds=False)
>>>
>>> newsched.view()
```

apply (*node*, *options=None*)

Switches constant loop bounds on or off for all loops in a GOInvokeSchedule. Default is ‘off’.

Parameters

- **node** (`psyclone.gocean1p0.GOInvokeSchedule`) – the GOInvokeSchedule of which all loops will get the constant loop bounds switched on or off.

- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["const_bounds"]** (*bool*) – whether the constant loop should be used (True) or not (False). Default is True.

Returns 2-tuple of new schedule and memento of transform.

Return type (psyclone.gocean1p0.GOInvokeSchedule, psyclone.undoredo.Memento)

property name

Return the name of the Transformation as a string.

class psyclone.transformations.GOLoopSwapTrans

Provides a loop-swap transformation, e.g.:

```
DO j=1, m
  DO i=1, n
```

becomes:

```
DO i=1, n
  DO j=1, m
```

This transform is used as follows:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast, invokeInfo = parse("shallow_alg.f90")
>>> psy = PSyFactory("gocean1.0").create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import GOLoopSwapTrans
>>> swap = GOLoopSwapTrans()
>>> new_schedule, memento = swap.apply(schedule.children[0])
>>> new_schedule.view()
```

apply (*outer, options=None*)

The argument *outer* must be a loop which has exactly one inner loop. This transform then swaps the outer and inner loop.

Parameters

- **outer** (psyclone.psyir.nodes.Loop) – the node representing the outer loop.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises TransformationError – if the supplied node does not allow a loop swap to be done.

Returns 2-tuple of new schedule and memento of transform.

Return type (psyclone.dynamo0p3.DynInvokeSchedule, psyclone.undoredo.Memento)

property name

Returns the name of this transformation as a string.

validate (*node_outer*, *options=None*)

Checks if the given node contains a valid Fortran structure to allow swapping loops. This means the node must represent a loop, and it must have exactly one child that is also a loop.

Parameters

- **node_outer** (py:class:psyclone.psyir.nodes.Loop) – a Loop node from an AST.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises TransformationError – if the supplied node does not allow a loop swap to be done.

NEMO API

In contrast to the other APIs supported by PSyclone, the NEMO API is designed to work with source code that does *not* follow the PSyKAl separation of concerns. Instead, the NEMO source code is treated as if it were a manually written PSy layer with all kernels in-lined. This approach relies upon the NEMO Coding Conventions [nem13] in order to reason about the code being processed. Rather than construct an InvokeSchedule for the PSy layer from scratch (as is done for other APIs), the InvokeSchedule is constructed by parsing the supplied Fortran code and generating a higher-level representation.

Note: the NEMO API is currently only a prototype. The known issues are listed in [Limitations](#).

13.1 Algorithm

Since NEMO source is treated as a pre-existing PSy layer, this API does not have the concept of an Algorithm layer.

13.2 Constructing the PSyIR

Transformations in PSyclone are applied to an Internal Representation, the “PSyIR.” In contrast to the other APIs where the PSyIR is constructed from scratch, for NEMO PSyclone must parse the existing Fortran and create a higher-level representation of it. This is done using rules based upon the NEMO Coding Conventions [nem13]. These rules are described in the following sections.

13.2.1 Loops

Explicit

PSyclone recognises the following loop types, based on the name of the loop variable:

Loop type	Loop variable
Vertical levels	jk
Latitude	ji
Longitude	jj
Tracer species	jn

PSyclone currently assumes that each of these loop types may be safely parallelised. In practice this will not always be the case (e.g. when performing a tri-diagonal solve) and this implementation will need to be refined.

Implicit

The use of Fortran array notation is encouraged in the NEMO Coding Conventions [nem13] (section 4.2) and is employed throughout the NEMO code base. The Coding Conventions mandate that the shape of every array in such expressions must be specified, e.g.:

```
onedarraya(:) = onedarrayb(:) + onedarrayc(:)
twodarray (:,:) = scalar * anotheartwodarray(:,:)
```

PSyclone therefore also recognises the loops implied by this notation.

Note, not all uses of Fortran array notation in NEMO imply a loop. For instance:

```
ascalar = afunc(twodarray(:,:))
```

is actually a function call which is passed a reference to `twodarray`. However, if the quantity being assigned to is actually an array, e.g.:

```
twodarray2(:,:) = afunc(twodarray(:,:))
```

then this does represent a loop. However, currently PSyclone does not recognise any occurrences of array notation that are themselves within an array access or function call. It is therefore not yet possible to transform such implicit loops into explicit loops. It is hoped that this limitation will be removed in future releases of PSyclone by adding the ability to discover the interface to functions such as `afunc` and thus determining whether they return scalar or array quantities.

13.3 Example

A typical fragment of NEMO source code (taken from the `traldf_iso` routine) is shown below:

```
DO jn = 1, kjpt
  zdit (1, :, :) = 0._wp      ;      zdit (jpi, :, :) = 0._wp
  zdjt (1, :, :) = 0._wp      ;      zdjt (jpi, :, :) = 0._wp

  DO jk = 1, jpkml
    DO jj = 1, jpjml
      DO ji = 1, fs_jpiml
        zdit(ji, jj, jk) = ( ptb(ji+1, jj, jk, jn) - ptb(ji, jj, jk, jn) ) * umask(ji, jj,
→ jk)
        zdjt(ji, jj, jk) = ( ptb(ji, jj+1, jk, jn) - ptb(ji, jj, jk, jn) ) * vmask(ji, jj,
→ jk)
      END DO
    END DO
  END DO
```

PSyclone uses `fparsr2` to parse such source code and then generates the PSy Internal Representation of it:

```
Loop[type='tracers', field_space='None', it_space='None']
  Loop[type='None', field_space='None', it_space='None']
    Loop[type='None', field_space='None', it_space='None']
      Loop[type='None', field_space='None', it_space='None']
        Loop[type='levels', field_space='None', it_space='None']
          Loop[type='lat', field_space='None', it_space='None']
            Loop[type='lon', field_space='None', it_space='None']
              CodedKern[]
```

13.4 Transformations

The following transformations are specific to the NEMO API.

Warning: the transformations below all operate on the PSyIR not the fparser2 ast. This means that any changes made by these transformations will not be observed in the code that is output by PSyclone. The changes will only become visible when the PSyIR back-ends replace the current fparser2 AST-based backend (see issue #435).

class `psyclone.domain.nemo.transformations.NemoArrayRange2LoopTrans`

Provides a transformation from a PSyIR ArrayReference Range to a PSyIR NemoLoop. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "nemo"
>>> filename = "tra_adv.F90" # examples/nemo/code
>>> ast, invoke_info = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invoke_info)
>>> schedule = psy.invokes.invoke_list[0].schedule
>>> schedule.view()
>>>
>>> from psyclone.psyir.nodes import Range
>>> from psyclone.domain.nemo.transformations import
↳ NemoArrayRange2LoopTrans
>>> from psyclone.transformations import TransformationError
>>>
>>> trans = NemoArrayRange2LoopTrans()
>>> for my_range in reversed(schedule.walk(Range)):
>>>     try:
>>>         trans.apply(my_range)
>>>     except TransformationError:
>>>         pass
>>> schedule.view()
```

The specified Range node must be the outermost Range (specifying an access to an array index) within an Array Reference and the array reference must be on the left-hand-side of an Assignment node. This is required for correctness and if not satisfied the transformation will raise an exception.

apply (*node*, *options=None*)

Apply the NemoArrayRange2Loop transformation if the supplied node is the outermost Range node (specifying an access to an array index) within an Array Reference that is on the left-hand-side of an Assignment node. These constraints are required for correctness and an exception will be raised if they are not satisfied. If the constraints are satisfied then the outermost Range nodes within array references within the Assignment node are replaced with references to a loop index. A NemoLoop loop (with the same loop index) is also placed around the modified assignment statement. If the array reference on the left-hand-side of the assignment only had one range node as an index (so now has none) then the assignment is also placed within a NemoKern.

The name of the loop index is taken from the PSyclone configuration file if a name exists for the particular array index, otherwise a new name is generated. The bounds of the loop are taken from the Range node if they are provided. If not, the loop bounds are taken from the PSyclone configuration file if bounds values are supplied. If not, the LBOUND or UBOUND intrinsics are used as appropriate. The type of the NemoLoop is also taken from the configuration file if it is supplied for that index, otherwise it is specified as being “unknown”.

Parameters

- **node** (*psyclone.psyir.nodes.Range*) – a Range node.
- **options** (*dict of string:values or None*) – a dictionary with options for transformations. No options are used in this transformation. This is an optional argument that defaults to None.

validate (*node, options=None*)

Perform various checks to ensure that it is valid to apply the NemoArrayRange2LoopTrans transformation to the supplied PSyIR Node.

Parameters

- **node** (*psyclone.psyir.nodes.Range*) – the node that is being checked.
- **options** (*dict of string:values or None*) – a dictionary with options for transformations. No options are used in this transformation. This is an optional argument that defaults to None.

Raises TransformationError – if the node argument is not a Range, if the Range node is not part of an ArrayReference, if the Range node is not the outermost Range node of the ArrayReference or if that ArrayReference does not constitute the left hand side of an Assignment node.

class `psyclone.domain.nemo.transformations.NemoOuterArrayRange2LoopTrans`

Provides a transformation from the outermost PSyIR ArrayReference Range to a PSyIR NemoLoop. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "nemo"
>>> filename = "tra_adv.F90" # examples/nemo/code
>>> ast, invoke_info = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invoke_info)
>>> schedule = psy.invokes.invoke_list[0].schedule
>>>
>>> from psyclone.psyir.nodes import Assignment
>>> from psyclone.domain.nemo.transformations import
↳ NemoOuterArrayRange2LoopTrans
>>> from psyclone.transformations import TransformationError
>>>
>>> schedule.view()
>>> trans = NemoOuterArrayRange2LoopTrans()
>>> for assignment in schedule.walk(Assignment):
>>>     while True:
>>>         try:
>>>             trans.apply(assignment)
>>>         except TransformationError:
>>>             break
>>> schedule.view()
```

apply (*node, options=None*)

Apply the NemoOuterArrayRange2Loop transformation to the specified node if the node is an Assignment and the left-hand-side of the assignment is an Array Reference containing at least one Range node specifying an access to an array index. If this is the case then the outermost Range nodes within array references within the assignment are replaced with references to a loop index. A NemoLoop loop (with the same loop index) is also placed around the modified assignment statement. If the array reference on the left-hand-side

of the assignment only had one range node as an index (so now has none) then the assignment is also placed within a NemoKern.

The name of the loop index is taken from the PSyclone configuration file if a name exists for the particular array index, otherwise a new name is generated. The bounds of the loop are taken from the Range node if they are provided. If not, the loop bounds are taken from the PSyclone configuration file if bounds values are supplied. If not, the LBOUND or UBOUND intrinsics are used as appropriate. The type of the NemoLoop is also taken from the configuration file if it is supplied for that index, otherwise it is specified as being “unknown”.

Parameters

- **node** (`psyclone.psyir.nodes.Assignment`) – an Assignment node.
- **options** (*dict of string:values or None*) – a dictionary with options for transformations. No options are used in this transformation. This is an optional argument that defaults to None.

validate (*node, options=None*)

Perform various checks to ensure that it is valid to apply the NemoOuterArrayRange2LoopTrans transformation to the supplied PSyIR Node.

Parameters

- **node** (`psyclone.psyir.nodes.Assignment`) – the node that is being checked.
- **options** (*dict of string:values or None*) – a dictionary with options for transformations. No options are used in this transformation. This is an optional argument that defaults to None.

Raises TransformationError – if the supplied node is not an Assignment node, if the Assignment node does not have an ArrayReference node on its left hand side or if the ArrayReference node does not contain at least one Range node.

class `psyclone.domain.nemo.transformations.NemoAllArrayRange2LoopTrans`

Provides a transformation for all PSyIR Array Ranges in an assignment to PSyIR NemoLoops. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "nemo"
>>> filename = "tra_adv.F90" # examples/nemo/code
>>> ast, invoke_info = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invoke_info)
>>> schedule = psy.invokes.invoke_list[0].schedule
>>>
>>> from psyclone.psyir.nodes import Assignment
>>> from psyclone.domain.nemo.transformations import
↳ NemoAllArrayRange2LoopTrans
>>>
>>> schedule.view()
>>> trans = NemoAllArrayRange2LoopTrans()
>>> for assignment in schedule.walk(Assignment):
>>>     trans.apply(assignment)
>>> schedule.view()
```

apply (*node, options=None*)

Apply the NemoAllArrayRange2Loop transformation to the specified node if the node is an Assignment and the left-hand-side of the assignment is an Array Reference containing at least one Range node specifying an access to an array index. If this is the case then all Range nodes within array references within

the assignment are replaced with references to the appropriate loop indices. The appropriate number of NemoLoop loops are also placed around the modified assignment statement and the assignment statement is placed within a NemoKern.

The name of each loop index is taken from the PSyclone configuration file if a name exists for the particular array index, otherwise a new name is generated. The bounds of each loop are taken from the Range node if they are provided. If not, the loop bounds are taken from the PSyclone configuration file if a bounds value is supplied. If not, the LBOUND or UBOUND intrinsics are used as appropriate. The type of the NemoLoop is also taken from the configuration file if it is supplied for that index, otherwise it is specified as being “unknown”.

Parameters

- **node** (`psyclone.psyir.nodes.Assignment`) – an Assignment node.
- **options** (*dict of string:values or None*) – a dictionary with options for transformations. No options are used in this transformation. This is an optional argument that defaults to None.

validate (*node, options=None*)

Perform various checks to ensure that it is valid to apply the NemoArrayRange2LoopTrans transformation to the supplied PSyIR Node.

Parameters

- **node** (`psyclone.psyir.nodes.Assignment`) – the node that is being checked.
- **options** (*dict of string:values or None*) – a dictionary with options for transformations. No options are used in this transformation. This is an optional argument that defaults to None.

Raises TransformationError – if the supplied node is not an Assignment.

13.5 Limitations

The NEMO API is currently only a prototype implementation. Here we list the current, known limitations/issues:

1. Scalar variables inside loops are not made private when parallelising using OpenMP;
2. All recognised loops (levels, latitude etc.) are assumed to be parallelisable. This will not always be the case (e.g. tridiagonal solve has a loop-carried dependence in the vertical);
3. Labelled do-loops are not handled (i.e. they will be put inside a ‘CodeBlock’ in the PSyIR);
4. Loops are currently only permitted to contain one kernel. This restriction will have to be lifted in order to permit loop fusion;
5. When generating new variable names, no attempt is made to avoid clashing with variables already present in the NEMO source. This needs to be resolved by querying the SymbolTable (#381);
6. The `psyir.nodes.Node` base class now has an `_ast` property to hold a pointer into the associated `fparser2` AST. However, the `psyGen.Kern` class already has an `_fp2_ast` property that points to the whole `fparser2` AST of the kernel code. This will be rationalised in #241;
7. The nemo-specific implicit-to-explicit transformations will not modify the code that is output until issue #435 has been addressed.

PSyIR : THE PSyclone INTERNAL REPRESENTATION

The PSyIR is at the heart of PSyclone, representing code (at both the PSy- and kernel-layer levels) in a language-agnostic form. A PSyIR may be constructed from scratch (in Python) or by processing existing source code using a frontend. Transformations act on the PSyIR and ultimately the generated code is produced by one of the PSyIR's backends.

14.1 PSyIR Nodes

The PSyIR consists of classes whose instances can be connected together to form a tree which represent computation in a language-independent way. These classes all inherit from the `Node` baseclass and, as a result, PSyIR instances are often referred to collectively as 'PSyIR nodes'.

At the present time PSyIR classes can be essentially split into two types. PSy-layer classes and Kernel-layer classes. PSy-layer classes make use of a `gen_code()` or an `update()` method to create Fortran code whereas Kernel-layer classes make use of PSyIR backends to create code.

Note: This separation will be removed in the future and eventually all PSyIR classes will make use of backends with the expectation that `gen_code()` and `update()` methods will be removed.

14.1.1 PSy-layer nodes

PSy-layer PSyIR classes are primarily used to create the PSy-layer. These tend to be relatively descriptive and do not specify how a particular PSyclone frontend would implement them. With the exception of `Loop`, these classes are currently not compatible with the PSyIR backends. The generic (non-api-specific) PSy-layer PSyIR nodes are: `InvokeSchedule`, `Directive`, `GlobalSum`, `HaloExchange`, `Loop` and `Kern`. The `Directive` class is subclassed into many directives associated with OpenMP and OpenACC. The `Kern` class is subclassed into `CodedKern`, `InlinedKern` and `BuiltinKern`.

14.1.2 Kernel-layer nodes

Kernel-layer PSyIR classes are currently used to describe existing code in a language independent way. Consequently these nodes are more prescriptive and are independent of a particular PSyclone frontend. These nodes are designed to be used with PSyIR backends. Two PSy-layer classes (`Loop` and `Schedule`) can also be used as Kernel-layer classes. Additionally, the `Schedule` class is further subclassed into a `Routine` and then a kernel-layer `KernelSchedule`. In addition to `KernelSchedule`, Kernel-layer PSyIR nodes are: `Loop`, `IfBlock`, `CodeBlock`, `Assignment`, `Range`, `Reference`, `Operation`, `Literal`, `Call`, `Return` and `Container`. The `Reference` class is further subclassed into `ArrayReference` and the `Operation` class is further subclassed into `UnaryOperation`, `BinaryOperation` and `NaryOperation`.

14.2 Node Descriptions

14.2.1 The Range node

class `psyclone.psyir.nodes.Range` (*ast=None*, *children=None*, *parent=None*, *annotations=None*)

The `Range` node is used to capture a range of integers via `start`, `stop` and `step` expressions. For example, `start=2`, `stop=6` and `step=2` indicates the values 2, 4 and 6.

At the moment the only valid use of `Range` in the PSyIR is to describe a set of accesses to an `Array` dimension (so-called array notation in Fortran). Therefore, the parent of a `Range` node should only be an `Array` node.

The `Range` node has three children nodes, the first child captures the start of the range, the second child captures the end of the range and the third captures the step within the range.

The nodes for each of the children must return an integer. Potentially valid nodes are therefore `Literal`, `Reference`, `Operation` and `CodeBlock`.

A common use case is to want to specify all the elements of a given array dimension without knowing the extent of that dimension. In the PSyIR this is achieved by using the `LBOUND`, and `UBOUND` binary operators:

```
>>> one = Literal("1", INTEGER_TYPE)
>>> # Declare a 1D real array called 'a' with 10 elements
>>> symbol = DataSymbol("a", ArrayType(REAL_TYPE, [10]))
>>> # Return the lower bound of the first dimension of array 'a'
>>> lbound = BinaryOperation.create(
    BinaryOperation.Operator.LBOUND,
    Reference(symbol), one)
>>> # Return the upper bound of the first dimension of array 'a'
>>> ubound = BinaryOperation.create(
    BinaryOperation.Operator.UBOUND,
    Reference(symbol), one)
>>> # Step defaults to 1 so no need to include it when creating range
>>> my_range = Range.create(lbound, ubound)
```

(continues on next page)

(continued from previous page)

```
>>> # Create an access to all elements in the 1st dimension of array 'a'
>>> array_access = Array.create(symbol, [my_range])
```

In Fortran the above access `array_access` can be represented by `a(:)`. The Fortran front-ends and back-ends are aware of array notation. Therefore the Fortran frontend is able to convert array notation to PSyIR and the Fortran backend is able to convert PSyIR back to array notation.

static create (*start*, *stop*, *step=None*, *parent=None*)

Create an internally-consistent Range object. If no step is provided then it defaults to an integer Literal with value 1.

Parameters

- **start** (`psyclone.psyGen.Node`) – the PSyIR for the start value.
- **stop** (`psyclone.psyGen.Node`) – the PSyIR for the stop value.
- **step** (`psyclone.psyGen.Node` or `NoneType`) – the PSyIR for the increment/step or `None`.
- **parent** (`psyclone.psyGen.Node` or `NoneType`) – the parent node of this Range in the PSyIR.

Returns a fully-populated Range object.

Return type `psyclone.psyir.nodes.ranges.Range`

property start

Checks that this Range is valid and then returns the PSyIR for the starting value of the range.

Returns the starting value of this range.

Return type `psyclone.psyGen.Node`

property step

Checks that this Range is valid and then returns the step (increment) value/expression.

Returns the increment used in this range.

Return type `psyclone.psyGen.Node`

property stop

Checks that this Range is valid and then returns the end value/expression.

Returns the end value of this range.

Return type `psyclone.psyGen.Node`

14.3 Text Representation

When developing a transformation script it is often necessary to examine the structure of the PSyIR. All nodes in the PSyIR have the `view` method that writes a text-representation of that node and all of its descendants to `stdout`. If the `termcolor` package is installed (see [Getting Going](#)) then colour highlighting is used for this output. For instance, part of the Schedule constructed for the second NEMO example is rendered as:

```

6: If[]
    BinaryOperation[operator:'OR']
        BinaryOperation[operator:'AND']
            BinaryOperation[operator:'EQ']
                Reference[name:'kpass']
                Literal[value:'1']
            Reference[name:'ln_traldf_lap']
        BinaryOperation[operator:'AND']
            BinaryOperation[operator:'EQ']
                Reference[name:'kpass']
                Literal[value:'2']
            Reference[name:'ln_traldf_blp']
    Schedule[]
        0: If[annotations='was_single_stmt']
            Reference[name:'l_ptr']
            Schedule[]
                0: CodeBlock[<class 'fparser.two.Fortran2003.Call Stmt'>]]
        1: If[annotations='was_single_stmt']
            Reference[name:'l_hst']
            Schedule[]
                0: CodeBlock[<class 'fparser.two.Fortran2003.Call Stmt'>]]
7: Loop[type='levels', field_space='None', it_space='None']
    Literal[value:'1']
    Reference[name:'jpk']
    Literal[value:'1']
    Schedule[]
        0: NemoImplicitLoop[type='None', field_space='None', it_space='None']
8: Loop[type='levels', field_space='None', it_space='None']
    Literal[value:'1']

```

Note that in this view, only those nodes which are children of Schedules have their indices shown. This means that nodes representing e.g. loop bounds or the conditional part of `if` statements are not indexed. For the example shown, the PSyIR node representing the `if(l_hst)` code would be reached by `schedule.children[6].if_body.children[1]` or, using the shorthand notation (see below), `schedule[6].if_body[1]` where `schedule` is the overall parent Schedule node (omitted from the above image).

14.4 Tree Navigation

Each PSyIR node provides several ways to navigate the AST:

The *children* and *parent* properties (available in all nodes) provide an homogeneous method to go up and down the tree hierarchy. This method is recommended when applying general operations or analysis to the tree, however, if one intends to navigate the tree in a way that depends on the type of node, the *children* and *parent* methods should be avoided. The structure of the tree may change in different versions of PSyclone and the encoded navigation won't be future-proof.

To solve this issue some Nodes also provide methods for semantic navigation:

- **Schedule:** subscript operator for indexing the statements (children) inside the Schedule, e.g. `sched[3]` or `sched[2:4]`.
- **Assignment:**

`Assignment.lhs()`

Returns the child node representing the Left-Hand Side of the assignment.

Return type `psyclone.psyir.nodes.Node`

Raises `InternalError` – Node has fewer children than expected.

`Assignment.rhs()`

Returns the child node representing the Right-Hand Side of the assignment.

Return type `psyclone.psyir.nodes.Node`

Raises `InternalError` – Node has fewer children than expected.

- **IfBlock:**

`IfBlock.condition()`

Return the PSyIR Node representing the conditional expression of this IfBlock.

Returns IfBlock conditional expression.

Return type `psyclone.psyir.nodes.Node`

Raises `InternalError` – If the IfBlock node does not have the correct number of children.

`IfBlock.if_body()`

Return the Schedule executed when the IfBlock evaluates to True.

Returns Schedule to be executed when IfBlock evaluates to True.

Return type `psyclone.psyir.nodes.Schedule`

Raises `InternalError` – If the IfBlock node does not have the correct number of children.

`IfBlock.else_body()`

If available return the Schedule executed when the IfBlock evaluates to False, otherwise return None.

Returns Schedule to be executed when IfBlock evaluates to False, if it doesn't exist returns None.

Return type `psyclone.psyir.nodes.Schedule` or `NoneType`

- **Directive:**

`Directive.dir_body()`

Returns the Schedule associated with this directive.

Return type `psyclone.psyir.nodes.Schedule`

Raises `InternalError` – if this node does not have a single Schedule as its child.

These are the recommended methods to navigate the tree for analysis or operations that depend on the Node type.

Additionally, the *walk* method (available in all nodes) is able to recurse through the tree and return objects of a given type. This is useful when the objective is to move down the tree to a specific node or list of nodes without information about the exact location.

`Node.walk(my_type, stop_type=None)`

Recurse through the PSyIR tree and return all objects that are an instance of 'my_type', which is either a single class or a tuple of classes. In the latter case all nodes are returned that are instances of any classes in the tuple. The recursion into the tree is stopped if an instance of 'stop_type' (which is either a single class or a tuple of classes) is found. This can be used to avoid analysing e.g. inlined kernels, or as performance optimisation to reduce the number of recursive calls.

Parameters

- **my_type** (either a single `psyclone.Node` class or a tuple of such classes) – the class(es) for which the instances are collected.
- **stop_type** (None or a single `psyclone.Node` class or a tuple of such classes) – class(es) at which recursion is halted (optional).

Returns list with all nodes that are instances of my_type starting at and including this node.

Return type list of `psyclone.Node` instances.

14.5 DataTypes

The PSyIR supports scalar, array, structure, deferred and unknown datatypes. These datatypes are used when creating instances of DataSymbol and Literal. The ‘deferred’ and ‘unknown’ types are both used when processing existing code. The former is used when a symbol is being imported from some other scope (e.g. via a USE statement in Fortran) that hasn’t yet been resolved and the latter is used when an unsupported form of declaration is encountered.

14.5.1 Scalar DataType

A Scalar datatype consists of an intrinsic and a precision.

The intrinsic can be one of INTEGER, REAL, BOOLEAN and CHARACTER.

The precision can be UNDEFINED, SINGLE, DOUBLE, an integer value specifying the precision in bytes, or a datasymbol (see Section *Symbols and Symbol Tables*) that contains precision information. Note that UNDEFINED, SINGLE and DOUBLE allow the precision to be set by the system so may be different for different architectures. For example:

```
> char_type = ScalarType(ScalarType.Intrinsic.CHARACTER,
>                        ScalarType.Precision.UNDEFINED)
>
> int_type = ScalarType(ScalarType.Intrinsic.INTEGER,
>                      ScalarType.Precision.SINGLE)
>
> bool_type = ScalarType(ScalarType.Intrinsic.BOOLEAN, 4)
>
> symbol = DataSymbol("rdef", int_type, constant_value=4)
> scalar_type = ScalarType(ScalarType.Intrinsic.REAL, symbol)
```

For convenience PSyclone predefines a number of scalar datatypes:

REAL_TYPE, INTEGER_TYPE, BOOLEAN_TYPE and CHARACTER_TYPE all have precision set to UNDEFINED;

REAL_SINGLE_TYPE, REAL_DOUBLE_TYPE, INTEGER_SINGLE_TYPE and INTEGER_DOUBLE_TYPE;

REAL4_TYPE, REAL8_TYPE, INTEGER4_TYPE and INTEGER8_TYPE.

14.5.2 Array DataType

An Array datatype itself has another datatype (or TypeSymbol) specifying the type of its elements and a shape. The shape can have an arbitrary number of dimensions. Each dimension captures what is known about its extent. It is necessary to distinguish between four cases:

Description	Entry in shape list
An array has a static extent known at compile time.	Integer Literal
An array has an extent defined by another symbol.	Symbol
An array has a definite extent which is not known at compile time but can be queried at runtime.	ArrayType.Extent.ATTRIBUTE
It is not known whether an array has memory allocated to it in the current scoping unit.	ArrayType.Extent.DEFERRED

The distinction between the last two cases is that in the former the extents are known but are kept internally with the array (for example an assumed shape array in Fortran) and in the latter the array has not yet been allocated any memory (for example the declaration of an allocatable array in Fortran) so the extents may have not been defined yet.

For example:

```
> array_type = ArrayType(REAL4_TYPE, [5, 10])
>
> n_var = DataSymbol("n", INTEGER_TYPE)
> array_type = ArrayType(INTEGER_TYPE, [n_var, n_var])
>
> array_type = ArrayType(REAL8_TYPE, [ArrayType.Extent.ATTRIBUTE,
>                                     ArrayType.Extent.ATTRIBUTE])
>
> array_type = ArrayType(LOGICAL_TYPE, [ArrayType.Extent.DEFERRED])
```

14.5.3 Structure Datatype

A Structure datatype consists of a dictionary of components where the name of each component is used as the corresponding key. Each component is stored as a named tuple with name, datatype and visibility members.

14.5.4 Unknown DataType

If a PSyIR frontend encounters an unsupported declaration then the corresponding Symbol is given `UnknownType`. The text of the original declaration is stored in the type object and is available via the `declaration` property.

14.6 Symbols and Symbol Tables

Some PSyIR nodes have an associated Symbol Table (*psyclone.psyir.symbols.SymbolTable*) which keeps a record of the Symbols (*psyclone.psyir.symbols.Symbol*) specified and used within them.

Symbol Tables can be nested (i.e. a node with an attached symbol table can be an ancestor or descendent of a node with an attached symbol table). If the same symbol name is used in a hierarchy of symbol tables then the symbol within the symbol table attached to the closest ancestor node is in scope. By default, symbol tables are aware of other symbol tables and will return information about relevant symbols from all symbol tables.

The `SymbolTable` has the following interface:

class `psyclone.psyir.symbols.SymbolTable` (*node=None*)

Encapsulates the symbol table and provides methods to add new symbols and look up existing symbols. Nested scopes are supported and, by default, the add and lookup methods take any ancestor symbol tables into consideration (ones attached to nodes that are ancestors of the node that this symbol table is attached to).

Parameters *node* (`psyclone.psyir.nodes.Schedule`, `psyclone.psyir.nodes.Container` or `NoneType`) – reference to the Schedule or Container to which this symbol table belongs.

Raises `TypeError` – if *node* argument is not a Schedule or a Container.

add (*new_symbol*, *tag=None*, *check_ancestors=True*)

Add a new symbol to the symbol table if the symbol name is not already in use.

Parameters

- **new_symbol** (*psyclone.psyir.symbols.Symbol*) – the symbol to add to the symbol table.
- **tag** (*str*) – a tag identifier for the new symbol, by default no tag is given.

- **check_ancestors** (*bool*) – optional logical flag indicating whether the symbol name and tag (if provided) should be unique in this symbol table (False) or in this and all ancestor symbol tables (True). Defaults to True.

Raises

- **InternalError** – if the new_symbol argument is not a symbol.
- **KeyError** – if the symbol name is already in use.
- **KeyError** – if a tag is supplied and it is already in use.

property argument_datasymbols

Returns list of symbols representing arguments.

Return type list of *psyclone.psyir.symbols.DataSymbol*

property argument_list

Checks that the contents of the SymbolTable are self-consistent and then returns the list of kernel arguments.

Returns ordered list of arguments.

Return type list of *psyclone.psyir.symbols.DataSymbol*

Raises **InternalError** – if the entries of the SymbolTable are not self-consistent.

property containersymbols

Returns a list of the ContainerSymbols present in the Symbol Table.

Return type list of *psyclone.psyir.symbols.ContainerSymbol*

copy_external_global (*globalvar*, *tag=None*)

Copy the given global variable (and its referenced ContainerSymbol if needed) into the SymbolTable.

Parameters

- **globalvar** (*psyclone.psyir.symbols.DataSymbol*) – the variable to be copied in.
- **tag** (*str*) – a tag identifier for the new copy, by default no tag is given.

Raises

- **TypeError** – if the given variable is not a global variable.
- **KeyError** – if the given variable name already exists in the symbol table.

property data_arguments

Returns list of symbols representing kernel data arguments.

Return type list of *psyclone.psyir.symbols.DataSymbol*

Raises **NotImplementedError** – this method is abstract.

property datasymbols

Returns list of symbols representing data variables.

Return type list of *psyclone.psyir.symbols.DataSymbol*

get_unresolved_datasymbols (*ignore_precision=False*)

Create a list of the names of all of the DataSymbols in the table that do not have a resolved interface. If ignore_precision is True then those DataSymbols that are used to define the precision of other DataSymbols are ignored. If no unresolved DataSymbols are found then an empty list is returned.

Parameters `ignore_precision` (*bool*) – whether or not to ignore DataSymbols that are used to define the precision of other DataSymbols.

Returns the names of those DataSymbols with unresolved interfaces.

Return type list of str

property `global_symbols`

Returns list of symbols that have ‘global’ interface (are associated with data that exists outside the current scope).

Return type list of `psyclone.psyir.symbols.Symbol`

imported_symbols (*csymbol*)

Examines the contents of this symbol table to see which DataSymbols (if any) are imported from the supplied ContainerSymbol (which must be present in the SymbolTable).

Parameters `csymbol` (`psyclone.psyir.symbols.ContainerSymbol`) – the ContainerSymbol to search for imports from.

Returns list of DataSymbols that are imported from the supplied ContainerSymbol. If none are found then the list is empty.

Return type list of `psyclone.psyir.symbols.DataSymbol`

Raises

- **TypeError** – if the supplied object is not a ContainerSymbol.
- **KeyError** – if the supplied object is not in this SymbolTable.

property `iteration_indices`

Returns list of symbols representing kernel iteration indices.

Return type list of `psyclone.psyir.symbols.DataSymbol`

Raises **NotImplementedError** – this method is abstract.

property `local_datasymbols`

Returns list of symbols representing local variables.

Return type list of `psyclone.psyir.symbols.DataSymbol`

property `local_typesymbols`

Returns the local TypeSymbols present in the Symbol Table.

Return type list of `psyclone.psyir.symbols.TypeSymbol`

lookup (*name*, *visibility=None*, *check_ancestors=True*)

Look up a symbol in the symbol table (if the *check_ancestors* argument is False) or in this or any ancestor symbol table (if the *check_ancestors* argument is True).

Parameters

- **name** (*str*) – name of the symbol.
- **visibilty** – the visibility or list of visibilities that the symbol must have.
- **check_ancestors** (*bool*) – optional logical flag indicating whether the symbol name should be unique in this symbol table (False) or in this and all ancestor symbol tables (True). Defaults to True.

Returns the symbol with the given name and, if specified, visibility.

Return type `psyclone.psyir.symbols.Symbol`

Raises

- **TypeError** – if the name argument is not a string.
- **SymbolError** – if the name exists in the Symbol Table but does not have the specified visibility.
- **TypeError** – if the visibility argument has the wrong type.
- **KeyError** – if the given name is not in the Symbol Table.

lookup_with_tag (*tag*, *check_ancestors=True*)

Look up a symbol using the supplied tag. If *check_ancestors* is *True* then this and any ancestor symbol tables are searched, otherwise only this symbol table is examined.

Parameters

- **tag** (*str*) – tag identifier.
- **check_ancestors** (*bool*) – optional logical flag indicating whether the tag should be from just this symbol table (*False*) or this and all ancestor symbol tables (*True*). Defaults to *True*.

Returns symbol with the given tag.

Return type `psyclone.psyir.symbols.Symbol`

Raises

- **TypeError** – if the tag argument is not a string.
- **KeyError** – if the given tag is not in the Symbol Table.

name_from_tag (*tag*, *root=None*, *check_ancestors=True*)

If the supplied tag exists in this symbol table (if the *check_ancestors* argument is *False*) or in this or any ancestor symbol table (if the *check_ancestors* argument is *True*), then return the symbol name associated with it, otherwise create a new symbol associated with this tag (using the tag as name or optionally the provided root) and return the name of the new symbol.

Note that this method creates generic Symbols without properties like datatypes and just returns the name string (not the Symbol object). This is commonly needed on the current psy-layer implementation but not recommended on new style PSyIR. This method may be deprecated in the future. (TODO #720)

There is no need to check the argument types as this method calls methods which check the argument types.

Parameters

- **tag** (*str*) – tag identifier.
- **root** (*str*) – optional name of the new symbols if this needs to be created.
- **check_ancestors** (*bool*) – optional logical flag indicating whether the tag should be from just this symbol table (*False*) or this and all ancestor symbol tables (*True*). Defaults to *True*.

Returns name associated with the given tag.

Return type `str`

new_symbol_name (*root_name=None*, *check_ancestors=True*)

Create a symbol name that is not in this symbol table (if the *check_ancestors* argument is *False*) or in this or any ancestor symbol table (if the *check_ancestors* argument is *True*). If the *root_name* argument is not

supplied or if it is an empty string then the name is generated internally, otherwise the *root_name* is used. If required, an additional integer is appended to avoid clashes.

Parameters

- **root_name** (*str* or *NoneType*) – optional name to use when creating a new symbol name. This will be appended with an integer if the name clashes with an existing symbol name.
- **check_ancestors** (*bool*) – optional logical flag indicating whether the name should be unique in this symbol table (False) or in this and all ancestor symbol tables (True). Defaults to True.

Returns the new unique symbol name.

Return type *str*

Raises **TypeError** – if the *root_name* argument is not a string or None.

property node

Returns reference to the Schedule or Container to which this symbol table belongs.

Return type *psyclone.psyir.nodes.Schedule*, *psyclone.psyir.nodes.Container* or *NoneType*

property parent

Returns the ‘parent’ SymbolTable of the current SymbolTable (i.e. the one that encloses this one in the PSyIR hierarchy).

Return type *psyclone.psyir.symbols.SymbolTable* or *NoneType*

property precision_datasymbols

Returns list of all symbols used to define the precision of other symbols within the table.

Return type list of *psyclone.psyir.symbols.DataSymbol*

remove (*symbol*)

Remove the supplied symbol from the Symbol Table. This has a high potential to leave broken links, so this method checks for some references to the removed symbol depending on the symbol type.

Currently, generic Symbols, ContainerSymbols and RoutineSymbols are supported. Support for removing other types of Symbol will be added as required.

TODO #898. This method should check for any references/uses of the target symbol.

Parameters **symbol** (*psyclone.psyir.symbols.ContainerSymbol*) – the container symbol to remove.

Raises

- **TypeError** – if the supplied symbol is not of type Symbol.
- **NotImplementedError** – the removal of this symbol type is not supported yet.
- **KeyError** – if the supplied symbol is not in the symbol table.
- **ValueError** – if the supplied container symbol is referenced by one or more DataSymbols.
- **InternalError** – if the supplied symbol is not the same as the entry with that name in this SymbolTable.

shallow_copy()

Create a copy of the symbol table where the top-level containers are new (symbols added to the new symbol table will not be added in the original but the existing objects are still the same).

Returns a copy of this symbol table.

Return type `psyclone.psyir.symbols.SymbolTable`

specify_argument_list(argument_symbols)

Sets-up the internal list storing the order of the arguments to this kernel.

Parameters **argument_symbols** (*list*) – ordered list of the DataSymbols representing the kernel arguments.

Raises **ValueError** – if the new argument_list is not consistent with the existing entries in the SymbolTable.

swap(old_symbol, new_symbol)

Remove the *old_symbol* from the table and replace it with the *new_symbol*.

Parameters

- **old_symbol** (`psyclone.psyir.symbols.Symbol`) – the symbol to remove from the table.
- **new_symbol** (`psyclone.psyir.symbols.Symbol`) – the symbol to add to the table.

Raises

- **TypeError** – if either old/new_symbol are not Symbols.
- **SymbolError** – if *old_symbol* and *new_symbol* don't have the same name.

swap_symbol_properties(symbol1, symbol2)

Swaps the properties of symbol1 and symbol2 apart from the symbol name. Argument list positions are also updated appropriately.

Parameters

- **symbol1** (`psyclone.psyir.symbols.Symbol`) – the first symbol.
- **symbol2** (`psyclone.psyir.symbols.Symbol`) – the second symbol.

Raises

- **KeyError** – if either of the supplied symbols are not in the symbol table.
- **TypeError** – if the supplied arguments are not symbols, or the names of the symbols are the same in the SymbolTable instance.

property symbols

Returns list of symbols.

Return type list of `psyclone.psyir.symbols.Symbol`

property symbols_dict

Returns ordered dictionary of symbols indexed by symbol name.

Return type `OrderedDict[str] = psyclone.psyir.symbols.Symbol`

property tags_dict

Returns ordered dictionary of symbols indexed by tag.

Return type `OrderedDict[str] = psyclone.psyir.symbols.Symbol`

view()

Print a representation of this Symbol Table to stdout.

Where each element is a Symbol with an immutable name:

```
class psyclone.psyir.symbols.Symbol (name, visibility=<Visibility.PUBLIC: 1>, interface=None)
```

Generic Symbol item for the Symbol Table. It always has a fixed name label that matches with the key in the SymbolTables that contain the symbol. If the symbol is not public then it is only visible to those nodes that are descendents of the Node to which its containing Symbol Table belongs.

Parameters

- **name** (*str*) – name of the symbol.
- **visibility** (*psyclone.psyir.symbols.Symbol.Visibility*) – the visibility of the symbol.
- **interface** (*psyclone.psyir.symbols.symbol.SymbolInterface*) – optional object describing the interface to this symbol (i.e. whether it is passed as a routine argument or accessed in some other way). Defaults to *psyclone.psyir.symbols.LocalInterface*

Raises *TypeError* – if the name is not a str or visibility is not an instance of Visibility.

```
class Visibility (value)
```

Enumeration of the different visibility attributes supported in the PSyIR. If no visibility information is supplied for a Symbol then it is given the DEFAULT_VISIBILITY value.

PUBLIC: the symbol is visible in any scoping region that has access to the SymbolTable containing it.

PRIVATE: the symbol is only visible inside the scoping region that contains the SymbolTable to which it belongs.

copy()

Create and return a copy of this object. Any references to the original will not be affected so the copy will not be referred to by any other object.

Returns A symbol object with the same properties as this symbol object.

Return type *psyclone.psyir.symbols.Symbol*

```
find_symbol_table (node)
```

Searches back up the PSyIR tree for the SymbolTable that contains this Symbol.

Parameters *node* (*psyclone.psyir.nodes.Node*) – the PSyIR node from which to search.

Returns the SymbolTable containing this Symbol or None.

Return type *psyclone.psyir.symbols.SymbolTable* or *NoneType*

Raises *TypeError* – if the supplied *node* argument is not a PSyIR Node.

```
get_external_symbol ()
```

Looks-up and returns the Symbol referred to by this Symbol's external (global) interface.

Raises

- **SymbolError** – if the module pointed to by the symbol interface does not contain the symbol (or the symbol is not public).
- **NotImplementedError** – if the this symbol does not have an external (global) interface.

property interface

Returns the an object describing the interface to this Symbol.

Return type Sub-class of `psyclone.psyir.symbols.symbol.SymbolInterface`

property is_argument

Returns whether the Symbol has an Argument interface.

Return type `bool`

property is_global

Returns whether the Symbol has a Global interface.

Return type `bool`

property is_local

Returns whether the Symbol has a Local interface.

Return type `bool`

property is_unresolved

Returns whether the Symbol has an unresolved interface.

Return type `bool`

property name

Returns name of the Symbol.

Return type `str`

resolve_deferred()

Search for the Container in which this Symbol is defined and create and return a symbol of the correct class and type. If the class and type of the looked-up symbol are the same as this one, some specialisations of this method update the differing properties in place rather than create a new symbol. If this symbol does not have a 'global' interface then there is no lookup needed and we just return this symbol.

Returns a symbol object with the class and type determined by examining the Container from which it is imported.

Return type subclass of `psyclone.psyir.symbols.Symbol`

property visibility

Returns the visibility of this Symbol.

Return type `psyclone.psyir.symbol.Symbol.Visibility`

There are several `Symbol` sub-classes to represent different labeled entities in the PSyIR. At the moment the available symbols are:

- **class** `psyclone.psyir.symbols.ContainerSymbol` (*name*, *visibility*=<`Visibility.PUBLIC`:
I>)

Symbol that represents a reference to a Container. The reference is lazy evaluated, this means that the Symbol will be created without parsing and importing the referenced container, but this can be imported when needed.

Parameters

- **name** (*str*) – name of the symbol.
- **visibility** – the visibility of the symbol.

- **class** `psyclone.psyir.symbols.DataSymbol` (*name*, *datatype*, *visibility*=<*Visibility.PUBLIC*: 1>, *constant_value*=None, *interface*=None)

Symbol identifying a data element. It contains information about: the datatype, the shape (in column-major order) and the interface to that symbol (i.e. Local, Global, Argument).

Parameters

- **name** (*str*) – name of the symbol.
- **datatype** (`psyclone.psyir.symbols.DataType`) – data type of the symbol.
- **visibility** (`psyclone.psyir.symbols.Visibility`) – the visibility of this symbol.
- **constant_value** (NoneType, item of `TYPE_MAP_TO_PYTHON` or `psyclone.psyir.nodes.Node`) – sets a fixed known expression as a permanent value for this `DataSymbol`. If the value is None then this symbol does not have a fixed constant. Otherwise it can receive PSyIR expressions or Python intrinsic types available in the `TYPE_MAP_TO_PYTHON` map. By default it is None.
- **interface** (`psyclone.psyir.symbols.symbol.SymbolInterface`) – object describing the interface to this symbol (i.e. whether it is passed as a routine argument or accessed in some other way).

- **class** `psyclone.psyir.symbols.RoutineSymbol` (*name*, *visibility*=<*Visibility.PUBLIC*: 1>, *interface*=None)

Symbol identifying a callable routine.

14.7 Creating PSyIR

14.7.1 Symbol names

PSyIR symbol names can be specified by a user. For example:

```
> var_name = "my_name"
> symbol_table = SymbolTable()
> data = DataSymbol(var_name, REAL_TYPE)
> symbol_table.add(data)
> reference = Reference(data)
```

However, the `SymbolTable.add()` method will raise an exception if a user tries to add a symbol with the same name as a symbol already existing in the symbol table. Therefore it is the responsibility of the user to determine unique names when manually specifying them.

The `SymbolTable` class also provides the `new_symbol_name()` method (see Section [Symbols and Symbol Tables](#) for more details) to avoid this problem. The method provides a name that is guaranteed to be distinct from any existing names in the symbol table. The name returned is the value of the `PSYIR_ROOT_NAME` variable specified in the `DEFAULT` section of the PSyclone config file, followed by an optional “_” and an integer. For example, the following code:

```
> from psyclone.psyir.symbols import DataSymbol, SymbolTable, REAL4_TYPE
> symbol_table = SymbolTable()
> for i in range(0, 3):
>     var_name = symbol_table.new_symbol_name()
>     symbol_table.add(DataSymbol(var_name, REAL4_TYPE))
>     print (var_name)
```

gives the following output:

```
psyir_tmp
psyir_tmp_0
psyir_tmp_1
```

As the root name (`psyir_tmp` in the example above) is specified in PSyclone's config file it can be set to whatever the user wants.

Note: The particular format used to create a unique name is the responsibility of the `SymbolTable` class and may change in the future.

A user might want to create a name that has some meaning in the context in which it is used e.g. `idx` for an index, `i` for an iterator, or `temp` for a temperature field. To support more readable names, the `new_symbol_name()` method allows the user to specify a root name as an argument to the method which then takes the place of the default root name. For example, the following code:

```
> from psyclone.psyir.symbols import DataSymbol, SymbolTable, REAL_SINGLE_TYPE
> symbol_table = SymbolTable()
> for i in range(0, 3):
>     var_name = symbol_table.new_symbol_name(root_name="something")
>     symbol_table.add(DataSymbol(var_name, REAL_SINGLE_TYPE))
>     print (var_name)
```

gives the following output:

```
something
something_0
something_1
```

An example of using the `new_symbol_name()` method can be found in the PSyclone `examples/psyir` directory.

14.7.2 Nodes

PSyIR nodes are connected together via parent and child methods provided by the `Node` baseclass.

These nodes can be created in isolation and then connected together. For example:

```
> ...
> assignment = Assignment()
> literal = Literal("0.0", REAL_TYPE)
> reference = Reference(symbol)
> literal.parent = assignment
> reference.parent = assignment
> assignment.children = [reference, literal]
```

However, as connections get more complicated, creating the correct connections can become difficult to manage and error prone. Further, in some cases children must be collected together within a `Schedule` (e.g. for `IfBlock` and for `Loop`).

To simplify this complexity, each of the Kernel-layer nodes which contain other nodes have a static `create` method which helps construct the PSyIR using a bottom up approach. Using this method, the above example then becomes:

```
> ...  
> literal = Literal("0.0", REAL_TYPE)  
> reference = Reference(symbol)  
> assignment = Assignment.create(reference, literal)
```

More complete examples of using this approach can be found in the PSyclone `examples/psyir` directory.

TRANSFORMATIONS

As discussed in the previous section, transformations can be applied to a schedule to modify it. Typically transformations will be used to optimise the PSy and/or Kernel layer(s) for a particular architecture, however transformations could be added for other reasons, such as to aid debugging or for performance monitoring.

15.1 Finding

Transformations can be imported directly, but the user needs to know what transformations are available. A helper class **TransInfo** is provided to show the available transformations

Note: The directory layout of PSyclone is currently being restructured. As a result of this some transformations are already in the new locations, while others have not been moved yet. Transformations in the new locations can at the moment not be found using the **TransInfo** approach, and need to be imported directly from the path indicated in the documentation.

class `psyclone.psyGen.TransInfo` (*module=None, base_class=None*)

This class provides information about, and access, to the available transformations in this implementation of PSyclone. New transformations will be picked up automatically as long as they subclass the abstract Transformation class.

For example:

```
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> print(t.list)
There is 1 transformation available:
  1: SwapTrans, A test transformation
>>> # accessing a transformation by index
>>> trans = t.get_trans_num(1)
>>> # accessing a transformation by name
>>> trans = t.get_trans_name("SwapTrans")
```

get_trans_name (*name*)

return the transformation with this name (use list() first to see available transformations)

get_trans_num (*number*)

return the transformation with this number (use list() first to see available transformations)

property list

return a string with a human readable list of the available transformations

property num_trans

return the number of transformations available

15.2 Standard Functionality

Each transformation must provide at least two functions for the user: one for validation, i.e. to verify that a certain transformation can be applied, and one to actually apply the transformation. They are described in detail in the *overview of all transformations*, but the following general guidelines apply.

15.2.1 Validation

Each transformation provides a function `validate`. This function can be called by the user, and it will raise an exception if the transformation can not be applied (and otherwise will return nothing). Validation will always be called when a transformation is applied. The parameters for `validate` can change from transformation to transformation, but each `validate` function accepts a parameter `options`. This parameter is either `None`, or a dictionary of string keys, that will provide additional parameters to the validation process. For example, some validation functions allow part of the validation process to be disabled in order to allow the HPC expert to apply a transformation that they know to be safe, even if the more general validation process might reject it. Those parameters are documented for each transformation, and will show up as a parameter, e.g.: `options["node-type-check"]`. As a simple example:

```
# The validation might reject the application, but in this
# specific case it is safe to apply the transformation,
# so disable the node type check:
my_transform.validate(node, {"node-type-check": False})
```

15.2.2 Application

Each transformation provides a function `apply` which will apply the transformation. It will first validate the transform by calling the `validate` function. Each `apply` function takes the same `options` parameter as the `validate` function described above. Besides potentially modifying the validation process, optional parameters for the transformation are also provided this way. A simple example:

```
kctrans = Dynamo0p3KernelConstTrans()
kctrans.apply(kernel, {"element_order": 0, "quadrature": True})
```

The same `options` dictionary will be used when calling `validate`.

15.3 Available transformations

Most transformations are generic as the schedule structure is independent of the API, however it often makes sense to specialise these for a particular API by adding API-specific errors checks. Some transformations are API-specific (or specific to a set of APIs e.g. `dynamo`). Currently these different types of transformation are indicated by their names.

The generic transformations currently available are listed in alphabetical order below (a number of these have specialisations which can be found in the API-specific sections).

Note: PSyclone currently only supports OpenCL and KernelGlobalsToArguments transformations for the GOcean 1.0 API, the OpenACC Data transformation is limited to the NEMO and GOcean 1.0 APIs and the OpenACC Kernels transformation is limited to the NEMO and Dynamo0.3 APIs.

Note: The directory layout of PSyclone is currently being restructured. As a result of this some transformations are already in the new locations, while others have not been moved yet.

class `psyclone.psyir.transformations.Abs2CodeTrans`

Provides a transformation from a PSyIR ABS Operator node to equivalent code in a PSyIR tree. Validity checks are also performed.

The transformation replaces

```
R = ABS(X)
```

with the following logic:

```
IF X < 0.0:
    R = X*-1.0
ELSE:
    R = X
```

apply (*node*, *options=None*)

Apply the ABS intrinsic conversion transformation to the specified node. This node must be an ABS UnaryOperation. The ABS UnaryOperation is converted to equivalent inline code. This is implemented as a PSyIR transform from:

```
R = ... ABS(X) ...
```

to:

```
tmp_abs = X
if tmp_abs < 0.0:
    res_abs = tmp_abs*-1.0
else:
    res_abs = tmp_abs
R = ... res_abs ...
```

where X could be an arbitrarily complex PSyIR expression and ... could be arbitrary PSyIR code.

This transformation requires the operation node to be a descendent of an assignment and will raise an exception if this is not the case.

Parameters

- **node** (`psyclone.psyGen.UnaryOperation`) – an ABS UnaryOperation node.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Warning: This transformation assumes that the ABS Operator acts on PSyIR Real scalar data and does not check that this is not the case. Once issue #658 is on master then this limitation can be fixed.

class `psyclone.transformations.ACCDataTrans`

Add an OpenACC data region around a list of nodes in the PSyIR. COPYIN, COPYOUT and COPY clauses are added as required.

For example:

```
>>> from psyclone.parse import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "NEMO"
>>> filename = "tra_adv.F90"
>>> ast, invokeInfo = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.transformations import ACCDataTrans
>>> dtrans = ACCDataTrans()
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> kernels = schedule.children[0].children[0].children[0:-1]
>>> # Enclose the kernels
>>> new_sched, _ = dtrans.apply(kernels)
```

apply (*nodes*, *options=None*)

Put the supplied node or list of nodes within an OpenACC data region.

Parameters

- **nodes** ((list of) `psyclone.psyir.nodes.Node`) – the PSyIR node(s) to enclose in the data region.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Returns (transformed schedule, memento of transformation)

Return type 2-tuple of (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`).

class `psyclone.transformations.ACCEnterDataTrans`

Adds an OpenACC “enter data” directive to a Schedule. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> dtrans = t.get_trans_name('ACCEnterDataTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> # Add an enter-data directive
>>> newschedule, _ = dtrans.apply(schedule)
>>> newschedule.view()
```

apply (*sched*, *options=None*)

Adds an OpenACC “enter data” directive to the invoke associated with the supplied Schedule. Any fields accessed by OpenACC kernels within this schedule will be added to this data region in order to ensure they remain on the target device.

Parameters

- **sched** (sub-class of `psyclone.psyir.nodes.Schedule`) – schedule to which to add an “enter data” directive.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Returns tuple of the modified schedule and a record of the transformation.

Return type (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`)

class `psyclone.transformations.ACCKernelsTrans`

Enclose a sub-set of nodes from a Schedule within an OpenACC kernels region (i.e. within “!\$acc kernels” ... “!\$acc end kernels” directives).

For example:

```
>>> from psyclone.parse import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "NEMO"
>>> filename = "tra_adv.F90"
>>> ast, invokeInfo = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.transformations import ACCKernelsTrans
>>> ktrans = ACCKernelsTrans()
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> kernels = schedule.children[0].children[0].children[0:-1]
>>> # Transform the kernel
>>> new_sched, _ = ktrans.apply(kernels)
```

apply (*nodes, options=None*)

Enclose the supplied list of PSyIR nodes within an OpenACC Kernels region.

Parameters

- **nodes** ((list of) `psyclone.psyir.nodes.Node`) – a node or list of nodes in the PSyIR to enclose.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["default_present"]** (*bool*) – whether or not the kernels region should have the ‘default present’ attribute (indicating that data is already on the accelerator). When using managed memory this option should be False.

Returns (transformed schedule, memento of transformation)

Return type 2-tuple of (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`).

class `psyclone.transformations.ACCLoopTrans`

Adds an OpenACC loop directive to a loop. This directive must be within the scope of some OpenACC Parallel region (at code-generation time).

For example:

```

>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory
>>> from psyclone.errors import GenerationError
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ltrans = t.get_trans_name('ACCLoopTrans')
>>> rtrans = t.get_trans_name('ACCPARALLELTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> new_schedule = schedule
>>>
# Apply the OpenACC Loop transformation to *every* loop
# in the schedule
>>> for child in schedule.children:
>>>     newschedule, memento = ltrans.apply(child, reprod=True)
>>>     schedule = newschedule
>>>
# Enclose all of these loops within a single OpenACC
# PARALLEL region
>>> rtrans.omp_schedule("dynamic,1")
>>> newschedule, memento = rtrans.apply(schedule.children)
>>>

```

apply (*node*, *options=None*)

Apply the ACCLoop transformation to the specified node. This node must be a Loop since this transformation corresponds to inserting a directive immediately before a loop, e.g.:

```

!$ACC LOOP
do ...
    ...
end do

```

At code-generation time (when `psyclone.psyGen.ACCLoopDirective.gen_code()` is called), this node must be within (i.e. a child of) a PARALLEL region.

Parameters

- **node** (`psyclone.psyir.nodes.Loop`) – the supplied node to which we will apply the Loop transformation.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["collapse"]** (*int*) – number of nested loops to collapse.
- **options["independent"]** (*bool*) – whether to add the “independent” clause to the directive (not strictly necessary within PARALLEL regions).

Returns 2-tuple of new schedule and memento of transform

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

class `psyclone.transformations.ACCParallelTrans`

Create an OpenACC parallel region by inserting directives. This parallel region *must* come after an enter-data directive (see `ACCEnterDataTrans`) or within a data region (see `ACCDDataTrans`). For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ptrans = t.get_trans_name('ACCParallelTrans')
>>> dtrans = t.get_trans_name('ACCDDataTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> # Enclose everything within a single OpenACC PARALLEL region
>>> newschedule, _ = ptrans.apply(schedule.children)
>>> # Add an enter-data directive
>>> newschedule, _ = dtrans.apply(newschedule)
>>> newschedule.view()
```

apply (*target_nodes*, *options=None*)

Apply this transformation to a subset of the nodes within a schedule - i.e. enclose the specified Loops in the schedule within a single parallel region.

Parameters

- **target_nodes** ((list of) `psyclone.psyir.nodes.Node`) – a single Node or a list of Nodes.
- **options** (dictionary of *string:values* or *None*) – a dictionary with options for transformations.
- **options["node-type-check"]** (*bool*) – this flag controls if the type of the nodes enclosed in the region should be tested to avoid using unsupported nodes inside a region.

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

class `psyclone.psyir.transformations.ArrayRange2LoopTrans`

Provides a transformation from a PSyIR Array Range to a PSyIR Loop. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "nemo"
>>> filename = "tra_adv_compute.F90"
>>> ast, invoke_info = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invoke_info)
>>> schedule = psy.invokes.invoke_list[0].schedule
>>>
```

(continues on next page)

(continued from previous page)

```

>>> from psyclone.psyir.nodes import Assignment
>>> from psyclone.psyir.transformations import ArrayRange2LoopTrans, >>>
↳ TransformationError
>>>
>>> schedule.view()
>>> trans = ArrayRange2LoopTrans()
>>> for assignment in schedule.walk(Assignment):
>>>     while True:
>>>         try:
>>>             trans.apply(assignment)
>>>         except TransformationError:
>>>             break
>>> schedule.view()

```

apply (*node*, *options*=None)

Apply the ArrayRange2Loop transformation to the specified node. The node must be an assignment. The rightmost range node in each array within the assignment is replaced with a loop index and the assignment is placed within a loop iterating over that index. The bounds of the loop are determined from the bounds of the array range on the left hand side of the assignment.

Parameters *node* (psyclone.psyir.nodes.Assignment) – an Assignment node.

Note: The ArrayRange2LoopTrans will have no effect when using the NEMO API until it is updated to use the PSyIR back-ends to generate code (see #435).

class psyclone.transformations.ColourTrans

Apply a colouring transformation to a loop (in order to permit a subsequent parallelisation over colours). For example:

```

>>> invoke = ...
>>> schedule = invoke.schedule
>>>
>>> ctrans = ColourTrans()
>>>
>>> # Colour all of the loops
>>> for child in schedule.children:
>>>     cschedule, _ = ctrans.apply(child)
>>>
>>> csched.view()

```

apply (*node*, *options*=None)

Converts the Loop represented by *node* into a nested loop where the outer loop is over colours and the inner loop is over cells of that colour.

Parameters

- **node** (psyclone.psyir.nodes.Loop) – the loop to transform.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Returns Tuple of modified schedule and record of transformation

Return type (psyclone.psyir.nodes.Schedule, :py:class:`psyclone.undoredo.Memento`)

```
class psyclone.psyir.transformations.extract_trans.ExtractTrans (node_class=<class
                                                                'psy-
                                                                clone.psyir.nodes.extract_node.ExtractNode)
```

This transformation inserts an ExtractNode or a node derived from ExtractNode into the PSyIR of a schedule. At code creation time this node will use the PSyData API to create code that can write the input and output parameters to a file. The node might also create a stand-alone driver program that can read the created file and then execute the instrumented region. Examples are given in the derived classes DynamoExtractTrans and GOceanExtractTrans.

After applying the transformation the Nodes marked for extraction are children of the ExtractNode. Nodes to extract can be individual constructs within an Invoke (e.g. Loops containing a Kernel or BuiltIn call) or entire Invokes. This functionality does not support distributed memory.

Parameters *node_class* (psyclone.psyir.nodes.ExtractNode or derived class) – The Node class of which an instance will be inserted into the tree (defaults to ExtractNode), but can be any derived class.

apply (*nodes, options=None*)

Apply this transformation to a subset of the nodes within a schedule - i.e. enclose the specified Nodes in the schedule within a single PSyData region.

Parameters

- **nodes** (psyclone.psyir.nodes.Node or list of psyclone.psyir.nodes.Node) – can be a single node or a list of nodes.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["prefix"]** (*str*) – a prefix to use for the PSyData module name (PREFIX_psy_data_mod) and the PSyDataType (PREFIX_PSYDATATYPE) - a “_” will be added automatically. It defaults to “”.
- **options["region_name"]** (*((str, str))*) – an optional name to use for this PSyData area, provided as a 2-tuple containing a location name followed by a local name. The pair of strings should uniquely identify a region unless aggregate information is required (and is supported by the runtime library).

Returns Tuple of the modified schedule and a record of the transformation.

Return type (psyclone.psyir.nodes.Schedule, psyclone.undoredo.Memento)

```
class psyclone.transformations.KernelGlobalsToArguments
```

Transformation that removes any accesses of global data from the supplied kernel and places them in the caller. The values/references are then passed by argument into the kernel.

apply (*node, options=None*)

Convert the global variables used inside the kernel into arguments and modify the InvokeSchedule to pass the same global variables to the kernel call.

This apply() method does not return anything, as agreed in #595. However, this change has yet to be applied to the other Transformation classes.

Parameters

- **node** (psyclone.psyGen.CodedKern) – a kernel call.

- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Note: This transformation modifies the PSyIR of both: the Invoke Schedule where the transformed CodedKernel is located and its associated Kernel Schedule.

Note: This transformation is only supported by the GOcean 1.0 API.

class `psyclone.transformations.KernelModuleInlineTrans`

Switches on, or switches off, the inlining of a Kernel subroutine into the PSy layer module. For example:

```
>>> invoke = ...
>>> schedule = invoke.schedule
>>>
>>> inline_trans = KernelModuleInlineTrans()
>>>
>>> ischedule, _ = inline_trans.apply(schedule.children[0].loop_body[0])
>>> ischedule.view()
```

Warning: For this transformation to work correctly, the Kernel subroutine must only use data that is passed in by argument, declared locally or included via use association within the subroutine. Two examples where in-lining will not work are:

1. A variable is declared within the module that contains the Kernel subroutine and is then accessed within that Kernel;
2. A variable is included via use association at the module level and accessed within the Kernel subroutine.

The transformation will reject attempts to in-line such kernels.

apply (*node, options=None*)

Checks that the node is of the correct type (a Kernel) then marks the Kernel to be inlined, or not, depending on the value of the inline option. If the inline option is not passed the Kernel is marked to be inlined.

Parameters

- **node** (`psyclone.psyir.nodes.Loop`) – the loop to transform.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["inline"]** (*bool*) – whether the kernel should be module inlined or not.

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

class `psyclone.transformations.LoopFuseTrans`

Provides a generic loop-fuse transformation to two Nodes in the PSyIR of a Schedule after performing validity checks for the supplied Nodes. Examples are given in the descriptions of any children classes.

apply (*node1*, *node2*, *options=None*)

Fuses two loops represented by *psyclone.psyir.nodes.Node* objects after performing validity checks.

Parameters

- **node1** (*psyclone.psyir.nodes.Node*) – the first Node that is being checked.
- **node2** (*psyclone.psyir.nodes.Node*) – the second Node that is being checked.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Returns two-tuple of the modified Schedule and a record of the transformation.

Return type (*psyclone.psyir.nodes.Schedule*, *psyclone.undoredo.Memento*).

class *psyclone.psyir.transformations.Matmul2CodeTrans*

Provides a transformation from a PSyIR MATMUL Operator node to equivalent code in a PSyIR tree. Validity checks are also performed. Currently only the matrix vector version of MATMUL is supported.

If the dimensions of R, A, and B are $R(N)$, $A(N, M)$, $B(M)$, The transformation replaces:

```
R=MATMUL(A, B)
```

with the following code:

```
do i=1,N
  R(i) = 0.0
  do j=1,M
    R(i) = R(i) + A(i, j) * B(j)
```

apply (*node*, *options=None*)

Apply the MATMUL intrinsic conversion transformation to the specified node. This node must be a MATMUL BinaryOperation. Currently only the matrix vector version of MATMUL is supported. The arguments are permitted to have additional dimensions (i.e. more than 2 for the matrix and more than 1 for the vector) but the matrix can only have two indices which are ranges and these must be the first two indices and the vector can only have one index that is a range and this must be the first index. Further, the ranges must be for the full index space for that dimension (i.e. array subsections are not supported). If the transformation is successful then an assignment which includes a MATMUL BinaryOperation node is converted to equivalent inline code.

Parameters

- **node** (*psyclone.psyGen.BinaryOperation*) – a MATMUL Binary-Operation node.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Note: This transformation is currently limited to translating the matrix vector form of MATMUL to equivalent PSyIR code.

class *psyclone.psyir.transformations.Min2CodeTrans*

Provides a transformation from a PSyIR MIN Operator node to equivalent code in a PSyIR tree. Validity checks are also performed.

The transformation replaces

```
R = MIN(A, B, C ...)
```

with the following logic:

```
R = A
if B < R:
    R = B
if C < R:
    R = C
...
```

apply (*node*, *options=None*)

Apply the MIN intrinsic conversion transformation to the specified node. This node must be an MIN NaryOperation. The MIN NaryOperation is converted to equivalent inline code. This is implemented as a PSyIR transform from:

```
R = ... MIN(A, B, C ...) ...
```

to:

```
res_min = A
tmp_min = B
IF tmp_min < res_min:
    res_min = tmp_min
tmp_min = C
IF tmp_min < res_min:
    res_min = tmp_min
...
R = ... res_min ...
```

where A, B, C ... could be arbitrarily complex PSyIR expressions and the ... before and after MIN (A, B, C ...) can be arbitrary PSyIR code.

This transformation requires the operation node to be a descendent of an assignment and will raise an exception if this is not the case.

Parameters

- **node** (`psyclone.psyGen.BinaryOperation` or `psyclone.psyGen.NaryOperation`) – a MIN Binary- or Nary-Operation node.
- **symbol_table** (`psyclone.psyir.symbols.SymbolTable`) – the symbol table.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Warning: This transformation assumes that the MIN Operator acts on PSyIR Real scalar data and does not check that this is not the case. Once issue #658 is on master then this limitation can be fixed.

class `psyclone.transformations.MoveTrans`

Provides a transformation to move a node in the tree. For example:

```

>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast, invokeInfo=parse("dynamo.F90")
>>> psy=PSyFactory("dynamo0.3").create(invokeInfo)
>>> schedule=psy.invokes.get('invoke_v3_kernel_type').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import MoveTrans
>>> trans=MoveTrans()
>>> new_schedule, memento = trans.apply(schedule.children[0],
>>>                                     schedule.children[2],
>>>                                     position="after")
>>> new_schedule.view()

```

Nodes may only be moved to a new location with the same parent and must not break any dependencies otherwise an exception is raised.

apply (*node*, *location*, *options=None*)

Move the node represented by *node* before *location* (which is also a node) by default and after if the optional *position* argument is set to 'after'.

Parameters

- **node** (`psyclone.psyir.nodes.Node`) – the node to be moved.
- **location** (`psyclone.psyir.nodes.Node`) – node before or after which the given node should be moved.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["position"]** (*str*) – either 'before' or 'after'.

Raises

- **TransformationError** – if the given node is not an instance of `psyclone.psyir.nodes.Node`
- **TransformationError** – if the location is not valid.

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

class `psyclone.transformations.OCLTrans`

Switches on/off the generation of an OpenCL PSy layer for a given `InvokeSchedule`. Additionally, it will generate OpenCL kernels for each of the kernels referenced by the `Invoke`. For example:

```

>>> invoke = ...
>>> schedule = invoke.schedule
>>>
>>> ocl_trans = OCLTrans()
>>> new_sched, _ = ocl_trans.apply(schedule)

```

apply (*sched*, *options=None*)

Apply the OpenCL transformation to the supplied `GOInvokeSchedule`. This causes PSyclone to generate an OpenCL version of the corresponding PSy-layer routine. The generated code makes use of the FortCL library (<https://github.com/stfc/FortCL>) in order to manage the OpenCL device directly from Fortran.

Parameters

- **sched** (psyclone.psyGen.GOInvokeSchedule) – the InvokeSchedule to transform.
- **options** (dictionary of string:values or *None*) – set of option to tune the OpenCL generation.
- **options["opencl"]** (*bool*) – whether or not to enable OpenCL generation.

Returns 2-tuple of new schedule and memento of transform.

Return type (psyclone.dynamo0p3.DynInvokeSchedule, psyclone.undoredo.Memento)

class psyclone.transformations.OMPLoopTrans (*omp_schedule='static'*)

Adds an orphaned OpenMP directive to a loop. i.e. the directive must be inside the scope of some other OMP Parallel REGION. This condition is tested at code-generation time. The optional ‘reprod’ argument in the apply method decides whether standard OpenMP reduction support is to be used (which is not reproducible) or whether a manual reproducible reproduction is to be used.

Parameters *omp_schedule* (*str*) – the OpenMP schedule to use.

For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory
>>> from psyclone.errors import GenerationError
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>> print psy.invokes.names
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ltrans = t.get_trans_name('OMPLoopTrans')
>>> rtrans = t.get_trans_name('OMPParallelTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> new_schedule = schedule
>>>
>>> # Apply the OpenMP Loop transformation to *every* loop
>>> # in the schedule
>>> for child in schedule.children:
>>>     newschedule, memento = ltrans.apply(child, reprod=True)
>>>     schedule = newschedule
>>>
>>> # Enclose all of these loops within a single OpenMP
>>> # PARALLEL region
>>> rtrans.omp_schedule("dynamic,1")
>>> newschedule, memento = rtrans.apply(schedule.children)
>>>
>>>
```

apply (*node*, *options=None*)

Apply the OMPLoopTrans transformation to the specified node in a Schedule. This node must be a Loop since this transformation corresponds to wrapping the generated code with directives like so:

```
!$OMP DO
do ...
...
end do
!$OMP END DO
```

At code-generation time (when `OMPLoopDirective.gen_code()` is called), this node must be within (i.e. a child of) an OpenMP PARALLEL region.

If the keyword “reprod” is specified in the options, it will cause a reproducible reduction to be generated if it is set to `True`, otherwise the default value (as read from the `psyclone.cfg` file) will be used. Note, reproducible in this case means obtaining the same results with the same number of OpenMP threads, not for different numbers of OpenMP threads.

Parameters

- **node** (`psyclone.psyir.nodes.Node`) – the supplied node to which we will apply the `OMPLoopTrans` transformation
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations and validation.
- **options["reprod"]** (*bool*) – indicating whether reproducible reductions should be used. By default the value from the config file will be used.

Returns (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`)

property `omp_schedule`

Returns the OpenMP schedule that will be specified by this transformation. The default schedule is ‘static’.

class `psyclone.transformations.OMPParallelLoopTrans` (*omp_schedule='static'*)

Adds an OpenMP PARALLEL DO directive to a loop.

For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast, invokeInfo = parse("dynamo.F90")
>>> psy = PSyFactory("dynamo0.1").create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_v3_kernel_type').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import OMPParallelLoopTrans
>>> trans = OMPParallelLoopTrans()
>>> new_schedule, memento = trans.apply(schedule.children[0])
>>> new_schedule.view()
```

apply (*node, options=None*)

Apply an `OMPParallelLoop` Transformation to the supplied node (which must be a `Loop`). In the generated code this corresponds to wrapping the Loop with directives:

```
!$OMP PARALLEL DO ...
do ...
...
end do
!$OMP END PARALLEL DO
```

Parameters

- **node** (`psyclone.f2pygen.DoGen`) – the node (loop) to which to apply the transformation.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations and validation.

Returns two-tuple of transformed schedule and a record of the transformation.

Return type (`psyclone.psyir.nodes.Schedule`, `:py:class:`psyclone.undoredo.Memento`)

class `psyclone.transformations.OMPParallelTrans`

Create an OpenMP PARALLEL region by inserting directives. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory
>>> from psyclone.errors import GenerationError
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ltrans = t.get_trans_name('GOceanOMPLoopTrans')
>>> rtrans = t.get_trans_name('OMPParallelTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> new_schedule = schedule
>>>
>>> # Apply the OpenMP Loop transformation to *every* loop
>>> # in the schedule
>>> for child in schedule.children:
>>>     newschedule, memento = ltrans.apply(child)
>>>     schedule = newschedule
>>>
>>> # Enclose all of these loops within a single OpenMP
>>> # PARALLEL region
>>> newschedule, _ = rtrans.apply(schedule.children)
>>> newschedule.view()
```

apply (*target_nodes*, *options=None*)

Apply this transformation to a subset of the nodes within a schedule - i.e. enclose the specified Loops in the schedule within a single parallel region.

Parameters

- **target_nodes** ((list of) `psyclone.psyir.nodes.Node`) – a single Node or a list of Nodes.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["node-type-check"]** (*bool*) – this flag controls if the type of the nodes enclosed in the region should be tested to avoid using unsupported nodes inside a region.

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

get_node_list (*nodes*)

This is a helper function for region based transformations. The parameter for any of those transformations is either a single node, a schedule, or a list of nodes. This function converts this into a list of nodes according to the parameter type. This function will always return a copy, to avoid issues e.g. if a child list of a node should be provided, and a transformation changes the order in this list (which would then also change the order of the nodes in the tree).

Parameters *nodes* (`psyclone.psyir.nodes.Node` or `psyclone.psyir.nodes.Schedule` or a list of `:py:obj:`psyclone.psyir.nodes.Node``) – can be a single node, a schedule or a list of nodes.

Returns a list of nodes.

Return type list of `psyclone.psyir.nodes.Node`

Raises **TransformationError** – if the supplied parameter is neither a single `Node`, nor a `Schedule`, nor a list of `Nodes`.

validate (*node_list*, *options=None*)

Perform OpenMP-specific validation checks.

Parameters

- **node_list** (list of `psyclone.psyir.nodes.Node`) – list of `Nodes` to put within parallel region.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["node-type-check"]** (*bool*) – this flag controls if the type of the nodes enclosed in the region should be tested to avoid using unsupported nodes inside a region.

Raises **TransformationError** – if the target `Nodes` are already within some OMP parallel region.

Note: PSyclone does not support (distributed-memory) halo swaps or global sums within OpenMP parallel regions. Attempting to create a parallel region for a set of nodes that includes halo swaps or global sums will produce an error. In such cases it may be possible to re-order the nodes in the `Schedule` such that the halo swaps or global sums are performed outside the parallel region. The [MoveTrans](#) transformation may be used for this.

class `psyclone.psyir.transformations.ProfileTrans`

Create a profile region around a list of statements. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory, GenerationError
>>> from psyclone.psyir.transformations import ProfileTrans
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> p_trans = ProfileTrans()
>>>
```

(continues on next page)

(continued from previous page)

```

>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> # Enclose all children within a single profile region
>>> newschedule, _ = p_trans.apply(schedule.children)
>>> newschedule.view()

```

This implementation relies completely on the base class PSyDataTrans for the actual work, it only adjusts the name etc, and the list of valid nodes.

apply (*nodes*, *options=None*)

Apply this transformation to a subset of the nodes within a schedule - i.e. enclose the specified Nodes in the schedule within a single PSyData region.

Parameters

- **nodes** (`psyclone.psyir.nodes.Node` or list of `psyclone.psyir.nodes.Node`) – can be a single node or a list of nodes.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["prefix"]** (*str*) – a prefix to use for the PSyData module name (`PREFIX_psy_data_mod`) and the PSyDataType (`PREFIX_PSYDATATYPE`) - a “_” will be added automatically. It defaults to “”.
- **options["region_name"]** (*(str, str)*) – an optional name to use for this PSyData area, provided as a 2-tuple containing a location name followed by a local name. The pair of strings should uniquely identify a region unless aggregate information is required (and is supported by the runtime library).

Returns Tuple of the modified schedule and a record of the transformation.

Return type (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`)

class `psyclone.psyir.transformations.ReadOnlyVerifyTrans` (*node_class=<class 'psyclone.psyir.nodes.read_only_verify_node.ReadOnlyVerifyNode'>*)

This transformation inserts a `ReadOnlyVerifyNode` or a node derived from `ReadOnlyVerifyNode` into the PSyIR of a schedule. At code creation time this node will use the PSyData API to create code that will verify that read-only quantities are not modified.

After applying the transformation the Nodes marked for verification are children of the `ReadOnlyVerifyNode`. Nodes to verify can be individual constructs within an Invoke (e.g. Loops containing a Kernel or BuiltIn call) or entire Invokes.

Parameters **node_class** (`psyclone.psyir.nodes.ReadOnlyVerifyNode` or derived class) – The class of Node which will be inserted into the tree (defaults to `ReadOnlyVerifyNode`), but can be any derived class.

apply (*nodes*, *options=None*)

Apply this transformation to a subset of the nodes within a schedule - i.e. enclose the specified Nodes in the schedule within a single PSyData region.

Parameters

- **nodes** (`psyclone.psyir.nodes.Node` or list of `psyclone.psyir.nodes.Node`) – can be a single node or a list of nodes.

- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["prefix"]** (*str*) – a prefix to use for the PSyData module name (PREFIX_psy_data_mod) and the PSyDataType (PREFIX_PSYDATATYPE) - a “_” will be added automatically. It defaults to “”.
- **options["region_name"]** (*((str, str))*) – an optional name to use for this PSyData area, provided as a 2-tuple containing a location name followed by a local name. The pair of strings should uniquely identify a region unless aggregate information is required (and is supported by the runtime library).

Returns Tuple of the modified schedule and a record of the transformation.

Return type (psyclone.psyir.nodes.Schedule, psyclone.undoredo.Memento)

class psyclone.psyir.transformations.Sign2CodeTrans

Provides a transformation from a PSyIR SIGN Operator node to equivalent code in a PSyIR tree. Validity checks are also performed.

The transformation replaces

```
R = SIGN(A, B)
```

with the following logic:

```
R = ABS(A)
if B < 0.0:
    R = R*-1.0
```

i.e. the value of A with the sign of B

apply (*node, options=None*)

Apply the SIGN intrinsic conversion transformation to the specified node. This node must be a SIGN BinaryOperation. The SIGN BinaryOperation is converted to equivalent inline code. This is implemented as a PSyIR transform from:

```
R = ... SIGN(A, B) ...
```

to:

```
tmp_abs = A
if tmp_abs < 0.0:
    res_abs = tmp_abs*-1.0
else:
    res_abs = tmp_abs
res_sign = res_abs
tmp_sign = B
if tmp_sign < 0.0:
    res_sign = res_sign*-1.0
R = ... res_sign ...
```

where A and B could be arbitrarily complex PSyIR expressions, ... could be arbitrary PSyIR code and where ABS has been replaced with inline code by the NemoAbsTrans transformation.

This transformation requires the operation node to be a descendent of an assignment and will raise an exception if this is not the case.

Parameters

- **node** (`psyclone.psyGen.BinaryOperation`) – a SIGN BinaryOperation node.
- **symbol_table** (`psyclone.psyir.symbols.SymbolTable`) – the symbol table.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Warning: This transformation assumes that the SIGN Operator acts on PSyIR Real scalar data and does not check whether or not this is the case. Once issue #658 is on master then this limitation can be fixed.

15.4 Kernels

PSyclone supports the transformation of Kernels as well as PSy-layer code. However, the transformation of kernels to produce new kernels brings with it additional considerations, especially regarding the naming of the resulting kernels. PSyclone supports two use cases:

1. the HPC expert wishes to optimise the same kernel in different ways, depending on where/how it is called;
2. the HPC expert wishes to transform the kernel just once and have the new version used throughout the Algorithm file.

The second case is really an optimisation of the first for the case where the same set of transformations is applied to every instance of a given kernel.

Since PSyclone is run separately for each Algorithm in a given application, ensuring that there are no name clashes for kernels in the application as a whole requires that some state is maintained between PSyclone invocations. This is achieved by requiring that the same kernel output directory is used for every invocation of PSyclone when building a given application. However, this is under the control of the user and therefore it is possible to use the same output directory for a subset of algorithms that require the same kernel transformation and then a different directory for another subset requiring a different transformation. Of course, such use would require care when building and linking the application since the differently-optimised kernels would have the same names.

By default, transformed kernels are written to the current working directory. Alternatively, the user may specify the location to which to write the modified code via the `-okern` command-line flag.

In order to support the two use cases given above, PSyclone supports two different kernel-renaming schemes: “multiple” and “single” (specified via the `--kernel-renaming` command-line flag). In the default, “multiple” scheme, PSyclone ensures that each transformed kernel is given a unique name (with reference to the contents of the kernel output directory). In the “single” scheme, it is assumed that any given kernel that is transformed is always transformed in the same way (or left unchanged) and thus just one transformed version of it is created. This assumption is checked by examining the Fortran code for any pre-existing transformed version of that kernel. If another transformed version of that kernel exists and does not match that created by the current transformation then PSyclone will raise an exception.

15.4.1 Rules

Kernel code that is to be transformed is subject to certain restrictions. These rules are intended to make kernel transformations as robust as possible, in particular by limiting the amount of code that must be parsed by PSyclone (via `fparser`). The rules are as follows:

- 1) Any variable or procedure accessed by a kernel must either be explicitly declared or named in the `only` clause of a module `use` statement within the scope of the subroutine containing the kernel implementation. This means that:
 - 1) Kernel subroutines are forbidden from accessing data using `COMMON` blocks;
 - 2) Kernel subroutines are forbidden from calling procedures declared via the `EXTERN` statement;
 - 3) Kernel subroutines must not access data or procedures made available via their parent (containing) module.
- 2) The full Fortran source of a kernel must be available to PSyclone. This includes the source of any modules from which it accesses either routines or data. (However, kernel routines are permitted to make use of Fortran intrinsic routines.)

For instance, consider the following Fortran module containing the `bc_ssh_code` kernel:

```

module boundary_conditions_mod
  real :: forbidden_var

contains

  subroutine bc_ssh_code(ji, jj, istep, ssh)
    use kind_params_mod, only: go_wp
    use model_mod, only: rdt
    integer,                intent(in)    :: ji, jj, istep
    real(go_wp), dimension(:, :), intent(inout) :: ssh
    real(go_wp) :: rtime

    rtime = real(istep, go_wp) * rdt
    ...
  end subroutine bc_ssh_code
end module boundary_conditions_mod

```

Since the kernel subroutine accesses data (the `rdt` variable) from the `model_mod` module, the source of that module must be available to PSyclone if a transformation is applied to this kernel. Should `rdt` not actually be defined in `model_mod` (i.e. `model_mod` itself imports it from another module) then the source containing its definition must also be available to PSyclone. Note that the rules forbid the `bc_ssh_code` kernel from accessing the `forbidden_var` variable that is available to it from the enclosing module scope.

Note: these rules *only* apply to kernels that are the target of PSyclone kernel transformations.

15.4.2 Available Kernel Transformations

PSyclone currently provides just one kernel-specific transformation (although there are a number that can be applied to either or both the PSy-layer and Kernel-layer PSyIR):

class `psyclone.transformations.ACCRoutineTrans`

Transform a kernel subroutine by adding a “!\$acc routine” directive (causing it to be compiled for the OpenACC accelerator device). For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.transformations import ACCRoutineTrans
>>> rtrans = ACCRoutineTrans()
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> kern = schedule.children[0].children[0].children[0]
>>> # Transform the kernel
>>> newkern, _ = rtrans.apply(kern)
```

apply (*kern*, *options=None*)

Modifies the AST of the supplied kernel so that it contains an ‘!\$acc routine’ OpenACC directive. This transformation affects the f2pygen and the PSyIR trees of this kernel.

Parameters

- **kern** (`psyclone.psyGen.Kern`) – the kernel object to transform.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises **TransformationError** – if we fail to find the subroutine corresponding to the kernel object.

Returns (transformed kernel, memento of transformation)

Return type 2-tuple of (`psyclone.psyGen.Kern`, `psyclone.undoredo.Memento`).

property name

Returns the name of this transformation class.

Return type `str`

validate (*kern*, *options=None*)

Perform checks that the supplied kernel can be transformed.

Parameters

- **kern** (`psyclone.psyGen.Kern`) – the kernel which is the target of the transformation.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises

- **TransformationError** – if the target kernel is a built-in.

- **TransformationError** – if any of the symbols in the kernel are accessed via a module use statement.

15.5 Applying

Transformations can be applied either interactively or through a script.

15.5.1 Interactive

To apply a transformation interactively we first parse and analyse the code. This allows us to generate a “vanilla” PSy layer. For example ...

```
from psyclone.parse.algorithm import parse
from psyclone.psyGen import PSyFactory

# This example uses version 0.1 of the Dynamo API
api = "dynamo0.1"

# Parse the file containing the algorithm specification and
# return the Abstract Syntax Tree and invokeInfo objects
ast, invokeInfo = parse("dynamo.F90", api=api)

# Create the PSy-layer object using the invokeInfo
psy = PSyFactory(api).create(invokeInfo)

# Optionally generate the vanilla PSy layer fortran
print psy.gen
```

We then extract the particular schedule we are interested in. For example ...

```
# List the various invokes that the PSy layer contains
print psy.invokes.names

# Get the required invoke
invoke = psy.invokes.get('invoke_0_v3_kernel_type')

# Get the schedule associated with the required invoke
schedule = invoke.schedule
schedule.view()
```

Now we have the schedule we can create and apply a transformation to it to create a new schedule and then replace the original schedule with the new one. For example ...

```
# Get the list of possible loop transformations
from psyclone.psyGen import TransInfo
t = TransInfo()
print t.list

# Create an OpenMPLoop-transformation
ol = t.get_trans_name('OMPParallelLoopTrans')

# Apply it to the loop schedule of the selected invoke
new_schedule, memento = ol.apply(schedule.children[0])
new_schedule.view()
```

(continues on next page)

(continued from previous page)

```
# Replace the original loop schedule of the selected invoke
# with the new, transformed schedule
invoke.schedule=new_schedule

# Generate the Fortran code for the new PSy layer
print psy.gen
```

More examples of use of the interactive application of transformations can be found in the `runme*.py` files within the `examples/lfric/eg1` and `examples/lfric/eg2` directories. Some simple examples of the use of transformations are also given in the previous section.

15.5.2 Script

PSyclone provides a Python script (**psyclone**) that can be used from the command line to generate PSy layer code and to modify algorithm layer code appropriately. By default this script will generate “vanilla” (unoptimised) PSy layer code. For example:

```
> psyclone algspec.f90
> psyclone -oalg alg.f90 -opsy psy.f90 -api dynamo0.3 algspec.f90
```

The **psyclone** script has an optional **-s** flag which allows the user to specify a script file to modify the PSy layer as required. Script files may be specified without a path. For example:

```
> psyclone -s opt.py algspec.f90
```

In this case the Python search path **PYTHONPATH** will be used to try to find the script file.

Alternatively, script files may be specified with a path. In this case the file is expected to be found in the specified location. For example ...

```
> psyclone -s ./opt.py algspec.f90
> psyclone -s ../scripts/opt.py algspec.f90
> psyclone -s /home/me/PSyclone/scripts/opt.py algspec.f90
```

PSyclone also provides the same functionality via a function (which is what the **psyclone** script calls internally).

```
psyclone.generator.generate(filename, api="", kernel_path="", script_name=None,
                             line_length=False, distributed_memory=None, kern_out_path="",
                             kern_naming='multiple')
```

Takes a PSyclone algorithm specification as input and outputs the associated generated algorithm and psy codes suitable for compiling with the specified kernel(s) and support infrastructure. Uses the `parse.algorithm.parse()` function to parse the algorithm specification, the `psyGen.PSy` class to generate the PSy code and the `alg_gen.Alg` class to generate the modified algorithm code.

Parameters

- **filename** (*str*) – The file containing the algorithm specification.
- **kernel_path** (*str*) – The directory from which to recursively search for the files containing the kernel source (if different from the location of the algorithm specification).
- **script_name** (*str*) – A script file that can apply optimisations to the PSy layer (can be a path to a file or a filename that relies on the **PYTHONPATH** to find the module).

- **line_length** (*bool*) – A logical flag specifying whether we care about line lengths being longer than 132 characters. If so, the input (algorithm and kernel) code is checked to make sure that it conforms. The default is `False`.
- **distributed_memory** (*bool*) – A logical flag specifying whether to generate distributed memory code. The default is set in the `config.py` file.
- **kern_out_path** (*str*) – Directory to which to write transformed kernel code.
- **kern_naming** (*bool*) – the scheme to use when re-naming transformed kernels.

Returns 2-tuple containing `fparser1` ASTs for the algorithm code and the psy code.

Return type (`fparser.one.block_statements.BeginSource`, `fparser.one.block_statements.Module`)

Raises

- **IOError** – if the filename or search path do not exist
- **GenerationError** – if an invalid API is specified.
- **GenerationError** – if an invalid kernel-renaming scheme is specified.

For example:

```
>>> from psyclone.generator import generate
>>> alg, psy = generate("algspec.f90")
>>> alg, psy = generate("algspec.f90", kernel_path="src/kernels")
>>> alg, psy = generate("algspec.f90", script_name="optimise.py")
>>> alg, psy = generate("algspec.f90", line_length=True)
>>> alg, psy = generate("algspec.f90", distributed_memory=False)
```

A valid script file must contain a **trans** function which accepts a **PSy** object as an argument and returns a **PSy** object, i.e.:

```
def trans(psy)
    ...
    return psy
```

It is up to the script what it does with the **PSy** object. The example below does the same thing as the example in the *Interactive* section.

```
def trans(psy):
    from psyclone.transformations import OMPParallelLoopTrans
    invoke = psy.invokes.get('invoke_0_v3_kernel_type')
    schedule = invoke.schedule
    ol = OMPParallelLoopTrans()
    new_schedule, _ = ol.apply(schedule.children[0])
    invoke.schedule = new_schedule
    return psy
```

Of course the script may apply as many transformations as is required for a particular schedule and may apply transformations to all the schedules (i.e. invokes and/or kernels) contained within the **PSy** layer.

Examples of the use of transformation scripts can be found in many of the examples, such as `examples/lfric/eg3` and `examples/lfric/scripts`. Please read the `examples/lfric/README` file first as it explains how to run the examples (and see also the `examples/check_examples` script).

15.6 OpenMP

OpenMP is added to a code by using transformations. The three transformations currently supported allow the addition of an **OpenMP Parallel** directive, an **OpenMP Do** directive and an **OpenMP Parallel Do** directive, respectively, to a code.

The generic versions of these three transformations (i.e. ones that theoretically work for all APIs) were given in the *Standard Functionality* section. The API-specific versions of these transformations are described in the API-specific sections of this document.

15.6.1 Reductions

PSyclone supports parallel scalar reductions. If a scalar reduction is specified in the Kernel metadata (see the API-specific sections for details) then PSyclone ensures the appropriate reduction is performed.

In the case of distributed memory, PSyclone will add **GlobalSum's** at the appropriate locations. As can be inferred by the name, only “summation” reductions are currently supported for distributed memory.

In the case of an OpenMP parallel loop the standard reduction support will be used by default. For example

```
!$omp parallel do, reduction(+:x)
!loop
!$omp end parallel do
```

OpenMP reductions do not guarantee to give bit reproducible results for different runs of the same problem even if the same problem is run using the same resources. The reason for this is that the order in which data is reduced is not mandated.

Therefore, an additional **reprod** option has been added to the **OpenMP Do** transformation. If the **reprod** option is set to “True” then the OpenMP reduction support is replaced with local per-thread reductions which are reduced serially after the loop has finished. This implementation guarantees to give bit-wise reproducible results for different runs of the same problem using the same resources, but will not bit-wise compare if the code is rerun with different numbers of OpenMP threads.

15.6.2 Restrictions

If two reductions are used within an OpenMP region and the same variable is used for both reductions then PSyclone will raise an exception. In this case the solution is to use a different variable for each reduction.

PSyclone does not support (distributed-memory) halo swaps or global sums within OpenMP parallel regions. Attempting to create a parallel region for a set of nodes that includes halo swaps or global sums will produce an error. In such cases it may be possible to re-order the nodes in the Schedule using the *MoveTrans* transformation.

15.7 OpenCL

OpenCL is added to a code by using the `OCLTrans` transformation (see the *Standard Functionality* Section above). Currently this transformation is only supported for the GOcean1.0 API and is applied to the whole `InvokeSchedule` of an `Invoke`. This transformation will add an OpenCL driver infrastructure to the PSy layer and generate an OpenCL kernel for each of the `Invoke` kernels. This means that all kernels in that `Invoke` will be executed on the OpenCL device. The PSy-layer OpenCL code generated by PSyclone is still Fortran and makes use of the FortCL library (<https://github.com/stfc/FortCL>) to access OpenCL functionality. It also relies upon the device acceleration support provided by the `dl_esm_inf` library (https://github.com/stfc/dl_esm_inf).

Note: The generated OpenCL files follow the `-kernel-renaming` argument conventions, but in addition to the `<module-name>` they also include the `<kernel-name>` as part of the filename in the format: `<module-name>_<kernel-name>_index.cl`

The `OCLTrans` transformation accepts an `options` argument with a map of optional parameters to tune the OpenCL host code in the PSy layer. These options will be attached to the transformed `InvokeSchedule`. The current available options are:

Option	Description	Default
<code>end_barrier</code>	Whether a synchronization barrier should be placed at the end of the <code>Invoke</code> .	True
<code>enable_profiling</code>	Enables the profiling of OpenCL Kernels.	False
<code>out_of_order</code>	Allows the OpenCL implementation to execute the enqueued kernels out-of-order.	False

Additionally, each individual kernel (inside the `Invoke` that is going to be transformed) also accepts a map of options which are provided by the `set_opencl_options()` method of the `Kern` object. This can affect both the driver layer and/or the OpenCL kernels. The current available options are:

Option	Description	De- fault
<code>local_size</code>	Number of work-items to group together in a work-group execution (kernel instances executed at the same time).	64
<code>queue_number</code>	The identifier of the OpenCL Command Queue to which the kernel should be submitted.	1

Below is an example of a PSyclone script that uses an `OCLTrans` with multiple `InvokeSchedule` and kernel-specific optimization options.

```

1  from psyclone.transformations import OCLTrans
2
3  # Get the Schedule associated with the first Invoke
4  invoke = psy.invokes.invoke_list[0]
5  sched = invoke.schedule
6
7  # Transform the Schedule
8  cltrans = OCLTrans()
9  cltrans.apply(sched, options={"end_barrier": True})
10
11 # Provide kernel-specific OpenCL optimization options
12 for kern in sched.kernels():
13     kern.set_opencl_options({"queue_number": 1, 'local_size': 4})
14
15 return psy

```

OpenCL delays the decision of which and where kernels will execute until run-time, therefore it is important to use the environment variables provided by FortCL and DL_ESM_INF to inform how things should execute. Specifically:

- `FORTCL_KERNELS_FILE`: Point to the file containing the kernels to execute, they can be compiled ahead-of-time or providing the source for JIT compilation. To link more than a single kernel, one must merge all the kernels generated by PSyclone in a single source file.
- `FORTCL_PLATFORM`: If the system has more than 1 OpenCL platform. This environment variable may be used to select which platform on which to execute the kernels.
- `DL_ESM_ALIGNMENT`: When using OpenCL <= 1.2 the `local_size` should be exactly divisible by the total size. If this is not the case some implementations fail silently. A way to solve this issue is to set the

DL_ESM_ALIGNMENT variable to be equal to the local size.

Note: The OpenCL generation can be combined with distributed memory generation. In the case where there is more than one accelerator available on each node, the PSyclone configuration file parameter `OCL_DEVICES_PER_NODE` has to be set to the appropriate value and the number of MPI-ranks-per-node set by the *mpirun* command has to match this value accordingly.

For instance if there are 2 accelerators per nodes, *psyclone.cfg* should have `OCL_DEVICES_PER_NODE=2` and the program must be executed with `mpirun -n <total_ranks> -ppn 2 ./application` (Note: *-ppn* is an Intel MPI specific parameter, use equivalent configuration parameters for other MPI implementations.)

For example, an execution of a PSyclone generated OpenCL code using all the mentioned run-time configuration options could look something like:

```
FORTCL_PLATFORM=3 FORTCL_KERNELS_FILE=allkernels.cl DL_ESM_ALIGNMENT=64 \  
mpirun -n 2 ./application.exe
```

15.8 OpenACC

PSyclone supports the generation of code targetting GPUs through the addition of OpenACC directives. This is achieved by a user applying various OpenACC transformations to the PSyIR before the final Fortran code is generated. The steps to parallelisation are very similar to those in OpenMP with the added complexity of managing the movement of data to and from the GPU device. For the latter task PSyclone provides the `ACCDataTrans` and `ACCEnterDataTrans` transformations, as described in the *Standard Functionality* Section above. These two transformations add statically- and dynamically-scoped data regions, respectively. The former manages what data is on the remote device for a specific section of code while the latter allows run-time control of data movement. This second option is essential for minimising data movement as, without it, PSyclone-generated code would move data to and from the device upon every entry/exit of an `Invoke`. The first option is mainly provided as an aid to incremental porting and/or debugging of an OpenACC application as it provides explicit control over what data is present on a device for a given (part of an) `Invoke` routine.

The PGI compiler provides an alternative approach to controlling data movement through its ‘unified memory’ option (`-ta=tesla:managed`). When this is enabled the compiler itself takes on the task of ensuring that data is copied to/from the GPU when required. (Note that this approach can struggle with Fortran code containing derived types however.)

As well as ensuring the correct data is copied to and from the remote device, OpenACC directives must also be added to a code in order to tell the compiler how it should be parallelised. PSyclone provides the `ACCKernelsTrans`, `ACCParallelTrans` and `ACCLoopTrans` transformations for this purpose. The simplest of these is `ACCKernelsTrans` (currently only supported for the NEMO and Dynamo0.3 APIs) which encloses the code represented by a sub-tree of the PSyIR within an OpenACC `kernels` region. This essentially gives free-reign to the compiler to automatically parallelise any suitable loops within the specified region. An example of the use of `ACCDataTrans` and `ACCKernelsTrans` may be found in `PSyclone/examples/nemo/eg3` and an example of `ACCKernelsTrans` may be found in `PSyclone/examples/lfric/eg14`.

However, as with any “automatic” approach, a more performant solution can almost always be obtained by providing the compiler with more explicit direction on how to parallelise the code. The `ACCParallelTrans` and `ACCLoopTrans` transformations allow the user to define thread-parallel regions and, within those, define which loops should be parallelised. For an example of their use please see `PSyclone/examples/gocean/eg2` or `PSyclone/examples/lfric/eg14`.

In order for a given section of code to be executed on a GPU, any routines called from within that section must also have been compiled for the GPU. This then requires either that any such routines are in-lined or that the `OpenACC routine` directive be added to any such routines. This situation will occur routinely in those PSyclone APIs

that use the PSyKAl separation of concerns since the user-supplied kernel routines are called from within PSyclone-generated loops in the PSy layer. PSyclone therefore provides the `ACCRoutineTrans` transformation which, given a `Kernel` node in the PSyIR, creates a new version of that kernel with the `routine` directive added. Again, please see `PSyclone/examples/gocean/eg2` for an example. This transformation is currently not supported for kernels in the `Dynamo0.3` API.

15.9 SIR

It is currently not possible for PSyclone to output SIR code without using a script. Two examples of such scripts are given in example 4 for the NEMO API, one of which includes transformations to remove PSyIR intrinsics (as the SIR does not support them).

DISTRIBUTED MEMORY

PSyclone supports the generation of code for distributed memory machines. When this option is switched on, PSyclone takes on responsibility for both performance and correctness, as described below.

16.1 Correctness

PSyclone is responsible for adding appropriate distributed memory communication calls to the PSy layer to ensure that the distributed memory code runs correctly. For example, a stencil operation will require halo exchanges between the different processes.

The burden of correctly placing distributed memory communication calls has traditionally been born by the user. However, PSyclone is able to determine the placing of these within the PSy-layer, thereby freeing the user from this responsibility. Thus, the Algorithm and Kernel code remain the same, irrespective of whether the target architecture does or does not require a distributed memory solution.

16.2 Performance

PSyclone adds **HaloExchange** and **GlobalSum** objects to the generated PSyIR **InvokeSchedule** at the required locations. The halo-exchange and global-sum objects are exposed here for the purposes of optimisation. For example the halo-exchange and/or global-sum objects may be moved in the schedule (via appropriate transformations) to enable overlap of computation with communication.

Note: When these optimisations are implemented, add a reference to the *Transformations* Section.

A halo exchange is required with distributed memory when a processor requires data from its halo and the halo information is out of date. One example is where a field is written to and then read using a stencil access. Halo exchanges have performance implications so should only be used where necessary.

A global sum is required with distributed memory when a scalar is written to. Global sums can have performance implications so should only be used where necessary. Global sums currently only occur in certain Built-in kernels. The description of Built-ins indicates when this is the case.

16.3 Implementation

Within the contents of an `invoke()` call, PSyclone is able to statically determine which communication calls are required and where they should be placed. However, PSyclone has no information on what happens outside `invoke()` calls and thus is unable to statically determine whether communication is required between these calls. The solution we use is to add run-time flags in the PSy layer to keep track of whether data has been written to and read from. These flags are then used to determine whether communication calls are required upon entry to an `invoke()`.

16.4 Control

Support for distributed memory can be switched on or off with the default being on. The default can be changed permanently by modifying the `DISTRIBUTED_MEMORY` variable in the `psyclone.cfg` configuration file to `false` (see [Configuration](#)).

Distributed memory can be switched on or off from the `psyclone` script using the `-dm/--dist_mem` or `-nodm/--no_dist_mem` flags, respectively.

For interactive access, the distributed memory option can be changed interactively from the `PSyFactory` class by setting the optional `distributed_memory` flag; for example:

```
psy = PSyFactory(api=api, distributed_memory=False)
```

Similarly the distributed memory option can be changed interactively from the `generate` function by setting the optional `distributed_memory` flag, e.g.:

```
psy, alg = generate("file.f90", distributed_memory=False).
```

16.5 Status

Distributed memory support is currently supported by the `dynamo0.3` and the `gocean1p0` APIs. The remaining APIs ignore the distributed memory flag and continue to produce code without any distributed memory functionality, irrespective of its value.

STUB GENERATION

17.1 Quick Start

- 1) Use an existing Kernel file or create a Kernel file containing a Kernel module with the required metadata and an empty Kernel subroutine with no arguments.
- 2) Run the following command

```
> genkernelstub <PATH>/my_file.f90
```

17.2 Introduction

PSyclone provides a kernel stub generator for the Dynamo0.3 API. The kernel stub generator takes a kernel file as input and outputs the kernel subroutine arguments and declarations. The word “stub” is used to indicate that it is only the subroutine arguments and their declarations that are generated; the subroutine has no content.

The primary reason the stub generator is useful is that it generates the correct Kernel subroutine arguments and declarations for the Dynamo0.3 API as specified by the Kernel metadata. As the number of arguments to Kernel subroutines can become large and the arguments have to follow a particular order, it can become burdensome, and potentially error prone, for the user to have to work out the appropriate argument list if written by hand.

The stub generator can be used when creating a new Kernel. A Kernel can first be written to specify the required metadata and then the generator can be used to create the appropriate (empty) Kernel subroutine. The user can then fill in the content of the subroutine.

The stub generator can also be used to check whether the arguments for an existing Kernel are correct i.e. whether the Kernel subroutine and Kernel metadata are consistent. One example would be where a Kernel is updated resulting in a change to the metadata and subroutine arguments.

The Dynamo0.3 API requires Kernels to conform to a set of rules which determine the required arguments and types for a particular Kernel. These rules are required as the generated PSy layer needs to know exactly how to call a Kernel. These rules are outlined in Section [Rules](#).

Therefore PSyclone has been coded with the Dynamo0.3 API rules which are then applied when reading the Kernel metadata to produce the require Kernel call and its arguments in the generated PSy layer. These same rules are used by the Kernel stub generator to produce Kernel subroutine stubs, thereby guaranteeing that Kernel calls from the PSy layer and the associated Kernel subroutines are consistent.

17.3 Use

Before using the stub generator, PSyclone must be installed. If you have not already done so, please follow the instructions for setting up PSyclone in Section [Getting Going](#).

PSyclone will be installed in a particular location on your machine, which will be referred to as the <PSYCLONEINSTALL> directory. The stub generator script `genkernelstub` comes with the PSyclone installation. A quick check `> which genkernelstub` should return the location of the <PSYCLONEINSTALL>/bin directory.

The easiest way to use the stub generator is to run the supplied `genkernelstub` script.

```
> genkernelstub
usage: genkernelstub [-h] [-o OUTFILE] [-api API] [-l] filename
genkernelstub: error: the following arguments are required: filename
```

You can get information about the `genkernelstub` arguments using `-h` or `--help`:

```
> genkernelstub -h
usage: genkernelstub [-h] [-o OUTFILE] [-api API] [-l] filename

Create Kernel stub code from Kernel metadata

positional arguments:
  filename              Kernel metadata

optional arguments:
  -h, --help            show this help message and exit
  -o OUTFILE, --outfile OUTFILE
                        filename of output
  -api API              choose a particular api from ['dynamo0.3'], default
                        dynamo0.3
  -l, --limit           limit the fortran line length to 132 characters
```

As is indicated when using the `-h` option, the `-api` option only accepts `dynamo0.3` at the moment and is redundant as this option is also the default. However the number of supported APIs is expected to expand in the future.

The `-o`, or `--outfile` option allows the user to specify that the output should be written to a particular file. If `-o` is not specified then the python `print` statement is used. Typically the `print` statement results in the output being printed to the terminal.

The `-l`, or `--limit` option utilises the PSyclone support for wrapping of lines within the 132 character limit (please see the [Line Length](#) chapter for more details).

17.4 Kernels

Any Dynamo0.3 kernel can be used as input to the stub generator. Example Kernels can be found in the `examples/lfric` repository or, for more simple cases, in the `tests/test_files/dynamo0p3` directory. These directories are located in the <PSYCLONEHOME>/src/psyclone directory where <PSYCLONEHOME> refers to the location where you download or clone PSyclone ([Getting Going](#)).

In the `tests/test_files/dynamo0p3` directory the majority of examples start with `testkern`. Amongst the exceptions are: `simple.f90`, `ru_kernel_mod.f90` and `matrix_vector_kernel_mod.F90`. The following test kernels can be used to generate kernel stub code (running stub generation from the <PSYCLONEHOME>/src/psyclone directory):

```
tests/test_files/dynamo0p3/testkern_chi_read_mod.F90
tests/test_files/dynamo0p3/testkern_coord_w0_mod.F90
tests/test_files/dynamo0p3/testkern_operator_mod.f90
tests/test_files/dynamo0p3/testkern_operator_nofield_mod.f90
tests/test_files/dynamo0p3/testkern_orientation_mod.F90
tests/test_files/dynamo0p3/testkern_operator_orient_mod.f90
tests/test_files/dynamo0p3/ru_kernel_mod.f90
tests/test_files/dynamo0p3/simple.f90
```

17.5 Example

A simple single field example of a kernel that can be used as input for the stub generator is found in `tests/test_files/dynamo0p3/simple.f90` and is shown below:

```
module simple_mod

  use argument_mod
  use fs_continuity_mod
  use kernel_mod
  use constants_mod

  implicit none

  type, extends(kernel_type) :: simple_type
    type(arg_type), dimension(1) :: meta_args = &
      (/ arg_type(gh_field, gh_real, gh_inc, wl) /)
    integer :: operates_on = cell_column
  contains
    procedure, nopass :: code => simple_code
  end type simple_type

contains

  subroutine simple_code()
  end subroutine

end module simple_mod
```

Note: The module name `simple_mod` and the type name `simple_type` share the same root `simple` and have the extensions `_mod` and `_type` respectively. This is a convention in Dynamo0.3 API and is required by the kernel stub generator as it needs to determine the name of the type containing the metadata and infers this by reading the module name. If this rule is not followed the kernel stub generator will return with an error message (see Section [Errors](#)).

Note: Whilst strictly the kernel stub generator only requires the Kernel metadata to generate the appropriate stub code, the parser that the generator relies on currently requires a dummy kernel subroutine to exist.

If we run the kernel stub generator on the `simple.f90` example:

```
> genkernelstub tests/test_files/dynamo0p3/simple.f90
```

we get the following kernel stub output:

```
MODULE simple_mod
  IMPLICIT NONE
  CONTAINS
  SUBROUTINE simple_code(nlayers, field_1_w1, ndf_w1, undf_w1, map_w1)
    USE constants_mod, ONLY: r_def, i_def
    IMPLICIT NONE
    INTEGER(KIND=i_def), intent(in) :: nlayers
    INTEGER(KIND=i_def), intent(in) :: ndf_w1
    INTEGER(KIND=i_def), intent(in), dimension(ndf_w1) :: map_w1
    INTEGER(KIND=i_def), intent(in) :: undf_w1
    REAL(KIND=r_def), intent(inout), dimension(undf_w1) :: field_1_w1
  END SUBROUTINE simple_code
END MODULE simple_mod
```

The subroutine content can then be copied into the required module, used as the basis for a new module, or checked with an existing subroutine for correctness.

Note: The output does not currently conform to Met Office coding standards so must be modified accordingly.

Note: The code will not compile without a) providing the `constants_mod`, `argument_mod` and `kernel_mod` modules in the compiler include path and b) adding in code that writes to any arguments declared as `intent out` or `inout`. For a quick check, the `USE` declaration and `KIND` declarations can be removed and the `field_1_w1` array can be initialised with some value in the subroutine. At this point the Kernel should compile successfully.

Note: Whilst there is only one field declared in the metadata there are 5 arguments to the Kernel. The first argument `nlayers` specifies the number of layers in a column for a field. The second argument is the array associated with the field. The field array is dimensioned as the *number of unique degrees of freedom* (hereafter `undf`) which is also passed into the kernel (the fourth argument). The naming convention is to call each field a `field`, followed by its position in the (algorithm) argument list (which is reflected in the metadata ordering). The third argument is the number of degrees of freedom for the particular column and is used to dimension the final argument which is the *degrees of freedom map* (`dofmap`) which indicates the location of the required values in the field array. The naming convention for the `dofmap`, `undf` and `ndf` is to append the name with the space that it is associated with.

We now take a look at a more complicated example. The metadata in this example is the same as an actual Dynamo0.3 kernel, however the subroutine content and various comments have been removed. The metadata specifies that there are four fields passed by the algorithm layer, the fourth of which is a vector field of size three. All three of the spaces require a basis function and the `W0` and `W2` function spaces additionally require a differential basis function. The content of the Kernel, excluding the subroutine body, is given below.

```
module ru_kernel_mod

  use argument_mod
  use fs_continuity_mod
  use kernel_mod
  use constants_mod

  implicit none

  private
```

(continues on next page)

(continued from previous page)

```

type, public, extends(kernel_type) :: ru_kernel_type
  private
  type(arg_type) :: meta_args(6) = (/\
    arg_type(GH_FIELD, GH_REAL, GH_INC, W2),
    arg_type(GH_FIELD, GH_REAL, GH_READ, W3),
    arg_type(GH_SCALAR, GH_INTEGER, GH_READ),
    arg_type(GH_SCALAR, GH_REAL, GH_READ),
    arg_type(GH_FIELD, GH_REAL, GH_READ, W0),
    arg_type(GH_FIELD*3, GH_REAL, GH_READ, W0)
  /\)
  type(func_type) :: meta_funcs(3) = (/\
    func_type(W2, GH_BASIS, GH_DIFF_BASIS),
    func_type(W3, GH_BASIS),
    func_type(W0, GH_BASIS, GH_DIFF_BASIS)
  /\)
  integer :: operates_on = CELL_COLUMN
  integer :: gh_shape = gh_quadrature_XYoZ
contains
  procedure, nopass :: ru_code
end type

public ru_code

contains

  subroutine ru_code()
  end subroutine ru_code

end module ru_kernel_mod

```

If we run the kernel stub generator on this example:

```
> genkernelstub tests/test_files/dynamo0p3/ru_kernel_mod.f90
```

we obtain the following output:

```

MODULE ru_mod
  IMPLICIT NONE
  CONTAINS
  SUBROUTINE ru_code(nlayers, field_1_w2, field_2_w3, iscalar_3, rscalar_4, &
    field_5_w0, field_6_w0_v1, field_6_w0_v2, field_6_w0_v3, &
    ndf_w2, undf_w2, map_w2, basis_w2_qr_xyoz, &
    diff_basis_w2_qr_xyoz, ndf_w3, undf_w3, map_w3, &
    basis_w3_qr_xyoz, ndf_w0, undf_w0, map_w0, &
    basis_w0_qr_xyoz, diff_basis_w0_qr_xyoz, &
    np_xy_qr_xyoz, np_z_qr_xyoz, weights_xy_qr_xyoz, weights_z_qr_
↪xyoz)
    USE constants_mod, ONLY: r_def, i_def
    IMPLICIT NONE
    INTEGER(KIND=i_def), intent(in) :: nlayers
    INTEGER(KIND=i_def), intent(in) :: ndf_w0
    INTEGER(KIND=i_def), intent(in), dimension(ndf_w0) :: map_w0
    INTEGER(KIND=i_def), intent(in) :: ndf_w2
    INTEGER(KIND=i_def), intent(in), dimension(ndf_w2) :: map_w2
    INTEGER(KIND=i_def), intent(in) :: ndf_w3
    INTEGER(KIND=i_def), intent(in), dimension(ndf_w3) :: map_w3
    INTEGER(KIND=i_def), intent(in) :: undf_w2, undf_w3, undf_w0

```

(continues on next page)

(continued from previous page)

```

REAL(KIND=r_def), intent(in) :: rscalar_4
INTEGER(KIND=i_def), intent(in) :: iscalar_3
REAL(KIND=r_def), intent(inout), dimension(undf_w2) :: field_1_w2
REAL(KIND=r_def), intent(in), dimension(undf_w3) :: field_2_w3
REAL(KIND=r_def), intent(in), dimension(undf_w0) :: field_5_w0
REAL(KIND=r_def), intent(in), dimension(undf_w0) :: field_6_w0_v1
REAL(KIND=r_def), intent(in), dimension(undf_w0) :: field_6_w0_v2
REAL(KIND=r_def), intent(in), dimension(undf_w0) :: field_6_w0_v3
INTEGER(KIND=i_def), intent(in) :: np_xy_qr_xyoz, np_z_qr_xyoz
REAL(KIND=r_def), intent(in), dimension(3,ndf_w2,np_xy_qr_xyoz,np_z_qr_xyoz) :: ↪
↪basis_w2_qr_xyoz
REAL(KIND=r_def), intent(in), dimension(1,ndf_w2,np_xy_qr_xyoz,np_z_qr_xyoz) :: ↪
↪diff_basis_w2_qr_xyoz
REAL(KIND=r_def), intent(in), dimension(1,ndf_w3,np_xy_qr_xyoz,np_z_qr_xyoz) :: ↪
↪basis_w3_qr_xyoz
REAL(KIND=r_def), intent(in), dimension(1,ndf_w0,np_xy_qr_xyoz,np_z_qr_xyoz) :: ↪
↪basis_w0_qr_xyoz
REAL(KIND=r_def), intent(in), dimension(3,ndf_w0,np_xy_qr_xyoz,np_z_qr_xyoz) :: ↪
↪diff_basis_w0_qr_xyoz
REAL(KIND=r_def), intent(in), dimension(np_xy_qr_xyoz) :: weights_xy_qr_xyoz
REAL(KIND=r_def), intent(in), dimension(np_z_qr_xyoz) :: weights_z_qr_xyoz
END SUBROUTINE ru_code
END MODULE ru_mod

```

The above example demonstrates that the argument list can get quite complex. Rather than going through an explanation of each argument you are referred to Section [Rules](#) for more details on the rules for argument types and argument ordering. Regarding naming conventions for arguments you can see that the arrays associated with the fields are labelled as 1-6 depending on their position in the metadata. For a vector field, each vector results in a different array. These are distinguished by appending `_vx` where `x` is the number of the vector.

The introduction of stencil operations on field arguments further complicates the argument list of a kernel. An example of the use of the stub generator for a kernel that performs stencil operations is provided in `examples/lfric/eg5`.

```
> genkernelstub ../../examples/lfric/eg5/conservative_flux_kernel_mod.F90
```

17.6 Errors

The stub generator has been written to provide useful errors if mistakes are found. If you run the generator and it does not produce a useful error - and in particular if it produces a stack trace - please contact the PSyclone developers.

The following tests do not produce stub kernel code either because they are invalid or because they contain functionality that is not supported in the stub generator.

```

tests/test_files/dynamo0p3/testkern_any_space_1_mod.f90
tests/test_files/dynamo0p3/testkern_any_space_4_mod.f90
tests/test_files/dynamo0p3/testkern_any_discontinuous_space_op_2_mod.f90
tests/test_files/dynamo0p3/testkern_dofs_mod.f90
tests/test_files/dynamo0p3/testkern_invalid_fortran.F90
tests/test_files/dynamo0p3/testkern_short_name.F90
tests/test_files/dynamo0p3/testkern_no_datatype.F90
tests/test_files/dynamo0p3/testkern_qr.F90

```

`testkern_invalid_fortran.F90`, `testkern_no_datatype.F90`, `testkern_short_name.F90` and `testkern_qr.F90` are designed to be invalid for PSyclone stub generation testing purposes and should produce appropriate errors. Two examples are below:

```
> genkernelstub tests/test_files/dynamo0p3/testkern_invalid_fortran.F90
Error: 'Parse Error: Code appears to be invalid Fortran'

> genkernelstub tests/test_files/dynamo0p3/testkern_no_datatype.F90
Error: 'Parse Error: Kernel type testkern_type does not exist'
```

testkern_dofs_mod.f90 is an example with an unsupported feature, as the `operates_on` metadata specifies `dof`. Currently only kernels with `operates_on=CELL_COLUMN` are supported by the stub generator.

Generic function space metadata `any_space` and `any_discontinuous_space` (see Section *Supported Function Spaces* for function-space identifiers) are currently only supported for *Dynamo0.3 fields* in the stub generator. Basis and differential basis functions on these generic function spaces, required for *quadrature* and *evaluators*, are not supported. Hence, `testkern_any_space_1_mod.f90`, `testkern_any_space_4_mod.f90` and `testkern_any_discontinuous_space_op_2_mod.f90` should fail with appropriate warnings because of that. For example:

```
> genkernelstub tests/test_files/dynamo0p3/testkern_any_space_1_mod.f90
Error: "Generation Error: Unsupported space for basis function, expecting
one of ['w3', 'wtheta', 'w2v', 'w2vtrace', 'w2broken', 'w0', 'w1', 'w2',
'w2trace', 'w2h', 'w2htrace', 'any_w2', 'wchi'] but found 'any_space_1'"
```

As noted above, if the Dynamo0.3 API naming convention for module and type names is not followed, the stub generator will return with an error message. For example:

```
> genkernelstub tests/test_files/dynamo0p3/testkern_qr.F90
Error: "Parse Error: Error, module name 'testkern_qr' does not have
'_mod' as an extension. This convention is assumed."
```


LINE LENGTH

By default PSyclone will generate Fortran code with no consideration of Fortran line-length limits. As the line-length limit for free-form Fortran is 132 characters, the code that is output may be non-conformant.

Line length is not an issue for many compilers as they allow compiler flags to be set which allow lines longer than the Fortran standard. However this is not the case for all compilers.

PSyclone therefore supports the wrapping of lines within the 132 character limit. The next two sections discuss how this is done when scripting and when working interactively respectively.

18.1 Script

The *psyclone* script provides the `-l` option to wrap lines. Please see the *Fortran line length* section for more details.

18.2 Interactive

When using PSyclone interactively the line lengths of the input algorithm and Kernel files can be checked by setting the `psyclone.parse.algorithm.parse()` function's `line_length` argument to `True`.

```
>>> from psyclone.parse.algorithm import parse
>>> ast, info = parse("argspec.F90", line_length=True)
```

Similarly the `line_length` argument can be set to `True` if calling the `generator.generate()` function. This function simply passes this argument on to the `psyclone.parse.algorithm.parse()` function.

```
>>> from psyclone.generator import generate
>>> alg, psy = generate("argspec.F90", line_length=True)
```

Line wrapping is performed as a post-processing step, i.e. after the code has been generated. This is done by an instance of the `line_length.FortLineLength` class. For example:

```
>>> from psyclone.generator import generate
>>> from psyclone.line_length import FortLineLength
>>> psy, alg = generate("algspec.f90", line_length=True)
>>> line_length = FortLineLength()
>>> psy_str = line_length.process(str(psy))
>>> print psy_str
>>> alg_str = line_length.process(str(alg))
>>> print alg_str
```

18.3 Limitations

The `line_length.FortLineLength` class is only partially aware of Fortran syntax. This awareness is required so that appropriate continuation characters can be used (for example `&` at the end of a line and `!$omp&` at the start of a line for OpenMP directives, `&` at the end of a line for statements and `&` at the end of a line and `&` at the beginning of a line for strings).

Whilst statements only require an `&` at the end of the line when line wrapping with free-form fortran they may optionally also have an `&` at the beginning of the subsequent line. In contrast, when splitting a string over multiple lines an `&` is required at both locations. Therefore an instance of the `line_length.FortLineLength` class will always add `&` at the beginning of a continuation line for a statement, in case the line is split within a string.

One known situation that could cause an instance of the `line_length.FortLineLength` class to fail is when an inline comment is used at the end of a line to make it longer than the 132 character limit. Whilst PSyclone does not generate such code for the PSy-layer, this might occur in Algorithm-layer code, even if the Algorithm-layer code conforms to the 132 line length limit. The reason for this is that PSyclone's internal parser concatenates lines together, thus a long line correctly split with continuation characters in the Algorithm-layer becomes a line that needs to be split by an instance of the `line_length.FortLineLength` class.

FORTTRAN NAMING CONVENTIONS

There is a convention in the kernel code for the Dynamo0.3 and GOcean1.0 APIs that if the name of the operation being performed is `<name>` then a kernel file is `<name>_mod.[fF90]`, the name of the module inside the kernel file is `<name>_mod`, the name of the kernel metadata in the module is `<name>_type` and the name of the kernel subroutine in the module is `<name>_code`.

PSyclone itself does not rely on this convention apart from in the stub generator (see the *Stub Generation* Section) where the name of the metadata to be parsed is determined from the module name.

20.1 The generator module

This module provides the PSyclone ‘main’ routine which is intended to be driven from the bin/psyclone executable script. ‘main’ takes an algorithm file as input and produces modified algorithm code and generated PSy code. A function, ‘generate’, is also provided which has the same functionality as ‘main’ but can be called from within another Python program.

```
psyclone.generator.generate(filename, api="", kernel_path="", script_name=None,  
                             line_length=False, distributed_memory=None, kern_out_path="",  
                             kern_naming='multiple')
```

Takes a PSyclone algorithm specification as input and outputs the associated generated algorithm and psy codes suitable for compiling with the specified kernel(s) and support infrastructure. Uses the `parse.algorithm.parse()` function to parse the algorithm specification, the `psyGen.PSy` class to generate the PSy code and the `alg_gen.Alg` class to generate the modified algorithm code.

Parameters

- **filename** (*str*) – The file containing the algorithm specification.
- **kernel_path** (*str*) – The directory from which to recursively search for the files containing the kernel source (if different from the location of the algorithm specification).
- **script_name** (*str*) – A script file that can apply optimisations to the PSy layer (can be a path to a file or a filename that relies on the PYTHONPATH to find the module).
- **line_length** (*bool*) – A logical flag specifying whether we care about line lengths being longer than 132 characters. If so, the input (algorithm and kernel) code is checked to make sure that it conforms. The default is False.
- **distributed_memory** (*bool*) – A logical flag specifying whether to generate distributed memory code. The default is set in the config.py file.
- **kern_out_path** (*str*) – Directory to which to write transformed kernel code.
- **kern_naming** (*bool*) – the scheme to use when re-naming transformed kernels.

Returns 2-tuple containing fparser1 ASTs for the algorithm code and the psy code.

Return type (`fparser.one.block_statements.BeginSource`, `fparser.one.block_statements.Module`)

Raises

- **IOError** – if the filename or search path do not exist
- **GenerationError** – if an invalid API is specified.
- **GenerationError** – if an invalid kernel-renaming scheme is specified.

For example:

```
>>> from psyclone.generator import generate
>>> alg, psy = generate("algspec.f90")
>>> alg, psy = generate("algspec.f90", kernel_path="src/kernels")
>>> alg, psy = generate("algspec.f90", script_name="optimise.py")
>>> alg, psy = generate("algspec.f90", line_length=True)
>>> alg, psy = generate("algspec.f90", distributed_memory=False)
```

20.2 The parse module

Module that uses the Fortran parser `fparser2` to parse PSyclone-conformant Algorithm code.

```
psyclone.parse.algorithm.parse(alg_filename, api="", invoke_name='invoke', kernel_path="",
                                line_length=False)
```

Takes a PSyclone conformant algorithm file as input and outputs a parse tree of the code contained therein and an object containing information about the ‘invoke’ calls in the algorithm file and any associated kernels within the invoke calls.

Parameters

- **alg_filename** (*str*) – The file containing the algorithm specification.
- **api** (*str*) – The PSyclone API to use when parsing the code.
- **invoke_name** (*str*) – The expected name of the invocation calls in the algorithm code.
- **kernel_path** (*str*) – The path to search for kernel source files (if different from the location of the algorithm source).
- **line_length** (*bool*) – A logical flag specifying whether we care about line lengths being longer than 132 characters. If so, the input (algorithm and kernel) code is checked to make sure that it conforms and an error raised if not. The default is False.

Returns 2-tuple consisting of the `fparser2` parse tree of the Algorithm file and an object holding details of the invokes found.

Return type (`fparser.two.Fortran2003.Program`, `psyclone.parse.FileInfo`)

For example:

```
>>> from psyclone.parse.algorithm import parse
>>> ast, info = parse("alg.f90")
```

20.3 The transformations module

This module provides the various transformations that can be applied to PSyIR nodes. There are both general and API-specific transformation classes in this module where the latter typically apply API-specific checks before calling the base class for the actual transformation.

class `psyclone.transformations.ACCTrans`

Add an OpenACC data region around a list of nodes in the PSyIR. COPYIN, COPYOUT and COPY clauses are added as required.

For example:

```

>>> from psyclone.parse import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "NEMO"
>>> filename = "tra_adv.F90"
>>> ast, invokeInfo = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.transformations import ACCDataTrans
>>> dtrans = ACCDataTrans()
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> kernels = schedule.children[0].children[0].children[0:-1]
>>> # Enclose the kernels
>>> new_sched, _ = dtrans.apply(kernels)

```

apply (*nodes*, *options=None*)

Put the supplied node or list of nodes within an OpenACC data region.

Parameters

- **nodes** ((list of) `psyclone.psyir.nodes.Node`) – the PSyIR node(s) to enclose in the data region.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Returns (transformed schedule, memento of transformation)

Return type 2-tuple of (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`).

property name

Returns the name of this transformation.

Return type `str`

validate (*nodes*, *options*)

Check that we can safely add a data region around the supplied list of nodes.

Parameters

- **nodes** ((list of) subclasses of `psyclone.psyir.nodes.Node`) – the proposed node(s) to enclose in a data region.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises

- **TransformationError** – if the Schedule to which the nodes belong already has an ‘enter data’ directive.
- **TransformationError** – if any of the nodes are themselves data directives.

class `psyclone.transformations.ACCEnterDataTrans`

Adds an OpenACC “enter data” directive to a Schedule. For example:

```

>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "gocean1.0"

```

(continues on next page)

(continued from previous page)

```

>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> dtrans = t.get_trans_name('ACCEnterDataTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> # Add an enter-data directive
>>> newschedule, _ = dtrans.apply(schedule)
>>> newschedule.view()

```

apply (*sched, options=None*)

Adds an OpenACC “enter data” directive to the invoke associated with the supplied Schedule. Any fields accessed by OpenACC kernels within this schedule will be added to this data region in order to ensure they remain on the target device.

Parameters

- **sched** (sub-class of `psyclone.psyir.nodes.Schedule`) – schedule to which to add an “enter data” directive.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Returns tuple of the modified schedule and a record of the transformation.

Return type (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation.

Return type `str`

validate (*sched, options=None*)

Check that we can safely apply the OpenACC enter-data transformation to the supplied Schedule.

Parameters

- **sched** (sub-class of `psyclone.psyir.nodes.Schedule`) – Schedule to which to add an “enter data” directive.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises

- **NotImplementedError** – for any API other than GOcean 1.0 or NEMO.
- **TransformationError** – if passed something that is not a (subclass of) `psyclone.psyir.nodes.Schedule`.

class `psyclone.transformations.ACCKernelsTrans`

Enclose a sub-set of nodes from a Schedule within an OpenACC kernels region (i.e. within “!\$acc kernels” ... “!\$acc end kernels” directives).

For example:


```

>>> from psyclone.parse import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "NEMO"
>>> filename = "tra_adv.F90"
>>> ast, invokeInfo = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.transformations import ACCKernelsTrans
>>> ktrans = ACCKernelsTrans()
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> kernels = schedule.children[0].children[0].children[0:-1]
>>> # Transform the kernel
>>> new_sched, _ = ktrans.apply(kernels)

```

apply (*nodes*, *options=None*)

Enclose the supplied list of PSyIR nodes within an OpenACC Kernels region.

Parameters

- **nodes** ((list of) `psyclone.psyir.nodes.Node`) – a node or list of nodes in the PSyIR to enclose.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["default_present"]** (*bool*) – whether or not the kernels region should have the ‘default present’ attribute (indicating that data is already on the accelerator). When using managed memory this option should be False.

Returns (transformed schedule, memento of transformation)

Return type 2-tuple of (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`).

property name

Returns the name of this transformation class.

Return type `str`

validate (*nodes*, *options*)

Check that we can safely enclose the supplied node or list of nodes within OpenACC kernels ... end kernels directives.

Parameters

- **nodes** ((list of) `psyclone.psyir.nodes.Node`) – the proposed PSyIR node or nodes to enclose in the kernels region.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises

- **NotImplementedError** – if the supplied Nodes do not belong to a NemoInvokeSchedule.
- **TransformationError** – if there are no Loops within the proposed region.

class `psyclone.transformations.ACCLoopTrans`

Adds an OpenACC loop directive to a loop. This directive must be within the scope of some OpenACC Parallel region (at code-generation time).

For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory
>>> from psyclone.errors import GenerationError
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ltrans = t.get_trans_name('ACCLoopTrans')
>>> rtrans = t.get_trans_name('ACCParallelTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> new_schedule = schedule
>>>
# Apply the OpenACC Loop transformation to *every* loop
# in the schedule
>>> for child in schedule.children:
>>>     newschedule, memento = ltrans.apply(child, reprod=True)
>>>     schedule = newschedule
>>>
# Enclose all of these loops within a single OpenACC
# PARALLEL region
>>> rtrans.omp_schedule("dynamic,1")
>>> newschedule, memento = rtrans.apply(schedule.children)
>>>
```

apply (*node*, *options=None*)

Apply the ACCLoop transformation to the specified node. This node must be a Loop since this transformation corresponds to inserting a directive immediately before a loop, e.g.:

```
!$ACC LOOP
do ...
...
end do
```

At code-generation time (when `psyclone.psyGen.ACCLoopDirective.gen_code()` is called), this node must be within (i.e. a child of) a PARALLEL region.

Parameters

- **node** (`psyclone.psyir.nodes.Loop`) – the supplied node to which we will apply the Loop transformation.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["collapse"]** (*int*) – number of nested loops to collapse.
- **options["independent"]** (*bool*) – whether to add the “independent” clause to the directive (not strictly necessary within PARALLEL regions).

Returns 2-tuple of new schedule and memento of transform

Return type (psyclone.dynamo0p3.DynInvokeSchedule, psyclone.undoredo.Memento)

excluded_node_types

alias of psyclone.psyir.nodes.psy_data_node.PSyDataNode

property name

Returns the name of this transformation as a string.

class psyclone.transformations.**ACCParallelTrans**

Create an OpenACC parallel region by inserting directives. This parallel region *must* come after an enter-data directive (see *ACCEnterDataTrans*) or within a data region (see *ACCDDataTrans*). For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ptrans = t.get_trans_name('ACCParallelTrans')
>>> dtrans = t.get_trans_name('ACCDDataTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> # Enclose everything within a single OpenACC PARALLEL region
>>> newschedule, _ = ptrans.apply(schedule.children)
>>> # Add an enter-data directive
>>> newschedule, _ = dtrans.apply(newschedule)
>>> newschedule.view()
```

property name

Returns the name of this transformation as a string.

Return type str

class psyclone.transformations.**ACCRoutineTrans**

Transform a kernel subroutine by adding a “!\$acc routine” directive (causing it to be compiled for the OpenACC accelerator device). For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.transformations import ACCRoutineTrans
>>> rtrans = ACCRoutineTrans()
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> kern = schedule.children[0].children[0].children[0]
```

(continues on next page)

(continued from previous page)

```
>>> # Transform the kernel
>>> newkern, _ = rtrans.apply(kern)
```

apply (*kern*, *options=None*)

Modifies the AST of the supplied kernel so that it contains an ‘!\$acc routine’ OpenACC directive. This transformation affects the f2pygen and the PSyIR trees of this kernel.

Parameters

- **kern** (`psyclone.psyGen.Kern`) – the kernel object to transform.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises TransformationError – if we fail to find the subroutine corresponding to the kernel object.

Returns (transformed kernel, memento of transformation)

Return type 2-tuple of (`psyclone.psyGen.Kern`, `psyclone.undoredo.Memento`).

property name

Returns the name of this transformation class.

Return type `str`

validate (*kern*, *options=None*)

Perform checks that the supplied kernel can be transformed.

Parameters

- **kern** (`psyclone.psyGen.Kern`) – the kernel which is the target of the transformation.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises

- **TransformationError** – if the target kernel is a built-in.
- **TransformationError** – if any of the symbols in the kernel are accessed via a module use statement.

class `psyclone.transformations.ColourTrans`

Apply a colouring transformation to a loop (in order to permit a subsequent parallelisation over colours). For example:

```
>>> invoke = ...
>>> schedule = invoke.schedule
>>>
>>> ctrans = ColourTrans()
>>>
>>> # Colour all of the loops
>>> for child in schedule.children:
>>>     cschedule, _ = ctrans.apply(child)
>>>
>>> csched.view()
```

apply (*node*, *options=None*)

Converts the Loop represented by *node* into a nested loop where the outer loop is over colours and the inner loop is over cells of that colour.

Parameters

- **node** (`psyclone.psyir.nodes.Loop`) – the loop to transform.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Returns Tuple of modified schedule and record of transformation

Return type (`psyclone.psyir.nodes.Schedule`, `:py:class:`psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

class `psyclone.transformations.Dynamo0p3AsyncHaloExchangeTrans`

Splits a synchronous halo exchange into a halo exchange start and halo exchange end. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "dynamo0.3"
>>> ast, invokeInfo = parse("file.f90", api=api)
>>> psy=PSyFactory(api).create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import Dynamo0p3AsyncHaloExchangeTrans
>>> trans = Dynamo0p3AsyncHaloExchangeTrans()
>>> new_schedule, memento = trans.apply(schedule.children[0])
>>> new_schedule.view()
```

apply (*node*, *options=None*)

Transforms a synchronous halo exchange, represented by a `HaloExchange` node, into an asynchronous halo exchange, represented by `HaloExchangeStart` and `HaloExchangeEnd` nodes.

Parameters

- **node** (`psyclone.psygen.HaloExchange`) – a synchronous haloexchange node.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Returns tuple of the modified schedule and a record of the transformation.

Return type (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

Return type `str`

validate (*node*, *options*)

Internal method to check whether the node is valid for this transformation.

Parameters

- **node** (`psyclone.psygen.HaloExchange`) – a synchronous Halo Exchange node

- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises TransformationError – if the node argument is not a HaloExchange (or subclass thereof)

class `psyclone.transformations.Dynamo0p3ColourTrans`

Split a Dynamo 0.3 loop over cells into colours so that it can be parallelised. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> import transformations
>>> import os
>>> import pytest
>>>
>>> TEST_API = "dynamo0.3"
>>> _, info=parse(os.path.join(os.path.dirname(os.path.abspath(__file__)),
>>>                             "tests", "test_files", "dynamo0p3",
>>>                             "4.6_multikernel_invokes.f90"),
>>>               api=TEST_API)
>>> psy = PSyFactory(TEST_API).create(info)
>>> invoke = psy.invokes.get('invoke_0')
>>> schedule = invoke.schedule
>>>
>>> ctrans = Dynamo0p3ColourTrans()
>>> otrans = DynamoOMPParallelLoopTrans()
>>>
>>> # Colour all of the loops
>>> for child in schedule.children:
>>>     cschedule, _ = ctrans.apply(child)
>>>
>>> # Then apply OpenMP to each of the colour loops
>>> schedule = cschedule
>>> for child in schedule.children:
>>>     newsched, _ = otrans.apply(child.children[0])
>>>
>>> newsched.view()
```

Colouring in the Dynamo 0.3 API is subject to the following rules:

- Only kernels with an iteration space of CELLS and which modify a continuous field require colouring. Any other type of loop will cause this transformation to raise an exception.
- A kernel may have at most one field with ‘INC’ access
- A separate colour map will be required for each field that is coloured (if an invoke contains >1 kernel call)

apply (*node, options=None*)

Performs Dynamo0.3-specific error checking and then uses the parent class to convert the Loop represented by *node* into a nested loop where the outer loop is over colours and the inner loop is over cells of that colour.

Parameters

- **node** (`psyclone.dynamo0p3.DynLoop`) – the loop to transform.
- **options** – a dictionary with options for transformations. :type options: dictionary of string:values or None

Returns 2-tuple of new schedule and memento of transform.

Return type (psyclone.dynamo0p3.DynInvokeSchedule, psyclone.undoredo.Memento)

property name

Returns the name of this transformation as a string.

class psyclone.transformations.Dynamo0p3KernelConstTrans

Modifies a kernel so that the number of dofs, number of layers and number of quadrature points are fixed in the kernel rather than being passed in by argument.

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "dynamo0.3"
>>> ast, invokeInfo = parse("file.f90", api=api)
>>> psy=PSyFactory(api).create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import Dynamo0p3KernelConstTrans
>>> trans = Dynamo0p3KernelConstTrans()
>>> for kernel in schedule.coded_kernels():
>>>     new_schedule, _ = trans.apply(kernel, number_of_layers=150)
>>>     kernel_schedule = kernel.get_kernel_schedule()
>>>     kernel_schedule.symbol_table.view()
```

apply (node, options=None)

Transforms a kernel so that the values for the number of degrees of freedom (if a valid value for the element_order arg is provided), the number of quadrature points (if the quadrature arg is set to True) and the number of layers (if a valid value for the number_of_layers arg is provided) are constant in a kernel rather than being passed in by argument.

The “cellshape”, “element_order” and “number_of_layers” arguments are provided to mirror the namelist values that are input into an LFRic model when it is run.

Quadrature support is currently limited to XYoZ in this transformation. In the case of XYoZ the number of quadrature points (for horizontal and vertical) are set to the element_order + 3 in the LFRic infrastructure so their value is derived.

Parameters

- **node** (psyclone.psygen.DynKern) – a kernel node.
- **options** (dictionary of string:values or None) – a dictionary with options for transformations.
- **options["cellshape"]** (str) – the shape of the cells. This is provided as it helps determine the number of dofs a field has for a particular function space. Currently only “quadrilateral” is supported which is also the default value.
- **options["element_order"]** (int) – the order of the cell. In combination with cellshape, this determines the number of dofs a field has for a particular function space. If it is set to None (the default) then the dofs values are not set as constants in the kernel, otherwise they are.
- **options["number_of_layers"]** (int) – the number of vertical layers in the LFRic model mesh used for this particular run. If this is set to None (the default) then the nlayers value is not set as a constant in the kernel, otherwise it is.
- **options["quadrature"]** (bool) – whether the number of quadrature points values are set as constants in the kernel (True) or not (False). The default is False.

Returns tuple of the modified schedule and a record of the transformation.

Return type (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

Return type `str`

validate (*node*, *options=None*)

This method checks whether the input arguments are valid for this transformation.

Parameters

- **node** (`psyclone.psygen.DynKern`) – a dynamo 0.3 kernel node.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["cellshape"]** (*str*) – the shape of the elements/cells.
- **options["element_order"]** (*int*) – the order of the elements/cells.
- **options["number_of_layers"]** (*int*) – the number of layers to use.
- **options["quadrature"]** (*bool*) – whether quadrature dimension sizes should or shouldn't be set as constants in a kernel.

Raises TransformationError – if the node argument is not a dynamo 0.3 kernel, the cell-shape argument is not set to “quadrilateral”, the element_order argument is not a 0 or a positive integer, the number of layers argument is not a positive integer, the quadrature argument is not a boolean, neither element order nor number of layers arguments are set (as the transformation would then do nothing), or the quadrature argument is True but the element order is not provided (as the former needs the latter).

class `psyclone.transformations.Dynamo0p3OMPLoopTrans` (*omp_schedule='static'*)

Dynamo 0.3 specific orphan OpenMP loop transformation. Adds Dynamo-specific validity checks. Actual transformation is done by *base class*.

apply (*node*, *options=None*)

Perform Dynamo 0.3 specific loop validity checks then call `OMPLoopTrans.apply()`.

Parameters

- **node** (`psyclone.psyir.nodes.Node`) – the Node in the Schedule to check
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations and validation.
- **options["reprod"]** (*bool*) – indicating whether reproducible reductions should be used. By default the value from the config file will be used.

Raises TransformationError – if an OMP loop transform would create incorrect code.

property name

Returns the name of this transformation as a string.

class `psyclone.transformations.Dynamo0p3RedundantComputationTrans`

This transformation allows the user to modify a loop's bounds so that redundant computation will be performed. Redundant computation can result in halo exchanges being modified, new halo exchanges being added or existing halo exchanges being removed.

- This transformation should be performed before any parallelisation transformations (e.g. for OpenMP) to the loop in question and will raise an exception if this is not the case.
- This transformation can not be applied to a loop containing a reduction and will again raise an exception if this is the case.
- This transformation can only be used to add redundant computation to a loop, not to remove it.
- This transformation allows a loop that is already performing redundant computation to be modified, but only if the depth is increased.

apply (*loop*, *options=None*)

Apply the redundant computation transformation to the loop *loop*. This transformation can be applied to loops iterating over ‘cells’ or ‘dofs’. If *depth* is set to a value then the value will be the depth of the field’s halo over which redundant computation will be performed. If *depth* is not set to a value then redundant computation will be performed to the full depth of the field’s halo.

Parameters

- **loop** (`psyclone.psyGen.DynLoop`) – the loop that we are transforming.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["depth"]** (*int*) – the depth of the stencil. Defaults to None.

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

validate (*node*, *options=None*)

Perform various checks to ensure that it is valid to apply the RedundantComputation transformation to the supplied node

Parameters

- **node** (`psyclone.psyir.nodes.Node`) – the supplied node on which we are performing validity checks
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["depth"]** (*int*) – the depth of the stencil if the value is provided and None if not.

Raises

- **TransformationError** – if the node is not a `psyclone.psyir.nodes.Loop`.
- **TransformationError** – if the parent of the loop is a `psyclone.psyGen.Directive`.
- **TransformationError** – if the parent of the loop is not a `psyclone.psyir.nodes.Loop` or a `psyclone.psyGen.DynInvokeSchedule`.
- **TransformationError** – if the parent of the loop is a `psyclone.psyir.nodes.Loop` but the original loop does not iterate over ‘colour’.
- **TransformationError** – if the parent of the loop is a `psyclone.psyir.nodes.Loop` but the parent does not iterate over ‘colours’.

- **TransformationError** – if the parent of the loop is a `psyclone.psyir.nodes.Loop` but the parent's parent is not a `psyclone.psyGen.DynInvokeSchedule`.
- **TransformationError** – if this transformation is applied when distributed memory is not switched on.
- **TransformationError** – if the loop does not iterate over cells, dofs or colour.
- **TransformationError** – if the transformation is setting the loop to the maximum halo depth but the loop already computes to the maximum halo depth.
- **TransformationError** – if the transformation is setting the loop to the maximum halo depth but the loop contains a stencil access (as this would result in the field being accessed beyond the halo depth).
- **TransformationError** – if the supplied depth value is not an integer.
- **TransformationError** – if the supplied depth value is less than 1.
- **TransformationError** – if the supplied depth value is not greater than 1 when a continuous loop is modified as this is the minimum valid value.
- **TransformationError** – if the supplied depth value is not greater than the existing depth value, as we should not need to undo existing transformations.
- **TransformationError** – if a depth value has been supplied but the loop has already been set to the maximum halo depth.

class `psyclone.transformations.DynamoLoopFuseTrans` (*same_space=False*)

Dynamo0.3 API specialisation of the *base class* in order to fuse two Dynamo loops after performing validity checks. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>>
>>> API = "dynamo0.3"
>>> FILENAME = "alg.x90"
>>> ast, invokeInfo = parse(FILENAME, api=API)
>>> psy = PSyFactory(API, distributed_memory=False).create(invoke_info)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>>
>>> from psyclone.transformations import DynamoLoopFuseTrans
>>> ftrans = DynamoLoopFuseTrans()
>>>
>>> new_schedule, memento = ftrans.apply(schedule[0], schedule[1])
>>> new_schedule.view()
```

The optional argument *same_space* can be set as

```
>>> ftrans.same_space = True
```

after the instance of the transformation is created.

apply (*node1*, *node2*, *options=None*)

Fuses two `psyclone.dynamo0p3.DynLoop` loops after performing validity checks by calling `LoopFuseTrans.apply()` method of the base class.

Parameters

- **node1** (`psyclone.dynamo0p3.DynLoop`) – the first Loop to fuse.
- **node2** (`psyclone.dynamo0p3.DynLoop`) – the second Loop to fuse.

- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Returns two-tuple of the modified Schedule and a record of the transformation.

Return type (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

property same_space

Returns the *same_space* flag that is specified when applying this transformation. The default value is *False*. This optional flag, set to *True*, asserts that an unknown iteration space (i.e. *ANY_SPACE*) matches the other iteration space. This is set at the user's own risk. If both iteration spaces are discontinuous the loops can be fused without having to use the *same_space* flag.

validate (*node1*, *node2*, *options=None*)

Performs various checks to ensure that it is valid to apply the `DynamoLoopFuseTrans` transformation to the supplied loops.

Parameters

- **node1** (`psyclone.dynamo0p3.DynLoop`) – the first Loop to fuse.
- **node2** (`psyclone.dynamo0p3.DynLoop`) – the second Loop to fuse.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises

- **TransformationError** – if either of the supplied loops contains an inter-grid kernel.
- **TransformationError** – if one or both function spaces have invalid names.
- **TransformationError** – if the *same_space* flag was set, but does not apply because neither field is on *ANY_SPACE* or the spaces are not the same.
- **TransformationError** – if one or more of the iteration spaces is unknown (*ANY_SPACE*) and the *same_space* flag is not set to *True*.
- **TransformationError** – if the loops are over different spaces that are not both discontinuous and the loops both iterate over cells.
- **TransformationError** – if the loops' upper bound names are not the same.
- **TransformationError** – if the halo-depth indices of two loops are not the same.
- **TransformationError** – if each loop already contains a reduction.
- **TransformationError** – if the first loop has a reduction and the second loop reads the result of the reduction.

class `psyclone.transformations.DynamoOMPParallelLoopTrans` (*omp_schedule='static'*)

Dynamo-specific OpenMP loop transformation. Adds Dynamo specific validity checks. Actual transformation is done by the *base class*.

apply (*node*, *options=None*)

Perform Dynamo specific loop validity checks then call the *apply()* method of the *base class*.

Parameters

- **node** (`psyclone.psyir.nodes.Node`) – the Node in the Schedule to check

- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises `TransformationError` – if the associated loop requires colouring.

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

class `psyclone.transformations.GOConstLoopBoundsTrans`

Switch on (or off) the use of constant loop bounds within a `GOInvokeSchedule`. In the absence of constant loop bounds, PSyclone will generate loops where the bounds are obtained by de-referencing a field object, e.g.:

```
DO j = my_field%grid%internal%ystart, my_field%grid%internal%ystop
```

Some compilers are able to produce more efficient code if they are provided with information on the relative trip-counts of the loops within an `Invoke`. With constant loop bounds switched on, PSyclone generates code like:

```
ny = my_field%grid%subdomain%internal%ystop
...
DO j = 1, ny-1
```

In practice, the application of the constant loop bounds looks something like, e.g.:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> import os
>>> TEST_API = "gocean1.0"
>>> _, info = parse(os.path.join("tests", "test_files", "gocean1p0",
>>>                               "single_invoke.f90"),
>>>                  api=TEST_API)
>>> psy = PSyFactory(TEST_API).create(info)
>>> invoke = psy.invokes.get('invoke_0_compute_cu')
>>> schedule = invoke.schedule
>>>
>>> from psyclone.transformations import GOConstLoopBoundsTrans
>>> clbtrans = GOConstLoopBoundsTrans()
>>>
>>> newsched, _ = clbtrans.apply(schedule)
>>> # or, to turn off const. loop bounds:
>>> # newsched, _ = clbtrans.apply(schedule, const_bounds=False)
>>>
>>> newsched.view()
```

apply (*node, options=None*)

Switches constant loop bounds on or off for all loops in a `GOInvokeSchedule`. Default is 'off'.

Parameters

- **node** (`psyclone.gocean1p0.GOInvokeSchedule`) – the `GOInvokeSchedule` of which all loops will get the constant loop bounds switched on or off.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

- **options["const_bounds"]** (*bool*) – whether the constant loop should be used (True) or not (False). Default is True.

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.gocean1p0.GOInvokeSchedule`, `psyclone.undoredo.Memento`)

property name

Return the name of the Transformation as a string.

class `psyclone.transformations.GOLoopSwapTrans`

Provides a loop-swap transformation, e.g.:

```
DO j=1, m
  DO i=1, n
```

becomes:

```
DO i=1, n
  DO j=1, m
```

This transform is used as follows:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast, invokeInfo = parse("shallow_alg.f90")
>>> psy = PSyFactory("gocean1.0").create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import GOLoopSwapTrans
>>> swap = GOLoopSwapTrans()
>>> new_schedule, memento = swap.apply(schedule.children[0])
>>> new_schedule.view()
```

apply (*outer*, *options=None*)

The argument *outer* must be a loop which has exactly one inner loop. This transform then swaps the outer and inner loop.

Parameters

- **outer** (`psyclone.psyir.nodes.Loop`) – the node representing the outer loop.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises TransformationError – if the supplied node does not allow a loop swap to be done.

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

validate (*node_outer*, *options=None*)

Checks if the given node contains a valid Fortran structure to allow swapping loops. This means the node must represent a loop, and it must have exactly one child that is also a loop.

Parameters

- **node_outer** (py:class:psyclone.psyir.nodes.Loop) – a Loop node from an AST.
- **options** (dictionary of string:values or None) – a dictionary with options for transformations.

Raises TransformationError – if the supplied node does not allow a loop swap to be done.

class psyclone.transformations.GOceanLoopFuseTrans

GOcean API specialisation of the *base class* in order to fuse two GOcean loops after performing validity checks (e.g. that the loops are over the same grid-point type). For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast, invokeInfo = parse("shallow_alg.f90")
>>> psy = PSyFactory("gocean1.0").create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import GOceanLoopFuseTrans
>>> ftrans = GOceanLoopFuseTrans()
>>> new_schedule, memento = ftrans.apply(schedule[0], schedule[1])
>>> new_schedule.view()
```

apply (node1, node2, options=None)

Fuses two *psyclone.gocean1p0.GOLoop* loops after performing validity checks by calling *LoopFuseTrans.apply()* method of the base class.

Parameters

- **node1** (psyclone.gocean1p0.GOLoop) – the first Node representing a GOLoop.
- **node2** (psyclone.gocean1p0.GOLoop) – the second Node representing a GOLoop.
- **options** (dictionary of string:values or None) – a dictionary with options for transformations.

Returns two-tuple of the modified Schedule and a record of the transformation.

Return type (psyclone.psyir.nodes.Schedule, psyclone.undoredo.Memento)

Raises

- **TransformationError** – if the supplied loops are over different grid-point types.
- **TransformationError** – if there is an unexpected exception.

property name

Returns the name of this transformation as a string.

validate (node1, node2, options=None)

Checks if it is valid to apply the GOceanLoopFuseTrans transform. It ensures that the fused loops are over the same grid-point types, before calling the normal LoopFuseTrans validation function.

Parameters

- **node1** (psyclone.gocean1p0.GOLoop) – the first Node representing a GOLoop.
- **node2** (psyclone.gocean1p0.GOLoop) – the second Node representing a GOLoop.

- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises

- **TransformationError** – if the supplied loops are over different grid-point types.
- **TransformationError** – if invalid parameters are passed in.

class `psyclone.transformations.GOceanOMPLoopTrans` (*omp_schedule='static'*)

GOcean-specific orphan OpenMP loop transformation. Adds GOcean specific validity checks (that the node is either an inner or outer Loop). Actual transformation is done by *base class*.

Parameters `omp_schedule` – the omp schedule to be created. Must be one of ‘runtime’, ‘static’, ‘dynamic’, ‘guided’ or ‘auto’.

apply (*node, options=None*)

Perform GOcean specific loop validity checks then call `OMPLoopTrans.apply()`.

Parameters

- **node** (`psyclone.psyir.nodes.Loop`) – the loop to parallelise using OMP Do.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

property name

Returns the name of this transformation as a string.

class `psyclone.transformations.GOceanOMPParallelLoopTrans` (*omp_schedule='static'*)

GOcean specific OpenMP Do loop transformation. Adds GOcean specific validity checks (that supplied Loop is an inner or outer loop). Actual transformation is done by *base class*.

Parameters `omp_schedule` – the omp schedule to be created. Must be one of ‘runtime’, ‘static’, ‘dynamic’, ‘guided’ or ‘auto’.

apply (*node, options=None*)

Perform GOcean-specific loop validity checks then call `OMPParallelLoopTrans.apply()`.

Parameters

- **node** (`psyclone.psyir.nodes.Loop`) – a Loop node from an AST.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations and validation.

Raises **TransformationError** – if the supplied node is not an inner or outer loop.

Returns 2-tuple of new schedule and memento of transform

Return type (`psyclone.dynamo0p3.GOInvokeSchedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

class `psyclone.transformations.KernelGlobalsToArguments`

Transformation that removes any accesses of global data from the supplied kernel and places them in the caller. The values/references are then passed by argument into the kernel.

apply (*node, options=None*)

Convert the global variables used inside the kernel into arguments and modify the InvokeSchedule to pass the same global variables to the kernel call.

This `apply()` method does not return anything, as agreed in #595. However, this change has yet to be applied to the other Transformation classes.

Parameters

- **node** (`psyclone.psyGen.CodedKern`) – a kernel call.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

property name

Returns the name of this transformation.

Return type `str`

validate (*node, options=None*)

Check that the supplied node is a valid target for this transformation.

Parameters

- **node** (`psyclone.psyGen.CodedKern`) – the PSyIR node to validate.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises

- **TransformationError** – if the supplied node is not a `CodedKern`.
- **TransformationError** – if this transformation is not applied to a Gocean API Invoke.
- **TransformationError** – if the supplied kernel contains wildcard imports of symbols from one or more containers (e.g. a USE without an ONLY clause in Fortran).

class `psyclone.transformations.KernelModuleInlineTrans`

Switches on, or switches off, the inlining of a Kernel subroutine into the PSy layer module. For example:

```
>>> invoke = ...
>>> schedule = invoke.schedule
>>>
>>> inline_trans = KernelModuleInlineTrans()
>>>
>>> ischedule, _ = inline_trans.apply(schedule.children[0].loop_body[0])
>>> ischedule.view()
```

Warning: For this transformation to work correctly, the Kernel subroutine must only use data that is passed in by argument, declared locally or included via use association within the subroutine. Two examples where in-lining will not work are:

1. A variable is declared within the module that contains the Kernel subroutine and is then accessed within that Kernel;
2. A variable is included via use association at the module level and accessed within the Kernel subroutine.

The transformation will reject attempts to in-line such kernels.

apply (*node, options=None*)

Checks that the node is of the correct type (a Kernel) then marks the Kernel to be inlined, or not, depending on the value of the inline option. If the inline option is not passed the Kernel is marked to be inlined.

Parameters

- **node** (`psyclone.psyir.nodes.Loop`) – the loop to transform.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["inline"]** (*bool*) – whether the kernel should be module inlined or not.

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

class `psyclone.transformations.KernelTrans`

Base class for all Kernel transformations.

static validate (*kern, options=None*)

Checks that the supplied node is a Kernel and that it is possible to construct the PSyIR of its contents.

Parameters

- **kern** (`psyclone.psyGen.Kern` or sub-class) – the kernel which is the target of the transformation.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises

- **TransformationError** – if the target node is not a sub-class of `psyGen.Kern`.
- **TransformationError** – if the subroutine containing the implementation of the kernel cannot be found in the `fparser2` Parse Tree.
- **TransformationError** – if the PSyIR cannot be constructed because there are symbols of unknown type.

class `psyclone.transformations.LoopFuseTrans`

Provides a generic loop-fuse transformation to two Nodes in the PSyIR of a Schedule after performing validity checks for the supplied Nodes. Examples are given in the descriptions of any children classes.

apply (*node1, node2, options=None*)

Fuses two loops represented by `psyclone.psyir.nodes.Node` objects after performing validity checks.

Parameters

- **node1** (`psyclone.psyir.nodes.Node`) – the first Node that is being checked.
- **node2** (`psyclone.psyir.nodes.Node`) – the second Node that is being checked.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Returns two-tuple of the modified Schedule and a record of the transformation.

Return type (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`).

property name

Returns the name of this transformation as a string.

validate (*node1*, *node2*, *options=None*)

Performs various checks to ensure that it is valid to apply the LoopFuseTrans transformation to the supplied Nodes.

Parameters

- **node1** (`psyclone.psyir.nodes.Node`) – the first Node that is being checked.
- **node2** (`psyclone.psyir.nodes.Node`) – the second Node that is being checked.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises

- **TransformationError** – if one or both of the Nodes is/are not a `psyclone.psyir.nodes.Loop`.
- **TransformationError** – if one or both Nodes are not fully-formed.
- **TransformationError** – if the Nodes do not have the same parent.
- **TransformationError** – if the Nodes are not next to each other in the tree.
- **TransformationError** – if the two Loops do not have the same iteration space.

class `psyclone.transformations.MoveTrans`

Provides a transformation to move a node in the tree. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast, invokeInfo=parse("dynamo.F90")
>>> psy=PSyFactory("dynamo0.3").create(invokeInfo)
>>> schedule=psy.invokes.get('invoke_v3_kernel_type').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import MoveTrans
>>> trans=MoveTrans()
>>> new_schedule, memento = trans.apply(schedule.children[0],
>>>                                     schedule.children[2],
>>>                                     position="after")
>>> new_schedule.view()
```

Nodes may only be moved to a new location with the same parent and must not break any dependencies otherwise an exception is raised.

apply (*node*, *location*, *options=None*)

Move the node represented by *node* before *location* *location* (which is also a node) by default and after if the optional *position* argument is set to ‘after’.

Parameters

- **node** (`psyclone.psyir.nodes.Node`) – the node to be moved.
- **location** (`psyclone.psyir.nodes.Node`) – node before or after which the given node should be moved.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["position"]** (*str*) – either ‘before’ or ‘after’.

Raises

- **TransformationError** – if the given node is not an instance of `psyclone.psyir.nodes.Node`
- **TransformationError** – if the location is not valid.

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

validate (*node*, *location*, *options=None*)

validity checks for input arguments.

Parameters

- **node** (`psyclone.psyir.nodes.Node`) – the node to be moved.
- **location** (`psyclone.psyir.nodes.Node`) – node before or after which the given node should be moved.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["position"]** (*str*) – either ‘before’ or ‘after’.

Raises

- **TransformationError** – if the given node is not an instance of `psyclone.psyir.nodes.Node`
- **TransformationError** – if the location is not valid.

class `psyclone.transformations.OCLTrans`

Switches on/off the generation of an OpenCL PSy layer for a given `InvokeSchedule`. Additionally, it will generate OpenCL kernels for each of the kernels referenced by the `Invoke`. For example:

```
>>> invoke = ...
>>> schedule = invoke.schedule
>>>
>>> ocl_trans = OCLTrans()
>>> new_sched, _ = ocl_trans.apply(schedule)
```

apply (*sched*, *options=None*)

Apply the OpenCL transformation to the supplied `GOInvokeSchedule`. This causes PSyclone to generate an OpenCL version of the corresponding PSy-layer routine. The generated code makes use of the FortCL library (<https://github.com/stfc/FortCL>) in order to manage the OpenCL device directly from Fortran.

Parameters

- **sched** (`psyclone.psyGen.GOInvokeSchedule`) – the `InvokeSchedule` to transform.
- **options** (*dictionary of string:values or None*) – set of option to tune the OpenCL generation.
- **options["opencl"]** (*bool*) – whether or not to enable OpenCL generation.

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation.

Return type `str`

validate (*sched*, *options=None*)

Checks that the supplied InvokeSchedule is valid and that an OpenCL version of it can be generated.

Parameters

- **sched** (`psyclone.psyGen.InvokeSchedule`) – the Schedule to check.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises

- **TransformationError** – if the InvokeSchedule is not for the GOcean1.0 API.
- **NotImplementedError** – if any of the kernels have arguments passed by value.

class `psyclone.transformations.OMPLoopTrans` (*omp_schedule='static'*)

Adds an orphaned OpenMP directive to a loop. i.e. the directive must be inside the scope of some other OMP Parallel REGION. This condition is tested at code-generation time. The optional ‘reprod’ argument in the apply method decides whether standard OpenMP reduction support is to be used (which is not reproducible) or whether a manual reproducible reproduction is to be used.

Parameters **omp_schedule** (*str*) – the OpenMP schedule to use.

For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory
>>> from psyclone.errors import GenerationError
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>> print psy.invokes.names
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ltrans = t.get_trans_name('OMPLoopTrans')
>>> rtrans = t.get_trans_name('OMPParallelTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> new_schedule = schedule
>>>
# Apply the OpenMP Loop transformation to *every* loop
# in the schedule
>>> for child in schedule.children:
>>>     newschedule, memento = ltrans.apply(child, reprod=True)
>>>     schedule = newschedule
>>>
# Enclose all of these loops within a single OpenMP
# PARALLEL region
>>> rtrans.omp_schedule("dynamic,1")
>>> newschedule, memento = rtrans.apply(schedule.children)
```

(continues on next page)

(continued from previous page)

```
>>>
>>>
```

apply (*node*, *options=None*)

Apply the OMPLoopTrans transformation to the specified node in a Schedule. This node must be a Loop since this transformation corresponds to wrapping the generated code with directives like so:

```
!$OMP DO
do ...
...
end do
!$OMP END DO
```

At code-generation time (when `OMPLoopDirective.gen_code()` is called), this node must be within (i.e. a child of) an OpenMP PARALLEL region.

If the keyword “reprod” is specified in the options, it will cause a reproducible reduction to be generated if it is set to True, otherwise the default value (as read from the `psyclone.cfg` file) will be used. Note, reproducible in this case means obtaining the same results with the same number of OpenMP threads, not for different numbers of OpenMP threads.

Parameters

- **node** (`psyclone.psyir.nodes.Node`) – the supplied node to which we will apply the OMPLoopTrans transformation
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations and validation.
- **options["reprod"]** (*bool*) – indicating whether reproducible reductions should be used. By default the value from the config file will be used.

Returns (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

property omp_schedule

Returns the OpenMP schedule that will be specified by this transformation. The default schedule is ‘static’.

class `psyclone.transformations.OMPParallelLoopTrans` (*omp_schedule='static'*)

Adds an OpenMP PARALLEL DO directive to a loop.

For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast, invokeInfo = parse("dynamo.F90")
>>> psy = PSyFactory("dynamo0.1").create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_v3_kernel_type').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import OMPParallelLoopTrans
>>> trans = OMPParallelLoopTrans()
>>> new_schedule, memento = trans.apply(schedule.children[0])
>>> new_schedule.view()
```

apply (*node*, *options=None*)

Apply an OMPParallelLoop Transformation to the supplied node (which must be a Loop). In the generated code this corresponds to wrapping the Loop with directives:

```
!$OMP PARALLEL DO ...
do ...
...
end do
!$OMP END PARALLEL DO
```

Parameters

- **node** (`psyclone.f2pygen.DoGen`) – the node (loop) to which to apply the transformation.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations and validation.

Returns two-tuple of transformed schedule and a record of the transformation.

Return type (`psyclone.psyir.nodes.Schedule`, `:py:class:`psyclone.undoredo.Memento`)

property name

Returns the name of this transformation as a string.

validate (*node, options=None*)

Validity checks for input arguments.

Parameters

- **node** (`psyclone.psyir.nodes.Node`) – the PSyIR node to validate.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.

Raises

- **TransformationError** – if the nodes is not a Loop.
- **TransformationError** – if the nodes is over colours.

class `psyclone.transformations.OMPParallelTrans`

Create an OpenMP PARALLEL region by inserting directives. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory
>>> from psyclone.errors import GenerationError
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ltrans = t.get_trans_name('GOceanOMPLoopTrans')
>>> rtrans = t.get_trans_name('OMPParallelTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> new_schedule = schedule
>>>
>>> # Apply the OpenMP Loop transformation to *every* loop
```

(continues on next page)

(continued from previous page)

```

>>> # in the schedule
>>> for child in schedule.children:
>>>     newschedule, memento = ltrans.apply(child)
>>>     schedule = newschedule
>>>
>>> # Enclose all of these loops within a single OpenMP
>>> # PARALLEL region
>>> newschedule, _ = rtrans.apply(schedule.children)
>>> newschedule.view()

```

property name

Returns the name of this transformation as a string.

Return type `str`

validate (*node_list*, *options=None*)

Perform OpenMP-specific validation checks.

Parameters

- **node_list** (list of `psyclone.psyir.nodes.Node`) – list of Nodes to put within parallel region.
- **options** (dictionary of *string:values* or *None*) – a dictionary with options for transformations.
- **options["node-type-check"]** (*bool*) – this flag controls if the type of the nodes enclosed in the region should be tested to avoid using unsupported nodes inside a region.

Raises **TransformationError** – if the target Nodes are already within some OMP parallel region.

class `psyclone.transformations.ParallelLoopTrans`

Adds an orphaned directive to a loop indicating that it should be parallelised. i.e. the directive must be inside the scope of some other Parallel REGION. This condition is tested at code-generation time.

apply (*node*, *options=None*)

Apply the Loop transformation to the specified node in a Schedule. This node must be a Loop since this transformation corresponds to wrapping the generated code with directives, e.g. for OpenMP:

```

!$OMP DO
do ...
...
end do
!$OMP END DO

```

At code-generation time (when `gen_code()` is called), this node must be within (i.e. a child of) a PARALLEL region.

Parameters

- **node** (`psyclone.psyir.nodes.Node`) – the supplied node to which we will apply the Loop transformation.
- **options** – a dictionary with options for transformations. :type options: dictionary of *string:values* or *None*
- **options["collapse"]** (*int*) – the number of loops to collapse into single iteration space or *None*.

Returns (`psyclone.psyir.nodes.Schedule`, `psyclone.undoredo.Memento`)

abstract property name

Returns the name of this transformation as a string.

Return type `str`

validate (*node*, *options=None*)

Perform validation checks before applying the transformation

Parameters

- **node** (`psyclone.psyir.nodes.Node`) – the node we are checking.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations. This transform supports “collapse”, which is the number of nested loops to collapse.
- **options["collapse"]** (*int*) – number of nested loops to collapse or None.

Raises

- **TransformationError** – if the node is not a `psyclone.psyir.nodes.Loop`.
- **TransformationError** – if the `psyclone.psyir.nodes.Loop` loop iterates over colours.
- **TransformationError** – if the target loop contains one of the node types specified in `self.excluded_node_types`.
- **TransformationError** – if ‘collapse’ is supplied with an invalid number of loops.

class `psyclone.transformations.ParallelRegionTrans`

Base class for transformations that create a parallel region.

apply (*target_nodes*, *options=None*)

Apply this transformation to a subset of the nodes within a schedule - i.e. enclose the specified Loops in the schedule within a single parallel region.

Parameters

- **target_nodes** ((list of) `psyclone.psyir.nodes.Node`) – a single Node or a list of Nodes.
- **options** (*dictionary of string:values or None*) – a dictionary with options for transformations.
- **options["node-type-check"]** (*bool*) – this flag controls if the type of the nodes enclosed in the region should be tested to avoid using unsupported nodes inside a region.

Returns 2-tuple of new schedule and memento of transform.

Return type (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

abstract property name

Returns the name of this transformation as a string.

validate (*node_list*, *options=None*)

Check that the supplied list of Nodes are eligible to be put inside a parallel region.

Parameters

- **node_list** (*list*) – list of nodes to put into a parallel region
- **options** – a dictionary with options for transformations. :type options: dictionary of string:values or None

- **options["node-type-check"]** (*bool*) – this flag controls whether or not the type of the nodes enclosed in the region should be tested to avoid using unsupported nodes inside a region.

Raises

- **TransformationError** – if the supplied node is an InvokeSchedule rather than being within an InvokeSchedule.
- **TransformationError** – if the supplied nodes are not all children of the same parent (siblings).

`psyclone.transformations.check_intergrid(node)`

Utility function to check that the supplied node does not have an intergrid kernel amongst its descendants.

This is used ensure any attempt to apply loop-fusion and redundant- computation transformations to loops containing inter-grid kernels is rejected (since support for those is not yet implemented).

Parameters `node` (`psyr.nodes.Node`) – the PSyIR node to check.

Raises **TransformationError** – if the supplied node has an inter-grid kernel as a descendant.

20.4 The psyGen module

This module provides generic support for PSyclone’s PSy code optimisation and generation. The classes in this method need to be specialised for a particular API and implementation.

class `psyclone.psyGen.PSy` (*invoke_info*)

Base class to help manage and generate PSy code for a single algorithm file. Takes the invocation information output from the function `parse.algorithm.parse()` as its input and stores this in a way suitable for optimisation and code generation.

Parameters `invoke_info` (`psyclone.parse.algorithm.FileInfo`) – An object containing the required invocation information for code optimisation and generation. Produced by the function `parse.algorithm.parse()`.

For example:

```
>>> from psyclone.parse.algorithm import parse
>>> ast, info = parse("argspec.F90")
>>> from psyclone.psyGen import PSyFactory
>>> api = "... "
>>> psy = PSyFactory(api).create(info)
>>> print(psy.gen)
```

property container

Returns the container associated with this PSy object

Return type `psyclone.psyir.nodes.Container`

abstract property gen

Abstract base class for code generation function.

Returns root node of generated Fortran AST.

Return type `psyclone.psyir.nodes.Node`

property invokes

Returns the list of invokes.

Return type `psyclone.psyGen.Invoke` or derived class

property name

Returns the name of the PSy object.

Return type `str`

20.5 The `alg_gen` module

This module provides the `Alg` class and supporting exception-handling to translate the original algorithm file into one that can be compiled and linked with the generated PSy code.

class `psyclone.alg_gen.Alg` (*parse_tree*, *psy*, *invoke_name*='invoke')

Generate a modified algorithm code for a single algorithm specification. Takes the parse tree of the algorithm specification output from the function `psyclone.parse.algorithm.parse()` and an instance of the `psyGen.PSy` class as input. The latter allows consistent names to be generated between the algorithm (calling) and psy (callee) layers.

For example:

```
>>> from psyclone.algorithm.parse import parse
>>> parse_tree, info = parse("argspec.F90")
>>> from psyclone.psyGen import PSy
>>> psy = PSy(info)
>>> from psyclone.alg_gen import Alg
>>> alg = Alg(parse_tree, psy)
>>> print(alg.gen)
```

Parameters

- **parse_tree** (`fparser.two.utils.Base`) – an object containing a parse tree of the algorithm specification which was produced by the function `psyclone.parse.algorithm.parse()`. Assumes the algorithm will be parsed by `fparser2` and expects a valid program unit, program, module, subroutine or function.
- **psy** (`psyclone.psyGen.PSy`) – an object containing information about the PSy layer.
- **invoke_name** (`str`) – the name that the algorithm layer uses to indicate an invoke call. This is an optional argument that defaults to the name “invoke”.

property `gen`

Return modified algorithm code.

Returns the modified algorithm specification as an `fparser2` parse tree.

Return type `fparser.two.utils.Base`

20.6 The `line_length` module

Provides support for breaking long fortran lines into smaller ones to allow the code to conform to the maximum line length limits (132 for f90 free format is the default)

class `psyclone.line_length.FortLineLength` (*line_length=132*)

This class take a free format fortran code as a string and line wraps any lines that are larger than the specified line length

property `length`

returns the maximum allowed line length

long_lines (*fortran_in*)

returns true if at least one of the lines in the input code is longer than the allowed length. Otherwise returns false

process (*fortran_in*)

takes fortran code as a string as input and output fortran code as a string with any long lines wrapped appropriately

PSYDATA API

PSyclone provides transformations that will insert callbacks to an external library at runtime. These callbacks allow third-party libraries to access data structures at specified locations in the code. Some example use cases are:

Profiling: By inserting callbacks before and after a region of code, performance measurements can be added. PSyclone provides wrapper libraries for some common performance profiling tools, see [Profiling](#) for details.

Kernel Data Extraction: PSyclone provides the ability to add callbacks that provide access to all input variables before, and output variables after a kernel invocation. This can be used to automatically create tests for a kernel, or to write a stand-alone driver that just calls one kernel, which can be used for performance tuning. An example library that extracts input- and output-data into a netcdf file is included with PSyclone (see [NetCDF Extraction Example](#)).

Access Verification: The callbacks can be used to make sure a field declared as read-only is not modified during a kernel call (either because of an incorrect declaration, or because memory is overwritten). The implementation included in PSyclone uses a simple 64-bit checksum to detect changes to a field (and scalar values). See [Read-Only Verification](#) for details.

NAN Test: The callbacks can be used to make sure that all floating point input and output parameters of a kernel are not a NAN (not-a-number) or infinite. See [NAN Test](#) for the full description.

In-situ Visualisation: By giving access to output fields of a kernel, an in-situ visualisation library can be used to plot fields while a (PSyclone-processed) application is running. There is no example library available at this stage, but the API has been designed with this application in mind.

The PsyData API should be general enough to allow these and other applications to be developed and used.

PSyclone provides transformations that will insert callbacks to the PSyData API, for example `ProfileTrans`, `GOceanExtractTrans` and `LFRicExtractTrans`. A user can develop additional transformations and corresponding runtime libraries for additional functionality. Refer to [PSyData API](#) for full details about the PSyData API.

21.1 Read-Only Verification

The PSyData interface is being used to verify that read-only variables in a kernel are not overwritten. The `ReadOnlyVerifyTrans` (in `psyir.transformations.read_only_verify_trans`, or the [Transformation Reference Guide](#)) uses the dependency analysis to determine all read-only variables (i.e. arguments declared to be read-only in metadata, most implicit arguments in LFRic, grid properties in GOcean). A simple 64-bit checksum is then computed for all these arguments before a kernel call, and compared with the checksum after the kernel call. Any change in the checksum causes a message to be printed at runtime, e.g.:

```
-----  
Double precision field b_fld has been modified in main : update  
Original checksum: 4611686018427387904
```

(continues on next page)

(continued from previous page)

```
New checksum:      4638355772470722560
-----
```

The transformation that adds read-only-verification to an application can be applied for both the LFRic and GOcean API - no API-specific transformations are required. Below is an example that searches for each loop in an invoke (which will always surround kernel calls) and applies the transformation to each one. This code has been successfully used as a global transformation with the LFRic gravity-wave miniapp:

```
def trans(psy):
    from psyclone.psyir.transformations import ReadOnlyVerifyTrans
    from psyclone.psyir.nodes import Loop
    read_only_verify = ReadOnlyVerifyTrans()

    for invoke in psy.invokes.invoke_list:
        schedule = invoke.schedule
        for node in schedule:
            if isinstance(node, Loop):
                schedule, _ = read_only_verify.apply(node)

    return psy
```

Besides the transformation, a library is required to do the actual verification at runtime. There are two implementations of the read-only-verification library included in PSyclone: one for LFRic, and one for GOcean. Both libraries support the environment variable `PSYDATA_VERBOSE`. This can be used to control how much output is generated by the read-only-verification library at runtime. If the variable is not specified or has the value '0', warnings will only be printed if checksums change. If it is set to '1', a message will be printed before and after each kernel call that is checked. If the variable is set to '2', it will additionally print the name of each variable that is checked.

21.1.1 Read-Only Verification Library for LFRic

This library is contained in `lib/read_only/lfric` and it must be compiled before compiling any LFRic-based application that uses read-only verification. Compiling this library requires access to the LFRic infrastructure library (since it must implement a generic interface for e.g. the LFRic `Field` class). The makefile uses the variable `LFRIC_DIR` to point to the location where LFRic's `field_mod` has been compiled. It defaults to `../../../../../../../../lfric/work/trunk/miniapps/gravity_wave/working/field` but this will certainly need to be changed for any user. The LFRic infrastructure library is not used in linking the verification library. The application which uses the read-only-verification library needs to link in the infrastructure library anyway.

Compilation of the library is done by invoking `make` and setting the required variables:

```
make LFRIC_DIR=some_path F90=ifort F90FLAGS="--some-flag"
```

This will create a library called `lib_read_only.a`.

21.1.2 Read-Only-Verification Library for GOcean

This library is contained in `lib/read_only/dl_esm_inf` and it must be compiled before linking any GOcean-based application that uses read-only verification. Compiling this library requires access to the Gocean infrastructure library (since it must implement a generic interface for e.g. the `dl_esm_inf r2d_field` class).

The makefile uses the variable `INF_DIR` to point to the location where `dl_esm_inf`'s `field_mod` has been compiled. It defaults to `../../../../external/dl_esm_inf/finite_difference/src`, which is the location of the `dl_esm_info` version that is included (as a git submodule) with PSyclone. It can be changed to a user-specified location if required. The `dl_esm_inf` library is not used in linking the verification library. The application which uses the read-only-verification library needs to link in the infrastructure library anyway.

Compilation of the library is done by invoking `make` and setting the required variables:

```
make INF_DIR=some_path F90=ifort F90FLAGS="--some-flag"
```

This will create a library called `lib_read_only.a`. An executable example for using the GOcean read-only-verification library is included in `examples/gocean/eg7`, see [Example 7: Read-only-verification](#).

21.2 NAN Test

This transformation can be used for both LFRic and GOcean APIs. It will test all input- and output-parameter of a kernel to make sure they are not NAN or infinite. If they are, an error message like the following is printed, but the program is not aborted:

```
PSYDATA: Variable perturbation has the invalid value
          NaN at index/indices          351
```

It uses the function `IEEE_IS_FINITE` from the `ieee_arithmetic` module for this test. Note that only floating point numbers will be tested. Integer numbers do not have a bit pattern for 'infinity' or NAN.

The runtime libraries for GOcean and LFRic are based on a jinja-template contained in the directory `<PSYCLONEHOME>/lib/nan_test`. The respective API-specific libraries map the internal field structures to Fortran basic types and call the functions from the base class to handle those.

PROFILING

PSyclone has the ability to define regions that can be profiled with various performance measurement tools. The profiling can be enabled automatically using command line parameters like:

```
psyclone --profile kernels ...
```

Or, for finer-grained control, it may be applied via a profiling transformation within a transformation script.

PSyclone can be used with a variety of existing profiling tools. It currently supports `dl_timer`, Dr Hook, the nvidia GPU profiling tools and it comes with a simple stand-alone timer library. The PSyData API ([PSyData API](#)) is utilised to implement wrapper libraries that connect the PSyclone application to the profiling libraries. Certain adjustments to the application's build environment are required:

- The compiler needs to be able to find the module files for the wrapper of the selected profiling library.
- The application needs to be linked with the wrapper library that interfaces between the PSyclone API and the tool-specific API.
- The tool-specific library also needs to be linked in.

It is the responsibility of the user to supply the corresponding compiler command line options when building the application that incorporates the PSyclone-generated code.

22.1 Interface to Third Party Profiling Tools

PSyclone comes with wrapper libraries to support usage of Dr Hook, `dl_timer`, NVTX (NVIDIA Tools Extension library), and a simple non-thread-safe timing library. Support for further profiling libraries will be added in the future. To compile the wrapper libraries, change into the directory `lib/profiling` of PSyclone and type `make` to compile all wrappers. If only some of the wrappers are required, you can either use `make wrapper-name` (e.g. `make drhook`), or change into the corresponding directory and use `make`. The corresponding README files contain additional parameters that can be set in order to find third party profiling tools. Below are short descriptions of each of the various wrapper libraries that come with PSyclone:

lib/profiling/template This is a simple library that just prints out the name as regions are entered and exited. It could act as a template to develop new wrapper libraries, hence its name.

lib/profiling/simple_timing This is a simple, stand-alone library that uses Fortran system calls to measure the execution time, and reports average, minimum and maximum execution time for all regions. It is not MPI aware (i.e. it will just report independently for each MPI process), and not thread-safe.

lib/profiling/dl_timer This wrapper uses the `apeg-dl_timer` library. In order to use this wrapper, you must download and install the `dl_timer` library from https://bitbucket.org/apeg/dl_timer. This library has various compile-time options and may be built with MPI or OpenMP support. Additional link options might therefore be required (e.g. enabling OpenMP, or linking with MPI).

lib/profiling/drhook This wrapper uses the DrHook library. You need to contact ECMWF to obtain a copy of DrHook.

lib/profiling/nvidia This is a wrapper library that maps the PSyclone profiling API to the NVIDIA Tools Extension library (NVTX). This library is available from <https://developer.nvidia.com/cuda-toolkit>.

Any user can create similar wrapper libraries for other profiling tools by providing a corresponding Fortran module. The functions that need to be implemented are described in the developer's guide ([PSyData API](#)).

Most libraries in `lib/profiling` need to be linked in with the corresponding 3rd party profiling tool. The exception is the `template-` and `simple_timing-`library, which are stand alone. The profiling example in `examples/gocean/eg5` can be used with any of the wrapper libraries (except `nvidia`) to see how they work.

22.2 Required Modifications to the Program

In order to guarantee that any profiling library is properly initialised, PSyclone's profiling wrappers utilise two additional function calls that the user must manually insert into the program:

22.2.1 `profile_PSyDataInit()`

This method needs to be called once to initialise the profiling tool. At this stage this call is not automatically inserted by PSyclone, so it is the responsibility of the user to add the call to an appropriate location in the application:

```
use profile_psy_data_mod, only : profile_PSyDataInit
...
call profile_PSyDataInit()
```

The 'appropriate' location might depend on the profiling library used. For example, it might be necessary to invoke this before or after a call to `MPI_Init()`.

22.2.2 `profile_PSyDataShutdown()`

At the end of the program the function `profile_PSyDataShutdown()` must be called. It will make sure that the measurements are printed, files are flushed, and that the profiling tool is closed correctly. Again at this stage it is necessary to manually insert the call at an appropriate location:

```
use profile_psy_data_mod, only : profile_PSyDataShutdown
...
call profile_PSyDataShutdown()
```

And again the appropriate location might depend on the profiling library used (e.g. before or after a call to `MPI_Finalize()`).

22.3 Profiling Command Line Options

PSyclone offers two command line options to automatically instrument code with profiling regions. It can create profile regions around a full invoke (including all kernel calls in this invoke), and/or around each individual kernel.

The option `--profile invoke` will automatically add calls to start and end a profile region at the beginning and end of every invoke subroutine created by PSyclone. All kernels called within this invoke subroutine will be included in the profiled region.

The option `--profile kernels` will surround each outer loop created by PSyclone with start and end profiling calls.

Note: In some APIs (for example dynamo when using distributed memory) additional minor code might get included in a profiled kernel section, for example `setDirty()` calls (expensive calls like `HaloExchange` are excluded).

Note: If the `kernels` option is used in combination with an optimisation script that introduces OpenACC then profiling calls are automatically excluded from within OpenACC regions (since the PSyData wrappers are not compiled for GPU execution).

Note: It is still the responsibility of the user to manually add the calls to `profile_PSyDataInit` and `profile_PSyDataShutdown` to the code base (see *Required Modifications to the Program*).

PSyclone will modify the schedule of each invoke to insert the profiling regions. Below we show an example of a schedule created when instrumenting invokes - all children of a Profile-Node will be part of the profiling region, including all loops created by PSyclone and all kernel calls (note that for brevity, the nodes holding the loop bounds have been omitted for all but the first loop):

```
GOInvokeSchedule[invoke='invoke_1',Constant loop bounds=True]
  0: [Profile]
    Schedule[]
      0: Loop[type='outer',field_space='go_cu',it_space='go_internal_pts']
        Literal[value:'2']
        Literal[value:'jstop']
        Literal[value:'1']
        Schedule[]
          0: Loop[type='inner',field_space='go_cu',
            it_space='go_internal_pts']
            ...
            Schedule[]
              0: CodedKern compute_unew_code(unew_fld,uold_fld,z_fld,
                cv_fld,h_fld,tdt,dy) [module_inline=False]
            1: Loop[type='outer',field_space='cv',it_space='internal_pts']
              ...
              Schedule[]
                0: Loop[type='inner',field_space='cv',it_space='internal_pts']
                  ...
                  Schedule[]
                    0: CodedKern compute_vnew_code(vnew_fld,vold_fld,z_fld,
                      cu_fld,h_fld,tdt,dy) [module_inline=False]
                2: Loop[type='outer',field_space='ct',it_space='internal_pts']
                  ...
                  Schedule[]
```

(continues on next page)

(continued from previous page)

```

0: Loop[type='inner',field_space='ct',it_space='internal_pts']
...
Schedule[]
0: CodedKern compute_pnew_code(pnew_fld,pold_fld,cu_fld,
cv_fld,tdt,dx,dy) [module_inline=False]

```

And now the same schedule when instrumenting kernels. In this case each loop nest and kernel call will be contained in a separate region:

```

GOInvokeSchedule[invoke='invoke_1',Constant loop bounds=True]
0: [Profile]
Schedule[]
0: Loop[type='outer',field_space='go_cu',it_space='go_internal_pts']
...
Schedule[]
0: Loop[type='inner',field_space='go_cu',
it_space='go_internal_pts']
...
Schedule[]
0: CodedKern compute_unew_code(unew_fld,uold_fld,z_fld,
cv_fld,h_fld,tdt,dy) [module_inline=False]
1: [Profile]
Schedule[]
0: Loop[type='outer',field_space='go_cv',it_space='go_internal_pts']
...
Schedule[]
0: Loop[type='inner',field_space='go_cv',
it_space='go_internal_pts']
...
Schedule[]
0: CodedKern compute_vnew_code(vnew_fld,vold_fld,z_
↪fld,
cu_fld,h_fld,tdt,dy) [module_inline=False]
2: [Profile]
Schedule[]
0: Loop[type='outer',field_space='go_ct',it_space='go_internal_pts']
...
Schedule[]
0: Loop[type='inner',field_space='go_ct',
it_space='go_internal_pts']
...
Schedule[]
0: CodedKern compute_pnew_code(pnew_fld,pold_fld,
cu_fld,cv_fld,tdt,dx,dy) [module_inline=False]

```

Both options can be specified at the same time:

```

GOInvokeSchedule[invoke='invoke_1',Constant loop bounds=True]
0: [Profile]
Schedule[]
0: [Profile]
Schedule[]
0: Loop[type='outer',field_space='go_cu',
it_space='go_internal_pts']
...
Schedule[]

```

(continues on next page)

(continued from previous page)

```

        0: Loop[type='inner',field_space='go_cu',
            it_space='go_internal_pts']
        ...
        Schedule[]
        0: CodedKern compute_unew_code(unew_fld,uold_fld,
            ...) [module_inline=False]

1: [Profile]
    Schedule[]
    0: Loop[type='outer',field_space='go_cv',
        it_space='go_internal_pts']
    ...
    Schedule[]
    0: Loop[type='inner',field_space='go_cv',
        it_space='go_internal_pts']
    ...
    Schedule[]
    0: CodedKern compute_vnew_code(vnew_fld,vold_
→fld,
        ...) [module_inline=False]

2: [Profile]
    Schedule[]
    0: Loop[type='outer',field_space='go_ct',
        it_space='go_internal_pts']
    ...
    Schedule[]
    0: Loop[type='inner',field_space='go_ct',
        it_space='go_internal_pts']
    ...
    Schedule[]
    0: CodedKern compute_pnew_code(pnew_fld,pold_fld,
        ...) [module_inline=False]

```

22.4 Profiling in Scripts - ProfileTrans

The greatest flexibility is achieved by using the profiler transformation explicitly in a transformation script. The script takes either a single PSyIR Node or a list of PSyIR Nodes as argument, and will insert a Profile Node into the PSyIR, with the specified nodes as children. At code creation time the listed children will all be enclosed in one profile region. As an example:

```

from psyclone.psyir.transformations import ProfileTrans

p_trans = ProfileTrans()
schedule = psy.invokes.get('invoke_0').schedule
schedule.view()

# Enclose some children within a single profile region
newschedule, _ = p_trans.apply(schedule.children[1:3])
newschedule.view()

```

The profiler transformation also allows the profile name to be set explicitly, rather than being automatically created (see [Naming Profiling Regions](#) for details). This allows for potentially more intuitive names or finer grain control over profiling (as particular regions could be provided with the same profile names). For example:

```

invoke = psy.invokes.invoke_list[0]
schedule = invoke.schedule
profile_trans = ProfileTrans()
# Use the actual psy-layer module and subroutine names.
options = {"region_name": (psy.name, invoke.name)}
profile_trans.apply(schedule.children, options=options)
# Use own names and repeat for different regions to aggregate profile.
options = {"region_name": ("my_location", "my_region")}
profile_trans.apply(schedule[0].children[1:2], options=options)
profile_trans.apply(schedule[0].children[5:7], options=options)

```

Warning: If “region_name” is misspelt in the options dictionary then the option will be silently ignored. This is true for all options. Issue #613 captures this problem.

Warning: It is the responsibility of the user to make sure that a profile region is only created inside a multi-threaded region if the profiling library used is thread-safe!

22.5 Naming Profiling Regions

A profile region derives its name from two components:

module_name A string identifying the psy-layer containing this profile node.

region_name A string identifying the invoke containing this profile node and its location within the invoke (where necessary).

By default PSyclone will generate appropriate names to uniquely determine a particular region. Since those names can be somewhat cryptic, alternative names can be specified by the user when adding profiling via a transformation script, see [Passing Parameters From the User to the Node Constructor](#).

The automatic name generation depends on the API according to the following rules:

For the *nemo* api,

- the *module_name* string is set to the name of the parent function/subroutine/program. This name is unique as Fortran requires these names to be unique within a program.
- the *region_name* is set to an *r* (standing for region) followed by an integer which uniquely identifies the profile within the parent function/subroutine/program (based on the profile node’s position in the PSyIR representation relative to any other profile nodes).

For the *dynamo* and *gocean* api’s,

- the *module_name* string is set to the module name of the generated PSy-layer. This name should be unique by design (otherwise module names would clash when compiling).
- the *region_name* is set to the name of the invoke in which it resides, followed by a *:* and a kernel name if the profile region contains a single kernel, and is completed by *:r* (standing for region) followed by an integer which uniquely identifies the profile within the invoke (based on the profile node’s position in the PSyIR representation relative to any other profile nodes). For example:

```

InvokeSchedule[invoke='invoke_0', dm=True]
0: Profile[]

```

(continues on next page)

(continued from previous page)

```

Schedule[]
  0: Profile[]
    Schedule[]
      0: HaloExchange[field='f2', type='region', depth=1,
        check_dirty=True]
      1: HaloExchange[field='m1', type='region', depth=1,
        check_dirty=True]
      2: HaloExchange[field='m2', type='region', depth=1,
        check_dirty=True]
    1: Profile[]
      Schedule[]
        0: Loop[type='', field_space='w1', it_space='cells',
          upper_bound='cell_halo(1)']
          Literal[value:'1', DataType.INTEGER]
          Literal[value:'mesh%get_last_halo_cell(1)',
            DataType.INTEGER]
          Literal[value:'1', DataType.INTEGER]
          Schedule[]
            0: CodedKern testkern_code(a,f1,f2,m1,m2)
              [module_inline=False]
          1: Profile[]
            Schedule[]
              0: Loop[type='', field_space='w1',
                it_space='cells',
                upper_bound='cell_halo(1)']
                Literal[value:'1', DataType.INTEGER]
                Literal[value:'mesh%get_last_halo_cell(1)',
                  DataType.INTEGER]
                Literal[value:'1', DataType.INTEGER]
                Schedule[]
                  0: CodedKern testkern_code(a,f1,f2,m1,m2)
                    [module_inline=False]
            2: Loop[type='', field_space='w1', it_space='cells',
              upper_bound='cell_halo(1)']
              Literal[value:'1', DataType.INTEGER]
              Literal[value:'mesh%get_last_halo_cell(1)', DataType.INTEGER]
              Literal[value:'1', DataType.INTEGER]
              Schedule[]
                0: CodedKern testkern_qr_code(f1,f2,m1,a,m2,istp)
                  [module_inline=False]

```

This is the code created for this example:

```

MODULE container
CONTAINS
SUBROUTINE invoke_0(a, f1, f2, m1, m2, istp, qr)
...
CALL psy_data_3%PreStart("multi_functions_multi_invokes_psy", "invoke_0:r0", &
  0, 0)
CALL psy_data%PreStart("multi_functions_multi_invokes_psy", "invoke_0:r1", 0, 0)
IF (f2_proxy%is_dirty(depth=1)) THEN
  CALL f2_proxy%halo_exchange(depth=1)
END IF
IF (m1_proxy%is_dirty(depth=1)) THEN
  CALL m1_proxy%halo_exchange(depth=1)
END IF
IF (m2_proxy%is_dirty(depth=1)) THEN

```

(continues on next page)

(continued from previous page)

```
        CALL m2_proxy%halo_exchange(depth=1)
    END IF
    CALL psy_data%PreEnd()
    CALL psy_data_1%PreStart("multi_functions_multi_invokes_psy", "invoke_0:r2", &
                           0, 0)
    DO cell=1,mesh%get_last_halo_cell(1)
        CALL testkern_code(...)
    END DO
    ...
    CALL psy_data_2%PreStart("multi_functions_multi_invokes_psy", &
                           "invoke_0:testkern_code:r3", 0, 0)
    DO cell=1,mesh%get_last_halo_cell(1)
        CALL testkern_code(...)
    END DO
    ...
    CALL psy_data_2%PostEnd()
    CALL psy_data_1%PostEnd()
    ...
    DO cell=1,mesh%get_last_halo_cell(1)
        CALL testkern_qr_code(...)
    END DO
    ...
    CALL psy_data_3%PostEnd()
    ...
END SUBROUTINE invoke_0
END MODULE container
```


PSY KERNEL EXTRACTOR (PSYKE)

23.1 Introduction

PSyclone has the ability to define regions of a PSyclone-conformant code to be extracted and run as a stand-alone application. This ability, called PSyKE (PSy Kernel Extractor), can be useful for benchmarking parts of a model, such as LFRic, without the need for using its infrastructure.

23.1.1 Mechanism

The code marked for extraction can be (subject to *Restrictions*):

- One or more Nodes in an Invoke (e.g. Loops containing Kernel or Built-In calls, a Directive enclosing one or more Loops) or
- The entire Invoke (extraction applied to all Nodes).

The basic mechanism of code extraction is through applying the `ExtractTrans` transformation to selected Nodes. This transformation is further sub-classed into API-specific implementations, `LFRicExtractTrans` and `GOceanExtractTrans`. Both sub-classed transformations insert an instance of the `ExtractNode` object into the Schedule of a specific Invoke. All Nodes marked for extraction become children of the `ExtractNode`.

The `ExtractNode` class uses the dependency analysis to detect which variables are input-, and which ones are output-parameters. The lists of variables are then passed to the `PSyDataNode`, which is the base class of any `ExtractNode` (details of the `PSyDataNode` can be found in [PSyData API](#)). This node then creates the actual code, as in the following LFRic example:

```
! ExtractStart
!
CALL extract_psy_data%PreStart("testkern_mod", "testkern_code", 4, 2)
CALL extract_psy_data%PreDeclareVariable("a", a)
CALL extract_psy_data%PreDeclareVariable("f2", f2)
CALL extract_psy_data%PreDeclareVariable("m1", m1)
CALL extract_psy_data%PreDeclareVariable("m2", m2)
CALL extract_psy_data%PreDeclareVariable("map_w1", map_w1)
...
CALL extract_psy_data%PreDeclareVariable("undf_w3", undf_w3)
CALL extract_psy_data%PreDeclareVariable("f1_post", f1)
CALL extract_psy_data%PreDeclareVariable("cell_post", cell)
CALL extract_psy_data%PreEndDeclaration
CALL extract_psy_data%ProvideVariable("a", a)
CALL extract_psy_data%ProvideVariable("f2", f2)
CALL extract_psy_data%ProvideVariable("m1", m1)
CALL extract_psy_data%ProvideVariable("m2", m2)
```

(continues on next page)

(continued from previous page)

```

CALL extract_psy_data%ProvideVariable("map_w1", map_w1)
...
CALL extract_psy_data%ProvideVariable("undf_w3", undf_w3)
CALL extract_psy_data%PreEnd
DO cell=1,f1_proxy%vspace%get_ncell()
  !
  CALL testkern_code(nlayers, a, f1_proxy%data, f2_proxy%data, &
    m1_proxy%data, m2_proxy%data, ndf_w1, undf_w1, &
    map_w1(:,cell), ndf_w2, undf_w2, map_w2(:,cell), ndf_w3, &
    undf_w3, map_w3(:,cell))
END DO
CALL extract_psy_data%PostStart
CALL extract_psy_data%ProvideVariable("cell_post", cell)
CALL extract_psy_data%ProvideVariable("f1_post", f1)
CALL extract_psy_data%PostEnd
!
! ExtractEnd

```

The PSyData API relies on generic Fortran interfaces to provide the field-type-specific implementations of the `ProvideVariable` for different types. This means that a different version of the external PSyData library that PSyKE uses must be supplied for each PSyclone API.

23.1.2 Restrictions

Code extraction can be applied to unoptimised or optimised code. There are restrictions that check for correctness of optimising transformations when extraction is applied, as well as restrictions that eliminate dependence on the specific model infrastructure.

General

This group of restrictions is enforced irrespective of whether optimisations are used or not.

- Extraction can be applied to a single Node or a list of Nodes in a Schedule. For the latter, Nodes in the list must be consecutive children of the same parent Schedule.
- Extraction cannot be applied to an `ExtractNode` or a Node list that already contains one (otherwise we would have an extract region within another extract region).
- A Kernel or a Built-In call cannot be extracted without its parent Loop.

Distributed memory

As noted in the *Distributed Memory* section, support for distributed memory in PSyclone is currently limited to the Dynamo0.3 API. Since the implementation generates calls to LFRic infrastructure (e.g. runtime checks for status of field halos), code extraction is not allowed when distributed memory is enabled.

Shared memory and API-specific

The ExtractTrans transformation cannot be applied to:

- A Loop without its parent Directive,
- An orphaned Directive (e.g. OMPDoDirective, ACCLoopDirective) without its parent Directive (e.g. ACC or OMP Parallel Directive),
- A Loop over cells in a colour without its parent Loop over colours in the Dynamo0.3 API,
- An inner Loop without its parent outer Loop in the GOcean1.0 API.

23.2 Use

The code extraction is currently enabled by utilising a transformation script (see *Script* section for more details).

For example, the transformation script which extracts the first Kernel call in Dynamo0.3 API test example 15.1.2_builtin_and_normal_kernel_invoke.f90 would be written as:

```
from psyclone.domain.lfric.transformations import LFRicExtractTrans

# Get instance of the ExtractRegionTrans transformation
ettrans = LFRicExtractTrans()

# Get Invoke and its Schedule
invoke = psy.invokes.get("invoke_0")
schedule = invoke.schedule

# Apply extract transformation to the selected Node
schedule, _ = ettrans.apply(schedule.children[2])
schedule.view()
```

and called as:

```
> psyclone -nodm -s ./extract_single_node.py \
    <path-to-example>/15.1.2_builtin_and_normal_kernel_invoke.f90
```

PSyclone modifies the Schedule of the selected invoke_0:

```
Schedule[invoke='invoke_0' dm=False]
  0: Loop[type='dofs',field_space='any_space_1',it_space='dofs',
        upper_bound='ndofs']
    Literal[value:'NOT_INITIALISED']
    Literal[value:'NOT_INITIALISED']
    Literal[value:'1']
    Schedule[]
      0: BuiltIn setval_c(f5,0.0)
  1: Loop[type='dofs',field_space='any_space_1',it_space='dofs',
        upper_bound='ndofs']
    ...
    Schedule[]
      0: BuiltIn setval_c(f2,0.0)
  2: Loop[type='',field_space='w2',it_space='cells', upper_bound='ncells']
    ...
    Schedule[]
      0: CodedKern testkern_code_w2_only(f3,f2) [module_inline=False]
```

(continues on next page)

(continued from previous page)

```

3: Loop[type='', field_space='wtheta', it_space='cells', upper_bound='ncells']
...
Schedule[]
0: CodedKern testkern_wtheta_code(f4,f5) [module_inline=False]
4: Loop[type='', field_space='w1', it_space='cells', upper_bound='ncells']
...
Schedule[]
0: CodedKern testkern_code(scalar,f1,f2,f3,f4) [module_inline=False]

```

to insert the extract region. As shown below, all children of an ExtractNode will be part of the region:

```

Schedule[invoke='invoke_0' dm=False]
0: Loop[type='dofs', field_space='any_space_1', it_space='dofs',
    upper_bound='ndofs']
...
Schedule[]
0: BuiltIn setval_c(f5,0.0)
1: Loop[type='dofs', field_space='any_space_1', it_space='dofs',
    upper_bound='ndofs']
...
Schedule[]
0: BuiltIn setval_c(f2,0.0)
2: Extract
Schedule[]
0: Loop[type='', field_space='w2', it_space='cells', upper_bound='ncells']
...
Schedule[]
0: CodedKern testkern_code_w2_only(f3,f2) [module_inline=False]
3: Loop[type='', field_space='wtheta', it_space='cells', upper_bound='ncells']
...
Schedule[]
0: CodedKern testkern_wtheta_code(f4,f5) [module_inline=False]
4: Loop[type='', field_space='w1', it_space='cells', upper_bound='ncells']
...
Schedule[]
0: CodedKern testkern_code(scalar,f1,f2,f3,f4) [module_inline=False]

```

To extract multiple Nodes, ExtractTrans can be applied to the list of Nodes (subject to *General* restrictions above):

```

# Apply extract transformation to the selected Nodes
schedule, _ = etrans.apply(schedule.children[1:3])

```

This modifies the above Schedule as:

```

...
Extract
Schedule[]
0: Loop[type='dofs', field_space='any_space_1', it_space='dofs',
    upper_bound='ndofs']
...
Schedule[]
0: BuiltIn setval_c(f2,0.0)
1: Loop[type='', field_space='w2', it_space='cells', upper_bound='ncells']
...
Schedule[]
0: CodedKern testkern_code_w2_only(f3,f2) [module_inline=False]
...

```

As said above, extraction can be performed on optimised code. For example, the following example transformation script first adds !\$OMP PARALLEL DO directive and then extracts the optimised code in Dynamo0.3 API test example 15.1.2_builtin_and_normal_kernel_invoke.f90:

```
from psyclone.domain.lfric.transformations import LFRicExtractTrans
from psyclone.transformations import DynamoOMPParallelLoopTrans

# Get instances of the transformations
etrans = LFRicExtractTrans()
otrans = DynamoOMPParallelLoopTrans()

# Get Invoke and its Schedule
invoke = psy.invokes.get("invoke_0")
schedule = invoke.schedule

# Add OMP PARALLEL DO directives
schedule, _ = otrans.apply(schedule.children[1])
schedule, _ = otrans.apply(schedule.children[2])
# Apply extract transformation to the selected Nodes
schedule, _ = etrans.apply(schedule.children[1:3])
schedule.view()
```

The generated code is now:

```
! ExtractStart
CALL extract_psy_data%PreStart("unknown-module", "setval_c", 1, 3)
CALL extract_psy_data%PreDeclareVariable("f2", f2)
CALL extract_psy_data%PreDeclareVariable("cell_post", cell)
CALL extract_psy_data%PreDeclareVariable("df_post", df)
CALL extract_psy_data%PreDeclareVariable("f3_post", f3)
...
CALL extract_psy_data%PreEndDeclaration
CALL extract_psy_data%ProvideVariable("f2", f2)
...
CALL extract_psy_data%PreEnd
!
!$omp parallel do default(shared), private(df), schedule(static)
DO df=1, undf_aspc1_f2
    f2_proxy%data(df) = 0.0
END DO
!$omp end parallel do
!$omp parallel do default(shared), private(cell), schedule(static)
DO cell=1, f3_proxy%vspace%get_ncell()
    !
    CALL testkern_code_w2_only(nlayers, f3_proxy%data, f2_proxy%data, ndf_w2, undf_w2, ↵
    ↵map_w2(:, cell))
END DO
!$omp end parallel do
CALL extract_psy_data%PostStart
CALL extract_psy_data%ProvideVariable("cell_post", cell)
CALL extract_psy_data%ProvideVariable("df_post", df)
CALL extract_psy_data%ProvideVariable("f3_post", f3)
CALL extract_psy_data%PostEnd
!
! ExtractEnd
```

Note: At this stage builtins are not fully supported, resulting in f2 being incorrectly detected as an input parameter,

and not as an output parameter. This issue is tracked in #637.

Examples in `examples/lfric/eg12` directory demonstrate how to apply code extraction by utilising PSyclone transformation scripts (see [Examples](#) section for more information).

23.3 NetCDF Extraction Example

PSyclone comes with an example NetCDF based extraction library in `lib/extract/netcdf/dl_esm_inf`. This library implements the full PSyData API for use with the GOcean 1.0 `dl_esm_inf` infrastructure library. In order to compile this library, you must have NetCDF installed. When running the code, it will create a NetCDF file for the instrumented code region. It includes all variables that are read before the code is executed, and all variables that have been modified. The output variables have the postfix `_post` attached to the NetCDF names, e.g. a variable `xyz` that is read and written will be stored with the name `xyz` containing the input values, and the name `xyz_post` containing the output values. Arrays have their size stored as NetCDF dimensions: again the variable `xyz` will have its sizes stored as `xyzdim1, xyzdim2` for the input values, and output arrays use the name `xyz_postdim1, xyz_postdim2`.

The output file contains the values of all variables used in the subroutine. The `GOceanExtractTrans` can automatically create a driver program which will read the netcdf file and then call the instrumented region. In order to create this driver program, the options parameter `create_driver` must be set to true:

```
extract = GOceanExtractTrans()
extract.apply(schedule.children,
              {"create_driver": True,
               "region_name": ("main", "init")})
```

This will create a Fortran file called `driver-main-init.f90`, which can then be compiled and executed.

CONFIGURATION

PSyclone reads various run-time configuration options from the `psyclone.cfg` file. As described in [Configuration](#), the default `psyclone.cfg` configuration file is installed in `<python-base-prefix>/share/psyclone/` during the installation process. The original version of this file is in the `PSyclone/config` directory of the PSyclone distribution.

At execution-time, the user can specify a custom configuration file to be used. This can either be done with the `--config` command line option, or by specifying the (full path to the) configuration file to use via the `PSYCLONE_CONFIG` environment variable. If the specified configuration file is not found then PSyclone will fall back to searching in a list of default locations.

The ordering of these locations depends upon whether PSyclone is being run within a Python virtual environment (such as `venv`). If no virtual environment is detected then the locations searched, in order, are:

1. `${PWD}/.psyclone/`
2. `${HOME}/.local/share/psyclone/`
3. `<python-base-dir>/share/psyclone/`

where `<python-base-dir>` is the path stored in Python's `sys.prefix`.

If a virtual environment is detected then it is assumed that the `share` directory will be a part of that environment. In order to maintain isolation of distinct virtual environments this directory is then checked *before* the user's home directory, i.e. the list of locations searched is now:

1. `${PWD}/.psyclone/`
2. `<python-base-dir>/share/psyclone/`
3. `${HOME}/.local/share/psyclone/`

Note that for developers a slightly different configuration handling is implemented, see [Module: configuration](#) for details.

24.1 Options

The configuration file is read by the Python `ConfigParser` class (<https://docs.python.org/3/library/configparser.html>) and must be formatted accordingly. It currently consists of a `DEFAULT` section e.g.:

```
[DEFAULT]
DEFAULTAPI = dynamo0.3
DEFAULTSTUBAPI = dynamo0.3
DISTRIBUTED_MEMORY = true
REPRODUCIBLE_REDUCTIONS = false
REPROD_PAD_SIZE = 8
```

(continues on next page)

(continued from previous page)

```
PSYIR_ROOT_NAME = psyir_tmp
VALID_PSY_DATA_PREFIXES = profile extract
```

and an optional API specific section, for example for the dynamo0.3 section:

```
[dynamo0.3]
access_mapping = gh_read: read, gh_write: write, gh_readwrite: readwrite,
                  gh_inc: inc, gh_sum: sum
COMPUTE_ANNEXED_DOFS = false
default_kind = real: r_def, integer: i_def, logical: l_def
RUN_TIME_CHECKS = false
```

or for gocean1.0:

```
[gocean1.0]
access_mapping = go_read:read, go_write:write, go_readwrite:readwrite
grid-properties = go_grid_xstop: {0}%%grid%%subdomain%%internal%%xstop: scalar,
                  go_grid_ystop: {0}%%grid%%subdomain%%internal%%ystop: scalar,
                  go_grid_data: {0}%%data: array,
                  ...
```

The meaning of the various entries is described in the following sub-sections.

Note that ConfigParser supports various forms of boolean entry including “true/false”, “yes/no” and “1/0”. See <https://docs.python.org/3/library/configparser.html#supported-datatypes> for more details.

24.1.1 DEFAULT Section

This section contains entries that are, in principle, applicable to all APIs supported by PSyclone.

Entry	Description
DEFAULTAPI	The API that PSyclone assumes an Algorithm/Kernel conforms to if no API is specified. Must be one of the APIs supported by PSyclone (dynamo0.1, dynamo0.3, gocean0.1, gocean1.0 and nemo). If there is no API specified and there is only one API-specific section in the config file loaded, this API will be used. This value can be overwritten by the command line option ‘-api’. If there is no API entry in the config file, and ‘-api’ is not specified on the command line, dynamo0.3 is used as default.
DEFAULTSTUBAPI	The API that the kernel-stub generator assumes by default. Must be one of the stub-APIs supported by PSyclone (dynamo0.3 only at this stage).
DISTRIBUTED_MEMORY	Whether or not to generate code for distributed-memory parallelism by default. Note that this is currently only supported for the dynamo0.3 API.
REPRODUCIBLE_REDUCTIONS	Whether or not to generate code for reproducible OpenMP reductions (see <i>Reductions</i>) by default.
REPROD_PAD_SIZE	If generating code for reproducible OpenMP reductions, this setting controls the amount of padding used between elements of the array in which each thread accumulates its local reduction. (This prevents false sharing of cache lines by different threads.)
PSYIR_ROOT_NAME	The root for generated PSyIR symbol names if one is not supplied when creating a symbol. Defaults to “psyir_tmp”.
VALID_PSY_DATA_PREFIXES	Which class prefixes are permitted in any PSyData-related transformations. See <i>PSyData API</i> for details.

24.1.2 Common Sections

The following entries must be defined for each API in order for PSyclone to work as expected:

Entry	Description
access_mapping	This field defines the strings that are used by a particular API to indicate write, read, ... access. Its value is a comma separated list of access-string:access pairs, e.g.: gh_read: read, gh_write: write, gh_readwrite: readwrite, gh_inc: inc, gh_sum: sum At this stage these 5 types are defined for read, write, read+write, increment and summation access by PSyclone. Sum is a form of reduction. The GOcean APIs do not support increment or sum, so they only define three mappings for read, write, and readwrite.

24.1.3 dynamo0.3 Section

This section contains configuration options that are only applicable when using the Dynamo 0.3 API.

Entry	Description
COMPUTE_ANNEXED_DOFS	Whether or not to perform redundant computation over annexed dofs in order to reduce the number of halo exchanges, see <i>Annexed DoFs</i> .
default_kind	Captures the default kinds (precisions) for the supported datatypes in LFRic (<i>real</i> , <i>integer</i> and <i>logical</i>).
RUN_TIME_CHECKS	Specifies whether to generate run-time validation checks, see <i>Run-time Checks</i> .

24.1.4 gocean1.0 Section

This section contains configuration options that are only applicable when using the Gocean 1.0 API.

Entry	Description
iteration-spaces	This contains definitions of additional iteration spaces used by PSyclone. A detailed description can be found in the <i>Iteration-spaces</i> section of the GOcean1.0 chapter.
grid-properties	This key contains definitions to access various grid properties. A detailed description can be found in the <i>Grid Properties</i> section of the GOcean1.0 chapter.

24.1.5 NEMO Section

This section contains configuration options that are only applicable when using the NEMO API.

Entry	Description
mapping-TYPE	This declares a mapping for a certain loop level, specified as TYPE. Each value must have three key:value pairs. A value can be empty if it is not required or not known, but the key must still be specified. The required keys are: var: the variable name that indicates the loop level, start: the first loop iteration, and stop: the last loop iteration. Each loop detected by the NEMO API will be given one of the TYPE values specified in the configuration file. See the example below for more details.
index-order	Specifies the order in which loops are created when converting an implicit loop to an explicit loop. All values in this comma-separated list must have a corresponding mapping-TYPE value defined.

Below we show an example of the NEMO section of a PSyclone configuration file. Note how the values in `index-order` have corresponding mapping entries, e.g. `mapping-lon`, `mapping-lat` etc.:

```
mapping-lon = var: ji, start: 1, stop: jpi
mapping-lat = var: jj, start: 1, stop: jpj
mapping-levels = var: jk, start: 1, stop: jpk
mapping-tracers = var: jt, start: 1, stop:
mapping-unknown = var: , start: 1, stop:

index-order = lon, lat, levels, tracers
```

If a NEMO loop then uses `Do jj=...`, PSyclone will give this loop the type 'lat', because the loop uses the variable name specified in the configuration file for a loop of type 'lat'. The loop type can be accessed using `loop.loop_type`, i.e. in this example it will be `loop.loop_type == 'lat'`.

The entry `mapping-unknown` has an empty value for the key 'var'. This means that the type 'unknown' will be used for any loop that can not be mapped using any of the other variable names in the configuration file.

BIBLIOGRAPHY

- [nem13] *NEMO Coding Conventions*. 2013. URL: https://forge.ipsl.jussieu.fr/nemo/attachment/wiki/Literature/NEMO_coding.conv_v3.pdf.
- [AFH+19] S. V. Adams, R. W. Ford, M. Hambley, J. M. Hobson, I. Kavčič, C. M. Maynard, T. Melvin, E. H. Müller, S. Mullerworth, A. R. Porter, M. Rezny, B. J. Shipway, and R. Wong. LFRic: Meeting the challenges of scalability and performance portability in Weather and Climate models. *Journal of Parallel and Distributed Computing*, 132:383–396, 2019. doi:<https://doi.org/10.1016/j.jpdc.2019.02.007>.

PYTHON MODULE INDEX

p

- `psyclone.alg_gen`, [212](#)
- `psyclone.generator`, [183](#)
- `psyclone.line_length`, [213](#)
- `psyclone.parse.algorithm`, [184](#)
- `psyclone.psyGen`, [211](#)
- `psyclone.transformations`, [184](#)

A

- ACCDDataTrans (class in *psyclone.transformations*), 184
- ACCEnterDataTrans (class in *psyclone.transformations*), 185
- ACCKernelsTrans (class in *psyclone.transformations*), 186
- ACCLoopTrans (class in *psyclone.transformations*), 187
- ACCParallelTrans (class in *psyclone.transformations*), 189
- ACCRoutineTrans (class in *psyclone.transformations*), 189
- add() (*psyclone.psyir.symbols.SymbolTable* method), 127
- Alg (class in *psyclone.alg_gen*), 212
- apply() (*psyclone.transformations.ACCDataTrans* method), 185
- apply() (*psyclone.transformations.ACCEnterDataTrans* method), 186
- apply() (*psyclone.transformations.ACCKernelsTrans* method), 187
- apply() (*psyclone.transformations.ACCLoopTrans* method), 188
- apply() (*psyclone.transformations.ACCRoutineTrans* method), 190
- apply() (*psyclone.transformations.ColourTrans* method), 190
- apply() (*psyclone.transformations.DynamoOp3AsyncHaloExchangeTrans* method), 191
- apply() (*psyclone.transformations.DynamoOp3ColourTrans* method), 192
- apply() (*psyclone.transformations.DynamoOp3KernelConstTrans* method), 193
- apply() (*psyclone.transformations.DynamoOp3OMPLoopTrans* method), 194
- apply() (*psyclone.transformations.DynamoOp3RedundantComputationTrans* method), 195
- apply() (*psyclone.transformations.DynamoLoopFuseTrans* method), 196
- apply() (*psyclone.transformations.DynamoOMPParallelLoopTrans* method), 197
- apply() (*psyclone.transformations.GOceanLoopFuseTrans* method), 200
- apply() (*psyclone.transformations.GOceanOMPLoopTrans* method), 201
- apply() (*psyclone.transformations.GOceanOMPParallelLoopTrans* method), 201
- apply() (*psyclone.transformations.GOConstLoopBoundsTrans* method), 198
- apply() (*psyclone.transformations.GOLoopSwapTrans* method), 199
- apply() (*psyclone.transformations.KernelGlobalsToArguments* method), 201
- apply() (*psyclone.transformations.KernelModuleInlineTrans* method), 202
- apply() (*psyclone.transformations.LoopFuseTrans* method), 203
- apply() (*psyclone.transformations.MoveTrans* method), 204
- apply() (*psyclone.transformations.OCLTrans* method), 205
- apply() (*psyclone.transformations.OMPLoopTrans* method), 207
- apply() (*psyclone.transformations.OMPParallelLoopTrans* method), 207
- apply() (*psyclone.transformations.ParallelLoopTrans* method), 209
- apply() (*psyclone.transformations.ParallelRegionTrans* method), 210
- argument_data_symbols() (*psyclone.psyir.symbols.SymbolTable* property), 128
- argument_list() (*psyclone.psyir.symbols.SymbolTable* property), 128
- check_integrity() (in module *psyclone.transformations*), 211
- ColourTrans (class in *psyclone.transformations*), 190
- condition() (*psyclone.psyir.nodes.IfBlock* method), 125
- container() (*psyclone.psyGen.PSy* property), 211

ContainerSymbol (class in *psyclone.psyir.symbols*), 134
 containersymbols() (*psyclone.psyir.symbols.SymbolTable* property), 128
 copy() (*psyclone.psyir.symbols.Symbol* method), 133
 copy_external_global() (*psyclone.psyir.symbols.SymbolTable* method), 128
 create() (*psyclone.psyir.nodes.Range* static method), 123

D

data_arguments() (*psyclone.psyir.symbols.SymbolTable* property), 128
 DataSymbol (class in *psyclone.psyir.symbols*), 135
 datasymbols() (*psyclone.psyir.symbols.SymbolTable* property), 128
 dir_body() (*psyclone.psyGen.Directive* method), 125
 Dynamo0p3AsyncHaloExchangeTrans (class in *psyclone.transformations*), 191
 Dynamo0p3ColourTrans (class in *psyclone.transformations*), 192
 Dynamo0p3KernelConstTrans (class in *psyclone.transformations*), 193
 Dynamo0p3OMPLoopTrans (class in *psyclone.transformations*), 194
 Dynamo0p3RedundantComputationTrans (class in *psyclone.transformations*), 194
 DynamoLoopFuseTrans (class in *psyclone.transformations*), 196
 DynamoOMPParallelLoopTrans (class in *psyclone.transformations*), 197

E

else_body() (*psyclone.psyir.nodes.IfBlock* method), 125
 excluded_node_types (*psyclone.transformations.ACCLoopTrans* attribute), 189

F

find_symbol_table() (*psyclone.psyir.symbols.Symbol* method), 133
 FortLineLength (class in *psyclone.line_length*), 213

G

gen() (*psyclone.alg_gen.Alg* property), 212
 gen() (*psyclone.psyGen.PSy* property), 211
 generate() (in module *psyclone.generator*), 183
 get_external_symbol() (*psyclone.psyir.symbols.Symbol* method), 133

get_trans_name() (*psyclone.psyGen.TransInfo* method), 139
 get_trans_num() (*psyclone.psyGen.TransInfo* method), 139
 get_unresolved_datasymbols() (*psyclone.psyir.symbols.SymbolTable* method), 128
 global_symbols() (*psyclone.psyir.symbols.SymbolTable* property), 129
 GOceanLoopFuseTrans (class in *psyclone.transformations*), 200
 GOceanOMPLoopTrans (class in *psyclone.transformations*), 201
 GOceanOMPParallelLoopTrans (class in *psyclone.transformations*), 201
 GOConstLoopBoundsTrans (class in *psyclone.transformations*), 198
 GOLoopSwapTrans (class in *psyclone.transformations*), 199

I

if_body() (*psyclone.psyir.nodes.IfBlock* method), 125
 imported_symbols() (*psyclone.psyir.symbols.SymbolTable* method), 129
 interface() (*psyclone.psyir.symbols.Symbol* property), 133
 invokes() (*psyclone.psyGen.PSy* property), 211
 is_argument() (*psyclone.psyir.symbols.Symbol* property), 134
 is_global() (*psyclone.psyir.symbols.Symbol* property), 134
 is_local() (*psyclone.psyir.symbols.Symbol* property), 134
 is_unresolved() (*psyclone.psyir.symbols.Symbol* property), 134
 iteration_indices() (*psyclone.psyir.symbols.SymbolTable* property), 129

K

KernelGlobalsToArguments (class in *psyclone.transformations*), 201
 KernelModuleInlineTrans (class in *psyclone.transformations*), 202
 KernelTrans (class in *psyclone.transformations*), 203

L

length() (*psyclone.line_length.FortLineLength* property), 213
 lhs() (*psyclone.psyir.nodes.Assignment* method), 124
 list() (*psyclone.psyGen.TransInfo* property), 139

`local_datasympols()` (`psyclone.psyir.symbols.SymbolTable` property), 129
`local_typesympols()` (`psyclone.psyir.symbols.SymbolTable` property), 129
`long_lines()` (`psyclone.line_length.FortLineLength` method), 213
`lookup()` (`psyclone.psyir.symbols.SymbolTable` method), 129
`lookup_with_tag()` (`psyclone.psyir.symbols.SymbolTable` method), 130
`LoopFuseTrans` (class in `psyclone.transformations`), 203

M

`module`
`psyclone.alg_gen`, 212
`psyclone.generator`, 183
`psyclone.line_length`, 213
`psyclone.parse.algorithm`, 184
`psyclone.psyGen`, 211
`psyclone.transformations`, 184
`MoveTrans` (class in `psyclone.transformations`), 204

N

`name()` (`psyclone.psyGen.PSy` property), 212
`name()` (`psyclone.psyir.symbols.Symbol` property), 134
`name()` (`psyclone.transformations.ACCEnterDataTrans` property), 185
`name()` (`psyclone.transformations.ACCEnterDataTrans` property), 186
`name()` (`psyclone.transformations.ACCKernelsTrans` property), 187
`name()` (`psyclone.transformations.ACCLoopTrans` property), 189
`name()` (`psyclone.transformations.ACCTrans` property), 189
`name()` (`psyclone.transformations.ACCTrans` property), 190
`name()` (`psyclone.transformations.ColourTrans` property), 191
`name()` (`psyclone.transformations.Dynamo0p3AsyncHaloExchangeTrans` property), 191
`name()` (`psyclone.transformations.Dynamo0p3ColourTrans` property), 193
`name()` (`psyclone.transformations.Dynamo0p3KernelConsTrans` property), 194
`name()` (`psyclone.transformations.Dynamo0p3OMPLoopTrans` property), 194
`name()` (`psyclone.transformations.Dynamo0p3RedundantComputationTrans` property), 195
`name()` (`psyclone.transformations.DynamoLoopFuseTrans` property), 197
`name()` (`psyclone.transformations.DynamoOMPParallelLoopTrans` property), 198
`name()` (`psyclone.transformations.GOceanLoopFuseTrans` property), 200
`name()` (`psyclone.transformations.GOceanOMPLoopTrans` property), 201
`name()` (`psyclone.transformations.GOceanOMPParallelLoopTrans` property), 201
`name()` (`psyclone.transformations.GOConstLoopBoundsTrans` property), 199
`name()` (`psyclone.transformations.GOLoopSwapTrans` property), 199
`name()` (`psyclone.transformations.KernelGlobalsToArguments` property), 202
`name()` (`psyclone.transformations.KernelModuleInlineTrans` property), 203
`name()` (`psyclone.transformations.LoopFuseTrans` property), 203
`name()` (`psyclone.transformations.MoveTrans` property), 205
`name()` (`psyclone.transformations.OCLTrans` property), 205
`name()` (`psyclone.transformations.OMPLoopTrans` property), 207
`name()` (`psyclone.transformations.OMPParallelLoopTrans` property), 208
`name()` (`psyclone.transformations.OMPParallelTrans` property), 209
`name()` (`psyclone.transformations.ParallelLoopTrans` property), 210
`name()` (`psyclone.transformations.ParallelRegionTrans` property), 210
`name_from_tag()` (`psyclone.psyir.symbols.SymbolTable` method), 130
`new_symbol_name()` (`psyclone.psyir.symbols.SymbolTable` method), 130
`node()` (`psyclone.psyir.symbols.SymbolTable` property), 131
`num_trans()` (`psyclone.psyGen.TransInfo` property), 139

O

`OCLTrans` (class in `psyclone.transformations`), 205
`ompt_schedule()` (`psyclone.transformations.OMPLoopTrans` property), 207
`OMPLoopTrans` (class in `psyclone.transformations`), 206
`OMPParallelLoopTrans` (class in `psyclone.transformations`), 207

OMPParallelTrans (class in psy-
clone.transformations), 208

P

ParallelLoopTrans (class in psy-
clone.transformations), 209

ParallelRegionTrans (class in psy-
clone.transformations), 210

parent() (psyclone.psyir.symbols.SymbolTable prop-
erty), 131

parse() (in module psyclone.parse.algorithm), 184

precision_datasympols() (psy-
clone.psyir.symbols.SymbolTable property),
131

process() (psyclone.line_length.FortLineLength
method), 213

PSy (class in psyclone.psyGen), 211

psyclone.alg_gen
module, 212

psyclone.generator
module, 183

psyclone.line_length
module, 213

psyclone.parse.algorithm
module, 184

psyclone.psyGen
module, 211

psyclone.transformations
module, 184

R

Range (class in psyclone.psyir.nodes), 122

remove() (psyclone.psyir.symbols.SymbolTable
method), 131

resolve_deferred() (psy-
clone.psyir.symbols.Symbol method), 134

rhs() (psyclone.psyir.nodes.Assignment method), 125

RoutineSymbol (class in psyclone.psyir.symbols), 135

S

same_space() (psy-
clone.transformations.DynamoLoopFuseTrans
property), 197

shallow_copy() (psy-
clone.psyir.symbols.SymbolTable method),
131

specify_argument_list() (psy-
clone.psyir.symbols.SymbolTable method),
132

start() (psyclone.psyir.nodes.Range property), 123

step() (psyclone.psyir.nodes.Range property), 123

stop() (psyclone.psyir.nodes.Range property), 123

swap() (psyclone.psyir.symbols.SymbolTable method),
132

swap_symbol_properties() (psy-
clone.psyir.symbols.SymbolTable method),
132

Symbol (class in psyclone.psyir.symbols), 133

Symbol.Visibility (class in psy-
clone.psyir.symbols), 133

symbols() (psyclone.psyir.symbols.SymbolTable prop-
erty), 132

symbols_dict() (psy-
clone.psyir.symbols.SymbolTable property),
132

SymbolTable (class in psyclone.psyir.symbols), 127

T

tags_dict() (psyclone.psyir.symbols.SymbolTable
property), 132

TransInfo (class in psyclone.psyGen), 139

V

validate() (psyclone.transformations.ACCEnterDataTrans
method), 185

validate() (psyclone.transformations.ACCEnterDataTrans
method), 186

validate() (psyclone.transformations.ACCKernelsTrans
method), 187

validate() (psyclone.transformations.ACCEnterDataTrans
method), 190

validate() (psyclone.transformations.Dynamo0p3AsyncHaloExchange
method), 191

validate() (psyclone.transformations.Dynamo0p3KernelConstTrans
method), 194

validate() (psyclone.transformations.Dynamo0p3RedundantComputati
method), 195

validate() (psyclone.transformations.DynamoLoopFuseTrans
method), 197

validate() (psyclone.transformations.GOOceanLoopFuseTrans
method), 200

validate() (psyclone.transformations.GOLoopSwapTrans
method), 199

validate() (psyclone.transformations.KernelGlobalsToArguments
method), 202

validate() (psyclone.transformations.KernelTrans
static method), 203

validate() (psyclone.transformations.LoopFuseTrans
method), 203

validate() (psyclone.transformations.MoveTrans
method), 205

validate() (psyclone.transformations.OCLTrans
method), 206

validate() (psyclone.transformations.OMPParallelLoopTrans
method), 208

validate() (psyclone.transformations.OMPParallelTrans
method), 209

`validate()` (*psyclone.transformations.ParallelLoopTrans*
method), [210](#)
`validate()` (*psyclone.transformations.ParallelRegionTrans*
method), [210](#)
`view()` (*psyclone.psyir.symbols.SymbolTable* *method*),
[132](#)
`visibility()` (*psyclone.psyir.symbols.Symbol* *prop-*
erty), [134](#)

W

`walk()` (*psyclone.psyir.nodes.Node* *method*), [125](#)