PSyclone Documentation

Release 1.7.0

Rupert Ford, Joerg Henrichs, Iva Kavcic, Andrew Porter and Sergi Siso

May 14, 2019

CONTENTS

PSyclone, the PSy code generator, is being developed for use in finite element, finite volume and finite difference codes. PSyclone development started with the aim to support the emerging API in the GungHo project for a finite element dynamical core.

The GungHo project was initiated in 2011 to address challenges of weather and climate prediction on the next generation of supercomputers. The project ran for 5 years as a collaboration between the Met Office, NERC (via NERC funded academics) and STFC. It laid a foundation for redesign of the heart of the Met Office's Unified Model, known as the dynamical core, from the choices of numerical methods and model grids to the implementation of parallel algorithms that will scale to millions of cores.

The software infrastructure based on the GungHo project recommendations is now being developed in the LFRic project and is expected to be operational in 2022. Its development is led by the requirements to support multiple meshes and element types, thus allowing for future model development.

GungHo also proposed a novel separation of concerns for the software implementation of the dynamical core. This approach distinguishes between three layers: the Algorithm layer, the Kernel layer and the Parallelisation System (PSy) layer. Together this separation is termed PSyKAl.

The Algorithm layer specifies the algorithm that the scientist would like to run (in terms of calls to kernel routines and built-in operations) and logically operates on full fields.

The Kernel layer provides the implementation of the code kernels as subroutines. These subroutines operate on local fields (a set of elements, a vertical column, or a set of vertical columns, depending on the kernel).

The PSy layer sits in-between the algorithm and kernel layers and its primary role is to provide node-based parallel performance for the target architecture. The PSy layer can be optimised for a particular hardware architecture, such as multi-core, many-core, GPGPUs, or some combination thereof with no change to the algorithm or kernel layer code. This approach therefore offers the potential for portable performance.

Rather than writing the PSy layer manually, the PSyclone code generation system can help a user to optimise the code for a particular architecture (by providing optimisations such as blocking, loop merging, inlining etc), or alternatively, generate the PSy layer automatically.

PSyclone is also being extended to support an API being developed in the GOcean project for two finite difference ocean model benchmarks, one of which is based on the NEMO ocean model.

GETTING GOING

1.1 Download

The following instructions are intended for a PSyclone user who wants to work with a released version of the code. If you are a developer or wish to test a specific branch of PSyclone from the GitHub repository please see *Installation* in the *Developers' guide*.

PSyclone is available on the Python Package Index (pypi.org) and is hosted on GitHub:

```
https://github.com/stfc/PSyclone
```

The latest release is 1.7.0 and the latest stable version is on the master branch.

PSyclone can be installed using pip:

```
> pip install psyclone
```

or downloaded from github - either see 1.7.0 in the `releases` tab on the PSyclone page or download and extract the latest release of PSyclone directly, e.g.

```
> wget https://github.com/stfc/PSyclone/archive/1.7.0.tar.gz
> tar zxf 1.7.0.tar.gz
> ls
PSyclone-1.7.0
```

Hereon the location where you download or clone PSyclone (including the PSyclone directory itself) will be referred to as `<PSYCLONEHOME>`.

1.2 Dependencies

PSyclone is written in Python so needs Python to be installed on the target machine. PSyclone has been tested under Python 2.7 and 3.6.

> **Warning:** Release 1.6.0 of PSyclone requires version 0.0.7 of fparser and will fail on more recent versions. However simply installing this version using pip will install a later version of fparser. The suggested solution is to use release 1.6.1 which fixes this problem and is otherwise identical to 1.6.0. However, if you want to use release 1.6.0 you need to ensure that you have fparser 0.0.7. This can be achieved by downgrading an existing fparser installation or by installing version 0.0.7 of fparser *before* installing PSyclone.

PSyclone immediately relies on four external Python packages; `six`, `configparser`, `fparser` and `pyparsing`. In order to run the test suite `py.test` is required. The easiest way to satisfy the Python dependencies is to use the Python Package Index (pypi.org) and `pip`. See https://packaging.python.org/installing/ for more information.

If everything is working correctly then using pip to install PSyclone:

```
> pip install psyclone
```

will automatically install the Python dependencies (`fparser` and `pyparsing`).

In addition to the mandatory dependencies just described, PSyclone also has optional dependencies on both `graphviz` and `termcolor`. PSyclone can use graphviz to produce a visualisation of a schedule's dependency graph. If this is desired then the Python package `graphviz` (for the Python bindings) as well as the graphviz package itself must be installed. If the `graphviz` package is not available then the associated PSyclone routines will return silently and no visualisations will be produced. The Python package `termcolor` is used for pretty-printing a schedule in terminals that support coloured text. If the package is not available then the schedule is simply printed in plain text without colour highlighting.

1.2.1 System-specific set-up

System-specific set-up instructions are available for Ubuntu 14.04.2 and OpenSUSE 42.2.

1.2.2 fparser

The fparser package (https://github.com/stfc/fparser) is a Fortran parser originally developed as a part of the f2py project.

The minimum version of fparser required by PSyclone is currently 0.0.8 but we strongly recommend you install the latest version to reduce the chance of encountering problems when parsing existing algorithm or kernel code. (Note that for older versions of PSyclone up to and including 1.5.1 you must use version 0.0.6 of fparser.)

fparser is available from the Python Package Index and thus may be installed using `pip` (https://packaging.python.org/installing/#requirements-for-installing-packages):

```
> pip install fparser
```

If you do not have sufficient permissions to perform a system-wide install then you can instruct pip to do a user-local install:

```
> pip install --user fparser
```

Should you wish to remove fparser then simply do:

```
> pip uninstall fparser
```

If you have already installed fparser and want to upgrade to the latest version simply do:

```
> pip install fparser --upgrade
```

1.2.3 pyparsing

PSyclone requires pyparsing, a library designed to allow parsers to be be built in Python. PSyclone uses pyparsing to parse fortran regular expressions as fparser does not fully parse these, (see http://pyparsing.wikispaces.com for more information).

PSyclone has been tested with pyparsing versions 1.5.2, 2.0.1 and 2.2.0.

You can test whether pyparsing is already installed on your machine by typing `import pyparsing` from the python command line. If pyparsing is installed, this command will complete succesfully. If pyparsing is installed you can check its version by typing `pyparsing.__version__` after succesfully importing it.

If pyparsing is not installed on your system then it may be installed from the Python Package Index using `pip`:

```
> pip install pyparsing
```

Should you wish to, uninstalling is simply performed by doing:

```
> pip uninstall pyparsing
```

If you do not have sufficient privileges for a system-wide install then you can instruct pip to do a user-local install:

```
> pip install --user pyparsing
```

Alternatively, you could follow the instructions here http://pyparsing.wikispaces.com/Download+and+Installation.

1.2.4 graphviz

The data dependencies of a PSyIR schedule determine the validity of changes to this schedule. PSyclone supports the visualisation of these dependencies as a graph using graphviz. This visualisation is not needed to use PSyclone.

If the Python bindings to graphviz are not installed on your system then it may be installed from the Python Package Index using `pip`:

```
> sudo pip install graphviz
```

Should you wish to, uninstalling is simply performed by doing:

```
> sudo pip uninstall graphviz
```

If you do not have sufficient privileges for a system-wide install then you can instruct pip to do a user-local install:

```
> pip install --user graphviz
```

If graphviz itself is not installed on your system and your system supports the `apt` package manager then see below, otherwise please refer to the download and install instructions which are available here http://www.graphviz.org/Download..php.

If your system supports the `apt` package manager then it can be installed and removed in the following way:

```
> sudo apt install graphviz
> sudo apt remove graphviz
```

1.2.5 termcolor

By default, the `view()` method available on any PSyIR (PSyclone Internal Representation) object prints a plain-text representation to standard-out. However, if the `termcolor` package is available then PSyclone uses this to add colour highlighting to the output text.

Installation (and uninstallation) of this package can be done via `pip` in exactly the same way as for graphviz, as described above.

1.2.6 py.test

The PSyclone test suite uses py.test. This is not needed to use PSyclone but is useful to check whether PSyclone is working correctly on your system. You can test whether it is already installed by simply typing `py.test` at a shell prompt. If it is present you will get output that begins with

```
======================= test session starts =================
```

If you do not have it then py.test can again be installed using `pip` or from here http://pytest.org/latest/ (or specifically here http://pytest.org/latest/getting-started.html).

1.3 Environment

In order to use PSyclone (including running the test suite and building documentation) you will need to install it. The simplest way to do this is to use pip with the supplied `setup.py` file:

```
> cd <PSYCLONEHOME>
> pip install .
```

By default pip will attempt a system-wide install. If you wish to do a user-local install instead then supply the `--user` flag:

```
> pip install --user .
```

This installs the PSyclone modules in ~/.local/lib/pythonX.Y/site-packages (where X.Y is the version of Python that you are using) and the 'psyclone' script in ~/.local/bin. Depending on your linux distribution, you may need to add the latter location to your $PATH.

If for some reason you'd rather not use pip then you can run the setup manually:

```
> python setup.py install
```

or, if you don't have root access:

```
> python setup.py install --user
```

or,

```
> python setup.py install --prefix /my/install/path
```

If using the latter method then it will be necessary to take further action to ensure PSyclone can find the configuration file installed as a part of this process - see below.

1.4 Configuration

Various aspects of PSyclone are configured through a configuration file, `psyclone.cfg`. The default version of this file is installed to `<python-base-prefix>/shared/psyclone/` during the installation process. If a system-wide installation is being performed then this will be something like `/usr/share/psyclone/`. If a user-local installation is performed (`--user` flag to `pip install`) then the location will be something like `~/.local/share/psyclone/`.

> **Warning:** if PSyclone is installed to a non-standard location (e.g. by specifying the `--prefix=/some/path` option to `pip install`) then PSyclone will not be able to find the configuration file at execution time. There are two solutions to this: 1. copy the configuration file to a location where PSyclone will find it (see *Configuration*) or 2. set the `PSYCLONE_CONFIG` environment variable to the full-path to the configuration file, e.g.:
>
> ```
> > export PSYCLONE_CONFIG=/some/path/PSyclone/config/psyclone.cfg
> ```

> **Warning:** when installing in 'editable' mode (`-e` flag to pip), pip does *not* install the configuration file. You will have to take one of the two actions described above.

See *Configuration* for details of the settings contained within the config file.

1.5 Test

You can install the necessary dependencies to run the PSyclone tests with:

```
> pip install psyclone[test]
```

The test dependencies are canonically documented in PSyclone's setup.py under the `extras_requires` section.

Once you have the necessary dependencies installed and your environment configured, you can check that things are working by using the PSyclone test suite. These tests are not required and can be skipped if preferred:

```
> cd <PSYCLONEHOME>/src/psyclone/tests
> py.test
```

If everything is working as expected then you should see output similar to:

```
=========================== test session starts ===============================
platform linux2 -- Python 2.6.5 -- py-1.4.29 -- pytest-2.7.2
rootdir: /home/rupert/proj/GungHoSVN/PSyclone_r3373_scripts/src/tests, inifile:
collected 175 items

alggen_test.py .......xxxxxxxxxxx.
dynamo0p1_transformations_test.py .
dynamo0p3_test.py .....................................x
generator_test.py ...................
gocean0p1_transformations_test.py .......
gocean1p0_test.py ....
gocean1p0_transformations_test.py .....................x........
parser_test.py ..........
psyGen_test.py ..............................

================== 160 passed, 15 xfailed in 13.59 seconds ====================
```

Most of the tests use Fortran source files in the `<PSYCLONEHOME>/src/psyclone/tests` directory and many of them can be compiled during the testing process. To enable compilation testing run:

```
> py.test --compile --f90="<compiler_name>" --f90flags="<compiler_flags_list>"
```

`"<compiler_name>"` and `"<compiler_flags_list>"` are optional arguments. The default value for `"<compiler_name>"` is `"gfortran"` and there are no defaults for the `"<compiler_flags_list>"`. Please note that the onus is on the user to provide correct values for these options. Full details are provided in the *Compilation testing* section.

1.6 Run

You are now ready to try running PSyclone on the examples. One way of doing this is to use the `psyclone` driver script. Assuming it is on your PATH:

```
> psyclone
usage: psyclone [-h] [-oalg OALG] [-opsy OPSY] [-okern OKERN] [-api API]
                [-s SCRIPT] [-d DIRECTORY] [-I INCLUDE] [-l] [-dm] [-nodm]
                [--kernel-renaming {multiple,single}]
                [--profile {invokes,kernels}]
                [--force-profile {invokes,kernels}] [-v]
                filename
psyclone: error: too few arguments
```

As indicated above, the psyclone script takes the name of the Fortran source file containing the algorithm specification (in terms of calls to invoke()). It parses this, finds the necessary kernel source files and produces two Fortran files. The first contains the PSy, middle layer and the second a re-write of the algorithm code to use that layer. These files are named according to the user-supplied arguments (options -oalg and -opsy). If those arguments are not supplied then the script writes the generated/re-written Fortran to the terminal. For details of the other command-line arguments please see the *PSyclone Script* Section.

Examples are provided in the examples directory. There are 3 subdirectories (dynamo, gocean and nemo) corresponding to different APIs that are supported by PSyclone. In this case we are going to use one of the dynamo examples:

```
> cd <PSYCLONEHOME>/examples/dynamo/eg1
> psyclone -api dynamo0.1 \
> -oalg dynamo_alg.f90 -opsy dynamo_psy.f90 dynamo.F90
```

You should see two new files created called dynamo_alg.f90 (containing the re-written algorithm layer) and dynamo_psy.f90 (containing the generated PSy- or middle-layer). Since this is a dynamo example the Fortran source code has dependencies on the dynamo system and therefore cannot be compiled stand-alone.

You can also use the runme.py example to see the interactive API in action. This script contains:

```python
from psyclone.parse.algorithm import parse
from psyclone.psyGen import PSyFactory

# This example uses version 0.1 of the Dynamo API
api = "dynamo0.1"

# Parse the file containing the algorithm specification and
# return the Abstract Syntax Tree and invokeInfo objects
ast, invokeInfo = parse("dynamo.F90", api=api)

# Create the PSy-layer object using the invokeInfo
psy = PSyFactory(api).create(invokeInfo)
# Generate the Fortran code for the PSy layer
print(psy.gen)

# List the invokes that the PSy layer has
print(psy.invokes.names)

# Examine the 'schedule' (e.g. loop structure) that each
# invoke has
schedule = psy.invokes.get('invoke_0_v3_kernel_type').schedule
schedule.view()

schedule = psy.invokes.get('invoke_1_v3_solver_kernel_type').schedule
schedule.view()
```

It can be run non-interactively as follows:

```
> cd <PSYCLONEHOME>/example/dynamo/eg1
> python runme.py
```

However, to understand this example in more depth it is instructive to cut-and-paste from the runme.py file into your own, interactive python session:

```
> cd <PSYCLONEHOME>/example/dynamo/eg1
> python
```

In addition to the runme.py script, there is also runme_openmp.py which illustrates how one applies an OpenMP transform to a loop schedule within the PSy layer. The initial part of this script is the same as that of runme.py (above) and is therefore omitted here:

```python
# List the various invokes that the PSy layer contains
print(psy.invokes.names)

# Get the loop schedule associated with one of these
# invokes
schedule = psy.invokes.get('invoke_v3_kernel_type').schedule
schedule.view()

# Get the list of possible loop transformations
from psyclone.psyGen import TransInfo
t = TransInfo()
print(t.list)

# Create an OpenMPLoop-transformation object
ol = t.get_trans_name('OMPLoopTrans')

# Apply it to the loop schedule of the selected invoke
new_schedule, memento = ol.apply(schedule.children[0])
new_schedule.view()

# Replace the original loop schedule of the selected invoke
# with the new, transformed schedule
psy.invokes.get('invoke_v3_kernel_type')._schedule = new_schedule
# Generate the Fortran code for the new PSy layer
print(psy.gen)
```

EXAMPLES

Various examples of the use of PSyclone are provided under the `examples` directory. All of these examples require that PSyclone be installed on the host system, see *Getting Going*. This document is intended to provide an overview of the various examples so that a user can find one that is appropriate to them. For details of how to run each example please see the README files in the associated directories.

2.1 GOcean

2.1.1 Example 1: Loop transformations

Examples of applying various transformations (loop fusion, OpenMP, OpenACC) to the semi-PSyKAl'd version of the Shallow benchmark. ("semi" because not all kernels are called from within invoke()'s.) Also includes an example of generating a DAG from an InvokeSchedule.

2.1.2 Example 2: OpenACC

This is a simple but complete example of using PSyclone to enable an application to run on a GPU by adding OpenACC directives. A Makefile is included which will use PSyclone to generate the PSy code and transformed kernels and then compile the application. This compilation requires that the dl_esm_inf library (github.com/stfc/dl_esm_inf) be installed.

2.1.3 Example 3: OpenCL

Example of the use of PSyclone to generate an OpenCL version of the PSy layer. The kernels are not yet transformed automatically (Issue #249).

2.1.4 Example 4: Kernels containing use statements

Transforming kernels for use with either OpenACC or OpenCL requires that we handle those that access data and/or routines via module *use* statements. This example shows the various forms for which support is planned (Issues #323 and #342).

2.2 Dynamo

Examples 1 and 2 are for the (deprecated) Dynamo 0.1 API. The remaining examples are all for the Dynamo 0.3 API.

2.2.1 Example 1

Basic operation of PSyclone with invoke()'s containing just one kernel for the Dynamo 0.1 API. Also includes an example of transforming for OpenMP.

2.2.2 Example 2

A more complex example for the Dynamo 0.1 API containing multi-kernel invokes. Provides examples of OpenMP and loop fusion transformations.

2.2.3 Example 3

Shows the use of colouring and OpenMP for the Dynamo 0.3 API. Includes multi-kernel, named invokes with both user-supplied and built-in kernels.

2.2.4 Example 4

Demonstrates the use of the special `enforce_bc_kernel` which PSyclone recognises as a boundary-condition kernel.

2.2.5 Example 5

Example of kernels which require stencil information.

2.2.6 Example 6

Example of applying OpenMP to an InvokeSchedule containing kernels that perform reduction operations. Two scripts are provided, one of which demonstrates how to request that PSyclone generate code for a reproducible OpenMP reduction. (The default OpenMP reduction is not guaranteed to be reproducible from one run to the next on the same number of threads.)

2.2.7 Example 7

Example of kernels requiring Column-Matrix Assembly operators.

2.2.8 Example 8

Example of the use of the redundant-computation and move transformations to eliminate and re-order halo exchanges.

2.2.9 Example 9

Demonstrates the behaviour of PSyclone for kernels that read and write quantities on horizontally-discontinuous function spaces. In addition, this example demonstrates how to write a PSyclone transformation script that only colours loops over continuous spaces.

2.2.10 Example 10

Demonstrates the use of "inter-grid" kernels that prolong or restrict fields (map between grids of different resolutions).

2.2.11 Example 11

Example of the use of transformations to introduce redundant computation, split synchronous halo exchanges into asynchronous exchanges (start and stop) and move the starts of those exchanges in order to overlap them with computation.

2.2.12 Example 12

Example of applying code extraction to Nodes in an Invoke Schedule:

```
> psyclone -nodm -s ./extract_nodes.py \
    gw_mixed_schur_preconditioner_alg_mod.x90
```

or to a Kernel in an Invoke after applying transformations:

```
> psyclone -nodm -s ./extract_kernel_with_transformations.py \
    gw_mixed_schur_preconditioner_alg_mod.x90
```

For now it only inserts comments in appropriate locations while the the full support for code extraction is being developed.

This example also contains a Python helper script `find_kernel.py` which displays the names and Schedules of Invokes containing call(s) to the specified Kernel:

```
> python find_kernel.py
```

2.3 NEMO

These examples may all be found in the `PSyclone/examples/nemo` directory.

2.3.1 Example 1: OpenMP parallelisation of tra_adv

Demonstrates the use of PSyclone to parallelise the loops over vertical levels in a NEMO tracer-advection benchmark using OpenMP.

2.3.2 Example 2: OpenMP parallelisation of traldf_iso

Demonstrates the use of PSyclone to parallelise the loops over vertical levels in some NEMO tracer-diffusion code using OpenMP.

2.3.3 Example 3: OpenACC parallelisation of tra_adv

Demonstrates the introduction of simple OpenACC parallelisation (using the `data` and `kernels` directives) for a NEMO tracer-advection benchmark.

PSYCLONE SCRIPT

The simplest way to run PSyclone is to use the `psyclone` script. If you installed PSyclone using `pip` then this script should be available on your PATH (see *Environment* for more details). Alternatively it can be found in the `<PSYCLONEHOME>/bin` directory. The script takes an algorithm file as input and outputs modified algorithm code and generated PSy code. This section walks through its functionality.

3.1 Running

The `psyclone` script is executable and designed to be run from the command line, e.g.:

```
> psyclone <args>
```

The optional `-h` argument gives a description of the options provided by the script:

```
> psyclone -h

usage: psyclone [-h] [-oalg OALG] [-opsy OPSY] [-okern OKERN] [-api API]
                [-s SCRIPT] [-d DIRECTORY] [-I INCLUDE] [-l] [-dm] [-nodm]
                [--kernel-renaming {multiple,single}]
                [--profile {invokes,kernels}]
                [--force-profile {invokes,kernels}] [-v] filename

Run the PSyclone code generator on a particular file

positional arguments:
  filename              algorithm-layer source code

optional arguments:
  -h, --help            show this help message and exit
  -oalg OALG            filename of transformed algorithm code
  -opsy OPSY            filename of generated PSy code
  -okern OKERN          directory in which to put transformed kernels
  -api API              choose a particular api from ['dynamo0.1',
                        'dynamo0.3', 'gocean0.1', 'gocean1.0', 'nemo'],
                        default 'dynamo0.3'.
  -s SCRIPT, --script SCRIPT
                        filename of a PSyclone optimisation script
  -d DIRECTORY, --directory DIRECTORY
                        path to root of directory structure containing kernel
                        source code
  -I INCLUDE, --include INCLUDE
                        path to Fortran INCLUDE files (nemo API only)
```

(continues on next page)

(continued from previous page)

```
-l, --limit           limit the fortran line length to 132 characters
-dm, --dist_mem       generate distributed memory code
-nodm, --no_dist_mem  do not generate distributed memory code
--kernel-renaming {single,multiple}
                      Naming scheme to use when re-naming transformed
                      kernels.
--profile {invokes,kernels}, -p {invokes,kernels}
                      Add profiling hooks for either 'kernels' or 'invokes'
--force-profile {invokes,kernels}
                      Add profiling hooks for either 'kernels' or 'invokes'
                      even if a transformation script is used. Use at your
                      own risk.
-v, --version         Display version information (1.6.0)
```

3.2 Basic Use

The simplest way to use `psyclone` is to provide it with an algorithm file:

```
> psyclone alg.f90
```

If the algorithm file is invalid for some reason, the script should return with an appropriate error. For example, if we use the Python `genkernelstub` script as an algorithm file we get the following:

```
> psyclone <PSYCLONEHOME>/bin/genkernelstub
...
    1:#!/usr/bin/env python <== no parse pattern found for "#" in 'BeginSource' block.
'Parse Error: Fatal error in external fparser tool'
```

If the algorithm file is valid then the modified algorithm code and the generated PSy code will be output to the terminal screen.

3.3 Choosing the API

In the previous section we relied on PSyclone using the default API. The default API, along with the supported APIs can be seen by running the `psyclone` script with the `-h` option.

If you use a particular API frequently and it is not the default then you can change the default by creating a copy of the default `psyclone.cfg` file and editing it. See *Configuration* for more details.

If your code uses an API that is different to the default then you can specify this as an argument to the `psyclone` script.

```
> psyclone -api dynamo0.1 alg.f90
```

3.4 File output

By default the modified algorithm code and the generated PSy code are output to the terminal. These can instead be output to files by using the `-oalg <file>` and `-opsy <file>` options, respectively. For example, the following will output the generated PSy code to the file 'psy.f90' but the algorithm code will be output to the terminal:

```
> psyclone -opsy psy.f90 alg.f90
```

If PSyclone is being used to transform Kernels then the location to write these to is specified using the `-okern <directory>` option. If this is not supplied then they are written to the current working directory. By default, PSyclone will overwrite any kernel of the same name in that directory. To change this behaviour, the user can use the `--no_kernel_clobber` option. This causes PSyclone to re-name any transformed kernel that would clash with any of those already present in the output directory.

3.5 Algorithm files with no invokes

If the `psyclone` script is provided with a file that contains no `invoke` calls then the script outputs a warning to `stdout` and copies the input file to `stdout`, or to the specified algorithm file (if the `-oalg <file>` option is used). No PSy code will be output. If a file is specified using the `-opsy <file>` option this file will not be created.

```
> psyclone -opsy psy.f90 -oalg alg_new.f90 empty_alg.f90
Warning: 'Algorithm Error: Algorithm file contains no invoke() calls: refusing to
generate empty PSy code'
```

3.6 Kernel search directory

When an algorithm file is parsed, the parser looks for the associated kernel files. The way in which this is done requires that any user-defined kernel routine (as opposed to *Built-ins*) called within an invoke must have an explicit use statement. For example, the following code gives an error:

```
> cat no_use.f90
program no_use
  call invoke(testkern_type(a,b,c,d,e))
end program no_use
> psyclone -api gocean1.0 no_use.f90
"Parse Error: kernel call 'testkern_type' must either be named in a use statement or
↪be a recognised built-in (one of '[]' for this API)"
```

(If the chosen API has any *Built-ins* defined then these will be listed within the `[]` in the above error message.) If the name of the kernel is provided in a use statement then the parser will look for a file with the same name as the module in the use statement. In the example below, the parser will look for a file called "testkern.f90" or "testkern.F90":

```
> cat use.f90
program use
  use testkern, only : testkern_type
  call invoke(testkern_type(a,b,c,d,e))
end program use
```

Therefore, for PSyclone to find kernel files, the module name of a kernel file must be the same as its filename. By default the parser looks for the kernel file in the same directory as the algorithm file. If this file is not found then an error is reported.

```
> psyclone use.f90
Kernel file 'testkern.[fF]90' not found in <location>
```

The `-d` option can be used to tell `psyclone` where to look for kernel files by supplying it with a directory. The script will recurse from the specified directory path to look for the required file. There must be only one instance of the specified file within (or below) the specified directory:

```
> cd <PSYCLONEHOME>/src/psyclone
> psyclone -d . use.f90
More than one match for kernel file 'testkern.[fF]90' found!
> psyclone -d tests/test_files/dynamo0p3 -api dynamo0.3 use.f90
[code output]
```

Note: The `-d` option is limited to a single directory. Therefore a current limitation in PSyclone is that all kernel files required by an algorithm file must exist within a directory hierarchy where their file names are unique.

3.7 Transformation script

By default the `psyclone` script will generate 'vanilla' PSy layer code. The -s option allows a Python script to be specified which can transform the PSy layer. This option is discussed in more detail in the *Script* section.

3.8 Fortran line length

By default the `psyclone` script will generate fortran code with no consideration of Fortran line-length limits. As the line-length limit for free-format Fortran is 132 characters, the code that is output may be non-conformant.

Line length is not an issue for many compilers as they allow compiler flags to be set which allow lines longer than the Fortran standard. However this is not the case for all compilers.

When the `-l` option is specified to the `psyclone` script, the output will be line wrapped so that the output lines are always within the 132 character limit.

The `-l` option also checks the parsed algorithm and kernel files for conformance and raises an error if they do not conform.

Line wrapping is not performed by default. There are two reasons for this. This first reason is that most compilers are able to cope with long lines. The second reason is that the line wrapping implementation could fail in certain pathological cases. The implementation and limitations of line wrapping are discussed in the *Limitations* section.

3.9 Distributed memory

By default the `psyclone` script will generate distributed memory (DM) code (i.e. parallelised using MPI). As with the choice of API, this default may be configured by editing `psyclone.cfg` - see *Configuration*. Alternatively, whether or not to generate DM code can be specified as an argument to the `psyclone` script using the `-dm`/`--dist_mem` or `-nodm`/`--no_dist_mem` flags, respectively.

For details of PSyclone's support for generating DM code see *Distributed Memory*.

3.10 Automatic Profiling Instrumentation

The `--profile` option allows the user to instruct PSyclone to automatically insert profiling calls within the generated PSy code. Two options are provided, `invokes` and `kernels`. The first of these causes PSyclone to insert profiling-start and -stop calls at the beginning and end of every generated invoke routine. The second puts profiling calls around every kernel call (including the associated loops). The generated code must be linked against the PSyclone profiling interface and the profiling tool itself. The application that calls the PSyclone-generated code is responsible for

initialising and finalising the profiling library that is being used. For full details on the use of this profiling functionality please see the *Profiling* section.

3.11 Outputting of Transformed Kernels

When transforming kernels there are two use-cases to consider:

1. a given kernel will be transformed only once and that version then used from multiple, different Invokes and Algorithms;

2. a given kernel is used from multiple, different Invokes and Algorithms and is transformed differently, depending on the Invoke.

Whenever PSyclone is used to transform a kernel, the new kernel must be re-named in order to avoid clashing with other possible calls to the original. By default (`--kernel-renaming multiple`), PSyclone generates a new, unique name for each kernel that is transformed. Since PSyclone is run on one Algorithm file at a time, it uses the chosen kernel output directory (`-okern`) to ensure that names created by different invocations do not clash. Therefore, when building a single application, the same kernel output directory must be used for each separate invocation of PSyclone.

Alternatively, in order to support use case 1, a user may specify `--kernel-renaming single`: now, before transforming a kernel, PSyclone will check the kernel output directory and if a transformed version of that kernel is already present then that will be used. Note, if the kernel file on disk does not match with what would be generated then PSyclone will raise an exception.

3.12 Fortran INCLUDE Files

For the NEMO API, if the source code to be processed by PSyclone contains INCLUDE statements (other than those for libraries such as MPI) then the location of any INCLUDE'd files must be supplied to PSyclone via the `-I` or `--include` option. (This is necessary because INCLUDE lines are a part of the Fortran language and must therefore be parsed - they are not handled by any pre-processing step.) Multiple locations may be specified by using multiple `-I` flags, e.g.:

```
> psyclone api "nemo" -I /some/path -I /some/other/path alg.f90
```

If no include paths are specified then the directory containing the source file currently being parsed is searched by default. If the specified include file is not found then ideally the INCLUDE line would be left unchanged. However, fparser currently treats any such INCLUDE lines as comments which results in them being lost (fparser issue #138). The workaround for this is to ensure that the location of *all* INCLUDE files is supplied to PSyclone.

Attempting to specify `-I`/`--include` for any API other than NEMO will be rejected by PSyclone.

KERNEL LAYER

In the PSyKAl separation of concerns, Kernel code (code which is created to run within the Kernel layer), works over a subset of a field (such as a column). The reason for doing this is that it gives the PSy layer the responsibility of calling the Kernel over the spatial domain which is where parallelism is typically exploited in finite element and finite difference codes. The PSy layer is therefore able to call the kernel layer in a flexible way (blocked and/or in parallel for example). Kernel code in the kernel layer is not allowed to include any parallelisation calls or directives and works on raw fortran arrays (to allow the compiler to optimise the code).

Since a Kernel is called over the spatial domain (by the PSy layer) it must take at least one field or operator as an argument.

4.1 API

Kernels in the kernel layer are implemented as subroutines within fortran modules. One or more kernel modules are allowed, each of which can contain one or more kernel subroutines. In the example below there is one module `integrate_one_module` which contains one kernel subroutine `integrate_one_code`. The kernel subroutines contain the code that operates over a subset of the field (such as a column).

Metadata describing the kernel subroutines is required by the PSyclone system to generate appropriate PSy layer code. The metadata is written by the kernel developer and is kept with the kernel code in the same module using a sub-type of the `kernel_type` type. In the example below the `w3_solver_kernel_type` type specifies the appropriate metadata information describing the kernel code for the `dynamo0.3` api:

```
module w3_solver_kernel_mod

  use kernel_mod,     only : kernel_type
  use constants_mod,  only : r_def
  use argument_mod,   only : arg_type, func_type,                &
                             GH_FIELD, GH_READ, GH_WRITE, W0, W3, &
                             GH_BASIS, GH_DIFF_BASIS, CELLS
  implicit none

  type, public, extends(kernel_type) :: w3_solver_kernel_type
    private
    type(arg_type) :: meta_args(4) = (/                &
         arg_type(GH_FIELD,   GH_WRITE, W3),           &
         arg_type(GH_FIELD,   GH_READ,  W3),           &
         arg_type(GH_FIELD*3, GH_READ,  W0),           &
         arg_type(GH_REAL,    GH_READ)                 &
         /)
    type(func_type) :: meta_funcs(2) = (/              &
         func_type(W3, GH_BASIS),                      &
```

```
        func_type(W0, GH_DIFF_BASIS)                       &
        /)
   integer :: gh_shape = gh_quadrature_XYoZ
   integer :: iterates_over = CELLS
 contains
   procedure, nopass :: solver_w3_code
 end type

contains

 subroutine solver_w3_code(nlayers,                                &
                   x, rhs,                                &
                   chi_1, chi_2, chi_3, ascalar,          &
                   ndf_w3, undf_w3, map_w3, w3_basis,     &
                   ndf_w0, undf_w0, map_w0, w0_diff_basis, &
                   nqp_h, nqp_v, wqp_h, wqp_v)
   ...
 end subroutine solver_w3_code

end module w3_solver_kernel_mod
```

4.2 Metadata

Kernel metadata is not required if the PSy layer is going to be written manually - its sole purpose is to let PSyclone know how to generate the PSy layer. The content of Kernel metadata differs depending on the particular API and this information can be found in the API-specific sections of this document.

In all APIs the kernel metadata is implemented as an extension of the *kernel_type* type. The reason for using a type to specify metadata is that it allows the metadata to be kept with the code and for it to be compilable. In addition, currently all APIs will contain information about the arguments in an array called `meta_args`, a specification of what the kernel code iterates over in a variable called `iterates_over` and a reference to the kernel code as a type bound procedure.

```
type, extends(kernel_type) :: integrate_one_kernel
  ...
  type(...) :: meta_args(...) = (/ ... /)
  ...
  integer :: ITERATES_OVER = ...
  ...
  contains
  ...
  procedure ...
  ...
end type integrate_one_kernel
```

ALGORITHM LAYER

In the PSyKAl separation of concerns, the Algorithm layer specifies the algorithm that the scientist would like to run (in terms of calls to kernel routines and Built-in operations) and logically operates on full fields. Algorithm code in the algorithm layer is not allowed to include any parallelisation calls or directives and passes datatypes specified by the particular API.

5.1 API

The Algorithm layer is forbidden from calling the Kernel layer directly. In PSyclone, if the programmer would like to call a Kernel routine or a Built-in operation from the algorithm layer they must use the `invoke` call (which is common to all APIs). The `invoke` call is not necessary (and indeed will not work) if the PSy layer is written manually.

To make an `invoke` call, the algorithm layer developer adds one or more `call invoke()` statements to their code and within the content of the `invoke` call they add a reference to the required Kernel/Built-in and the data to pass to it. For example,

```
...
call invoke(integrate_one_kernel(arg1,arg2))
...
```

For more information on the concept of Built-in operations see the *Built-ins* Section. Details of which operations are supported for a specific API are given in the documentation of that API.

The algorithm layer can consist of an arbitrary number of files containing fortran code, any of which may contain as many `invoke()` calls as is required. PSyclone is applied to an individual algorithm layer file and must therefore be run multiple times if multiple files containing `invoke()` calls exist in the algorithm layer.

The algorithm developer is also able to reference more than one Kernel/Built-in within an invoke. In fact this feature is encouraged for performance reasons. **As a general guideline the developer should aim to use as few invokes as possible with as many Kernel references within them as is possible**. The reason for this is that it allows for greater freedom for optimisation in the PSy layer as PSy layer optimisations are limited to the contents of individual invoke calls - PSyclone currently does not attempt to optimise the PSy layer over multiple invoke calls.

As well as generating the PSy layer code, PSyclone modifies the Algorithm layer code, replacing `invoke` calls with calls to the generated PSy layer so that the algorithm code is compilable and linkable to the PSy layer and adding in the appropriate `use` statement. For example, the above `integrate_one_kernel` invoke is translated into something like the following:

```
...
use psy, only : invoke_0_integrate_one_kernel
...
call invoke_0_integrate_one_kernel(arg1,arg2)
...
```

You may have noticed from other examples in this guide that an algorithm specification in an invoke call references the metadata `type` in an invoke call, not the `code` directly; this is by design.

For example, in the invoke call below, `integrate_one_kernel` is used.

```
...
call invoke(integrate_one_kernel(arg1,arg2))
...
```

`integrate_one_kernel` is the name of the metadata type in the module, not the name of the subroutine in the Kernel . . .

```
module integrate_one_module
  ...
  type, extends(kernel_type) :: integrate_one_kernel
    ...
  end type
  ...
contains
  ...
  subroutine integrate_one_code(...)
  ...
  end subroutine integrate_one_code
  ...
end module integrate_one_module
```

5.1.1 Named Invokes

PSyclone permits the user to optionally specify a label for an invoke call like so:

```
...
call invoke(integrate_one_kernel(arg1,arg2), &
            name="compute something")
...
```

The `name` argument to the invoke call is optional. If supplied it must be a string literal. The content of this string (with any spaces replaced by '_' characters) is used in naming the corresponding PSy-layer routine generated by PSyclone. So, for the above example, the generated PSy-layer subroutine will be named "invoke_compute_something." Each invoke label must currently be unique within an Algorithm source file. Note that, in keeping with the Fortran language, labels are not case sensitive and, after having any spaces replaced by underscores, must be valid Fortran names (e.g. `name="compute(1)"` is invalid). In the future it is intended that the labelling of invokes will help to support invoke-specific optimisations to be applied as well as enabling more readable profiling output. It may also be used to instruct PSyclone to just generate a single subroutine to implement all invokes that share the same label.

5.2 Limitations

In order to re-write the Algorithm layer, as just described, PSyclone must obviously be able to parse the invoke calls. Since the Fortran expression parser used by PSyclone is relatively simple, this means there are limitations on what Fortran may be used when specifying kernel arguments in an invoke call. Since these limitations can have a direct impact on the natural science code, the PSyclone developers endeavour to keep them to a minimum.

The current list of known limitations/restrictions on the form of kernel arguments within an invoke is:

- No arithmetic expressions (e.g. `kernel_type(a+b)` or `kernel_type(-a)`)

- No named (optional) arguments (e.g. `kernel_type(fn(my_arg=a))`)

If you encounter any other limitations (or have a burning desire to use one of the above forms) then please contact the PSyclone developers.

PSY LAYER

In the PSyKAl separation of concerns, the PSy layer is responsible for linking together the Algorithm and Kernel layers and for providing the implementation of any Built-in operations used. Its functional responsibilities are to

1. map the arguments supplied by an Algorithm `invoke` call to the arguments required by a Built-in or Kernel call (as these will not have a one-to-one correspondance).

2. call any Kernel routines such that they cover the required iteration space and

3. perform any Built-in operations (either by including the necessary code directly in the PSy layer or by e.g. calling a maths library) and

4. include any required distributed memory operations such as halo swaps and reductions.

Its other role is to allow the optimisation expert to optimise any required distributed memory operations, include and optimise any shared memory parallelism and optimise for single node (e.g. cache and vectorisation) performance.

6.1 Code Generation

The PSy layer can be written manually but this is error prone and potentially complex to optimise. The PSyclone code generation system generates the PSy layer so there is no need to write the code manually.

To generate correct PSy layer code, PSyclone needs to understand the arguments and datatypes passed by the algorithm layer and the arguments and datatypes expected by the Kernel layer; it needs to know the name of the Kernel subroutine(s); it needs to know the iteration space that the Kernel(s) is/are written to iterate over; it also needs to know the ordering of Kernels and Built-ins as specified in the algorithm layer. Finally, it needs to know where to place any distributed memory operations.

PSyclone determines the above information by being told the API in question (by the user), by reading the appropriate Kernel and Built-in metadata and by reading the order of Kernels and Built-ins in an invoke call (as specified in the algorithm layer).

PSyclone has an API-specific parsing stage which reads the algorithm layer and all associated Kernel metadata. This information is passed to a PSy-generation stage which creates a high level view of the PSy layer. From this high level view the PSy-generation stage can generate the required PSy code.

For example, the following Python code shows a code being parsed, a PSy-generation object being created using the output from the parser and the PSy layer code being generated by the PSy-generation object.

```python
from psyclone.parse.algorithm import parse
from psyclone.psyGen import PSyFactory

# This example uses version 0.1 of the Dynamo API
api = "dynamo0.1"
```

```
# Parse the file containing the algorithm specification and
# return the Abstract Syntax Tree and invokeInfo objects
ast, invokeInfo = parse("dynamo.F90", api=api)

# Create the PSy-layer object using the invokeInfo
psy = PSyFactory(api).create(invokeInfo)
# Generate the Fortran code for the PSy layer
print psy.gen
```

6.2 API

The PSy-layer of a single algorithm file is represented by the **PSy** class. The PSy class has an **Invokes** object which contain one or more **Invoke** instances (one for each invoke in the algorithm layer). Each **Invoke** has an **InvokeSchedule** object with the PSyIR tree that describes the PSy layer invoke subroutine. This subroutine is called by the Algorithm layer and itself calls one or more kernels and/or implements any required Built-in operations.

All this classes can be specialised in each PSyclone API to support the specific features of the APIs. The class diagram for the above base classes is shown below using the dynamo0.3 API as an illustration. This class diagram was generated from the source code with pyreverse and edited with inkscape.

The InvokeSchedule can currently contain nodes of type: **Loop**, **Kernel**, **Built-in** (see the *Built-ins* section), **Directive** (of various types), **HaloExchange**, or **GlobalSum** (the latter two are only used if distributed memory is supported and is switched on; see the *Distributed Memory* section). The order of the tree (depth first) indicates the order of the associated Fortran code.

PSyclone will initially create a "vanilla" (functionally correct but not optimised) InvokeSchedule. This "vanilla" InvokeSchedule can be modified by changing the objects within it. For example, the order that two Kernel calls appear in the generated code can be changed by changing their order in the tree. The ability to modify this high level view of a InvokeSchedule allows the PSy layer to be optimised for a particular architecture (by applying optimisations such as blocking, loop merging, inlining, OpenMP parallelisation etc.). The tree could be manipulated directly, however, to simplify optimisation, a set of transformations are supplied. These transformations are discussed in the next section.

6.2.1 InvokeSchedule visualisation

PSyclone supports visualising an InvokeSchedule (or any other PSyIR node) in two ways. Firstly the *view()* method outputs textual information about the contents of a PSyIR node. If we were to look at the dynamo eg6 example we would see the following output:

```
>>> schedule.view()
InvokeSchedule[invoke='invoke_0', dm=True]
    Directive[OMP parallel do]
        Loop[type='dofs',field_space='any_space_1',it_space='dofs']
            Call setval_X_code(p,z)
            Call X_innerproduct_Y_code(rs_old,res,z)
    GlobalSum[scalar='rs_old']
```

The above output tells us that the invoke name for the InvokeSchedule we are looking at is *invoke_0* and that the distributed_memory option has been switched on. Within the InvokeSchedule is an OpenMP parallel directive containing a loop which itself contains two builtin calls. As the latter of the two builtin calls requires a reduction and distributed memory is switched on, PSyclone has added a GlobalSum call for the appropriate scalar.

Secondly, the *dag()* method (standing for directed acyclic graph), outputs the PSyIR nodes and its data dependencies. By default a file in dot format is output with the name `dag` and a file in svg format is output with the name `dag.svg`. The file name can be changed using the `file_name` optional argument and the output file format can be changed using the `file_format` optional argument. The file_format value is simply passed on to graphviz so the graphviz documentation should be consulted for valid formats if svg is not required.

```
>>> schedule.dag(file_name="lovely", file_format="png")
```

Note: The dag method can be called from any node and will output the dag for that node and all of its children.

If we were to look at the dynamo eg6 example we would see the following image:

schedule_start

OMP_parallel_do_1_start

loop_[dofs]_0_start

builtin_3   builtin_4

loop_[dofs]_0_end

OMP_parallel_do_1_end

globalsum(rs_old)_1

schedule_end

In the image, all PSyIR nodes with children are split into a start vertex and an end vertex (for example the InvokeSchedule node has both *schedule_start* and *schedule_end* vertices). Blue arrows indicate that there is a parent to child relationship (from a start node) or a child to parent relationship (to an end node). Green arrows indicate that a Node depends on another Node later in the schedule (which we call a forward dependence). Therefore the OMP parallel loop must complete before the globalsum is performed. Red arrows indicate that a Node depends on another Node that is earlier in the schedule (which we call a backward dependence). However the direction of the red arrows are reversed to improve the flow of the dag layout. In this example the forward and backward dependence is the same, however this is not always the case. The two built-ins do not depend on each other, so they have no associated green or red arrows.

The dependence graph output gives an indication of whether nodes can be moved within the InvokeSchedule. In this case it is valid to run the builtin's in either order. The underlying dependence analysis used to create this graph is used to determine whether a transformation of a Schedule is valid from the perspective of data dependencies.

TRANSFORMATIONS

As discussed in the previous section, transformations can be applied to a schedule to modify it. Typically transformations will be used to optimise the PSy layer for a particular architecture, however transformations could be added for other reasons, such as to aid debugging or for performance monitoring.

7.1 Finding

Transformations can be imported directly, but the user needs to know what transformations are available. A helper class **TransInfo** is provided to show the available transformations

class psyclone.psyGen.**TransInfo**(*module=None*, *base_class=None*)

> This class provides information about, and access, to the available transformations in this implementation of PSyclone. New transformations will be picked up automatically as long as they subclass the abstract Transformation class.
>
> For example:

```
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> print(t.list)
There is 1 transformation available:
  1: SwapTrans, A test transformation
>>> # accessing a transformation by index
>>> trans = t.get_trans_num(1)
>>> # accessing a transformation by name
>>> trans = t.get_trans_name("SwapTrans")
```

> **get_trans_name**(*name*)
>
>> return the transformation with this name (use list() first to see available transformations)
>
> **get_trans_num**(*number*)
>
>> return the transformation with this number (use list() first to see available transformations)
>
> **list**
>
>> return a string with a human readable list of the available transformations
>
> **num_trans**
>
>> return the number of transformations available

7.2 Available

Most transformations are generic as the schedule structure is independent of the API, however it often makes sense to specialise these for a particular API by adding API-specific errors checks. Some transformations are API-specific (or specific to a set of APIs e.g. dynamo). Currently these different types of transformation are indicated by their names.

The generic transformations currently available are listed in alphabetical order below (a number of these have specialisations which can be found in the API-specific sections).

Note: PSyclone currently only supports OpenCL transformations for the GOcean 1.0 API and OpenACC transformations for the GOcean 1.0 and NEMO APIs. Attempts to apply these transformations to (members of) schedules from other APIs will be rejected.

---

class psyclone.transformations.**ACCDataTrans**

Add an OpenACC data region around a list of nodes in the PSyIR. COPYIN, COPYOUT and COPY clauses are added as required.

For example:

```
>>> from psyclone.parse import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "NEMO"
>>> filename = "tra_adv.F90"
>>> ast, invokeInfo = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.transformations import ACCDataTrans
>>> dtrans = ACCDataTrans()
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> kernels = schedule.children[0].children[0].children[0:-1]
>>> # Enclose the kernels
>>> new_sched, _ = dtrans.apply(kernels)
```

apply(*node_list*)

Put the supplied list of nodes within an OpenACC data region.

> **Parameters node_list** (list of *psyclone.psyGen.Node*) – The list of PSyIR nodes to enclose in the data region.
>
> **Returns** (transformed schedule, memento of transformation)
>
> **Return type** 2-tuple of (*psyclone.psyGen.Schedule*, psyclone.undoredo.Memento).

---

class psyclone.transformations.**ACCEnterDataTrans**

Adds an OpenACC "enter data" directive to a Schedule. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
```

```
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> dtrans = t.get_trans_name('ACCEnterDataTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> # Add an enter-data directive
>>> newschedule, _ = dtrans.apply(schedule)
>>> newschedule.view()
```

apply(*sched*)

Adds an OpenACC "enter data" directive to the invoke associated with the supplied Schedule. Any fields accessed by OpenACC kernels within this schedule will be added to this data region in order to ensure they remain on the target device.

> **Parameters** **sched** (sub-class of *psyclone.psyGen.Schedule*.) – Schedule to which to add an "enter data" directive.

> **Returns** Tuple of the modified schedule and a record of the transformation.

> **Return type** (*psyclone.psyGen.Schedule*, psyclone.undoredo.Memento)

---

class psyclone.transformations.**ACCKernelsTrans**

Enclose a sub-set of nodes from a Schedule within an OpenACC kernels region (i.e. within "!$acc kernels" ... "!$acc end kernels" directives). Currently only supported for the NEMO API.

For example:

```
>>> from psyclone.parse import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "NEMO"
>>> filename = "tra_adv.F90"
>>> ast, invokeInfo = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.transformations import ACCKernelsTrans
>>> ktrans = ACCKernelsTrans()
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> kernels = schedule.children[0].children[0].children[0:-1]
>>> # Transform the kernel
>>> new_sched, _ = ktrans.apply(kernels)
```

apply(*node_list*, *default_present=False*)

Enclose the supplied list of PSyIR nodes within an OpenACC Kernels region.

> **Parameters**

> - **node_list** (list of *psyclone.psyGen.Node*) – The list of nodes in the PSyIR to enclose.

- **default_present** (*bool*) – whether or not the kernels region should have the 'default present' attribute (indicating that data is already on the accelerator). When using managed memory this option should be False.

> **Returns** (transformed schedule, memento of transformation)
>
> **Return type** 2-tuple of (*psyclone.psyGen.Schedule*, psyclone.undoredo. Memento).

---

class psyclone.transformations.**ACCLoopTrans**

Adds an OpenACC loop directive to a loop. This directive must be within the scope of some OpenACC Parallel region (at code-generation time).

For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory, GenerationError
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ltrans = t.get_trans_name('ACCLoopTrans')
>>> rtrans = t.get_trans_name('ACCParallelTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> new_schedule = schedule
>>>
# Apply the OpenACC Loop transformation to *every* loop
# in the schedule
>>> for child in schedule.children:
>>>     newschedule, memento = ltrans.apply(child, reprod=True)
>>>     schedule = newschedule
>>>
# Enclose all of these loops within a single OpenACC
# PARALLEL region
>>> rtrans.omp_schedule("dynamic,1")
>>> newschedule, memento = rtrans.apply(schedule.children)
>>>
```

apply (*node*, *collapse=None*, *independent=True*, *sequential=False*)

Apply the ACCLoop transformation to the specified node in a GOInvokeSchedule. This node must be a Loop since this transformation corresponds to inserting a directive immediately before a loop, e.g.:

```
!$ACC LOOP
do ...
   ...
end do
```

At code-generation time (when psyclone.psyGen.ACCLoopDirective.gen_code() is called), this node must be within (i.e. a child of) a PARALLEL region.

> **Parameters**

- **node** (*psyclone.psyGen.Loop*.) – the supplied node to which we will apply the Loop transformation.

- **collapse** (*int*) – number of loops to collapse into single iteration space or None.

- **independent** (*bool*) – whether to add the "independent" clause to the directive (not strictly necessary within PARALLEL regions).

Returns (psyclone.psyGen.GOInvokeSchedule, psyclone.undoredo. Memento)

---

class psyclone.transformations.**ACCParallelTrans**

Create an OpenACC parallel region by inserting directives. This parallel region *must* come after an enter-data directive (see *ACCEnterDataTrans*) or within a data region (see *ACCDataTrans*). For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ptrans = t.get_trans_name('ACCParallelTrans')
>>> dtrans = t.get_trans_name('ACCDataTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> # Enclose everything within a single OpenACC PARALLEL region
>>> newschedule, _ = ptrans.apply(schedule.children)
>>> # Add an enter-data directive
>>> newschedule, _ = dtrans.apply(newschedule)
>>> newschedule.view()
```

apply(*nodes*)

Apply this transformation to a subset of the nodes within a schedule - i.e. enclose the specified Loops in the schedule within a single parallel region.

Parameters nodes ((list of) *psyclone.psyGen.Node*.) – a single Node or a list of Nodes.

Raises *TransformationError* – if the nodes argument is not of the correct type.

---

class psyclone.transformations.**ColourTrans**

Apply a colouring transformation to a loop (in order to permit a subsequent parallelisation over colours). For example:

```
>>> invoke = ...
>>> schedule = invoke.schedule
>>>
>>> ctrans = ColourTrans()
>>>
>>> # Colour all of the loops
>>> for child in schedule.children:
>>>     cschedule, _ = ctrans.apply(child)
```

(continues on next page)

```
>>>
>>> csched.view()
```

apply(*node*)

Converts the Loop represented by `node` into a nested loop where the outer loop is over colours and the inner loop is over cells of that colour. :param node: The loop to transform. :type node: *psyclone.psyGen.Loop* :returns: Tuple of modified schedule and record of transformation :rtype: (psyclone.psyGen.Schedule, :py:class:`psyclone.undoredo.Memento)

class psyclone.transformations.**ExtractRegionTrans**

Provides a transformation to extract code represented by a subset of the Nodes in the PSyIR of a Schedule into a stand-alone program. Examples are given in descriptions of children classes DynamoExtractRegionTrans and GOceanExtractRegionTrans.

After applying the transformation the Nodes marked for extraction are children of the ExtractNode. Nodes to extract can be individual constructs within an Invoke (e.g. Loops containing a Kernel or BuiltIn call) or entire Invokes. This functionality does not support distributed memory.

apply(*nodes*)

Apply this transformation to a subset of the Nodes within a Schedule - i.e. enclose the specified Nodes in the Schedule within a single Extract region.

> **Parameters nodes** ((list of) *psyclone.psyGen.Node*.) – a single Node or a list of Nodes.
>
> **Returns** tuple of the modified Schedule and a record of the transformation.
>
> **Return type** (*psyclone.psyGen.Schedule*, psyclone.undoredo.Memento).
>
> **Raises** *TransformationError* – if the *nodes* argument is not of the correct type.

class psyclone.transformations.**KernelModuleInlineTrans**

Switches on, or switches off, the inlining of a Kernel subroutine into the PSy layer module. For example:

```
>>> invoke = ...
>>> schedule = invoke.schedule
>>>
>>> inline_trans = KernelModuleInlineTrans()
>>>
>>> ischedule, _ = inline_trans.apply(schedule.children[0].children[0])
>>> ischedule.view()
```

> **Warning:** For this transformation to work correctly, the Kernel subroutine must only use data that is passed in by argument, declared locally or included via use association within the subroutine. Two examples where in-lining will not work correctly are:
>
> 1. A variable is declared within the module that `contains` the Kernel subroutine and is then accessed within that Kernel;
>
> 2. A variable is included via use association at the module level and accessed within the Kernel subroutine.
>
> *There are currently no checks that these rules are being followed when in-lining so the onus is on the user to ensure correctness.*

apply(*node*, *inline=True*)

> Checks that the node is of the correct type (a Kernel) then marks the Kernel to be inlined, or not, depending on the value of the inline argument. If the inline argument is not passed the Kernel is marked to be inlined.

Note: PSyclone does not currently permit module-inlining of transformed kernels (issue #229).

class psyclone.transformations.**LoopFuseTrans**

> Provides a loop-fuse transformation. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast,invokeInfo=parse("dynamo.F90")
>>> psy=PSyFactory("dynamo0.1").create(invokeInfo)
>>> schedule=psy.invokes.get('invoke_v3_kernel_type').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import LoopFuseTrans
>>> trans=LoopFuseTrans()
>>> new_schedule,memento=trans.apply(schedule.children[0],
                                     schedule.children[1])
>>> new_schedule.view()
```

apply(*node1*, *node2*)

> Fuse the loops represented by `node1` and `node2`.

class psyclone.transformations.**MoveTrans**

> Provides a transformation to move a node in the tree. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast,invokeInfo=parse("dynamo.F90")
>>> psy=PSyFactory("dynamo0.3").create(invokeInfo)
>>> schedule=psy.invokes.get('invoke_v3_kernel_type').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import MoveTrans
>>> trans=MoveTrans()
>>> new_schedule, memento = trans.apply(schedule.children[0],
                                        schedule.children[2],
                                        position="after")
>>> new_schedule.view()
```

Nodes may only be moved to a new location with the same parent and must not break any dependencies otherwise an exception is raised.

apply(*node*, *location*, *position='before'*)

> Move the node represented by `node` before location `location` (which is also a node) by default and after if the optional *position* argument is set to 'after'. An exception is raised if the move is invalid.

class psyclone.transformations.**OCLTrans**

> Switches on/off the generation of an OpenCL PSy layer for a given InvokeSchedule. For example:

```
>>> invoke = ...
>>> schedule = invoke.schedule
>>>
>>> ocl_trans = OCLTrans()
>>> new_sched, _ = ocl_trans.apply(schedule)
```

apply(*sched*, *opencl=True*)

Apply the OpenCL transformation to the supplied GOInvokeSchedule. This causes PSyclone to generate an OpenCL version of the corresponding PSy-layer routine. The generated code makes use of the FortCL library (https://github.com/stfc/FortCL) in order to manage the OpenCL device directly from Fortran.

> **Parameters**
>
> • **sched** (psyclone.psyGen.GOInvokeSchedule) – InvokeSchedule to transform.
>
> • **opencl** (*bool*) – whether or not to enable OpenCL generation.

Note: OpenCL support is still under development. See *OpenCL* for more details.

---

class psyclone.transformations.**OMPLoopTrans**(*omp_schedule='static'*)

Adds an orphaned OpenMP directive to a loop. i.e. the directive must be inside the scope of some other OMP Parallel REGION. This condition is tested at code-generation time. The optional 'reprod' argument in the apply method decides whether standard OpenMP reduction support is to be used (which is not reproducible) or whether a manual reproducible reproduction is to be used.

> **Parameters** **omp_schedule** (*str*) – The OpenMP schedule to use.

For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory, GenerationError
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>> print psy.invokes.names
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ltrans = t.get_trans_name('OMPLoopTrans')
>>> rtrans = t.get_trans_name('OMPParallelTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> new_schedule = schedule
>>>
# Apply the OpenMP Loop transformation to *every* loop
# in the schedule
>>> for child in schedule.children:
>>>     newschedule, memento = ltrans.apply(child, reprod=True)
>>>     schedule = newschedule
>>>
# Enclose all of these loops within a single OpenMP
# PARALLEL region
```

(continues on next page)

```
>>> rtrans.omp_schedule("dynamic,1")
>>> newschedule, memento = rtrans.apply(schedule.children)
>>>
>>>
```

apply(*node*, *reprod=None*)

Apply the OMPLoopTrans transformation to the specified node in a Schedule. This node must be a Loop since this transformation corresponds to wrapping the generated code with directives like so:

```
!$OMP DO
do ...
    ...
end do
!$OMP END DO
```

At code-generation time (when `OMPLoopDirective.gen_code()` is called), this node must be within (i.e. a child of) an OpenMP PARALLEL region.

The optional reprod argument will cause a reproducible reduction to be generated if it is set to True, otherwise the default value (as read from the psyclone.cfg file) will be used. Note, reproducible in this case means obtaining the same results with the same number of OpenMP threads, not for different numbers of OpenMP threads.

> **Parameters**
>
> - **node** (*psyclone.psyGen.Node*) – the supplied node to which we will apply the OMPLoopTrans transformation
>
> - **reprod** (*Boolean or None*) – optional argument to determine whether to generate reproducible OpenMP reductions (True) or not (False). The default value is None which will cause PSyclone to look up a default value
>
> **Returns** (*psyclone.psyGen.Schedule*, psyclone.undoredo.Memento)

omp_schedule

Returns the OpenMP schedule that will be specified by this transformation. The default schedule is 'static'.

---

class psyclone.transformations.**OMPParallelLoopTrans**(*omp_schedule='static'*)

Adds an OpenMP PARALLEL DO directive to a loop.

For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast, invokeInfo = parse("dynamo.F90")
>>> psy = PSyFactory("dynamo0.1").create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_v3_kernel_type').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import OMPParallelLoopTrans
>>> trans = OMPParallelLoopTrans()
>>> new_schedule, memento = trans.apply(schedule.children[0])
>>> new_schedule.view()
```

apply(*node*)

Apply an OMPParallelLoop Transformation to the supplied node (which must be a Loop). In the generated code this corresponds to wrapping the Loop with directives:

```
!$OMP PARALLEL DO ...
do ...
  ...
end do
!$OMP END PARALLEL DO
```

> **Parameters node** (psyclone.f2pygen.DoGen) – the node (loop) to which to apply the transformation.
>
> **Returns** Two-tuple of transformed schedule and a record of the transformation.
>
> **Return type** (psyclone.psyGen.Schedule, :py:class:`psyclone.undoredo.Memento)

class psyclone.transformations.**OMPParallelTrans**

> Create an OpenMP PARALLEL region by inserting directives. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory, GenerationError
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ltrans = t.get_trans_name('GOceanOMPLoopTrans')
>>> rtrans = t.get_trans_name('OMPParallelTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> new_schedule = schedule
>>>
>>> # Apply the OpenMP Loop transformation to *every* loop
>>> # in the schedule
>>> for child in schedule.children:
>>>     newschedule, memento = ltrans.apply(child)
>>>     schedule = newschedule
>>>
>>> # Enclose all of these loops within a single OpenMP
>>> # PARALLEL region
>>> newschedule, _ = rtrans.apply(schedule.children)
>>> newschedule.view()
```

> **apply** (*nodes*)
>
> > Apply this transformation to a subset of the nodes within a schedule - i.e. enclose the specified Loops in the schedule within a single parallel region.
> >
> > > **Parameters nodes** ((list of) *psyclone.psyGen.Node*.) – a single Node or a list of Nodes.
> > >
> > > **Raises** *TransformationError* – if the nodes argument is not of the correct type.

Note: PSyclone does not support (distributed-memory) halo swaps or global sums within OpenMP parallel regions. Attempting to create a parallel region for a set of nodes that includes halo swaps or global sums will produce an error.

In such cases it may be possible to re-order the nodes in the Schedule such that the halo swaps or global sums are performed outside the parallel region. The *MoveTrans* transformation may be used for this.

---

class psyclone.transformations.**ProfileRegionTrans**

Create a profile region around a list of statements. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory, GenerationError
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> p_trans = t.get_trans_name('ProfileRegionTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> # Enclose all children within a single profile region
>>> newschedule, _ = p_trans.apply(schedule.children)
>>> newschedule.view()
```

apply(*nodes*)

Apply this transformation to a subset of the nodes within a schedule - i.e. enclose the specified Nodes in the schedule within a single profiler region.

> **Parameters nodes** (psyclone.psygen.Node or list of psyclone.psygen.Node.) – Can be a single node or a list of nodes.

7.3 Kernels

PSyclone supports the transformation of Kernels as well as PSy-layer code. However, the transformation of kernels to produce new kernels brings with it additional considerations, especially regarding the naming of the resulting kernels. PSyclone supports two use cases:

1. the HPC expert wishes to optimise the same kernel in different ways, depending on where/how it is called;

2. the HPC expert wishes to transform the kernel just once and have the new version used throughout the Algorithm file.

The second case is really an optimisation of the first for the case where the same set of transformations is applied to every instance of a given kernel.

Since PSyclone is run separately for each Algorithm in a given application, ensuring that there are no name clashes for kernels in the application as a whole requires that some state is maintained between PSyclone invocations. This is achieved by requiring that the same kernel output directory is used for every invocation of PSyclone when building a given application. However, this is under the control of the user and therefore it is possible to use the same output directory for a subset of algorithms that require the same kernel transformation and then a different directory for another subset requiring a different transformation. Of course, such use would require care when building and linking the application since the differently-optimised kernels would have the same names.

By default, transformed kernels are written to the current working directory. Alternatively, the user may specify the location to which to write the modified code via the `-okern` command-line flag.

In order to support the two use cases given above, PSyclone supports two different kernel-renaming schemes: "multiple" and "single" (specified via the `--kernel-renaming` command-line flag). In the default, "multiple" scheme, PSyclone ensures that each transformed kernel is given a unique name (with reference to the contents of the kernel output directory). In the "single" scheme, it is assumed that any given kernel that is transformed is always transformed in the same way (or left unchanged) and thus just one transformed version of it is created. This assumption is checked by examining the Fortran code for any pre-existing transformed version of that kernel. If another transformed version of that kernel exists and does not match that created by the current transformation then PSyclone will raise an exception.

7.3.1 Rules

Kernel code that is to be transformed is subject to certain restrictions. These rules are intended to make kernel transformations as robust as possible, in particular by limiting the amount of code that must be parsed by PSyclone (via fparser). The rules are as follows:

1) Any variable or procedure accessed by a kernel must either be explicitly declared or named in the `only` clause of a module `use` statement within the scope of the subroutine containing the kernel implementation. This means that:

    1) Kernel subroutines are forbidden from accessing data using COMMON blocks;

    2) Kernel subroutines are forbidden from calling procedures declared via the EXTERN statement;

    3) Kernel subroutines must not access data or procedures made available via their parent (containing) module.

2) The full Fortran source of a kernel must be available to PSyclone. This includes the source of any modules from which it accesses either routines or data. (However, kernel routines are permitted to make use of Fortran intrinsic routines.)

For instance, consider the following Fortran module containing the `bc_ssh_code` kernel:

```fortran
module boundary_conditions_mod
  real :: forbidden_var

contains

  subroutine bc_ssh_code(ji, jj, istep, ssha)
    use kind_params_mod, only: go_wp
    use model_mod, only: rdt
    integer,                        intent(in)    :: ji, jj, istep
    real(go_wp), dimension(:,:), intent(inout) :: ssha
    real(go_wp) :: rtime

    rtime = real(istep, go_wp) * rdt
    ...
  end subroutine bc_ssh_code

end module boundary_conditions_mod
```

Since the kernel subroutine accesses data (the `rdt` variable) from the `model_mod` module, the source of that module must be available to PSyclone if a transformation is applied to this kernel. Should `rdt` not actually be defined in `model_mod` (i.e. `model_mod` itself imports it from another module) then the source containing its definition must also be available to PSyclone. Note that the rules forbid the `bc_ssh_code` kernel from accessing the `forbidden_var` variable that is available to it from the enclosing module scope.

Note: these rules *only* apply to kernels that are the target of PSyclone kernel transformations.

7.3.2 Available Kernel Transformations

PSyclone currently provides just one kernel transformation:

class psyclone.transformations.**ACCRoutineTrans**
> Transform a kernel subroutine by adding a "!$acc routine" directive (causing it to be compiled for the OpenACC accelerator device). For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.transformations import ACCRoutineTrans
>>> rtrans = ACCRoutineTrans()
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> kern = schedule.children[0].children[0].children[0]
>>> # Transform the kernel
>>> newkern, _ = rtrans.apply(kern)
```

> **apply**(*kern*)
> > Modifies the AST of the supplied kernel so that it contains an '!$acc routine' OpenACC directive.
> >
> > > **Parameters kern** (psyclone.psyGen.Call) – The kernel object to transform.
> > >
> > > **Returns** (transformed kernel, memento of transformation)
> > >
> > > **Return type** 2-tuple of (psyclone.psyGen.Kern, psyclone.undoredo.Memento).
> > >
> > > **Raises** *TransformationError* – if we fail to find the subroutine corresponding to the kernel object.
>
> **name**
> > > **Returns** the name of this transformation class.
> > >
> > > **Return type** str
>
> **validate**(*kern*)
> > Perform checks that the supplied kernel can be transformed.
> >
> > > **Parameters kern** (psyclone.psyGen.Call) – the kernel which is the target of the transformation.
> > >
> > > **Raises** *TransformationError* – if the target kernel is a built-in.

Note: PSyclone does not currently permit transformed kernels to be module-inlined. (Issue #229.)

7.4 Applying

Transformations can be applied either interactively or through a script.

7.4.1 Interactive

To apply a transformation interactively we first parse and analyse the code. This allows us to generate a "vanilla" PSy layer. For example . . .

```python
from psyclone.parse.algorithm import parse
from psyclone.psyGen import PSyFactory

# This example uses version 0.1 of the Dynamo API
api = "dynamo0.1"

# Parse the file containing the algorithm specification and
# return the Abstract Syntax Tree and invokeInfo objects
ast, invokeInfo = parse("dynamo.F90", api=api)

# Create the PSy-layer object using the invokeInfo
psy = PSyFactory(api).create(invokeInfo)

# Optionally generate the vanilla PSy layer fortran
print psy.gen
```

We then extract the particular schedule we are interested in. For example . . .

```python
# List the various invokes that the PSy layer contains
print psy.invokes.names

# Get the required invoke
invoke = psy.invokes.get('invoke_0_v3_kernel_type')

# Get the schedule associated with the required invoke
schedule = invoke.schedule
schedule.view()
```

Now we have the schedule we can create and apply a transformation to it to create a new schedule and then replace the original schedule with the new one. For example . . .

```python
# Get the list of possible loop transformations
from psyclone.psyGen import TransInfo
t = TransInfo()
print t.list

# Create an OpenMPLoop-transformation
ol = t.get_trans_name('OMPParallelLoopTrans')

# Apply it to the loop schedule of the selected invoke
new_schedule,memento = ol.apply(schedule.children[0])
new_schedule.view()

# Replace the original loop schedule of the selected invoke
# with the new, transformed schedule
invoke.schedule=new_schedule
```

```
# Generate the Fortran code for the new PSy layer
print psy.gen
```

More examples of use of the interactive application of transformations can be found in the runme*.py files within the examples/dynamo/eg1 and examples/dynamo/eg2 directories. Some simple examples of the use of transformations are also given in the previous section.

7.4.2 Script

PSyclone provides a Python script (**psyclone**) that can be used from the command line to generate PSy layer code and to modify algorithm layer code appropriately. By default this script will generate "vanilla" (unoptimised) PSy layer code. For example:

```
> psyclone algspec.f90
> psyclone -oalg alg.f90 -opsy psy.f90 -api dynamo0.3 algspec.f90
```

The **psyclone** script has an optional **-s** flag which allows the user to specify a script file to modify the PSy layer as required. Script files may be specified without a path. For example:

```
> psyclone -s opt.py algspec.f90
```

In this case the Python search path **PYTHONPATH** will be used to try to find the script file.

Alternatively, script files may be specified with a path. In this case the file is expected to be found in the specified location. For example . . .

```
> psyclone -s ./opt.py algspec.f90
> psyclone -s ../scripts/opt.py algspec.f90
> psyclone -s /home/me/PSyclone/scripts/opt.py algspec.f90
```

PSyclone also provides the same functionality via a function (which is what the **psyclone** script calls internally).

psyclone.generator.**generate**(*filename*, *api=''*, *kernel_path=''*, *script_name=None*, *line_length=False*, *distributed_memory=None*, *kern_out_path=''*, *kern_naming='multiple'*)

Takes a PSyclone algorithm specification as input and outputs the associated generated algorithm and psy codes suitable for compiling with the specified kernel(s) and support infrastructure. Uses the parse.algorithm.parse() function to parse the algorithm specification, the psyGen.PSy class to generate the PSy code and the algGen.Alg class to generate the modified algorithm code.

> **Parameters**
>
> - **filename** (*str*) – The file containing the algorithm specification.
> - **kernel_path** (*str*) – The directory from which to recursively search for the files containing the kernel source (if different from the location of the algorithm specification).
> - **script_name** (*str*) – A script file that can apply optimisations to the PSy layer (can be a path to a file or a filename that relies on the PYTHONPATH to find the module).
> - **line_length** (*bool*) – A logical flag specifying whether we care about line lengths being longer than 132 characters. If so, the input (algorithm and kernel) code is checked to make sure that it conforms. The default is False.
> - **distributed_memory** (*bool*) – A logical flag specifying whether to generate distributed memory code. The default is set in the config.py file.
> - **kern_out_path** (*str*) – Directory to which to write transformed kernel code.

- **kern_naming** ([*bool*](bool)) – the scheme to use when re-naming transformed kernels.

Returns 2-tuple containing fparser1 ASTs for the algorithm code and the psy code.

Return type (fparser.one.block_statements.BeginSource,   fparser.one.
block_statements.Module)

Raises

- **IOError** – if the filename or search path do not exist

- **GenerationError** – if an invalid API is specified.

- **GenerationError** – if an invalid kernel-renaming scheme is specified.

For example:

```
>>> from psyclone.generator import generate
>>> alg, psy = generate("algspec.f90")
>>> alg, psy = generate("algspec.f90", kernel_path="src/kernels")
>>> alg, psy = generate("algspec.f90", script_name="optimise.py")
>>> alg, psy = generate("algspec.f90", line_length=True)
>>> alg, psy = generate("algspec.f90", distributed_memory=False)
```

A valid script file must contain a **trans** function which accepts a **PSy** object as an argument and returns a **PSy** object, i.e.:

```
def trans(psy)
    ...
    return psy
```

It is up to the script what it does with the PSy object. The example below does the same thing as the example in the *Interactive* section.

```
def trans(psy):
    from psyclone.transformations import OMPParallelLoopTrans
    invoke = psy.invokes.get('invoke_0_v3_kernel_type')
    schedule = invoke.schedule
    ol = OMPParallelLoopTrans()
    new_schedule, _ = ol.apply(schedule.children[0])
    invoke.schedule = new_schedule
    return psy
```

Of course the script may apply as many transformations as is required for a particular schedule and may apply transformations to all the schedules (i.e. invokes) contained within the PSy layer.

Examples of the use of transformation scripts can be found in the examples/dynamo/eg3 and examples/dynamo/scripts directories. Please read the examples/dynamo/README file first as it explains how to run the examples (and see also the examples/check_examples script).

7.5 OpenMP

OpenMP is added to a code by using transformations. The three transformations currently supported allow the addition of an **OpenMP Parallel** directive, an **OpenMP Do** directive and an **OpenMP Parallel Do** directive, respectively, to a code.

The generic versions of these three transformations (i.e. ones that theoretically work for all APIs) were given in the *Available* section. The API-specific versions of these transformations are described in the API-specific sections of this document.

7.5.1 Reductions

PSyclone supports parallel scalar reductions. If a scalar reduction is specified in the Kernel metadata (see the API-specific sections for details) then PSyclone ensures the appropriate reduction is performed.

In the case of distributed memory, PSyclone will add **GlobalSum's** at the appropriate locations. As can be inferred by the name, only "summation" reductions are currently supported for distributed memory.

In the case of an OpenMP parallel loop the standard reduction support will be used by default. For example

```
!$omp parallel do, reduction(+:x)
!loop
!$omp end parallel do
```

OpenMP reductions do not guarantee to give bit reproducible results for different runs of the same problem even if the same problem is run using the same resources. The reason for this is that the order in which data is reduced is not mandated.

Therefore, an additional **reprod** option has been added to the **OpenMP Do** transformation. If the reprod option is set to "True" then the OpenMP reduction support is replaced with local per-thread reductions which are reduced serially after the loop has finished. This implementation guarantees to give bit-wise reproducible results for different runs of the same problem using the same resources, but will not bit-wise compare if the code is rerun with different numbers of OpenMP threads.

7.5.2 Restrictions

If two reductions are used within an OpenMP region and the same variable is used for both reductions then PSyclone will raise an exception. In this case the solution is to use a different variable for each reduction.

PSyclone does not support (distributed-memory) halo swaps or global sums within OpenMP parallel regions. Attempting to create a parallel region for a set of nodes that includes halo swaps or global sums will produce an error. In such cases it may be possible to re-order the nodes in the Schedule using the *MoveTrans* transformation.

7.6 OpenCL

In common with OpenMP, the conversion of the generated code to use OpenCL is performed by a transformation (`OCLTrans` - see the *Available* Section above). Currently this transformation is only supported for the GOcean1.0 API and is applied to the whole InvokeSchedule of an Invoke. This means that all kernels in that Invoke will be executed on the OpenCL device. At present the `OCLTrans` transformation only alters the generated PSy-layer code. It is currently the user's responsibility to convert the actual kernel code from Fortran into OpenCL. Work is underway to extend PSyclone in order to perform this translation automatically.

The OpenCL code generated by PSyclone is still Fortran and makes use of the FortCL library (https://github.com/stfc/FortCL) to access OpenCL functionality. It also relies upon the OpenCL support provided by the dl_esm_inf library (https://github.com/stfc/dl_esm_inf).

The introduction of OpenCL code generation in PSyclone has been largely motivated by the need to target Field Programmable Gate Array (FPGA) accelerator devices. It is not currently designed to target the other compute devices that OpenCL supports (such as GPUs and multi-core CPUs) but this is a potentially fruitful area for future work.

7.7 OpenACC

PSyclone supports the generation of code targetting GPUs through the addition of OpenACC directives. This is achieved by a user applying various OpenACC transformations to the PSyIR before the final Fortran code is generated. The steps to parallelisation are very similar to those in OpenMP with the added complexity of managing the movement of data to and from the GPU device. For the latter task PSyclone provides the `ACCDataTrans` and `ACCEnterDataTrans` transformations, as described in the *Available* Section above. These two transformations add statically- and dynamically-scoped data regions, respectively. The former manages what data is on the remote device for a specific section of code while the latter allows run-time control of data movement. This second option is essential for minimising data movement as, without it, PSyclone-generated code would move data to and from the device upon every entry/exit of an Invoke. The first option is mainly provided as an aid to incremental porting and/or debugging of an OpenACC application as it provides explicit control over what data is present on a device for a given (part of an) Invoke routine.

The PGI compiler provides an alternative approach to controlling data movement through its 'unified memory' option (`-ta=tesla:managed`). When this is enabled the compiler itself takes on the task of ensuring that data is copied to/from the GPU when required. (Note that this approach can struggle with Fortran code containing derived types however.)

As well as ensuring the correct data is copied to and from the remote device, OpenACC directives must also be added to a code in order to tell the compiler how it should be parallelised. PSyclone provides the `ACCKernelsTrans`, `ACCParallelTrans` and `ACCLoopTrans` transformations for this purpose. The simplest of these is `ACCKernelsTrans` (currently only supported for the NEMO API) which encloses the code represented by a sub-tree of the PSyIR within an OpenACC `kernels` region. This essentially gives free-reign to the compiler to automatically parallelise any suitable loops within the specified region. An example of the use of `ACCDataTrans` and `ACCKernelsTrans` may be found in PSyclone/examples/nemo/eg3.

However, as with any "automatic" approach, a more performant solution can almost always be obtained by providing the compiler with more explicit direction on how to parallelise the code. The `ACCParallelTrans` and `ACCLoopTrans` transformations allow the user to define thread-parallel regions and, within those, define which loops should be parallelised. (Note that these two transformations are currently only supported for the GOcean1.0 and NEMO APIs.) For an example of their use please see PSyclone/examples/gocean/eg2.

In order for a given section of code to be executed on a GPU, any routines called from within that section must also have been compiled for the GPU. This then requires either that any such routines are in-lined or that the OpenACC `routine` directive be added to any such routines. This situation will occur routinely in those PSyclone APIs that use the PSyKAl separation of concerns since the user-supplied kernel routines are called from within PSyclone-generated loops in the PSy layer. PSyclone therefore provides the `ACCRoutineTrans` transformation which, given a Kernel node in the PSyIR, creates a new version of that kernel with the `routine` directive added. Again, please see PSyclone/examples/gocean/eg2 for an example.

DISTRIBUTED MEMORY

PSyclone supports the generation of code for distributed memory machines. When this option is switched on, PSyclone takes on responsibility for both performance and correctness, as described below.

8.1 Correctness

PSyclone is responsible for adding appropriate distributed memory communication calls to the PSy layer to ensure that the distributed memory code runs correctly. For example, a stencil operation will require halo exchanges between the different processes.

The burden of correctly placing distributed memory communication calls has traditionally been born by the user. However, PSyclone is able to determine the placing of these within the PSy-layer, thereby freeing the user from this responsibility. Thus, the Algorithm and Kernel code remain the same, irrespective of whether the target architecture does or does not require a distributed memory solution.

8.2 Performance

PSyclone adds **HaloExchange** and **GlobalSum** objects to the generated PSyIR **InvokeSchedule** at the required locations. The halo-exchange and global-sum objects are exposed here for the purposes of optimisation. For example the halo-exchange and/or global-sum objects may be moved in the schedule (via appropriate transformations) to enable overlap of computation with communication.

Note: When these optimisations are implemented, add a reference to the *Transformations* Section.

A halo exchange is required with distributed memory when a processor requires data from its halo and the halo information is out of date. One example is where a field is written to and then read using a stencil access. Halo exchanges have performance implications so should only be used where necessary.

A global sum is required with distributed memory when a scalar is written to. Global sums can have performance implications so should only be used where necessary. Global sums currently only occur in certain Built-in kernels. The description of Built-ins indicates when this is the case.

8.3 Implementation

Within the contents of an `invoke()` call, PSyclone is able to statically determine which communication calls are required and where they should be placed. However, PSyclone has no information on what happens outside `invoke()` calls and thus is unable to statically determine whether communication is required between these calls. The solution

we use is to add run-time flags in the PSy layer to keep track of whether data has been written to and read from. These flags are then used to determine whether communication calls are required upon entry to an `invoke()`.

8.4 Control

Support for distributed memory can be switched on or off with the default being on. The default can be changed permanently by modifying the `DISTRIBUTED_MEMORY` variable in the `psyclone.cfg` configuration file to `false` (see *Configuration*).

Distributed memory can be switched on or off from the `psyclone` script using the `-dm`/`--dist_mem` or `-nodm`/`--no_dist_mem` flags, respectively.

For interactive access, the distributed memory option can be changed interactively from the `PSyFactory` class by setting the optional `distributed_memory` flag; for example:

```
psy = PSyFactory(api=api, distributed_memory=False)
```

Similarly the distributed memory option can be changed interactively from the `generate` function by setting the optional `distributed_memory` flag, e.g.:

```
psy, alg = generate("file.f90", distributed_memory=False).
```

8.5 Status

Distributed memory support is currently limited to the `dynamo0.3` API. The remaining APIs ignore the distributed memory flag and continue to produce code without any distributed memory functionality, irrespective of its value.

BUILT-INS

Built-ins (named by analogy with the native functionality provided by Python) are operations which can be specified within an invoke call in the algorithm layer but do not require an associated kernel to be implemented as they are provided directly by the infrastructure.

One use of Built-ins is for commonly used operations. In this case Built-ins simplify the use of the system as users do not need to write kernel routines. Built-ins also offer a potential performance advantage as they provide a specification of what is required without an implementation. Therefore the PSy layer is free to implement these operations in whatever way it chooses.

Note: In general, PSyclone will need to know the types of the arguments being passed to any Built-ins. The parser obtains this information from an API-specific file that contains the metadata for all Built-in operations supported for that API.

9.1 Example

In the following example, the invoke call includes a call to a Built-in (`setval_c`) and a user-supplied kernel (`matrix_vector_kernel_mm_type`). The Built-in sets all values in the field `Ax` to `0.0`. Notice that, unlike the kernel call, no use association is required for the Built-in since it is provided as part of the environment (*c.f.* Fortran intrinsics such as `sin()`).

```fortran
subroutine jacobi_solver_algorithm(lhs, rhs, mm, mesh, n_iter)
  use matrix_vector_mm_mod, only: matrix_vector_kernel_mm_type
  integer,              intent(in)    :: n_iter
  type(field_type),     intent(inout) :: lhs, rhs
  type(operator_type),  intent(inout) :: mm
  type(mesh_type),      intent(in)    :: mesh
  type(field_type)                    :: Ax, lumped_weight, res

  real(kind=r_def), parameter :: MU = 0.9_r_def
  ...

  do iter = 1,n_iter
    call invoke( setval_c(Ax, 0.0) )
    call invoke( matrix_vector_kernel_mm_type(Ax,lhs,mm) )
    ...
  end do

end subroutine jacobi_solver_algorithm
```

Below is an example of a kernel that is consistent with the `matrix_vector_kernel_mm_type` kernel specified in the example above.

```fortran
module matrix_vector_mm_mod
  type, public, extends(kernel_type) :: matrix_vector_kernel_mm_type
    private
    type(arg_type) :: meta_args(3) = (/                            &
         arg_type(GH_FIELD,    GH_INC,  ANY_SPACE_1),              &
         arg_type(GH_FIELD,    GH_READ, ANY_SPACE_1),              &
         arg_type(GH_OPERATOR, GH_READ, ANY_SPACE_1, ANY_SPACE_1) &
         /)
    integer :: iterates_over = CELLS
  contains
    procedure, nopass ::matrix_vector_mm_code
  end type
contains
  subroutine matrix_vector_mm_code(cell,         &
                                   nlayers,      &
                                   lhs, x,       &
                                   ncell_3d,     &
                                   mass_matrix, &
                                   ndf,undf,map)
  end subroutine matrix_vector_mm_code
end module matrix_vector_mm_mod
```

We now translate the algorithm layer code and generate the psy layer code. The algorithm code is assumed to be in a file call *solver_mod.x90*. In this case we use the top level python interface. See the *API* section for different ways to translate/generate code.

```python
>>> from psyclone.generator import generate
>>> alg, psy = generate("solver_mod.x90")
>>> print alg
>>> print psy
```

The resultant generated algorithm code is given below.

Ignoring the difference in case (which is due to the output format of the code parser) the differences between the original algorithm code and the translated algorithm code are:

- the generic calls to `invoke` have been replaced by specific `CALL invoke_xx`. The calls within the invoke are removed, as are duplicate arguments and any literals leaving the three fields being passed in.

- a use statement is added for the each of the new `CALL invoke_xx` which will call the generated PSy layer code.

The existance of a call to a Built-in has made no difference at this point:

```fortran
SUBROUTINE jacobi_solver_algorithm(lhs, rhs, mm, mesh, n_iter)
  USE solver_mod_psy, ONLY: invoke_5_matrix_vector_kernel_mm_type
  USE solver_mod_psy, ONLY: invoke_4
  INTEGER, intent(in) :: n_iter
  TYPE(field_type), intent(inout) :: lhs, rhs
  TYPE(operator_type), intent(inout) :: mm
  TYPE(mesh_type), intent(in) :: mesh
  TYPE(field_type) ax, lumped_weight, res

  REAL(KIND=r_def), parameter :: mu = 0.9_r_def

  INTEGER iter
```

(continues on next page)

```fortran
    INTEGER rhs_fs
    TYPE(function_space_type) fs
    ...
    DO iter = 1,n_iter
      CALL invoke_4(ax)
      CALL invoke_5_matrix_vector_kernel_mm_type(ax, lhs, mm)
      ...
    END DO
END SUBROUTINE jacobi_solver_algorithm
```

A vanilla (not optimised) version of the generated PSy layer is given below. As expected the kernel code is called from the PSy layer. However, in the case of the *setval_c* Built-in, the code for this has been written directly into the PSy layer (the loop setting *ax_proxy%data(df) = 0.0*). This example illustrates that Built-ins may be implemented in whatever way the generator sees fit with no change to the algorithm and kernel layers.

```fortran
MODULE solver_mod_psy
  ...
  SUBROUTINE invoke_4(ax)
    USE mesh_mod, ONLY: mesh_type
    TYPE(field_type), intent(inout) :: ax
    INTEGER df
    INTEGER undf_any_space_1
    TYPE(field_proxy_type) ax_proxy
    !
    ! Initialise field proxies
    !
    ax_proxy = ax%get_proxy()
    !
    ! Initialise sizes and allocate any basis arrays for any_space_1
    !
    undf_any_space_1 = ax_proxy%vspace%get_undf()
    !
    ...
    ! Call our kernels
    !
    DO df=1,undf_any_space_1
      ax_proxy%data(df) = 0.0
    END DO
    !
    ...
    !
  END SUBROUTINE invoke_4
  SUBROUTINE invoke_5_matrix_vector_kernel_mm_type(ax, lhs, mm)
    USE matrix_vector_mm_mod, ONLY: matrix_vector_mm_code
    ...
    TYPE(field_type), intent(inout) :: ax, lhs
    TYPE(operator_type), intent(inout) :: mm
    ...
    !
    ! Initialise field proxies
    !
    ax_proxy = ax%get_proxy()
    lhs_proxy = lhs%get_proxy()
    mm_proxy = mm%get_proxy()
    !
    ! Initialise number of layers
```

```
    !
    nlayers = ax_proxy%vspace%get_nlayers()
    !
    ! Initialise sizes and allocate any basis arrays for any_space_1
    !
    ndf_any_space_1 = ax_proxy%vspace%get_ndf()
    undf_any_space_1 = ax_proxy%vspace%get_undf()
    !
    ...
    DO cell=1,mesh%get_last_halo_cell(1)
      !
      map_any_space_1 => ax_proxy%vspace%get_cell_dofmap(cell)
      !
      CALL matrix_vector_mm_code(cell, nlayers, ax_proxy%data,          &
                                 lhs_proxy%data, mm_proxy%ncell_3d,      &
                                 mm_proxy%local_stencil, ndf_any_space_1, &
                                 undf_any_space_1, map_any_space_1)
      ...
      !
    END DO
    !
    ...
    !
  END SUBROUTINE invoke_5_matrix_vector_kernel_mm_type
  ...
END MODULE solver_mod_psy
```

This example is distributed with PSyclone and can be found in `<PSYCLONEHOME>/examples/dynamo/eg4`.

9.2 Supported Built-in operations

The list of supported Built-ins is API-specific and therefore is described under the documentation of each API.

9.3 Adding new additional Built-in operations

1. Identify the PSyclone source file for the API to be extended. *e.g.* for dynamo0.3 it is `src/dynamo0p3_builtins.py`.

2. Edit this source file to create the class for this new call. It must inherit from the API-specific parent class for Built-in operations (`DynBuiltInKern` for dynamo0.3).

3. Implement `__str__` and `gen_code()` methods for this new class.

4. Add the name of the new Built-in operation and its corresponding class to the `BUILTIN_MAP` dictionary in that source file.

5. Add metadata describing this call to the appropriate file specified in the `BUILTIN_DEFINITIONS_FILE` in that source file. For dynamo0.3 this is `dynamo0p3_builtins_mod.f90`.

6. Add relevant tests to the PSyclone test file for the API to be extended. *e.g.* for dynamo0.3 it is `src/tests/dynamo0p3_builtins_test.py`. The tests rely on `single_invoke` Fortran examples in the relevant `src/tests/test_files/` subfolder.

7. Add an appropriate Fortran `single_invoke` example for the new Built-in in the relevant `src/tests/test_files/` subfolder. *e.g.* for dynamo0.3 it is `src/tests/test_files/dynamo0p3/`. Names of examples follow the template `<category.number>.<subcategory.number>_<single_invoke_name>.f90`. *e.g.* for dynamo0.3 `<category.number>` is 15.

8. Document the new Built-in in the documentation of the relevant API (*e.g.* `doc/dynamo0p3.rst`).

If the API being extended does not currently support any Built-ins then the `BUILTIN_MAP` and `BUILTIN_DEFINITIONS_FILE` module variables must be added to the source file for the API. A Fortran module file must be created in the PSyclone src directory (with the name specified in `BUILTIN_DEFINITIONS_FILE`) containing metadata describing the Built-in operations. Finally, `parse.get_builtin_defs()` must be extended to import `BUILTIN_MAP` and `BUILTIN_DEFINITIONS_FILE` for this API.

DYNAMO0.3 API

This section describes the Dynamo0.3 application programming interface (API). This API explains what a user needs to write in order to make use of the Dynamo0.3 API in PSyclone.

As with the majority of PSyclone APIs, the Dynamo0.3 API specifies how a user needs to write the algorithm layer and the kernel layer to allow PSyclone to generate the PSy layer. These algorithm and kernel APIs are discussed separately in the following sections.

The Dynamo0.3 API supports the Met Office's finite element (hereafter FEM) based GungHo dynamical core (see *Introduction*). The Met Office Collaboration Wiki (login required) provides more detailed information about the dynamical core's formulation and data model.

10.1 Algorithm

The general requirements for the structure of an Algorithm are explained in the *Algorithm layer* section. This section explains the Dynamo0.3-specific specialisations and extensions.

10.1.1 Example

An example Dynamo0.3 API invoke call is given below with various different types of objects supported by the API. These different objects and their use are discussed in the following sections.

```
real(kind=r_def)              :: scalar1
integer(kind=i_def)           :: stencil_extent
type(field_type)              :: field1, field2, field3
type(field_type)              :: field5(3), field6(3)
type(quadrature_type)         :: qr
type(operator_type)           :: operator1
type(columnwise_operator_type) :: cma_op1
...
call invoke( kernel1(field1, field2, operator1, qr),          &
             builtin1(scalar1, field2, field3),               &
             kernel2(field1, stencil_extent, field3, scalar1), &
             assembly_kernel(cma_op1, operator1),             &
             name="some_calculation"                          &
           )
call invoke( prolong_kernel_type(field1, field4),            &
             restrict_kernel_type(field5, field6)
           )
```

Please see the *Algorithm layer* section for a description of the `name` argument.

Objects in the Dynamo0.3 API can be categorised by their functionality as data types and information that specifies supported operations on a particular data type. The above example introduces four of five data types supported by the Dynamo0.3 API: field, scalar, operator and column-wise operator (field vector is the fifth). `qr` represents a quadrature object which provides information required by a kernel to operate on fields (see section *Quadrature* for more details).

10.1.2 Field

Dynamo0.3 API fields, identified with `GH_FIELD` metadata, represent FEM discretisations of various dynamical core prognostic and diagnostic variables. In FEM, variables are discretised by placing them into a function space (see *Supported Function Spaces*) from which they inherit a polynomial expansion via the basis functions of that space. Field values at points within a cell are evaluated as the sum of a set of basis functions multiplied by coefficients which are the data points. Points of evaluation are determined by a quadrature object (*Quadrature*) and are independent of the function space the field is on. Placement of field data points, also called degrees of freedom ("dof"), is determined by the function space the field is on.

10.1.3 Field Vector

Depending on the function space a field lives on, the field data at a point can be a scalar or a vector (see *Supported Function Spaces* for the list of scalar and vector function spaces). There is an additional option which specifies whether the data itself is vector valued and it usually refers to bundles of scalar valued fields. Field vectors are represented as `GH_FIELD*N` where `N` is the size of the vector. The 3D coordinate field, for example, has (`x, y, z`) scalar values at the nodes and therefore has a vector size of 3.

10.1.4 Scalar

In Dynamo 0.3 API a scalar is a single value variable that can be either real or integer. Real scalars are identified with `GH_REAL` and integer scalars are identified with `GH_INTEGER` metadata.

10.1.5 Operator

Represents a matrix constructed on a per-cell basis using Local Matrix Assembly (LMA) and is identified with `GH_OPERATOR` metadata.

10.1.6 Column-Wise Operator

The Dynamo 0.3 API has support for the construction and use of column-wise/Column Matrix Assembly (CMA) operators whose metadata identifier is `GH_COLUMNWISE_OPERATOR`. As the name suggests, these are operators constructed for a whole column of the mesh. These are themselves constructed from the Local Matrix Assembly (LMA) operators of each cell in the column. The rules governing Kernels that have CMA operators as arguments are given in the *Kernel* section below.

There are three recognised Kernel types involving CMA operations; construction, application (including inverse application) and matrix-matrix. The following example sketches-out what the use of such kernels might look like in the Algorithm layer:

```
use field_mod, only: field_type
use operator_mod, only : operator_type, columnwise_operator_type
type(field_type) :: field1, field2, field3
type(operator_type) :: lma_op1, lma_op2
type(columnwise_operator_type) :: cma_op1, cma_op2, cma_op3
```

<div align="right">(continues on next page)</div>

```
real(kind=r_def) :: alpha
...
call invoke(                                                   &
        assembly_kernel(cma_op1, lma_op1, lma_op2),            &
        assembly_kernel2(cma_op2, lma_op1, lma_op2, field3),   &
        apply_kernel(field1, field2, cma_op1),                 &
        matrix_matrix_kernel(cma_op3, cma_op1, alpha, cma_op2), &
        apply_kernel(field3, field1, cma_op3),                 &
        name="cma_example")
```

The above invoke uses two LMA operators to construct the CMA operator `cma_op1`. A second CMA operator, `cma_op2`, is assembled from the same two LMA operators but also uses a field. The first of these CMA operators is then applied to `field2` and the result stored in `field1` (assuming that the metadata for `apply_kernel` specifies that it is the first field argument that is written to). The two CMA operators are then combined to produce a third, `cma_op3`. This is then applied to `field1` and the result stored in `field3`.

Note that PSyclone identifies the type of kernels performing Column-Wise operations based on their arguments as described in metadata (see *Rules for Kernels that work with CMA Operators* below). The names of the kernels in the above example are purely illustrative and are not used by PSyclone when determining kernel type.

A full example of CMA operator construction is available in `examples/dynamo/eg7`.

10.1.7 Quadrature

Kernels conforming to the Dynamo 0.3 API may require quadrature information (specified using e.g. `gh_shape = gh_quadrature_XYoZ` in the kernel metadata - see Section *gh_shape and gh_evaluator_targets*). This information must be passed to the kernel from the Algorithm layer in the form of a *quadrature_type* object. This must be the last argument passed to the kernel, e.g.:

```
type( quadrature_type )   :: qr
...
qr = quadrature_type(element_order+2, GAUSSIAN)
call invoke(pressure_gradient_kernel_type(rhs_tmp(igh_u), rho, theta, qr),   &
            kinetic_energy_gradient_kernel_type(rhs_tmp(igh_u), u, chi, qr), &
            geopotential_gradient_kernel_type(rhs_tmp(igh_u), geopotential, qr))
```

This quadrature object specifies the set of points at which the basis/differential-basis functions required by the kernel are to be evaluated.

10.1.8 Stencils

Kernel metadata may specify that a Kernel performs a stencil operation on a field. Any such metadata must provide a stencil type. See the *meta_args* section for more details. The supported stencil types are `X1D`, `Y1D`, `XORY1D` or `CROSS`.

If a stencil operation is specified by the Kernel metadata the algorithm layer must provide the `extent` of the stencil (the maximum distance from the central cell that the stencil extends). The Dynamo0.3 API expects this information to be added as an additional `integer` argument immediately after the relevant field when specifying the Kernel via an `invoke`.

For example:

```
integer :: extent = 2
call invoke(kernel(field1, field2, extent))
```

where `field2` has kernel metadata specifying that it has a stencil access.

`extent` may also be passed as a literal. For example:

```
call invoke(kernel(field1, field2, 2))
```

where, again, `field2` has kernel metadata specifying that it has a stencil access.

Note: The stencil extent specified in the Algorithm layer is not the same as the stencil size passed in to the Kernel. The latter contains the number of cells in the stencil which is dependent on both the stencil type and extent.

If the Kernel metadata specifies that the stencil is of type `XORY1D` (which means `X1D` or `Y1D`) then the algorithm layer must specify whether the stencil is `X1D` or `Y1D` for that particular kernel call. The Dynamo0.3 API expects this information to be added as an additional argument immediately after the relevant stencil extent argument. The argument should be an `integer` with valid values being `x_direction` or `y_direction`, both being supplied by the `LFRic` infrastructure via the `flux_direction_mod` fortran module

For example:

```
use flux_direction_mod, only : x_direction
integer :: direction = x_direction
integer :: extent = 2
! ...
call invoke(kernel(field1, field2, extent, direction))
```

`direction` may also be passed as a literal. For example:

```
use flux_direction_mod, only : x_direction
integer :: extent = 2
! ...
call invoke(kernel(field1, field2, extent, x_direction))
```

If certain fields use the same value of extent and/or direction then the same variable, or literal value can be provided.

For example:

```
call invoke(kernel1(field1, field2, extent,  field3, extent, direction), &
            kernel2(field1, field2, extent2, field4, extent, direction))
```

In the above example `field2` and `field3` in `kernel1` and `field4` in `kernel2` will have the same `extent` value but `field2` in `kernel2` may have a different value. Similarly, `field3` in `kernel1` and `field4` in `kernel2` will have the same `direction` value.

An example of the use of stencils is available in `examples/dynamo/eg5`.

There is currently no attempt to perform type checking in PSyclone so any errors in the type and/or position of arguments will not be picked up until compile time. However, PSyclone does check for the correct number of algorithm arguments. If the wrong number of arguments is provided then an exception is raised.

For example, running test 19.2 from the Dynamo0.3 API test suite gives:

```
cd <PSYCLONEHOME>/src/psyclone/tests
psyclone test_files/dynamo0p3/19.2_single_stencil_broken.f90
"Generation Error: error: expected '5' arguments in the algorithm layer but found '4'.
Expected '4' standard arguments, '1' stencil arguments and '0' qr_arguments'"
```

10.1.9 Inter-grid

From the Algorithm layer, an Invoke for inter-grid kernels (those that map fields between grids of different resolution) looks much like an Invoke containing general-purpose kernels. The only restrictions to be aware of are that inter-grid kernels accept only field or field-vectors as arguments and that an Invoke may not mix inter-grid kernels with any other kernel type. (Hence the second, separate Invoke in the example Algorithm code given at the beginning of this Section.)

10.2 PSy-layer

The general details of the PSy-layer are explained in the *PSy layer* section. This section describes any dynamo0p3 specific issues.

10.2.1 Module name

The PSy-layer code is contained within a Fortran module. The name of the module is determined from the algorithm-layer name with "_psy" appended. The algorithm-layer name is the algorithm's module name if it is a module, its subroutine name if it is a subroutine that is not within a module, or the program name if it is a program.

So, for example, if the algorithm code is contained within a module called "fred" then the PSy-layer module name will be "fred_psy".

10.3 Kernel

The general requirements for the structure of a Kernel are explained in the *Kernel layer* section. In the Dynamo API there are four different Kernel types; general purpose, CMA, inter-grid and *Built-ins*. In the case of built-ins, PSyclone generates the source of the kernels. This section explains the rules for the other three, user-supplied kernel types and then goes on to describe their metadata and subroutine arguments.

10.3.1 Rules for all User-Supplied Kernels

In the following, 'operator' refers to both LMA and CMA operator types.

1) A Kernel must have at least one argument that is a field, field vector, or operator. This rule reflects the fact that a Kernel iterates over a space and therefore must have some representation over that space.

2) The continuity of the iteration space of the Kernel is determined from the function space of the modified argument (see *Supported Function Spaces*). If more than one argument is modified then the iteration space is taken to be the largest required by any of those arguments. e.g. if a Kernel writes to two fields, the first on W3 (discontinuous) and the second on W1 (continuous), then the iteration space of that Kernel will be determined by the field on the continuous space.

3) If the function space of the modified argument(s) cannot be determined then they are assumed to be continuous. This is the case if any of the modified arguments are declared as `ANY_SPACE` and their actual space cannot be determined statically. This assumption is always safe but leads to additional computation if the quantities being updated are actually on discontinuous function spaces.

4) Operators do not have halo operations operating on them as they are either cell- (LMA) or column-based (CMA) and therefore act like discontinuous fields.

5) Any Kernel that writes to an operator will have its iteration space expanded such that valid values for the operator are computed in the level-1 halo.

6) Any Kernel that reads from an operator must not access halos beyond level 1. In this case PSyclone will check that the Kernel does not require values beyond the level-1 halo. If it does then PSyclone will abort.

10.3.2 Rules specific to General-Purpose Kernels without CMA Operators

1) General-purpose kernels accept arguments of any of the following types: field, field vector, LMA operator, scalar integer, scalar real.

2) A Kernel is permitted to write to more than one quantity (field or operator) and these quantities may be on the same or different function spaces.

3) A Kernel may not write to a scalar argument. (Only *Built-ins* are permitted to do this.) Any scalar aguments must therefore be declared in the metadata as `GH_READ` - see below.

10.3.3 Rules for Kernels that work with CMA Operators

The Dynamo 0.3 API has support for kernels that assemble, apply (or inverse-apply) column-wise/Column Matrix Assembly (CMA) operators. Such operators may also be used by matrix-matrix kernels. There are thus three types of CMA-related kernels. Since, by definition, CMA operators only act on data within a column, they have no horizontal dependencies. Therefore, kernels that write to them may be parallelised without colouring.

All three CMA-related kernel types must obey the following rules:

1) Since a CMA operator only acts within a single column of data, stencil operations are not permitted.

2) No vector quantities (e.g. "GH_FIELD*3" - see below) are permitted as arguments.

There are then additional rules specific to each of the three CMA kernel types. These are described below.

Assembly

CMA operators are themselves constructed from Local-Matrix-Assembly (LMA) operators. Therefore, any kernel which assembles a CMA operator must obey the following rules:

1) Have one or more LMA operators as read-only arguments.

2) Have exactly one CMA operator argument which must have write access.

3) Other types of argument (e.g. scalars or fields) are permitted but must be read-only.

Application and Inverse Application

Column-wise operators can only be applied to fields. CMA-Application kernels must therefore:

1) Have a single CMA operator as a read-only argument.

2) Have exactly two field arguments, one read-only and one that is written to.

3) The function spaces of the read and written fields must match the from and to spaces, respectively, of the supplied CMA operator.

Matrix-Matrix

A kernel that has just column-wise operators as arguments and zero or more read-only scalars is identified as performing a matrix-matrix operation. In this case:

1) Arguments must be CMA operators and, optionally, one or more scalars.

2) Exactly one of the CMA arguments must be written to while all other arguments must be read-only.

10.3.4 Rules for Inter-Grid Kernels

1) An inter-grid kernel is identified by the presence of a field or field-vector argument with the optional *mesh_arg* metadata element (see *Inter-Grid Metadata*).

2) An invoke that contains one or more inter-grid kernels must not contain any other kernel types. (This restriction is an implementation decision and could be lifted in future if there is a need.)

3) An inter-grid kernel is only permitted to have field or field-vector arguments.

4) All inter-grid kernel arguments must have the *mesh_arg* metadata entry.

5) An inter-grid kernel (and metadata) must have at least one field on each of the fine and coarse meshes. Specifying all fields as coarse or fine is forbidden.

6) Fields on different meshes must always live on different function spaces.

7) All fields on a given mesh must be on the same function space.

A consequence of Rules 5-7 is that an inter-grid kernel will only involve two function spaces.

10.3.5 Metadata

The code below outlines the elements of the Dynamo0.3 API kernel metadata, 1) 'meta_args', 2) 'meta_funcs', 3) 'gh_shape', 4) 'iterates_over' and 5) 'procedure'.

```
type, public, extends(kernel_type) :: my_kernel_type
  type(arg_type) :: meta_args(...) = (/ ... /)
  type(func_type) :: meta_funcs(...) = (/ ... /)
  integer :: gh_shape = gh_quadrature_XYoZ
  integer :: iterates_over = cells
contains
  procedure, nopass :: my_kernel_code
end type
```

These five metadata elements are discussed in order in the following sections.

meta_args

The `meta_args` array specifies information about data that the kernel code expects to be passed to it via its argument list. There is one entry in the `meta_args` array for each **scalar**, **field**, or **operator** passed into the Kernel and the order that these occur in the `meta_args` array must be the same as they are expected in the kernel code argument list. The entry must be of `arg_type` which itself contains metadata about the associated argument. The size of the meta_args array must correspond to the number of **scalars**, **fields** and **operators** passed into the Kernel.

Note: it makes no sense for a Kernel to have only **scalar** arguments (because the PSy layer will call a Kernel for each point in the spatial domain) and PSyclone will reject such Kernels.

For example, if there are a total of 2 **scalar / field / operator** entities being passed to the Kernel then the meta_args array will be of size 2 and there will be two `arg_type` entries:

```
type(arg_type) :: meta_args(2) = (/                                   &
    arg_type( ... ),                                                  &
    arg_type( ... )                                                   &
    /)
```

Argument-metadata (metadata contained within the brackets of an `arg_type` entry), describes either a **scalar**, a **field** or an **operator** (either LMA or CMA).

The first argument-metadata entry describes whether the data that is being passed is for a real scalar (`GH_REAL`), an integer scalar (`GH_INTEGER`), a field (`GH_FIELD`) or an operator (either `GH_OPERATOR` for LMA or `GH_COLUMNWISE_OPERATOR` for CMA). This information is mandatory.

Additionally, argument-metadata can be used to describe a vector of fields (see the *Field Vector* section for more details).

As an example, the following `meta_args` metadata describes 4 entries, the first is a real scalar, the next two are fields and the fourth is an operator. The third entry is a field vector of size 3.

```
type(arg_type) :: meta_args(4) = (/                                   &
    arg_type(GH_REAL, ...),                                           &
    arg_type(GH_FIELD, ... ),                                         &
    arg_type(GH_FIELD*3, ... ),                                       &
    arg_type(GH_OPERATOR, ...)                                        &
    /)
```

The second entry to argument-metadata (information contained within the brackets of an `arg_type`) describes how the Kernel makes use of the data being passed into it (the way it is accessed within a Kernel). This information is mandatory. There are currently 5 possible values of this metadata `GH_WRITE`, `GH_READ`, `GH_INC`, `GH_READWRITE` and `GH_SUM`. However, not all combinations of metadata entries are valid and PSyclone will raise an exception if an invalid combination is specified. Valid combinations are specified later in this section (see *Valid Access Modes*).

- `GH_WRITE` indicates the data is modified in the Kernel before (optionally) being read.

- `GH_READ` indicates that the data is read and is unmodified.

- `GH_INC` indicates that different iterations of a Kernel make contributions to shared values. For example, values at cell faces may receive contributions from cells on either side of the face. This means that such a Kernel needs appropriate synchronisation (or colouring) to run in parallel.

- `GH_READWRITE` indicates that different iterations of a Kernel update quantitites which do not share dofs, such as operators and fields over discontinuous function spaces. If a Kernel modifies only discontinuous fields and/or operators there is no need for synchronisation or colouring when running such Kernels in parallel. However, modifying another field with a `GH_INC` access in a Kernel means that synchronisation or colouring is required for parallel runs.

- `GH_SUM` is an example of a reduction and is the only reduction currently supported in PSyclone. This metadata indicates that values are summed over calls to Kernel code.

For example:

```
type(arg_type) :: meta_args(4) = (/                                   &
    arg_type(GH_REAL,   GH_SUM),                                      &
    arg_type(GH_FIELD, GH_INC, ... ),                                 &
    arg_type(GH_FIELD*3, GH_WRITE, ... ),                             &
    arg_type(GH_OPERATOR, GH_READ, ...)                              &
    /)
```

Note: In the Dynamo 0.3 API only *Built-ins* are permitted to write to scalar arguments (and hence perform reductions). Furthermore, this permission is currently restricted to real scalars (`GH_REAL`) as the LFRic infrastructure does not yet

support integer reductions.

For a scalar the argument metadata contains only these two entries. However, fields and operators require further entries specifying function-space information. The meaning of these further entries differs depending on whether a field or an operator is being described.

In the case of an operator, the 3rd and 4th arguments describe the `to` and `from` function spaces respectively. In the case of a field the 3rd argument specifies the function space that the field lives on. More details about the supported function spaces are in subsection *Supported Function Spaces*.

For example, the metadata for a kernel that applies a Column-wise operator to a field might look like:

```
type(arg_type) :: meta_args(3) = (/                            &
     arg_type(GH_FIELD, GH_INC, W1),                           &
     arg_type(GH_FIELD, GH_READ, W2H),                         &
     arg_type(GH_COLUMNWISE_OPERATOR, GH_READ, W1, W2H) &
     /)
```

In some cases a Kernel may be written so that it works for fields and/or operators from any type of w2 space i.e. one of `w2`, `w2h` or `w2v`. In this case the metadata should be specified as being `any_w2`.

> **Warning:** in the current implementation it is assumed that all fields and/or operators specifying `any_w2` within a kernel will use the **same** function space. It is up to the user to ensure this is the case as otherwise invalid code would be generated.

It may be that a Kernel is written such that a field and/or operators may be on/map-between any function space(s). In this case the metadata should be specified as being one of `any_space_1`, `any_space_2`, ..., `any_space_9`. The reason for having different names is that a Kernel might be written to allow 2 or more arguments to be able to support any function space but for a particular call the function spaces may have to be the same as each other.

In the example below, the first field entry supports any function space but it must be the same as the operator's `to` function space. Similarly, the second field entry supports any function space but it must be the same as the operator's `from` function space. Note, the metadata does not forbid `ANY_SPACE_1` and `ANY_SPACE_2` from being the same.

```
type(arg_type) :: meta_args(3) = (/                            &
     arg_type(GH_FIELD, GH_INC, ANY_SPACE_1),                  &
     arg_type(GH_FIELD*3, GH_WRITE, ANY_SPACE_2),              &
     arg_type(GH_OPERATOR, GH_READ, ANY_SPACE_1, ANY_SPACE_2) &
     /)
```

Note also that the scope of this naming of any-space function spaces is restricted to the argument list of individual kernels. i.e. if an Invoke contains say, two kernel calls that each support arguments on any function space, e.g. `ANY_SPACE_1`, there is no requirement that these two function spaces be the same. Put another way, if an Invoke contained two calls of a kernel with arguments described by the above metadata then the first field argument passed to each kernel call need not be on the same space.

> **Note:** A `GH_FIELD` argument that specifies `GH_WRITE` or `GH_READWRITE` as its access pattern must be a discontinuous function in the horizontal (see *Valid Access Modes* below). That means it must belong to `w3`, `wtheta` or `w2v` function spaces (see *Supported Function Spaces*). A `GH_FIELD` that specifies `GH_INC` as its access pattern may be continuous in the vertical (and discontinuous in the horizontal), continuous in the horizontal (and discontinuous in the vertical), or continuous in both. In each case the code is the same. However, if a field is discontinuous in the horizontal then it will not need colouring and, if is described as being on any space, there is currently no way to determine this from the metadata (unless we can statically determine the space of the field being passed in). At the moment this type

of Kernel is always treated as if it is continuous in the horizontal, even if it is not.

Valid Access Modes

As mentioned earlier, not all combinations of metadata are valid. Valid combinations are summarised here. All types of data (`GH_INTEGER`, `GH_REAL`, `GH_FIELD`, `GH_OPERATOR` and `GH_COLUMNWISE_OPERATOR`) may be read within a Kernel and this is specified in metadata using `GH_READ`. At least one kernel argument must be listed as being modified. When data is *modified* in a Kernel then the permitted access modes depend on the type of data it is and the function space it is on. Valid values are given in the table below.

Argument Type	Function space	Access type
GH_INTEGER	*n/a*	*GH_SUM (Built-ins only)*
GH_REAL	n/a	GH_SUM (Built-ins only)
GH_FIELD	Discontinuous	GH_WRITE, GH_READWRITE
GH_FIELD	Continuous	GH_INC
GH_OPERATOR	Any for both 'to' and 'from'	GH_WRITE, GH_READWRITE
GH_COLUMNWISE_OPERATOR	Any for both 'to' and 'from'	GH_WRITE, GH_READWRITE

Note: As mentioned above, note that only Built-ins may modify scalar arguments. *Since the LFRic infrastructure does not currently support integer reductions, integer scalar arguments are restricted to having read-only access.*

There is no restriction on the number and function-spaces of other quantities that a general-purpose kernel can modify other than that it must modify at least one. The rules for kernels involving CMA operators, however, are stricter and only one argument may be modified (the CMA operator itself for assembly, a field for CMA-application and a CMA operator for matrix-matrix kernels). If a kernel writes to quantities on different function spaces then PSyclone generates loop bounds appropriate to the largest iteration space. This means that if a single kernel updates one quantity on a continuous function space and one on a discontinuous space then the resulting loop will include cells in the level 1 halo since they are required for a quantity on a continuous space. As a consequence, any quantities on a discontinuous space will then be computed redundantly in the level 1 halo. Currently PSyclone makes no attempt to take advantage of this (by e.g. setting the appropriate level-1 halo to 'clean').

PSyclone ensures that both CMA and LMA operators are computed (redundantly) out to the level-1 halo cells. This permits their use in kernels which modify quantities on continuous function spaces and also in subsequent redundant computation of other quantities on discontinuous function spaces. In conjunction with this, PSyclone also checks (when generating the PSy layer) that any kernels which read operator values do not do so beyond the level-1 halo. If any such accesses are found then PSyclone aborts.

Supported Function Spaces

As mentioned in the *Field* and *Field Vector* sections, the function space of an argument specifies how it maps onto the underlying topology and, additionally, whether the data at a point is a vector.

Function spaces can share dofs between cells in the horizontal, vertical or both directions. Depending on the function space and FEM order, the shared dofs can lie on one or more cell entities (faces, edges and vertices) in each direction. This property is referred to as the **continuity** of a function space (horizontal, vertical or full). Alternatively, if there are no shared dofs a function space is described as **discontinuous** (fully or in a particular direction).

The mixed FEM formulation is built on a foundation set of four function spaces described below.

- `w0` is the space of scalar functions with full continuity. The shared dofs lie on cell vertices in the lowest order FEM and on all three entities in higher order FEM.

- `w1` is the space of vector functions with full continuity in the tangential direction only. In the lowest order FEM the shared dofs lie on cell edges for each component, whereas in higher order they also lie on cell faces.

- `w2` is the space of vector functions with full continuity in the normal direction only. The shared dofs lie on cell faces for each component.

- `w3` is the space of scalar functions with full discontinuity. All dofs lie within the cell volume and are not shared across the cell boundaries.

Additional function spaces required for representation of scalar or component-wise vector variables are:

- `wtheta` is the space of scalar functions based on the vertical part of `w2`, discontinuous in the horizontal and continuous in the vertical;

- `w2v` is the space of vector functions based on the vertical part of `w2`, discontinuous in the horizontal and continuous in the vertical;

- `w2h` is the space of vector functions based on the horizontal part of `w2`, continuous in the horizontal and discontinuous in the vertical.

Since the Dynamo0.3 API operates on columns of data, function spaces are categorised as continuous or discontinuous with regard to their horizontal continuity.

- **Continuous** function spaces are `w0`, `w1`, `w2` and `w2h`;

- **Discontinuous** function spaces are `w3`, `wtheta` and `w2v`.

Two additonal function space metadata descriptors as mentioned in sections above are:

- `ANY_W2` for any type of `w2` function spaces;

- `ANY_SPACE` for when the function space of the modified argument(s) cannot be determined and/or for when a Kernel has been written so that it works with fields on any of the available spaces.

As mentioned previously, both `ANY_W2` and `ANY_SPACE` function space types are treated as continuous.

Horizontally discontinuous function spaces and fields over them will not need colouring so PSyclone does not perform it. If such attempt is made, PSyclone will raise a `Generation Error` in the **Dynamo0p3ColourTrans** transformation (see *Transformations* for more details on transformations). An example of fields iterating over a discontinuous function space `wtheta` is given in `examples/dynamo/eg9`, with the `GH_READWRITE` access descriptor denoting an update to the relevant fields. This example also demonstrates how to only colour loops over continuous function spaces when transformations are applied.

Optional Field Metadata

A field entry in the meta_args array may have an optional fourth element. This element describes either a stencil access or, for inter-grid kernels, which mesh the field is on. Since an inter-grid kernel is not permitted to have stencil accesses, these two options are mutually exclusive. The metadata for each case is described in the following sections.

Stencil Metadata

Stencil metadata specifies that the corresponding field argument is accessed as a stencil operation within the Kernel. Stencil metadata only makes sense if the associated field is read within a Kernel i.e. it only makes sense to specify stencil metadata if the first entry is `GH_FIELD` and the second entry is `GH_READ`.

Stencil metadata is written in the following format:

```
STENCIL(type)
```

where `type` may be one of `X1D`, `Y1D`, `XORY1D` or `CROSS`. As the stencil `extent` (the maximum distance from the central cell that the stencil extends) is not provided in the metadata, it is expected to be provided by the algorithm writer as part of the `invoke` call (see Section *Stencils*). As there is currently no way to specify a fixed extent value for stencils in the Kernel metadata, Kernels must therefore be written to support different values of extent (i.e. stencils with a variable number of cells).

The `XORY1D` stencil type indicates that the Kernel can accept either `X1D` or `Y1D` stencils. In this case it is up to the algorithm developer to specify which of these it is from the algorithm layer as part of the `invoke` call (see Section *Stencils*).

For example, the following stencil (with `extent=2`):

```
| 4 | 2 | 1 | 3 | 5 |
```

would be declared as

```
STENCIL(X1D)
```

and the following stencil (with `extent=2`)

```
|   |   | 9 |   |   |
|   |   | 5 |   |   |
| 6 | 2 | 1 | 3 | 7 |
|   |   | 4 |   |   |
|   |   | 8 |   |   |
```

would be declared as

```
STENCIL(CROSS)
```

Below is an example of stencil information within the full kernel metadata.

```
type(arg_type) :: meta_args(3) = (/                          &
    arg_type(GH_FIELD, GH_INC, W1),                          &
    arg_type(GH_FIELD, GH_READ, W2H, STENCIL(CROSS)),        &
    arg_type(GH_OPERATOR, GH_READ, W1, W2H)                  &
    /)
```

There is a full example of this distributed with PSyclone. It may be found in `examples/dynamo/eg5`.

Inter-Grid Metadata

The alternative form of the optional fourth metadata argument for a field specifies which mesh the associated field is on. This is required for inter-grid kernels which perform prolongation or restriction operations on fields (or field vectors) existing on grids of different resolutions.

Mesh metadata is written in the following format:

```
mesh_arg=type
```

where `type` may be one of `GH_COARSE` or `GH_FINE`. Any kernel having a field argument with this metadata is assumed to be an inter-grid kernel and, as such, all of its other arguments (which must also be fields) must have it specified too. An example of the metadata for such a kernel is given below:

```
type(arg_type) :: meta_args(2) = (/                          &
    arg_type(GH_FIELD, GH_INC,  ANY_SPACE_1, mesh_arg=GH_COARSE), &
```

```
    arg_type(GH_FIELD, GH_READ, ANY_SPACE_2, mesh_arg=GH_FINE ) &
    /)
```

Note that an inter-grid kernel must have at least one field (or field- vector) argument on each mesh type. Fields that are on different meshes cannot be on the same function space while those on the same mesh must also be on the same function space.

Column-wise Operators (CMA)

In this section we provide example metadata for each of the three recognised kernel types involving CMA operators.

Column-wise operators are constructed from cell-wise (local) operators. Therefore, in order to **assemble** a CMA operator, a kernel must have at least one read-only LMA operator, e.g.:

```
type(arg_type) :: meta_args(2) = (/                                      &
    arg_type(GH_OPERATOR,          GH_READ,  ANY_SPACE_1, ANY_SPACE_2),&
    arg_type(GH_COLUMNWISE_OPERATOR, GH_WRITE, ANY_SPACE_1, ANY_SPACE_2) &
    /)
```

CMA operators (and their inverse) are **applied** to fields. Therefore any kernel of this type must have one read-only CMA operator, one read-only field and a field that is updated, e.g.:

```
type(arg_type) :: meta_args(3) = (/                                      &
    arg_type(GH_FIELD,    GH_INC,  ANY_SPACE_1),                         &
    arg_type(GH_FIELD,    GH_READ, ANY_SPACE_2),                         &
    arg_type(GH_COLUMNWISE_OPERATOR, GH_READ, ANY_SPACE_1, ANY_SPACE_2) &
    /)
```

Matrix-matrix kernels compute the product/linear combination of CMA operators. They must therefore have one such operator that is updated while the rest are read-only. They may also have read-only scalar arguments, e.g.:

```
type(arg_type) :: meta_args(3) = (/                                      &
    arg_type(GH_COLUMNWISE_OPERATOR, GH_WRITE, ANY_SPACE_1, ANY_SPACE_2), &
    arg_type(GH_COLUMNWISE_OPERATOR, GH_READ, ANY_SPACE_1, ANY_SPACE_2),  &
    arg_type(GH_COLUMNWISE_OPERATOR, GH_READ, ANY_SPACE_1, ANY_SPACE_2),  &
    arg_type(GH_REAL, GH_READ) /)
```

Note: The order with which arguments are specified in metadata for CMA kernels does not affect the process of identifying the type of kernel (whether it is assembly, matrix-matrix etc.)

---

meta_funcs

The (optional) second component of kernel metadata specifies whether any quadrature or evaluator data is required for a given function space. (If no quadrature or evaluator data is required then this metadata should be omitted.) Consider the following kernel metadata:

```
type, extends(kernel_type) :: testkern_operator_type
  type(arg_type), dimension(3) :: meta_args =   &
      (/ arg_type(gh_operator,gh_write,w0,w0),  &
         arg_type(gh_field*3,gh_read,w1),       &
         arg_type(gh_integer,gh_read)           &
```

```
      /)
  type(func_type) :: meta_funcs(2) =            &
      (/ func_type(w0, gh_basis, gh_diff_basis) &
         func_type(w1, gh_basis)                &
      /)
  integer :: gh_shape = gh_quadrature_XYoZ
  integer :: iterates_over = cells
contains
  procedure, nopass :: code => testkern_operator_code
end type testkern_operator_type
```

The `arg_type` component of this metadata describes a kernel that takes three arguments (an operator, a field and an integer scalar). Following the `meta_args` array we now have a `meta_funcs` array. This allows the user to specify that the kernel requires basis functions (`gh_basis`) and/or the differential of the basis functions (`gh_diff_basis`) on one or more of the function spaces associated with the arguments listed in `meta_args`. In this case we require both for the W0 function space but only basis functions for W1.

gh_shape and gh_evaluator_targets

If a kernel requires basis or differential-basis functions then the metadata must also specify the set of points on which these functions are required. This information is provided by the `gh_shape` component of the metadata. Currently PSyclone supports two shapes; `gh_quadrature_XYoZ` for Gaussian quadrature points and `gh_evaluator` for evaluation (of the basis/differential-basis functions) at nodal points. For the latter, there are two options: if an evaluator is required for multiple function spaces then these can be specified using the additional `gh_evaluator_targets` metadata entry. This entry is a one-dimensional, integer array containing the desired function spaces. For example, to request basis/differential-basis functions evaluated on both W0 and W1 the metadata would be:

```
integer :: gh_shape = gh_evaluator
integer :: gh_evaluator_targets(2) = (/W0, W1/)
```

The kernel must have an argument (field or operator) on each of the function spaces listed in `gh_evaluator_targets`. The default behaviour if `gh_evaluator_targets` is not specified is to provide evaluators for each function space associated with the quantities that the kernel is updating. All necessary data is extracted in the PSy layer and passed to the kernel(s) as required - nothing is required from the Algorithm layer. If a kernel requires quadrature on the other hand, the Algorithm writer must supply a `quadrature_type` object as the last argument to the kernel (see Section *Quadrature*).

Note that it is an error for kernel metadata to specify a value for `gh_shape` if no basis or differential-basis functions are required. It is also an error to specify `gh_evaluator_targets` if the kernel does not require an evaluator (i.e. `gh_shape != gh_evaluator`).

iterates over

The fourth type of metadata provided is `ITERATES_OVER`. This specifies that the Kernel has been written with the assumption that it is iterating over the specified entity. For user-supplied kernels this currently only has one valid value which is `CELLS`.

Procedure

The fifth and final type of metadata is `procedure` metadata. This specifies the name of the Kernel subroutine that this metadata describes.

For example:

```
procedure, nopass :: my_kernel_subroutine
```

10.3.6 Subroutine

Rules for General-Purpose Kernels

The arguments to general-purpose kernels (those that do not involve either CMA operators or prolongation/restriction operations) follow a set of rules which have been specified for the Dynamo0.3 API. These rules are encoded in the `generate()` method within the `ArgOrdering` abstract class in the `dynamo0p3.py` file. The rules, along with PSyclone's naming conventions, are:

1) If an LMA operator is passed then include the `cells` argument. `cells` is an integer and has intent `in`.

2) Include `nlayers`, the number of layers in a column. `nlayers` is an integer and has intent `in`.

3) For each scalar/field/vector_field/operator in the order specified by the meta_args metadata:

   1) if the current entry is a scalar quantity then include the Fortran variable in the argument list. The intent is determined from the metadata (see *meta_args* for an explanation).

   2) if the current entry is a field then include the field array. The field array name is currently specified as being `"field_"<argument_position>"_"<field_function_space>`. A field array is a real array of type `r_def` and dimensioned as the unique degrees of freedom for the space that the field is on. This value is passed in separately. Again, the intent is determined from the metadata (see *meta_args*).

      1) If the field entry has a stencil access then add an integer stencil-size argument with intent `in`. This will supply the number of cells in the stencil.

      2) If the field entry stencil access is of type `XORY1D` then add an integer direction argument with intent `in`.

   3) if the current entry is a field vector then for each dimension of the vector, include a field array. The field array name is specified as being using `"field_"<argument_position>"_"<field_function_space>"_v"<vector_position>`. A field array in a field vector is declared in the same way as a field array (described in the previous step).

   4) if the current entry is an operator then first include a dimension size. This is an integer. The name of this size is `<operator_name>"_ncell_3d"`. Next include the operator. This is a real array of type `r_def` and is 3 dimensional. The first two dimensions are the local degrees of freedom for the `to` and `from` function spaces respectively. The third dimension is the dimension size mentioned before. The name of the operator is `"op_"<argument_position>`. Again the intent is determined from the metadata (see *meta_args*).

4) For each function space in the order they appear in the metadata arguments (the `to` function space of an operator is considered to be before the `from` function space of the same operator as it appears first in lexicographic order)

   1) Include the number of local degrees of freedom (i.e. number per-cell) for the function space. This is an integer and has intent `in`. The name of this argument is `"ndf_"<field_function_space>`.

   2) If there is a field on this space

      1) Include the unique number of degrees of freedom for the function space. This is an integer and has intent `in`. The name of this argument is `"undf_"<field_function_space>`.

      2) Include the dofmap for this function space. This is an integer array with intent `in`. It has one dimension sized by the local degrees of freedom for the function space.

   3) For each operation on the function space (`basis`, `diff_basis`, `orientation`) in the order specified in the metadata

---

10.3. Kernel   **73**

1) if it is a basis or differential basis function then we must pass real arrays of kind `r_def` with intent `in`:

    1) if `gh_shape` is `gh_quadrature_xyoz` then the arrays are of rank 4 with extent (`dimension`, `number_of_dofs`, `np_xy`, `np_z`). The name of the argument is `"basis_"<field_function_space>` or `"diff_basis_"<field_function_space>`, as appropriate.

    2) if `gh_shape` is `gh_evaluator` then we pass one array for each target function space (e.g. as specified by `gh_evaluator_targets`). Each of these arrays are of rank 3 with extent (`dimension`, `number_of_dofs`, `ndf_<target_function_space>`). The name of the argument is `"basis_"<field_function_space>"_on_"<target_function_space>` or `"diff_basis_"<field_function_space>"_on_"<target_function_space>`, as appropriate.

    where `dimension` is 1 or 3 and depends upon the function space and whether or not it is a basis or a differential basis function. For the former it is (w0=1, w1=3, w2=3, w3=1, wtheta=1, w2h=3, w2v=3, any_w2=3). For the latter it is (w0=3, w1=3, w2=1, w3=3, wtheta=3, w2h=1, w2v=1, any_w2=3). `number_of_dofs` is the number of degrees of freedom (dofs) associated with the function space and `np_*` are the number of points to be evaluated: i) `*_xyz` in all directions (3D); ii) `*_xy` in the horizontal plane (2D); iii) `*_x`, `*_y` in the horizontal (1D); and iv) `*_z` in the vertical (1D).

2) If it is an orientation array, include the associated argument. The argument is an integer array with intent `in`. There is one dimension of size the local degrees of freedom for the function space. The name of the array is `"orientation_"<field_function_space>`.

5) If Quadrature is required (`gh_shape = gh_quadrature_*`)

    1) include integer, scalar arguments with intent `in` that specify the extent of the basis/diff-basis arrays:

        1) if `gh_shape` is `gh_quadrature_XYoZ` then pass `np_xy` and `np_z`.

    2) include weights which are real arrays of kind `r_def`:

        1) If `gh_quadrature_XYoZ` pass in `w_XZ(np_xy)` and `w_Z(np_z)`

Examples

For instance, if a kernel has only one written argument and requires an evaluator then its metadata might be:

```
type, extends(kernel_type) :: testkern_operator_type
   type(arg_type), dimension(2) :: meta_args =      &
       (/ arg_type(gh_operator, gh_write, w0, w1), &
          arg_type(gh_field*3, gh_read, w0) /)
   type(func_type) :: meta_funcs(1) =               &
       (/ func_type(w0, gh_basis) /)
   integer :: iterates_over = cells
   integer :: gh_shape = gh_evaluator
 contains
   procedure, nopass :: code => testkern_operator_code
end type testkern_operator_type
```

then we only pass the basis functions evaluated on W0 (the space of the written kernel argument). The subroutine arguments will therefore be:

```
subroutine testkern_operator_code(cell, nlayers, ncell_3d,       &
    local_stencil, xdata, ydata, zdata, ndf_w0, undf_w0, map_w0, &
    basis_w0_on_w0, ndf_w1)
```

where `local_stencil` is the operator, `xdata`, `ydata` etc. are the three components of the field vector and `map_w0` is the dof map for the W0 function space.

If instead, `gh_evaluator_targets` is specified in the metadata:

```
type, extends(kernel_type) :: testkern_operator_type
  type(arg_type), dimension(2) :: meta_args =       &
       (/ arg_type(gh_operator, gh_write, w0, w1), &
          arg_type(gh_field*3, gh_read, w0) /)
  type(func_type) :: meta_funcs(1) =                &
       (/ func_type(w0, gh_basis) /)
  integer :: iterates_over = cells
  integer :: gh_shape = gh_evaluator
  integer :: gh_evaluator_targets(2) = (/W0, W1/)
 contains
  procedure, nopass :: code => testkern_operator_code
end type testkern_operator_type
```

then we will need to pass two sets of basis functions (evaluated at W0 and at W1):

```
subroutine testkern_operator_code(cell, nlayers, ncell_3d,        &
    local_stencil, xdata, ydata, zdata, ndf_w0, undf_w0, map_w0, &
    basis_w0_on_w0, basis_w0_on_w1, ndf_w1)
```

Rules for CMA Kernels

Kernels involving CMA operators are restricted to just three types; assembly, application/inverse-application and matrix-matrix. We give the rules for each of these in the sections below.

Assembly

An assembly kernel requires the column-banded dofmap for both the to- and from-function spaces of the CMA operator being assembled as well as the number of dofs for each of the dofmaps. The full set of rules is:

1) Include the `cell` argument. `cell` is an integer and has intent `in`.

2) Include `nlayers`, the number of layers in a column. `nlayers` is an integer and has intent `in`.

3) Include the number of cells in the 2D mesh, `ncell_2d`, which is an integer with intent `in`.

4) Include the total number of cells, `ncell_3d`, which is an integer with intent `in`.

5) For each argument in the `meta_args` metadata array:

   1) If it is a LMA operator, include a real, 3-dimensional array of type `r_def`. The first two dimensions are the local degrees of freedom for the `to` and `from` spaces, respectively. The third dimension is `ncell_3d`.

   2) If it is a CMA operator, include a real, 3-dimensional array of type `r_def`. The first dimension is is `"bandwidth_"<operator_name>`, the second is `"nrow_"<operator_name>`, and the third is `ncell_2d`.

      1) Include the number of rows in the banded matrix. This is an integer with intent `in` and is named as `"nrow_"<operator_name>`.

      2) If the from-space of the operator is *not* the same as the to-space then include the number of columns in the banded matrix. This is an integer with intent `in` and is named as `"ncol_"<operator_name>`.

      3) Include the bandwidth of the banded matrix. This is an integer with intent `in` and is named as `"bandwidth_"<operator_name>`.

4) Include banded-matrix parameter `alpha`. This is an integer with intent `in` and is named as `"alpha_"<operator_name>`.

5) Include banded-matrix parameter `beta`. This is an integer with intent `in` and is named as `"beta_"<operator_name>`.

6) Include banded-matrix parameter `gamma_m`. This is an integer with intent `in` and is named as `"gamma_m_"<operator_name>`.

7) Include banded-matrix parameter `gamma_p`. This is an integer with intent `in` and is named as `"gamma_p_"<operator_name>`.

3) If it is a field or scalar argument then include arguments following the same rules as for general-purpose kernels.

6) For each unique function space in the order they appear in the metadata arguments (the `to` function space of an operator is considered to be before the `from` function space of the same operator as it appears first in lexicographic order):

1) Include the number of degrees of freedom per cell for the space. This is an integer with intent `in`. The name of this argument is `"ndf_"<arg_function_space>`.

2) If there is a field on this space then:

1) Include the unique number of degrees of freedom for the function space. This is an integer and has intent `in`. The name of this argument is `"undf_"<field_function_space>`.

2) Include the dofmap for this space. This is an integer array with intent `in`. It has one dimension sized by the local degrees of freedom for the function space.

3) If the CMA operator has this space as its to/from space then include the column-banded dofmap, the list of offsets for the to/from-space. This is an integer array of rank 2. The first dimension is `"ndf_"<arg_function_space>``` and the second is `nlayers`.

Application/Inverse-Application

A kernel applying a CMA operator requires the column-indirection dofmap for both the to- and from-function spaces of the CMA operator. Since it does not have any LMA operator arguments it does not require the `ncell_3d` and `nlayers` scalar arguments. (Since a column-wise operator is, by definition, assembled for a whole column, there is no loop over levels when applying it.) The full set of rules is then:

1) Include the `cell` argument. `cell` is an integer and has intent `in`.

2) Include the number of cells in the 2D mesh, `ncell_2d`, which is an integer with intent `in`.

3) For each argument in the `meta_args` metadata array:

1) If it is a field, include the field array. This is a real array of kind `r_def` and is of rank 1. The field array name is currently specified as being `"field_"<argument_position>"_"<field_function_space>`. The extent of the array is the number of unique degrees of freedom for the function space that the field is on. This value is passed in separately. The intent of the argument is determined from the metadata (see *meta_args*).

2) If it is a CMA operator, include it and its associated parameters (see Rule 5 of CMA Assembly kernels).

4) For each of the unique function spaces encountered in the metadata arguments (the `to` function space of an operator is considered to be before the `from` function space of the same operator as it appears first in lexicographic order):

1) Include the number of degrees of freedom per cell for the associated function space. This is an integer with intent `in`. The name of this argument is `"ndf_"<field_function_space>`.

2) Include the number of unique degrees of freedom for the associated function space. This is an integer with intent `in`. The name of this argument is `"undf_"<field_function_space>`.

3) Include the dofmap for this function space. This is a rank-1 integer array with extent equal to the number of degrees of freedom of the space (`"ndf_"<field_function_space>`).

5) Include the indirection map for the to-space of the CMA operator. This is a rank-1 integer array with extent `nrow`.

6) If the from-space of the operator is *not* the same as the to-space then include the indirection map for the from-space of the CMA operator. This is a rank-1 integer array with extent `ncol`.

Matrix-Matrix

Does not require any dofmaps and also does not require the `nlayers` and `ncell_3d` scalar arguments. The full set of rules are then:

1) Include the `cell` argument. `cell` is an integer and has intent `in`.

2) Include the number of cells in the 2D mesh, `ncell_2d`, which is an integer with intent `in`.

3) For each CMA operator or scalar argument specifed in metadata:

   1) If it is a CMA operator, include it and its associated parameters (see Rule 5 of CMA Assembly kernels).

   2) If it is a scalar argument include the corresponding Fortran variable in the argument list with intent `in`.

Rules for Inter-Grid Kernels

As already specified, inter-grid kernels are only permitted to take fields and/or field-vectors as arguments. Fields (and field-vectors) that are on different meshes must be on different function spaces. Fields on the same mesh must also be on the same function space.

Argument ordering follows the general pattern used for 'normal' kernels with field data being followed by dofmap data. The rules for arguments to inter-grid kernels are as follows:

1) Include `nlayers`, the number of layers in a column. `nlayers` is an integer and has intent `in`.

2) Include the `cell_map` for the current cell (column). This is an integer array of rank one and intent `in` which provides the mapping from the coarse to the fine mesh. It has extent *ncell_f_per_c*.

3) Include `ncell_f_per_c`, the number of fine cells per coarse cell. This is an integer and has intent `in`.

4) Include `ncell_f`, the number of cells (columns) in the fine mesh. This is an integer and has intent `in`.

5) For each argument in the `meta_args` metadata array (which must be a field or field-vector):

   1) Pass in field data as done for a regular kernel.

6) For each unique function space (of which there will currently be two) in the order in which they are encountered in the `meta_args` metadata array, include dofmap information:

   If the dofmap is associated with an argument on the fine mesh:

   1) Include `ndf_fine`, the number of DoFs per cell for the FS of the field on the fine mesh.

   2) Include `undf_fine`, the number of unique DoFs per cell for the FS of the field on the fine mesh.

   3) Include `dofmap_fine`, the *whole* dofmap for the fine mesh. This is an integer array of rank two with intent `in`. The extent of the first dimension is `ndf_fine` and that of the second is `ncell_f`.

   else, the dofmap is associated with an argument on the coarse mesh:

1) Include `undf_coarse`, the number of unique DoFs for the coarse field. This is an integer with intent `in`.

2) Include `dofmap_coarse`, the dofmap for the current cell (column) in the coarse mesh. This is an integer array of rank one and has intent `in`.

10.4 Built-ins

The basic concept of a PSyclone Built-in is described in the *Built-ins* section. In the Dynamo 0.3 API, calls to Built-ins generally follow a convention that the field/scalar written to comes first in the argument list. Dynamo 0.3 Built-ins must conform to the following four rules:

1) Built-in kernels must have one and only one modified (i.e. written to) argument.

2) There must be at least one field in the argument list. This is so that we know the number of dofs to iterate over.

3) Kernel arguments must be either fields or scalars.

4) All field arguments to a given Built-in must be on the same function space. This is because all current Built-ins iterate over dofs and therefore all fields should have the same number. It also means that we can determine the number of dofs uniquely when a scalar is written to.

The Built-ins supported for the Dynamo 0.3 API are listed in the related subsections, grouped by the mathematical operation they perform. For clarity, the calculation performed by each Built-in is described using Fortran array syntax; this does not necessarily reflect the actual implementation of the Built-in (*e.g.* it could be implemented by PSyclone generating a call to an optimised maths library).

10.4.1 Naming scheme

The supported Built-ins in the Dynamo 0.3 API are named according to the scheme presented below. Any new Built-in needs to comply with these rules.

1) Ordering of arguments in Built-ins calls follows *LHS (result) <- RHS (operation on arguments)* direction, except where a Built-in returns the *LHS* result to one of the *RHS* arguments. In that case ordering of arguments remains as in the *RHS* expression, with the returning *RHS* argument written as close to the *LHS* as it can be without affecting the mathematical expression.

2) Field names begin with upper case in short form (e.g. **X**, **Y**, **Z**) and any case in long form (e.g. **Field1**, **field**).

3) Scalar names begin with lower case: e.g. **a**, **b**, are **scalar1**, **scalar2**. Special names for scalars are: **constant** (or **c**), **innprod** (inner/scalar product of two fields) and **sumfld** (sum of a field).

4) Arguments in Built-ins variable declarations and constructs (PSyclone Fortran and Python definitions):

   a) Are always written in long form and lower case (e.g. **field1**, **field2**, **scalar1**, **scalar2**);

   b) *LHS* result arguments are always listed first;

   c) *RHS* arguments are listed in order of appearance in the mathematical expression, except when one of them is the *LHS* result.

5) Built-ins names in Fortran consist of:

   1) *RHS* arguments in short form (e.g. **X**, **Y**, **a**, **b**) only;

   2) Descriptive name of mathematical operation on *RHS* arguments in the form `<operationname>_<RHSarg>` for one *RHS* argument or `<RHSargs>_<operationname>_<RHSargs>` for more;

3) Prefix `"inc_"` where the result is returned to one of the *RHS* arguments (i.e. `"inc_"<RHSargs>_<operationname>_<RHSargs>`).

6) Built-ins names in Python definitions are similar to their Fortran counterparts, with a few differences:

   1) Operators and *RHS* arguments are all in upper case (e.g. **X**, **Y**, **A**, **B**, **Plus**, **Minus**);

   2) There are no underscores;

   3) Common prefix is `"Dyn"`, common suffix is `"Kern"`.

10.4.2 Addition

Built-ins which add (scaled) fields are denoted with the keyword **plus**.

X_plus_Y

X_plus_Y (*field3*, *field1*, *field2*)

Sums two fields (Z = X + Y):

```
field3(:) = field1(:) + field2(:)
```

where:

   • type(field_type), intent(out) :: *field3*

   • type(field_type), intent(in) :: *field1*

   • type(field_type), intent(in) :: *field2*

inc_X_plus_Y

inc_X_plus_Y (*field1*, *field2*)

Adds the second field to the first and returns it (X = X + Y):

```
field1(:) = field1(:) + field2(:)
```

where:

   • type(field_type), intent(inout) :: *field1*

   • type(field_type), intent(in) :: *field2*

aX_plus_Y

aX_plus_Y (*field3*, *scalar*, *field1*, *field2*)

Performs Z = aX + Y:

```
field3(:) = scalar*field1(:) + field2(:)
```

where:

   • real(r_def), intent(in) :: *scalar*

   • type(field_type), intent(out) :: *field3*

   • type(field_type), intent(in) :: *field1*, *field2*

inc_aX_plus_Y

inc_aX_plus_Y (*scalar*, *field1*, *field2*)

Performs X = aX + Y (increments the first field):

```
field1(:) = scalar*field1(:) + field2(:)
```

where:

- real(r_def), intent(in) :: *scalar*
- type(field_type), intent(inout) :: *field1*
- type(field_type), intent(in) :: *field2*

inc_X_plus_bY

inc_X_plus_bY (*field1*, *scalar*, *field2*)

Performs X = X + bY (increments the first field):

```
field1(:) = field1(:) + scalar*field2(:)
```

where:

- real(r_def), intent(in) :: *scalar*
- type(field_type), intent(inout) :: *field1*
- type(field_type), intent(in) :: *field2*

aX_plus_bY

aX_plus_bY (*field3*, *scalar1*, *field1*, *scalar2*, *field2*)

Performs Z = aX + bY:

```
field3(:) = scalar1*field1(:) + scalar2*field2(:)
```

where:

- real(r_def), intent(in) :: *scalar1*, *scalar2*
- type(field_type), intent(out) :: *field3*
- type(field_type), intent(in) :: *field1*, *field2*

inc_aX_plus_bY

inc_aX_plus_bY (*scalar1*, *field1*, *scalar2*, *field2*)

Performs X = aX + bY (increments the first field):

```
field1(:) = scalar1*field1(:) + scalar2*field2(:)
```

where:

- real(r_def), intent(in) :: *scalar1*, *scalar2*

- type(field_type), intent(inout) :: *field1*
- type(field_type), intent(in) :: *field2*

10.4.3 Subtraction

Built-ins which subtract (scaled) fields are denoted with the keyword **minus**.

X_minus_Y

X_minus_Y (*field3*, *field1*, *field2*)

Subtracts the second field from the first and stores the result in the third (Z = X - Y):

```
field3(:) = field1(:) - field2(:)
```

where:

- type(field_type), intent(out) :: *field3*
- type(field_type), intent(in) :: *field1*
- type(field_type), intent(in) :: *field2*

inc_X_minus_Y

inc_X_minus_Y (*field1*, *field2*)

Subtracts the second field from the first and returns it (X = X - Y):

```
field1(:) = field1(:) - field2(:)
```

where:

- type(field_type), intent(inout) :: *field1*
- type(field_type), intent(in) :: *field2*

aX_minus_Y

aX_minus_Y (*field3*, *scalar*, *field1*, *field2*)

Performs Z = aX - Y:

```
field3(:) = scalar*field1(:) - field2(:)
```

where:

- real(r_def), intent(in) :: *scalar*
- type(field_type), intent(out) :: *field3*
- type(field_type), intent(in) :: *field1*, *field2*

X_minus_bY

X_minus_bY (*field3*, *field1*, *scalar*, *field2*)

Performs Z = X - bY:

```
field3(:) = field1(:) - scalar*field2(:)
```

where:

- real(r_def), intent(in) :: *scalar*
- type(field_type), intent(out) :: *field3*
- type(field_type), intent(in) :: *field1*, *field2*

inc_X_minus_bY

inc_X_minus_bY (*field1*, *scalar*, *field2*)

Performs X = X - bY (increments the first field):

```
field1(:) = field1(:) - scalar*field2(:)
```

where:

- real(r_def), intent(in) :: *scalar*
- type(field_type), intent(inout) :: *field1*
- type(field_type), intent(in) :: *field2*

10.4.4 Multiplication

Built-ins which multiply (scaled) fields are denoted with the keyword **times**.

X_times_Y

X_times_Y (*field3*, *field1*, *field2*)

Multiplies two fields together and returns the result in a third field (Z = X*Y):

```
field3(:) = field1(:)*field2(:)
```

where:

- type(field_type), intent(out) :: *field3*
- type(field_type), intent(in) :: *field1*, *field2*

inc_X_times_Y

inc_X_times_Y (*field1*, *field2*)

Multiplies the first field by the second and returns it (X = X*Y):

```
field1(:) = field1(:)*field2(:)
```

where:

- type(field_type), intent(inout) :: *field1*
- type(field_type), intent(in) :: *field2*

inc_aX_times_Y

inc_aX_times_Y (*scalar*, *field1*, *field2*)

Performs X = a*X*Y (increments the first field):

```
field1(:) = scalar*field1(:)*field2(:)
```

where:

- real(r_def), intent(in) :: *scalar*
- type(field_type), intent(inout) :: *field1*
- type(field_type), intent(in) :: *field2*

10.4.5 Scaling

Built-ins which scale fields are technically cases of multiplying a field by a scalar and are hence also denoted with the keyword **times**.

a_times_X

a_times_X (*field2*, *scalar*, *field1*)

Multiplies a field by a scalar and stores the result in a second field (Y = a*X):

```
field2(:) = scalar*field1(:)
```

where:

- real(r_def), intent(in) :: *scalar*
- type(field_type), intent(out) :: *field2*
- type(field_type), intent(in) :: *field1*

inc_a_times_X

inc_a_times_X (*scalar*, *field*)

Multiplies a field by a scalar value and returns the field (X = a*X):

```
field(:) = scalar*field(:)
```

where:

- real(r_def), intent(in) :: *scalar*
- type(field_type), intent(inout) :: *field*

10.4.6 Division

Built-ins which divide (scaled) fields are denoted with the keyword **divideby**.

X_divideby_Y

X_divideby_Y (*field3*, *field1*, *field2*)

Divides the first field by the second and returns the result in the third (Z = X/Y):

```
field3(:) = field1(:)/field2(:)
```

where:

- type(field_type), intent(out) :: *field3*
- type(field_type), intent(in) :: *field1*, *field2*

inc_X_divideby_Y

inc_X_divideby_Y (*field1*, *field2*)

Divides the first field by the second and returns it (X = X/Y):

```
field1(:) = field1(:)/field2(:)
```

where:

- type(field_type), intent(inout) :: *field1*
- type(field_type), intent(in) :: *field2*

10.4.7 Setting to value

Built-ins which set field elements to some value and hence are denoted with the keyword **setval**.

setval_c

setval_c (*field*, *constant*)

Sets all elements of the field *field* to the value *constant* (X = c):

```
field(:) = constant
```

where:

- type(field_type), intent(out) :: *field*
- real(r_def), intent(in) :: *constant*

Note: The field may be on any function space.

---

setval_X

setval_X (*field2*, *field1*)

Sets a field *field2* equal to field *field1* (Y = X):

```
field2(:) = field1(:)
```

where:

- type(field_type), intent(out) :: *field2*
- type(field_type), intent(in) :: *field1*

10.4.8 Raising to power

Built-ins which raise field elements to an exponent are denoted with the keyword **powreal** for real exponent or **powint** for integer exponent.

inc_X_powreal_a

inc_X_powreal_a (*field*, *rscalar*)

Raises a field to a real scalar value and returns the field (X = X**a):

```
field(:) = field(:)**rscalar
```

where:

- type(field_type), intent(inout) :: *field*
- real(r_def), intent(in) :: *rscalar*

inc_X_powint_n

inc_X_powint_n (*field*, *iscalar*)

Raises a field to an integer scalar value and returns the field (X = X**n):

```
field(:) = field(:)**iscalar
```

where:

- type(field_type), intent(inout) :: *field*
- integer(i_def), intent(in) :: *iscalar*

10.4.9 Inner product

Built-ins which calculate the inner product of two fields or of a field with itself are denoted with the keyword **inner-product**.

X_innerproduct_Y

X_innerproduct_Y (*innprod*, *field1*, *field2*)

Computes the inner product of the fields *field1* and *field2*, *i.e.*:

```
innprod = SUM(field1(:)*field2(:))
```

where:

- real(r_def), intent(out) :: *innprod*
- type(field_type), intent(in) :: *field1*, *field2*

Note: When used with distributed memory this Built-in will trigger the addition of a global sum which may affect the performance and/or scalability of the code.

X_innerproduct_X

X_innerproduct_X (*innprod*, *field*)

Computes the inner product of the field *field1* by itself, *i.e.*:

```
innprod = SUM(field(:)*field(:))
```

where:

- real(r_def), intent(out) :: *innprod*
- type(field_type), intent(in) :: *field*

Note: When used with distributed memory this Built-in will trigger the addition of a global sum which may affect the performance and/or scalability of the code.

10.4.10 Sum of elements

Built-in which sums the elements of a field is denoted with the keyword *sum*.

sum_X

sum_X (*sumfld*, *field*)

Sums all of the elements of the field *field* and returns the result in the scalar variable *sumfld*:

```
sumfld = SUM(field(:))
```

where:

- real(r_def), intent(out) :: sumfld
- type(field_type), intent(in) :: field

Note: When used with distributed memory this Built-in will trigger the addition of a global sum which may affect the performance and/or scalability of the code.

10.5 Boundary Conditions

In the Dynamo0.3 API, boundary conditions for a field or LMA operator can be enforced by the algorithm developer by calling the Kernels `enforce_bc_type` or `enforce_operator_bc_type`, respectively. These kernels take a field or operator as input and apply boundary conditions. For example:

```fortran
call invoke( kernel_type(field1, field2),      &
             enforce_bc_type(field1),          &
             kernel_with_op_type(field1, op1), &
             enforce_operator_bc_type(op1)     &
           )
```

The particular boundary conditions that are applied are not known by PSyclone, PSyclone simply recognises these kernels by their names and passes pre-specified dofmap and boundary_value arrays into the kernel implementations, the contents of which are set by the LFRic infrastructure.

Up to and including version 1.4.0 of PSyclone, boundary conditions were applied automatically after a call to `matrix_vector_type` if the field arguments were on a vector function space (one of `w1`, `w2`, `w2h` or `w2v`). With the subsequent introduction of the ability to apply boundary conditions to operators this functionality is no longer required and has been removed.

Example `eg4` in the `examples/dynamo` directory includes a call to `enforce_bc_kernel_type` so can be used to see the boundary condition code that is added by PSyclone. See the `README` in the `examples/dynamo` directory for instructions on how to run this example.

An example of applying boundary conditions to an operator is the kernel `enforce_operator_bc_kernel_mod.F90` in the `<PSYCLONEHOME>/src/psyclone/tests/test_files/dynamo0p3` directory. Since operators are discontinuous quantities, updating their values can be safely performed in parallel (see Section *Kernel*). The `GH_READWRITE` access is used for updating discontinuous operators (see subsection *Valid Access Modes* for more details).

10.6 Conventions

The naming of Dynamo0.3 API kernels and associated entities (types, subroutines and modules) follows the PSyclone Fortran naming conventions (see *Fortran Naming Conventions*). However, PSyclone does not need this convention to be followed apart from the stub generator (see the *Stub Generation* Section) where the name of the metadata to be parsed is determined from the module name.

The contents of the metadata is also usually declared private but this does not affect PSyclone.

Finally, the `procedure` metadata (located within the kernel metadata) usually has `nopass` specified but again this is ignored by PSyclone.

10.7 Configuration

10.7.1 Annexed DoFs

When a kernel iterates over dofs (rather than cells) for a continuous field using distributed memory (see the *Distributed Memory* Section), then PSyclone need only ensure that dofs owned by a processor are computed. However, for continuous fields, shared dofs at the boundary between processors must be replicated (as different cells share the same dof). Only one processor can own a dof, therefore processors will have continuous fields which contain dofs that the processor does not own. These unowned dofs are called *annexed* in the dynamo0.3 api and are a separate, but related, concept to field halos.

When a kernel that iterates over cells needs to read a continuous field then the annexed dofs must be up-to-date on all processors. If they are not then a halo exchange must be added. Currently PSyclone defaults, for kernels which iterate over dofs, to iterating over only owned dofs. This behaviour can be changed by setting *COMPUTE_ANNEXED_DOFS* to `true` in the *dynamo0.3* section of the configuration file (see the *Configuration* section). PSyclone will then generate code to iterate over both owned and annexed dofs, thereby reducing the number of halo exchanges required (at the expense of redundantly computing annexed dofs). For more details please refer to the *Dynamo0.3* developers section.

10.8 Transformations

This section describes the dynamo-api-specific transformations. In all cases, excepting **Dynamo0p3RedundantComputationTrans** and **Dynamo0p3AsyncHaloExchangeTrans**, these transformations are specialisations of generic transformations described in the *Transformations* section. The difference between these transformations and the generic ones is that these perform dynamo-api-specific checks to make sure the transformations are valid. In practice these transformations perform the required checks then call the generic ones internally.

The use of the dynamo-api-specific transformations is exactly the same as the equivalent generic ones in all cases excepting **DynamoLoopFuseTrans**. In this case an additional optional argument **same_space** has been added to the **apply** method. The reason for this is to allow loop fusion when one or more of the iteration-spaces is determined by a function space that is unknown by PSyclone at compile time. This is the case when the **ANY_SPACE** function space is specified in the Kernel metadata. By default PSyclone will not allow loop fusion if it does not know the spaces are the same. The **same_space** option allows the user to specify that the spaces are the same. This option should therefore be used with caution. Note, if PSyclone knows the spaces are different this option has no effect and the transformation will always raise an exception.

The **Dynamo0p3RedundantComputationTrans** and **Dynamo0p3AsyncHaloExchange** transformations are only valid for the "Dynamo0p3" API. This is because this API is currently the only one that supports distributed memory. An example of redundant computation can be found in `examples/dynamo/eg8` and an example of asynchronous halo exchanges can be found in `examples/dynamo/eg11`.

The Dynamo-specific transformations currently available are given below. If the name of a transformation includes "Dynamo0p3" it means that the transformation is only valid for this particular API. If the name of the transformation includes "Dynamo" then it should work with all versions of the Dynamo API.

Note: Only the loop-colouring and OpenMP transformations are currently supported for loops that contain inter-grid kernels. Attempting to apply other transformation types will result in PSyclone raising an error.

---

class `psyclone.transformations.`**`DynamoExtractRegionTrans`**
    Dynamo0.3 API application of ExtractRegionTrans transformation to extract code into a stand-alone program. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>>
>>> API = "dynamo0.3"
>>> FILENAME = "solver_alg.x90"
>>> ast, invokeInfo = parse(FILENAME, api=API)
>>> psy = PSyFactory(API, distributed_memory=False).create(invoke_info)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>>
>>> from psyclone.transformations import DynamoExtractRegionTrans
>>> etrans =  DynamoExtractRegionTrans()
>>>
>>> # Apply DynamoExtractRegionTrans transformation to selected Nodes
>>> newsched, _ = etrans.apply(schedule.children[0:3])
>>> newsched.view()
```

name
> Returns the name of this transformation as a string.

class psyclone.transformations.**DynamoLoopFuseTrans**
> Performs error checking before calling the *apply()* method of the *base class* in order to fuse two Dynamo loops.

> **apply**(*node1*, *node2*, *same_space=False*)
> > Fuse the two Dynamo loops represented by node1 and node2. The optional same_space flag asserts that an unknown iteration space (i.e. any_space) matches the other iteration space. This is set at the users own risk.

> > **Parameters**
> > > - **node1** (psyclone.dynamo0p3.DynLoop) – First Loop to fuse.
> > > - **node2** (psyclone.dynamo0p3.DynLoop) – Second Loop to fuse.

> > **Returns** two-tuple of modified Schedule and Memento

> > **Return type** *psyclone.psyGen.Schedule*, psyclone.undoredo.Memento

> > **Raises** *TransformationError* – if either of the supplied loops contains an inter-grid kernel.

name
> Returns the name of this transformation as a string.

class psyclone.transformations.**DynamoOMPParallelLoopTrans**(*omp_schedule='static'*)
> Dynamo-specific OpenMP loop transformation. Adds Dynamo specific validity checks. Actual transformation is done by the *base class*.

> **apply**(*node*)
> > Perform Dynamo specific loop validity checks then call the *apply()* method of the *base class*.

> > **Parameters** **node** (*psyclone.psyGen.Node*) – the Node in the Schedule to check

> > **Raises** *TransformationError* – if the associated loop requires colouring.

name
> Returns the name of this transformation as a string.

class psyclone.transformations.**Dynamo0p3AsyncHaloExchangeTrans**
> Splits a synchronous halo exchange into a halo exchange start and halo exchange end. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "dynamo0.3"
>>> ast, invokeInfo = parse("file.f90", api=api)
>>> psy=PSyFactory(api).create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import Dynamo0p3AsyncHaloExchangeTrans
>>> trans = Dynamo0p3AsyncHaloExchangeTrans()
>>> new_schedule, memento = trans.apply(schedule.children[0])
>>> new_schedule.view()
```

apply(*node*)

> Transforms a synchronous halo exchange, represented by a HaloExchange node, into an asynchronous halo exchange, represented by HaloExchangeStart and HaloExchangeEnd nodes.
>
>> **Parameters** **node** (`psyclone.psygen.HaloExchange`) – A synchronous haloexchange node
>>
>> **Returns** Tuple of the modified schedule and a record of the transformation.
>>
>> **Return type** (*psyclone.psyGen.Schedule*, `psyclone.undoredo.Memento`)

name

>> **Returns** the name of this transformation as a string.
>>
>> **Return type** str

class psyclone.transformations.**Dynamo0p3ColourTrans**

> Split a Dynamo 0.3 loop over cells into colours so that it can be parallelised. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> import transformations
>>> import os
>>> import pytest
>>>
>>> TEST_API = "dynamo0.3"
>>> _,info=parse(os.path.join(os.path.dirname(os.path.abspath(__file__)),
>>>              "tests", "test_files", "dynamo0p3",
>>>              "4.6_multikernel_invokes.f90"),
>>>              api=TEST_API)
>>> psy = PSyFactory(TEST_API).create(info)
>>> invoke = psy.invokes.get('invoke_0')
>>> schedule = invoke.schedule
>>>
>>> ctrans = Dynamo0p3ColourTrans()
>>> otrans = DynamoOMPParallelLoopTrans()
>>>
>>> # Colour all of the loops
>>> for child in schedule.children:
>>>     cschedule, _ = ctrans.apply(child)
>>>
>>> # Then apply OpenMP to each of the colour loops
>>> schedule = cschedule
>>> for child in schedule.children:
>>>     newsched, _ = otrans.apply(child.children[0])
```

```
>>>
>>> newsched.view()
```

Colouring in the Dynamo 0.3 API is subject to the following rules:

- Only kernels with an iteration space of CELLS and which modify a continuous field require colouring. Any other type of loop will cause this transformation to raise an exception.

- A kernel may have at most one field with 'INC' access

- A separate colour map will be required for each field that is coloured (if an invoke contains >1 kernel call)

apply(*node*)
    Performs Dynamo0.3-specific error checking and then uses the parent class to convert the Loop represented by `node` into a nested loop where the outer loop is over colours and the inner loop is over cells of that colour.

> **Parameters node** (`psyclone.dynamo0p3.DynLoop`) – the loop to transform.
>
> **Returns** 2-tuple of new schedule and memento of transform
>
> **Return type** (`psyclone.dynamo0p3.DynInvokeSchedule`,   `psyclone.undoredo.Memento`)

name
    Returns the name of this transformation as a string.

class psyclone.transformations.**Dynamo0p3OMPLoopTrans**(*omp_schedule='static'*)
    Dynamo 0.3 specific orphan OpenMP loop transformation. Adds Dynamo-specific validity checks. Actual transformation is done by *base class*.

apply(*node*, *reprod=None*)
    Perform Dynamo 0.3 specific loop validity checks then call *OMPLoopTrans.apply()*.

> **Parameters**
>
> - **node** (*psyclone.psyGen.Node*) – the Node in the Schedule to check
>
> - **reprod**   (*bool or None (default, which indicates to use the default from the config file)*) – if reproducible reductions should be used.
>
> **Raises** ***TransformationError*** – if an OMP loop transform would create incorrect code.

name
    Returns the name of this transformation as a string.

class psyclone.transformations.**Dynamo0p3RedundantComputationTrans**
    This transformation allows the user to modify a loop's bounds so that redundant computation will be performed. Redundant computation can result in halo exchanges being modified, new halo exchanges being added or existing halo exchanges being removed.

- This transformation should be performed before any parallelisation transformations (e.g. for OpenMP) to the loop in question and will raise an exception if this is not the case.

- This transformation can not be applied to a loop containing a reduction and will again raise an exception if this is the case.

- This transformation can only be used to add redundant computation to a loop, not to remove it.

- This transformation allows a loop that is already performing redundant computation to be modified, but only if the depth is increased.

apply (*loop*, *depth=None*)

Apply the redundant computation tranformation to the loop `loop`. This transformation can be applied to loops iterating over 'cells or 'dofs'. if `depth` is set to a value then the value will be the depth of the field's halo over which redundant computation will be performed. If `depth` is not set to a value then redundant computation will be performed to the full depth of the field's halo.

> **Parameters**
>
> - **loop** (`psyclone.psyGen.DynLoop`) – the loop that we are transforming
>
> - **depth** (*int or None*) – the depth of the stencil. Defaults to None if a depth is not provided.

name

Returns the name of this transformation as a string.

GOCEAN1.0 API

11.1 Introduction

The GOcean 1.0 application programming interface (API) was originally designed to support ocean models that use the finite-difference scheme for two-dimensional domains. However, the approach is not specific to ocean models and can potentially be applied to any finite-difference code.

As with all PSyclone APIs, the GOcean 1.0 API specifies how a user must write the Algorithm Layer and the Kernel Layer to allow PSyclone to generate the PSy Layer. These Algorithm and Kernel APIs are discussed separately in the sections below. Before these we describe the functionality provided by the GOcean Library.

11.2 The GOcean Library

The use of PSyclone and the GOcean 1.0 API implies the use of a standard set of data types and associated infrastructure. This is provided by version 1.0 of the GOcean Library (GOLib v.1.0). Currently this library is distributed separately from PSyclone and is available from http://puma.nerc.ac.uk/trac/GOcean.

11.2.1 Grid

The GOLib contains a `grid_mod` module which defines a `grid_type` and associated constructor:

```fortran
use grid_mod
...
!> The grid on which our fields are defined
type(grid_type), target :: model_grid
...
! Create the model grid
model_grid = grid_type(GO_ARAKAWA_C,                              &
                       (/GO_BC_EXTERNAL,GO_BC_EXTERNAL,GO_BC_NONE/), &
                       GO_OFFSET_NE)
```

Note: The grid object itself must be declared with the `target` attribute. This is because each field object will contain a pointer to it.

The `grid_type` constructor takes three arguments:

1. The type of grid (only GO_ARAKAWA_C is currently supported)

2. The boundary conditions on the domain for the *x*, *y* and *z* dimensions (see below). The value for the *z* dimension is currently ignored.

3. The 'index offset' - the convention used for indexing into offset fields.

Three types of boundary condition are currently supported:

Name	Description
GO_BC_NONE	No boundary conditions are applied.
GO_BC_EXTERNAL	Some external forcing is applied. This must be implemented by a kernel. The domain must be defined with a T-point mask (see *The grid_init Routine*).
GO_BC_PERIODIC	Periodic boundary conditions are applied.

The infrastructure requires this information in order to determine the extent of the model grid.

The index offset is required because a model (kernel) developer has choice in how they actually implement the staggering of variables on a grid. This comes down to a choice of which grid points in the vicinity of a given T point have the same array (i, j) indices. In the diagram below, the image on the left corresponds to choosing those points to the South and West of a T point to have the same (i, j) index. That on the right corresponds to choosing those points to the North and East of the T point (this is the offset scheme used in the NEMO ocean model):

The GOcean 1.0 API supports these two different offset schemes, which we term `GO_OFFSET_SW` and `GO_OFFSET_NE`.

Note that the constructor does not specify the extent of the model grid. This is because this information is normally obtained by reading a file (a namelist file, a netcdf file etc.) which is specific to an application. Once this information has been obtained, a second routine, `grid_init`, is provided with which to 'load' a grid object with state. This is discussed below.

The `grid_init` Routine

Once an application has determined the details of the model configuration, it must use this information to populate the grid object. This is done via a call to the `grid_init` subroutine:

```fortran
subroutine grid_init(grid, m, n, dxarg, dyarg, tmask)
  !> The grid object to configure
  type(grid_type), intent(inout) :: grid
  !> Dimensions of the model grid
  integer,         intent(in)    :: m, n
  !> The (constant) grid spacing in x and y (m)
  real(wp),        intent(in)    :: dxarg, dyarg
  !> Optional T-point mask specifying whether each grid point is
  !! wet (1), dry (0) or external (-1).
  integer, dimension(m,n), intent(in), optional :: tmask
```

If no T-mask is supplied then this routine configures the grid appropriately for an all-wet domain with periodic boundary conditions in both the *x*- and *y*-dimensions. It should also be noted that currently only grids with constant resolution in *x* and *y* are supported by this routine.

11.2.2 Fields

Once a model has a grid defined it will require one or more fields. The GOLib contains a `field_mod` module which defines an `r2d_field` type (real, 2-dimensional field) and associated constructor:

```fortran
use field_mod
...
!> Current ('now') sea-surface height at different grid points
type(r2d_field) :: sshn_u_fld, sshn_v_fld, sshn_t_fld
...

! Sea-surface height now (current time step)
sshn_u = r2d_field(model_grid, GO_U_POINTS)
sshn_v = r2d_field(model_grid, GO_V_POINTS)
sshn_t = r2d_field(model_grid, GO_T_POINTS)
```

The constructor takes two arguments:

1. The grid on which the field exists

2. The type of grid point at which the field is defined (`GO_U_POINTS`, `GO_V_POINTS`, `GO_T_POINTS` or `GO_F_POINTS`)

Note that the grid object need not have been fully configured (by a call to `grid_init` for instance) before it is passed into this constructor.

11.2.3 Example

PSyclone is distributed with a full example of the use of the GOcean Library. See `<PSYCLONEHOME>/examples/gocean/shallow_alg.f90`. In what follows we will walk through a slightly cut-down example for a different program.

The following code illustrates the use of the GOLib in constructing an application:

```fortran
program gocean2d
  use grid_mod  ! From dl_esm_inf
  use field_mod ! From dl_esm_inf
  use model_mod
  use boundary_conditions_mod

  !> The grid on which our fields are defined. Must have the 'target'
  !! attribute because each field object contains a pointer to it.
  type(grid_type), target :: model_grid

  !> Current ('now') velocity component fields
  type(r2d_field) :: un_fld, vn_fld
  !> 'After' velocity component fields
  type(r2d_field) :: ua_fld, va_fld
  ...

  ! time stepping index
  integer :: istp

  ! Create the model grid. We use a NE offset (i.e. the U, V and F
  ! points immediately to the North and East of a T point all have the
  ! same i,j index).  This is the same offset scheme as used by NEMO.
  model_grid = grid_type(GO_ARAKAWA_C,                               &
```

```
                        (/GO_BC_EXTERNAL,GO_BC_EXTERNAL,GO_BC_NONE/), &
                        GO_OFFSET_NE)

  !! read in model parameters and configure the model grid
  CALL model_init(model_grid)

  ! Create fields on this grid

  ! Velocity components now (current time step)
  un_fld = r2d_field(model_grid, GO_U_POINTS)
  vn_fld = r2d_field(model_grid, GO_V_POINTS)

  ! Velocity components 'after' (next time step)
  ua_fld = r2d_field(model_grid, GO_U_POINTS)
  va_fld = r2d_field(model_grid, GO_V_POINTS)


  ...

  !! time stepping
  do istp = nit000, nitend, 1

    call step(istp,                                  &
              ua_fld, va_fld, un_fld, vn_fld,     &
              ...)
  end do
  ...
end program gocean2d
```

The `model_init` routine is application specific since it must determine details of the model configuration being run, *e.g.* by reading a namelist file. An example might look something like:

```
subroutine model_init(grid)
  type(grid_type), intent(inout) :: grid

  !> Problem size, read from namelist
  integer :: jpiglo, jpjglo
  real(wp) :: dx, dy
  integer, dimension(:,:), allocatable :: tmask

  ! Read model configuration from namelist
  call read_namelist(jpiglo, jpjglo, dx, dy, &
                     nit000, nitend, irecord, &
                     jphgr_msh, dep_const, rdt, cbfr, visc)

  ! Set-up the T mask. This defines the model domain.
  allocate(tmask(jpiglo,jpjglo))

  call setup_tpoints_mask(jpiglo, jpjglo, tmask)

  ! Having specified the T points mask, we can set up mesh parameters
  call grid_init(grid, jpiglo, jpjglo, dx, dy, tmask)

  ! Clean-up. T-mask has been copied into the grid object.
  deallocate(tmask)

end subroutine model_init
```

Here, only `grid_type` and the `grid_init` routine come from the GOLib. The remaining code is all application specific.

Once the grid object is fully configured and all fields have been constructed, a simulation will proceed by performing calculations with those fields. In the example program given above, this calculation is performed in the time-stepping loop within the `step` subroutine. The way in which this routine uses Invoke calls is described in the *Invokes* Section.

11.3 Algorithm

The Algorithm is the top-level specification of the natural science implemented in the software. Essentially it consists of mesh setup, field declarations, initialisation of fields and (a series of) Kernel calls. Infrastructure to support these tasks is provided in version 1.0 of the GOcean library (see *The GOcean Library*).

11.3.1 Invokes

The Kernels to call are specified through the use of Invokes, e.g.:

```
call invoke( kernel1(field1, field2),                     &
             kernel2(field1, field3)                      &
           )
```

The location and number of these `call invoke(...)` statements within the source code is entirely up to the user. The only requirement is that PSyclone must be run on every source file that contains one or more Invokes. The body of each Invoke specifies the kernels to be called, the order in which they are to be applied and the fields (and scalars) that they work with.

Note that the kernel names specified in an Invoke are the names of the corresponding kernel *types* defined in the kernel metadata (see the *Kernel* Section). These are not the same as the names of the Fortran subroutines which contain the actual kernel code. The kernel arguments are typically field objects, as described in the *Fields* Section, but they may also be scalar quantities (real or integer).

In the example `gocean2d` program shown earlier, there is only one Invoke call and it is contained within the `step` subroutine:

```
subroutine step(istp,                     &
                ua, va, un, vn,           &
                sshn_t, sshn_u, sshn_v, &
                ssha_t, ssha_u, ssha_v, &
                hu, hv, ht)
  use kind_params_mod ! From dl_esm_inf
  use grid_mod        ! From dl_esm_inf
  use field_mod       ! From dl_esm_inf
  use model_mod, only: rdt ! The model time-step
  use continuity_mod,  only: continuity
  use momentum_mod,    only: momentum_u, momentum_v
  use boundary_conditions_mod, only: bc_ssh, bc_solid_u
  !> The current time step
  integer,        intent(inout) :: istp
  type(r2d_field), intent(inout) :: un, vn, sshn_t, sshn_u, sshn_v
  type(r2d_field), intent(inout) :: ua, va, ssha_t, ssha_u, ssha_v
  type(r2d_field), intent(inout) :: hu, hv, ht

  call invoke(                                             &
              continuity(ssha_t, sshn_t, sshn_u, sshn_v,     &
```

(continued from previous page)
```
                            hu, hv, un, vn, rdt),                    &
                   momentum_u(ua, un, vn, hu, hv, ht,                &
                              ssha_u, sshn_t, sshn_u, sshn_v),       &
                   momentum_v(va, un, vn, hu, hv, ht,                &
                              ssha_v, sshn_t, sshn_u, sshn_v),       &
                   bc_ssh(istp, ssha_t),                             &
                   bc_solid_u(ua),                                   &
                   ...
                   )
end subroutine step
```

Note that in this example the grid was constructed for a model with 'external' boundary conditions. These boundary conditions are applied through several user-supplied kernels, two of which (`bc_ssh` and `bc_solid_u`) are include in the above code fragment.

11.4 Kernel

The general requirements for the structure of a Kernel are explained in the *Kernel layer* section. This section explains the metadata and subroutine arguments that are specific to the GOcean 1.0 API.

11.4.1 Metadata

The metadata for a GOcean 1.0 API kernel has four components:

1) 'meta_args',

2) 'iterates_over',

3) 'index_offset' and

4) 'procedure':

These are illustrated in the code below:

```
type, extends(kernel_type) :: my_kernel_type
   type(go_arg), dimension(...) :: meta_args = (/ ... /)
   integer :: iterates_over = ...
   integer :: index_offset = ...
contains
  procedure, nopass :: code => my_kernel_code
end type my_kernel_type
```

These four metadata elements are discussed in order in the following sections.

Argument Metadata: meta_args

The `meta_args` array specifies information about data that the kernel code expects to be passed to it via its argument list. There is one entry in the `meta_args` array for each **scalar**, **field**, or **grid-property** passed into the Kernel. Their ordering in the `meta_args` array must be the same as that in the kernel code argument list. The entry must be of type `go_arg` which itself contains metadata about the associated argument. The size of the meta_args array must correspond to the total number of **scalars**, **fields** and **grid properties** passed into the Kernel.

For example, if there are a total of two **field** entities being passed to the Kernel then the meta_args array will be of size 2 and there will be two entries of type `GO_arg`:

```fortran
type(GO_arg) :: meta_args(2) = (/                                    &
     go_arg( ... ),                                                  &
     go_arg( ... )                                                   &
     /)
```

Argument-metadata (metadata contained within the brackets of an `go_arg` entry), describes either a **scalar**, a **field** or a **grid property**.

The first argument-metadata entry describes how the kernel will access the corresponding argument. As an example, the following `meta_args` metadata describes four entries, the first one is written to by the kernel while the remaining three are only read.

```fortran
type(go_arg) :: meta_args(4) = (/                                    &
     go_arg(GO_WRITE, ... ),                                         &
     go_arg(GO_READ, ... ),                                          &
     go_arg(GO_READ, ... ),                                          &
     go_arg(GO_READ, ...)                                            &
     /)
```

The second entry to argument-metadata (information contained within the brackets of an `go_arg` type) describes the type of data represented by the argument. This type falls into three categories; field data, scalar data and grid properties. For field data the metadata entry consists of the type of grid-point that field values are defined on. Since the GOcean API supports fields on an Arakawa C grid, the possible grid-point types are `GO_CU`, `GO_CV`, `GO_CF` and `GO_CT`. GOcean Kernels can also take scalar quantities as arguments. Since these do not live on grid-points they are specified as either `GO_R_SCALAR` or `GO_I_SCALAR` depending on whether the corresponding Fortran variable is a real or integer quantity. Finally, grid-property entries are used to specify any properties of the grid required by the kernel (*e.g.* the area of cells at U points or whether T points are wet or dry).

For example:

```fortran
type(go_arg) :: meta_args(4) = (/                                    &
     go_arg(GO_WRITE, GO_CT, ... ),                                  &
     go_arg(GO_READ,  GO_CU, ... ),                                  &
     go_arg(GO_READ,  GO_R_SCALAR, ... ),                            &
     go_arg(GO_READ,  GO_GRID_AREA_U)                                &
     /)
```

Here, the first argument is a field on T points, the second is a field on U points, the fourth is a real scalar and the fifth is a property of the grid (cell area at U points).

The full list of supported grid properties in the GOcean 1.0 API is:

Name	Description	Type
go_grid_area_t	Cell area at T point	Real array, rank=2
go_grid_area_u	Cell area at U point	Real array, rank=2
go_grid_area_v	Cell area at V point	Real array, rank=2
go_grid_mask_t	T-point mask (1=wet, 0=dry)	Integer array, rank=2
go_grid_dx_t	Grid spacing in x at T points	Real array, rank=2
go_grid_dx_u	Grid spacing in x at U points	Real array, rank=2
go_grid_dx_v	Grid spacing in x at V points	Real array, rank=2
go_grid_dy_t	Grid spacing in y at T points	Real array, rank=2
go_grid_dy_u	Grid spacing in y at U points	Real array, rank=2
go_grid_dy_v	Grid spacing in y at V points	Real array, rank=2
go_grid_lat_u	Latitude of U points (gphiu)	Real array, rank=2
go_grid_lat_v	Latitude of V points (gphiv)	Real array, rank=2
go_grid_dx_const	Grid spacing in x if constant	Real, scalar
go_grid_dy_const	Grid spacing in y if constant	Real, scalar
go_grid_x_min_index	Minimum X index	Integer, scalar
go_grid_x_max_index	Maximum X index	Integer, scalar
go_grid_y_min_index	Minimum Y index	Integer, scalar
go_grid_y_max_index	Maximum Y index	Integer, scalar

These are stored in a dictionary named `GRID_PROPERTY_DICT` at the top of the `gocean1p0.py` file. All of the rank-two arrays have the first rank as longitude (*x*) and the second as latitude (*y*).

Scalars and fields contain a third argument-metadata entry which describes whether the kernel accesses the corresponding argument with a stencil. The value `GO_POINTWISE` indicates that there is no stencil access. Metadata for a scalar field is limited to this value. Grid-property arguments have no third metadata argument. If there are no stencil accesses then the full argument metadata for our previous example will be:

```fortran
type(go_arg) :: meta_args(4) = (/                           &
     go_arg(GO_WRITE, GO_CT,       GO_POINTWISE),            &
     go_arg(GO_READ,  GO_CU,       GO_POINTWISE),            &
     go_arg(GO_READ,  GO_R_SCALAR, GO_POINTWISE),            &
     go_arg(GO_READ,  GO_GRID_AREA_U)                        &
     /)
```

If a kernel accesses a field using a stencil then the third argument metadata entry should take the form `go_stencil(...)`. Note, a stencil access is only allowed for a field that is `READ` by a kernel.

In the GOcean API, fields are implemented as two-dimensional arrays. In Fortran, a standard 5-point stencil would look something like the following:

```fortran
a(i,j) + a(i+1,j) + a(i-1,j) + a(i,j+1) + a(i,j-1)
```

If we view the above accesses as co-ordinates relative to the `a(i,j)` access we get `(0,0)`, `(1,0)`, `(-1,0)`, `(0,1)`, `(0,-1)`. If we then view these accesses in graphical form with `i` being in the horizontal direction and `j` in the vertical and with a `1` indicating a (depth-1) access and a `0` indicating there is no access we get the following:

```
010
111
010
```

In the GOcean API a stencil access is captured as a triplet of integers (one row at a time from top to bottom) using the above view i.e.

```
go_stencil(010,111,010)
```

So far we have only considered depth-1 stencils. In our notation the depth of access is captured by the integer value (`0` for no access, `1` for depth 1, `2` for depth 2 etc). For example:

```
a(i,j) + a(i,j+1) + a(i,j+2)
```

would be captured as:

```
go_stencil(020,010,000)
```

All forms of stencil can be **summarised** using this triplet notation up to a depth of `9` apart from the central `a(i,j)` value which can either be `0` (not accessed) or `1` (accessed). Note, the central value is not currently used by PSyclone. The notation is a **summary** in two ways

1) it only captures the depth of the stencil in a particular direction, not the actual accesses. Therefore, there is no way to distinguish between the stencil `a(i+2,j)` and the stencil `a(i+1,j) + a(i+2,j)`.

2) when there are offsets for both `i` and `j` e.g. `a(i+1,j+1)` it only captures whether there is an access in that direction at a particular depth, not the details of the access. For example, there is no way to distinguish between `a(i+2,j+2)` and `a(i+2,j+2) + a(i+1,j+2) + a(i+2,j+1)`.

Whilst the description is a summary, it is accurate enough for PSyclone as this information is primarily used to determine which grid partitions must communicate with which for the purposes of placing halo exchange calls. In this case, it is the depth and direction information that is most important.

Iterates Over

The second element of kernel metadata is `ITERATES_OVER`. This specifies that the Kernel has been written with the assumption that it is iterating over grid points of the specified type. By default the supported values are: `GO_INTERNAL_PTS`, `GO_EXTERNAL_PTS` and `GO_ALL_PTS`. These may be understood by considering the following diagram of an example model configuration:

GO_INTERNAL_PTS are then those points that are within the Model domain (fuscia box), GO_EXTERNAL_PTS are

those outside the domain and `GO_ALL_PTS` encompasses all grid points in the model. The chosen value is specified in the kernel-meta data like so:

```
integer :: iterates_over = GO_INTERNAL_PTS
```

A user can use a config file (see *Configuration*) to add additional iteration spaces to PSyclone.

Index Offset

The third element of kernel metadata, `INDEX_OFFSET`, specifies the index-offset that the kernel uses. This is the same quantity as supplied to the grid constructor (see the *Grid* Section for a description).

The GOcean 1.0 API supports two different offset schemes; `GO_OFFSET_NE`, `GO_OFFSET_SW`. The scheme used by a kernel is specified in the metadata as, e.g.:

```
integer :: index_offset = GO_OFFSET_NE
```

Currently all kernels used in an application must use the same offset scheme which must also be the same as passed to the grid constructor.

Procedure

The fourth and final type of metadata is `procedure` metadata. This specifies the name of the Kernel Fortran subroutine that this metadata describes.

For example:

```
procedure :: my_kernel_code
```

11.4.2 Subroutine

Rules

Kernel arguments follow a set of rules which have been specified for the GOcean 1.0 API. These rules are encoded in the `gen_code()` method of the `GOKern` class in the `gocean1p0.py` file. The rules, along with PSyclone's naming conventions, are:

1) Every kernel has the indices of the current grid point as the first two arguments, `i` and `j`. These are integers and have intent `in`.

2) For each field/scalar/grid property in the order specified by the meta_args metadata:

    1) For a field; the field array itself. A field array is a real array of kind `go_wp` and rank two. The first rank is longitude (*x*) and the second latitude (*y*).

    2) For a scalar; the variable itself. A real scalar is of kind `go_wp`.

    3) For a grid property; the array or variable (see the earlier table) containing the specified property.

Note: Grid properties are not passed from the Algorithm Layer. PSyclone generates the necessary lookups in the PSy Layer and includes the resulting references in the arguments passed to the kernel.

As an example, consider the `bc_solid_u` kernel that is used in the `gocean2d` program shown earlier. The metadata for this kernel is:

```fortran
type, extends(kernel_type) :: bc_solid_u
  type(go_arg), dimension(2) :: meta_args =  &
        (/ go_arg(GO_WRITE, GO_CU, GO_POINTWISE),      &
           go_arg(GO_READ,       GO_GRID_MASK_T)     &
         /)

  !> This is a boundary-conditions kernel and therefore
  !! acts on all points of the domain rather than just
  !! those that are internal
  integer :: ITERATES_OVER = GO_ALL_PTS

  integer :: index_offset = GO_OFFSET_NE

contains
  procedure, nopass :: code => bc_solid_u_code
end type bc_solid_u
```

The interface to the subroutine containing the implementation of this kernel is:

```fortran
subroutine bc_solid_u_code(ji, jj, ua, tmask)
  integer,                    intent(in)    :: ji, jj
  integer,  dimension(:,:), intent(in)    :: tmask
  real(wp), dimension(:,:), intent(inout) :: ua
```

As described above, the first two arguments to this subroutine specify the grid-point at which the computation is to be performed. The third argument is the field that this kernel updates and the fourth argument is the T-point mask. The latter is a property of the grid and is provided to the kernel call from the PSy Layer.

Comparing this interface definition with the use of the kernel in the Invoke call:

```fortran
call invoke ( ...,             &
              bc_solid_u(ua), &
              ... )
```

we see that in the Algorithm Layer the user need only provide the field(s) (and possibly scalars) that a kernel operates on. The index of the grid point and any grid properties are provided in the (generated) PSy Layer where the kernel subroutine proper is called.

11.5 Built-ins

The GOcean 1.0 API does not support any built-in operations.

11.6 Conventions

The GOcean 1.0 API kernel code conforms to the PSyclone Fortran naming conventions (see *Fortran Naming Conventions*). However, PSyclone's support for the GOcean 1.0 API does not rely on this convention.

The contents of the kernel metadata is usually declared private but this does not affect PSyclone.

Finally, the `procedure` metadata (located within the kernel metadata) usually has `nopass` specified but again this is ignored by PSyclone.

11.7 Configuration

The configuration file (see *Configuration*) used by PSyclone can contain GOcean 1.0 specific options. For example, after the default section the GOcean 1.0 specific section looks like this:

```
[gocean1.0]
iteration-spaces=offset_sw:ct:test_only:1:2:3:4
                 offset_sw:ct:internal_ns_halo:{start}-1:{stop}+1:{start}:{stop}
```

The supported keys are listed in the next section.

11.7.1 Iteration-spaces

This section lists additional iteration spaces that can be used in a kernel metadata declaration to allow PSyclone to create a loop with different loop boundaries. Each line of the `iteration-spaces` declaration contains 7 values, separated by ':'. The fields are:

Field	Description	Details
1	Index Offset	See *Index Offset*.
2	grid-point types	See *Grid point types*.
3	Iterates Over	See *Iterates Over*.
4	Start index of outer loop	Start index of North-South loop.
5	End index of outer loop	End index of North-South loop.
6	Start index of inner loop	Start index of East-West loop.
7	End index of inner loop	End index of East-West loop.

Two special variables can be used in an iteration space: `{start}` and `{stop}`. These values will be replaced by PSyclone with the correct loop boundaries for the inner points of a grid (i.e. the non-halo area). This means that the depth-1 halo region can be specified using `{start}-1` and `{stop}+1`.

For example, given the iteration-spaces declaration above, a kernel declared with `iterates_over=internal_ns_halo` for a field type `ct` and index offset `offset_sw` would create the following loop boundaries:

```
DO j=2-1,jstop+1
  DO i=2,istop
    CALL (i, j, ...)
  END DO
END DO
```

> **Warning:** With user defined iteration spaces it is possible that PSyclone will create code that does not compile: if you specify syntactically correct, but semantically incorrect boundary definitions, the PSyclone internal tests will accept the new iteration space, but the compiler will not. For example if one of the loop boundaries contains the name of a variable that is not defined, compilation will fail. It is the responsibility of the user to make sure that valid loop boundaries are specified in a new iteration space definition.

11.8 Transformations

In this section we describe the transformations that are specific to the GOcean 1.0 API. For an overview of transformations in general see *Transformations*.

class psyclone.transformations.**GOceanExtractRegionTrans**
> GOcean1.0 API application of ExtractRegionTrans transformation to extract code into a stand-alone program. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>>
>>> API = "gocean1.0"
>>> FILENAME = "shallow_alg.f90"
>>> ast, invokeInfo = parse(FILENAME, api=API)
>>> psy = PSyFactory(API, distributed_memory=False).create(invoke_info)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>>
>>> from psyclone.transformations import GOceanExtractRegionTrans
>>> etrans = GOceanExtractRegionTrans()
>>>
>>> # Apply GOceanExtractRegionTrans transformation to selected Nodes
>>> newsched, _ = etrans.apply(schedule.children[0])
>>> newsched.view()
```

> **name**
> > Returns the name of this transformation as a string.

---

class psyclone.transformations.**GOceanLoopFuseTrans**
> Performs error checking (that the loops are over the same grid-point type) before calling the *LoopFuseTrans.apply()* method of the *base class* in order to fuse two GOcean loops.

> **apply**(*node1*, *node2*)
> > Fuse the two GOcean loops represented by node1 and node2.

> > **Parameters**
> > > • **node1** (psyclone.gocean1p0.GOLoop) – A node representing a GOLoop.
> > >
> > > • **node2** (psyclone.gocean1p0.GOLoop) – A node representing a GOLoop.

> > **Raises** *TransformationError* – if the supplied node2 can not be fused, e.g. not all nodes are loops, don't have the same parent, are not next to each other or have different iteration spaces.

> **name**
> > Returns the name of this transformation as a string.

---

class psyclone.transformations.**GOceanOMPParallelLoopTrans**(*omp_schedule='static'*)
> GOcean specific OpenMP Do loop transformation. Adds GOcean specific validity checks (that supplied Loop is an inner or outer loop). Actual transformation is done by *base class*.

> > **Parameters** **omp_schedule** – The omp schedule to be created. Must be one of 'runtime', 'static', 'dynamic', 'guided' or 'auto'.

> **apply**(*node*)
> > Perform GOcean-specific loop validity checks then call *OMPParallelLoopTrans.apply()*.

> > **Parameters** **node** (*psyclone.psyGen.Loop*) – A Loop node from an AST.

> > **Raises** *TransformationError* – if the supplied node is not an inner or outer loop.

> **name**
> > Returns the name of this transformation as a string.

---

class `psyclone.transformations.`**`GOceanOMPLoopTrans`**(*omp_schedule='static'*)

GOcean-specific orphan OpenMP loop transformation. Adds GOcean specific validity checks (that the node is either an inner or outer Loop). Actual transformation is done by *base class*.

> **Parameters** `omp_schedule` – The omp schedule to be created. Must be one of 'runtime', 'static', 'dynamic', 'guided' or 'auto'.

apply(*node*)

Perform GOcean specific loop validity checks then call *OMPLoopTrans.apply()*.

> **Parameters** `node` (*psyclone.psyGen.Loop.*) – The loop to parallelise using OMP Do.

name

Returns the name of this transformation as a string.

---

class `psyclone.transformations.`**`GOConstLoopBoundsTrans`**

Switch on (or off) the use of constant loop bounds within a GOInvokeSchedule. In the absence of constant loop bounds, PSyclone will generate loops where the bounds are obtained by de-referencing a field object, e.g.:

```
DO j = my_field%grid%internal%ystart, my_field%grid%internal%ystop
```

Some compilers are able to produce more efficient code if they are provided with information on the relative trip-counts of the loops within an Invoke. With constant loop bounds switched on, PSyclone generates code like:

```
ny = my_field%grid%subdomain%internal%ystop
...
DO j = 1, ny-1
```

In practice, the application of the constant loop bounds looks something like, e.g.:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> import os
>>> TEST_API = "gocean1.0"
>>> _, info = parse(os.path.join("tests", "test_files", "gocean1p0",
>>>                              "single_invoke.f90"),
>>>                  api=TEST_API)
>>> psy = PSyFactory(TEST_API).create(info)
>>> invoke = psy.invokes.get('invoke_0_compute_cu')
>>> schedule = invoke.schedule
>>>
>>> from psyclone.transformations import GOConstLoopBoundsTrans
>>> clbtrans = GOConstLoopBoundsTrans()
>>>
>>> newsched, _ = clbtrans.apply(schedule)
>>> # or, to turn off const. looop bounds:
>>> # newsched, _ = clbtrans.apply(schedule, const_bounds=False)
>>>
>>> newsched.view()
```

apply(*node*, *const_bounds=True*)

Switches constant loop bounds on or off for all loops in a GOInvokeSchedule. Default is 'on'.

> **Parameters**

- **node** (psyclone.gocean1p0.GOInvokeSchedule) – The GOInvokeSchedule of which all loops will get the constant loop bounds switched on or off.

- **const_bounds** (*bool*) – If the constant loop should be used (True) or not (False). Default is True.

name
> Return the name of the Transformation as a string.

---

class psyclone.transformations.**GOLoopSwapTrans**
> Provides a loop-swap transformation, e.g.:

```
DO j=1, m
    DO i=1, n
```

becomes:

```
DO i=1, n
    DO j=1, m
```

This transform is used as follows:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast, invokeInfo = parse("shallow_alg.f90")
>>> psy = PSyFactory("gocean1.0").create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import GOLoopSwapTrans
>>> swap = GOLoopSwapTrans()
>>> new_schedule, memento = swap.apply(schedule.children[0])
>>> new_schedule.view()
```

apply (*outer*)
> The argument outer must be a loop which has exactly one inner loop. This transform then swaps the outer and inner loop.

> > **Parameters** **outer** (*psyclone.psyGen.Loop*) – The node representing the outer loop.

> > **Returns** A tuple consisting of the new schedule, and a Memento.

> > **Raises** *TransformationError* – if the supplied node does not allow a loop swap to be done.

name
> Returns the name of this transformation as a string.

NEMO API

In contrast to the other APIs supported by PSyclone, the NEMO API is designed to work with source code that does *not* follow the PSyKAl separation of concerns. Instead, the NEMO source code is treated as if it were a manually written PSy layer with all kernels in-lined. This approach relies upon the NEMO Coding Conventions [nem13] in order to reason about the code being processed. Rather than construct an InvokeSchedule for the PSy layer from scratch (as is done for other APIs), the InvokeSchedule is constructed by parsing the supplied Fortran code and generating a higher-level representation.

Note: the NEMO API is currently only a prototype. The known issues are listed in *Limitations*.

12.1 Algorithm

Since NEMO source is treated as a pre-existing PSy layer, this API does not have the concept of an Algorithm layer.

12.2 Constructing the PSyIR

Transformations in PSyclone are applied to an Internal Representation, the "PSyIR." In contrast to the other APIs where the PSyIR is constructed from scratch, for NEMO PSyclone must parse the existing Fortran and create a higher-level representation of it. This is done using rules based upon the NEMO Coding Conventions [nem13]. These rules are described in the following sections.

12.2.1 Loops

Explicit

PSyclone recognises the following loop types, based on the name of the loop variable:

Loop type	Loop variable
Vertical levels	jk
Latitude	ji
Longitude	jj
Tracer species	jn

PSyclone currently assumes that each of these loop types may be safely parallelised. In practice this will not always be the case (e.g. when performing a tri-diagonal solve) and this implementation will need to be refined.

Implicit

The use of Fortran array notation is encouraged in the NEMO Coding Conventions [nem13] (section 4.2) and is employed throughout the NEMO code base. The Coding Conventions mandate that the shape of every array in such expressions must be specified, e.g.:

```
onedarraya(:) = onedarrayb(:) + onedarrayc(:)
twodarray (:,:) = scalar * anothertwodarray(:,:)
```

PSyclone therefore also recognises the loops implied by this notation.

It has been found by CMCC and Intel that the Intel compiler is better able to OpenMP-parallelise explicit loops. Therefore, PSyclone provides a transformation to allow the user to convert implicit loops into explicit loop nests (see *Transformations*). However, as with all compiler-specific optimisations, whether or not this transformation is beneficial will depend on the precise details of the compiler being used.

Note, not all uses of Fortran array notation in NEMO imply a loop. For instance:

```
ascalar = afunc(twodarray(:,:))
```

is actually a function call which is passed a reference to `twodarray`. However, if the quantity being assigned to is actually an array, e.g.:

```
twodarray2(:,:) = afunc(twodarray(:,:))
```

then this does represent a loop. However, currently PSyclone does not recognise any occurrences of array notation that are themselves within an array access or function call. It is therefore not yet possible to transform such implicit loops into explicit loops. It is hoped that this limitation will be removed in future releases of PSyclone by adding the ability to discover the interface to functions such as `afunc` and thus determining whether they return scalar or array quantities.

12.3 Example

A typical fragment of NEMO source code (taken from the traldf_iso routine) is shown below:

```
DO jn = 1, kjpt
 zdit (1,:,:) = 0._wp    ;    zdit (jpi,:,:) = 0._wp
 zdjt (1,:,:) = 0._wp    ;    zdjt (jpi,:,:) = 0._wp

 DO jk = 1, jpkm1
    DO jj = 1, jpjm1
       DO ji = 1, fs_jpim1
          zdit(ji,jj,jk) = ( ptb(ji+1,jj  ,jk,jn) - ptb(ji,jj,jk,jn) ) * umask(ji,jj,
↪jk)
          zdjt(ji,jj,jk) = ( ptb(ji  ,jj+1,jk,jn) - ptb(ji,jj,jk,jn) ) * vmask(ji,jj,
↪jk)
       END DO
    END DO
 END DO
```

PSyclone uses fparser2 to parse such source code and then generates the PSy Internal Representation of it:

```
Loop[type='tracers',field_space='None',it_space='None']
    Loop[type='None',field_space='None',it_space='None']
    Loop[type='None',field_space='None',it_space='None']
```
(continues on next page)

```
    Loop[type='None',field_space='None',it_space='None']
    Loop[type='None',field_space='None',it_space='None']
    Loop[type='levels',field_space='None',it_space='None']
        Loop[type='lat',field_space='None',it_space='None']
            Loop[type='lon',field_space='None',it_space='None']
                KernCall[]
```

12.4 Transformations

This section describes the transformations that are specific to the NEMO API. For an overview of transformations in general see *Transformations*.

class psyclone.transformations.**NemoExplicitLoopTrans**

Transforms the outermost array slice in an implicit loop in a NEMOInvokeSchedule into an explicit loop. For example, if "implicit_loop.f90" contained:

```
my_array(:, :, :) = 1.0
```

then doing:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "nemo"
>>> filename = "implicit_loop.f90"
>>> ast, invokeInfo = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.transformations import NemoExplicitLoopTrans
>>> rtrans = NemoExplicitLoopTrans()
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> loop = schedule.children[0]
>>> newloop, _ = rtrans.apply(loop)
```

will create a new NemoLoop object for an explicit loop over levels (the outermost slice) that then contains an implicit loop:

```
DO jk = 1, jpk
  my_array(:, :, jk) = 1.0
END DO
```

Subsequently applying *rtrans* to *newloop* will create:

```
DO jk = 1, jpk
  DO jj = 1, jpj
    my_array(:, jj, jk) = 1.0
  END DO
END DO
```

apply(*loop*)

Transform the outermost array slice in the supplied implicit loop into an explicit loop.

> **Parameters** **loop** (psyclone.nemo.NemoImplicitLoop) – the NemoImplicitLoop to transform.

> **Returns** a new PSyIR loop object and a memento of the transformation.
>
> **Return type** (`psyclone.nemo.NemoLoop`, `psyclone.undoredo.Memento`)
>
> **Raises**
>
> - **`NotImplementedError`** – if the array slice has explicit bounds.
>
> - **`TransformationError`** – if an array slice is not in dimensions 1-3 of the array.

name

> **Returns** the name of this transformation class.
>
> **Return type** str

validate(*loop*)

> Check that the supplied loop is a valid target for this transformation.
>
> **Parameters loop** (`psyclone.nemo.NemoImplicitLoop`) – the loop node to validate.
>
> **Raises** *TransformationError* – if the supplied loop is not a NemoImplicitLoop.

The type of the loop being transformed is inferred from the position of the colon(s) in the array subscripts. Since NEMO uses *(ji,jj,jk)* index ordering (longitude, latitude, levels) this means that loop types are determined according to the following table:

Array index	Loop type	Loop limits
1	Longitude	1, jpi
2	Latitude	1, jpj
3	Vertical levels	1, jpk

If this transformation encounters an implicit loop in an array index other than 1-3 then currently PSyclone will raise an error.

12.5 Limitations

The NEMO API is currently only a prototype implementation. Here we list the current, known limitations/issues:

1. When transforming implicit loops into explicit loops, the declaration of the loop variables can be repeated (there is an x-failing test for this);

2. Scalar variables inside loops are not made private when parallelising using OpenMP;

3. All recognised loops (levels, latitude etc.) are assumed to be parallelisable. This will not always be the case (e.g. tridiagonal solve has a loop-carried dependence in the vertical);

4. Labelled do-loops are not handled (i.e. they will be put inside a 'CodeBlock' in the PSyIR);

5. Loops are currently only permitted to contain one kernel. This restriction will have to be lifted in order to permit loop fusion;

6. Array slices with specified bounds (e.g. umask(1:10)) are not yet supported and will raise a TransformationError when attempting to transform them into explicit loops;

7. When generating new variable names, no attempt is made to avoid clashing with variables already present in the NEMO source. This needs to be resolved by querying the SymbolTable (#381);

8. The psyGen.Node base class now has an _ast property to hold a pointer into the associated fparser2 AST. However, the psyGen.Kern class already has an _fp2_ast property that points to the whole fparser2 AST of the kernel code. This will be rationalised in #241.

STUB GENERATION

13.1 Quick Start

1) Use an existing Kernel file or create a Kernel file containing a Kernel module with the required metadata and an empty Kernel subroutine with no arguments.

2) Run the following command

```
> genkernelstub <PATH>/my_file.f90
```

13.2 Introduction

PSyclone provides a kernel stub generator for the Dynamo0.3 API. The kernel stub generator takes a kernel file as input and outputs the kernel subroutine arguments and declarations. The word "stub" is used to indicate that it is only the subroutine arguments and their declarations that are generated; the subroutine has no content.

The primary reason the stub generator is useful is that it generates the correct Kernel subroutine arguments and declarations for the Dynamo0.3 API as specified by the Kernel metadata. As the number of arguments to Kernel subroutines can become large and the arguments have to follow a particular order, it can become burdensome, and potentially error prone, for the user to have to work out the appropriate argument list if written by hand.

The stub generator can be used when creating a new Kernel. A Kernel can first be written to specify the required metadata and then the generator can be used to create the appropriate (empty) Kernel subroutine. The user can then fill in the content of the subroutine.

The stub generator can also be used to check whether the arguments for an existing Kernel are correct i.e. whether the Kernel subroutine and Kernel metadata are consistent. One example would be where a Kernel is updated resulting in a change to the metadata and subroutine arguments.

The Dynamo0.3 API requires Kernels to conform to a set of rules which determine the required arguments and types for a particular Kernel. These rules are required as the generated PSy layer needs to know exactly how to call a Kernel. These rules are outlined in Section *Rules*.

Therefore PSyclone has been coded with the Dynamo0.3 API rules which are then applied when reading the Kernel metadata to produce the require Kernel call and its arguments in the generated PSy layer. These same rules are used by the Kernel stub generator to produce Kernel subroutine stubs, thereby guaranteeing that Kernel calls from the PSy layer and the associated Kernel subroutines are consistent.

13.3 Use

Before using the stub generator, PSyclone must be installed. If you have not already done so, please follow the instructions for setting up PSyclone in Section *Getting Going*.

PSyclone will be installed in a particular location on your machine, which will be referred to as the `<PSYCLONEINSTALL>` directory. The stub generator script `genkernelstub` comes with the PSyclone installation. A quick check `> which genkernelstub` should return the location of the `<PSYCLONEINSTALL>/bin` directory.

The easiest way to use the stub generator is to run the supplied `genkernelstub` script.

```
> genkernelstub
usage: genkernelstub [-h] [-o OUTFILE] [-api API] [-l] filename
genkernelstub: error: too few arguments
```

You can get information about the `genkernelstub` arguments using `-h` or `--help`:

```
>  genkernelstub -h
usage: genkernelstub [-h] [-o OUTFILE] [-api API] [-l] filename

Create Kernel stub code from Kernel metadata

positional arguments:
  filename              Kernel metadata

optional arguments:
  -h, --help            show this help message and exit
  -o OUTFILE, --outfile OUTFILE
                        filename of output
  -api API              choose a particular api from ['dynamo0.3'], default
                        dynamo0.3
  -l, --limit           limit the fortran line length to 132 characters
```

As is indicated when using the `-h` option, the `-api` option only accepts `dynamo0.3` at the moment and is redundant as this option is also the default. However the number of supported APIs is expected to expand in the future.

The `-o`, or `--outfile` option allows the user to specify that the output should be written to a particular file. If `-o` is not specified then the python `print` statement is used. Typically the print statement results in the output being printed to the terminal.

The `-l`, or `--limit` option utilises the PSyclone support for wrapping of lines within the 132 character limit (please see the *Line Length* chapter for more details).

13.4 Kernels

Any Dynamo0.3 kernel can be used as input to the stub generator. Example Kernels can be found in the `examples/dynamo` repository or, for more simple cases, in the `tests/test_files/dynamo0p3` directory. These directories are located in the `<PSYCLONEHOME>/src/psyclone` directory where `<PSYCLONEHOME>` refers to the location where you download or clone PSyclone (*Getting Going*).

In the `tests/test_files/dynamo0p3` directory the majority of examples start with `testkern`. The exceptions are: `simple.f90`, `ru_kernel_mod.f90` and `matrix_vector_mod.F90`. The following test kernels can be used to generate kernel stub code (running stub generation from the `<PSYCLONEHOME>/src/psyclone` directory):

```
tests/test_files/dynamo0p3/testkern_chi_2.F90
tests/test_files/dynamo0p3/testkern_chi.F90
tests/test_files/dynamo0p3/testkern_operator_mod.f90
tests/test_files/dynamo0p3/testkern_operator_nofield_mod.f90
tests/test_files/dynamo0p3/testkern_orientation.F90
tests/test_files/dynamo0p3/testkern_operator_orient_mod.f90
tests/test_files/dynamo0p3/ru_kernel_mod.f90
tests/test_files/dynamo0p3/simple.f90
```

13.5 Example

A simple single field example of a kernel that can be used as input for the stub generator is found in `tests/test_files/dynamo0p3/simple.f90` and is shown below:

```
module simple_mod
type, extends(kernel_type) :: simple_type
   type(arg_type), dimension(1) :: meta_args =  &
       (/ arg_type(gh_field,gh_write,w1) /)
   integer :: iterates_over = cells
 contains
   procedure, nopass :: code => simple_code
end type simple_type
contains
subroutine simple_code()
end subroutine
end module simple_mod
```

Note: The module name `simple_mod` and the type name `simple_type` share the same root `simple` and have the extensions `_mod` and `_type` respectively. This is a convention in Dynamo0.3 and is required by the kernel stub generator as it needs to determine the name of the type containing the metadata and infers this by reading the module name. If this rule is not followed the kernel stub generator will return with an error message (see Section *Errors*).

Note: Whilst strictly the kernel stub generator only requires the Kernel metadata to generate the appropriate stub code, the parser that the generator relies on currently requires a dummy kernel subroutine to exist.

If we run the kernel stub generator on the `simple.f90` example:

```
> genkernelstub tests/test_files/dynamo0p3/simple.f90
```

we get the following kernel stub output:

```
MODULE simple_mod
  IMPLICIT NONE
  CONTAINS
  SUBROUTINE simple_code(nlayers, field_1_w1, ndf_w1, undf_w1, map_w1)
    USE constants_mod, ONLY: r_def
    IMPLICIT NONE
    INTEGER, intent(in) :: nlayers
    INTEGER, intent(in) :: ndf_w1
    INTEGER, intent(in) :: undf_w1
    REAL(KIND=r_def), intent(out), dimension(undf_w1) :: field_1_w1
```

(continues on next page)

```
    INTEGER, intent(in), dimension(ndf_w1) :: map_w1
  END SUBROUTINE simple_code
END MODULE simple_mod
```

The subroutine content can then be copied into the required module, used as the basis for a new module, or checked with an existing subroutine for correctness.

Note: The output does not currently conform to Met Office coding standards so must be modified accordingly.

Note: The code will not compile without a) providing the `constants_mod`, `argument_mod` and `kernel_mod` modules in the compiler include path and b) adding in code that writes to any arguments declared as intent `out` or `inout`. For a quick check, the `USE` declaration and `KIND` declarations can be removed and the `field_1_w1` array can be initialised with some value in the subroutine. At this point the Kernel should compile successfully.

Note: Whilst there is only one field declared in the metadata there are 5 arguments to the Kernel. The first argument `nlayers` specifies the number of layers in a column for a field. The second argument is the array associated with the field. The field array is dimensioned as the number of unique degrees of freedom (undf) which is also passed into the kernel (the fourth argument). The naming convention is to call each field a `field`, followed by its position in the (algorithm) argument list (which is reflected in the metadata ordering). The third argument is the number of degrees of freedom for the particular column and is used to dimension the final argument which is the degrees of freedom map (dofmap) which indicates the location of the required values in the field array. The naming convention for the `dofmap`, `undf` and `ndf` is to append the name with the space that it is associated with.

We now take a look at a more complicated example. The metadata in this example is the same as an actual Dynamo kernel, however the subroutine content and various comments have been removed. The metadata specifies that there are four fields passed by the algorithm layer, the fourth of which is a vector field of size three. All three of the spaces require a basis function and the w0 and w2 function spaces additionally require a differential basis function. The content of the Kernel is given below.

```
module ru_kernel_mod
type, public, extends(kernel_type) :: ru_kernel_type
  private
  type(arg_type) :: meta_args(6) = (/                        &
      arg_type(GH_FIELD,   GH_INC,  W2),                     &
      arg_type(GH_FIELD,   GH_READ, W3),                     &
      arg_type(GH_INTEGER, GH_READ),                         &
      arg_type(GH_REAL,    GH_READ),                         &
      arg_type(GH_FIELD,   GH_READ, W0),                     &
      arg_type(GH_FIELD*3, GH_READ, W0)                      &
      /)
  type(func_type) :: meta_funcs(3) = (/                      &
      func_type(W2, GH_BASIS, GH_DIFF_BASIS),                &
      func_type(W3, GH_BASIS),                               &
      func_type(W0, GH_BASIS, GH_DIFF_BASIS)                 &
      /)
  integer :: iterates_over = CELLS
  integer :: gh_shape = gh_quadrature_XYoZ
contains
  procedure, nopass ::ru_code
end type
```

```
contains
subroutine ru_code()
end subroutine ru_code
end module ru_kernel_mod
```

If we run the kernel stub generator on this example:

```
> genkernelstub tests/test_files/dynamo0p3/ru_kernel_mod.f90
```

we obtain the following output:

```
MODULE ru_mod
  IMPLICIT NONE
  CONTAINS
  SUBROUTINE ru_code(nlayers, field_1_w2, field_2_w3, iscalar_3, rscalar_4, field_5_
→w0, field_6_w0_v1, field_6_w0_v2, field_6_w0_v3, ndf_w2, undf_w2, map_w2, basis_w2,␣
→diff_basis_w2, ndf_w3, undf_w3, map_w3, basis_w3, ndf_w0, undf_w0, map_w0, basis_w0,
→ diff_basis_w0, np_xy, np_z, weights_xy, weights_z)
    USE constants_mod, ONLY: r_def
    IMPLICIT NONE
    INTEGER, intent(in) :: nlayers
    INTEGER, intent(in) :: ndf_w2
    INTEGER, intent(in) :: undf_w2
    INTEGER, intent(in) :: ndf_w3
    INTEGER, intent(in) :: undf_w3
    INTEGER, intent(in) :: ndf_w0
    INTEGER, intent(in) :: undf_w0
    REAL(KIND=r_def), intent(inout), dimension(undf_w2) :: field_1_w2
    REAL(KIND=r_def), intent(in), dimension(undf_w3) :: field_2_w3
    INTEGER, intent(in) :: iscalar_3
    REAL(KIND=r_def), intent(in) :: rscalar_4
    REAL(KIND=r_def), intent(in), dimension(undf_w0) :: field_5_w0
    REAL(KIND=r_def), intent(in), dimension(undf_w0) :: field_6_w0_v1
    REAL(KIND=r_def), intent(in), dimension(undf_w0) :: field_6_w0_v2
    REAL(KIND=r_def), intent(in), dimension(undf_w0) :: field_6_w0_v3
    INTEGER, intent(in), dimension(ndf_w2) :: map_w2
    REAL(KIND=r_def), intent(in), dimension(3,ndf_w2,np_xy,np_z) :: basis_w2
    REAL(KIND=r_def), intent(in), dimension(1,ndf_w2,np_xy,np_z) :: diff_basis_w2
    INTEGER, intent(in), dimension(ndf_w3) :: map_w3
    REAL(KIND=r_def), intent(in), dimension(1,ndf_w3,np_xy,np_z) :: basis_w3
    INTEGER, intent(in), dimension(ndf_w0) :: map_w0
    REAL(KIND=r_def), intent(in), dimension(1,ndf_w0,np_xy,np_z) :: basis_w0
    REAL(KIND=r_def), intent(in), dimension(3,ndf_w0,np_xy,np_z) :: diff_basis_w0
    INTEGER, intent(in) :: np_xy, np_z
    REAL(KIND=r_def), intent(in), dimension(np_xy) :: weights_xy
    REAL(KIND=r_def), intent(in), dimension(np_z) :: weights_z
  END SUBROUTINE ru_code
END MODULE ru_mod
```

The above example demonstrates that the argument list can get quite complex. Rather than going through an explanation of each argument you are referred to Section *Rules* for more details on the rules for argument types and argument ordering. Regarding naming conventions for arguments you can see that the arrays associated with the fields are labelled as 1-6 depending on their position in the metadata. For a vector field, each vector results in a different array. These are distinguished by appending _vx where x is the number of the vector.

The introduction of stencil operations on field arguments futher complicates the argument list of a kernel. An example

---

13.5. Example 117

of the use of the stub generator for a kernel that performs stencil operations is provided in `examples/dynamo/eg5`.

```
> genkernelstub ../../examples/dynamo/eg5/conservative_flux_kernel_mod.F90
```

13.6 Errors

The stub generator has been written to provide useful errors if mistakes are found. If you run the generator and it does not produce a useful error - and in particular if it produces a stack trace - please contact the PSyclone developers.

The following tests do not produce stub kernel code either because they are invalid or because they contain functionality that is not supported in the stub generator.

```
tests/test_files/dynamo0p3/matrix_vector_mod.F90
tests/test_files/dynamo0p3/testkern_any_space_1_mod.f90
tests/test_files/dynamo0p3/testkern_any_space_4_mod.f90
tests/test_files/dynamo0p3/testkern_invalid_fortran.F90
tests/test_files/dynamo0p3/testkern_short_name.F90
tests/test_files/dynamo0p3/testkern_no_datatype.F90
tests/test_files/dynamo0p3/testkern.F90
tests/test_files/dynamo0p3/testkern_qr.F90
```

`testkern_invalid_fortran.F90`, `testkern_no_datatype.F90`, `testkern_short_name.F90`, `testkern.F90` and `matrix_vector_mod.F90` are designed to be invalid for PSyclone testing purposes and should produce appropriate errors. For example:

```
> genkernelstub tests/test_files/dynamo0p3/testkern_invalid_fortran.F90
Error: 'Parse Error: Code appears to be invalid Fortran'
```

`any_space` is not currently supported in the stub generator so `testkern_any_space_1_mod.f90` and `testkern_any_space_4_mod.f90` should fail with appropriate warnings because of that. For example:

```
> genkernelstub tests/test_files/dynamo0p3/testkern_any_space_1_mod.f90
Error: "Generation Error: Unsupported space for basis function, expecting
one of ['w3', 'wtheta', 'w2v', 'w0', 'w1', 'w2', 'w2h', 'any_w2'] but
found 'any_space_1'"
```

As noted above, if the Dynamo0.3 naming convention for module and type names is not followed, the stub generator will return with an error message. For example:

```
> genkernelstub tests/test_files/dynamo0p3/testkern.F90
Error: "Parse Error: Error, module name 'testkern' does not have
'_mod' as an extension. This convention is assumed."
```

FOURTEEN

LINE LENGTH

By default PSyclone will generate fortran code with no consideration of fortran line length limits. As the line length limit for free-form fortran is 132 characters, the code that is output may be non-conformant.

Line length is not an issue for many compilers as they allow compiler flags to be set which allow lines longer than the fortran standard. However this is not the case for all compilers.

PSyclone therefore supports the wrapping of lines within the 132 character limit. The next two sections discuss how this is done when scripting and when working interactively respectively.

14.1 Script

The *psyclone* script provides the -l option to wrap lines. Please see the *Fortran line length* section for more details.

14.2 Interactive

When using PSyclone interactively the line lengths of the input algorithm and Kernel files can be checked by setting the *psyclone.parse.algorithm.parse()* function's "line_length" argument to "True".

```
>>> from psyclone.parse.algorithm import parse
>>> ast, info = parse("argspec.F90", line_length=True)
```

Similarly the "line_length" argument can be set to "True" if calling the generator.generate() function. This function simply passes this argument on to the *psyclone.parse.algorithm.parse()* function.

```
>>> from psyclone.generator import generate
>>> alg, psy = generate("argspec.F90", line_length=True)
```

Line wrapping is performed as a post processing step, i.e. after the code has been generated. This is done by an instance of the line_length.FortLineLength class. For example:

```
>>> from psyclone.generator import generate
>>> from psyclone.line_length import FortLineLength
>>> psy, alg = generate("algspec.f90", line_length=True)
>>> line_length = FortLineLength()
>>> psy_str = line_length.process(str(psy))
>>> print psy_str
>>> alg_str = line_length.process(str(alg))
>>> print alg_str
```

14.3 Limitations

The `line_length.FortLineLength` class is only partially aware of fortran syntax. This awareness is required so that appropriate continuation characters can be used (for example `&` at the end of a line and `!$omp&` at the start of a line for OpenMP directives, `&` at the end of a line for statements and `&` at the end of a line and `&` at the beginning of a line for strings).

Whilst statements only require an `&` at the end of the line when line wrapping with free-form fortran they may optionally also have an `&` at the beginning of the subsequent line. In contrast, when splitting a string over multiple lines an `&` is required at both locations. Therefore an instance of the `line_length.FortLineLength` class will always add `&` at the beginning of a continuation line for a statement, in case the line is split within a string.

One known situation that could cause an instance of the `line_length.FortLineLength` class to fail is when an inline comment is used at the end of a line to make it longer than the 132 character limit. Whilst PSyclone does not generate such code for the PSy-layer, this might occur in Algorithm-layer code, even if the Algorithm-layer code conforms to the 132 line length limit. The reason for this is that PSyclone's internal parser concatenates lines together, thus a long line correctly split with continuation characters in the Algorithm-layer becomes a line that needs to be split by an instance of the `line_length.FortLineLength` class.

FORTRAN NAMING CONVENTIONS

There is a convention in the kernel code for the Dynamo0.3 and GOcean1.0 APIs that if the name of the operation being performed is `<name>` then a kernel file is `<name>_mod.[fF90]`, the name of the module inside the kernel file is `<name>_mod`, the name of the kernel metadata in the module is `<name>_type` and the name of the kernel subroutine in the module is `<name>_code`.

PSyclone itself does not rely on this convention apart from in the stub generator (see the *Stub Generation* Section) where the name of the metadata to be parsed is determined from the module name.

API

16.1 The generator module

This module provides the PSyclone 'main' routine which is intended to be driven from the bin/psyclone executable script. 'main' takes an algorithm file as input and produces modified algorithm code and generated PSy code. A function, 'generate', is also provided which has the same functionality as 'main' but can be called from within another Python program.

`psyclone.generator.`**`generate`**(*filename*, *api=''*, *kernel_path=''*, *script_name=None*, *line_length=False*, *distributed_memory=None*, *kern_out_path=''*, *kern_naming='multiple'*)

Takes a PSyclone algorithm specification as input and outputs the associated generated algorithm and psy codes suitable for compiling with the specified kernel(s) and support infrastructure. Uses the `parse.algorithm.parse()` function to parse the algorithm specification, the `psyGen.PSy` class to generate the PSy code and the `algGen.Alg` class to generate the modified algorithm code.

> **Parameters**
>
> - **filename** (`str`) – The file containing the algorithm specification.
>
> - **kernel_path** (`str`) – The directory from which to recursively search for the files containing the kernel source (if different from the location of the algorithm specification).
>
> - **script_name** (`str`) – A script file that can apply optimisations to the PSy layer (can be a path to a file or a filename that relies on the PYTHONPATH to find the module).
>
> - **line_length** (`bool`) – A logical flag specifying whether we care about line lengths being longer than 132 characters. If so, the input (algorithm and kernel) code is checked to make sure that it conforms. The default is False.
>
> - **distributed_memory** (`bool`) – A logical flag specifying whether to generate distributed memory code. The default is set in the config.py file.
>
> - **kern_out_path** (`str`) – Directory to which to write transformed kernel code.
>
> - **kern_naming** (`bool`) – the scheme to use when re-naming transformed kernels.
>
> **Returns** 2-tuple containing fparser1 ASTs for the algorithm code and the psy code.
>
> **Return type** (fparser.one.block_statements.BeginSource, fparser.one.block_statements.Module)
>
> **Raises**
>
> - **IOError** – if the filename or search path do not exist
>
> - **GenerationError** – if an invalid API is specified.
>
> - **GenerationError** – if an invalid kernel-renaming scheme is specified.

For example:

```
>>> from psyclone.generator import generate
>>> alg, psy = generate("algspec.f90")
>>> alg, psy = generate("algspec.f90", kernel_path="src/kernels")
>>> alg, psy = generate("algspec.f90", script_name="optimise.py")
>>> alg, psy = generate("algspec.f90", line_length=True)
>>> alg, psy = generate("algspec.f90", distributed_memory=False)
```

16.2 The parse module

Module that uses the Fortran parser fparser2 to parse PSyclone-conformant Algorithm code.

psyclone.parse.algorithm.**parse**(*alg_filename*, *api=''*, *invoke_name='invoke'*, *kernel_path='''*, *line_length=False*)

Takes a PSyclone conformant algorithm file as input and outputs a parse tree of the code contained therein and an object containing information about the 'invoke' calls in the algorithm file and any associated kernels within the invoke calls.

> **Parameters**
>
> - **alg_filename** (*str*) – The file containing the algorithm specification.
> - **api** (*str*) – The PSyclone API to use when parsing the code.
> - **invoke_name** (*str*) – The expected name of the invocation calls in the algorithm code.
> - **kernel_path** (*str*) – The path to search for kernel source files (if different from the location of the algorithm source).
> - **line_length** (*bool*) – A logical flag specifying whether we care about line lengths being longer than 132 characters. If so, the input (algorithm and kernel) code is checked to make sure that it conforms and an error raised if not. The default is False.
>
> **Returns** 2-tuple consisting of the fparser2 parse tree of the Algorithm file and an object holding details of the invokes found.
>
> **Return type** (fparser.two.Fortran2003.Program, psyclone.parse.FileInfo)

For example:

```
>>> from psyclone.parse.algorithm import parse
>>> ast, info = parse("alg.f90")
```

16.3 The transformations module

This module provides the various transformations that can be applied to PSyIR nodes. There are both general and API-specific transformation classes in this module where the latter typically apply API-specific checks before calling the base class for the actual transformation.

class psyclone.transformations.**ACCDataTrans**

Add an OpenACC data region around a list of nodes in the PSyIR. COPYIN, COPYOUT and COPY clauses are added as required.

For example:

```
>>> from psyclone.parse import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "NEMO"
>>> filename = "tra_adv.F90"
>>> ast, invokeInfo = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.transformations import ACCDataTrans
>>> dtrans = ACCDataTrans()
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> kernels = schedule.children[0].children[0].children[0:-1]
>>> # Enclose the kernels
>>> new_sched, _ = dtrans.apply(kernels)
```

apply(*node_list*)

> Put the supplied list of nodes within an OpenACC data region.

> > **Parameters node_list** (list of *psyclone.psyGen.Node*) – The list of PSyIR nodes to enclose in the data region.

> > **Returns** (transformed schedule, memento of transformation)

> > **Return type** 2-tuple of (*psyclone.psyGen.Schedule*, psyclone.undoredo.Memento).

name

> > **Returns** the name of this transformation.

> > **Return type** str

psyGen = <module 'psyclone.psyGen' from '/home/rupert/proj/PSyclone/src/psyclone/psyGe

class psyclone.transformations.**ACCEnterDataTrans**

> Adds an OpenACC "enter data" directive to a Schedule. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> dtrans = t.get_trans_name('ACCEnterDataTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> # Add an enter-data directive
>>> newschedule, _ = dtrans.apply(schedule)
>>> newschedule.view()
```

apply(*sched*)

> Adds an OpenACC "enter data" directive to the invoke associated with the supplied Schedule. Any fields accessed by OpenACC kernels within this schedule will be added to this data region in order to ensure they remain on the target device.

> > **Parameters** **sched** (sub-class of *psyclone.psyGen.Schedule*.) – Schedule to which to add an "enter data" directive.
>
> > **Returns** Tuple of the modified schedule and a record of the transformation.
>
> > **Return type** (*psyclone.psyGen.Schedule*, psyclone.undoredo.Memento)

> **name**
>
> > **Returns** the name of this transformation.
>
> > **Return type** str

class psyclone.transformations.**ACCKernelsTrans**

> Enclose a sub-set of nodes from a Schedule within an OpenACC kernels region (i.e. within "!$acc kernels" . . . "!$acc end kernels" directives). Currently only supported for the NEMO API.
>
> For example:

```
>>> from psyclone.parse import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "NEMO"
>>> filename = "tra_adv.F90"
>>> ast, invokeInfo = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.transformations import ACCKernelsTrans
>>> ktrans = ACCKernelsTrans()
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> kernels = schedule.children[0].children[0].children[0:-1]
>>> # Transform the kernel
>>> new_sched, _ = ktrans.apply(kernels)
```

> **apply**(*node_list*, *default_present=False*)
>
> > Enclose the supplied list of PSyIR nodes within an OpenACC Kernels region.
>
> > **Parameters**
> >
> > - **node_list** (list of *psyclone.psyGen.Node*) – The list of nodes in the PSyIR to enclose.
> >
> > - **default_present** (*bool*) – whether or not the kernels region should have the 'default present' attribute (indicating that data is already on the accelerator). When using managed memory this option should be False.
>
> > **Returns** (transformed schedule, memento of transformation)
>
> > **Return type** 2-tuple of (*psyclone.psyGen.Schedule*, psyclone.undoredo. Memento).

> **name**
>
> > **Returns** the name of this transformation class.
>
> > **Return type** str

nemo = <module 'psyclone.nemo' from '/home/rupert/proj/PSyclone/src/psyclone/nemo.py'>

psyGen = <module 'psyclone.psyGen' from '/home/rupert/proj/PSyclone/src/psyclone/psyGen

class psyclone.transformations.**ACCLoopTrans**

Adds an OpenACC loop directive to a loop. This directive must be within the scope of some OpenACC Parallel region (at code-generation time).

For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory, GenerationError
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ltrans = t.get_trans_name('ACCLoopTrans')
>>> rtrans = t.get_trans_name('ACCParallelTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> new_schedule = schedule
>>>
# Apply the OpenACC Loop transformation to *every* loop
# in the schedule
>>> for child in schedule.children:
>>>     newschedule, memento = ltrans.apply(child, reprod=True)
>>>     schedule = newschedule
>>>
# Enclose all of these loops within a single OpenACC
# PARALLEL region
>>> rtrans.omp_schedule("dynamic,1")
>>> newschedule, memento = rtrans.apply(schedule.children)
>>>
```

apply (*node*, *collapse=None*, *independent=True*, *sequential=False*)

Apply the ACCLoop transformation to the specified node in a GOInvokeSchedule. This node must be a Loop since this transformation corresponds to inserting a directive immediately before a loop, e.g.:

```
!$ACC LOOP
do ...
    ...
end do
```

At code-generation time (when psyclone.psyGen.ACCLoopDirective.gen_code() is called), this node must be within (i.e. a child of) a PARALLEL region.

Parameters

- **node** (*psyclone.psyGen.Loop.*) – the supplied node to which we will apply the Loop transformation.

- **collapse** (*int*) – number of loops to collapse into single iteration space or None.

- **independent** (*bool*) – whether to add the "independent" clause to the directive (not strictly necessary within PARALLEL regions).

Returns (psyclone.psyGen.GOInvokeSchedule,   psyclone.undoredo. Memento)

> **name**
>> Returns the name of this transformation as a string.

class psyclone.transformations.**ACCParallelTrans**

> Create an OpenACC parallel region by inserting directives. This parallel region *must* come after an enter-data directive (see *ACCEnterDataTrans*) or within a data region (see *ACCDataTrans*). For example:

```python
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ptrans = t.get_trans_name('ACCParallelTrans')
>>> dtrans = t.get_trans_name('ACCDataTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> # Enclose everything within a single OpenACC PARALLEL region
>>> newschedule, _ = ptrans.apply(schedule.children)
>>> # Add an enter-data directive
>>> newschedule, _ = dtrans.apply(newschedule)
>>> newschedule.view()
```

> **gocean1p0 = <module 'psyclone.gocean1p0' from '/home/rupert/proj/PSyclone/src/psyclone**

> **name**

>> **Returns** The name of this transformation as a string.

>> **Return type** str

> **nemo = <module 'psyclone.nemo' from '/home/rupert/proj/PSyclone/src/psyclone/nemo.py'>**

> **psyGen = <module 'psyclone.psyGen' from '/home/rupert/proj/PSyclone/src/psyclone/psyGen**

class psyclone.transformations.**ACCRoutineTrans**

> Transform a kernel subroutine by adding a "!$acc routine" directive (causing it to be compiled for the OpenACC accelerator device). For example:

```python
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.transformations import ACCRoutineTrans
>>> rtrans = ACCRoutineTrans()
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> kern = schedule.children[0].children[0].children[0]
>>> # Transform the kernel
>>> newkern, _ = rtrans.apply(kern)
```

apply(*kern*)

> Modifies the AST of the supplied kernel so that it contains an '!$acc routine' OpenACC directive.

> > **Parameters** **kern** (`psyclone.psyGen.Call`) – The kernel object to transform.

> > **Returns** (transformed kernel, memento of transformation)

> > **Return type** 2-tuple of (`psyclone.psyGen.Kern`, `psyclone.undoredo.Memento`).

> > **Raises** *`TransformationError`* – if we fail to find the subroutine corresponding to the kernel object.

name

> > **Returns** the name of this transformation class.

> > **Return type** str

validate(*kern*)

> Perform checks that the supplied kernel can be transformed.

> > **Parameters** **kern** (`psyclone.psyGen.Call`) – the kernel which is the target of the transformation.

> > **Raises** *`TransformationError`* – if the target kernel is a built-in.

class psyclone.transformations.**ColourTrans**

> Apply a colouring transformation to a loop (in order to permit a subsequent parallelisation over colours). For example:

```
>>> invoke = ...
>>> schedule = invoke.schedule
>>>
>>> ctrans = ColourTrans()
>>>
>>> # Colour all of the loops
>>> for child in schedule.children:
>>>     cschedule, _ = ctrans.apply(child)
>>>
>>> csched.view()
```

apply(*node*)

> Converts the Loop represented by `node` into a nested loop where the outer loop is over colours and the inner loop is over cells of that colour. :param node: The loop to transform. :type node: *`psyclone.psyGen.Loop`* :returns: Tuple of modified schedule and record of transformation :rtype: (`psyclone.psyGen.Schedule`, :py:class:`psyclone.undoredo.Memento`)

name

> Returns the name of this transformation as a string.

class psyclone.transformations.**Dynamo0p3AsyncHaloExchangeTrans**

> Splits a synchronous halo exchange into a halo exchange start and halo exchange end. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "dynamo0.3"
>>> ast, invokeInfo = parse("file.f90", api=api)
>>> psy=PSyFactory(api).create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import Dynamo0p3AsyncHaloExchangeTrans
```

(continues on next page)

```
>>> trans = Dynamo0p3AsyncHaloExchangeTrans()
>>> new_schedule, memento = trans.apply(schedule.children[0])
>>> new_schedule.view()
```

apply(*node*)

Transforms a synchronous halo exchange, represented by a HaloExchange node, into an asynchronous halo exchange, represented by HaloExchangeStart and HaloExchangeEnd nodes.

> **Parameters node** (psyclone.psygen.HaloExchange) – A synchronous haloexchange node
>
> **Returns** Tuple of the modified schedule and a record of the transformation.
>
> **Return type** (*psyclone.psyGen.Schedule*, psyclone.undoredo.Memento)

name

> **Returns** the name of this transformation as a string.
>
> **Return type** str

class psyclone.transformations.**Dynamo0p3ColourTrans**

Split a Dynamo 0.3 loop over cells into colours so that it can be parallelised. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> import transformations
>>> import os
>>> import pytest
>>>
>>> TEST_API = "dynamo0.3"
>>> _,info=parse(os.path.join(os.path.dirname(os.path.abspath(__file__)),
>>>              "tests", "test_files", "dynamo0p3",
>>>              "4.6_multikernel_invokes.f90"),
>>>              api=TEST_API)
>>> psy = PSyFactory(TEST_API).create(info)
>>> invoke = psy.invokes.get('invoke_0')
>>> schedule = invoke.schedule
>>>
>>> ctrans = Dynamo0p3ColourTrans()
>>> otrans = DynamoOMPParallelLoopTrans()
>>>
>>> # Colour all of the loops
>>> for child in schedule.children:
>>>     cschedule, _ = ctrans.apply(child)
>>>
>>> # Then apply OpenMP to each of the colour loops
>>> schedule = cschedule
>>> for child in schedule.children:
>>>     newsched, _ = otrans.apply(child.children[0])
>>>
>>> newsched.view()
```

Colouring in the Dynamo 0.3 API is subject to the following rules:

- Only kernels with an iteration space of CELLS and which modify a continuous field require colouring. Any other type of loop will cause this transformation to raise an exception.

- A kernel may have at most one field with 'INC' access

- A separate colour map will be required for each field that is coloured (if an invoke contains >1 kernel call)

apply(*node*)

Performs Dynamo0.3-specific error checking and then uses the parent class to convert the Loop represented by `node` into a nested loop where the outer loop is over colours and the inner loop is over cells of that colour.

> **Parameters** **node** (`psyclone.dynamo0p3.DynLoop`) – the loop to transform.
>
> **Returns** 2-tuple of new schedule and memento of transform
>
> **Return type** (`psyclone.dynamo0p3.DynInvokeSchedule`, `psyclone.undoredo.Memento`)

name

Returns the name of this transformation as a string.

class psyclone.transformations.**Dynamo0p3KernelConstTrans**

Modifies a kernel so that the number of dofs, number of layers and number of quadrature points are fixed in the kernel rather than being passed in by argument.

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "dynamo0.3"
>>> ast, invokeInfo = parse("file.f90", api=api)
>>> psy=PSyFactory(api).create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import Dynamo0p3KernelConstTrans
>>> trans = Dynamo0p3KernelConstTrans()
>>> for kernel in schedule.kern_calls():
>>>     new_schedule, _ = trans.apply(kernel)
>>>     kernel_schedule = kernel.get_kernel_schedule()
>>>     kernel_schedule.symbol_table.view()
```

apply(*node*, *cellshape='quadrilateral'*, *element_order=None*, *number_of_layers=None*, *quadrature=False*)

Transforms a kernel so that the values for the number of degrees of freedom (if a valid value for the element_order arg is provided), the number of quadrature points (if the quadrature arg is set to True) and the number of layers (if a valid value for the number_of_layers arg is provided) are constant in a kernel rather than being passed in by argument.

The "cellshape", "element_order" and "number_of_layers" arguments are provided to mirror the namelist values that are input into an LFRic model when it is run.

Quadrature support is currently limited to XYoZ in ths transformation. In the case of XYoZ the number of quadrature points (for horizontal and vertical) are set to the element_order + 3 in the LFRic infrastructure so their value is derived.

> **Parameters** **node** (`psyclone.psygen.DynKern`) – A kernel node
>
> **Returns** Tuple of the modified schedule and a record of the transformation.
>
> **Return type** (*psyclone.psyGen.Schedule*, `psyclone.undoredo.Memento`)

name

> **Returns** the name of this transformation as a string.
>
> **Return type** str

---

class psyclone.transformations.**Dynamo0p3OMPLoopTrans**(*omp_schedule='static'*)
  Dynamo 0.3 specific orphan OpenMP loop transformation. Adds Dynamo-specific validity checks. Actual
  transformation is done by *base class*.

  apply(*node*, *reprod=None*)
    Perform Dynamo 0.3 specific loop validity checks then call *OMPLoopTrans.apply()*.

      Parameters

        • **node** (*psyclone.psyGen.Node*) – the Node in the Schedule to check

        • **reprod**   (*bool or None (default, which indicates to use the
          default from the config file*)) – if reproducible reductions should be used.

      Raises ***TransformationError*** – if an OMP loop transform would create incorrect code.

  name
    Returns the name of this transformation as a string.

class psyclone.transformations.**Dynamo0p3RedundantComputationTrans**
  This transformation allows the user to modify a loop's bounds so that redundant computation will be performed.
  Redundant computation can result in halo exchanges being modified, new halo exchanges being added or exist-
  ing halo exchanges being removed.

    • This transformation should be performed before any parallelisation transformations (e.g. for OpenMP) to
      the loop in question and will raise an exception if this is not the case.

    • This transformation can not be applied to a loop containing a reduction and will again raise an exception
      if this is the case.

    • This transformation can only be used to add redundant computation to a loop, not to remove it.

    • This transformation allows a loop that is already performing redundant computation to be modified, but
      only if the depth is increased.

  apply(*loop*, *depth=None*)
    Apply the redundant computation tranformation to the loop loop. This transformation can be applied to
    loops iterating over 'cells or 'dofs'. if depth is set to a value then the value will be the depth of the field's
    halo over which redundant computation will be performed. If depth is not set to a value then redundant
    computation will be performed to the full depth of the field's halo.

      Parameters

        • **loop** (*psyclone.psyGen.DynLoop*) – the loop that we are transforming

        • **depth** (*int or None*) – the depth of the stencil. Defaults to None if a depth is not
          provided.

  name
    Returns the name of this transformation as a string.

class psyclone.transformations.**DynamoExtractRegionTrans**
  Dynamo0.3 API application of ExtractRegionTrans transformation to extract code into a stand-alone program.
  For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>>
>>> API = "dynamo0.3"
>>> FILENAME = "solver_alg.x90"
>>> ast, invokeInfo = parse(FILENAME, api=API)
>>> psy = PSyFactory(API, distributed_memory=False).create(invoke_info)
```

(continues on next page)

```
>>> schedule = psy.invokes.get('invoke_0').schedule
>>>
>>> from psyclone.transformations import DynamoExtractRegionTrans
>>> etrans =  DynamoExtractRegionTrans()
>>>
>>> # Apply DynamoExtractRegionTrans transformation to selected Nodes
>>> newsched, _ = etrans.apply(schedule.children[0:3])
>>> newsched.view()
```

> **name**
>> Returns the name of this transformation as a string.

class psyclone.transformations.**DynamoLoopFuseTrans**
> Performs error checking before calling the *apply()* method of the *base class* in order to fuse two Dynamo loops.
>
> **apply**(*node1*, *node2*, *same_space=False*)
>> Fuse the two Dynamo loops represented by node1 and node2. The optional same_space flag asserts that an unknown iteration space (i.e. any_space) matches the other iteration space. This is set at the users own risk.
>>
>> **Parameters**
>>
>>> • **node1** (psyclone.dynamo0p3.DynLoop) – First Loop to fuse.
>>>
>>> • **node2** (psyclone.dynamo0p3.DynLoop) – Second Loop to fuse.
>>
>> **Returns** two-tuple of modified Schedule and Memento
>>
>> **Return type** *psyclone.psyGen.Schedule*, psyclone.undoredo.Memento
>>
>> **Raises** *TransformationError* – if either of the supplied loops contains an inter-grid kernel.
>
> **name**
>> Returns the name of this transformation as a string.

class psyclone.transformations.**DynamoOMPParallelLoopTrans**(*omp_schedule='static'*)
> Dynamo-specific OpenMP loop transformation. Adds Dynamo specific validity checks. Actual transformation is done by the *base class*.
>
> **apply**(*node*)
>> Perform Dynamo specific loop validity checks then call the *apply()* method of the *base class*.
>>
>> **Parameters node** (*psyclone.psyGen.Node*) – the Node in the Schedule to check
>>
>> **Raises** *TransformationError* – if the associated loop requires colouring.
>
> **name**
>> Returns the name of this transformation as a string.

class psyclone.transformations.**ExtractRegionTrans**
> Provides a transformation to extract code represented by a subset of the Nodes in the PSyIR of a Schedule into a stand-alone program. Examples are given in descriptions of children classes DynamoExtractRegionTrans and GOceanExtractRegionTrans.
>
> After applying the transformation the Nodes marked for extraction are children of the ExtractNode. Nodes to extract can be individual constructs within an Invoke (e.g. Loops containing a Kernel or BuiltIn call) or entire Invokes. This functionality does not support distributed memory.

apply(*nodes*)

> Apply this transformation to a subset of the Nodes within a Schedule - i.e. enclose the specified Nodes in the Schedule within a single Extract region.

> > **Parameters** **nodes** ((list of) *psyclone.psyGen.Node.*) – a single Node or a list of Nodes.

> > **Returns** tuple of the modified Schedule and a record of the transformation.

> > **Return type** (*psyclone.psyGen.Schedule*, psyclone.undoredo.Memento).

> > **Raises** *TransformationError* – if the *nodes* argument is not of the correct type.

name

> Returns the name of this transformation as a string.

psyGen = <module 'psyclone.psyGen' from '/home/rupert/proj/PSyclone/src/psyclone/psyGen

class psyclone.transformations.**GOConstLoopBoundsTrans**

> Switch on (or off) the use of constant loop bounds within a GOInvokeSchedule. In the absence of constant loop bounds, PSyclone will generate loops where the bounds are obtained by de-referencing a field object, e.g.:

```
DO j = my_field%grid%internal%ystart, my_field%grid%internal%ystop
```

> Some compilers are able to produce more efficient code if they are provided with information on the relative trip-counts of the loops within an Invoke. With constant loop bounds switched on, PSyclone generates code like:

```
ny = my_field%grid%subdomain%internal%ystop
...
DO j = 1, ny-1
```

> In practice, the application of the constant loop bounds looks something like, e.g.:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> import os
>>> TEST_API = "gocean1.0"
>>> _, info = parse(os.path.join("tests", "test_files", "gocean1p0",
>>>                              "single_invoke.f90"),
>>>                  api=TEST_API)
>>> psy = PSyFactory(TEST_API).create(info)
>>> invoke = psy.invokes.get('invoke_0_compute_cu')
>>> schedule = invoke.schedule
>>>
>>> from psyclone.transformations import GOConstLoopBoundsTrans
>>> clbtrans = GOConstLoopBoundsTrans()
>>>
>>> newsched, _ = clbtrans.apply(schedule)
>>> # or, to turn off const. looop bounds:
>>> # newsched, _ = clbtrans.apply(schedule, const_bounds=False)
>>>
>>> newsched.view()
```

> **apply**(*node*, *const_bounds=True*)

> > Switches constant loop bounds on or off for all loops in a GOInvokeSchedule. Default is 'on'.

> > **Parameters**

> > > • **node** (psyclone.gocean1p0.GOInvokeSchedule) – The GOInvokeSchedule of which all loops will get the constant loop bounds switched on or off.

- **const_bounds** (*bool*) – If the constant loop should be used (True) or not (False). Default is True.

name
> Return the name of the Transformation as a string.

class psyclone.transformations.**GOLoopSwapTrans**
> Provides a loop-swap transformation, e.g.:

```
DO j=1, m
   DO i=1, n
```

becomes:

```
DO i=1, n
   DO j=1, m
```

This transform is used as follows:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast, invokeInfo = parse("shallow_alg.f90")
>>> psy = PSyFactory("gocean1.0").create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import GOLoopSwapTrans
>>> swap = GOLoopSwapTrans()
>>> new_schedule, memento = swap.apply(schedule.children[0])
>>> new_schedule.view()
```

apply(*outer*)
> The argument outer must be a loop which has exactly one inner loop. This transform then swaps the outer and inner loop.

> > **Parameters outer** (*psyclone.psyGen.Loop*) – The node representing the outer loop.

> > **Returns** A tuple consisting of the new schedule, and a Memento.

> > **Raises** *TransformationError* – if the supplied node does not allow a loop swap to be done.

name
> Returns the name of this transformation as a string.

class psyclone.transformations.**GOceanExtractRegionTrans**
> GOcean1.0 API application of ExtractRegionTrans transformation to extract code into a stand-alone program. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>>
>>> API = "gocean1.0"
>>> FILENAME = "shallow_alg.f90"
>>> ast, invokeInfo = parse(FILENAME, api=API)
>>> psy = PSyFactory(API, distributed_memory=False).create(invoke_info)
>>> schedule = psy.invokes.get('invoke_0').schedule
>>>
>>> from psyclone.transformations import GOceanExtractRegionTrans
>>> etrans = GOceanExtractRegionTrans()
```

(continues on next page)

```
>>>
>>> # Apply GOceanExtractRegionTrans transformation to selected Nodes
>>> newsched, _ = etrans.apply(schedule.children[0])
>>> newsched.view()
```

name
Returns the name of this transformation as a string.

class psyclone.transformations.**GOceanLoopFuseTrans**
Performs error checking (that the loops are over the same grid-point type) before calling the *LoopFuseTrans.apply()* method of the *base class* in order to fuse two GOcean loops.

apply(*node1*, *node2*)
Fuse the two GOcean loops represented by node1 and node2.

   Parameters

   • **node1** (psyclone.gocean1p0.GOLoop) – A node representing a GOLoop.

   • **node2** (psyclone.gocean1p0.GOLoop) – A node representing a GOLoop.

   Raises *TransformationError* – if the supplied node2 can not be fused, e.g. not all nodes are loops, don't have the same parent, are not next to each other or have different iteration spaces.

name
Returns the name of this transformation as a string.

class psyclone.transformations.**GOceanOMPLoopTrans**(*omp_schedule='static'*)
GOcean-specific orphan OpenMP loop transformation. Adds GOcean specific validity checks (that the node is either an inner or outer Loop). Actual transformation is done by *base class*.

   Parameters **omp_schedule** – The omp schedule to be created. Must be one of 'runtime', 'static', 'dynamic', 'guided' or 'auto'.

apply(*node*)
Perform GOcean specific loop validity checks then call *OMPLoopTrans.apply()*.

   Parameters **node** (*psyclone.psyGen.Loop.*) – The loop to parallelise using OMP Do.

name
Returns the name of this transformation as a string.

class psyclone.transformations.**GOceanOMPParallelLoopTrans**(*omp_schedule='static'*)
GOcean specific OpenMP Do loop transformation. Adds GOcean specific validity checks (that supplied Loop is an inner or outer loop). Actual transformation is done by *base class*.

   Parameters **omp_schedule** – The omp schedule to be created. Must be one of 'runtime', 'static', 'dynamic', 'guided' or 'auto'.

apply(*node*)
Perform GOcean-specific loop validity checks then call *OMPParallelLoopTrans.apply()*.

   Parameters **node** (*psyclone.psyGen.Loop*) – A Loop node from an AST.

   Raises *TransformationError* – if the supplied node is not an inner or outer loop.

name
Returns the name of this transformation as a string.

class psyclone.transformations.**KernelModuleInlineTrans**
Switches on, or switches off, the inlining of a Kernel subroutine into the PSy layer module. For example:

```
>>> invoke = ...
>>> schedule = invoke.schedule
>>>
>>> inline_trans = KernelModuleInlineTrans()
>>>
>>> ischedule, _ = inline_trans.apply(schedule.children[0].children[0])
>>> ischedule.view()
```

> **Warning:** For this transformation to work correctly, the Kernel subroutine must only use data that is passed in by argument, declared locally or included via use association within the subroutine. Two examples where in-lining will not work correctly are:
>
> 1. A variable is declared within the module that `contains` the Kernel subroutine and is then accessed within that Kernel;
>
> 2. A variable is included via use association at the module level and accessed within the Kernel subroutine.
>
> *There are currently no checks that these rules are being followed when in-lining so the onus is on the user to ensure correctness.*

apply(*node*, *inline=True*)

Checks that the node is of the correct type (a Kernel) then marks the Kernel to be inlined, or not, depending on the value of the inline argument. If the inline argument is not passed the Kernel is marked to be inlined.

name

Returns the name of this transformation as a string.

validate(*node*, *inline*)

Check that the supplied kernel is eligible to be module inlined.

> **Parameters**
>
> - **node** (sub-class of *psyclone.psyGen.Node*) – the node in the PSyIR that is to be module inlined.
>
> - **inline** (*bool*) – whether or not the kernel is to be inlined.
>
> **Raises**
>
> - *TransformationError* – if the supplied node is not a kernel.
>
> - *TransformationError* – if the supplied kernel has itself been transformed (Issue #229).

class psyclone.transformations.**LoopFuseTrans**

Provides a loop-fuse transformation. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast,invokeInfo=parse("dynamo.F90")
>>> psy=PSyFactory("dynamo0.1").create(invokeInfo)
>>> schedule=psy.invokes.get('invoke_v3_kernel_type').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import LoopFuseTrans
>>> trans=LoopFuseTrans()
>>> new_schedule,memento=trans.apply(schedule.children[0],
```

```
                                        schedule.children[1])
>>> new_schedule.view()
```

apply(*node1*, *node2*)
>   Fuse the loops represented by `node1` and `node2`.

name
>   Returns the name of this transformation as a string.

class psyclone.transformations.**MoveTrans**
>   Provides a transformation to move a node in the tree. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast,invokeInfo=parse("dynamo.F90")
>>> psy=PSyFactory("dynamo0.3").create(invokeInfo)
>>> schedule=psy.invokes.get('invoke_v3_kernel_type').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import MoveTrans
>>> trans=MoveTrans()
>>> new_schedule, memento = trans.apply(schedule.children[0],
                                        schedule.children[2],
                                        position="after")
>>> new_schedule.view()
```

> Nodes may only be moved to a new location with the same parent and must not break any dependencies otherwise an exception is raised.

apply(*node*, *location*, *position='before'*)
>   Move the node represented by `node` before location `location` (which is also a node) by default and after if the optional *position* argument is set to 'after'. An exception is raised if the move is invalid.

name
>   Returns the name of this transformation as a string.

class psyclone.transformations.**NemoExplicitLoopTrans**
>   Transforms the outermost array slice in an implicit loop in a NEMOInvokeSchedule into an explicit loop. For example, if "implicit_loop.f90" contained:

```
my_array(:, :, :) = 1.0
```

> then doing:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "nemo"
>>> filename = "implicit_loop.f90"
>>> ast, invokeInfo = parse(filename, api=api)
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.transformations import NemoExplicitLoopTrans
>>> rtrans = NemoExplicitLoopTrans()
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> loop = schedule.children[0]
>>> newloop, _ = rtrans.apply(loop)
```

will create a new NemoLoop object for an explicit loop over levels (the outermost slice) that then contains an implicit loop:

```
DO jk = 1, jpk
  my_array(:, :, jk) = 1.0
END DO
```

Subsequently applying *rtrans* to *newloop* will create:

```
DO jk = 1, jpk
  DO jj = 1, jpj
    my_array(:, jj, jk) = 1.0
  END DO
END DO
```

apply(*loop*)

> Transform the outermost array slice in the supplied implicit loop into an explicit loop.
>
> > **Parameters loop** (`psyclone.nemo.NemoImplicitLoop`) – the NemoImplicitLoop to transform.
> >
> > **Returns** a new PSyIR loop object and a memento of the transformation.
> >
> > **Return type** (`psyclone.nemo.NemoLoop, psyclone.undoredo.Memento`)
> >
> > **Raises**
> >
> > - **NotImplementedError** – if the array slice has explicit bounds.
> >
> > - **TransformationError** – if an array slice is not in dimensions 1-3 of the array.

name

> > **Returns** the name of this transformation class.
> >
> > **Return type** str

validate(*loop*)

> Check that the supplied loop is a valid target for this transformation.
>
> > **Parameters loop** (`psyclone.nemo.NemoImplicitLoop`) – the loop node to validate.
> >
> > **Raises** **TransformationError** – if the supplied loop is not a NemoImplicitLoop.

class `psyclone.transformations.`**OCLTrans**

> Switches on/off the generation of an OpenCL PSy layer for a given InvokeSchedule. For example:

```
>>> invoke = ...
>>> schedule = invoke.schedule
>>>
>>> ocl_trans = OCLTrans()
>>> new_sched, _ = ocl_trans.apply(schedule)
```

apply(*sched*, *opencl=True*)

> Apply the OpenCL transformation to the supplied GOInvokeSchedule. This causes PSyclone to generate an OpenCL version of the corresponding PSy-layer routine. The generated code makes use of the FortCL library (https://github.com/stfc/FortCL) in order to manage the OpenCL device directly from Fortran.
>
> > **Parameters**
> >
> > - **sched** (`psyclone.psyGen.GOInvokeSchedule`) – InvokeSchedule to transform.
> >
> > - **opencl** (*bool*) – whether or not to enable OpenCL generation.

name

> **Returns** the name of this transformation.
>
> **Return type** str

class psyclone.transformations.**OMPLoopTrans**(*omp_schedule='static'*)

> Adds an orphaned OpenMP directive to a loop. i.e. the directive must be inside the scope of some other OMP Parallel REGION. This condition is tested at code-generation time. The optional 'reprod' argument in the apply method decides whether standard OpenMP reduction support is to be used (which is not reproducible) or whether a manual reproducible reproduction is to be used.
>
> > **Parameters** **omp_schedule** (*str*) – The OpenMP schedule to use.
>
> For example:

```python
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory, GenerationError
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>> print psy.invokes.names
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ltrans = t.get_trans_name('OMPLoopTrans')
>>> rtrans = t.get_trans_name('OMPParallelTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> new_schedule = schedule
>>>
# Apply the OpenMP Loop transformation to *every* loop
# in the schedule
>>> for child in schedule.children:
>>>     newschedule, memento = ltrans.apply(child, reprod=True)
>>>     schedule = newschedule
>>>
# Enclose all of these loops within a single OpenMP
# PARALLEL region
>>> rtrans.omp_schedule("dynamic,1")
>>> newschedule, memento = rtrans.apply(schedule.children)
>>>
>>>
```

apply(*node*, *reprod=None*)

> Apply the OMPLoopTrans transformation to the specified node in a Schedule. This node must be a Loop since this transformation corresponds to wrapping the generated code with directives like so:

```fortran
!$OMP DO
do ...
   ...
end do
!$OMP END DO
```

> At code-generation time (when OMPLoopDirective.gen_code() is called), this node must be within (i.e. a child of) an OpenMP PARALLEL region.

The optional reprod argument will cause a reproducible reduction to be generated if it is set to True, otherwise the default value (as read from the psyclone.cfg file) will be used. Note, reproducible in this case means obtaining the same results with the same number of OpenMP threads, not for different numbers of OpenMP threads.

> Parameters
>
> - **node** (*psyclone.psyGen.Node*) – the supplied node to which we will apply the OMPLoopTrans transformation
>
> - **reprod** (*Boolean or None*) – optional argument to determine whether to generate reproducible OpenMP reductions (True) or not (False). The default value is None which will cause PSyclone to look up a default value
>
> Returns (*psyclone.psyGen.Schedule*, psyclone.undoredo.Memento)

name
> Returns the name of this transformation as a string.

omp_schedule
> Returns the OpenMP schedule that will be specified by this transformation. The default schedule is 'static'.

class psyclone.transformations.**OMPParallelLoopTrans**(*omp_schedule='static'*)
> Adds an OpenMP PARALLEL DO directive to a loop.

> For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> ast, invokeInfo = parse("dynamo.F90")
>>> psy = PSyFactory("dynamo0.1").create(invokeInfo)
>>> schedule = psy.invokes.get('invoke_v3_kernel_type').schedule
>>> schedule.view()
>>>
>>> from psyclone.transformations import OMPParallelLoopTrans
>>> trans = OMPParallelLoopTrans()
>>> new_schedule, memento = trans.apply(schedule.children[0])
>>> new_schedule.view()
```

apply(*node*)
> Apply an OMPParallelLoop Transformation to the supplied node (which must be a Loop). In the generated code this corresponds to wrapping the Loop with directives:

```
!$OMP PARALLEL DO ...
do ...
   ...
end do
!$OMP END PARALLEL DO
```

> Parameters **node** (psyclone.f2pygen.DoGen) – the node (loop) to which to apply the transformation.

> Returns Two-tuple of transformed schedule and a record of the transformation.

> Return type (psyclone.psyGen.Schedule, :py:class:`psyclone.undoredo.Memento)

name
> Returns the name of this transformation as a string.

class psyclone.transformations.**OMPParallelTrans**

Create an OpenMP PARALLEL region by inserting directives. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory, GenerationError
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> ltrans = t.get_trans_name('GOceanOMPLoopTrans')
>>> rtrans = t.get_trans_name('OMPParallelTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>> new_schedule = schedule
>>>
>>> # Apply the OpenMP Loop transformation to *every* loop
>>> # in the schedule
>>> for child in schedule.children:
>>>     newschedule, memento = ltrans.apply(child)
>>>     schedule = newschedule
>>>
>>> # Enclose all of these loops within a single OpenMP
>>> # PARALLEL region
>>> newschedule, _ = rtrans.apply(schedule.children)
>>> newschedule.view()
```

name

> **Returns** The name of this transformation as a string.
>
> **Return type** str

psyGen = <module 'psyclone.psyGen' from '/home/rupert/proj/PSyclone/src/psyclone/psyGen

class psyclone.transformations.**ParallelLoopTrans**

Adds an orphaned directive to a loop indicating that it should be parallelised. i.e. the directive must be inside the scope of some other Parallel REGION. This condition is tested at code-generation time.

apply(*node*, *collapse=None*)

Apply the Loop transformation to the specified node in a Schedule. This node must be a Loop since this transformation corresponds to wrapping the generated code with directives, e.g. for OpenMP:

```
!$OMP DO
do ...
    ...
end do
!$OMP END DO
```

At code-generation time (when gen_code()' is called), this node must be within (i.e. a child of) a PARALLEL region.

> **Parameters**
>
> - **node** (*psyclone.psyGen.Node.*) – the supplied node to which we will apply the Loop transformation.

- **collapse** (*int*) – number of loops to collapse into single iteration space or None.

> **Returns** (*psyclone.psyGen.Schedule*, psyclone.undoredo.Memento)

name

> **Returns** the name of this transformation as a string.

> **Return type** str

class psyclone.transformations.**ParallelRegionTrans**
Base class for transformations that create a parallel region.

apply(*nodes*)
Apply this transformation to a subset of the nodes within a schedule - i.e. enclose the specified Loops in the schedule within a single parallel region.

> **Parameters nodes** ((list of) *psyclone.psyGen.Node*.) – a single Node or a list of Nodes.

> **Raises** *TransformationError* – if the nodes argument is not of the correct type.

name
Returns the name of this transformation as a string.

class psyclone.transformations.**ProfileRegionTrans**
Create a profile region around a list of statements. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.parse.utils import ParseError
>>> from psyclone.psyGen import PSyFactory, GenerationError
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> p_trans = t.get_trans_name('ProfileRegionTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> # Enclose all children within a single profile region
>>> newschedule, _ = p_trans.apply(schedule.children)
>>> newschedule.view()
```

apply(*nodes*)
Apply this transformation to a subset of the nodes within a schedule - i.e. enclose the specified Nodes in the schedule within a single profiler region.

> **Parameters nodes** (psyclone.psygen.Node or list of psyclone.psygen.Node.) – Can be a single node or a list of nodes.

name
Returns the name of this transformation as a string

profiler = <module 'psyclone.profiler' from '/home/rupert/proj/PSyclone/src/psyclone/p

psyGen = <module 'psyclone.psyGen' from '/home/rupert/proj/PSyclone/src/psyclone/psyGen

class psyclone.transformations.**RegionTrans**
This abstract class is a base class for all transformations that act on a list of nodes. It gives access to a _validate function that makes sure that the nodes in the list are in the same order as in the original AST, no node is

---

duplicated, and that all nodes have the same parent. We also check that all nodes to be enclosed are valid for this transformation - this requires that the sub-class populate the *valid_node_types* tuple.

exception `psyclone.transformations.`**`TransformationError`**(*value*)
> Provides a PSyclone-specific error class for errors found during code transformation operations.

`psyclone.transformations.`**`check_intergrid`**(*node*)
> Utility function to check that the supplied node does not have an intergrid kernel as a child.

> This is used to ensure that we reject any attempt to apply transformations to loops containing inter-grid kernels. (This restriction will be lifted in Issue #134 and this routine can then be removed.)

> # TODO remove this routine once #134 is complete.

> > **Parameters node** (`psyGen.Node`) – The PSyIR node to check.

> > **Raises** [`TransformationError`](#) – if the supplied node has an inter-grid kernel as a child

16.4 The psyGen module

This module provides generic support for PSyclone's PSy code optimisation and generation. The classes in this method need to be specialised for a particular API and implementation.

class `psyclone.psyGen.`**`PSy`**(*invoke_info*)
> Base class to help manage and generate PSy code for a single algorithm file. Takes the invocation information output from the function `parse.algorithm.parse()` as its input and stores this in a way suitable for optimisation and code generation.

> > **Parameters invoke_info** (`psyclone.parse.algorithm.FileInfo`) – An object containing the required invocation information for code optimisation and generation. Produced by the function `parse.algorithm.parse()`.

> For example:

```
>>> from psyclone.parse.algorithm import parse
>>> ast, info = parse("argspec.F90")
>>> from psyclone.psyGen import PSyFactory
>>> api = "..."
>>> psy = PSyFactory(api).create(info)
>>> print(psy.gen)
```

> **inline**(*module*)
> > inline all kernel subroutines into the module that are marked for inlining. Avoid inlining the same kernel more than once.

16.5 The algGen module

This module provides the Alg class and supporting exception-handling to translate the original algorithm file into one that can be compiled and linked with the generated PSy code.

class `psyclone.algGen.`**`Alg`**(*parse_tree*, *psy*, *invoke_name='invoke'*)
> Generate a modified algorithm code for a single algorithm specification. Takes the parse tree of the algorithm specification output from the function [`psyclone.parse.algorithm.parse()`](#) and an instance of the `psyGen.PSy` class as input. The latter allows consistent names to be generated between the algorithm (callng) and psy (callee) layers.

> For example:

```
>>> from psyclone.algorithm.parse import parse
>>> parse_tree, info = parse("argspec.F90")
>>> from psyclone.psyGen import PSy
>>> psy = PSy(info)
>>> from psyclone.algGen import Alg
>>> alg = Alg(parse_tree, psy)
>>> print(alg.gen)
```

Parameters

- **parse_tree** (fparser.two.utils.Base) – An object containing a parse tree of the algorithm specification which was produced by the function *psyclone.parse. algorithm.parse()*. Assumes the algorithm will be parsed by fparser2 and expects a valid program unit, program, module, subroutine or function.

- **psy** (*psyclone.psyGen.PSy*) – An object containing information about the PSy layer.

- **invoke_name** (*str*) – The name that the algorithm layer uses to indicate an invoke call. This is an optional argument that defaults to the name "invoke".

gen

Return modified algorithm code.

Returns The modified algorithm specification as an fparser2 parse tree.

Return type fparser.two.utils.Base

16.6 The line_length module

Provides support for breaking long fortran lines into smaller ones to allow the code to conform to the maximum line length limits (132 for f90 free format is the default)

class psyclone.line_length.**FortLineLength**(*line_length=132*)

This class take a free format fortran code as a string and line wraps any lines that are larger than the specified line length

length

returns the maximum allowed line length

long_lines(*fortran_in*)

returns true if at least one of the lines in the input code is longer than the allowed length. Otherwise returns false

process(*fortran_in*)

takes fortran code as a string as input and output fortran code as a string with any long lines wrapped appropriately

PROFILING

PSyclone has the ability to define regions that can be profiled with various performance measurement tools. The profiling can be enabled automatically using command line parameters like:

```
psyclone --profile kernels ...
```

Or more fine grained by applying a profiling transformation in a transformation script.

PSyclone can be used with a variety of existing profiling tools. It currently supports dl_timer, and comes with a simple stand-alone timer library. An application needs to be able to find the module files for the profile wrapper, and needs to be linked with the included wrapper library that interfaces between the PSyclone API and the tool-specific API. It is the responsibility of the user to supply the corresponding compiler command line options when building the application that incorporates the PSyclone-generated code.

17.1 Profiling API

In order to be used with different profiling tools, PSyclone supports a simple profiling API. For each existing profiling tool a simple interface library needs to be implemented that maps the PSyclone profiling calls to the corresponding call for the profiling tool.

PSyclone utilises 4 profiling calls which are described in the following sub-sections.

17.1.1 ProfileInit()

This method needs to be called once to initialise the profiling tool. At this stage this call is not automatically inserted by PSyclone, so it is the responsibility of the user to add the call to an appropriate location in the application:

```
use profile_mod, only : ProfileInit
...
call ProfileInit()
```

The 'appropriate' location might depend on the profiling library used. For example, it might be necessary to invoke this before or after a call to `MPI_Init()`.

17.1.2 ProfileFinalise()

At the end of the program the function `ProfileFinalise()` must be called. It will make sure that the measurements are printed or written to file correctly, and that the profiling tool is closed correctly. Again at this stage it is necessary to manually insert the call at an appropriate location:

```
use profile_mod, only : ProfileFinalise
...
call ProfileFinalise()
```

And again the appropriate location might depend on the profiling library used (e.g. before or after a call to `MPI_Finalize()`).

17.1.3 ProfileStart()/ProfileEnd()

The `ProfileStart` and `ProfileEnd` functions define the beginning and end of a region to be measured. In general it is up to the user what exactly a region is, arbitrary code can be enclosed, as long as it is guaranteed that each call of `ProfileStart` is matched with exactly one corresponding call to `ProfileEnd`. PSyclone supplies one saved (static) variable of type `ProfileData` to each matching Start/End pair.

This is the code sequence which is created by PSyclone:

```
use profile_mod, only : ProfileData, ProfileStart, ProfileEnd
...
type(ProfileData), save :: profiler_data
...
call ProfileStart("Module", "Region", profiler_data)
...
call ProfileEnd(profiler_data)
```

PSyclone guarantees that different ProfileStart/End pairs have different `ProfileData` variables.

17.2 Profiling Command Line Options

PSyclone offers two command line options to automatically instrument code with profiling regions. It can create profile regions around a full invoke (including all kernel calls in this invoke), and/or around each individual kernel.

The option `--profile invokes` will automatically add a call to `ProfileStart` and `ProfileEnd` at the beginning and end of every invoke subroutine created by PSyclone. All kernels called within this invoke subroutine will be included in the profiled region.

The option `--profile kernels` will add a call to `ProfileStart` before any loops created by PSyclone, and a `ProfileEnd` call at the end of the loop. Two caveats:

1. In some APIs (for example dynamo when using distributed memory) additional minor code might get included in a profiled kernel section, for example setDirty() calls (expensive calls like HaloExchange are excluded).

2. If loop transforms are applied using a script, the profiling nodes added to the AST will very likely cause errors in the script or in the generated output. As example consider a case where an OMPLoop transform is applied to a loop. With profiling enabled instead of the expected loop there could be a profile node in the AST (with the loop as child). Since an OMP DO directive can only have loops inside, and it now has a call to `ProfileStart`, the generated code is incorrect and will not compile.

In order to avoid the second issue, automatic profiling using `--profile` is not allowed together with a transformation script. On the other hand, since it is possible to write scripts that are more flexible in handling a modified AST, you can use the command line option `--force-profile`. It takes the same parameters as `--profile`, and will allow you to combine a transformation script together with automatic profiling. Use this option at your own risk!

It is also the responsibility of the user to manually add the calls to `ProfileInit` and `ProfileFinalise` to the code base.

PSyclone will modify the schedule of each invoke to insert the profiling regions. Below we show an example of a schedule created when instrumenting invokes - all children of a Profile-Node will be part of the profiling region, including all loops created by PSyclone and all kernel calls:

```
GOInvokeSchedule[invoke='invoke_1',Constant loop bounds=True]
    [Profile]
        Loop[type='outer',field_space='cu',it_space='internal_pts']
            Loop[type='inner',field_space='cu',it_space='internal_pts']
                KernCall compute_unew_code(unew_fld,uold_fld,z_fld,cv_fld,h_fld,tdt,
↪dy) [module_inline=False]
        Loop[type='outer',field_space='cv',it_space='internal_pts']
            Loop[type='inner',field_space='cv',it_space='internal_pts']
                KernCall compute_vnew_code(vnew_fld,vold_fld,z_fld,cu_fld,h_fld,tdt,
↪dy) [module_inline=False]
        Loop[type='outer',field_space='ct',it_space='internal_pts']
            Loop[type='inner',field_space='ct',it_space='internal_pts']
                KernCall compute_pnew_code(pnew_fld,pold_fld,cu_fld,cv_fld,tdt,dx,dy)␣
↪[module_inline=False]
```

And now the same schedule when instrumenting kernels. In this case each loop nest and kernel call will be contained in a separate region:

```
GOInvokeSchedule[invoke='invoke_1',Constant loop bounds=True]
    [Profile]
        Loop[type='outer',field_space='cu',it_space='internal_pts']
            Loop[type='inner',field_space='cu',it_space='internal_pts']
                KernCall compute_unew_code(unew_fld,uold_fld,z_fld,cv_fld,h_fld,tdt,
↪dy) [module_inline=False]
    [Profile]
        Loop[type='outer',field_space='cv',it_space='internal_pts']
            Loop[type='inner',field_space='cv',it_space='internal_pts']
                KernCall compute_vnew_code(vnew_fld,vold_fld,z_fld,cu_fld,h_fld,tdt,
↪dy) [module_inline=False]
    [Profile]
        Loop[type='outer',field_space='ct',it_space='internal_pts']
            Loop[type='inner',field_space='ct',it_space='internal_pts']
                KernCall compute_pnew_code(pnew_fld,pold_fld,cu_fld,cv_fld,tdt,dx,dy)␣
↪[module_inline=False]
```

Both options can be specified at the same time:

```
GOInvokeSchedule[invoke='invoke_1',Constant loop bounds=True]
    [Profile]
        [Profile]
            Loop[type='outer',field_space='cu',it_space='internal_pts']
                Loop[type='inner',field_space='cu',it_space='internal_pts']
                    KernCall compute_unew_code(unew_fld,uold_fld,z_fld,cv_fld,h_fld,
↪tdt,dy) [module_inline=False]
        [Profile]
            Loop[type='outer',field_space='cv',it_space='internal_pts']
                Loop[type='inner',field_space='cv',it_space='internal_pts']
                    KernCall compute_vnew_code(vnew_fld,vold_fld,z_fld,cu_fld,h_fld,
↪tdt,dy) [module_inline=False]
        [Profile]
            Loop[type='outer',field_space='ct',it_space='internal_pts']
                Loop[type='inner',field_space='ct',it_space='internal_pts']
                    KernCall compute_pnew_code(pnew_fld,pold_fld,cu_fld,cv_fld,tdt,dx,
↪dy) [module_inline=False]
```

17.3 Profiling in Scripts - ProfileRegionTransform

The greatest flexibility is achieved by using the profiler transformation explicitly in a transformation script. The script takes either a single AST Node or a list of AST Nodes as argument, and will insert a Profile Node into the AST, with the specified nodes as children. At code creation time the listed children will all be enclosed in one profile region. As an example:

```python
from psyclone.transformations import ProfileRegionTrans

t=TransInfo()
p_trans= ProfileRegionTrans()
schedule=psy.invokes.get('invoke_0').schedule
schedule.view()

# Enclose all children within a single profile region
newschedule, _ = p_trans.apply(schedule.children[1:3])
newschedule.view()
```

> **Warning:** It is the responsibility of the user to make sure that a profile region is only created inside a multi-threaded region if the profiling library used is thread-safe!

17.4 Interface to Third Party Profiling Tools

PSyclone comes with wrapper libraries to support usage of Dr Hook, dl_timer and a simple non-thread-safe timing library. Support for further profiling libraries will be added in the future. To compile the wrapper libraries, change into the directory `lib/profiling` of PSyclone and type `make` to compile all wrappers. If only some of the wrapper are required, you can either use `make wrapper-name` (e.g. `make drhook`), or change into the corresponding directory and use `make`. The corresponding README files contain additional parameters that can be set in order to find third party profiling tools.

Any user can create similar wrapper libraries for other profiling tools by providing a corresponding Fortran module. The four profiling calls described in the section about the *ProfilingAPI* must be implemented, and an opaque, user-defined type `ProfileData` needs to be provided in the module.

Note that the `ProfileEnd` call does not have the module or region name as an argument. If this is required by the profiling library, this data must be stored in the `ProfileData` object so that it is available in the `ProfileEnd` call.

The examples in the lib/profiling directory show various ways in which the opaque data type can be used to interface with existing profiling tools - for example by storing an index used by the profiling tool in `ProfileData`, or by storing pointers to the profiling data to be able to print all results in a ProfileFinalise() subroutine.

PSY KERNEL EXTRACTOR (PSYKE)

18.1 Introduction

PSyclone has the ability to define regions of a PSyclone-conformant code to be extracted and run as a stand-alone application. This ability, called PSyKE (PSy Kernel Extractor), can be useful for benchmarking parts of a model, such as LFRic, without the need for using its infrastructure.

18.1.1 Mechanism

The code marked for extraction can be (subject to *Restrictions*):

- One or more Nodes in an Invoke (e.g. Loops containing Kernel or Built-In calls, a Directive enclosing one or more Loops) or

- The entire Invoke (extraction applied to all Nodes).

The basic mechanism of code extraction is through applying the **ExtractRegionTrans** transformation to selected Nodes. This transformation is further sub-classed into API-specific implementations, **DynamoExtractRegionTrans** and **GOceanExtractRegionTrans**. Both sub-classed transformations insert an instance of the **ExtractNode** object into the Schedule of a specific Invoke. All Nodes marked for extraction become children of the **ExtractNode**.

For now, the **ExtractNode** class simply adds comments at the beginning and the end of the extract region, that is at the position of the extract Node and after all its children. For example, marking a single Loop containing a Kernel call in Dynamo0.3 API generates:

```
! ExtractStart
! CALL write_extract_arguments(argument_list)
!
DO cell=1,f3_proxy%vspace%get_ncell()
  !
  CALL testkern_code_w2_only(nlayers, f3_proxy%data, f2_proxy%data, ...)
END DO
!
! ExtractEnd
```

The `! CALL write_extract_arguments(argument_list)` comment will be replaced by appropriate calls to write out arguments of extracted Node(s) or Kernel(s) in Issue #234. The **ExtractNode** base class can be sub-classed by API-specific code. This will be required for constructing the list of quantities required by Kernel (and Built-In) calls in the extracted Nodes.

18.1.2 Restrictions

Code extraction can be applied to unoptimised or optimised code. There are restrictions that check for correctness of optimising transformations when extraction is applied, as well as restrictions that eliminate dependence on the specific model infrastructure.

General

This group of restrictions is enforced irrespective of whether optimisations are used or not.

- Extraction can be applied to a single Node or a list of Nodes in a Schedule. For the latter, Nodes in the list must be consecutive children of the same parent Schedule.

- Extraction cannot be applied to an **ExtractNode** or a Node list that already contains one (otherwise we would have an extract region within another extract region).

- A Kernel or a Built-In call cannot be extracted without its parent Loop.

Distributed memory

As noted in the *Distributed Memory* section, support for distributed memory in PSyclone is currently limited to the Dynamo0.3 API. Since the implementation generates calls to LFRic infrastructure (e.g. runtime checks for status of field halos), code extraction is not allowed when distributed memory is enabled.

Shared memory and API-specific

The **ExtractRegionTrans** transformation cannot be applied to:

- A Loop without its parent Directive,

- An orphaned Directive (e.g. **OMPDoDirective**, **ACCLoopDirective**) without its parent Directive (e.g. ACC or OMP Parallel Directive),

- A Loop over cells in a colour without its parent Loop over colours in the Dynamo0.3 API,

- An inner Loop without its parent outer Loop in the GOcean1.0 API.

18.2 Use

The code extraction is currently enabled by utilising a transformation script (see *Script* section for more details).

For example, the transformation script which extracts the first Kernel call in Dynamo0.3 API test example `15.1.2_builtin_and_normal_kernel_invoke.f90` would be written as:

```python
from psyclone.transformations import DynamoExtractRegionTrans

# Get instance of the ExtractRegionTrans transformation
etrans = DynamoExtractRegionTrans()

# Get Invoke and its Schedule
invoke = psy.invokes.get("invoke_0")
schedule = invoke.schedule

# Apply extract transformation to the selected Node
```

(continues on next page)

```
schedule, _ = etrans.apply(schedule.children[2])
schedule.view()
```

and called as:

```
> psyclone -nodm -s ./extract_single_node.py \
    <path-to-example>/15.1.2_builtin_and_normal_kernel_invoke.f90
```

PSyclone modifies the Schedule of the selected invoke_0:

```
Schedule[invoke='invoke_0' dm=False]
    Loop[type='dofs',field_space='any_space_1',it_space='dofs', upper_bound='ndofs']
        Call setval_c(f5,0.0)
    Loop[type='dofs',field_space='any_space_1',it_space='dofs', upper_bound='ndofs']
        Call setval_c(f2,0.0)
    Loop[type='',field_space='w2',it_space='cells', upper_bound='ncells']
        KernCall testkern_code_w2_only(f3,f2) [module_inline=False]
    Loop[type='',field_space='wtheta',it_space='cells', upper_bound='ncells']
        KernCall testkern_wtheta_code(f4,f5) [module_inline=False]
    Loop[type='',field_space='w1',it_space='cells', upper_bound='ncells']
        KernCall testkern_code(scalar,f1,f2,f3,f4) [module_inline=False]
```

to insert the extract region. As shown below, all children of an **ExtractNode** will be part of the region:

```
Schedule[invoke='invoke_0' dm=False]
    Loop[type='dofs',field_space='any_space_1',it_space='dofs', upper_bound='ndofs']
        Call setval_c(f5,0.0)
    Loop[type='dofs',field_space='any_space_1',it_space='dofs', upper_bound='ndofs']
        Call setval_c(f2,0.0)
    Extract
        Loop[type='',field_space='w2',it_space='cells', upper_bound='ncells']
            KernCall testkern_code_w2_only(f3,f2) [module_inline=False]
    Loop[type='',field_space='wtheta',it_space='cells', upper_bound='ncells']
        KernCall testkern_wtheta_code(f4,f5) [module_inline=False]
    Loop[type='',field_space='w1',it_space='cells', upper_bound='ncells']
        KernCall testkern_code(scalar,f1,f2,f3,f4) [module_inline=False]
```

To extract multiple Nodes, **ExtractRegionTrans** can be applied to the list of Nodes (subject to *General* restrictions above):

```
# Apply extract transformation to the selected Nodes
schedule, _ = etrans.apply(schedule.children[1:3])
```

This modifies the above Schedule as:

```
...
    Extract
        Loop[type='dofs',field_space='any_space_1',it_space='dofs', upper_bound='ndofs
↪']
            Call setval_c(f2,0.0)
        Loop[type='',field_space='w2',it_space='cells', upper_bound='ncells']
            KernCall testkern_code_w2_only(f3,f2) [module_inline=False]
...
```

As said above, extraction can be performed on optimised code. For example, the following example transformation script first adds !$OMP PARALLEL DO directive and then extracts the optimised code in Dynamo0.3 API test example 15.1.2_builtin_and_normal_kernel_invoke.f90:

```python
from psyclone.transformations import DynamoOMPParallelLoopTrans, \
    DynamoExtractRegionTrans

# Get instances of the transformations
etrans = DynamoExtractRegionTrans()
otrans = DynamoOMPParallelLoopTrans()

# Get Invoke and its Schedule
invoke = psy.invokes.get("invoke_0")
schedule = invoke.schedule

# Add OMP PARALLEL DO directives
schedule, _ = otrans.apply(schedule.children[1])
schedule, _ = otrans.apply(schedule.children[2])
# Apply extract transformation to the selected Nodes
schedule, _ = etrans.apply(schedule.children[1:3])
schedule.view()
```

The generated code is now:

```fortran
! ExtractStart
! CALL write_extract_arguments(argument_list)
!
!$omp parallel do default(shared), private(df), schedule(static)
DO df=1,undf_any_space_1_f2
  f2_proxy%data(df) = 0.0
END DO
!$omp end parallel do
!$omp parallel do default(shared), private(cell), schedule(static)
DO cell=1,f3_proxy%vspace%get_ncell()
  !
  CALL testkern_code_w2_only(nlayers, f3_proxy%data, f2_proxy%data, ndf_w2, undf_w2,␣
→map_w2(:,cell))
END DO
!$omp end parallel do
!
! ExtractEnd
```

Examples in `examples/dynamo/eg12` directory demonstrate how to apply code extraction by utilising PSyclone transformation scripts (see *Examples* section for more information).

SYSTEM-SPECIFIC SET-UP

This chapter is split into two sections: *User set-up* describes the setup for a user of PSyclone. It includes all steps necessary to be able to use PSyclone. And while you could obviously do some development, none of the required tools for testing or documentation creation will be installed.

Developer set-up describes the additional installation of all required tools to run tests and create documentation.

Both sections have detailed instructions for Ubuntu 16.04.2 and OpenSUSE 42.2 - if you are working with a different Linux distribution some adjustments will be necessary.

19.1 User set-up

This section provides system-specific information on how to set-up your system to use PSyclone for users. It has been tested with a vanilla installation of Ubuntu 16.04.2 and OpenSUSE 42.2. You need a terminal window for entering the commands into.

19.1.1 Installing dependencies

Most required dependencies are installed from the Python Package Index (https://packaging.python.org/installing/) using the program pip ("PIP Installs Packages"). Besides `pip` it is also recommended to install the graphviz package to be able to visualise dependency graphs. This is optional and the associated routine will silently return if the graphviz bindings are not installed.

Installing dependencies on Ubuntu

On Ubuntu `pip` and `graphviz` are installed using `apt-get`. Remember that graphviz is optional and that you'll need to install the graphviz package in addition to the Python bindings.

```
> sudo apt-get install python-pip graphviz
```

Installing dependencies on OpenSUSE

The vanilla OpenSUSE installation includes `pip` for Python 3. Note that the graphviz package is installed by default.

```
> sudo zypper install python-pip
```

19.1.2 Installing PSyclone

Change your working directory to where you would like to place the code and download the latest stable release of PSyclone.

```
> cd <PSYCLONEHOME>
> wget https://github.com/stfc/PSyclone/archive/1.7.0.tar.gz
> gunzip 1.7.0.tar.gz
> tar xf 1.7.0.tar
> rm 1.7.0.tar
> cd PSyclone-1.7.0
> export PYTHONPATH=`pwd`/src:${PYTHONPATH}
> export PATH=`pwd`/bin:${PATH}
```

This sets up your python path and path appropriately. You may want to set these paths permanently (e.g. by editing your ${HOME}/.bashrc file if you run the BASH shell). You can also use the latest version using git, as described in *Installing PSyclone From Git*.

19.1.3 Common installation

To avoid warnings during the dependency installation, it is recommended to update `pip` to the latest version:

```
> sudo pip install --upgrade pip
```

Next you need to install the `fparser` and `pyparsing` packages:

```
> sudo pip install fparser pyparsing
```

Tip: With `pip` it is possible to install packages either system-wide (which requires root privileges) as above, or for a single user only (in ~/.local). While the latter is only useful for one particular user, it means that PSyclone can be installed using `pip` without needing root privileges. In order to install a package for a user, add the –user command line option to all pip commands. This flag requests that the packages be installed locally for the current user rather than requiring root access:

```
> pip install --user fparser pyparsing
```

You may remove the use of `sudo` and add the `--user` option to all pip commands described in this document.

Uninstalling is simply a matter of doing:

```
> sudo pip uninstall fparser pyparsing
```

PSyclone supports the ability to output a schedule dependency graph using the graphviz package. This is optional and the associated routine will silently return if the graphviz bindings are not installed. If you have the graphviz package installed (see especially section *Installing dependencies on Ubuntu* if you are on Ubuntu), you also need to install the python bindings to the graphviz package:

```
> sudo pip install graphviz
```

If you just want to use PSyclone then you've installed all you need and you are ready to go to the getting-going *Run* section.

19.2 Developer set-up

This section adds software that is used to develop and test PSyclone. It includes all packages for testing and creation of documentation in html and pdf. We assume you have already installed the software described in the *User set-up* section.

19.2.1 Installing PSyclone From Git

For development it is recommended to get a copy of PSyclone using git to get access to the latest development version.

Installing git for Ubuntu

You need to install the git package:

```
> sudo apt-get install git
```

Installing git on OpenSUSE

You need to install the git package:

```
>> sudo zypper --no-recommends install git
```

Cloning PSyclone using git

Cloning PSyclone from git and setting up your environment is done as follows:

```
> cd <PSYCLONEHOME>
> git clone https://github.com/stfc/PSyclone.git
> cd PSyclone
> pip install --user -e .
```

Note that the "-e" flag causes the project to be installed in 'editable' mode so that any changes to the PSyclone source take effect immediately.

> **Warning:** On OpenSUSE it is necessary to add *$HOME/.local/bin* to your *$PATH* if you have done a user-local install.

19.2.2 Installing documentation tools

Install Sphinx along with bibtex support for creating PSyclone documentation:

```
> sudo pip install sphinx sphinxcontrib.bibtex
```

You can now build html documentation:

```
> cd doc
> make html
```

The latex package is required to create the pdf documentation for PSyclone. Installing the full dependencies can take up several GB, the instructions for Ubuntu and OpenSUSE only install a minimal subset.

Installing LaTeX on Ubuntu

The following three packages need to be installed to create the pdf documentation. It is recommended to install the packages in one `apt-get` command, since otherwise, depending on your filesystem, unnecessary snapshots might be created that take up additional space. The `--no-install-recommends` option significantly reduces the number of installed packages:

```
> sudo apt-get install --no-install-recommends texlive \
texlive-latex-extra latexmk
```

Installing LaTeX on OpenSUSE

The following command installs the minimum number of packages in order to create the pdf documentation - around 130 packages all in all, requiring approximately 300 MB.

> **Warning:** It is important to install the packages in one `zypper` command, since otherwise, depending on your filesystem, unnecessary snapshots might be created after each package, which can add up to several GB of data.

```
> sudo zypper install --no-recommends texlive-latex texlive-latexmk \
texlive-babel-english texlive-cmap texlive-psnfss texlive-fncychap  \
texlive-fancyhdr texlive-titlesec texlive-tabulary texlive-varwidth \
texlive-framed texlive-fancyvrb texlive-float texlive-wrapfig       \
texlive-parskip texlive-upquote texlive-capt-of texlive-needspace   \
texlive-metafont texlive-makeindex texlive-times texlive-helvetic   \
texlive-courier texlive-dvips
```

Creating pdf documentation

You can now build the pdf documentation using

```
> cd doc
> make latexpdf
```

19.2.3 Installing testing tools

The following modules are recommended to get access to testing and formatting tools:

```
> sudo pip install pytest pep8 pylint==1.6.5 pytest-cov pytest-pep8 \
pytest-pylint pytest-flakes pytest-pep257
```

> **Warning:** It appears that the 1.7 series of `pylint` has a bug (at least up to 1.7.2) and does not work properly with PSyclone - it aborts with a "maximum recursion depth exceeded" error message. It is therefore recommended to use version 1.6.5, as specified in the above `pip` command.

You can now run the PSyclone python tests:

```
> cd PSyclone.git
> py.test
```

In order to see whether the Python code conforms to the pep8 standards, use:

```
> pep8 code.py
```

> **Warning:** From release 1.7.1 `pep8` returns a runtime user warning to install and use `pycodestyle` instead.

Verifying the pylint standards is done with:

```
> pylint code.py
```

OK, you're all set up.

CONFIGURATION

PSyclone reads various run-time configuration options from the `psyclone.cfg` file. As described in *Configuration*, the default `psyclone.cfg` configuration file is installed in `<python-base-prefix>/share/psyclone/` during the installation process. The original version of this file is in the `PSyclone/config` directory of the PSyclone distribution.

At execution-time, the user can specify a custom configuration file to be used. This can either be done with the `--config` command line option, or by specifying the (full path to the) configuration file to use via the `PSYCLONE_CONFIG` environment variable. If the specified configuration file is not found then PSyclone will fall back to searching in a list of default locations.

The ordering of these locations depends upon whether PSyclone is being run within a Python virtual environment (such as `venv`). If no virtual environment is detected then the locations searched, in order, are:

1. `${PWD}/.psyclone/`

2. `${HOME}/.local/share/psyclone/`

3. `<python-base-dir>/share/psyclone/`

where `<python-base-dir>` is the path stored in Python's `sys.prefix`.

If a virtual environment is detected then it is assumed that the `share` directory will be a part of that environment. In order to maintain isolation of distinct virtual environments this directory is then checked *before* the user's home directory. i.e. the list of locations searched is now:

1. `${PWD}/.psyclone/`

2. `<python-base-dir>/share/psyclone/`

3. `${HOME}/.local/share/psyclone/`

Note that for developers a slightly different configuration handling is implemented, see *Module: configuration* for details.

20.1 Options

The configuration file is read by the Python ConfigParser class (https://docs.python.org/3/library/configparser.html) and must be formatted accordingly. It currently consists of a `DEFAULT` section e.g.:

```
[DEFAULT]
API = dynamo0.3
DEFAULTSTUBAPI = dynamo0.3
DISTRIBUTED_MEMORY = true
REPRODUCIBLE_REDUCTIONS = false
REPROD_PAD_SIZE = 8
```

and a an optional API specific section, for example for `dynamo0.3` section:

```
[dynamo0.3]
access_mapping = gh_read:read, gh_write:write, gh_rw:readwrite,
                 gh_inc:inc, gh_sum:sum

COMPUTE_ANNEXED_DOFS = false
```

or for `gocean1.0`:

```
[gocean1.0]
access_mapping = go_read:read, go_write:write, go_readwrite:readwrite

iteration-spaces = offset_sw:ct:internal_we_halo:1:2:3:4
                   offset_sw:ct:internal_ns_halo:1:{stop}:1:{stop}+1
```

The meaning of the various entries is described in the following sub-sections.

Note that ConfigParser supports various forms of boolean entry including "true/false", "yes/no" and "1/0". See https://docs.python.org/3/library/configparser.html#supported-datatypes for more details.

20.1.1 `DEFAULT` Section

This section contains entries that are, in principle, applicable to all APIs supported by PSyclone.

Entry	Description
API	The API that PSyclone assumes an Algorithm/Kernl conforms to if no API is specified. Must be one of the APIs supported by PSyclone (dynamo0.1, dynamo0.3, gocean0.1, gocean1.0 and nemo). If there is no API specified and there is only one API-specific section in the config file loaded, this API will be used. This value can be overwritten by the command line option '-api'. If there is no API entry in the config file, and '-api' is not specified on the command line, dynamo0.3 is used as default.
DEFAULTSTUBAPI	The API that the kernel-stub generator assumes by default. Must be one of the stub-APIs supported by PSyclone (dynamo0.3 only at this stage).
DISTRIBUTED_MEMORY	Whether or not to generate code for distributed-memory parallelism by default. Note that this is currently only supported for the dynamo0.3 API.
REPRODUCIBLE_REDUCTIONS	Whether or not to generate code for reproducible OpenMP reductions (see *Reductions*) by default.
REPROD_PAD_SIZE	If generating code for reproducible OpenMP reductions, this setting controls the amount of padding used between elements of the array in which each thread accumulates its local reduction. (This prevents false sharing of cache lines by different threads.)

20.1.2 Common Sections

The following entries must be defined for each API in order for PSyclone to work as expected:

Entry	Description
access_mapping	This field defines the strings that are used by a particular API to indicate write, read, . . . access. Its value is a comma separated list of access-string:access pairs, e.g.: `gh_read:read, gh_write:write, gh_rw:readwrite, gh_inc:inc, gh_sum:sum` At this stage these 5 types are defined for read, write, read+write, increment and summation access by PSyclone. Sum is a form of reduction. The gocean APIs do not support increment or sum, so they only define three mappings for read, write, and readwrite.

20.1.3 `dynamo0.3` Section

This section contains configuration options that are only applicable when using the Dynamo 0.3 API.

Entry	Description
COMPUTE_ANNEXED_DOFS	Whether or not to perform redundant computation over annexed dofs in order to reduce the number of halo exchanges. See *Dof iterators* in the Developers' guide.

20.1.4 `gocean1.0` Section

This section contains configuration options that are only applicable when using the Gocean 1.0 API.

Entry	Description
iteration-spaces	This contains definitions of additional iteration spaces used by PSyclone. A detailed description can be found in the *Configuration* section of the GOcean1.0 chapter.

DEVELOPERS' GUIDE

21.1 Working With PSyclone from GitHub

A PSyclone developer will, by definition, be working with the GitHub PSyclone repository rather than installing a
released version from pypi (using e.g. `pip install psyclone`). This section describes the set-up necessary
when using PSyclone in this way. It also describes some of the development practises of the PSyclone project.

21.1.1 Installation

Although PSyclone releases always work with a released version of fparser, the same is not always true of other
versions (e.g. the HEAD of the master branch). For those versions of PSyclone requiring fparser functionality that is
not yet in a release, we use the git submodule feature such that the PSyclone repository always has a link to the correct
version of fparser. In order to obtain this version the PSyclone repository must be cloned with the `--recursive`
flag:

```
> git clone --recursive https://github.com/stfc/PSyclone.git
```

Alternatively, if you already have a local clone of the PSyclone github repository then doing:

```
> cd <PSYCLONEHOME>
> git submodule init
> git submodule update --init --recursive
```

will fetch the fparser submodule. Failure to do this will mean that for example `<PSYCLONEHOME>/external/`
`fparser` directory will be empty. The `--recursive` option is necessary because dl_esm_inf uses submodules as
well.

Note that after cloning the repository from GitHub, the local copy will be on the master branch. If you are working
with some other branch then this must be checked out by doing:

```
> cd <PSYCLONEHOME>
> git checkout <BRANCH_NAME>
```

Once the above steps have been performed, the `<PSYCLONEHOME>/external/fparser` directory will contain
the correct version of the fparser code. This can then be installed using `pip`:

```
> cd <PSYCLONEHOME>/external/fparser
> pip install --user .
```

Once you have the correct version of fparser installed you are ready to install PSyclone itself. Again, the simplest way
of doing this is to use `pip`:

```
> cd <PSYCLONEHOME>
> pip install -e --user .
```

where `-e` requests an 'editable' installation so that changes to the PSyclone source are immediately reflected in the installed package. (For alternatives to using pip please see the *Getting Going* section.)

21.1.2 Test Suite

The PSyclone test suite is integral to the development process and all new code must be covered (i.e. executed) by one or more tests. As described in *Getting Going*, the test suite is written for use with `pytest`.

Tests should be run from the `<PSYCLONEHOME>/src/psyclone` directory, from which all tests in subdirectories (e.g. `tests`, `core/tests`) will be automatically found and started. If only a subset of all tests need to be run, `pytest` can be invoked from the corresponding subdirectory or with that subdirectory as an argument.

Coverage

The easiest and most user-friendly way of checking the coverage of any new code is to use CodeCov (https://codecov.io/gh/stfc/PSyclone) which is integrated with GitHub. Coverage for Pull Requests is automatically reported and will appear as a comment on the Pull Request. This comment is then automatically updated whenever new code is pushed to the associated branch.

For checking test coverage on your local machine you will need to install the `cov` plugin (`pip install pytest-cov`). You can then request various types of coverage report when running the test suite. e.g. to ask for a terminal report of missed lines for the `dynamo0p3` module you would do:

```
> cd <PSYCLONEHOME>
> py.test --cov-report term-missing --cov psyclone.dynamo0p3
```

Note that you specify the python module name, and not the file name. This will produce output along the lines of:

```
----------- coverage: platform linux, python 3.5.4-final-0 -----------
Name                       Stmts   Miss  Cover   Missing
----------------------------------------------------------
src/psyclone/dynamo0p3.py   2540     23    99%   558, 593, 777, 2731, 2972, 3865,
↪4132-4133, 4135-4136, 4139-4140, 4143-4144, 4149-4151, 4255, 4270, 4488, 5026, 6540,
↪ 6658, 6768
```

showing the line numbers which are not covered. By using `--cov` more than once you can report on more than one file. You can also request only selected tests to be run by specifying the file names on the command line. Additionally html output can be created by adding the option `--cov-report html`:

```
> cd <PSYCLONEHOME>/src/psyclone/tests
> py.test --cov-report term-missing --cov-report html --cov psyclone.dynamo0p3 ./
↪dynamo0p3_basis_test.py ./parse_test.py
```

The html output can be viewed with a browser at `file:///.../tests/htmlcov/index.html` and it highlights all source lines in red that are not covered by at least one test.

Parallel execution

The size of the test suite is such that running all of it in serial can take many minutes, especially if you have requested a coverage report. It is therefore very helpful to run it in parallel and pytest provides support for this via the `xdist`

plugin (`pip install pytest-xdist`). Once you have this plugin, the test suite may be run in parallel simply by providing the number of cores to use via the `-n` flag:

```
> cd <PSYCLONEHOME>
> py.test -n 4
```

Running the test suite in parallel also changes the order in which tests are run which can reveal any problems resulting from tests not being sufficiently isolated from one another.

Gotchas

Note that pytest will not complain if two tests (within a module) have the same name - it will just silently ignore one of them! The best way of checking for this is to run pylint on any modified test modules. (This needs to be done anyway as one of the requirements of the *Code Review* is that all new code be pylint-clean.)

Compilation testing

The test suite provides support for testing that the code generated by PSyclone is valid Fortran. This is performed by writing the generated code to file and then invoking a Fortran compiler. This testing is not performed by default since it requires a Fortran compiler and significantly increases the time taken to run the test suite.

Note: Compilaton testing is currently only supported for the "dynamo0.3" and "gocean1.0" APIs.

The Gnu Fortran compiler (gfortran) is used by default. If you wish to use a different compiler and/or supply specific flags then these are specified by further command-line flags:

```
> py.test --compile --f90=ifort --f90flags="-O3"
```

If you want to test OpenCL code created by PSyclone, you must use the command line option –compileopencl (which can be used together with –compile, and –f90 and –f90flags), e.g.:

```
> py.test --compileopencl --f90=<opencl-compiler> --f90flags="<opencl-specific flags>"
```

Infrastructure libraries

Since the code generated by PSyclone makes calls to an infrastructure library, compilation tests must have access to compiler specific .mod files. For dynamo0.3 a stub implementation of the required functions from the LFRic infrastructure is included in `tests/test_files/dynamo0p3/infrastructure`. When compilation tests are requested, the stub files are automatically compiled to create the required .mod files.

For the gocean1.0 API a complete copy of the dl_esm_inf library is included as a submodule in `<PSYCLONEHOME>/external/dl_esm_inf`. Before running tests with compilation, make sure this submodule is up-to-date (see *Installation*). The test process will compile dl_esm_inf automatically, and all PSyclone gocean1.0 compilation tests will reference these files.

If you run the tests in parallel (see *Parallel execution* section) each process will compile its own version of the wrapper files and infrastructure library to avoid race conditions. This happens only once per process in each test session.

21.1.3 Continuous Integration

The PSyclone project uses Travis (https://travis-ci.org/stfc/PSyclone) for continuous integration. GitHub triggers Travis to execute the test suite whenever there is a push to the repository. The work performed by Travis is configured by the `.travis.yml` file in the root directory of the repository. Currently Travis is configured to run the test suite for both Python 2.7 and 3.6. It also runs all of the examples using the `check_examples` script in the `examples` directory.

By default, the Travis configuration uses `pip` to install the dependencies required by PSyclone before running the test suite. This works well when PSyclone only depends upon released versions of other packages. However, PSyclone relies heavily upon fparser which is also under development. Occasionally it may be that a given branch of PSyclone requires a version of fparser that is not yet released. As described in *Installation*, PSyclone has fparser as a git submodule. In order to configure Travis to use that version of fparser instead of a release, the `.travis.yml` file must be edited and the line executing the "install_optional.sh" script (in the `before_install` section) must be edited to pass in the "fparser_submodule" argument:

```
- ./bin/install_optional.sh fparser_submodule
```

Note that this functionality is only for development purposes. Any release of PSyclone must work with a released version of fparser and therefore the "fparser_submodule" argument must be removed before making a release.

Given that a run of the test-suite on Travis uses approximately 45 minutes of CPU time, it is good practise to avoid triggering it unnecessarily (e.g. if you know that a certain commit won't pass). This can be achieved by appending "[skip ci]" (without the quotes) to the end of the associated git commit message.

21.1.4 Code Review

Before a branch can be merged to master it must pass code review. The guidelines for performing a review (i.e. what is expected from the developer) are available on the GitHub PSyclone wiki pages: https://github.com/stfc/PSyclone/wiki.

21.2 The PSyclone Internal Representation (PSyIR)

The PSyclone Internal Representation (PSyIR) is a language independent AST that PSyclone uses to represent the PSy layer and the kernel code. It can be constructed from scratch or produced from existing code using one of the ASTProcessors provided in PSyclone.

All nodes in the AST are sub-classes of the abstract *Node* base class, which provides the following common interface:

class psyclone.psyGen.**Node**(*ast=None*, *children=None*, *parent=None*)
　　Base class for a node in the PSyIR (schedule).

　　　　Parameters

　　　　　　• **ast** (sub-class of `fparser.two.Fortran2003.Base`) – reference into the fparser2 AST corresponding to this node.

　　　　　　• **children** (list of *psyclone.psyGen.Node*) – the PSyIR nodes that are children of this node.

　　　　　　• **parent** (*psyclone.psyGen.Node*) – that parent of this node in the PSyIR tree.

　　abs_position
　　　　Find a Node's absolute position in the tree (starting with 0 if it is the root). Needs to be computed dynamically from the starting position (0) as its position may change. :returns: absolute position of a Node in the tree :raises InternalError: if the absolute position cannot be found :rtype: int

ancestor(*my_type*, *excluding=None*)
Search back up tree and check whether we have an ancestor that is an instance of the supplied type. If we do then we return it otherwise we return None. A list of (sub-) classes to ignore may be provided via the *excluding* argument.

> **Parameters**
>
> - **my_type** (`type`) – Class to search for.
> - **excluding** (`list`) – list of (sub-)classes to ignore or None.
>
> **Returns** First ancestor Node that is an instance of the requested class or None if not found.

annotations
Return the list of annotations attached to this Node.

> **Returns** List of anotations
>
> **Return type** list of str

args
Return the list of arguments associated with this Node. The default implementation assumes the Node has no directly associated arguments (i.e. is not a Call class or subclass). Arguments of any of this nodes descendants are considered to be associated.

ast

> **Returns** a reference to that part of the fparser2 parse tree that this node represents or None.
>
> **Return type** sub-class of `fparser.two.utils.Base`

ast_end

> **Returns** a reference to the last node in the fparser2 parse tree that represents a child of this PSyIR node or None.
>
> **Return type** sub-class of `fparser.two.utils.Base`

backward_dependence()
Returns the closest preceding Node that this Node has a direct dependence with or None if there is not one. Only Nodes with the same parent as self are returned. Nodes inherit their descendants' dependencies. The reason for this is that for correctness a node must maintain its parent if it is moved. For example a halo exchange and a kernel call may have a dependence between them but it is the loop body containing the kernel call that the halo exchange must not move beyond i.e. the loop body inherits the dependencies of the routines within it.

calls()
Return all calls that are descendants of this node.

dag(*file_name='dag'*, *file_format='svg'*)
Create a dag of this node and its children.

dag_gen(*graph*)
Output my node's graph (dag) information and call any children. Nodes with children are represented as two vertices, a start and an end. Forward dependencies are represented as green edges, backward dependencies are represented as red edges (but their direction is reversed so the layout looks reasonable) and parent child dependencies are represented as blue edges.

dag_name
Return the base dag name for this node.

depth
Returns this Node's depth in the tree: 1 for the Schedule and increasing for its descendants at each level. :returns: depth of the Node in the tree :rtype: int

---

following()
> Return all *psyclone.psyGen.Node* nodes after me in the schedule. Ordering is depth first.
>
> > **Returns** a list of nodes
> >
> > **Return type** list() of *psyclone.psyGen.Node*

following_calls
> Return all calls after me in the schedule.

forward_dependence()
> Returns the closest following Node that this Node has a direct dependence with or None if there is not one. Only Nodes with the same parent as self are returned. Nodes inherit their descendants' dependencies. The reason for this is that for correctness a node must maintain its parent if it is moved. For example a halo exchange and a kernel call may have a dependence between them but it is the loop body containing the kernel call that the halo exchange must not move beyond i.e. the loop body inherits the dependencies of the routines within it.

gen_c_code(*indent=0*)
> Abstract method for the generation of C source code
>
> > **Parameters** **indent** (*int*) – Depth of indent for the output string.
> >
> > **Raises** **NotImplementedError** – is an abstract method.

static indent(*count*, *indent=' '*)
> Helper function to produce indentation strings.
>
> > **Parameters**
> >
> > > * **count** (*int*) – Number of indentation levels.
> > > * **indent** (*str*) – String representing one indentation level.
> >
> > **Returns** Complete indentation string.
> >
> > **Return type** str

is_openmp_parallel()
> Returns true if this Node is within an OpenMP parallel region.

is_valid_location(*new_node*, *position='before'*)
> If this Node can be moved to the new_node (where position determines whether it is before of after the new_node) without breaking any data dependencies then return True, otherwise return False.

kern_calls()
> Return all user-supplied kernel calls in this schedule.

loops()
> Return all loops currently in this schedule.

position
> Find a Node's position relative to its parent Node (starting with 0 if it does not have a parent). :returns: relative position of a Node to its parent :rtype: int

preceding(*reverse=None*)
> Return all *psyclone.psyGen.Node* nodes before me in the schedule. Ordering is depth first. If the *reverse* argument is set to *True* then the node ordering is reversed i.e. returning the nodes closest to me first
>
> > **Param** reverse: An optional, default *False*, boolean flag
> >
> > **Type** reverse: bool
> >
> > **Returns** A list of nodes

> **Return type** `list()` of *[psyclone.psyGen.Node](#)*

preceding_calls
> Return all calls before me in the schedule.

reductions(*reprod=None*)
> Return all calls that have reductions and are decendents of this node. If reprod is not provided, all reductions are returned. If reprod is False, all builtin reductions that are not set to reproducible are returned. If reprod is True, all builtins that are set to reproducible are returned.

update()
> By default we assume there is no need to update the existing fparser2 AST which this Node represents. We simply call the update() method of any children.

walk(*children*, *my_type*)
> Recurse through tree and return objects of 'my_type'.

21.2.1 Schedule

The Schedule node represents a sequence of statments. It is a important node in PSyclone because two of its specialisations: InvokeSchedule and KernelSchedule (described below), are used as the root nodes of PSy-layer invokes and kernel subroutines. This makes them the starting points for any walking of the PSyIR tree in PSyclone transformation scripts and a common target for the application of transformations.

class `psyclone.psyGen.`**Schedule**(*sequence=None*, *parent=None*)
> Stores schedule information for a sequence of statements.

> > **Parameters**
> >
> > - **sequence** (list of *[psyclone.psyGen.Node](#)*) – the sequence of PSyIR nodes that make up the schedule.
> >
> > - **parent** (*[psyclone.psyGen.Node](#)*) – that parent of this node in the PSyIR tree.

coloured_text
> Returns the name of this node with appropriate control codes to generate coloured output in a terminal that supports it.

> > **Returns** Text containing the name of this node, possibly coloured.

> > **Return type** str

dag_name

> > **Returns** The name of this node in the dag.

> > **Return type** str

view(*indent=0*)
> Print a text representation of this node to stdout and then call the view() method of any children.

> > **Parameters** **indent** (*[int](#)*) – Depth of indent for output text.

InvokeSchedule

The InvokeSchedule is a PSyIR node that represents an invoke subroutine in the PSy-layer. It extends the *psyclone.psyGen.Schedule* functionality with a *psyclone.psyGen.NameSpace* and a reference to its associated *psyclone.psyGen.Invoke* object.

class `psyclone.psyGen.`**InvokeSchedule**(*KernFactory*, *BuiltInFactory*, *alg_calls=None*)
> Stores schedule information for an invocation call. Schedules can be optimised using transformations.

```
>>> from psyclone.parse.algorithm import parse
>>> ast, info = parse("algorithm.f90")
>>> from psyclone.psyGen import PSyFactory
>>> api = "..."
>>> psy = PSyFactory(api).create(info)
>>> invokes = psy.invokes
>>> invokes.names
>>> invoke = invokes.get("name")
>>> schedule = invoke.schedule
>>> schedule.view()
```

Parameters

- **KernFactory** (*type*) – class instance of the factory to use when creating Kernels. e.g. `psyclone.dynamo0p3.DynKernCallFactory`.

- **BuiltInFactory** (*type*) – class instance of the factory to use when creating built-ins. e.g. `psyclone.dynamo0p3_builtins.DynBuiltInCallFactory`.

- **alg_calls** (list of `psyclone.parse.algorithm.KernelCall`) – list of Kernel calls in the schedule.

coloured_text
Returns the name of this node with appropriate control codes to generate coloured output in a terminal that supports it.

> **Returns** Text containing the name of this node, possibly coloured
>
> **Return type** string

gen_code (*parent*)
Generate the Nodes in the f2pygen AST for this schedule.

> **Parameters parent** (`psyclone.f2pygen.SubroutineGen`) – the parent Node (i.e. the enclosing subroutine) to which to add content.

opencl

> **Returns** Whether or not we are generating OpenCL for this InvokeSchedule.
>
> **Return type** bool

view (*indent=0*)
Print a text representation of this node to stdout and then call the view() method of any children.

> **Parameters indent** (*integer*) – Depth of indent for output text

KernelSchedule

The KernelSchedule is a PSyIR node that represents a kernel subroutine. It extends the *psyclone.psyGen.Schedule* functionality with a Symbol Table (*psyclone.psyGen.SymbolTable*) that keeps a record of the Symbols (*psyclone.psyGen.Symbol*) used in the kernel scope. A Symbol is defined as:

class `psyclone.psyGen.`**Symbol** (*name*, *datatype*, *shape=[]*, *scope='local'*, *is_input=False*, *is_output=False*)
Symbol item for the Symbol Table. It contains information about: the name, the datatype, the shape (in column-major order), the scope and for global-scoped symbols whether the data is already defined and/or survives after the kernel.

Parameters

- **name** (`str`) – Name of the symbol.

- **datatype** (`str`) – Data type of the symbol.

- **shape** (`list`) – Shape of the symbol in column-major order (leftmost index is contiguous in memory). Each entry represents an array dimension. If it is 'None' the extent of that dimension is unknown, otherwise it holds an integer literal or a reference to an integer symbol with the extent. If it is an empty list then the symbol represents a scalar.

- **scope** (`str`) – It is 'local' if the symbol just exists inside the kernel scope or 'global_*' if the data survives outside of the kernel scope. Note that global-scoped symbols also have postfixed information about the sharing mechanism, at the moment just 'global_argument' is available for variables passed in/out of the kernel by argument.

- **is_input** (`bool`) – Whether the symbol represents data that exists before the kernel is entered and that is passed into the kernel.

- **is_output** (`bool`) – Whether the symbol represents data that is passed outside the kernel upon exit.

Raises

- **NotImplementedError** – Provided parameters are not supported yet.

- **TypeError** – Provided parameters have invalid error type.

- **ValueError** – Provided parameters contain invalid values.

The Symbol Table has the following interface:

class psyclone.psyGen.**SymbolTable**(*kernel=None*)

Encapsulates the symbol table and provides methods to declare new symbols and look up existing symbols. It is implemented as a single scope symbol table (nested scopes not supported).

> **Parameters kernel** (`psyclone.psyGen.KernelSchedule` or NoneType) – Reference to the KernelSchedule to which this symbol table belongs.

argument_list

> **Returns** Ordered list of arguments.

> **Return type** list of *psyclone.psyGen.Symbol*

declare(*name*, *datatype*, *shape=[]*, *scope='local'*, *is_input=False*, *is_output=False*)

Declare a new symbol in the symbol table.

Parameters

- **name** (`str`) – Name of the symbol.

- **datatype** (`str`) – Datatype of the symbol.

- **shape** (`list`) – Shape of the symbol in column-major order (leftmost index is contiguous in memory). Each entry represents an array dimension. If it is 'None' the extent of that dimension is unknown, otherwise it holds an integer literal or a reference to an integer symbol with the extent. If it is an empty list then the symbol represents a scalar.

- **scope** (`str`) – It is 'local' if the symbol just exists inside the kernel scope or 'global_*' if the data survives outside of the kernel scope. Note that global-scoped symbols also have postfixed information about the sharing mechanism, at the moment just 'global_argument' is available for variables passed in/out of the kernel by argument.

- **is_input** (`bool`) – Whether the symbol represents data that exists before the kernel is entered and that is passed into the kernel.

- **is_output** (`bool`) – Whether the symbol represents data that is survives outside the kernel upon exit.

> **Raises** `KeyError` – The provided name can not be used as key in the table.

gen_c_local_variables(*indent=0*)
> Generate C code that defines all local symbols in the Symbol Table.

> > **Parameters** **indent** (`int`) – Indentation level

> > **Returns** C languague definition of the local symbols.

> > **Return type** str

gen_ocl_argument_list(*indent=0*)
> Generate OpenCL argument list.

> > **Raises** `NotImplementedError` – is an abstract method.

gen_ocl_array_length(*indent=0*)
> Generate OpenCL array length variable declarations.

> > **Raises** `NotImplementedError` – is an abstract method.

gen_ocl_iteration_indices(*indent=0*)
> Generate OpenCL iteration indices declaration.

> > **Raises** `NotImplementedError` – is an abstract method.

local_symbols

> > **Returns** List of local symbols.

> > **Return type** list of `psyclone.psyGen.Symbol`

lookup(*name*)
> Look up a symbol in the symbol table.

> > **Parameters** **name** (`str`) – Name of the symbol

> > **Raises** `KeyError` – If the given name is not in the Symbol Table.

specify_argument_list(*argument_name_list*)
> Keep track of the order of the arguments and provide the scope, is_input and is_ouput information if it was not available on the variable declaration.

> > **Parameters** **argument_name_list** (`list`) – Ordered list of the argument names.

view()
> Print a representation of this Symbol Table to stdout.

21.2.2 Control Flow Nodes

The PSyIR has two control flow nodes: IfBlock and Loop. These nodes represent the canonical structure with which conditional branching constructs and iteration constructs are built. Additional language-specific syntax for branching and iteration will be normalized to use these same constructs. For example, Fortran has the additional branching constructs *ELSE IF* and *CASE*: when a Fortran code is translated into the PSyIR, Psyclone will build a semantically equivalent implementation using IfBlocks. However, the necessary nodes in the new tree structure will be annotated with information to enable the original language-specific syntax to be recreated if required.

class psyclone.psyGen.**IfBlock**(*parent=None*, *annotation=None*)
> Class representing an if-block within the PSyIR. It has two mandatory children: the first one represents the if-condition and the second one the if-body; and an optional third child representing the else-body.

Parameters

- **parent** (`psyclone.psyGen.Node`) – the parent of this node within the PSyIR tree.

- **annotation** (`str`) – Tags that provide additional information about the node. The node should still be functionally correct when ignoring these tags. Currently, it includes: 'was_elseif' to tag nested ifs originally written with the 'else if' languague syntactic constructs, 'was_single_stmt' to tag ifs with a 1-statement body which were originally written in a single line, and 'was_case' to tag an conditional structure which was originally written with the Fortran 'case' or C 'switch' syntactic constructs.

Raises **InternalError** – when initialised with invalid parameters.

coloured_text

Return text containing the (coloured) name of this node type.

Returns the name of this node type, possibly with control codes for colour.

Return type str

condition

Return the PSyIR Node representing the conditional expression of this IfBlock.

Returns IfBlock conditional expression.

Return type `psyclone.psyGen.Node`

Raises **InternalError** – If the IfBlock node does not have the correct number of children.

else_body

Return children of the Schedule executed when the IfBlock evaluates to False.

Returns Statements to be executed when IfBlock evaluates to False.

Return type list of `psyclone.psyGen.Node`

gen_c_code(*indent=0*)

Generate a string representation of this node using C language.

Parameters **indent** (`int`) – Depth of indent for the output string.

Returns C language code representing the node.

Return type str

Raises **InternalError** – If any mandatory children of the IfBlock node are missing.

if_body

Return children of the Schedule executed when the IfBlock evaluates to True.

Returns Statements to be executed when IfBlock evaluates to True.

Return type list of `psyclone.psyGen.Node`

Raises **InternalError** – If the IfBlock node does not have the correct number of children.

view(*indent=0*)

Print representation of this node to stdout.

Parameters **indent** (`int`) – the level to which to indent the output.

class psyclone.psyGen.**Loop**(*parent=None*, *variable_name=''*, *valid_loop_types=None*)

Represents a loop in the PSyIR.

Parameters

- **parent** (sub-class of `psyclone.psyGen.Node`) – Parent of this node in the PSyIR.

- **variable_name** (`str`) – Optional name of the loop iterator variable. Defaults to an empty string.

- **valid_loop_types** (`list of str`) – A list of loop types that are specific to a particular API.

args_filter(*arg_types=None*, *arg_accesses=None*, *unique=False*)
> Return all arguments of type arg_types and arg_accesses. If these are not set then return all arguments. If unique is set to True then only return uniquely named arguments

coloured_text
> Returns a string containing the name of this node along with control characters for colouring in terminals that support it.

> > **Returns** The name of this node, possibly with control codes for colouring

> > **Return type** string

dag_name
> Return the name to use in a dag for this node

> > **Returns** Return the dag name for this loop

> > **Return type** string

gen_code(*parent*)
> Generate the Fortran Loop and any associated code.

> > **Parameters** **parent** (`psyclone.f2pygen.SubroutineGen`) – the node in the f2pygen AST to which to add content.

has_inc_arg()
> Returns True if any of the Kernels called within this loop have an argument with INC access. Returns False otherwise

kernel
> > **Returns** the kernel object associated with this Loop (if any).

> > **Return type** `psyclone.psyGen.Kern`

unique_modified_args(*arg_type*)
> Return all unique arguments of type arg_type from Kernels in this loop that are modified. :param str arg_type: the type of kernel argument (e.g. field, operator) to search for. :returns: all unique arguments of type arg_type from Kernels in this loop that are modified. :rtype: List of `psyclone.psyGen.DynKernelArgument`.

variable_name
> > **Returns** the name of the control variable for this loop.

> > **Return type** str

view(*indent=0*)
> Write out a textual summary of this Loop node to stdout and then call the view() method of any children.

> > **Parameters** **indent** (`integer`) – Depth of indent for output text

21.2.3 Dependence Analysis

Dependence Analysis in PSyclone produces ordering constraints between instances of the *Argument* class within a PSyIR.

The *Argument* class is used to specify the data being passed into and out of instances of the *Call* class, *HaloExchange* Class and *GlobalSum* class (and their subclasses).

As an illustration consider the following invoke:

```
invoke(              &
    kernel1(a,b),  &
    kernel2(b,c))
```

where the metadata for *kernel1* specifies that the 2nd argument is written to and the metadata for *kernel2* specifies that the 1st argument is read.

In this case the PSyclone dependence analysis will determine that there is a flow dependence between the second argument of *Kernel1* and the first argument of *Kernel2* (a read after a write).

Information about arguments is aggregated to the PSyIR node level (*kernel1* and *kernel2* in this case) and then on to the parent *loop* node resulting in a flow dependence (a read after a write) between a loop containing *kernel1* and a loop containing *kernel2*. This dependence is used to ensure that a transformation is not able to move one loop before or after the other in the PSyIR schedule (as this would cause incorrect results).

Dependence analysis is implemented in PSyclone to support functionality such as adding and removing halo exchanges, parallelisation and moving nodes in a PSyIR schedule. Dependencies between nodes in a PSyIR schedule can be viewed as a DAG using the *dag()* method within the *Node* base class.

DataAccess Class

The *DataAccess* class is at the core of PSyclone data dependence analysis. It takes an instance of the *Argument* class on initialisation and provides methods to compare this instance with other instances of the *Argument* class. The class is used to determine 2 main things, called *overlap* and *covered*.

Overlap

Overlap specifies whether accesses specified by two instances of the *Argument* class access the same data or not. If they do access the same data their accesses are deemed to *overlap*. The best way to explain the meaning of *overlap* is with an example:

Consider a one dimensional array called *A* of size 4 (*A(4)*). If one instance of the *Argument* class accessed the first two elements of array *A* and another instance of the *Argument* class accessed the last two elements of array *A* then they would both be accessing array *A* but their accesses would *not overlap*. However, if one instance of the *Argument* class accessed the first three elements of array *A* and another instance of the *Argument* class accessed the last two elements of array *A* then their accesses would *overlap* as they are both accessing element *A(3)*.

Having explained the idea of *overlap* in its general sense, in practice PSyclone currently assumes that *any* two instances of the *Argument* class that access data with the same name will always *overlap* and does no further analysis (apart from halo exchanges and vectors, which are discussed below). The reason for this is that nearly all accesses to data, associated with an instance of the *Argument* class, start at index 1 and end at the number of elements, dofs or some halo depth. The exceptions to this are halo exchanges, which only access the halo and boundary conditions, which only access a subset of the data. However these subset accesses are currently not captured in metadata so PSyclone must assume subset accesses do not exist.

If there is a field vector associated with an instance of an *Argument* class then all of the data in its vector indices are assumed to be accessed when the argument is part of a *Call* or a *GlobalSum*. However, in contrast, a *HaloExchange* only acts on a single index of a field vector. Therefore there is one halo exchange per field vector index. For example:

```
InvokeSchedule[invoke='invoke_0_testkern_stencil_vector_type', dm=True]
... HaloExchange[field='f1', type='region', depth=1, check_dirty=True]
... HaloExchange[field='f1', type='region', depth=1, check_dirty=True]
... HaloExchange[field='f1', type='region', depth=1, check_dirty=True]
... Loop[type='',field_space='w0',it_space='cells', upper_bound='cell_halo(1)']
... ... KernCall testkern_stencil_vector_code(f1,f2) [module_inline=False]
```

In the above PSyIR schedule, the field *f1* is a vector field and the *Call testkern_stencil_vector_code* is assumed to access data in all of the vector components. However, there is a separate *HaloExchange* for each component. This means that halo exchanges accessing the same field but different components do not *overlap*, but each halo exchange does overlap with the loop node. The current implementation of the *overlaps()* method deals with field vectors correctly.

Coverage

The concept of *coverage* naturally follows from the discussion in the previous section.

Again consider a one dimensional array called *A* of size 4 (*A(4)*). If one instance (that we will call the *source*) of the *Argument* class accessed the first 3 elements of array *A* (i.e. elements 1 to 3) and another instance of the *Argument* class accessed the first two elements of array *A* then their accesses would *overlap* as they are both accessing elements *A(1) and A(2)* and elements *A(1) and A(2)* would be *covered*. However, access *A(3)* for the *source Argument* class would not yet be *covered*. If a subsequent instance of the *Argument* class accessed the 2nd and 3rd elements of array *A* then all of the accesses (*A(1), A(2) and A(3)*)) would now be *covered* so the *source argument* would be deemed to be covered.

In PSyclone the above situation occurs when a vector field is accessed in a kernel and also requires halo exchanges e.g.:

```
InvokeSchedule[invoke='invoke_0_testkern_stencil_vector_type', dm=True]
    HaloExchange[field='f1', type='region', depth=1, check_dirty=True]
    HaloExchange[field='f1', type='region', depth=1, check_dirty=True]
    HaloExchange[field='f1', type='region', depth=1, check_dirty=True]
    Loop[type='',field_space='w0',it_space='cells', upper_bound='cell_halo(1)']
        KernCall testkern_stencil_vector_code(f1,f2) [module_inline=False]
```

In this case the PSyIR loop node needs to know about all 3 halo exchanges before its access is fully *covered*. This functionality is implemented by passing instances of the *Argument* class to the *DataAccess* class *update_coverage()* method and testing the access.covered property until it returns *True*.

```
# this example is for a field vector 'f1' of size 3
# f1_index[1,2,3] are halo exchange accesses to vector indices [1,2,3] respectively
access = DataAccess(f1_loop)
access.update_coverage(f1_index1)
result = access.covered  # will be False
access.update_coverage(f1_index2)
result = access.covered  # will be False
access.update_coverage(f1_index3)
result = access.covered  # will be True
access.reset_coverage()
```

Note the *reset_coverage()* method can be used to reset internal state so the instance can be re-used (but this is not used by PSyclone at the moment).

The way in which halo exchanges are placed means that it is not possible for two halo exchange with the same index to depend on each other in a schedule. As a result an exception is raised if this situation is found.

Notice there is no concept of read or write dependencies here. Read or write dependencies are handled by classes that make use of the DataAccess class i.e. the *_field_write_arguments()* and *_field_read_arguments()* methods, both

of which are found in the *Arguments* class.

21.3 Parsing Code

The PSyclone *parse* module is responsible for parsing science (algorithm and kernel) code and extracting the required information for the algorithm translation and PSy generation phases.

The *parse* module contains modules for parsing algorithm (*algorithm.py*) and kernel (*kernel.py*) code as well as a utility module (*utils.py*) for common functionality.

21.3.1 Parsing Algorithm Code

The first thing PSyclone typically does is parse an input file. This input file is expected to contain Fortran source code conforming to the particular API in question. For the *nemo* API this is standard code, for the other API's this is algorithm code (conforming to the PSyKAl separation of concerns). The PSyclone code to do this is found in *parse/algorithm.py*.

An input file can be parsed via the *parse* function or via an instance of the *Parser* class. In practice the *parse* function simply calls the *Parser* class so we will concentrate on the latter in this section. The *parse* function could be removed from PSyclone but it is simple and is used in existing PSyclone scripts and examples.

The *Parser* class is initialised with a number of optional arguments. A particular *api* can be specified (this is required so the parser knows what sort of code and metadata to expect, how to parse it and which *builtins* are supported). The name used to specify an invoke (defaulting to *invoke*) can be changed, a path to where to look for associated kernel files can be provided and a particular maximum line length can be specified.

class psyclone.parse.algorithm.**Parser**(*api=''*,   *invoke_name='invoke'*,   *kernel_path=''*,
                                    line_length=False)
   Supports the parsing of PSyclone conformant algorithm code within a file and extraction of relevant information for any 'invoke' calls contained within the code.

   Parameters

   - **api** (*str*) – The PSyclone API to use when parsing the code.

   - **invoke_name** (*str*) – The expected name of the invocation calls in the algorithm code.

   - **kernel_path** (*str*) – The path to search for kernel source files (if different from the location of the algorithm source).

   - **line_length** (*bool*) – A logical flag specifying whether we care about line lengths being longer than 132 characters. If so, the input (algorithm and kernel) code is checked to make sure that it conforms and an error raised if not. The default is False.

   For example:

   ```
   >>> from psyclone.parse.algorithm import Parser
   >>> parser = Parser()
   >>> ast, info = parser.parse("alg.f90")
   ```

Once an instance of the *Parser()* class is created and configured with required values for the optional arguments, then the parse method can be called. This takes the name of the input code as its argument and returns a parse tree of the input code and a FileInfo object that captures the required invoke information found in the input code and in the associated kernel codes.

If the *nemo* API is specified then the *parse* method of the *Parser* instance simply parses the code with *fparser2* and returns the resultant *fparser2* parse tree.

For all other APIs the *parse* method of the *Parser* instance returns the resultant *fparser2* parse tree and a *FileInfo* instance which captures the invoke and kernel metadata in a hierarchy of classes.

When the *Parser* instance parses the code it expects to find Fortran code containing a program, module, subroutine or function (and it aborts if not). Currently the first of these (there may be more than one subroutine for example) is assumed to be the one that is required. This limititation is captured in issue #307.

The native *fparser2* tree of the Fortran code is then walked and all use statement names are captured and stored in a map (called *_arg_name_to_module_name*). This map allows the module name to be found given a particular argument name. Also, if an *invoke* call is found then an *InvokeCall* object is created and added to a list of such instances. Once all invokes have been found the *FileInfo* instance is created.

Note: In the future we want to be able to simply replace one parse tree with another. So how should we do this? One option would be to try to minimise the parser-specific parts and create some form of interface. The question is really what level of interface we should use.

An *InvokeCall* object is created in the *create_invoke_call* method by first parsing each of the kernels specified in an invoke call and creating a list of *kernel* objects which are then used to create an *InvokeCall*.

A *kernel* object is created in the *create_kernel_call* method which extracts the kernel name and kernel arguments, then creates either a *BuiltInCall* instance (via the *create_builtin_kernel_call* method) or a *KernelCall* instance (via the *create_coded_kernel_call* method). *BuiltInCalls* are created if the kernel name is the same as one of those specified in the builtin names for this particular API (see the variable *_builtin_name_map* which is initialised by the *get_builtin_defs* function.

The *create_kernel_call* method uses the *get_kernel* function to find out the kernel name and create a list of *Arg* instances representing the arguments. The *get_kernel* function parses each kernel argument using the fparser2 AST and determines the required argument information. An advantage of *fparser2* when compared with *fparser1* is that it parses all of a code, so we can use the parse tree to determine the type of each kernel argument appearing in the *invoke* call (e.g. scalar variable, array reference, literal constant) and create the appropriate *Arg* instance. Previously we relied on the *expression* module to do this (which has limitations).

Note: the analysis in the *get_kernel* function is the place to extend if we were to support arithmetic operations in an invoke call.

21.3.2 Parsing Kernel Code

A *BuiltInCall* instance is created by being passed a *BuiltinKernelType* instance for the particular API via the *BuiltInKernelTypeFactory* class which is found in the *kernels.py* file. This class parses the Fortran module file which specifies builtin description metadata. Currently *fparser1* is used but we will be migrating to *fparser2* in the future. The builtin metadata is specified in the same form as coded kernel metadata so the same logic can be used (i.e. the *KernelTypeFactory create* method is called) which is why *BuiltInKernelTypeFactory* subclasses *KernelTypeFactory*.

A *KernelCall* instance is created by being passed the module name of the kernel and a *KernelType* instance for the particular API via the *KernelTypeFactory* class which is also found in the *kernels.py* file. This class is given the parsed kernel module (via the *get_kernel_ast* function - which searches for the kernel file using the kernel path information introduced earlier). Again, currently *fparser1* is used but we will be migrating to *fparser2* in the future.

The *KernelTypeFactory create* method is used for both coded kernels and builtin kernels to specify the API-specific class to use. As an example, in the case of the *dynamo0.3* API the class is *DynKernMetadata* which is found in *psyclone.dynamo0p3*. Once this instance has been created (by passing it an *fparser1* AST) it can return information about the metadata contained therein. Moving from *fparser1* to *fparser2* would required changing the parse code logic in each of the API-specific classes.

21.4 Generic Code

PSyclone is designed to be configurable so that new front-ends (called APIs) can be built, re-using as much existing code as possible. The generic code is kept in the *psyGen.py* file for psy-code generation.

21.5 New APIs

TBD

21.6 Existing APIs

21.6.1 Dynamo0.3

Mesh

The Dynamo0.3 API supports meshes that are unstructured in the horizontal and structured in the vertical. This is often thought of as a horizontal 2D unstructured mesh which is extruded into the vertical. The LFRic infrastructure represents this mesh as a list of 2D cells with a scalar value capturing the number of levels in the vertical "column".

Cells

The Dynamo0.3 API currently assumes that all kernels which support iterating over cells work internally on a column of cells. This means that PSyclone need only be concerned with iterating over cell-columns in the horizontal. As a result the LFRic infrastructure presents the mesh information to PSyclone as if the mesh were 2-dimensional. From now on this 2D view will be assumed i.e. a cell will actually be a column of cells. The LFRic infrastructure provides a global 2D cell index from 1 to the number of cells.

For example, a simple quadrilateral element mesh with 4 cells might be indexed in the following way.

When the distributed memory option is switched on in the Dynamo0.3 API (see the *Distributed Memory* Section) the cells in the model are partitioned amongst processors and halo cells are added at the boundaries to a depth determined by the LFRic infrastructure. In this case the LFRic infrastructure maintains the global cell index and adds a unique local cell index from 1 to the number of cells in each partition, including any halo cells.

An example for a depth-1 halo implementation with the earlier mesh split into 2 partitions is given below, with the halo cells being coloured red. An example local indexing scheme is also provided below the cells. Notice the local indexing scheme is set up such that owned cells have lower indices than halo cells.

Dofs

In the LFRic infrastracture the degrees-of-freedom (dofs) are indexed from 1 to the total number of dofs. The infrastructure also indexes dofs so that the values in a column are contiguous and their values increase in the vertical. Thus, given the dof indices for the "bottom" cell, the rest of the dof indices can be determined for the column. This set of dof indices for the bottom cell is called a dofmap.

Dofs represent a field's values at various locations in the mesh. Fields can either be continuous or discontinuous. Continuous fields are so named because their values are continuous across cell boundaries. Dofs that represent continuous fields are shared between neighbouring cells. Discontinuous fields have values that are not necessarily related between neighbouring cells (there can be discontinuities across cell boundaries). Dofs that represent discontinuous fields are local to a cell.

Discontinuous Dofs

A simple example of discontinuous dofs is given below. In this case each cell contains 1 dof and there are 10 cells in a column. We only show the bottom cells and their corresponding dof indices. As explained earlier, the dof indices increase contiguously up the column, so the cell above the cell containing dof index 1 contains dof index 2 and the cell above that contains dof index 3 etc.

As discussed in the previous section, when the distributed memory option is switched on in the Dynamo0.3 API (see the *Distributed Memory* Section) the cells in the model are partitioned amongst processors and halo cells are added at the boundaries to a depth determined by the LFRic infrastructure. This results in the dofs being replicated in the halo cells, leading to a dof halo. As for cells, the LFRic infrastructure maintains the global dof indexing scheme and adds a local dof indexing scheme from 1 to the number of dofs in each partition, including any halo dofs.

An example for a depth-1 halo implementation with the earlier mesh split into 2 partitions is given below, with the halo cells drawn in grey and halo dofs coloured red. An example local partition indexing scheme is also provided below the dofs. As with cells, notice the local indexing scheme ensures that owned dofs have lower indices than halo dofs.

Continuous Dofs

A simple continuous dof example is given below for the same mesh as before. In this case dofs are on cell edges in the horizontal and there are 10 cells in a column. Again we only show the bottom cells and their corresponding dof indices. As explained earlier, the dof indices increase contiguously up the column, so the cell above the cell containing dof index 1 contains dof index 2 and the cell above that contains dof index 3 etc.

As already explained, when the distributed memory option is switched on in the Dynamo0.3 API (see the *Distributed Memory* Section) the cells in the model are partitioned amongst processors and halo cells are added at the boundaries.

In the example below we ignore the additional halo cells and just look at the partitioning of cells amongst processors (with the same mesh and 2 partitions as shown earlier). It can be seen that the dofs shared between cells which are on

different partitions now need to be replicated if fields on continuous dofs are going to be able to be computed locally on each partition. This concept is different to halos as there are no halo cells here, the fact that the cells are partitioned has meant that continuous dofs on the edge of the partition are replicated. The convention used in Dynamo0.3 is that the cell with the lowest global id determines which partition owns a dof and which has the copy. Dofs which are copies are called *annexed*. Annexed dofs are coloured blue in the example:

If we now extend the above example to include the halo cells (coloured grey) then we get:

An example for a depth-1 halo implementation with the earlier mesh split into 2 partitions is given below, with the halo cells drawn in grey and halo dofs coloured red. An example local indexing scheme is also provided below the dofs. Notice the local indexing scheme ensures that owned dofs have lower indices than annexed dofs, which in turn have lower indices than halo dofs.

Cell and Dof Ordering

Cells in a partition are sequentially indexed by the LFRic infrastructure, starting at 1, so that local cells occur first, then level-1 halo cells, then level-2 halo cells etc. A benefit of this layout is that it makes it easy for PSyclone to specify the required iteration space for cells as a single range, allowing a single Fortran do loop (or other language construct as required) to be generated. The LFRic infrastructure provides an API that returns the index of the last owned cell, the index of the last halo cell at a particular depth and the index of the last halo cell, to support PSyclone code generation.

Dofs on a partition are also sequentially indexed by the LFRic infrastructure, starting at 1, so that local dofs occur first, then annexed dofs (if the field is continuous), then level-1 halo dofs, then level-2 halo dofs etc. Again, this layout makes it easy for PSyclone to specify the required iteration space for dofs as a single range. As before, the LFRic infrastructure provides an API that returns the index of the last owned dof, the index of the last annexed dof, the index of the last halo dof at a particular depth and the index of the last halo dof, to support PSyclone code generation.

Multi-grid

The Dynamo 0.3 API supports kernels that map fields between meshes of different horizontal resolutions; these are termed "inter-grid" kernels. As indicated in Fig. 21.1 below, the change in resolution between each level is always a factor of two in both the x and y dimensions.

Inter-grid kernels are only permitted to deal with fields on two, neighbouring levels of the mesh hierarchy. In the context of a single inter-grid kernel we term the coarser of these meshes the "coarse" mesh and the other the "fine" mesh.

There are two types of inter-grid operation; the first is "prolongation" where a field on a coarse mesh is mapped onto a fine mesh. The second is "restriction" where a field on a fine mesh is mapped onto a coarse mesh. Given the factor of two difference in resolution between the fine and coarse meshes, the depth of any halo accesses for the field on the fine mesh must automatically be double that of those on the coarse mesh.

Fig. 21.1: The arrangement of cells in the multi-grid hierarchy used by LFRic. (Courtesy of R. Wong, Met Office.)

Loop iterators

In the current implementation of the Dynamo0.3 API it is possible to iterate (loop) either over cells or dofs. At the moment all coded kernels are written to iterate over cells and all builtin kernels are written to iterate over dofs, but that does not have to be the case.

The loop iteration information is specified in the kernel metadata. In the case of builtin's there is kernel metadata but it is part of PSyclone and is specified in *src/psyclone/dynamo0p3_builtins_mod.f90*.

For inter-grid kernels, it is the coarse mesh that provides the iteration space. (The kernel is passed a list of the cells in the fine mesh that are associated with the current coarse cell.)

Cell iterators: Continuous

When a kernel is written to iterate over cells and modify a continuous field, PSyclone always computes dofs on owned cells and redundantly computes dofs in the level-1 halo (or to depth 2 if the field is on the fine mesh of an inter-grid kernel - see *Multi-grid*). Users can apply a redundant computation transformation to increase the halo depth for additional redundant computation but it must always at least compute the level-1 halo. The reason for this is to ensure that the shared dofs on cells on the edge of the partition (both owned and annexed) are always correctly computed. Note that the outermost halo dofs are not correctly computed and therefore the outermost halo of the modified field is dirty after redundant computation. Also note that if we do not know whether a modified field is discontinuous or continuous then we must assume it is continuous.

An alternative solution could have been adopted in Dynamo0.3 whereby no redundant computation is performed and partial-sum results are shared between processors in a communication pattern similar to halo exchanges. However, a decision was made to always perform redundant computation.

A downside of performing redundant computation in the level-1 halo is that any fields being read by the kernel must have their level-1 halo clean (up-to-date), which can result in halo exchanges. Note that this is not the case for the modified field, it does not need its halo to be clean.

Cell iterators: Discontinuous

When a kernel is written to iterate over cells and modify a discontinuous field, PSyclone only needs to compute dofs on owned cells. Users can apply a redundant computation transformation to redundantly compute into the halo but this is not done by default.

Dof iterators

When a kernel that is written to iterate over dofs modifies a field, PSyclone must ensure that all dofs in that field are updated. If the distributed memory flag is set to `false` then PSyclone must iterate over all dofs. PSyclone simply needs to create a loop that iterates from 1 to the total number of dofs. The latter value is provided by the LFRic API.

If the distributed memory flag is set to `true` then PSyclone must ensure that each partition only iterates over owned dofs. Again PSyclone just needs to create a loop that iterates from 1 to the total number of owned dofs on that partition. The latter value is provided by the LFRic API.

When the distributed memory flag is set to `true` an aditional configuration option can be set which makes PSyclone always create loops which iterate over both owned and annexed dofs. Whilst this is not necessary for correctness, it can improve performance by reducing the number of halo exchanges required (at the expense of computing annexed dofs redundantly). The only change for PSyclone is that it calls a different LFRic routine which returns the index of the last annexed dof. This iteration space will necessarily also include all owned dofs due to the ordering of dof indices discussed earlier.

The configuration variable is called *COMPUTE_ANNEXED_DOFS* and is found in the *dynamo0.3* section of the *psyclone.cfg* configuration file (see *Configuration*). If it is `true` then annexed dofs are always computed in loops that iterate over dofs and if it is `false` then annexed dofs are not computed. The default in PSyclone is `false`.

The computation of annexed dofs could have been added as a transformation optimisation. The reason for using a configuration switch is that it is then guaranteed that annexed dofs are always computed for loops that iterate over dofs which then allows us to always remove certain halo exchanges without needing to add any new ones.

If we first take the situation where annexed dofs are not computed for loops that iterate over dofs i.e. (*COMPUTE_ANNEXED_DOFS* is `false`), then a field's annexed dofs will be dirty (out-of-date) after the loop has completed. If a following kernel needs to read the field's annexed dofs, then PSyclone will need to add a halo exchange to make them clean.

There are 4 cases to consider:

1) the field is read in a loop that iterates over dofs,

2) the field is read in a loop that iterates over owned cells and level-1 halo cells,

3) the field is incremented in a loop that iterates over owned cells and level-1 halo cells, and

4) the field is read in a loop that iterates over owned cells

In case 1) the annexed dofs will not be read as the loop only iterates over owned dofs so a halo exchange is not required. In case 2) the full level-1 halo will be read (including annexed dofs) so a halo exchange is required. In case 3) the annexed dofs will be updated so a halo exchange is required. In case 4) the annexed dofs will be read so a halo exchange will be required.

If we now take the case where annexed dofs are computed for loops that iterate over dofs (*COMPUTE_ANNEXED_DOFS* is `true`) then a field's annexed dofs will be clean after the loop has completed. If a following kernel needs to read the field's annexed dofs, then PSyclone will no longer need a halo exchange.

We can now guarantee that annexed dofs will always be clean after a continuous field has been modified by a kernel. This is because loops that iterate over either dofs or cells now compute annexed dofs and there are no other ways for a continuous field to be updated.

We now consider the same four cases. In case 1) the annexed dofs will now be read, but annexed dofs are guaranteed to be clean, so no halo exchange is required. In case 2) the full level-1 halo is read so a halo exchange is still required. Note, as part of this halo exchange we will update annexed dofs that are already clean. In case 3) the annexed dofs will be updated but a halo exchange is not required as the annexed dofs are guaranteed to be clean. In case 4) the annexed dofs will be read but a halo exchange is not required as the annexed dofs are guaranteed to be clean.

Furthermore, in the 3rd and 4th cases (in which annexed dofs are read or updated but the rest of the halo does not have to be clean), where the previous writer is unknown (as it comes from a different invoke call) we need to add a speculative halo exchange (one that makes use of the runtime clean and dirty flags) when *COMPUTE_ANNEXED_DOFS* is *False*, as the previous writer *may* have iterated over dofs, leaving the annexed dofs dirty. In contrast, when *COMPUTE_ANNEXED_DOFS* is *True*, we do not require a speculative halo exchange as we know that annexed dofs are always clean.

Therefore no additional halo exchanges are required when *COMPUTE_ANNEXED_DOFS* is changed from `false` to `true` i.e. case 1) does not require a halo exchange in either situation and case 2) requires a halo exchange in both situations. We also remove halo exchanges for cases 3) and 4) so the number of halo exchanges may be reduced.

If a switch were not used and it were possible to use a transformation to selectively perform computation over annexed dofs for loops that iterate over dofs, then we would no longer be able to guarantee that annexed dofs would always be clean. In this situation, if the dofs were known to be dirty then PSyclone would need to add a halo exchange and if it were unknown whether the dofs were dirty or not, then a halo exchange would need to be added that uses the run-time flags to determine whether a halo exchange is required. As run-time flags are based on whether the halo is dirty or not (not annexed dofs) then a halo exchange would be performed if the halo were dirty, even if the annexed dofs were clean, potentially resulting in more halo exchanges than are necessary.

Halo Exchange Logic

Halo exchanges are required when the *DISTRIBUTED_MEMORY* flag is set to `true` in order to make sure any accesses to a field's halo or to its annexed dofs receive the correct value.

Operators and Halo Exchanges

Halo Exchanges are only created for fields. This causes an issue for operators. If a loop iterates over halos to a given depth and the loop includes a kernel that reads from an operator then the operator must have valid values in the halos to that depth. In the current implementation of PSyclone all loops which write to, or update an operator are computed redundantly in the halo up to depth-1 (see the *load()* method in the *DynLoop* class). This implementation therefore requires a check that any loop which includes a kernel that reads from an operator is limited to iterating in the halo up to depth-1. PSyclone will raise an exception if an optimisation attempts to increase the iteration space beyond this (see the *gen_code()* method in the *DynKern* class).

To alleviate the above restriction one could add a configurable depth with which to compute operators e.g. operators are always computed up to depth-2, or perhaps up to the maximum halo depth. An alternative would be to halo exchange operators as required in the same way that halo exchanges are used for fields.

First Creation

When first run, PSyclone creates a separate InvokeSchedule for each of the invokes found in the algorithm layer. This schedule includes all required loops and kernel calls that need to be generated in the PSy layer for the particular invoke call. Once the loops and kernel calls have been created then (if the *DISTRIBUTED_MEMORY* flag is set to `true`) PSyclone adds any required halo exchanges and global sums. This work is all performed in the *DynInvoke* constructor (*__init__*) method.

In PSyclone we apply a lazy halo exchange approach (as opposed to an eager one), adding a halo exchange just before it is required.

It is simple to determine where halo exchanges should be added for the initial schedule. There are three cases:

1) loops that iterate over cells and modify a continuous field will access the level-1 halo. This means that any field that is read within such a loop must have its level-1 halo clean (up-to-date) and therefore requires a halo exchange. A modified field (specified as *GH_INC* which involves a read before a write) will require a halo exchange if its annexed dofs are not clean, or if their status is unknown. Whilst it is only the annexed dofs that need to be made clean in this case, the only way to acheive this is via a halo exchange (which updates the halo i.e. more than is required). Note, if the *COMPUTE_ANNEXED_DOFS* configuration variable is set to `true` then no halo exchange is required as annexed dofs will always be clean.

2) continuous fields that are read from within loops that iterate over cells and modify a discontinuous field will access their annexed dofs. If the annexed dofs are known to be dirty (because the previous modification of the field is known to be from within a loop over dofs) or their status is unknown (because the previous modification to the field is outside of the current invoke) then a halo exchange will be required (As already mentioned, currently the only way to make annexed dofs clean is to perform a halo swap. Note, if the *COMPUTE_ANNEXED_DOFS* configuration variable is set to `true` then no halo exchange is required as annexed dofs will always be clean.

3) fields that have a stencil access will access the halo and need halo exchange calls added.

Halo exchanges are created separately (for fields with halo reads) for each loop by calling the *create_halo_exchanges()* method within the *DynLoop* class.

In the situation where a field's halo is read in more than one kernel in different loops, we do not want to add too many halo exchanges - one will be enough as long as it is placed correctly. To avoid this problem we add halo exchange calls for loops in the order in which they occur in the schedule. A halo exchange will be added before the first loop for a

field but the same field in the second loop will find that there is a dependence on the previously inserted halo exchange so no additional halo exchange will be added.

The algorithm for adding the necessary halo exchanges is as follows: For each loop in the schedule, the *create_halo_exchanges()* method iterates over each field that reads from its halo (determined by the *unique_fields_with_halo_reads()* method in the *DynLoop* class).

For each field we then look for its previous dependencies (the previous writer(s) to that field) using PSyclone's dependence analysis. Three cases can occur: 1) there is no dependence, 2) there are multiple dependencies and 3) there is one dependence.

1) If no previous dependence is found then we add a halo exchange call before the loop (using the internal helper method *_add_halo_exchange()*). If the field is a vector field then a halo exchange is added for each component. The internal helper method *_add_halo_exchange* itself uses the internal helper method *_add_halo_exchange_code()*. This method creates an instance of the *DynHaloExchange* class for the field in question and adds it to the schedule before the loop. You might notice that this method then checks that the halo exchange is actually required and removes it again if not. In our current situation the halo exchange will always be needed so this check is not required but in more complex situations after transformations have been applied to the schedule this may not be the case. We discuss this type of situation later.

2) If multiple previous dependencies are found then the field must be a vector field as this is the only case where this can occur. We then choose the closest one and treat it as a single previous dependency (see 3).

3) If a single previous dependency is found and it is a halo exchange then we do nothing, as it is already covered. This will only happen when more than one reader depends on a writer, as discussed earlier. If the dependence is not a halo exchange then we add one.

After completing the above we have all the halo exchanges required for correct execution.

Note that we do not need to worry about halo depth or whether a halo is definitely required, or whether it might be required, as this is determined by the halo exchange itself at code generation time. The reason for deferring this information is that it can change as transformations are applied.

Asynchronous Halo Exchanges

The Dynamo0p3AsynchronousHaloExchange transformation allows the default synchronous halo exchange to be split into a halo exchange start and a halo exhange end which are represented separately as nodes in the schedule. These can then be moved in the schedule to allow overlapping of communication and computation, as long as data dependencies are honoured.

A halo exchange both reads and modifies a field so has a readwrite access for dependence analysis purposes. An asynchronous halo exchange start reads the field and an asynchronous halo exchange end writes to the field. Therefore the obvious thing to do would be to have the associated field set to read and write access respectively. However, the way the halo exchange logic works means that it is simplest to set the halo exchange end access to readwrite. The reason for this is that the logic to determine whether a halo exchange is required (*_required()*) needs information from all fields that read from the halo after the halo exchange has been called (and therefore must be treated as a write with following reads for dependence analysis) and it needs information from all fields that write to the field before the halo exchange has been called (and therefore must be treated as a read with previous writes for dependence analysis). An alternative would be to make the *_required()* method use the halo exchange start for previous writes and the halo exchange end for following reads. However, it was decided that this would be more complicated than the solution chosen.

Both halo exchange start and halo exchange end inherit from halo exchange. However, the halo exchange start and end are really two parts of the same thing and need to have consistent properties including after transformations have been performed. This is achieved by having the halo exchange start find and use the methods from the halo exchange end, rather than implement them independently. The actual methods needed are *_compute_stencil_type()*, *_compute_halo_depth()* and *_required()*. It is unclear how much halo exhange start really benefits from inheriting

from halo exchange and this could probably be removed at the expense of returning appropriate names for the dag, colourmap, declaration etc.

Note: The dependence analysis for halo exchanges for field vectors is currently over zealous. It does not allow halo exchanges for independent vector components to be moved past one another. For example, a halo exchange for vector component 2, if placed after a halo exchange for component 1 could not be moved before the halo exchange for component 1, even though the accesses are independent of each other. This is also the case for asynchronous halo exchanges. See issue #220.

Evaluators

Evaluators consist of basis and/or differential basis functions for a given function space, evaluated at the nodes of another, 'target', function space. A kernel can request evaluators on multiple target spaces through the use of the *gh_evaluator_targets* metadata entry. Every evaluator used by that kernel will then be provided on all of the target spaces.

When constructing a *DynKernMetadata* object from the parsed kernel metadata, the list of target function-space names (as they appear in the meta-data) is stored in *DynKernMetadata._eval_targets*. This information is then used in the *DynKern._setup()* method which populates *DynKern._eval_targets*. This is an *OrderedDict* which has the (mangled) names of the target function spaces as keys and 2-tuples consisting of *FunctionSpace* and *DynKernelArgument* objects as values. The *DynKernelArgument* object provides the kernel argument from which to extract the function space and the *FunctionSpace* object holds full information on the target function space.

The *DynInvokeBasisFunctions* class is responsible for managing the evaluators required by all of the kernels called from an Invoke. *DynInvokeBasisFunctions._eval_targets* collects all of the unique target function spaces from the *DynKern._eval_targets* of each kernel.

DynInvokeBasisFunctions._basis_fns is a list holding information on each basis/differential basis function required by a kernel within the invoke. Each entry in this list is a *dict* with keys:

Key	Entry	Type
shape	Shape of the evaluator	*str*
type	Whether basis or differential basis	*str*
fspace	Function space	*FunctionSpace*
arg	Associated kernel argument	*DynKernelArgument*
qr_var	Quadrature argument name	*str*
nodal_fspaces	Target function spaces	list of *(FunctionSpace, DynKernelArgument)*

Modifying the Schedule

Transformations modify the schedule. At the moment only one of these transformations - the *Dynamo0p3RedundantComputationTrans* class in *transformations.py* - affects halo exchanges. This transformation can mean there is a requirement for new halo exchanges, it can mean existing halo exchanges are no longer required and it can mean that the properties of a halo exchange (e.g. depth) can change.

The redundant computation transformation is applied to a loop in a schedule. When this is done the *update_halo_exchanges()* method for that loop is called - see the *apply()* method in *Dynamo0p3RedundantComputationTrans*.

The first thing that the *update_halo_exchanges()* method does is call the *create_halo_exchanges()* method to add in any new halo exchanges that are required before this loop, due to any fields that now have a read access to their halo when they previously did not. For example, a loop containing a kernel that writes to a certain field might previously have iterated up to the number of owned cells in a partition (*ncells*) but now iterates up to halo depth 1.

However, a field that has its halo read no longer guarantees that a halo exchange is required, as the previous dependence may now compute redundantly to halo depth 2, for example. The solution employed in *create_halo_exchanges()* is to add a halo exchange speculatively and then remove it if it is not required. The halo exchange itself determines whether it is required or not via the *required()* method. The removal code is found at the end of the *_add_halo_exchange_code()* method in the *DynLoop()* class.

The second thing that the *update_halo_exchanges()* method does is check that any halo exchanges after this loop are still required. It finds all relevant halo exchanges, asks them if they are required and if they are not it removes them.

We only need to consider adding halo exchanges before the loop and removing halo exchanges after the loop. This is because redundant computation can only increase the depth of halo to which a loop computes so can not remove existing halo exchanges before a loop (as an increase in depth will only increase the depth of an existing halo exchange before the loop) or add existing halo exchanges after a loop (as an increase in depth will only make it more likely that a halo exchange is no longer required after the loop).

Kernel Transformations

Since PSyclone is invoked separately for each Algorithm file in an application, the naming of the new, transformed kernels is done with reference to the kernel output directory. All transformed kernels (and the modules that contain them) are re-named following the PSyclone Fortran naming conventions (*Fortran Naming Conventions*). This enables the reliable identification of transformed versions of any given kernel within the output directory.

If the "multiple" kernel-renaming scheme is in use, PSyclone simply appends an integer to the original kernel name, checks whether such a kernel is present in the output directory and if not, creates it. If a kernel with the generated name is present then the integer is incremented and the process repeated. If the "single" kernel-renaming scheme is in use, the same procedure is followed but if a matching kernel is already present in the output directory then the new kernel is not written (and we check that the contents of the existing kernel are the same as the one we would create).

If an application is being built in parallel then it is possible that different invocations of PSyclone will happen simultaneously and therefore we must take care to avoid race conditions when querying the filesystem. For this reason we use `os.open`:

```
fd = os.open(<filename>, os.O_CREAT | os.O_WRONLY | os.O_EXCL)
```

The `os.O_CREATE` and `os.O_EXCL` flags in combination mean that `open()` raises an error if the file in question already exists.

Colouring

If a loop contains one or more kernels that write to a field on a continuous function space then it cannot be safely executed in parallel on a shared-memory device. This is because fields on a continuous function space share dofs between neighbouring cells. One solution to this is to 'colour' the cells in a mesh so that all cells of a given colour may be safely updated in parallel (Fig. 21.2).

The loop over colours must then be performed sequentially but the loop over cells of a given colour may be done in parallel. A loop that requires colouring may be transformed using the `Dynamo0p3ColourTrans` transformation.

Each mesh in the multi-grid hierarchy is coloured separately (https://code.metoffice.gov.uk/trac/lfric/wiki/LFRicInfrastructure/MeshColouring) and therefore we cannot assume any relationship between the colour maps of meshes of differing resolution.

However, the iteration space for inter-grid kernels (that map a field from one mesh to another) is always determined by the coarser of the two meshes. Consequently, it is always the colouring of this mesh that must be used. Due to the set-up of the mesh hierarchy (see Fig. 21.1), this guarantees that there will not be any race conditions when updating shared quantities on either the fine or coarse mesh.

Fig. 21.2: Example of the colouring of the horizontal cells used to ensure the thread-safe update of shared dofs (black circles). (Courtesy of S. Mullerworth, Met Office.)

21.6.2 GOcean1.0

TBD

21.6.3 NEMO

Implicit Loops

When constructing the PSyIR of NEMO source code, PSyclone identifies loops that are implied by the use of Fortran array notation. Such use of array notation is encouraged in the NEMO Coding Conventions [nem13] and identifying these loops can be important when introducing, e.g. OpenMP.

However, not all uses of Fortran array notation in NEMO imply a loop. For instance,

```
ascalar = afunc(twodarray1(:,:))
```

means that the function `afunc` is passed the (whole of the) `twodarray1` and returns a scalar value. (The requirement for explicit array shapes in the NEMO Coding Convention means that any quantity without such a shape must therefore be a scalar.)

Alternatively, a statement that assigns to an array must imply a loop:

```
twodarray2(:,:) = bfunc(twodarray1(:,:))
```

but it can only be converted into an explicit loop by PSyclone if the function `bfunc` returns a scalar. Since PSyclone does not currently attempt to fully resolve all symbols when parsing NEMO code, this information is not available and therefore such statements are not identified as loops (issue https://github.com/stfc/PSyclone/issues/286). This may then mean that opportunities for optimisation are missed.

21.7 Modules

This section describes the functionality of the various Python modules that make up PSyclone.

21.7.1 Module: f2pygen

f2pygen provides functionality for generating Fortran code from scratch and supports the addition of a use statement to an existing parse tree.

Variable Declarations

Three different classes are provided to support the creation of variable declarations (for intrinsic, character and derived-type variables). An example of their use might be:

```
>>> from psyclone.f2pygen import ModuleGen, SubroutineGen, DeclGen, \
CharDeclGen, TypeDeclGen
>>> module = ModuleGen(name="testmodule")
>>> sub = SubroutineGen(module, name="testsubroutine")
>>> module.add(sub)
>>> sub.add(DeclGen(sub, datatype="integer", entity_decls=["my_int"]))
>>> sub.add(CharDeclGen(sub, length="10", entity_decls=["my_char"]))
>>> sub.add(TypeDeclGen(sub, datatype="field_type", entity_decls=["ufld"]))
>>> gen = str(module.root)
>>> print(gen)
  MODULE testmodule
    IMPLICIT NONE
    CONTAINS
    SUBROUTINE testsubroutine()
      TYPE(field_type) ufld
      CHARACTER(LEN=10) my_char
      INTEGER my_int
    END SUBROUTINE testsubroutine
  END MODULE testmodule
```

The full interface to each of these classes is detailed below:

class psyclone.f2pygen.**DeclGen**(*parent,   datatype=",   entity_decls=None,   intent=",   pointer=False,   kind=",   dimension=",   allocatable=False,   save=False, target=False, initial_values=None*)

Generates a Fortran declaration for variables of various intrinsic types (integer, real and logical). For character variables CharDeclGen should be used.

> **Parameters**
>
> - **parent** (psyclone.f2pygen.BaseGen.) – node to which to add this declaration as a child.
>
> - **datatype** (`str`) – the (intrinsic) type for this declaration.
>
> - **entity_decls** (`list`) – list of variable names to declare.
>
> - **intent** (`str`) – the INTENT attribute of this declaration.
>
> - **pointer** (`bool`) – whether or not this is a pointer declaration.
>
> - **kind** (`str`) – the KIND attribute to use for this declaration.
>
> - **dimension** (`str`) – the DIMENSION specifier (i.e. the xx in DIMENSION(xx)).

- **allocatable** (`bool`) – whether this declaration is for an ALLOCATABLE quantity.

- **save** (`bool`) – whether this declaration has the SAVE attribute.

- **target** (`bool`) – whether this declaration has the TARGET attribute.

- **initial_values** (`list of str with same no. of elements as entity_decls.`) – Initial value to give each variable.

    Raises **`RuntimeError`** – if datatype is not one of DeclGen.SUPPORTED_TYPES.

class psyclone.f2pygen.**CharDeclGen**(*parent*, *entity_decls=None*, *intent=''*, *pointer=False*, *kind=''*, *dimension=''*, *allocatable=False*, *save=False*, *target=False*, *length=''*, *initial_values=None*)

Generates a Fortran declaration for character variables.

Parameters

- **parent** (`psyclone.f2pygen.BaseGen.`) – node to which to add this declaration as a child.

- **entity_decls** (`list`) – list of variable names to declare.

- **intent** (`str`) – the INTENT attribute of this declaration.

- **pointer** (`bool`) – whether or not this is a pointer declaration.

- **kind** (`str`) – the KIND attribute to use for this declaration.

- **dimension** (`str`) – the DIMENSION specifier (i.e. the xx in DIMENSION(xx)).

- **allocatable** (`bool`) – whether this declaration is for an ALLOCATABLE quantity.

- **save** (`bool`) – whether this declaration has the SAVE attribute.

- **target** (`bool`) – whether this declaration has the TARGET attribute.

- **length** (`str`) – expression to use for the (len=xx) selector.

- **initial_values** (`list of str with same no. of elements as entity_decls. Each of these can be either a variable name or a literal, quoted string (e.g. "'hello'")`) – Initial value to give each variable.

class psyclone.f2pygen.**TypeDeclGen**(*parent*, *datatype=''*, *entity_decls=None*, *intent=''*, *pointer=False*, *dimension=''*, *allocatable=False*, *save=False*, *target=False*)

Generates a Fortran declaration for variables of a derived type.

Parameters

- **parent** (`psyclone.f2pygen.BaseGen`) – node to which to add this declaration as a child

- **datatype** (`str`) – the type for this declaration

- **entity_decls** (`list`) – list of variable names to declare

- **intent** (`str`) – the INTENT attribute of this declaration

- **pointer** (`bool`) – whether or not this is a pointer declaration

- **dimension** (`str`) – the DIMENSION specifier (i.e. the xx in DIMENSION(xx))

- **allocatable** (`bool`) – whether this declaration is for an ALLOCATABLE quantity

- **save** (`bool`) – whether this declaration has the SAVE attribute

- **target** (`bool`) – whether this declaration has the TARGET attribute

Adding code

f2pygen supports the addition of use statements to an existing *fparser1* parse tree:

psyclone.f2pygen.**adduse**(*name*, *parent*, *only=False*, *funcnames=None*)
> Adds a use statement with the specified name to the supplied object. This routine is required when modifying an existing AST (e.g. when modifying a kernel). The classes are used when creating an AST from scratch (for the PSy layer).
>
> **Parameters**
>
> - **name** (`str`) – name of module to USE
>
> - **parent** (`fparser.one.block_statements.*`) – node in fparser1 AST to which to add this USE as a child
>
> - **only** (`bool`) – whether this USE has an "ONLY" clause
>
> - **funcnames** (`list`) – list of quantities to follow the "ONLY" clause
>
> **Returns** an fparser1 Use object
>
> **Return type** `fparser.one.block_statements.Use`

The PSyclone code where the *adduse* function was used has recently been migrated from using *fparser1* to using *fparser2*. In recognition of this change a new version of *adduse* has been developed which adds use statements to an existing *fparser2* parse tree. For the timebeing this new version is located in the same file it is used - *algGen.py* - but will be migrated to *f2pygen* (or equivalent) in the future:

psyclone.algGen.**adduse**(*parse_tree*, *location*, *name*, *only=None*, *funcnames=None*)
> Add a Fortran 'use' statement to an existing fparser2 parse tree. This will be added at the first valid location before the current location.
>
> This function should be part of the fparser2 replacement for f2pygen (which uses fparser1) but is kept here until this is developed, see issue #240.
>
> The 'parse_tree' argument is only required as fparser2 currently does not connect a child to a parent. This will be addressed in issue fparser:#102.
>
> **Parameters**
>
> - **parse_tree** (`fparser.two.utils.Base`) – The full parse tree of the associated code
>
> - **location** (`fparser.two.utils.Base`) – The current location (node) in the parse tree provided in the parse_tree argument
>
> - **name** (`str`) – The name of the use statement
>
> - **only** (`bool`) – Whether to include the 'only' clause in the use statement or not. Defaults to None which will result in only being added if funcnames has content and not being added otherwise.
>
> - **funcnames** (`list of str`) – A list of names to include in the use statement's only list. If the list is empty or None then nothing is added. Defaults to None.
>
> **Raises**
>
> - **GenerationError** – if the location is not part of the parse tree.
>
> - **GenerationError** – if the location is not a valid location to add a use statement.

- **`NotImplementedError`** – if the type of parent node is not supported.

- **`InternalError`** – if the type of parent node does not have the expected structure.

21.7.2 Module: configuration

PSyclone uses the Python `ConfigParser` class (https://docs.python.org/3/library/configparser.html) for reading the configuration file. This is managed by the `psyclone.configuration` module which provides a `Config` class. This class is a singleton, which can be (created and) accessed using `Config.get()`. Only one such instance will ever exist:

`class` `psyclone.configuration.`**`Config`**
>   Handles all configuration management. It is implemented as a singleton using a class _instance variable and a get() function.

> **`api`**
> >   Getter for the API selected by the user.

> > > **Returns** The name of the selected API.

> > > **Return type** str

> **`api_conf`** (*api=None*)
> >   Getter for the object holding API-specific configuration options.

> > > **Parameters** **`api`** (`str`) – Optional, the API for which configuration details are required. If none is specified, returns the config for the default API.

> > > **Returns** object containing API-specific configuration

> > > **Return type** One of `psyclone.configuration.DynConfig`, `psyclone.configuration.GOceanConfig` or None.

> > > **Raises**

> > > > - **`ConfigurationError`** – if api is not in the list of supported APIs.

> > > > - **`ConfigurationError`** – if the config file did not contain a section for the requested API.

> **`default_api`**
> >   Getter for the default API used by PSyclone.

> > > **Returns** default PSyclone API

> > > **Return type** str

> **`default_stub_api`**
> >   Getter for the default API used by the stub generator.

> > > **Returns** default API for the stub generator

> > > **Return type** str

> **`distributed_memory`**
> >   Getter for whether or not distributed memory is enabled

> > > **Returns** True if DM is enabled, False otherwise

> > > **Return type** bool

> **`filename`**
> >   Getter for the full path and name of the configuration file used to initialise this configuration object.

> > > **Returns** full path and name of configuration file

Return type str

static find_file()
Static method that searches various locations for a configuration file. If the full path to an existing file has been provided in the PSYCLONE_CONFIG environment variable then that is returned. Otherwise, we search the following locations, in order:

- ${PWD}/.psyclone/

- **if inside-a-virtual-environment:** <base-dir-of-virtual-env>/share/psyclone/

- ${HOME}/.local/share/psyclone/

- <system-install-prefix>/share/psyclone/

Returns the fully-qualified path to the configuration file

Return type str

Raises **ConfigurationError** – if no config file is found

static get(*do_not_load_file=False*)
Static function that if necessary creates and returns the singleton config instance.

Parameters **do_not_load_file** (*bool*) – If set it will not load the default config file. This is used when handling the command line so that the user can specify the file to load.

get_default_keys()
Returns all keys from the default section. :returns list: List of all keys of the default section as strings.

static get_repository_config_file()
This function returns the absolute path to the config file included in the PSyclone repository. It is used by the testing framework to make sure all tests get the same config file (see tests/config_tests for the only exception). :return str: Absolute path to the config file included in the PSyclone repository.

include_paths

Returns the list of paths to search for Fortran include files.

Return type list of str.

kernel_naming

Returns what naming scheme to use when writing transformed kernels to file.

Return type str

kernel_output_dir

Returns the directory to which to write transformed kernels.

Return type str

load(*config_file=None*)
Loads a configuration file.

Parameters **config_file** (*str*) – Override default configuration file to read.

Raises **ConfigurationError** – if there are errors or inconsistencies in the specified config file.

reprod_pad_size
Getter for the amount of padding to use for the array required for reproducible OpenMP reductions

Returns padding size (no. of array elements)

> > > **Return type** int

> **reproducible_reductions**
> > Getter for whether reproducible reductions are enabled.

> > > **Returns** True if reproducible reductions are enabled, False otherwise.

> > > **Return type** bool

> **supported_apis**
> > Getter for the list of APIs supported by PSyclone.

> > > **Returns** list of supported APIs

> > > **Return type** list of str

> **supported_stub_apis**
> > Getter for the list of APIs supported by the stub generator.

> > > **Returns** list of supported APIs.

> > > **Return type** list of str

The `Config` class is responsible for finding the configuration file (if no filename is passed to the constructor), parsing it and then storing the various configuration options. If PSyclone is started via `pytest`, the environment variable `PSYCLONE_CONFIG` is set to `<PSYCLONEHOME/config>`. This will guarantee that all tests use the config file provided in the PSyclone repository, and not a (potentially modified) user installed version.

The `Config` class also stores the list of supported APIs (`Config._supported_api_list`) and the default API to use if none is specified in either a config file or the command line (`Config._default_api`). Additionally, it performs some basic consistency checks on the values it obtains from the configuration file.

Since the PSyclone API to use can be read from the configuration file, it is not possible to have API-specifc sub-classes of `Config` as we don't know which API is in use before we read the file. However, the configuration file can contain API-specific settings. These are placed in separate sections, named for the API to which they apply, e.g.:

```
[dynamo0.3]
COMPUTE_ANNEXED_DOFS = false
```

Having parsed and stored the options from the default section of the configuration file, the `Config` constructor then creates a dictionary using the list of supported APIs to provide the keys. The configuration file is then checked for API-specific sections (again using the API names from the default section) and, if any are found, an API-specifc sub-class is created using the parsed entries from the corresponding section. The resulting object is stored in the dictionary under the appropriate key. The API-specific values may then be accessed as, e.g.:

```
Config.get().api_conf("dynamo0.3").compute_annexed_dofs
```

The API-specific sub-classes exist to provide validation/type-checking and encapsulation for API-specific options. They do not sub-class `Config` directly but store a reference back to the `Config` object to which they belong.

21.7.3 Module: transformations

As one might expect, the transformations module holds the various transformation classes that may be used to modify the Schedule of an Invoke and/or the kernels called from within it.

The base class for any transformation must be the class `Transformation`:

class psyclone.psyGen.**Transformation**
> Abstract baseclass for a transformation. Uses the abc module so it can not be instantiated.

_validate(**args*)
> Method that validates that the input data is correct. It will raise exceptions if the input data is incorrect. This function needs to be implemented by each transformation.

> > **Parameters args** (*Type depends on actual transformation.*) – Arguments for the applying the transformation - specific to the actual transform used.

apply(**args*)
> Abstract method that applies the transformation. This function must be implemented by each transform.

> > **Parameters args** (*Type depends on actual transformation.*) – Arguments for the transformation - specific to the actual transform used.

> > **Returns** A tuple of the new schedule, and a momento.

> > **Return type** Tuple.

name
> Returns the name of the transformation.

Those transformations that work on a region of code (e.g. enclosing multiple kernel calls within an OpenMP region) must sub-class the RegionTrans class:

class psyclone.transformations.**RegionTrans**
> This abstract class is a base class for all transformations that act on a list of nodes. It gives access to a _validate function that makes sure that the nodes in the list are in the same order as in the original AST, no node is duplicated, and that all nodes have the same parent. We also check that all nodes to be enclosed are valid for this transformation - this requires that the sub-class populate the *valid_node_types* tuple.

> **_validate**(*node_list*)
> > Test if the nodes in node_list are in the original order.

> > **Parameters node_list** (*[list](#)*) – List of PSyIR nodes.

> > **Raises**

> > - **[*TransformationError*](#)** – If the nodes in the list are not in the original order in which they are in the AST, a node is duplicated or the nodes have different parents.

> > - **[*TransformationError*](#)** – if any of the nodes to be enclosed in the region are of an unsupported type.

> > - **[*TransformationError*](#)** – if the condition part of an IfBlock is erroneously included in the region.

21.7.4 Module: psyGen

Provides the base classes for PSy-layer code generation.

21.7.5 Module: dynamo0p3

Specialises various classes from the psyclone.psyGen module in order to support the Dynamo 0.3 API.

When constructing the Fortran subroutine for either an Invoke or Kernel stub (see *Stub Generation*), there are various groups of related quantities for which variables must be declared and (for Invokes) initialised. Each of these groupings is managed by a distinct sub-class of the DynCollection abstract class:

class psyclone.dynamo0p3.**DynCollection**(*node*)
> Base class for managing the declaration and initialisation of a group of related entities within an Invoke or Kernel stub

> > **Parameters node** (psyclone.dynamo0p3.DynInvoke or psyclone.dynamo0p3. DynKern) – the Kernel or Invoke for which to manage variable declarations and initialisation.
>
> > **Raises InternalError** – if the supplied node is not a DynInvoke or a DynKern.

> **_invoke_declarations**(*parent*)
> Add all necessary declarations for an Invoke.
>
> > **Parameters parent** (psyclone.f2pygen.SubroutineGen) – node in the f2pygen AST representing the Invoke to which to add declarations.

> **_stub_declarations**(*parent*)
> Add all necessary declarations for a Kernel stub. Not abstract because not all entities need representing within a Kernel.
>
> > **Parameters parent** (psyclone.f2pygen.SubroutineGen) – node in the f2pygen AST representing the Kernel stub to which to add declarations.

> **declarations**(*parent*)
> Insert declarations for all necessary variables into the AST of the generated code. Simply calls either _invoke_declarations() or _stub_declarations() depending on whether we're handling an Invoke or a Kernel stub.
>
> > **Parameters parent** (psyclone.f2pygen.SubroutineGen) – the node in the f2pygen AST representing the routine in which to insert the declarations.
> >
> > **Raises InternalError** – if neither self._invoke or self._kernel are set.

> **initialise**(*parent*)
> Add code to initialise the entities being managed by this class. We do nothing by default - it is up to the sub-class to override this method if initialisation is required.
>
> > **Parameters parent** (psyclone.f2pygen.SubroutineGen) – the node in the f2pygen AST to which to add initialisation code.

(A single base class is used for both Invokes and Kernel stubs since it allows the code dealing with variable declarations to be shared.) A concrete sub-class of DynCollection must provide an implementation of the _invoke_declarations method. If the quantities associated with the collection require initialisation within the PSy layer then the initialise method must also be implemented. If stub-generation is to be supported for kernels that make use of the collection type then an implementation must also be provided for _stub_declarations.

Although instances of (sub-classes of) DynCollection handle all declarations and initialisation, there remains the problem of constructing the list of arguments for a kernel (or kernel stub). The psyclone.dynamo0p3. ArgOrdering base class provides support for this:

class psyclone.dynamo0p3.**ArgOrdering**(*kern*)
Base class capturing the arguments, type and ordering of data in a Kernel call.

> **banded_dofmap**(*function_space*)
> Add banded dofmap (required for CMA operator assembly)
>
> > **Raises NotImplementedError** – because this is an abstract method

> **basis**(*function_space*)
> Add basis function information for this function space
>
> > **Raises NotImplementedError** – because this is an abstract method

> **cell_map**()
> Add cell-map information (for inter-grid kernels)
>
> > **Raises NotImplementedError** – because this is an abstract method

cell_position()
> Add cell position information

> > **Raises** `NotImplementedError` – because this is an abstract method

cma_operator(*arg*)
> Add information on the CMA operator

> > **Raises** `NotImplementedError` – because this is an abstract method

diff_basis(*function_space*)
> Add differential basis function information for this function space

> > **Raises** `NotImplementedError` – because this is an abstract method

field(*arg*)
> Add field information for this field argument

> > **Raises** `NotImplementedError` – because this is an abstract method

field_bcs_kernel(*function_space*)
> Add boundary condition information for a field on this function space

> > **Raises** `NotImplementedError` – because this is an abstract method

field_vector(*arg*)
> Add field-vector information for this field-vector argument

> > **Raises** `NotImplementedError` – because this is an abstract method

fs_common(*function_space*)
> Add information common to LMA operators and fields for this function space

> > **Raises** `NotImplementedError` – because this is an abstract method

fs_compulsory_field(*function_space*)
> Add compulsory information for this function space

> > **Raises** `NotImplementedError` – because this is an abstract method

generate()
> Specifies which arguments appear in an argument list, their type and their ordering. Calls methods for each type of argument that can be specialised by a child class for its particular need.

> > **Raises** `GenerationError` – if the kernel arguments break the rules for the Dynamo 0.3 API.

indirection_dofmap(*function_space*, *operator=None*)
> Add indirection dofmap required when applying a CMA operator

> > **Raises** `NotImplementedError` – because this is an abstract method

mesh_height()
> Add height information (i.e. no. of layers)

> > **Raises** `NotImplementedError` – because this is an abstract method

mesh_ncell2d()
> Add the number of columns in the mesh

> > **Raises** `NotImplementedError` – because this is an abstract method

operator(*arg*)
> Add operator information for this operator argument

> > **Raises** `NotImplementedError` – because this is an abstract method

operator_bcs_kernel(*function_space*)
>   Add boundary condition information for an operator on this function space

>>   **Raises** **NotImplementedError** – because this is an abstract method

orientation(*function_space*)
>   Add orientation information for this function space

>>   **Raises** **NotImplementedError** – because this is an abstract method

quad_rule()
>   Add qr information

>>   **Raises** **NotImplementedError** – because this is an abstract method

scalar(*arg*)
>   Add scalar information for this scalar argument

>>   **Raises** **NotImplementedError** – because this is an abstract method

stencil(*arg*)
>   Add stencil information for this stencil argument

>>   **Raises** **NotImplementedError** – because this is an abstract method

stencil_unknown_direction(*arg*)
>   Add stencil direction information for this stencil argument

>>   **Raises** **NotImplementedError** – because this is an abstract method

stencil_unknown_extent(*arg*)
>   Add stencil extent information for this stencil argument

>>   **Raises** **NotImplementedError** – because this is an abstract method

This class is then sub-classed in order to support the generation of argument lists when *calling* kernels (KernCallArgList) and when *creating* kernel stubs (KernStubArgList). KernCallArgList is only used in DynKernelArguments.raw_arg_list(). KernStubArgList is only used in DynKern. gen_stub(). These classes make use of DynCollection sub-classes in order to ensure that argument naming is consistent.

Kernel Transformations

PSyclone is able to perform kernel transformations. Currently it has two ways to apply transformations: by directly manipulating the language AST or by translating the language AST to PSyIR, apply the transformation, and producing the resulting language AST or code.

For now, both methods only support fparser2 AST for kernel code. This AST is obtained by converting the fparser1 AST (stored when the kernel code was originally parsed to process the meta-data) back into a Fortran string and then parsing that with fparser2. (Note that in future we intend to adopt fparser2 throughout PSyclone so that this translation between ASTs will be unnecessary.) The *ast* property of the *psyclone.psyGen.Kern* class is responsible for performing this translation the first time it is called. It also stores the resulting AST in *Kern._fp2_ast* for return by future calls.

See *psyclone.transformations.ACCRoutineTrans* for an example of directly manipulating the fparser2 AST.

Alternatively, one can call the *psyclone.psyGen.Kern.get_kernel_schedule()* to generate the PSyIR representation of the kernel code.

Kern.**get_kernel_schedule**()
>   Returns a PSyIR Schedule representing the kernel code. The Schedule is just generated on first invocation, this allows us to retain transformations that may subsequently be applied to the Schedule (but will not adapt to transformations applied to the fparser2 AST).

Returns Schedule representing the kernel code.

Return type `psyclone.psyGen.KernelSchedule`

The AST to AST transformation is done using an ASTProcessor. At the moment, *psyclone.psyGen.Fparser2ASTProcessor* and its specialised version for Nemo *psyclone.nemo.NemoFparser2ASTProcessor* are available. (In the future we aim to have a generic ASTProcessor class, specialized for different language parsers: <parser>ASTProcessor, and specialized again for specific APIs when additional functionality is reqiuered <API><parser>ASTProcessor.)

class `psyclone.psyGen.`**`Fparser2ASTProcessor`**
> Class to encapsulate the functionality for processing the fparser2 AST and convert the nodes to PSyIR.

> **`generate_schedule`**(*name*, *module_ast*)
>> Create a KernelSchedule from the supplied fparser2 AST.

>> **Parameters**

>>> • **`name`** (`str`) – Name of the subroutine represented by the kernel.

>>> • **`module_ast`** (`fparser.two.Fortran2003.Program`) – fparser2 AST of the full module where the kernel code is located.

>> **Raises `GenerationError`** – Unable to generate a kernel schedule from the provided fpaser2 parse tree.

> **static `get_inputs_outputs`**(*nodes*)
>> Identify variables that are inputs and outputs to the section of Fortran code represented by the supplied list of nodes in the fparser2 parse tree. Loop variables are ignored.

>> **Parameters `nodes`** (list of `fparser.two.utils.Base`) – list of Nodes in the fparser2 AST to analyse.

>> **Returns** 3-tuple of list of inputs, list of outputs, list of in-outs

>> **Return type** (list of str, list of str, list of str)

> **static `nodes_to_code_block`**(*parent*, *statements*)
>> Create a CodeBlock for the supplied list of statements and then wipe the list of statements. A CodeBlock is a node in the PSyIR (Schedule) that represents a sequence of one or more Fortran statements which PSyclone does not attempt to handle.

>> **Parameters**

>>> • **`parent`** (`psyclone.psyGen.Node`) – Node in the PSyclone AST to which to add this code block.

>>> • **`statements`** (`list`) – List of fparser2 AST nodes constituting the code block.

>> **Return type** `psyclone.CodeBlock`

> **`process_declarations`**(*parent*, *nodes*, *arg_list*)
>> Transform the variable declarations in the fparser2 parse tree into symbols in the PSyIR parent node symbol table.

>> **Parameters**

>>> • **`parent`** (`psyclone.psyGen.KernelSchedule`) – PSyIR node in which to insert the symbols found.

>>> • **`nodes`** (list of `fparser.two.utils.Base`) – fparser2 AST nodes to search for declaration statements.

>>> • **`arg_list`** (`fparser.Fortran2003.Dummy_Arg_List`) – fparser2 AST node containing the argument list.

> > **Raises** `NotImplementedError` – The provided declarations contain attributes which are not
> > supported yet.

> **process_nodes**(*parent*, *nodes*, *nodes_parent*)
> > Create the PSyIR of the supplied list of nodes in the fparser2 AST. Currently also inserts parent information
> > back into the fparser2 AST. This is a workaround until fparser2 itself generates and stores this information.

> > **Parameters**

> > > - **parent** (*psyclone.psyGen.Node*) – Parent node in the PSyIR we are constructing.

> > > - **nodes** (list of fparser.two.utils.Base) – List of sibling nodes in fparser2 AST.

> > > - **nodes_parent** (fparser.two.utils.Base) – the parent of the supplied list of
> > > nodes in the fparser2 AST.

The results of *psyclone.psyGen.Kern.get_kernel_schedule* is a *psyclone.psyGen.KernelSchedule* which has the same
functionality as a PSyIR Schedule but with the addition of a Symbol Table (see *KernelSchedule*).

21.8 Transformations

21.8.1 OpenACC

PSyclone is able to generate code for execution on a GPU through the use of OpenACC. Support for generating
OpenACC code is implemented via *Transformations*. The specification of parallel regions and loops is very similar
to that in OpenMP and does not require any special treatment. However, a key feature of GPUs is the fact that they
have their own, on-board memory which is separate from that of the host. Managing (i.e. minimising) data movement
between host and GPU is then a very important part of obtaining good performance.

Since PSyclone operates at the level of Invokes, it has no information about when an application starts and thus no
single place in which to initiate data transfers to a GPU. (We assume that the host is responsible for model I/O and
therefore for populating fields with initial values.) Fortunately, OpenACC provides support for this kind of situation
with the `enter data` directive. This may be used to "define scalars, arrays and subarrays to be allocated in the cur-
rent device memory for the remaining duration of the program" [ope17]. The `ACCEnterDataTrans` transformation
adds an `enter data` directive to an Invoke:

class psyclone.transformations.**ACCEnterDataTrans**
> Adds an OpenACC "enter data" directive to a Schedule. For example:

```
>>> from psyclone.parse.algorithm import parse
>>> from psyclone.psyGen import PSyFactory
>>> api = "gocean1.0"
>>> filename = "nemolite2d_alg.f90"
>>> ast, invokeInfo = parse(filename, api=api, invoke_name="invoke")
>>> psy = PSyFactory(api).create(invokeInfo)
>>>
>>> from psyclone.psyGen import TransInfo
>>> t = TransInfo()
>>> dtrans = t.get_trans_name('ACCEnterDataTrans')
>>>
>>> schedule = psy.invokes.get('invoke_0').schedule
>>> schedule.view()
>>>
>>> # Add an enter-data directive
>>> newschedule, _ = dtrans.apply(schedule)
>>> newschedule.view()
```

The resulting generated code will then contain an `enter data` directive protected by an `IF(this is the first time in this Invoke)` block, e.g. (for the GOcean1.0 API):

```
! Ensure all fields are on the device and
! copy them over if not.
IF (first_time) THEN
  !$acc enter data  &
  !$acc& copyin(sshn_t,sshn_t%data,un%grid,un%grid%tmask,...)
  first_time = .false.
  ssha_t%data_on_device = .true.
  ...
```

Note that the `IF` block is not strictly required as the OpenACC run-time identifies when a reference is already on the device and does not copy it it over again. However, when profiling an application, it was seen that there was a small overhead associated with doing the `enter data`, even when the data was already on the device. The `IF` block eliminates this.

Of course, a given field may already be on the device (and have been updated) due to a previous Invoke. In this case, the fact that the OpenACC run-time does not copy over the now out-dated host version of the field is essential for correctness.

In order to support the incremental porting and/or debugging of an application, PSyclone also supports the OpenACC `data` directive that creates a statically-scoped data region. See the description of the `ACCDataTrans` transformation in the *Available* section for more details.

21.8.2 OpenCL

PSyclone is able to generate an OpenCL [ope18] version of PSy-layer code for the GOcean 1.0 API. Such code may then be executed on devices such as GPUs and FPGAs (Field-Programmable Gate Arrays). Since OpenCL code is very different to that which PSyclone normally generates, its creation is handled by `gen_ocl` methods instead of the normal `gen_code`. Which of these to use is determined by the value of the `InvokeSchedule.opencl` flag. In turn, this is set at a user level by the `transformations.OCLTrans` transformation.

The PSyKAl model of calling kernels for pre-determined iteration spaces is a natural fit to OpenCL's concept of an `NDRangeKernel`. However, the kernels themselves must be created or loaded at runtime, their arguments explicitly set and any arrays copied to the compute device. All of this 'boilerplate' code is generated by PSyclone. In order to minimise the changes required, the generated code is still Fortran and makes use of the FortCL library (https://github.com/stfc/FortCL) to access OpenCL functionality. We could of course generate the PSy layer in C instead but this would require further extension of PSyclone.

Consider the following invoke:

```
call invoke( compute_cu(CU_fld, p_fld, u_fld) )
```

When creating the OpenCL PSy layer for this invoke, PSyclone creates three subroutines instead of the usual one. The first, `psy_init` is responsible for ensuring that a valid kernel object is created for each kernel called by the invoke, e.g.:

```
use fortcl, only: ocl_env_init, add_kernels
...
! Initialise the OpenCL environment/device
CALL ocl_env_init
! The kernels this PSy layer module requires
kernel_names(1) = "compute_cu_code"
! Create the OpenCL kernel objects. Expects to find all of the
! compiled kernels in PSYCLONE_KERNELS_FILE.
CALL add_kernels(1, kernel_names)
```

As indicated in the comment, the `FortCL::add_kernels` routine expects to find all kernels in a pre-compiled file pointed to by the PSYCLONE_KERNELS_FILE environment variable. (A pre-compiled file is used instead of run-time kernel compilation in order to support execution on FPGAs.)

The second routine created by PSyclone sets the kernel arguments, e.g.:

```
SUBROUTINE compute_cu_code_set_args(kernel_obj, nx, cu_fld, p_fld, u_fld)
  USE clfortran, ONLY: clSetKernelArg
  USE iso_c_binding, ONLY: c_sizeof, c_loc, c_intptr_t
  ...
  INTEGER(KIND=c_intptr_t), target :: cu_fld, p_fld, u_fld
  INTEGER(KIND=c_intptr_t), target :: kernel_obj
  INTEGER, target :: nx
  ! Set the arguments for the compute_cu_code OpenCL Kernel
  ierr = clSetKernelArg(kernel_obj, 0, C_SIZEOF(nx), C_LOC(nx))
  ierr = clSetKernelArg(kernel_obj, 1, C_SIZEOF(cu_fld), C_LOC(cu_fld))
  ...
END SUBROUTINE compute_cu_code_set_args
```

The third routine generated is the ususal psy-layer routine that is responsible for calling all of the kernels. However, it must now also call `psy_init`, create buffers on the compute device (if they are not already present) and copy data over:

```
SUBROUTINE invoke_compute_cu(...)
  ...
  IF (first_time) THEN
    first_time = .false.
    CALL psy_init
    num_cmd_queues = get_num_cmd_queues()
    cmd_queues => get_cmd_queues()
    kernel_compute_cu_code = get_kernel_by_name("compute_cu_code")
  END IF
  globalsize = (/p_fld%grid%nx, p_fld%grid%ny/)
  ! Ensure field data is on device
  IF (.NOT. cu_fld%data_on_device) THEN
    size_in_bytes = int(p_fld%grid%nx*p_fld%grid%ny, 8)* &
                    c_sizeof(cu_fld%data(1,1))
    ! Create buffer on device
    cu_fld%device_ptr = create_rw_buffer(size_in_bytes)
    ierr = clEnqueueWriteBuffer(cmd_queues(1), cu_fld%device_ptr,  &
                                CL_TRUE, 0_8, size_in_bytes,       &
                                C_LOC(cu_fld%data), 0, C_NULL_PTR, &
                                C_LOC(write_event))
    cu_fld%data_on_device = .true.
  END IF
  ...
```

Note that we use the `data_on_device` member of the field derived type (implemented in github.com/stfc/dl_esm_inf) to keep track of whether a given field has been copied to the compute device. Once all of this setup is done, the kernel itself is launched by calling `clEnqueueNDRangeKernel`:

```
ierr = clEnqueueNDRangeKernel(cmd_queues(1), kernel_compute_cu_code, &
                              2, C_NULL_PTR, C_LOC(globalsize),      &
                              C_NULL_PTR, 0, C_NULL_PTR, C_NULL_PTR)
```

Limitations

Currently PSyclone can only generate the OpenCL version of the PSy layer. Execution of the resulting code requires that the kernels themselves be converted from Fortran to OpenCL (a dialect of C) and at present this must be done manually. Since all data accessed by an OpenCL kernel must be passed as an argument, this conversion must also convert any accesses to module data into routine arguments. Work is in progress to support kernel transformation and this will be made available in a future PSyclone release.

In OpenCL, all tasks to be performed (whether copying data or kernel execution) are associated with a command queue. Tasks submitted to different command queues may then be executed concurrently, potentially giving greater performance. The OpenCL PSy code currently generated by PSyclone makes use of just one command queue but again, this could be extended in the future.

The current implementation only supports the conversion of a whole Invoke to use OpenCL. In the future we may refine this functionality so that it may be applied to just a subset of kernels within an Invoke.

Since PSyclone knows nothing about the I/O performed by a model, the task of ensuring that the correct data is written out by a model (including when doing halo exchanges for distributed memory) is left to the dl_esm_inf library since that has the information on whether field data is local or on a remote compute device.

BIBLIOGRAPHY

[nem13] *NEMO Coding Conventions*. 2013. URL: https://forge.ipsl.jussieu.fr/nemo/attachment/wiki/Literature/ NEMO_coding.conv_v3.pdf.

[ope17] *The OpenACC Application Programming Interface, Version 2.6*. 2017. URL: http://www.openacc.org.

[ope18] *The OpenCL 2.2 Reference Guide*. 2018. URL: https://www.khronos.org/files/opencl22-reference-guide.pdf.

PYTHON MODULE INDEX

p

Symbols

A

S

T

U

V

W